

2ND EDITION

Web Development with Django

A definitive guide to building modern Python
web applications using Django 4



BEN SHAW | SAURABH BADHWAR
CHRIS GUEST | BHARATH CHANDRA K S

Web Development with Django

A definitive guide to building modern Python web applications using Django 4

Ben Shaw

Saurabh Badhwar

Chris Guest

Bharath Chandra K S



BIRMINGHAM—MUMBAI

Web Development with Django

Copyright © 2023 Packt Publishing

All rights reserved. No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, without the prior written permission of the publisher, except in the case of brief quotations embedded in critical articles or reviews.

Every effort has been made in the preparation of this book to ensure the accuracy of the information presented. However, the information contained in this book is sold without warranty, either express or implied. Neither the authors, nor Packt Publishing or its dealers and distributors, will be held liable for any damages caused or alleged to have been caused directly or indirectly by this book.

Packt Publishing has endeavored to provide trademark information about all of the companies and products mentioned in this book by the appropriate use of capitals. However, Packt Publishing cannot guarantee the accuracy of this information.

Group Product Manager: Pavan Ramchandani

Publishing Product Manager: Bhavya Rao

Senior Editor: Aamir Ahmed

Senior Content Development Editor: Rakhi Patel

Technical Editor: Simran Ali

Copy Editor: Safis Editing

Project Coordinator: Manthan Patel

Proofreader: Safis Editing

Indexer: Pratik Shirodkar

Production Designer: Prashant Ghare

Marketing Coordinators: Anamika Singh, Nivedita Pandey, and Namita Velgekar

First published: January 2021

Second edition: May 2023

Production reference: 1140423

Published by Packt Publishing Ltd.

Livery Place

35 Livery Street

Birmingham

B3 2PB, UK.

ISBN 978-1-80323-060-3

www.packtpub.com

To my family, who have supported me in every step and have provided that space where I can explore my creative sides.

— *Saurabh Badhwar*

To Annette Wilson, my mentor, who believed in me and had the patience to teach me to code for the real world.

— *Chris Guest*

To my wife Sneha and daughter Sharanya, for their support and encouragement throughout my journey.

— *Bharath Chandra K S*

Contributors

About the authors

Ben Shaw is a software engineer based in Auckland, New Zealand. He has worked as a developer for over 16 years and has been building websites with Django since 2007. In that time, his experience has helped many different types of companies, ranging in size from start-ups to large enterprises. He is also interested in machine learning, data science, automating deployments, and DevOps. When not programming, Ben enjoys outdoor sports and spending time with his partner and son.

Saurabh Badhwar is an infrastructure engineer who works on building tools and frameworks that enhance developer productivity. A major part of his work involves using Python to develop services that scale to thousands of concurrent users. He is currently employed at LinkedIn and works on infrastructure performance tools and services.

Chris Guest is based in Melbourne and started programming in Python 24 years ago when it was an obscure academic language. He has since used his Python knowledge in the publishing, hospitality, medical, academic, and financial sectors. Throughout his career, he has worked with many Python web development frameworks, including Zope, TurboGears, web2py, and Flask, although he still prefers Django.

Bharath Chandra K S is a passionate software developer with over 14 years of experience in the industry, currently residing in Sydney, Australia. He specializes in Python stack development, including frameworks such as Flask and Django, and has extensive experience with both monolithic and micro-service architectures. Bharath has developed various public-facing applications and data processing backend systems. When not creating software, he enjoys cooking delicious food.

I am grateful to my family for their constant support while working on this book. Thanks to my wife, Sneha, daughter, Sharanya, parents, Manjula and Dr. Srikanta Rao, brother, Dr.

Sharath, in-laws, Sudha and Ramesh. Special thanks to my grandfather, Javagal Nagesh Bhat, whose inspiration and guidance have been instrumental in my academic journey.

About the reviewer

While in college in the mid '80s, **Wayne Tooley** began programming as part of his program. In the early '90s, he began a career in IT and wrote software utilities to assist with various tasks. After 10 years working in large corporations, Wayne began a freelancing career offering business solutions for business as well as multimedia production.

Table of Contents

An Introduction to Django

Technical requirements	2	Working with GET, POST, and QueryDict objects	36
Scaffolding a Django project and app	2	Exercise 1.04 – exploring GET values and QueryDict objects	38
Exercise 1.01 – creating a project and app, and starting the development server	4		
Understanding the model-view-template paradigm	7	Exploring Django settings	40
Models	7	Referring to Django Settings in Your Code	42
Views	8		
Templates	8	Finding HTML templates in app directories	43
MVT in practice	8	Exercise 1.05 – creating a templates directory and a base template	43
An introduction to HTTP	11	Rendering a template with the render function	47
Processing a request	15	Exercise 1.06 – rendering a template in a view	47
Exploring the Django project structure	16	Rendering variables in templates	49
The myproject directory	18	Exercise 1.07 – using variables in templates	49
Django development server	19		
Django apps	19	Debugging and dealing with errors	52
PyCharm setup	21	Exceptions	52
Exercise 1.02 – project setup in PyCharm	21	Exercise 1.08 – generating and viewing exceptions	54
Introducing Django views	30	Debugging	55
Exploring URL mapping detail	31	Exercise 1.09 – debugging your code	56
Exercise 1.03 – writing a view and mapping a URL to it	32	Activity 1.01 – creating a site welcome screen	60

Activity 1.02 – a book search scaffold	61	Summary	63
--	----	---------	----

2

Models and Migrations		65	
Technical requirements	66	Exercise 2.04 – creating records for a many-to-one relationship	98
Understanding and using databases	66	Exercise 2.05 – creating records with many-to-many relationships	100
Relational databases	67	Exercise 2.06 – a many-to-many relationship using the add() method	101
Non-relational databases	67	Using the create() and set() methods for many-to-many relationships	102
Database operations using SQL	68	Read operations	103
Data types in relational databases	68	Exercise 2.07 – using the get() method to retrieve an object	103
Exercise 2.01 – creating a book database	68	Returning an object using the get() method	104
Understanding CRUD operations using SQL	72	Exercise 2.08 – using the all() method to retrieve a set of objects	105
SQL create operations	73	Retrieving objects by filtering	105
SQL read operations	73	Exercise 2.09 – using the filter() method to retrieve objects	106
SQL update operations	74	Filtering by field lookups	106
SQL delete Operations	75	Using pattern matching for filtering operations	107
Exploring Django ORM	75	Retrieving objects by using the exclude() method	108
Database configuration and creating Django applications	76	Retrieving objects using the order_by() method	108
Django apps	77	Querying across relationships	111
Django migration	78	Querying using foreign keys	111
Creating Django models and migrations	80	Querying using the model name	111
Field options	82	Querying across foreign key relationships using the object instance	112
Primary keys	84	Exercise 2.10 – querying across a many-to-many relationship using the field lookup	112
Relationships	86	Exercise 2.11 – a many-to-many query using objects	113
Adding the Review model	90	Exercise 2.12 – a many-to-many query using the set() method	113
Model methods	92	Exercise 2.13 – using the update() method	114
Migrating the reviews app	93	Exercise 2.14 – using the delete() method	115
Django's database CRUD operations	95		
Exercise 2.02 – creating an entry in the book database	96		
Exercise 2.03 – using the create() method to create an entry	97		
Creating an object with a foreign key	98		

Bulk create and bulk update operations	115	Exercise 2.18 – verifying whether a queryset contains a given object	118
Exercise 2.15 – creating multiple records using <code>bulk_create</code>	116	Activity 2.01 – creating models for a project management application	120
Exercise 2.16 – updating multiple records using <code>bulk_update</code>	117	Populating the Bookr project's database	120
Performing complex lookups using Q objects	117	Summary	121
Exercise 2.17 – performing a complex query using a Q object	118		

3

URL Mapping, Views, and Templates	123		
Technical requirements	124	Django's template language	134
Understanding function-based views	124	Exercise 3.03 – displaying a list of books and reviews	136
Understanding class-based views	124	Template inheritance	139
URL configuration	125	Template styling with Bootstrap	140
Exercise 3.01 – implementing a simple function-based view	127	Exercise 3.04 – adding template inheritance and a Bootstrap navigation bar	142
Working with Django templates	130	Activity 3.01 – implementing the book details view	144
Exercise 3.02 – using templates to display a greeting message	132	Summary	146

4

An Introduction to Django Admin	147		
Technical requirements	149	Managing Django users and groups	159
Creating a superuser account	149	Exercise 4.02 – adding and modifying users and groups through the admin app	160
Exercise 4.01 – creating a superuser account	149	Registering models with the admin app	165
CRUD operations using the Django admin app	151	Registering the reviews model	166
Create	152	Change lists	167
Retrieve	154	The Change publisher page	168
Update	156	The Book change page	172
Delete	158		

Exercise 4.03 – foreign keys and deletion behavior in the admin app	175	The filter	195
Customizing the admin interface	177	Exercise 4.04 – adding the date list_filter and date_hierarchy filters	196
Site-wide Django admin customizations	177	The search bar	199
Customizing the ModelAdmin classes	187	Excluding and grouping fields	202
The list display fields	188	Activity 4.03 – customizing the Model admins	206
The display decorator	194	Summary	209

5

Serving Static Files **211**

Technical requirements	213	Static file finders – use during collectstatic	233
Static file serving	213	Exercise 5.04 – collecting static files for production	235
Introduction to Static Files Finder	214	STATICFILES_DIRS prefixed mode	236
Static file finders – use during a request	215	The findstatic command	239
AppDirectoriesFinder	215	Exercise 5.05 – finding files using findstatic	240
Static file namespacing	216	Serving the latest files (for cache invalidation)	243
Exercise 5.01 – serving a file from an app directory	218	Exercise 5.06 – exploring the ManifestFilesStorage storage engine	245
Generating static URLs with the static template tag	222	Custom storage engines	249
Exercise 5.02 – using the static template tag	226	Activity 5.01 – adding a Reviews logo	252
FileSystemFinder	229	Activity 5.02 – CSS enhancements	254
Exercise 5.03 – serving from a project static directory	230	Activity 5.03 – adding a global logo	256
		Summary	258

6

Forms **259**

Technical requirements	259	Types of inputs	263
What is a form?	260	Exercise 6.01 – building a form in HTML	264
The form element	262	Form security with Cross-Site Request	

Forgery Protection	272	Exercise 6.03 – building and rendering a	
Accessing data in the View	276	Django form	300
Exercise 6.02 – working with POST data in a view	276	Validating forms and retrieving	
Choosing between GET or POST	280	Python values	305
Why use GET when we can put parameters in the URL mappings?	282	Exercise 6.04 – validating forms in a view	308
The Django Forms library	283	Built-in field validation	311
Defining a form	284	Exercise 6.05 – adding extra field validation	312
Rendering a form in a template	295	Activity 1 – Book Search	314
		Summary	318

7

Advanced Form Validation and Model Forms 319

Technical requirements	320	Creating or editing Django models	340
Custom field validation and cleaning	320	The ModelForm class	342
Custom validators	320	Exercise 7.03 – creating and editing a publisher	346
Cleaning methods	322	Activity 7.01 – styling and integrating the	
Multi-field validation	323	publisher form	352
Exercise 7.01 – custom clean and validation methods	327	Activity 7.02 – Review Creation UI	356
Adding placeholders and initial values	336	Summary	363
Exercise 7.02 – placeholders and initial values	338		

8

Media Serving and File Uploads 365

Technical requirements	366	Exercise 8.02 – template settings and using MEDIA_URL in templates	373
Settings for media uploads and serving	366	File uploads using HTML forms	376
Serving media files in development	366	Working with uploaded files in a view	378
Exercise 8.01 – configuring media storage and serving	367	Exercise 8.03 – file upload and download	381
Context processors and using MEDIA_URL in templates	371	File uploads with Django forms	385
		Exercise 8.04 – file uploads with a Django form	386

Image uploads with Django forms	391	ModelForms and file uploads	414
Resizing an image with Pillow	393	Exercise 8.07 – file and image uploads using	
Exercise 8.05 – image uploads using Django	394	ModelForm	416
forms		Handling file saving	419
Serving uploaded (and other) files	397	Activity 8.01 – image and PDF	
using Django		upload of books	421
Storing files on model instances	398	Activity 8.02 – displaying the cover	
Storing images on model instances	402	and sample link	425
Working with FieldFile	403	Summary	427
Referring to media in templates	408		

9

Sessions and Authentication		429	
Technical requirements	430	Enhancing templates with	
Middleware	430	authentication data	450
Middleware modules	431	Exercise 9.04 – toggling login and logout	
Implementing authentication views and		links in the base template	450
templates	433	Activity 9.01 – authentication-based content	
Exercise 9.01 – repurposing the Admin app		using conditional blocks in templates	452
login template	438	Sessions	454
Password storage in Django	441	The session engine	454
The profile page and the request.user object	441	Do you need to flag cookie content?	455
Exercise 9.02 – adding a profile page	442	Pickle or JSON storage?	456
Authentication decorators and		Exercise 9.05 – examining the session key	457
redirection		Storing data in sessions	461
Exercise 9.03 – adding authentication		Exercise 9.06 – storing recently viewed books	
decorators to the views	447	in sessions	462
		Activity 9.02 – using session storage for the	
		Book Search page	467
		Summary	469

10

Advanced Django Admin and Customizations	471
Technical requirements	472
Customizing the admin site	472
Discovering admin files in Django	472
Django's AdminSite class	473
Exercise 10.01 – creating a custom admin site for Bookr	475
Overriding the default admin.site	478
Exercise 10.02 – overriding the default admin site	478
Customizing admin site text using AdminSite attributes	480
Customizing admin site templates	481
Exercise 10.03 – customizing the logout	
template for the Bookr admin site	483
Adding views to the admin site	485
Creating the view function	485
Accessing common template variables	486
Mapping URLs for the custom view	486
Restricting custom views to the admin site	487
Exercise 10.04 – adding custom views to the admin site	488
Passing additional keys to templates using template variables	491
Activity 10.01 – building a custom admin dashboard with a built-in search	492
Summary	494

11

Advanced Templating and Class-Based Views	495
Technical requirements	496
Template filters	496
Custom template filters	497
Creating custom template filters	498
Implementing the custom filter function	499
Using custom filters inside templates	499
Exercise 11.01 – Creating a custom template filter	500
Django views	515
Class-based views	515
Exercise 11.02 – Creating a custom simple tag	507
Inclusion tags	510
Exercise 11.03 – Building a custom inclusion tag	512
Exercise 11.04 – Creating a book catalog using a CBV	517
CRUD operations with CBVs	523
The Read view	524
Activity 11.01 – Rendering details on the user profile page using inclusion tags	528
Summary	530

12

Building a REST API	531
Technical requirements	532
Understanding REST APIs	532
Django REST framework	532
Installation and configuration	533
Functional API views	533
Understanding serializers	536
Class-based API views and generic views	539
Activity 12.01 – creating an API endpoint for a top contributors page	544
Simplifying the code using ViewSets	545
URL configuration using routers and Viewsets	546
Exercise 12.04 – using ViewSets and routers	546
Implementing authentication	550
Token-based authentication	551
Exercise 12.05 – implementing token-based authentication for Bookr APIs	551
Summary	556

13

Generating CSV, PDF, and Other Binary Files	557
Technical requirements	558
Working with CSV files inside Python	558
Working with Python's csv module	558
Reading data from a CSV file	558
Exercise 13.0 – reading a CSV file with Python	559
Writing to CSV files using Python	562
Exercise 13.02 – generating a CSV file using Python's csv module	563
A better way to read and write CSV files	566
Working with Excel files in Python	568
Binary file formats for data exports	569
Working with XLSX files using the XlsxWriter package	569
XLSX files	569
Exercise 13.03 – creating XLSX files in Python	572
Working with PDF files in Python	575
Converting web pages into PDFs	575
Exercise 13.04 – generating a PDF version of a web page in Python	575
Playing with graphs in Python	578
Integrating plotly with Django	583
Integrating visualizations with Django	583
Exercise 13.06 – visualizing a user's reading history on the user's profile page	583
Activity 13.01 – exporting the books read by a user as an XLSX file	589
Summary	590

14

Testing Your Django Applications	591		
Technical requirements	592	Testing views with authentication	607
Importance of testing	592	Exercise 14.04 – writing test cases to validate authenticated users	608
Automation testing	592		
Testing in Django	593	Django RequestFactory	612
Implementing test cases	594	Exercise 14.05 – using RequestFactory to test views	612
Unit testing in Django	594	Testing class-based views	615
Utilizing assertions	594		
Exercise 14.01 – writing a simple unit test	596	Test case classes in Django	615
Performing pre-test setup and cleanup after every test case run	598	The SimpleTestCase class	616
Testing Django models	599	The TransactionTestCase class	616
Exercise 14.02 – testing Django models	599	The LiveServerTestCase class	616
Testing Django views	603	Modularizing test code	617
Exercise 14.03 – writing unit tests for Django views	604	Activity 14.01 – testing models and views in Bookr	618
		Summary	619

15

Django Third-Party Libraries	621		
Technical requirements	622	django-crispy-forms	658
Environment variables	622	The crispy filter	659
django-configurations	625	The crispy template Tag	660
manage.py changes	627	Exercise 15.04 – using Django Crispy Forms with SearchForm	662
Configuration from environment variables	628		
Exercise 15.01 – Django configurations setup	629	django-allauth	667
dj-database-url	633	django-allauth installation and setup	670
Exercise 15.02 – dj-database-url and setup	637	Initiating authentication with django-allauth	673
The Django Debug Toolbar	638	Other django-allauth features	673
Exercise 15.03 – setting up the Django Debug Toolbar	655	Activity 15.01 – using FormHelper to update forms	674
		Summary	677

16

Using a Frontend JavaScript Library with Django	679
Technical requirements	679
JavaScript frameworks	680
An introduction to JavaScript	682
Loading JavaScript	682
Variables and constants	683
Working with React	688
Components	689
Exercise 16.01 – setting up a React example	694
JSX – a JavaScript syntax extension	697
Exercise 16.02 – JSX and Babel	699
JSX properties	700
Exercise 16.03 – React component properties	701
JavaScript promises	703
The fetch function	704
Exercise 16.04 – fetching and rendering books	708
The verbatim template tag	712
Activity 16.01 – reviews preview	713
Summary	719
Index	721
Other Books You May Enjoy	736

Preface

Do you want to develop reliable and secure applications that stand out from the crowd without spending hours on boilerplate code? You've made the right choice by trusting the Django framework, and this book will tell you why. Often referred to as a “batteries included” web development framework, Django comes with all the core features needed to build a standalone application. *Web Development with Django* will take you through all the essential concepts and help you explore its power to build real-world applications using Python. Throughout the book, you'll get to grips with the major features of Django by building a website called Bookr – a repository for book reviews. This end-to-end case study is split into a series of bitesize projects presented as exercises and activities, allowing you to challenge yourself in an enjoyable and attainable way. As you advance, you'll acquire various practical skills, including how to serve static files to add CSS, JavaScript, and images to your application, implement forms to accept user input, and manage sessions to ensure a reliable user experience. You'll cover everyday tasks that are part of the development cycle of a real-world web application. By the end of this Django book, you'll have the skills and confidence to creatively develop and deploy your own projects.

Who this book is for

This book is for programmers looking to enhance their web development skills using the Django framework. To fully understand the concepts explained in this book, basic knowledge of Python programming as well as familiarity with JavaScript, HTML, and CSS, is assumed.

What this book covers

Chapter 1, An Introduction to Django, starts by getting a Django project set up almost immediately. You'll learn how to bootstrap a Django project, respond to web requests, and use HTML templates.

Chapter 2, Models and Migrations, introduces Django data models, the method of persisting data to a SQL database.

Chapter 3, URL Mapping, Views, and Templates, builds on the techniques that were introduced in *Chapter 1, An Introduction to Django*, and explains in greater depth how to route web requests to Python code and render HTML templates.

Chapter 4, An Introduction to Django Admin, shows how to use Django's built-in Admin GUI to create, update, and delete data stored by your models.

Chapter 5, Serving Static Files, explains how to enhance your website with styles and images and how Django makes managing these files easier.

Chapter 6, Forms, shows you how to collect user input through your website by using Django's Forms module.

Chapter 7, Advanced Form Validation and Model Forms, builds upon *Chapter 6, Forms*, by adding more advanced validation logic to make your forms more powerful.

Chapter 8, Media Serving and File Uploads, shows how to further enhance sites by allowing your users to upload files and serve them with Django.

Chapter 9, Sessions and Authentication, introduces the Django session and shows you how to use it to store user data and authenticate users.

Chapter 10, Advanced Django Admin and Customization, continues from *Chapter 4, An Introduction to Django Admin*. Now that you know more about Django, you can customize the Django admin with advanced features.

Chapter 11, Advanced Templating and Class-Based Views, lets you see how to reduce the amount of code you need to write using some of Django's advanced templating features and classes.

Chapter 12, Building a REST API, looks at how to add a REST API to Django to allow programmatic access to your data from different applications.

Chapter 13, Generating CSV, PDF, and Other Binary Files, further expands the capabilities of Django by showing how you can use it to generate more than just HTML.

Chapter 14, Testing Your Django Applications, is an important part of real-world development. This chapter shows how to use the Django and Python testing frameworks to validate your code.

Chapter 15, Django Third-Party Libraries, exposes you to some of the many community-built Django libraries, showing how to use existing third-party code to add functionality to your project quickly.

Chapter 16, Using a Frontend JavaScript Library with Django, brings interactivity to your website by integrating with React and the REST API created in *Chapter 12, Building a REST API*.

Chapter 17, Deploying a Django Application (Part 1 – Server Setup), begins the process of deploying the application by setting up your own server. This is a bonus chapter and is downloadable from the GitHub repository for this book.

Chapter 18, Deploying a Django Application (Part 2 – Configuration and Code Deployment), finishes the project by showing you how to deploy your project to a virtual server. This is also a bonus chapter and is downloadable from the GitHub repository for this book.

To get the most out of this book

Software/hardware covered in the book	Operating system requirements
Python 3.8	Windows, macOS, or Linux
Django 4.0	
React 16	

If you are using the digital version of this book, we advise you to type the code yourself or access the code from the book's GitHub repository (a link is available in the next section). Doing so will help you avoid any potential errors related to the copying and pasting of code.

Download the example code files

You can download the example code files for this book from GitHub at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition>. If there's an update to the code, it will be updated in the GitHub repository.

We also have other code bundles from our rich catalog of books and videos available at <https://github.com/PacktPublishing/>. Check them out!

Download the color images

We also provide a PDF file that has color images of the screenshots and diagrams used in this book. You can download it here: <https://packt.link/5pZtF>.

Conventions used

There are a number of text conventions used throughout this book.

Code in text: Indicates code words in text, database table names, folder names, filenames, file extensions, pathnames, dummy URLs, user input, and Twitter handles. Here is an example: “First, create a new file named `inclusion_tag.py` under the `filter_demo/templatetags` directory.”

A block of code is set as follows:

```
from django.http import HttpResponseRedirect
from django.views import View

class IndexView(View):

    def get(self, request):
        return HttpResponseRedirect("Hey there!")
```

When we wish to draw your attention to a particular part of a code block, the relevant lines or items are set in **bold**:

```
from django.shortcuts import render

def index(request):
    names = "john,doe,mark,swain"
    return render(request, "index.html", {'names': names})
```

Any command-line input or output is written as follows:

```
POST /form-submit HTTP/1.1
Host: www.example.com
Content-Length: 31
Content-Type: application/x-www-form-urlencoded

first_name=Joe&last_name=Bloggs
```

Bold: Indicates a new term, an important word, or words that you see onscreen. For instance, words in menus or dialog boxes appear in **bold**. Here is an example: “On filling the data and clicking **Save record**, Django will save the data to the database.”

Tips or important notes

Appear like this.

Get in touch

Feedback from our readers is always welcome.

General feedback: If you have questions about any aspect of this book, email us at customercare@packtpub.com and mention the book title in the subject of your message.

Errata: Although we have taken every care to ensure the accuracy of our content, mistakes do happen. If you have found a mistake in this book, we would be grateful if you would report this to us. Please visit www.packtpub.com/support/errata and fill in the form.

Piracy: If you come across any illegal copies of our works in any form on the internet, we would be grateful if you would provide us with the location address or website name. Please contact us at copyright@packt.com with a link to the material.

If you are interested in becoming an author: If there is a topic that you have expertise in and you are interested in either writing or contributing to a book, please visit authors.packtpub.com.

Share Your Thoughts

Once you've read *Web Development with Django – Second Edition*, we'd love to hear your thoughts! Please click [here](#) to go straight to the Amazon review page for this book and share your feedback.

Your review is important to us and the tech community and will help us make sure we're delivering excellent quality content.

Download a free PDF copy of this book

Thanks for purchasing this book!

Do you like to read on the go but are unable to carry your print books everywhere? Is your eBook purchase not compatible with the device of your choice?

Don't worry, now with every Packt book you get a DRM-free PDF version of that book at no cost.

Read anywhere, any place, on any device. Search, copy, and paste code from your favorite technical books directly into your application.

The perks don't stop there, you can get exclusive access to discounts, newsletters, and great free content in your inbox daily

Follow these simple steps to get the benefits:

1. Scan the QR code or visit the link below



<https://packt.link/free-ebook/9781803230603>

2. Submit your proof of purchase
3. That's it! We'll send your free PDF and other benefits to your email directly

An Introduction to Django

“The web framework for perfectionists with deadlines.” It’s a tagline that aptly describes Django, a framework that has been around for over 10 years now. It is battle-tested and widely used, with more and more people using it every day. All this might make you think that Django is old and no longer relevant. On the contrary – its longevity has proved that its API is reliable and consistent, and even those who learned Django v1.0 in 2007 can mostly write the same code for Django 4 today. Django is still in active development, with bug fixes and security patches being released monthly.

Like Python, the language in which it is written, Django is easy to learn, yet powerful and flexible enough to grow with your needs. It is a “batteries-included” framework – in other words, you do not have to find and install many other libraries or components to get your application up and running. Other frameworks, such as **Flask** or **Pylons**, require manually installing third-party frameworks for database connections or template rendering. Instead, Django has built-in support for database querying, URL mapping, and template rendering (we’ll go into detail on what these mean soon). However, just because Django is easy to use doesn’t mean it is limited. Django is used by many large sites, including Disqus (<https://disqus.com/>), Instagram (<https://www.instagram.com/>), Mozilla (<https://www.mozilla.org/>), Pinterest (<https://www.pinterest.com/>), OpenStack (<https://www.openstack.org/>), and National Geographic (<http://www.nationalgeographic.com/>).

Where does Django fit into the web? When talking about web frameworks, you might think of frontend JavaScript frameworks such as ReactJS, Angular, or Vue. These frameworks are used to enhance or add interactivity to already-generated web pages. Django sits at the layer beneath these tools and instead is responsible for routing a URL, fetching data from databases, rendering templates, and handling form input from users. However, this does not mean you must pick one or the other; JavaScript frameworks can be used to enhance the output from Django or interact with a REST API generated by Django.

In this book, we will build a Django project using the same methods that professional Django developers use every day. The application is called **Bookr** and allows you to browse and add books and book reviews. This workshop is divided into four sections. In the first section, we’ll start with the basics of scaffolding a Django app, quickly build some pages, and serve them with the Django development server. You’ll be able to add data to the database using the Django admin site.

This chapter introduces you to Django and its role in web development. You will begin by understanding how the **model-view-template** paradigm works and how Django processes HTTP **requests** and **responses**. Equipped with the basic concepts, you'll create your first Django project called Bookr, an application for adding, viewing, and managing book reviews. It's an application you'll keep enhancing and adding features to throughout this book. You will then learn about the `manage.py` command (used to orchestrate Django actions). You will use this command to start the Django development server and test whether the code you wrote so far works as expected. You will also learn how to work with **PyCharm**, a popular Python IDE, which you'll be using throughout this book. You will use it to write code that returns a **response** to your web browser. Finally, you'll learn how to use PyCharm's debugger to troubleshoot problems with your code. By the end of this chapter, you'll have the necessary skills to start creating projects using Django.

We will cover the following topics in this chapter:

- Scaffolding a Django project and app
- Understanding the model-view-template paradigm
- Exploring the Django project structure
- Introducing Django views
- Exploring URL mapping detail
- Exploring Django settings
- Finding HTML templates in app directories
- Debugging and dealing with errors
- Activity 1.01 – creating a site welcome screen
- Activity 1.02 – a book search scaffold

By the end of the book, you will have enough experience to design and build your own Django project from start to finish.

Technical requirements

Throughout this book, you will be writing code. If you need to refer to the complete code for this chapter, you can find it in this GitHub repository: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter01>.

Scaffolding a Django project and app

Before diving deeply into the theory behind Django paradigms and HTTP requests, we'll show you how easy it is to get a Django project up and running. After this first section and exercise, you will have created a Django project, made a request to it with your browser, and seen the response.

A Django project is a directory that contains all the data for your project – code, settings, templates, and assets. It is created and scaffolded by running the `django-admin` command on the command line with the `startproject` argument and providing the project name. For example, to create a Django project with the name `myproject`, the command that is run is as follows:

```
django-admin startproject myproject
```

This will create the `myproject` directory, which Django populates with the necessary files to run the project. Inside the `myproject` directory are two files (shown in *Figure 1.1*):

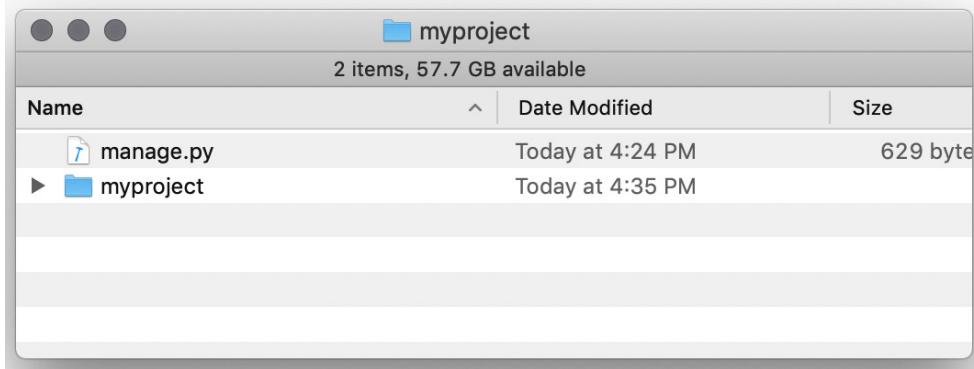


Figure 1.1 – The project directory for `myproject`

`manage.py` is a Python script that is executed at the command line to interact with your project. We will use it to start the **Django development server**, a development web server you will use to interact with your Django project on your local computer. Like `django-admin`, commands are passed in on the command line. Unlike `django-admin`, this script is not mapped in your system path, so we must execute it using Python. We will need to use the command line to do that. For example, inside the project directory, we run the following command:

```
python3 manage.py runserver
```

This passes the `runserver` command to the `manage.py` script, which starts the Django Dev Server. We will examine more of the commands that `manage.py` accepts later. When interacting with `manage.py` in this way, we call these management commands. For example, we might say that we are executing the `runserver` management command.

The `startproject` command also created a directory with the same name as the project – in this case, `myproject`. This is a Python package that contains settings and some other configuration files that your project needs to run. We will examine its contents later.

After starting the Django project, the next thing to do is to start a Django app. We should try to segregate our Django project into different apps, grouped by functionality. For example, with Bookr, we will have a *reviews* app. This will hold all the code, HTML, assets, and database classes specific to working with book reviews. If we decide to expand Bookr to sell books as well, we might add a *store* application, containing the files for the bookstore. Apps are created with the `startapp` management command, passing in the application name. Here is an example:

```
python3 manage.py startapp myapp
```

This creates the app directory (`myapp`) inside the project directory. Django automatically populates this with files for the app that are ready to be filled in when you start developing. We'll examine these files and discuss what makes a good app in the *Django apps* section.

Now that we've introduced the basic commands to scaffold a Django project and application, let's put them into practice by starting the Bookr project in the first exercise of this book.

Exercise 1.01 – creating a project and app, and starting the development server

Throughout this book, we will be building a book review website named Bookr. It will allow you to add fields for publishers, contributors, books, and reviews. A publisher will publish one or more books, and each book will have one or more contributors (author, editor, co-author, and so on). Only admin users will be allowed to modify these models. Once a user has signed up for an account on the site, they will be able to start adding reviews for a book.

In this exercise, you will scaffold the *bookr* Django project, test that Django is working by running the development server, and then create the *reviews* Django app.

You should already have a virtual environment set up with Django installed. To know how to do that, you can refer to the *Preface*. Once you're ready, let's start by creating the Bookr project:

1. Open the terminal and run the command to create the `bookr` project directory and the default subfolders:

```
DJANGO - ADMIN  STARTPROJECT  BOOKR
```

This command does not generate any output but will create a folder called `bookr` inside the directory in which you ran the command. You can look inside this directory and see the items we described before for the `myproject` example – the `bookr` package directory and the `manage.py` file.

We can now test that the project and Django are set up correctly by running the Django development server. Starting the server is done with the `manage.py` script.

2. In your terminal (or Command Prompt), change to the bookr project directory (using the `cd` command), and then run the `manage.py runserver` command as follows:

```
PYTHON3 MANAGE.PY RUNSERVER
```

Note

On Windows, you may need to run and replace `python3` with just `python` to make the command work every time you run it.

This command starts the Django development server. You should get output similar to the following:

```
WATCHING FOR FILE CHANGES WITH STATRELOADER
PERFORMING SYSTEM CHECKS...

SYSTEM CHECK IDENTIFIED NO ISSUES (0 SILENCED).

YOU HAVE 18 UNAPPLIED MIGRATION(S). YOUR PROJECT MAY NOT WORK PROPERLY UNTIL YOU
APPLY THE MIGRATIONS FOR APP(S): ADMIN, AUTH, CONTENTTYPES, SESSIONS.

RUN 'PYTHON MANAGE.PY MIGRATE' TO APPLY THEM.

JANUARY 01, 2022 - 02:54:03
DJANGO VERSION 4.0, USING SETTINGS 'BOOKR.SETTINGS'
STARTING DEVELOPMENT SERVER AT HTTP://127.0.0.1:8000/
QUIT THE SERVER WITH CONTROL-C.
```

You will probably receive some warnings about unapplied migrations, but that's okay for now.

3. Open up a web browser and go to `http://127.0.0.1:8000/`, which will show you the Django welcome screen (*Figure 1.2*). If you see this, you will know that your Django project was created successfully and all is working fine for now:

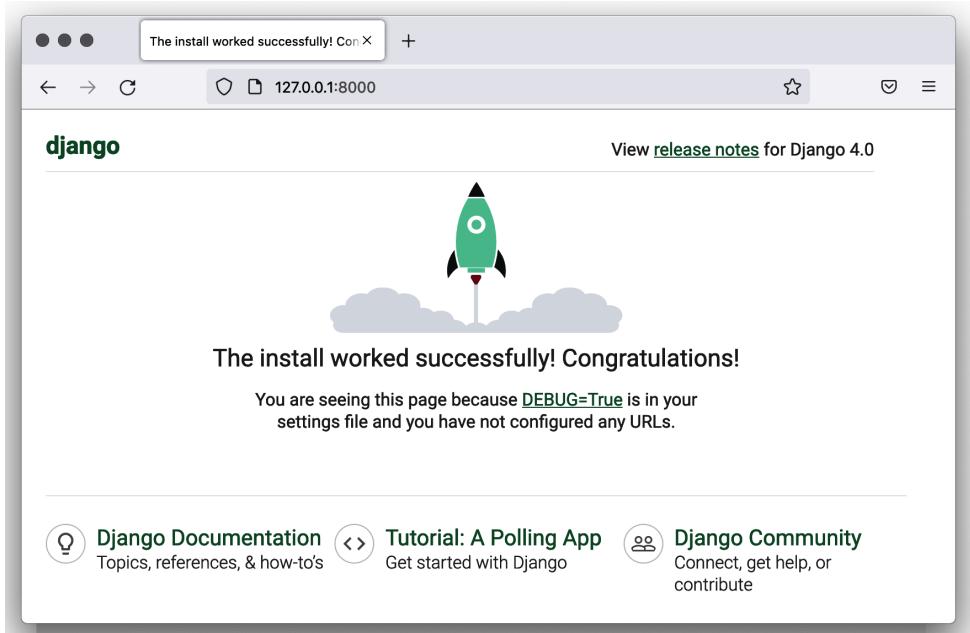


Figure 1.2 – The Django welcome screen

4. Go back to your terminal and stop the development server running by typing *Ctrl + C*.
5. We'll now create the *reviews* app for the Bookr project. In your terminal, make sure you are in the *bookr* project directory, and then execute the following command to create the *reviews* app:

```
PYTHON3 MANAGE.PY STARTAPP REVIEWS
```

Note

After creating the *reviews* app, the files in your Bookr project directory will look like this: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter01/Exercise1.01/bookr>.

There is no output if the command was successful, but a *reviews* app directory has been created. You can look inside this directory to see the files that were created – the *migrations* directory, *admin.py*, *models.py*, and so on. We'll examine these in detail in the *Django apps* section.

In this exercise, we created the Bookr project, tested that the project was working by starting the Django development server, and then created the *reviews* app for the project. Now that you've had some hands-on time with a Django project, we'll return to some of the theory behind Django's design, and HTTP requests and responses.

Understanding the model-view-template paradigm

A common design pattern in application design is the **Model View Controller (MVC)**, where the model of an application (its data) is displayed in one or more views, and a controller marshals interaction between the model and view. Django follows a different, yet similar, paradigm called the **Model-View-Template (MVT)**.

Like MVC, MVT also uses models for storing data. However, with MVT, a view will query a model and then render it with a template. Usually, with MVC languages, all three components need to be developed with the same language. With MVT, the template can be in a different language. In the case of Django, models and views are written in Python, and the template is written in HTML. This means that a Python developer could work on the models and views, while a specialist HTML developer works on the HTML. We'll first explain models, views, and templates in more detail and then look at some example scenarios where they are used.

Models

Django models define data for your application and provide an abstraction layer to a SQL database access through an **Object Relational Mapper (ORM)**. An ORM lets you define your data schema (classes, fields, and their relationships) using Python code, without needing an understanding of the underlying database. Basically, this means you can define your database layer in Python code and Django will take care of generating SQL queries for you. ORM will be discussed in detail in *Chapter 2, Models and Migrations*.

Note

SQL stands for **Structured Query Language** and is a way of describing a type of database that stores its data in tables, with each table having several rows. Think of each table being like an individual spreadsheet. Unlike a spreadsheet though, relationships can be defined between the data in each table. You can interact with data by executing SQL queries (often referred to as just queries when talking about databases). Queries allow you to retrieve data (SELECT), add or change data (INSERT and UPDATE respectively), and remove data (DELETE). There are many SQL database servers to choose from, such as SQLite, PostgreSQL, MySQL, or Microsoft SQL Server. Much of the SQL syntax is similar between each database, but there can be some differences in dialect. Django's ORM takes care of these differences for you – when we start coding, we will use the SQLite database to store data on disk, but later, when we deploy to a server, we will switch to PostgreSQL but won't need to make any code changes.

Normally, when querying a database, the results come back as primitive Python objects, (for example, lists of strings, ints, floats, or bytes). When using the ORM, results are automatically converted into instances of the model classes you have defined. Using an ORM means that you are automatically protected from a type of vulnerability known as a SQL injection attack.

If you're more familiar with databases and SQL, you always have the option of writing your own queries too.

Views

A Django view is where most of the logic for your application is defined. When a user visits your site, their web browser will send a request to retrieve data from your site (we will go into more detail on what an HTTP request is and what information it contains in the next section). A view is a function that you write that will receive this request in the form of a Python object (specifically, a Django `HttpRequest` object). It is up to your view to decide how it should respond to the request and what it should send back to the user. Your view must return an `HttpResponse` object that encapsulates all the information being provided to the client – content, HTTP status, and other headers.

The view can also optionally receive information from the URL of the request – for example, an ID number. A common design pattern of a view is to query a database via the Django ORM, using an ID that is passed into your view. Then, the view can render a template (there'll be more on this shortly) by providing it with data from the model retrieved from the database. The rendered template becomes the content of the `HttpResponse` object and is returned from the view function. Django takes care of the communication of the data back to the browser.

Templates

A template is a **HyperText Markup Language (HTML)** file (usually – any text file can be a template) that contains special placeholders that are replaced by variables your application provides. For example, your application could render a list of items in either a gallery layout or a table layout. Your view would fetch the same models for either one but would be able to render a different HTML file with the same information to present the data differently. Django emphasizes safety, so it will take care of automatically escaping variables for you. For example, the `<` and `>` symbols (among others) are special characters in HTML. If you try to use them in a variable, then Django automatically encodes them so that they render correctly in a browser.

MVT in practice

We'll now look at some examples to illustrate how MVT works in practice. In the examples, we have a `Book` model that stores information about different books, and a `Review` model that stores information about different reviews of books.

In the first example, we want to be able to edit the information about `Book` or `Review`. Take the first scenario of editing a book's details. We would have a view to fetch the `Book` data from the database and provide the `Book` model. Then, we would pass context information containing the `Book` object (and other data) to a template that would show a form to capture the new information. The second scenario (editing a review) is similar – fetch a `Review` model from the database, and then pass the

Review object and other data to a template to display an edit form. These scenarios might be so similar that we can reuse the same template for both, as shown in *Figure 1.3*.

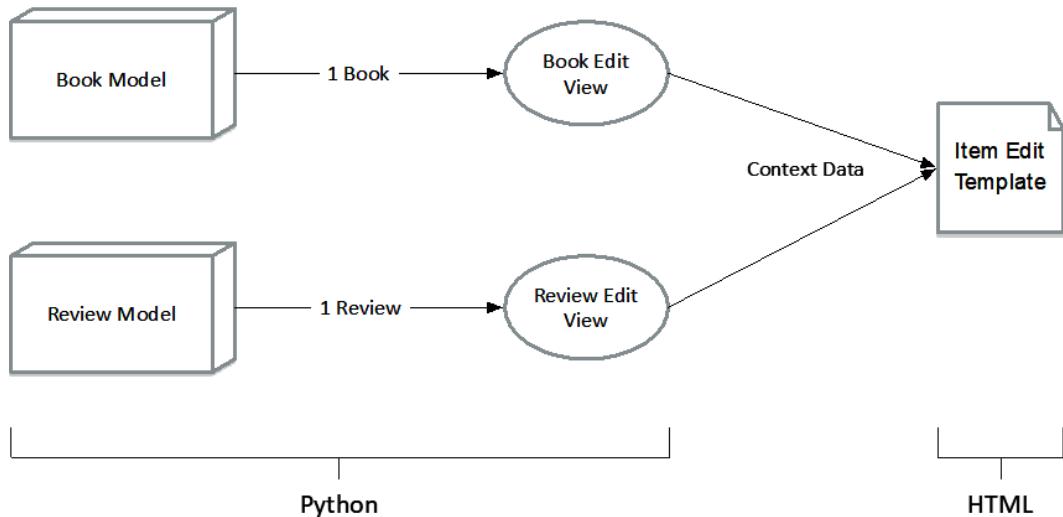


Figure 1.3 – Editing a single book or review

You can see here that we use two models, two views, and one template. Each view fetches a single instance of its associated model, but they can both use the same template, which is a generic HTML page to display a form. Views can provide extra context data to slightly alter the display of the template for each model type. Also illustrated in the diagram are the parts of the code that are written in Python and those that are written in HTML.

In the second example, we want to be able to show a user a list of the books or reviews that are stored in the application. Furthermore, we want to allow the user to search for books and get a list of all that match their criteria. We will use the same two models as the previous example (Book and Review) but we will create new views and templates. Since there are three scenarios, we'll use three views this time – the first fetches all books, the second fetches all reviews, and the last searches for books based on some search criteria. Once again, if we write a template well, we might be able to just use a single HTML template again, as shown in *Figure 1.4*.

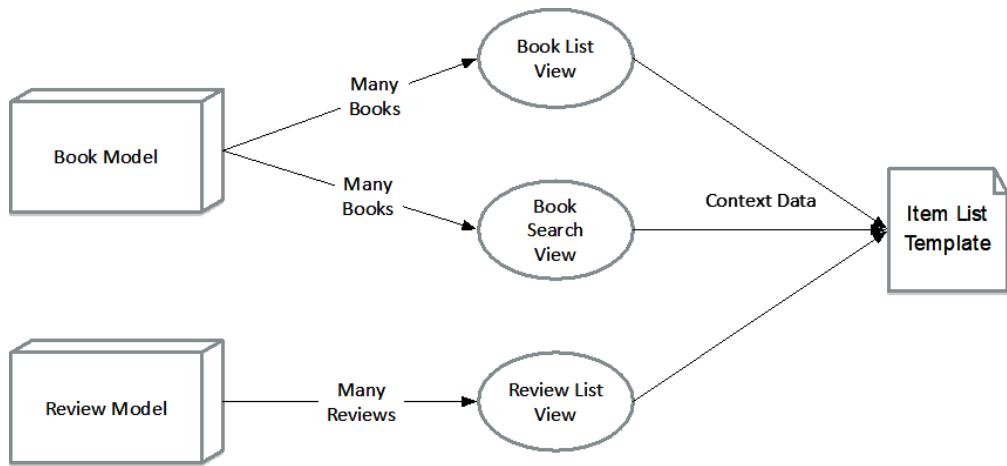


Figure 1.4 – Viewing multiple books or reviews

The Book and Review models remain unchanged from the previous example; the three views will fetch many (zero or more) books or reviews. Then, each view can use the same template, which is a generic HTML file that iterates over a list of objects it is provided and renders them. Once again, the views can send extra data in the context to alter how the template behaves, but the majority of the template will be as generic as possible.

In Django, a model does not always need to be used to render an HTML template. A view can generate the context data itself and render a template with it, without requiring any model data. *Figure 1.5* shows a view sending data straight to a template.

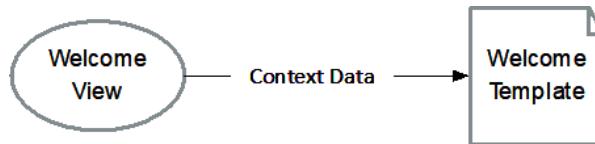


Figure 1.5 – A view sending data to a template without a model

In this example, there is a welcome view to welcome a user to the site. It doesn't need any information from the database, so it can just generate the context data itself. The context data depends on the type of information you want to display – for example, you could pass the user information to greet them by name if they are logged in.

It is also possible for a view to render a template without any context data. This can be useful if you have static information in an HTML file that you want to serve.

Now that you have been introduced to MVT in Django, we can look at how Django processes an HTTP request and generates an HTTP response. However, we first need to explain in more detail what HTTP requests and responses are and what information they contain. We will see this in the next section.

An introduction to HTTP

Let's say someone wants to visit your web page. They type in its URL or click a link to your site from a page they are already on. Their web browser creates an HTTP request, which is sent to the server hosting your website. Once a web server receives the HTTP request from your browser, it can interpret it and then send back a response. The response that the server sends might be simple, such as just reading an HTML or image file from disk and sending it. Alternatively, the response might be more complex, maybe using server-side software (such as Django) to dynamically generate the content before sending it.

The following diagram shows the direction of the transmission of HTTP requests and HTTP responses, between a browser and a web server.

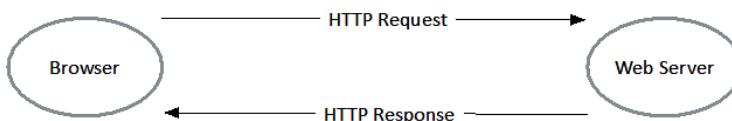


Figure 1.6 – An HTTP request and an HTTP response

The request is made up of four main parts – the method, the path, headers, and the body. Some types of requests don't have a body. If you just visit a web page, your browser will not send a body, whereas if you are submitting a form (for example, by logging into a site or performing a search), then your request will have a body containing the data you're submitting. We'll look at two example requests now to illustrate this.

The first request will be to an example page with the following URL: `https://www.example.com/page`. When your browser visits that page, this is what it's sending behind the scenes:

```
GET /page HTTP/1.1
Host: www.example.com
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:15.0)
Firefox/15.0.1
Cookie: sessid=abc123def456
```

The first line contains the method (GET) and the path (/page). It also contains the HTTP version – in this case, 1.1, although you don't have to worry about this. There are many different HTTP methods that can be used, depending on how you want to interact with the remote page. Some common ones are GET (to retrieve the remote page), POST (to send data to the remote page), PUT (to create a remote page) and DELETE (to delete the remote page). Note that the descriptions of the actions are somewhat simplified – the remote server can choose how it responds to different methods, and even experienced developers can disagree on the correct method to implement for a particular action. It's also important to note that even if a server supports a particular method, you will probably need the correct permissions to perform that action – you can't just delete a web page you don't like, for example.

When writing a web application, the vast majority of the time you will only deal with GET requests. When you start accepting forms, you'll also have to use POST requests. It is only when you are working with advanced features, such as creating REST APIs, that you will have to worry about PUT, DELETE, and other methods.

Referring to the example request again, line 2 onwards features the headers of the request. The headers contain extra metadata about the request. Each header is on its own line, with the header name and its value separated by a colon. Most are optional (except for `Host` – there'll be more on that soon). Header names are not case-sensitive. For the sake of this example, we're only showing three common headers here. Let's look at the example headers in order:

- `Host`: As mentioned, this is the only header that is required (for HTTP 1.1 or later). It is needed for the web server to know which website or application should respond to the request if there are multiple sites hosted on a single server.
- `User-Agent`: Your browser usually sends to the server a string identifying its version and operating system. Your server application could use this to serve different pages to different devices (for example, a mobile-specific page for smartphones).
- `Cookie`: You have probably seen a message when visiting a web page that lets you know that it is storing a cookie in the browser. These are small pieces of information that a website can store in your browser and use to identify you or save settings for when you return to the site. If you were wondering about how your browser sends these cookies back to the server, it is through this header.

There are many other standard headers defined, and it would take up too much space to list them all. They can be used to authenticate to the server (`Authorization`), tell the server what kind of data you can receive (`Accept`), or even what language you'd like for the page (`Accept-Language`, although this will only work if the page creator has made the content available in the particular language you request). You can even define your own headers that only your application knows how to respond to.

Now, let's look at a slightly more advanced request – one that sends some information to a server and thus (unlike the previous example) contains a body. In this example, we are logging into a web page by sending a username and password – for example, you visit `https://www.example.com/login`, and it displays a form to enter a username and password. After you click the **Login** button, this is the request that is sent to the server:

```
POST /login HTTP/1.1
Host: www.example.com
Content-Type: application/x-www-form-urlencoded
Content-Length: 32

username=user&password=password1
```

As you can see, this looks similar to the first example, but there are a few differences. The method is now `POST`, and two new headers have been introduced (you can assume that your browser will still be sending the other headers that were in the previous example too):

- `Content-Type`: This tells the server the type of data that is included in the body. In the case of `application/x-www-form-urlencoded`, the body is a set of key-value pairs. An HTTP client could set this header to tell the server if it was sending other types of data, such as JSON or XML.
- `Content-Length`: For the server to know how much data to read, the client must tell it how much data is being sent. The content length header contains the length of the body. If you count the length of the body in this example, you'll see that it's 32 characters.

The headers are always separated from the body by a blank line. By looking at the example, you should be able to tell how the form data is encoded in the body – `username` has the `user1` value and `password` has the `password1` value.

These requests were quite simple, but most requests don't get much more complicated. They might have different methods and headers but should follow the same format. Now that we've seen requests, we'll take a look at HTTP responses that come back from the server.

An HTTP response looks similar to a request and consists of three main parts – a status, headers, and a body. Like a request though, depending on the type of response, it might not have a body. The first response example is a simple successful response:

```
HTTP/1.1 200 OK
Server: nginx
Content-Length: 18132
Content-Type: text/html
Set-Cookie: sessid=abc123def46

<!DOCTYPE html><html><head>...
```

The first line contains the HTTP version, a numeric status code (200), and then a text description of what the code means (OK – the request was a success). We'll show some more statuses after the next example. Lines 2–5 contain headers, similar to a request. Some headers you may have seen before; we will explain them all in this context:

- `Server`: This is similar but the opposite of the `User-Agent` header – this is the server telling the client what software it is running.
- `Content-Length`: The client uses this value to determine how much data to read from the server to get the body.

- **Content-Type:** The server uses this header to indicate to the client what type of data it is sending. The client can then choose how it will display the data – an image must be displayed differently from HTML, for example.
- **Set-Cookie:** We saw in the first request example how a client sends a cookie to the server. This is the corresponding header that a server sends to set that cookie in the browser.

After the headers is a blank line, and then the body of the response. We haven't shown it all here, just the first few characters of the HTML that is being received, out of the 18,132 that the server has sent.

Next, we'll show an example of a response that is returned if a requested page is not found:

```
HTTP/1.1 404 Not Found
Server: nginx
Content-Length: 55
Content-Type: text/html

<!DOCTYPE html><html><body>Page Not Found</body></html>
```

It is similar to the previous example, but the status is now 404 Not Found. If you've ever been browsing the internet and received a 404 error, this is the type of response your browser received. The various status codes are grouped by the type of success or failure they indicate:

- 100–199: The server sends codes in this range to indicate protocol changes or that more data is required. You don't have to worry about these.
- 200–299: A status code in this range indicates successful handling of a response. As we saw, the most common one you will deal with is 200 OK.
- 300–399: A status code in this range means the page you are requesting has moved to another address. An example of this is a URL shortening service that would redirect you from the short URL to the full one when you visit it. Common responses are 301 Moved Permanently or 302 Found. When sending a redirect response, the server will also include a Location header that contains the URL you should be redirected to.
- 400–499: A status code in this range means that the request could not be handled because there was a problem with what the client sent. This is in contrast to a request not being able to be handled due to a problem on the server (we will discuss those soon). We've already seen a 404 Not Found response; this is due to a bad request because the client is requesting a document that does not exist. Some other common responses are 401 Unauthorized (the client should log in) or 403 Forbidden (the client is not allowed to access the specific resource). Both of these problems could be avoided by getting the client to log in, hence them being considered client-side (request) problems.

- 500–599: Status codes in this range indicate an error on the server’s side. The client shouldn’t expect to be able to adjust a request to fix the problem. When working with Django, the most common server error status you will see is 500 Internal Server Error. This will be generated if your code raises an exception. Another common one is 504 Gateway Timeout, which might occur if your code is taking too long to run. The other variants that are common to see are 502 Bad Gateway and 503 Service Unavailable, which generally mean there is a problem with your application’s hosting in some way.

These are only some of the most common HTTP statuses. You can find a more complete list at <https://developer.mozilla.org/en-US/docs/Web/HTTP/Status>. Like HTTP headers though, statuses are arbitrary, and an application can return custom statuses. It is up to the server and clients to decide what these custom statuses and codes mean.

If this is your first time being introduced to the HTTP protocol, there’s quite a lot of information to take in. Luckily, Django does all the hard work and encapsulates the incoming data into an `HttpRequest` object. Most of the time, you don’t need to know about most of the information coming in, but it’s available if you need it. Likewise, when sending a response, Django encapsulates your data in an `HttpResponse` object. Normally, you just set the content to return, but you also have the freedom to set HTTP status codes and headers. We will discuss how to access and set the information in `HttpRequest` and `HttpResponse` later in this chapter. In the next section, *Processing a request*, we will look at how Django receives, parses, and responds to an HTTP request.

Processing a request

This is a basic timeline of the request and response flows so that you can get an idea of what the code you’ll be writing does at each stage. In terms of writing code, the first part you will write is your view. The view you create will perform some actions, such as querying a database for data. Then, the view will pass this data to another function to render a template, finally returning the `HttpResponse` object that encompasses the data you want to send back to the client.

Next, Django needs to know how to map a specific URL to your view so that it can load the correct view for the URL it receives as part of a request. You will write this URL mapping in a URL configuration Python file.

When Django receives a request, it parses the URL config file and then finds the corresponding view. It calls the view, passing in the `HttpRequest` object that represents the request. Your view will return its `HttpResponse`, and then Django takes over again to send this data to its host web server and back out to the client that requested it.

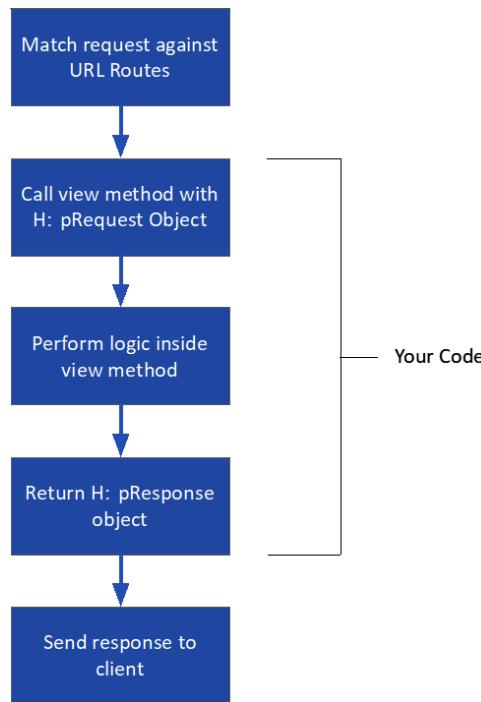


Figure 1.7 – The request and response flow

The request and response flow is illustrated in *Figure 1.7*; the sections indicated as **Your Code** are for the code that you write, and the first and last steps are taken care of by Django. Django does the URL matching for you, calls your view code, and then handles passing the response back to the client.

In this section, we learned about the structure of HTTP requests and responses, including what different types of requests are used for. We also saw how HTTP status codes can be used to denote errors in HTTP requests or responses. We also got an overview of how Django can process an HTTP request. In the next section, we'll explore the Django project structure and learn what the various files and directories are used for.

Exploring the Django project structure

We already introduced Django projects in the *Scaffolding a Django project and app* section. Let's remind ourselves of what happens when we run `startproject` (for a project named `myproject`) – the command creates a `myproject` directory with one file called `manage.py`, and a directory called `myproject` (this matches the project name; in *Exercise 1.01 – creating a project and app, and starting the development server*, this folder was called `bookr`, the same as the project). The directory layout is shown in *Figure 1.8*. We'll now examine the `manage.py` file and `myproject` package contents in more detail.

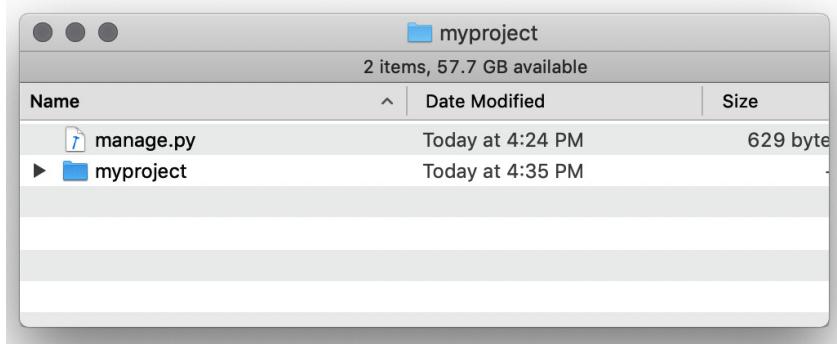


Figure 1.8 – The project directory for myproject

As the name suggests, `manage.py` is a script that is used to manage your Django project. Most of the commands that are used to interact with your project will be supplied to this script on the command line. The commands are supplied as an argument to this script – for example, if we want to run the `manage.py runserver` command, it would mean running the `manage.py` script like this:

```
python3 manage.py runserver
```

There are a number of useful commands that `manage.py` provides. You will be introduced to them in more detail throughout the book; some of the common ones are as follows:

- `runserver`: This starts the Django development HTTP server to serve the Django app on your local computer.
- `startapp`: This creates a new Django app in your project. We'll talk about what apps are in more depth soon.
- `shell`: This starts a Python interpreter with the Django settings pre-loaded. This is useful for interacting with your application without having to manually load your Django settings.
- `dbshell`: This starts an interactive shell connected to your database, using the default parameters from your Django settings. You can run manual SQL queries this way.
- `makemigrations`: This generates database change instructions from your model definitions. You will learn what this means and how to use this command in *Chapter 2, Models and Migrations*.
- `migrate`: This applies migrations generated by the `makemigrations` command. You will use this in *Chapter 2, Models and Migrations*, as well.
- `test`: This runs automated tests that you have written. You'll use this command in *Chapter 14, Testing Your Django Applications*.

A full list of all commands is available at <https://docs.djangoproject.com/en/3.0/ref/django-admin/>.

In the following sections, we will explore the contents of the project directory, learn what Django app directories contain, and see how to load up your project in PyCharm.

The myproject directory

As well as the `manage.py` file, the other file item created by `startproject` is the `myproject` directory. This is the actual Python package for your project. It contains settings for the project, some configuration files for your web server, and the global URL maps. Inside the `myproject` directory are five files:

- `__init__.py`
- `asgi.py`
- `settings.py`
- `urls.py`
- `wsgi.py`

They can also be seen in the following screenshot:

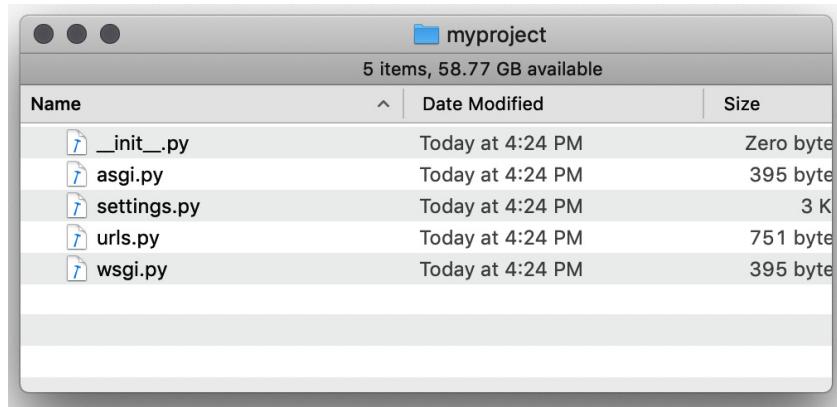


Figure 1.9 – The `myproject` package (inside the `myproject` project directory)

Now let's see what these files contain:

- `__init__.py`: This is an empty file that lets Python know that the `myproject` directory is a Python module. You'll be familiar with these files if you've worked with Python before.
- `settings.py`: This contains all the Django settings for your application. We will explain the contents soon.
- `urls.py`: This has the global URL mappings that Django will initially use to locate views or other child URL mappings. You will add a URL map to this file soon.

- `asgi.py` and `wsgi.py`: These files are what ASGI or WSGI web servers use to communicate with your Django app when you deploy it to a production web server. You normally don't need to edit these at all, and they aren't used in day-to-day development. Their use will be discussed more in *Chapter 17, Deployment a Django Application (Part 1 – Server Setup)*, which is hosted on GitHub repository of this book.

Django development server

You already started the Django development server in *Exercise 1.01 – creating a project and app, and starting the development server*. As we mentioned previously, it is a web server intended to only be run on a developer's machine during development. It is not intended for use in production.

By default, the server listens on port 8000 on `localhost` (`127.0.0.1`), but this can be changed by adding a port number, or address and port number, after the `runserver` argument:

```
python3 manage.py runserver 8001
```

This will have the server listen on port 8001 on `localhost` (`127.0.0.1`).

You can also have it listen on a specific address if your computer has more than one, or `0.0.0.0` for all addresses:

```
python3 manage.py runserver 0.0.0.0:8000
```

This will have the server listen on all your computer's addresses on port 8000, which can be useful if you want to test your application from another computer or your smartphone.

The development server watches your Django project directory and will restart automatically every time you save a file so that any code changes you make are automatically reloaded into the server. However, you still have to manually refresh your browser to see changes there.

When you want to stop the `runserver` command, it can be done in the usual way for stopping processes in the terminal – by typing `Ctrl + C`.

Django apps

Now that we've covered most of the theory about apps, we can be more specific about their purpose. An app directory contains all the models, views, templates (and more) that they need to provide application functionality. A Django project will contain at least one app (unless it has been heavily customized to not rely on a lot of Django functionality). If well-designed, an app should be able to be removed from a project and moved to another project without modification. Usually, an app will contain models for a single design domain, and this can be a useful way of determining whether your app should be split into multiple apps.

Your app can have any name as long as it is a valid Python module name (i.e., only letters, numbers, and underscores) and does not conflict with other files in your project directory. For example, as we have seen, there is already a directory called `myproject` in the project directory (containing the `settings.py` file), so you could not have an app called `myproject`. As we saw in *Exercise 1.01 – creating a project and app, and starting the development server*, creating an app uses the `manage.py startapp appname` command, as shown here:

```
python3 manage.py startapp myapp
```

The `startapp` command creates a directory within your project with the name of the app specified. It also scaffolds files for the app. Inside the app directory are a number of files and a folder, as shown in *Figure 1.10*:

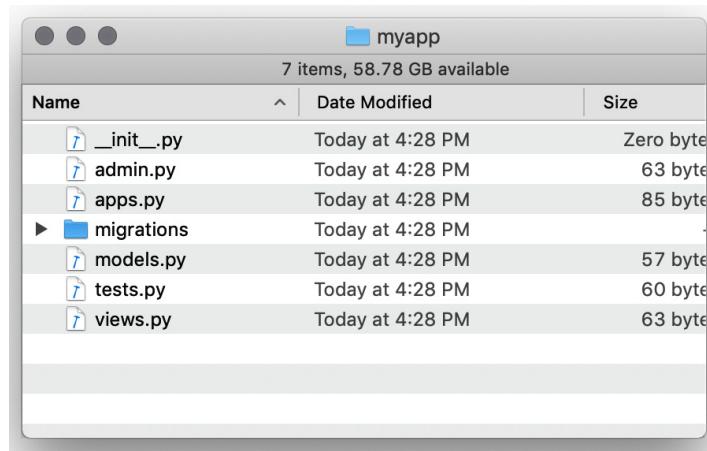


Figure 1.10 – The contents of the `myapp` app directory

Let's go over these files and a folder:

- `__init__.py`: An empty file, indicating that this directory is a Python module.
- `admin.py`: Django has a built-in admin site for viewing and editing data with a GUI. In this file, you will define how your app's models are exposed in the Django admin site. We'll cover this in more detail in *Chapter 4, Introduction to Django Admin*.
- `apps.py`: This contains some configuration for the metadata of your app. You won't need to edit this file.
- `models.py`: This is where you will define the models for your application. You'll read about this in more detail in *Chapter 2, Models and Migrations*.

- `migrations`: Django uses migration files to automatically record changes to your underlying database as models change. They are generated by Django when you run the `manage.py makemigrations` command and stored in this directory. They do not get applied to the database until you run `manage.py migrate`. They will also be covered in *Chapter 2, Models and Migrations*.
- `tests.py`: In order to test that your code is behaving correctly, Django supports writing tests (unit, functional, or integration) and will look for them inside this file. We will write some tests throughout this book and cover testing in detail in *Chapter 14, Testing Your Django Applications*.
- `views.py`: Your Django views (the code that responds to HTTP requests) will go in here. You will create a basic view soon, and views will be covered in more detail in *Chapter 3, URL Mapping, Views, and Templates*.

We will examine the contents of these files in more detail later, but for now, we'll get PyCharm up and running.

PyCharm setup

We confirmed in *Exercise 1.01 – creating a project and app, and starting the development server* that the Bookr project has been set up properly (since the development server runs successfully), so we can now start using PyCharm to run and edit our project. PyCharm is an IDE for Python development and includes features such as code completion, automatic style formatting, and a built-in debugger. We will then use PyCharm to start writing our own URL maps, views, and templates. It will also be used to start and stop the development server, which will allow debugging of your code by setting breakpoints.

Exercise 1.02 – project setup in PyCharm

In this exercise, we will open the Bookr project in PyCharm and set up the project interpreter so that PyCharm can run and debug the project. Follow these steps:

1. Open PyCharm. When you first open PyCharm, you will be shown the **Welcome to PyCharm** screen that asks you what you want to do:

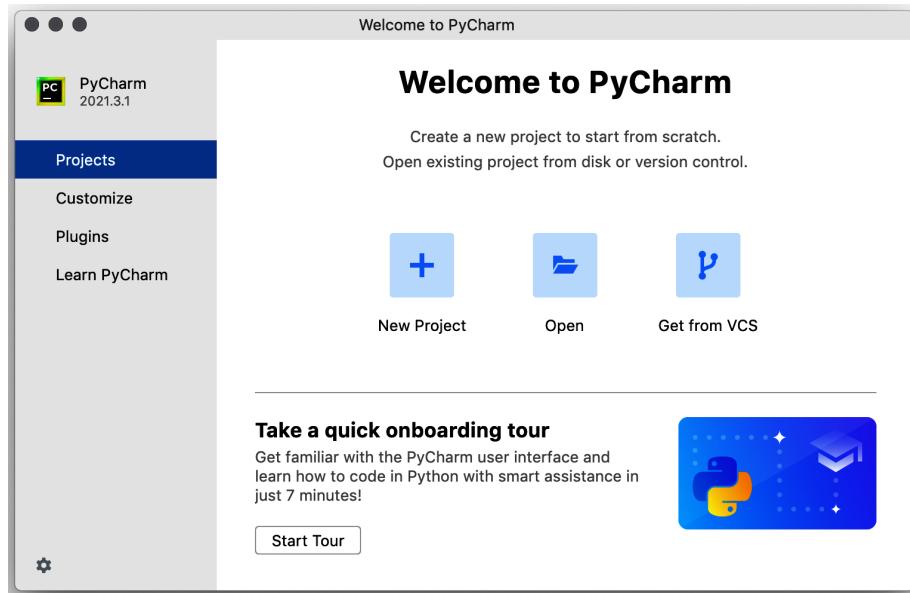


Figure 1.11 – The PyCharm welcome screen

2. Click on **Open**, and then browse to the `bookr` project you just created, then open it. Make sure you are opening the `bookr` project directory and not the `bookr` package directory inside.
3. You will be asked if you want to trust the `bookr` project. Since you just created it, it's safe, so click on **Trust Project**:

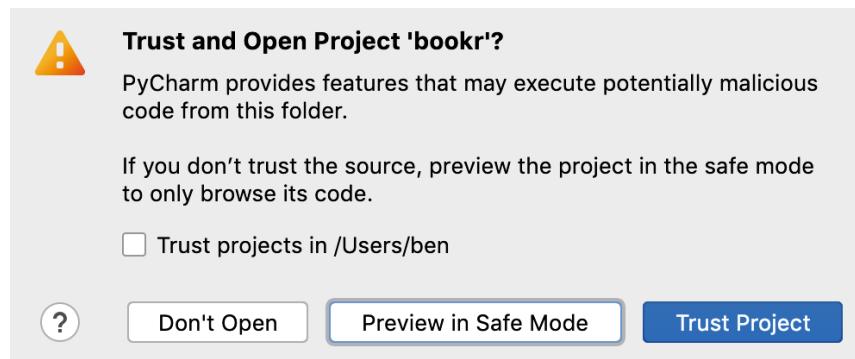


Figure 1.12 – PyCharm's Trust Project dialog

If you haven't used PyCharm before, it will ask you about what settings and themes you want to use, and once you have answered all those questions, you will see your `bookr` project structure open in the **Project** pane on the left of the window:

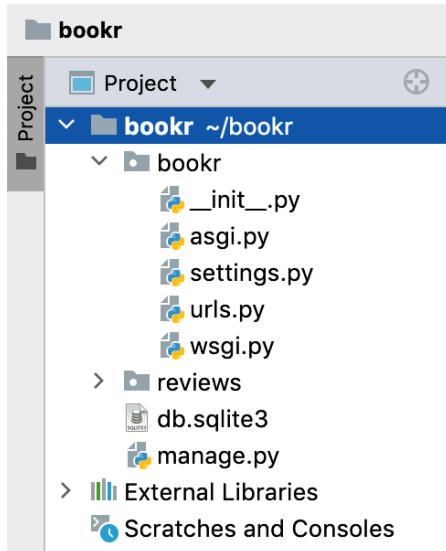


Figure 1.13 – The PyCharm project pane

Your project pane should look like *Figure 1.13* and show the `bookr` and `reviews` directories, and the `manage.py` file. If you do not see these and instead see `asgi.py`, `settings.py`, `urls.py`, and `wsgi.py`, then you have opened the `bookr` package directory instead. Select **File | Open**, and then browse and open the `bookr` project directory instead.

Before PyCharm knows how to execute your project to start the Django development server, the interpreter must be set to the Python binary inside your virtual environment. This is done first by adding the interpreter to the global interpreter settings.

4. Open the **Preferences** (macOS) or **Settings** (Windows/Linux) window inside PyCharm:
 - macOS: **PyCharm** menu | **Preferences**
 - Windows and Linux: **File | Settings**

5. In the preferences list pane on the left, open the **Project: bookr** item, and then click **Project Interpreter**:

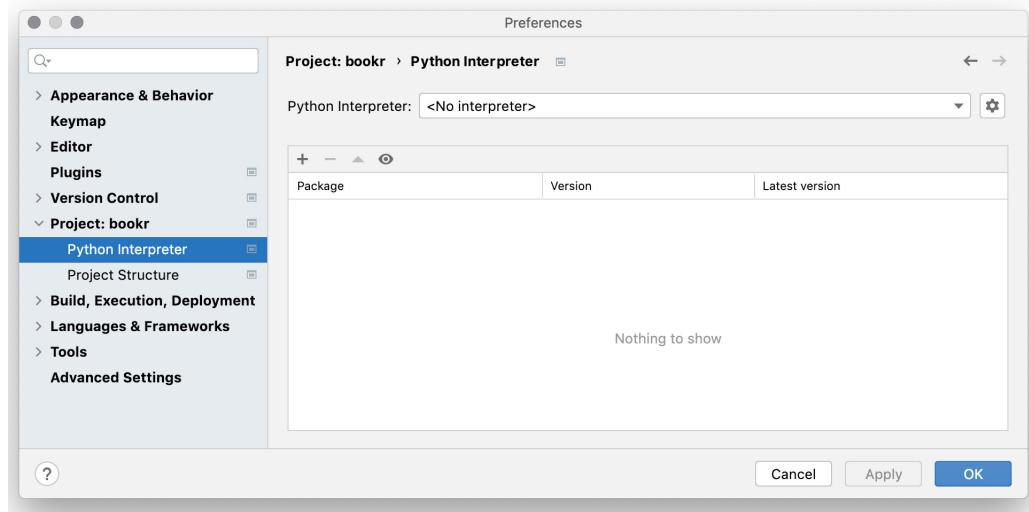


Figure 1.14 – The project interpreter settings

6. Sometimes, PyCharm is able to automatically determine the virtual environments, so in this case, the project interpreter may already be populated with the correct interpreter. If it has, and you see Django in the list of packages, click **OK** to close the window and complete this exercise.
7. In most cases though, the Python interpreter must be set manually. Click the cog icon next to the **Python Interpreter** dropdown, and then click **Add....**
8. The **Add Python Interpreter** window is now displayed. Select the **Existing environment** radio button, and then click the ellipses (...) next to the **Interpreter** selection. You should then browse and select the Python interpreter for your virtual environment.

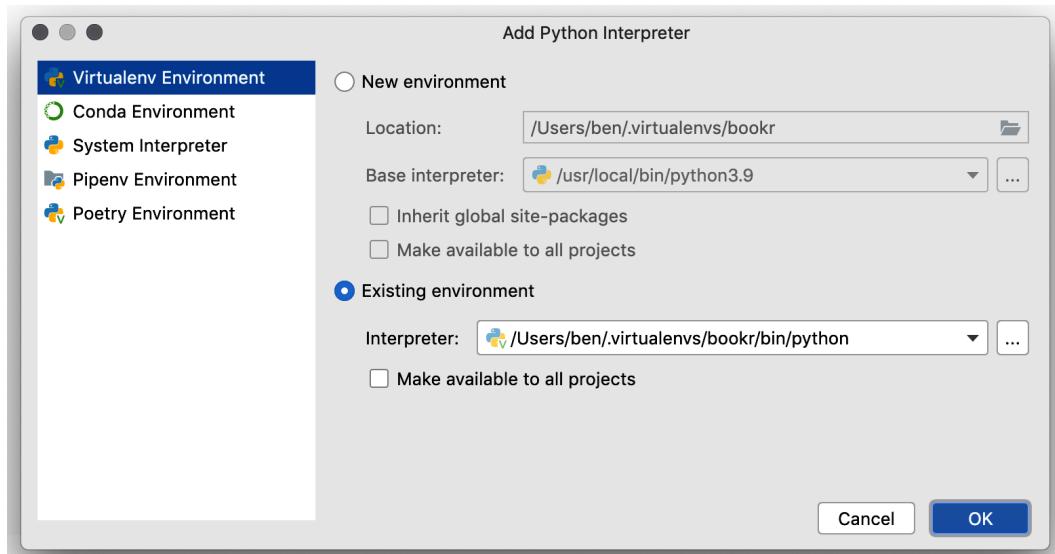


Figure 1.15 – The Add Python Interpreter window

9. On macOS (assuming you called the virtual environment `bookr`), the path is usually `/Users/<yourusername>/.virtualenvs/bookr/bin/python3`. Similar, in Linux, it should be in `/home/<yourusername>/.virtualenvs/bookr/bin/python3`.
10. If you're unsure, you can run the `which python3` command in the terminal where you previously ran the `python manage.py` command, and it will tell you the path to the Python interpreter:

```
WHICH PYTHON3
/USERS/BEN/.VIRTUALENVS/BOOKR/BIN/PYTHON3
```

On Windows, it will be wherever you created your virtual environment with the `virtualenv` command.

After selecting the interpreter, your **Add Python Interpreter** window should look like *Figure 1.15*.

11. Click **OK** to close the **Add Python Interpreter** window.

You should now see the main preferences window, and Django (and other packages in your virtual environment) will be listed (see *Figure 1.16*).

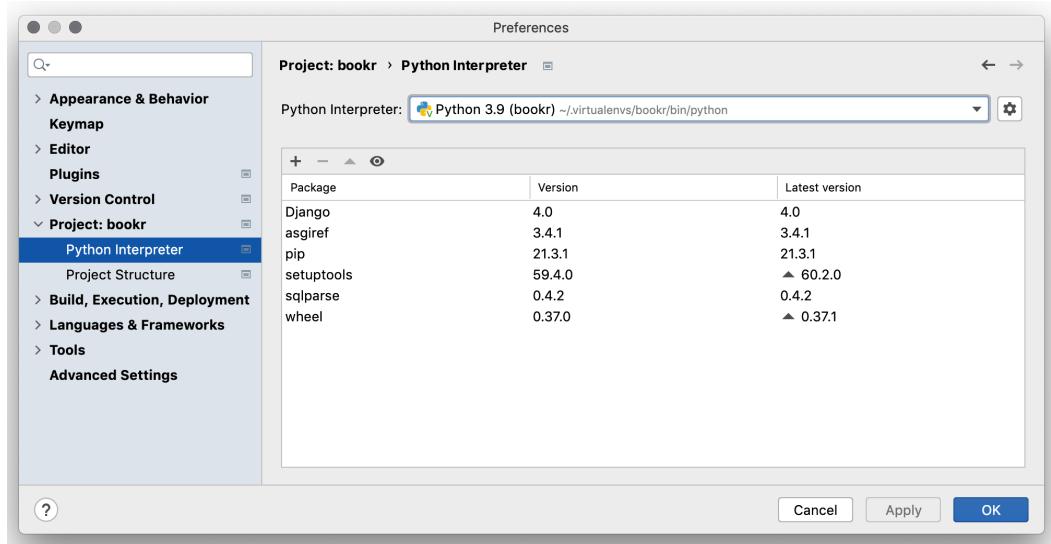


Figure 1.16 – Packages in the virtual environment are listed

12. Click **OK** in the main **Preferences** window to close it. PyCharm will now take a few seconds to index your environment and the libraries installed. You can see the process in its bottom-right status bar. Wait for this process to finish, and the progress bar will disappear.

In order to run the Django development server, Python needs to be configured with a run configuration. We will set this up now:

1. Click **Add Configuration...** in the top right of the PyCharm project window to open the **Run/Debug Configurations** window:

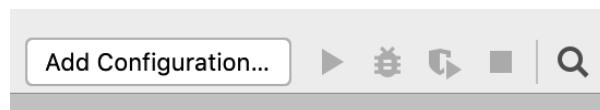


Figure 1.17 – The Add Configuration... button in the top right of the PyCharm window

2. Click the **+** button in the top left of this window, and select **Python** from the drop-down menu.

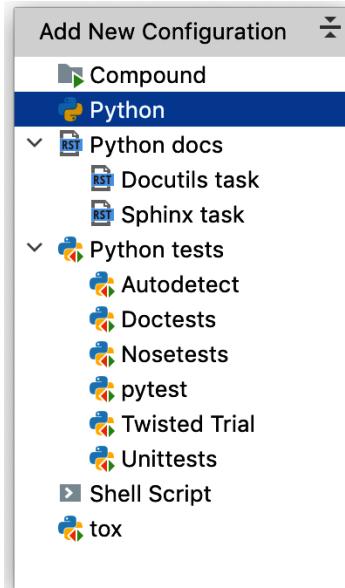


Figure 1.18 – Adding a new Python configuration in the Run/Debug Configurations window

3. A new configuration panel with fields regarding how to run your project will display on the right of the window. Go ahead and fill out the fields as follows:
 - The **Name** field can be anything but should be something understandable. Enter `Django Dev Server`.
 - **Script path** is the path to your `manage.py` file. If you click the folder icon in this field, you can browse your filesystem to select the `manage.py` file inside the `bookr` project directory.
 - **Parameters** refers to the arguments that come after the `manage.py` script, the same as if we were running it from the command line. We will use the same argument here to start the server, so enter `runserver`.

Note

As mentioned earlier, the `runserver` command can also accept an argument for the port or address to listen to. If you want to, you can add this argument after `runserver` in the same **Parameters** field.

4. The **Python interpreter** setting should have been automatically set to the one that was set in *Exercise 1.02*. If not, click the arrow dropdown on the right to select it.

The **Working directory** setting should be the `bookr` project directory. This has probably already been set correctly.

5. Make sure that **Add content roots to PYTHONPATH** and **Add source roots to PYTHONPATH** are checked. This will ensure PyCharm adds your `bookr` project directory to PYTHONPATH (the list of paths that the Python interpreter searches when loading a module). Without these checked, the imports from your project will not work correctly.

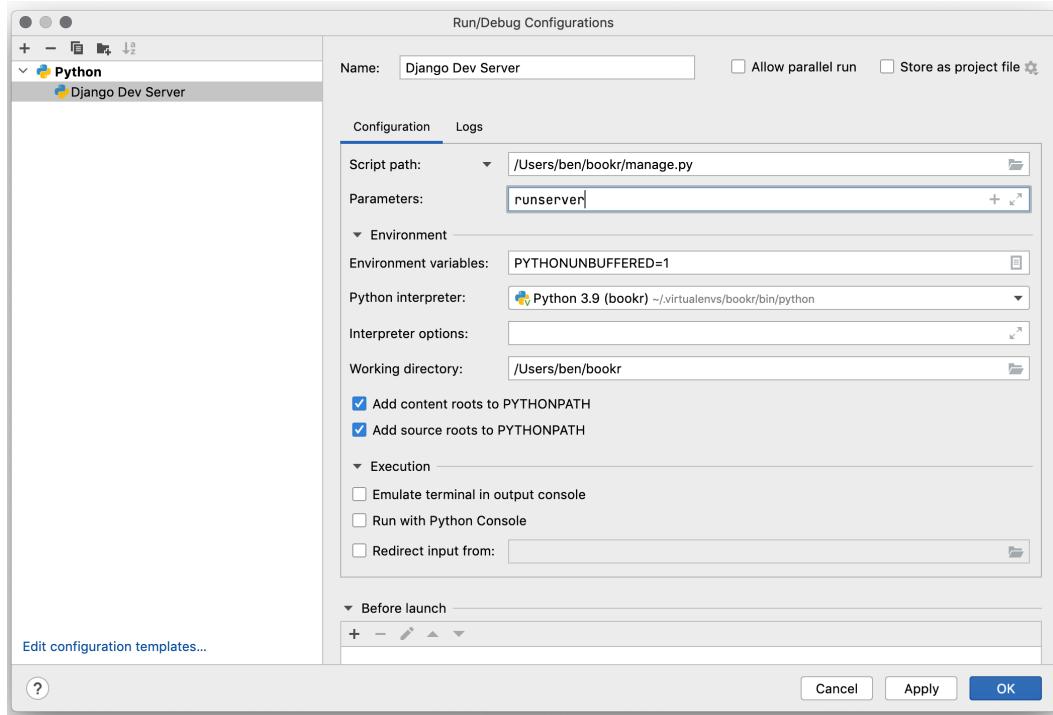


Figure 1.19 – The configuration settings

6. Check your **Run/Debug Configurations** window looks similar to *Figure 1.19*, and then click **OK** to save the configuration.
7. Now, instead of starting the Django development server in a terminal, click the **Play** icon in the top right of the project window to start it (see *Figure 1.20*).

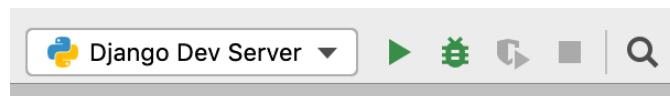


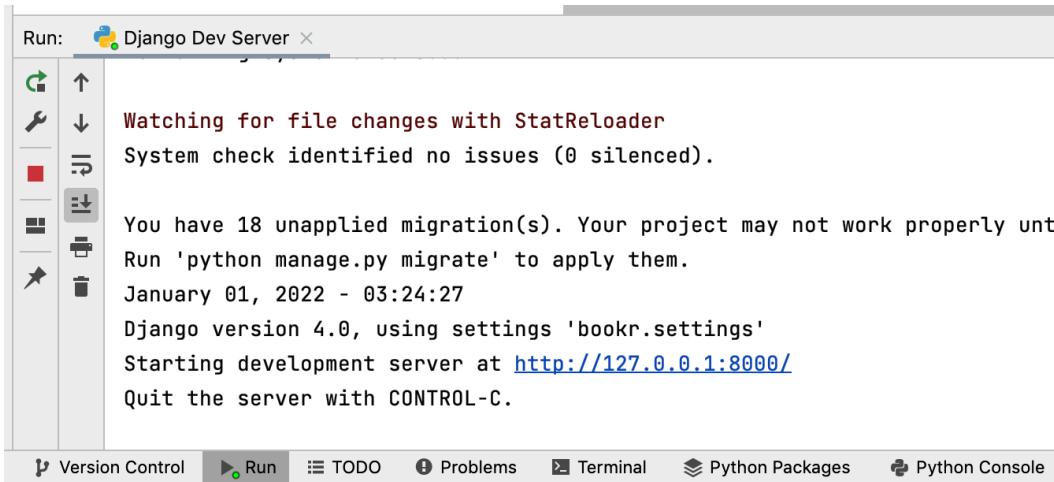
Figure 1.20 – Django development server configuration with the Play, Debug, and Stop buttons

8. Click the **Play** icon to start the Django development server.

Note

Make sure you stop any other instances of the Django development server that are running (such as in a terminal); otherwise, the one you are starting will not be able to bind to port 8000 and will fail to start.

A console will open at the bottom of the PyCharm window, which will show output indicating that the development server has started (*Figure 1.21*).



```
Run: Django Dev Server
Watching for file changes with StatReloader
System check identified no issues (0 silenced).

You have 18 unapplied migration(s). Your project may not work properly until
Run 'python manage.py migrate' to apply them.
January 01, 2022 - 03:24:27
Django version 4.0, using settings 'bookr.settings'
Starting development server at http://127.0.0.1:8000/
Quit the server with CONTROL-C.
```

Version Control Run TODO Problems Terminal Python Packages Python Console

Figure 1.21 – The console with the Django development server running

9. Open a web browser and navigate to `http://127.0.0.1:8000`. You should see the same Django example screen that you saw earlier, in *Exercise 1.01 – creating a project and app, and starting the development server* (*Figure 1.21*), which will confirm that, once again, everything is set up correctly.

In this exercise, we opened the `bookr` project in PyCharm, and then set the Python interpreter for our project. We then added a run configuration in PyCharm, which allows us to start and stop the Django development server from within PyCharm. We will also be able to debug our project later by running it inside PyCharm's debugger.

In this section, we explored the Django folder structure and saw the purpose of the directories and folders that were scaffolded. We also saw how to create a Django app and the files that it contains. Finally, we got a Django project up and running in PyCharm.

Introducing Django views

You now have everything set up to start writing your own Django views and configure the URLs that will map to them. As we saw earlier in this chapter, a view is simply a function that takes an `HttpRequest` instance (built by Django) and (optionally) some parameters from the URL. It will then perform some operations, such as fetching data from a database. Finally, it returns `HttpResponse`.

To use our Bookr app as an example, we might have a view that receives a request for a certain book. It queries the database for this book, and then returns a response containing an HTML page, showing information about the book. Another view could receive a request to list all the books, and then return a response with another HTML page containing this list. Views can also create or modify data; another view could receive a request to create a new book, and it would then add the book to the database and return a response with HTML that displays the new book's information.

In this chapter, we will only be using functions as views, but Django also supports class-based views that allow you to leverage object-oriented paradigms (such as inheritance). This allows you to simplify code used in multiple views that have the same business logic. For example, you might want to show all books or just books by a certain publisher. Both views need to query a list of books from the database and render them to a book list template. One view class could inherit from another and just implement the data fetching differently and leave the rest of the functionality (such as rendering) identical. Class-based views can be more powerful but also harder to learn. They will be introduced later, in *Chapter 11, Advanced Templates and Class-Based Views*, when you have more experience with Django.

The `HttpRequest` instance that is passed to the view contains all the data related to the request, with attributes such as the following:

- `method`: A string containing the HTTP method the browser used to request the page, usually `GET`, but it will be `POST` if a user has submitted a form. You can use this to change the flow of the view – for example, show an empty form on `GET`, or validate and process a form submission on `POST`.
- `GET`: This is a `QueryDict` class containing the parameters used in the URL query string. This is the part of the URL after `?`, if it contains one. We will go further into `QueryDict` classes soon. Note that this attribute is always available even if the request was not `GET`.
- `POST`: This is another `QueryDict` containing the parameters sent to the view in a `POST` request, such as from a form submission. Usually, you would use this in conjunction with a Django form, which will be covered in *Chapter 6, Forms*.
- `headers`: This is a case-insensitive key dictionary with the HTTP headers from the request. For example, you could vary the response with different content for different browsers based on the `User-Agent` header. We discussed some HTTP headers that are sent by the client earlier in this chapter.

- `path`: This is the path used in the request. Normally, you don't need to examine this because Django will automatically parse the path and pass it to the view function as parameters, but it can be useful in some instances.

We won't be using all these attributes yet, and there are others that will be introduced later, but you can now see what role the `HttpRequest` argument plays in your view.

Next, we'll see how Django determines what view to use for a specific URL.

Exploring URL mapping detail

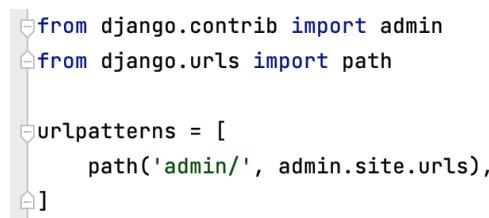
We briefly mentioned URL maps earlier in the *Processing a request* section. Django does not automatically know which view function should be executed when it receives a request for a particular URL. The role of URL mapping is to build a link between a URL and a view. For example, in Bookr, you might want to map the `/books/` URL to a `books_list` view that you have created.

The URL-to-view mapping is defined in the file that Django automatically created called `urls.py`, inside the `bookr` package directory (although a different file can be set in `settings.py`; there'll be more on that later).

This file contains a variable, `urlpatterns`, which is a list of paths that Django evaluates in turn until it finds a match for the URL being requested. The match will either resolve to a view function or another `urls.py` file, also containing a `urlpatterns` variable, which will be resolved in the same manner. URL files can be chained in this manner for as long as you want. This way, you can split URL maps into separate files (such as one or more per app) so that they don't become too large. Once a view has been found, Django calls it with `HttpRequest` and any parameters parsed from the URL.

Rules are set by calling the `path` function, which takes the path of the URL as the first argument. The path can contain named parameters that will be passed to a view as function parameters. Its second argument is either a view or another file also containing `urlpatterns`.

There is also the `re_path` function, which is similar to `path`, except that it takes a regular expression as the first argument for more advanced configuration. There is much more to URL mapping, however; it will be covered in *Chapter 3, URL Mapping, Views, and Templates*. To illustrate these concepts, *Figure 1.22* shows the default `urls.py` file that Django generates:



```
from django.contrib import admin
from django.urls import path

urlpatterns = [
    path('admin/', admin.site.urls),
]
```

Figure 1.22 – The default `urls.py` file

You can see the `urlpatterns` variable that lists all the URLs that are set up. Currently there is only one rule set up, which maps any path starting with `admin/` to the admin URL maps (the `admin.site.urls` module). This is not a mapping to a view; instead, it is an example of chaining URL maps together – the `admin.site.urls` module will define the remainder of the paths (after `admin/`) that map to the admin views. We will cover the Django admin site later in the book, starting in *Chapter 4, Introduction to Django Admin*.

We will now write a view and set up a URL map to it, to see these in action.

Exercise 1.03 – writing a view and mapping a URL to it

Our first view will be very simple and will just return some static text content. In this exercise, we will see how to write a view and set up a URL map to resolve to a view:

Note

As you make changes to files in your project and save them, you might see the Django development server automatically restarting in the terminal or console in which it is running. This is normal – it automatically restarts to load any code changes that you make. Please also note that it won't automatically apply changes to the database if you edit models or migrations – there'll be more on this in *Chapter 2, Models and Migrations*.

1. In PyCharm, expand the `reviews` folder in the **Project** browser on the left, and then double-click on the `views.py` file inside to open it. In the right (editor) pane in PyCharm, you should see the placeholder text automatically generated by Django:

```
from django.shortcuts import render

# Create your views here.
```

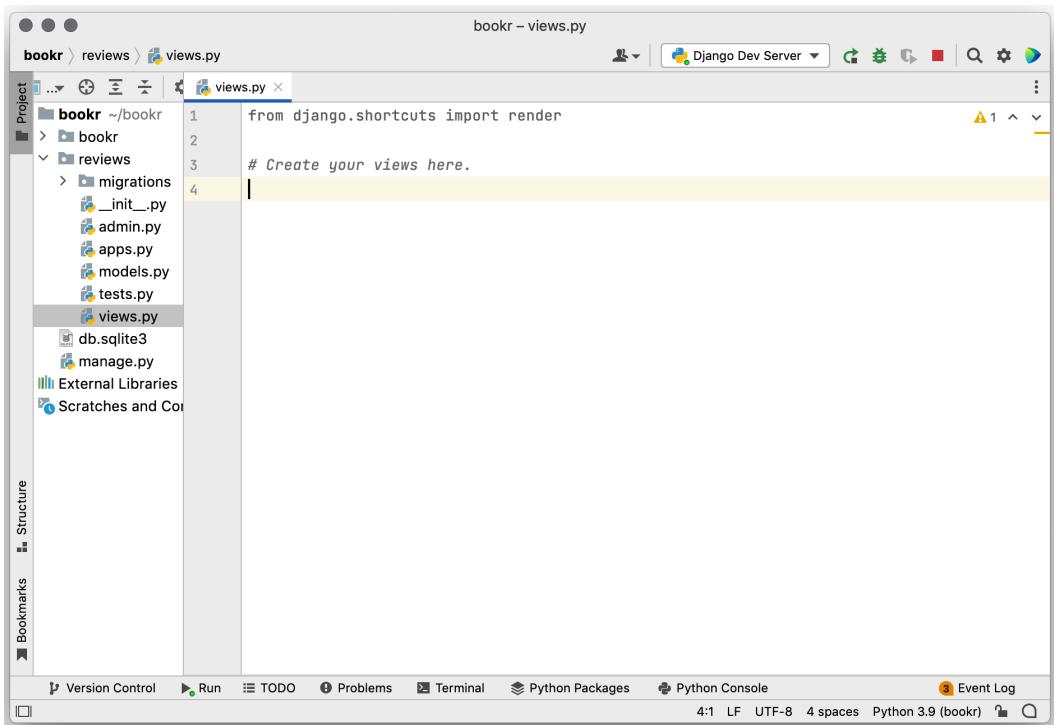


Figure 1.23 – The views.py default content

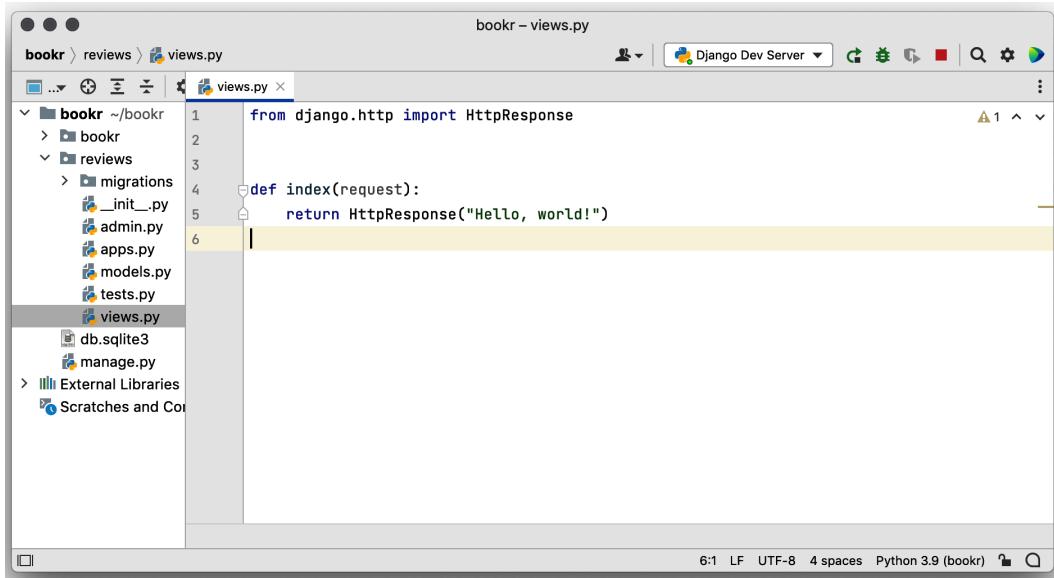
2. Remove this placeholder text from `views.py` and instead insert this content:

```
from django.http import HttpResponse

def index(request):
    return HttpResponse("Hello, world!")
```

First, the `HttpResponse` class needs to be imported from `django.http`. This is what is used to create the response that goes back to the web browser. You can also use it to control things such as the HTTP headers or status code. For now, it will just use the default headers and the `200 Success` status code. Its first argument is the string content to send as the body of the response.

Then, the view function returns an `HttpResponse` instance with the content we defined (`Hello, world!`).



```
bookr - views.py
bookr > reviews > views.py
views.py x
bookr ~/bookr
> bookr
> reviews
  > migrations
    __init__.py
    admin.py
    apps.py
    models.py
    tests.py
  views.py
  db.sqlite3
  manage.py
> External Libraries
> Scratches and Co

from django.http import HttpResponse
def index(request):
    return HttpResponse("Hello, world!")

6:1 LF UTF-8 4 spaces Python 3.9 (bookr)
```

Figure 1.24 – The contents of `views.py` after editing

3. We will now set up a URL map to the `index` view. This will be very simple and not contain any parameters. Expand the `bookr` directory in the project pane, and then open `urls.py`. Django has automatically generated this file.

For now, we'll just add a simple URL to replace the default `index` that Django provides.

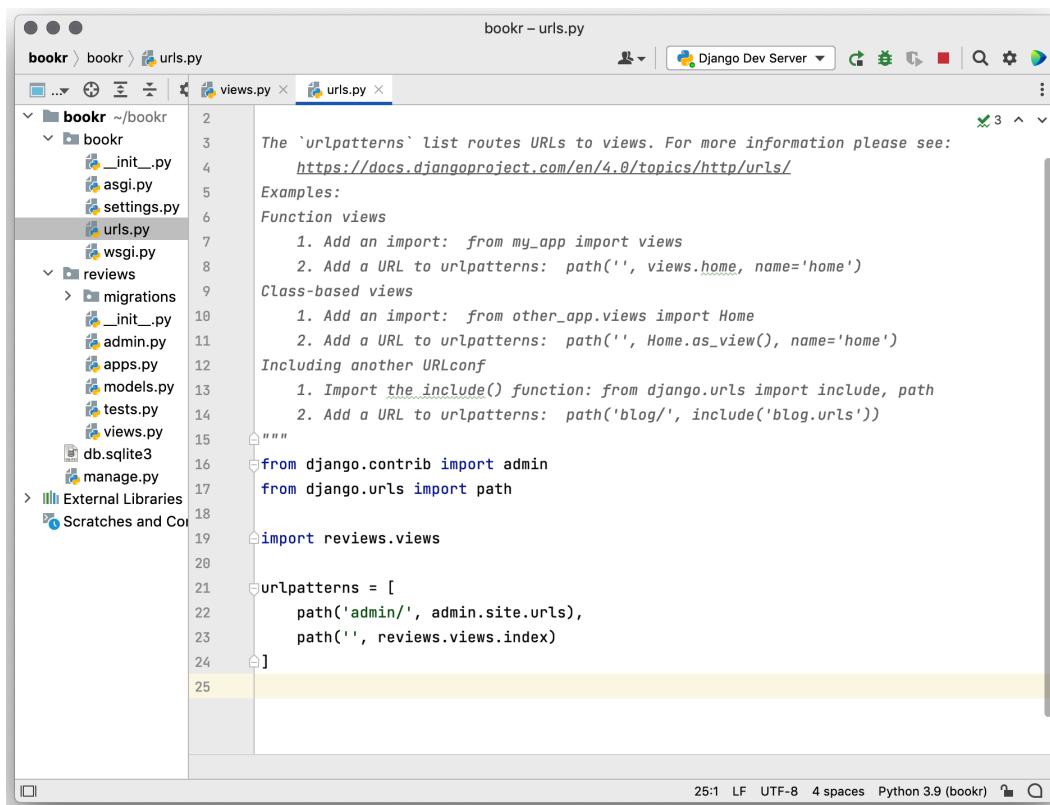
4. Import your views into the `urls.py` file by adding this line after the other existing imports:

```
import reviews.views
```

5. Add a map to the `index` view to the `urlpatterns` list by adding a call to the `path` function, with an empty string and a reference to the `index` function:

```
urlpatterns = [
    path('admin/', admin.site.urls),
    path('', reviews.views.index)
]
```

Make sure you don't add brackets after the `index` function (i.e., it should be `reviews.views.index` and not `reviews.views.index()`), as we are passing a reference to a function rather than calling it. When you're finished, your `urls.py` file should look like *Figure 1.25*.



```
bookr - urls.py
bookr / bookr / urls.py
views.py x urls.py x

bookr ~ /bookr
bookr
  __init__.py
  asgi.py
  settings.py
  urls.py
  wsgi.py
reviews
  migrations
    __init__.py
    admin.py
    apps.py
    models.py
    tests.py
    views.py
  db.sqlite3
  manage.py
External Libraries
Scratches and Code

The 'urlpatterns' list routes URLs to views. For more information please see:
    https://docs.djangoproject.com/en/4.0/topics/http/urls/
Examples:
Function views
1. Add an import: from my_app import views
2. Add a URL to urlpatterns: path('', views.home, name='home')
Class-based views
1. Add an import: from other_app.views import Home
2. Add a URL to urlpatterns: path('', Home.as_view(), name='home')
Including another URLconf
1. Import the include() function: from django.urls import include, path
2. Add a URL to urlpatterns: path('blog/', include('blog.urls'))
```
from django.contrib import admin
from django.urls import path

import reviews.views

urlpatterns = [
 path('admin/', admin.site.urls),
 path('', reviews.views.index)
]
```
25:1 LF UTF-8 4 spaces Python 3.9 (bookr) 
```

Figure 1.25 – urls.py after editing

6. Switch back to your web browser and refresh. The Django default welcome screen should be replaced with the text defined in the View, **Hello, world!**.



Figure 1.26 – The web browser should now display the Hello, world! message

We just saw how to write a view function and map a URL to it. We then tested the view by loading it in a web browser.

Next, we'll see how to work with data that is sent in an HTTP request, but not in the URL path.

Working with GET, POST, and QueryDict objects

Data can come through an HTTP request as parameters on a URL or inside the body of a POST request. You might have noticed parameters in a URL when browsing the web – the text after ? – for example, `http://www.example.com/?parameter1=value1¶meter2=value2`. We also saw earlier in this chapter an example of form data in a POST request to log in a user (the request body was `username=user&password=password1`).

Django automatically parses these parameter strings into `QueryDict` objects. The data is then available on the `HttpRequest` object that is passed to your view – specifically, in the `HttpRequest.GET` and `HttpRequest.POST` attributes for URL parameters and body parameters respectively. `QueryDict` objects mostly behave like dictionaries, except that they can contain multiple values for a key.

To show different methods of accessing items in, we'll use a simple `QueryDict` (the `qd` variable) with only one key (`k`) as an example. The `k` item has three values in a list – the `a`, `b`, and `c` strings. The following code snippets show output from a Python interpreter.

First, the `QueryDict` `qd` variable is constructed from a parameter string:

```
>>> qd = QueryDict("k=a&k=b&k=c")
```

When accessing items with square bracket notation or the `get` method, the last value for that key is returned:

```
>>> qd["k"]
'c'
>>> qd.get("k")
'c'
```

To access all the values for a key, the `getlist` method should be used, as follows:

```
>>> qd.getlist("k")
['a', 'b', 'c']
```

`getlist` will always return a list; it will be empty if the key does not exist:

```
>>> qd.getlist("bad key")
[]
```

While `getlist` does not raise an exception for keys that do not exist, accessing a key that does not exist with square bracket notation will raise `KeyError`, like a normal dictionary. Use the `get` method to avoid this error.

The `QueryDict` objects for GET and POST are immutable (they cannot be changed), so the `copy` method should be used to get a mutable copy if you need to change its values, as follows:

```
>>> qd["k"] = "d"
AttributeError: This QueryDict instance is immutable
>>> qd2 = qd.copy()
>>> qd2
<QueryDict: {'k': ['a', 'b', 'c']}>
>>> qd2["k"] = "d"
>>> qd2["k"]
"d"
```

To give an example of how `QueryDict` is populated from a URL, let's imagine an example URL: `http://127.0.0.1:8000?val1=a&val2=b&val2=c&val3`.

Behind the scenes, Django passes the query from the URL (everything after `?`) to instantiate a `QueryDict` object and attach it to `request` that is passed to the view function. It goes something like this:

```
request.GET = QueryDict("val1=a&val2=b&val2=c&val3")
```

Remember, this is done to `request` before you receive it inside your view function, so you do not need to do this.

In the case of the preceding example URL, we can access the parameters inside the view function as follows:

```
request.GET["val1"]
```

Using standard dictionary access, it would return the "a" value:

```
request.GET["val2"]
```

Again, using standard dictionary access, there are two values set for the `val2` key, so it would return the last value, "c":

```
request.GET.getlist("val2")
```

This would return a list of all the values for `val2` - `["b", "c"]`:

```
request.GET["val3"]
```

This key is in the query string but has no value set, so this returns an empty string:

```
request.GET["val4"]
```

This key is not set, so `KeyError` will be raised. Use `request.GET.get("val4")` instead, which will return `None`:

```
request.GET.getlist("val4")
```

Since this key is not set as an empty list, `([])` will be returned.

We will now look at `QueryDict` in action using the GET parameters. We will examine POST parameters further in *Chapter 6, Forms*.

Exercise 1.04 – exploring GET values and QueryDict objects

We will now make some changes to our `index` view from the previous exercise to read values from the URL in the `GET` attribute, and then we will experiment with passing different parameters to see the result:

1. Open the `views.py` file in PyCharm. Add a new variable called `name`, which reads the user's name from the `GET` parameters. Add this line after the `index` function definition:

```
name = request.GET.get("name") or "world"
```

2. Change the return value so that the `name` is used as part of the content that is returned:

```
return HttpResponseRedirect("Hello, {}!".format(name))
```

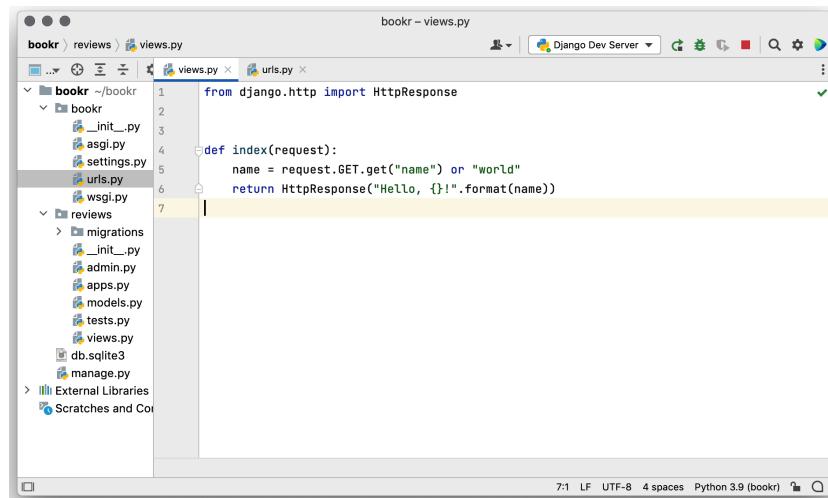


Figure 1.27 – The updated `views.py` reading `name` parameter

3. Visit `http://127.0.0.1:8000` in your browser. Note that the page still says **Hello, world!**. This is because we have not supplied a name parameter. You can add your name to the URL – for example, `http://127.0.0.1:8000?name=Ben`. This can be seen in the following screenshot:

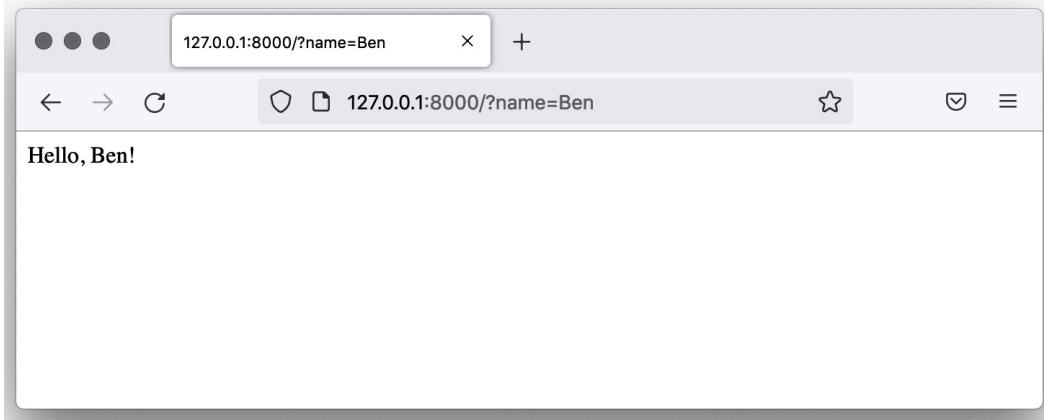


Figure 1.28 – Setting the name in the URL

4. Try adding two names – for example, `http://127.0.0.1:8000?name=Ben&name=John`. As we mentioned, the last value for the parameter is retrieved with the `get` function, so you should see **Hello, John!**:

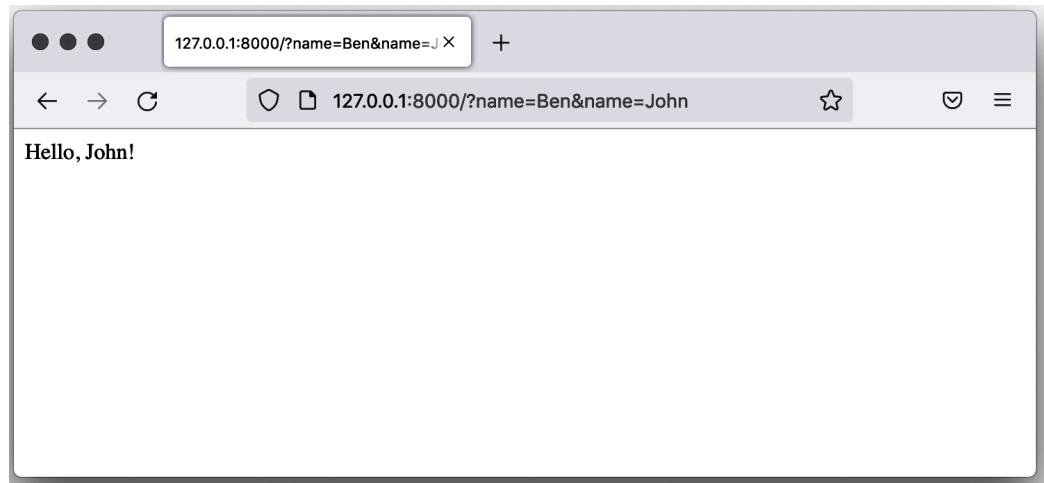


Figure 1.29 – Setting multiple names in the URL

5. Try setting no name, like this – `http://127.0.0.1:8000?name=`. The page should go back to displaying **Hello, world!**:

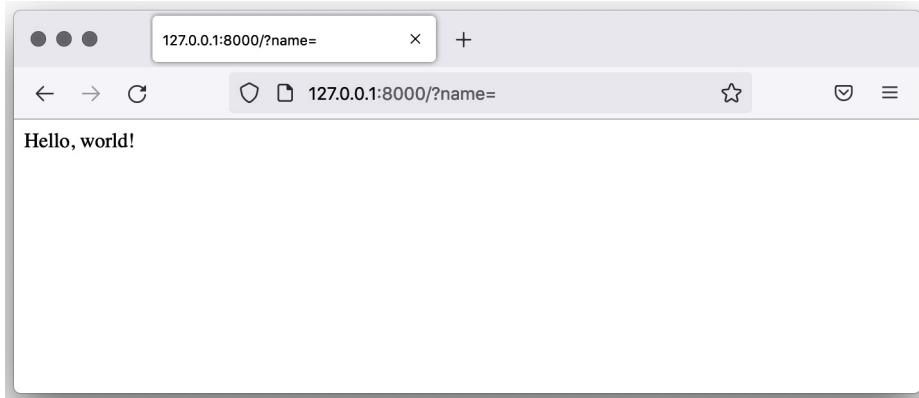


Figure 1.30 – No name set in the URL

Note

You might wonder why we set `name` to the default **world** by using `or`, instead of passing '`world`' as the default value to `get`. Consider what happened in *step 6*, when we passed in a blank value for the `name` parameter. If we had passed '`world`' as a default value for `get`, then the `get` function would still have returned an empty string. This is because a value *is* set for `name`; it's just that it's blank. Keep this in mind when developing your views, as there is a difference between no value being set and a blank value being set. Depending on your use case, you might choose to pass the default to `get`.

In this exercise, we retrieved values from the URL in our view using the `GET` attribute of the incoming request. We saw how to set default values and which value is retrieved if multiple values are set for the same parameter. The view returned `HttpResponse`, containing the message we wanted to send to the browser.

In the next section, we'll learn about how to update and use Django's settings.

Exploring Django settings

We haven't yet looked at how Django stores its settings. Now that we've seen the different parts of Django, it is a good time to examine the `settings.py` file. There are many settings Django contains that can be used to customize it. A default `settings.py` file was created for you when you started the Bookr project. We will discuss some of the more important settings in the file now, and a few others that might be useful as you become more fluent with Django. You should open your `settings.py` file in PyCharm and follow along so that you can see where and what the values are for your project.

Each setting in this file is just a file-global variable. The order in which we will discuss the settings is the same order in which they appear in this file, although we may skip over some – for example, there is the `ALLOWED_HOSTS` setting between `DEBUG` and `INSTALLED_APPS`, which we won’t cover in this part of the book (you’ll see it in *Chapter 16, Deployment of a Django Application*):

```
SECRET_KEY = '...'
```

This is an automatically generated value that shouldn’t be shared with anyone. It is used for hashing, tokens, and other cryptographic functions. If you had existing sessions in a cookie and changed this value, the sessions would no longer be valid:

```
DEBUG = True
```

With this value set to `True`, Django will automatically display exceptions to the browser to allow you to debug any problems you encounter. It should be set to `False` when deploying your app to production:

```
INSTALLED_APPS = [...]
```

When you write your own Django apps (such as the `reviews` app) or install third-party applications (which will be covered in *Chapter 15, Plugging in Third-Party Libraries*), they should be added to this list. As we’ve seen, it is not strictly necessary to add them here (our `index` view worked without our `reviews` app being in this list). However, for Django to be able to automatically find the app’s templates, static files, migrations, and other configurations, it must be listed here:

```
ROOT_URLCONF = 'bookr.urls'
```

This is the Python module that Django will first load to find URLs. Note that it is the file we added our `index` view URL map to previously:

```
TEMPLATES = [...]
```

Right now, it’s not too important to understand everything in this setting, as you won’t be changing it; the important line to point out is this one:

```
'APP_DIRS': True,
```

This tells Django that it should look in a `templates` directory inside each `INSTALLED_APP` when loading a template to render. We don’t have a `templates` directory for `reviews` yet, but we will add one in the next exercise.

Django has more settings available that aren’t listed in the `settings.py` file, so it will use its built-in defaults in these cases. You can also use the file to set arbitrary settings that you make up for your application. Third-party applications might want settings to be added here as well. In later chapters, we will add settings here for other applications. You can find a list of all settings, and their defaults, at <https://docs.djangoproject.com/en/3.0/ref/settings/>.

Referring to Django Settings in Your Code

It can sometimes be useful to refer to settings from `settings.py` in your own code, whether they are Django's built-in settings or ones you have defined yourself. You might be tempted to write code such as this import settings:

```
from bookr import settings

if settings.DEBUG:  # check if running in DEBUG mode
    do_some_logging()
```

This method is incorrect, for several reasons:

- It is possible to run Django and specify a different settings file to read from, in which case the previous code would cause an error, as it would not be able to find that file. Alternatively, if the file exists, the import would succeed but contain the wrong settings.
- Django has settings that might not be listed in the `settings.py` file, and if they aren't, it will use its own internal default. For example, if you removed the `DEBUG = True` line from your `settings.py` file, Django would fall back to using its internal value for `DEBUG` (which is `False`). You would get an error if you tried to access it using `settings.DEBUG` directly though.
- Third-party libraries can change how your settings are defined, so your `settings.py` file would look completely different. None of the expected variables may exist at all. The behavior of all these applications is beyond the scope of this book, but it is something to be aware of.

The preferred way is to use the `django.conf` module instead, like this:

```
from django.conf import settings  # import settings from here instead

if settings.DEBUG:
    do_some_logging()
```

When importing settings from `django.conf`, Django mitigates the three issues we just discussed:

- Settings are read from whatever Django settings file has been specified
- Any default settings values are interpolated
- Django takes care of parsing any settings defined by a third-party library

In our new, shorter example code snippet, even if `DEBUG` is missing from the `settings.py` file, it will fall back to the default value that Django has internally (which is `False`). The same is true for all other settings that Django defines; however, if you define your own custom settings in this file, Django will not have internal values for them, so in your code, you should have some provision for them not existing – how your code behaves is your choice and beyond the scope of this book.

Next, we'll see how Django locates template files.

Finding HTML templates in app directories

There are many options that are available to tell Django how to find templates, which can be set in the `TEMPLATES` setting of `settings.py`, but the easiest one (for now) is to create a `templates` directory inside the `reviews` directory. Django will look in this (and in other apps' `templates` directories) because of `APP_DIRS` being `True` in the `settings.py` file, as we saw in the previous section. However, for Django to know that the `reviews` directory is an app, we need to configure it in the `settings`. We'll do that in the next exercise.

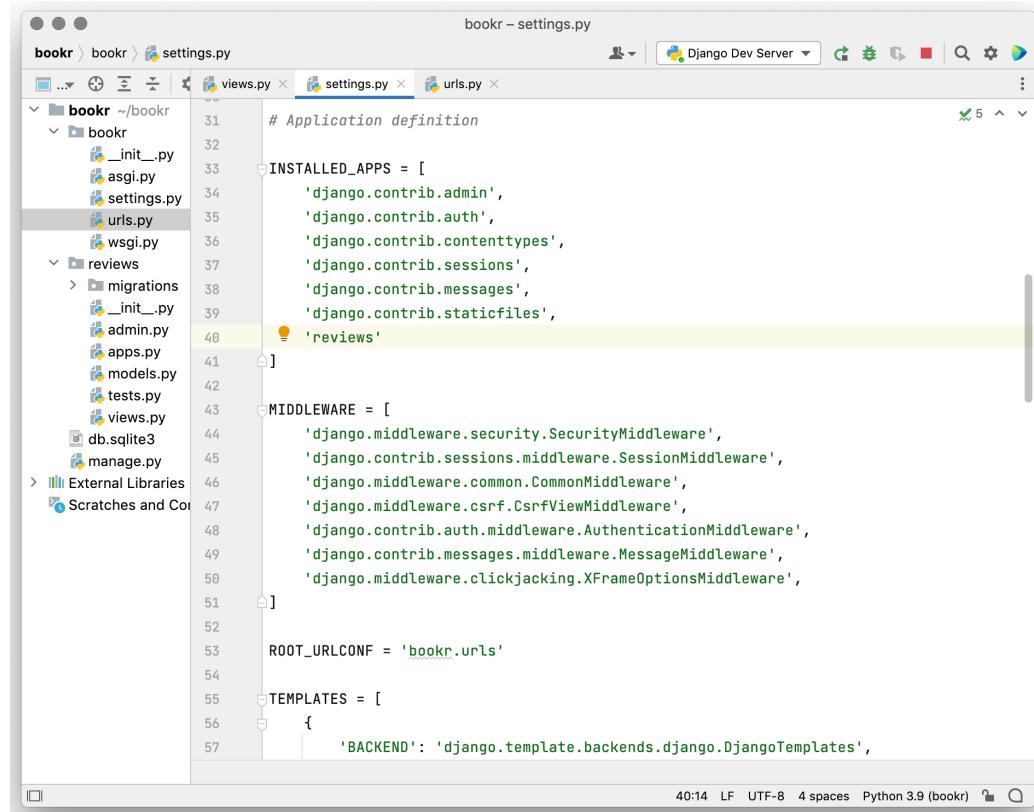
Exercise 1.05 – creating a templates directory and a base template

In this exercise, you will create a `templates` directory for the `reviews` app. Then, you will add an HTML template file that Django will be able to render to an HTTP response.

We discussed `settings.py` and its `INSTALLED_APPS` setting in the *Exploring Django settings* section. We need to add the `reviews` app to `INSTALLED_APPS` for Django to be able to find templates. Follow these steps:

1. Open `settings.py` in PyCharm. Update the `INSTALLED_APPS` setting and add `reviews` to the end. It should look like this:

```
INSTALLED_APPS = [  
    'django.contrib.admin',  
    'django.contrib.auth',  
    'django.contrib.contenttypes',  
    'django.contrib.sessions',  
    'django.contrib.messages',  
    'django.contrib.staticfiles',  
    'reviews'  
]
```



```

bookr - settings.py
bookr > bookr > settings.py
  views.py <--> settings.py <--> urls.py
  book
    __init__.py
    asgi.py
    settings.py
    urls.py
    wsgi.py
  reviews
    migrations
      __init__.py
    admin.py
    apps.py
    models.py
    tests.py
    views.py
  db.sqlite3
  manage.py
> External Libraries
> Scratches and Co

# Application definition

INSTALLED_APPS = [
    'django.contrib.admin',
    'django.contrib.auth',
    'django.contrib.contenttypes',
    'django.contrib.sessions',
    'django.contrib.messages',
    'django.contrib.staticfiles',
    'reviews'
]

MIDDLEWARE = [
    'django.middleware.security.SecurityMiddleware',
    'django.contrib.sessions.middleware.SessionMiddleware',
    'django.middleware.common.CommonMiddleware',
    'django.middleware.csrf.CsrfViewMiddleware',
    'django.contrib.auth.middleware.AuthenticationMiddleware',
    'django.contrib.messages.middleware.MessageMiddleware',
    'django.middleware.clickjacking.XFrameOptionsMiddleware',
]

ROOT_URLCONF = 'bookr.urls'

TEMPLATES = [
    {
        'BACKEND': 'django.template.backends.django.DjangoTemplates',
    }
]

```

Figure 1.31 – The reviews app added to settings.py

2. Save and close `settings.py`.
3. In the PyCharm **Project** browser, right-click the `reviews` directory and select **New | Directory**.

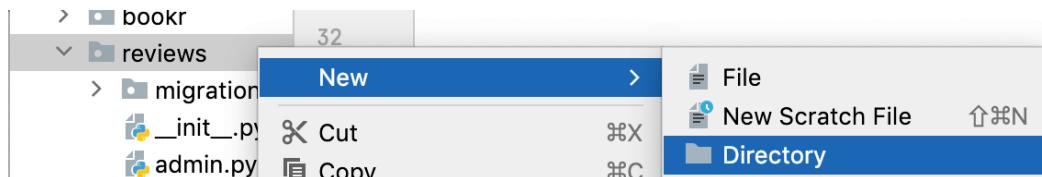


Figure 1.32 – Creating a new directory inside the reviews directory

4. Enter the name `templates` and press *Enter/Return* to create it:

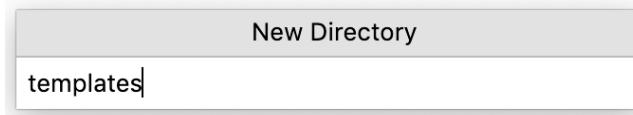


Figure 1.33 – Naming the directory templates

5. Right-click the newly created `templates` directory and select **New | HTML File**.

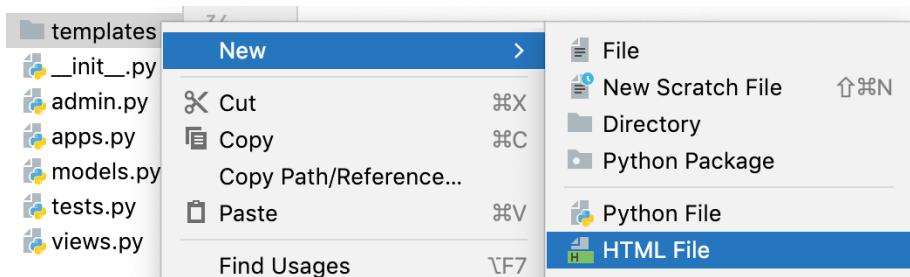


Figure 1.34 – Creating a new HTML file in the templates directory

6. In the window that appears, enter the name `base.html`, leave **HTML 5 file** selected, and then press *Enter/Return* to create the file:

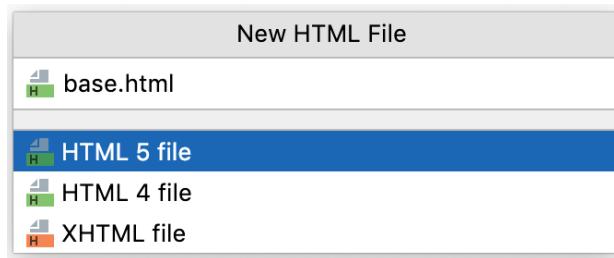


Figure 1.35 – The new HTML file window

7. After PyCharm creates the file, it will automatically open it too. It will have this content:

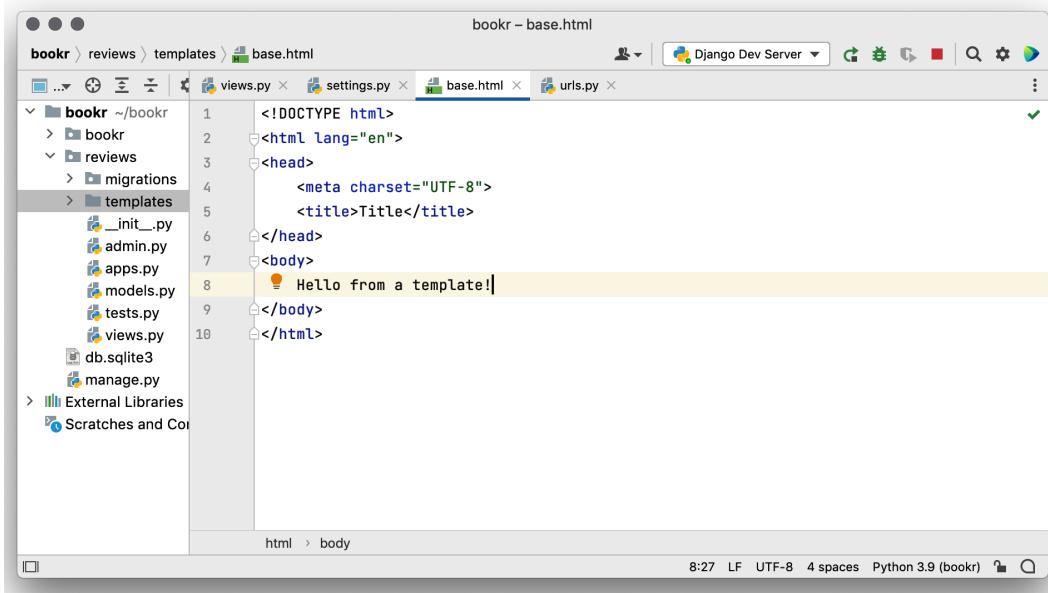
```
<!DOCTYPE html>
<html lang="en">
<head>
    <meta charset="UTF-8">
    <title>Title</title>
</head>
<body>

</body>
</html>
```

8. Between the `<body>` and `</body>` tags, add a short message to validate that the template is being rendered:

```
<body>
    Hello from a template!
</body>
```

Here's a screenshot of what your full code should look like:



The screenshot shows the PyCharm interface with the 'bookr' project open. The 'base.html' file is the active tab in the top navigation bar. The code editor displays the following HTML template:

```
<!DOCTYPE html>
<html lang="en">
<head>
    <meta charset="UTF-8">
    <title>Title</title>
</head>
<body>
    Hello from a template!!
</body>
</html>
```

The line `Hello from a template!!` is highlighted with a yellow background, indicating it is the current selection. The PyCharm status bar at the bottom shows the file is saved, with the last edit time as 8:27, line endings as LF, encoding as UTF-8, 4 spaces indentation, and Python 3.9 (bookr) as the language.

Figure 1.36 – The base.html template with some example text

In this exercise, we created a `templates` directory for the `reviews` app and added an HTML template to it. The HTML template will be rendered once we implement the use of the `render` function on our view.

Rendering a template with the `render` function

We now have a template to use, but we need to update our `index` view so that it renders the template, instead of returning the *Hello (name)!* text that it is currently displaying (refer to *Figure 1.30* to see how it currently looks). We will do this by using the `render` function and providing the name of the template. `render` is a shortcut function that returns `HttpResponse`. There are other ways to render a template that provide more control over how it is rendered, but for now, this function is fine for our needs. `render` takes at least two arguments – the first is always the request that was passed to the view, and the second is the name/relative path of the template being rendered. We will also call it with a third argument, the `render` context that contains all the variables that will be available in the template – there'll be more on this in *Exercise 1.07 – using variables in templates*.

Exercise 1.06 – rendering a template in a view

In this exercise, you will update your `index` view function to render the HTML template you created in *Exercise 1.05 – creating a templates directory and a base template*. You will make use of the `render` function, which loads your template from disk, renders it, and sends it to the browser. This will replace the static text you are currently returning from the `index` view function:

1. In PyCharm, open `views.py` in the `reviews` directory.
2. We no longer manually create `HttpResponse`, so remove the `HttpResponse` import line:

```
from django.http import HttpResponse
```

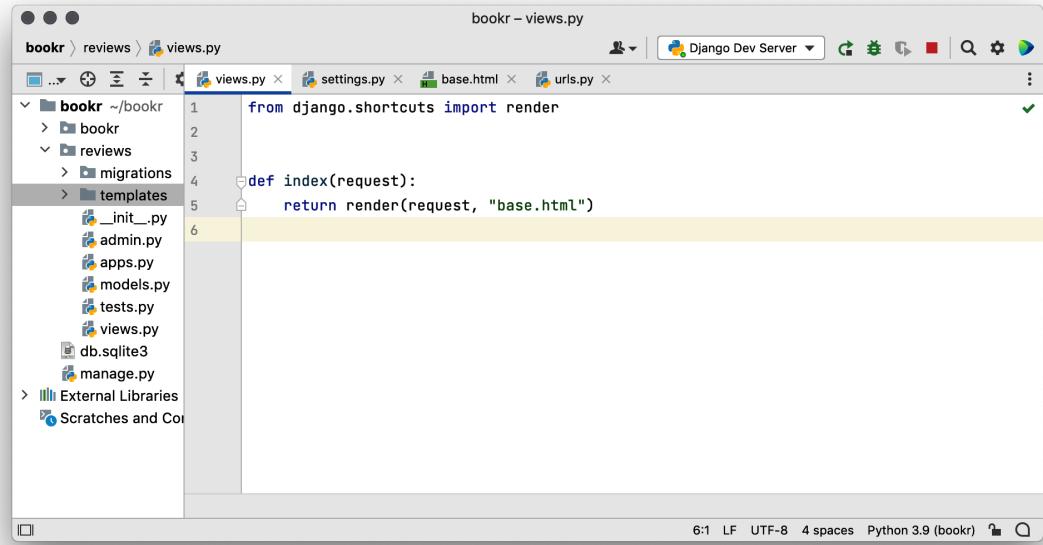
3. Replace it with an import of the `render` function from `django.shortcuts`:

```
from django.shortcuts import render
```

4. Update the `index` function so that, instead of returning `HttpResponse`, it's returning a call to `render`, passing in `request` and the template name:

```
def index(request):  
    return render(request, "base.html")
```

The following screenshot shows the whole `views.py` file:



```
bookr - views.py
bookr > reviews > views.py
bookr ~/bookr
bookr
reviews
migrations
templates
__init__.py
admin.py
apps.py
models.py
tests.py
views.py
db.sqlite3
manage.py
External Libraries
Scratches and Code

from django.shortcuts import render

def index(request):
    return render(request, "base.html")
```

Figure 1.37 – The completed `view.py` with the template render

5. Start the development server if it's not already running. Then, open your web browser and refresh `http://127.0.0.1:8000`. You should see the **Hello from a template!** message rendered, as shown in *Figure 1.38*:

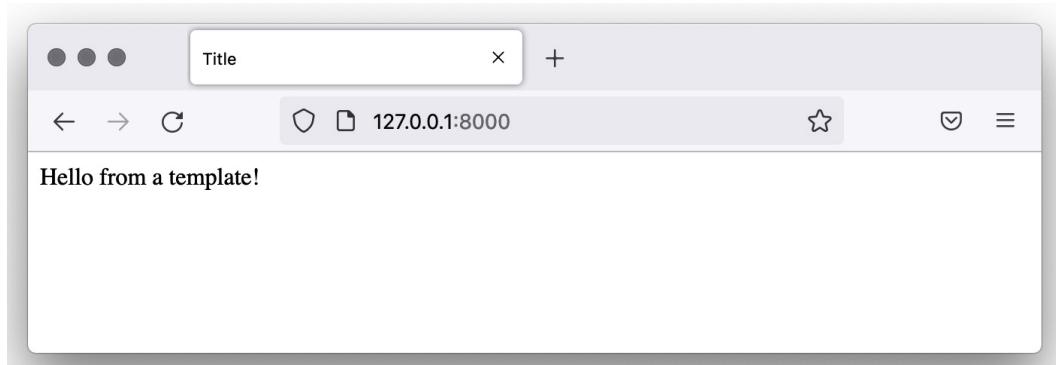


Figure 1.38 – Your first rendered HTML template

Rendering variables in templates

Templates aren't just static HTML. Most of the time, they will contain variables that are interpolated as part of the rendering process. They are passed from a view to a template using a context – a dictionary (or dictionary-like object) that contains names of all the variables a template can use. Let's take Bookr again as an example. Without variables in your template, you would need a different HTML file for each book you wanted to display. Instead, we use a variable such as `book_name` inside the template, and then the view provides the template with a `book_name` variable, set to the title of the book model it has loaded. When displaying a different book, the HTML does not need to change; the view just passes a different book to it. You can see how models, views, and templates are all now coming together.

Unlike some other languages, such as PHP, variables must be explicitly passed to a template, and variables in a view aren't automatically available to the template. This is for security reasons as well as to avoid accidentally polluting the template's namespace (we don't want any unexpected variables in the template).

Inside a template, variables are denoted by double braces – `{ { } }`. While not strictly a standard, this style is quite common and used in other templating tools such as Vue.js and Mustache. Symfony (a PHP framework) also uses double braces in its Twig templating language, so you might have seen them used similarly there.

To render a variable in a template, simply wrap it with braces – `{ { book_name } }`. Django will automatically escape HTML in output so that you can include special characters (such as `<` or `>`) in your variable without worrying about it garbling your output. If a variable is not passed to a template, Django will simply render nothing at that location, instead of throwing an exception.

There are many more ways to render a variable differently using filters, but these will be covered in *Chapter 3, URL Routers, Views, and Templates*.

Exercise 1.07 – using variables in templates

We'll put a simple variable inside the `base.html` file to demonstrate how Django's variable interpolation works:

1. In PyCharm, open the `base.html` file.
2. Update the `<body>` element so that it contains a place to render the `name` variable:

```
<body>
Hello, {{ name }}!
</body>
```

3. Go back to your web browser and refresh (you should still be at `http://127.0.0.1:8000`). You will see that the page now displays **Hello, !**. This is because we have not set the `name` variable in the rendering context.



Figure 1.39 – No value is rendered in the template because no context was set

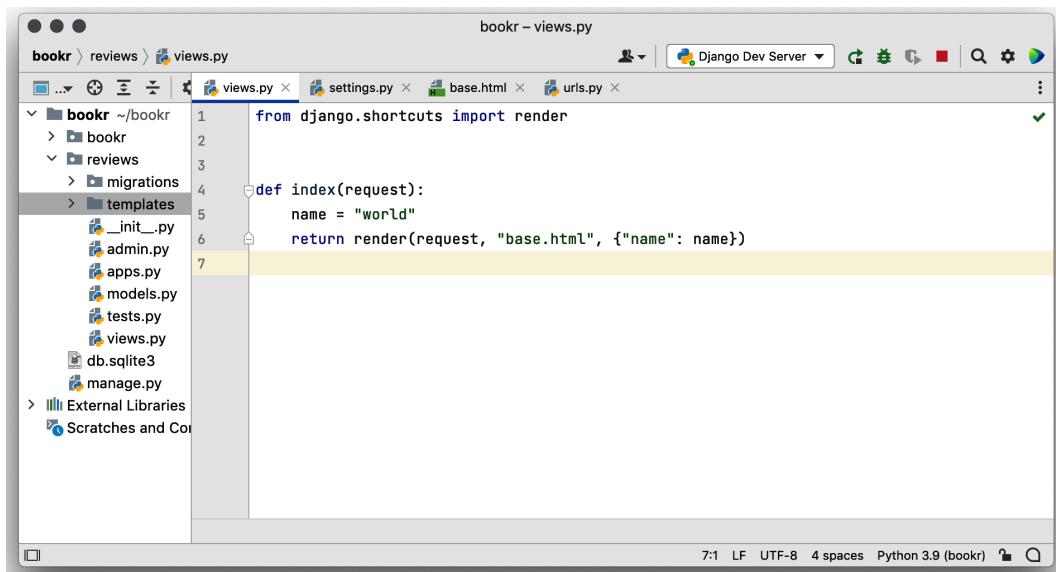
4. Open `views.py` and add a variable called `name`, set to the "world" value, inside the `index` function:

```
def index(request):  
    name = "world"  
    return render(request, "base.html")
```

5. Refresh your browser again. Note that nothing has changed – anything we want to render must be explicitly passed to the `render` function as **context**. This is the dictionary of variables that are made available when rendering.
6. Add the context dictionary as the third argument to the `render` function. Change your `render` line to this:

```
return render(request, "base.html", {"name": name})
```

The entire `views.py` file is shown in the following screenshot:



```
bookr - views.py
bookr > reviews > views.py
bookr > bookr > reviews > migrations > templates > views.py
bookr > bookr > reviews > migrations > __init__.py
bookr > bookr > reviews > migrations > admin.py
bookr > bookr > reviews > migrations > apps.py
bookr > bookr > reviews > migrations > models.py
bookr > bookr > reviews > migrations > tests.py
bookr > bookr > reviews > migrations > views.py
bookr > bookr > db.sqlite3
bookr > bookr > manage.py
bookr > External Libraries
bookr > Scratches and Code

from django.shortcuts import render

def index(request):
    name = "world"
    return render(request, "base.html", {"name": name})
```

Figure 1.40 – `views.py` with the `name` variable set in the `render` context

7. Refresh your browser again, and you'll see it now says **Hello, world!**.

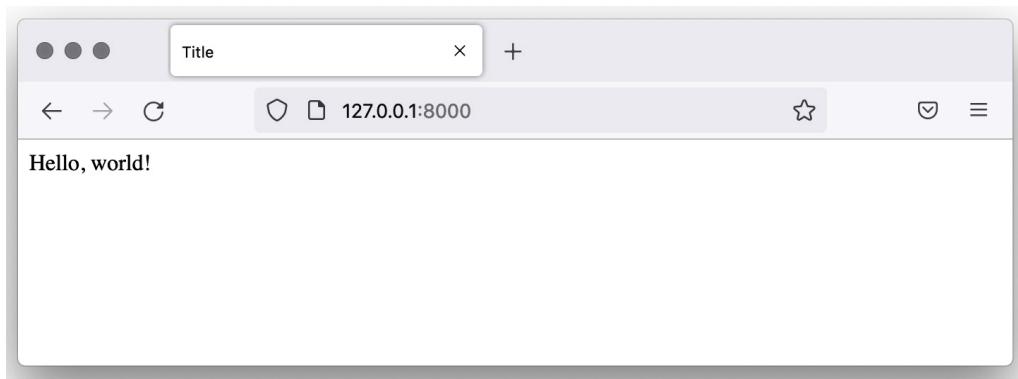


Figure 1.41 – A template rendered with a variable

In this exercise, we combined the template we created in the previous exercise with the `render` function to render an HTML page, with the `name` variable that was passed to it inside a context dictionary.

In the next section, we'll look at some ways you can debug your code and handle errors.

Debugging and dealing with errors

When programming, unless you're the perfect programmer who never makes mistakes, you'll probably have to deal with errors or debug your code at some point. When there is an error in your program, there are usually two ways to tell – either your code will raise an exception, or you will get unexpected output or results when viewing the page. You will probably see exceptions more often, as there are many accidental ways to cause them. If your code is generating unexpected output but not raising any exceptions, you will probably want to use the PyCharm debugger to find out why.

We'll start with an overview of some Python exceptions and how to handle them, along with an exercise that demonstrates how Django shows exceptions. Then, we'll look at running Django inside the PyCharm debugger so that you can peek inside your program while it executes.

Exceptions

If you have worked with Python or other programming languages before, you have probably come across exceptions. If not, here's a quick introduction. Exceptions are *raised* (or *thrown* in other languages) when an error occurs. The execution of a program stops at that point in the code, and the exception travels back up the function call chain until it is caught. If it is not caught, then the program will crash, sometimes with an error message describing the exception and where it occurred. There are exceptions that are raised by Python itself, and your code can raise exceptions to quickly stop execution at any point. Some common exceptions that you might see when programming Python are as follows:

- `IndentationError`: Python will raise this if your code is not correctly indented or has mixed tabs and spaces
- `SyntaxError`: Python raises this error if your code has invalid syntax:

```
>>> a === 1
      File "<stdin>", line 1
      a === 1
      ^
SyntaxError: invalid syntax
```

- `ImportError`: This is raised when an import fails – for example, if trying to import from a file that does not exist or trying to import a name that is not set in a file:

```
>>> import missing_file
      Traceback (most recent call last):
      File "<stdin>", line 1, in <module>
      ImportError: No module named missing_file
```

- **NameError**: This is raised when trying to access a variable that has not yet been set:

```
>>> a = b + 5
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'b' is not defined
```

- **KeyError**: This is raised when accessing a key that is not set in a dict (or a dict-like object):

```
>>> d = {'a': 1}
>>> d['b']
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'b'
```

- **IndexError**: This is raised when accessing an index outside the length of a list:

```
>>> l = ['a', 'b']
>>> l[3]
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
IndexError: list index out of range
```

- **TypeError**: This is raised when trying to perform an operation on an object that does not support it, or when using two objects of the wrong type – for example, trying to add a string to an int:

```
>>> 1 + '1'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: unsupported operand type(s) for +: 'int' and 'str'
```

Django also raises its own custom exceptions, and you will be introduced to them throughout the book.

When running the Django development server with `DEBUG = True` in your `settings.py` file, Django will automatically capture exceptions that occur in your code (instead of crashing). It will then generate an HTTP response, showing you a stack trace and other information to help you debug the problem. When running in production, `DEBUG` should be set to `False`. Django will then return a standard internal server error page, without any sensitive information. You also have the option to display a custom error page.

Exercise 1.08 – generating and viewing exceptions

Let's create a simple exception in our view so that you are familiar with how Django displays them. In this case, we'll try to use a variable that doesn't exist, which will raise `NameError`:

1. In PyCharm, open `views.py`. In the `index` view function, change the context being sent to the `render` function so that it's using a variable that doesn't exist. We'll try to send `invalid_name` in the context dictionary, instead of `name`. Don't change the context dictionary key, just its value:

```
return render(request, "base.html", {"name": invalid_name})
```

2. Go back to your browser and refresh the page. You should see a screen like *Figure 1.42*:

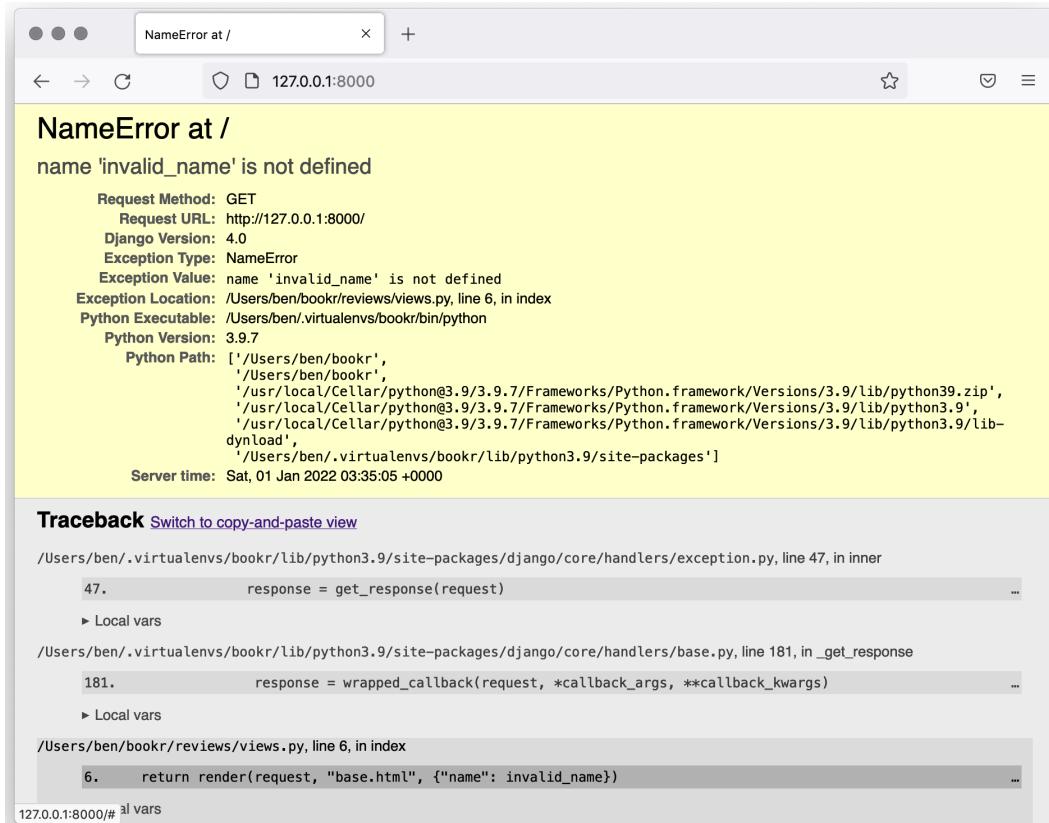
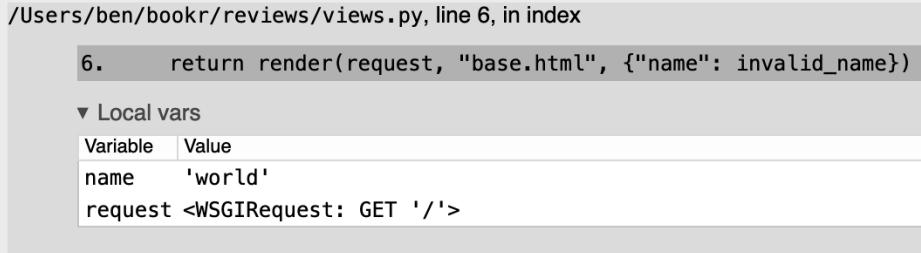


Figure 1.42 – A Django exception screen

The first couple of header lines on the page inform you of the error that occurred:

```
NameError at /  
name 'invalid_name' is not defined
```

3. Shown next the header is a traceback to where the exception occurred. Click on the various lines of code to expand them and see the surrounding code, or click **Local vars** for each frame to expand them and see what the values of the variables are:



The screenshot shows a debugger interface. At the top, it says "/Users/ben/bookr/reviews/views.py, line 6, in index". Below that is a code line: "6. return render(request, "base.html", {"name": invalid_name})". Underneath the code is a "Local vars" section with a "▼" icon. A table shows variables and their values: "name" is 'world', and "request" is a WSGIRequest object with a GET method. The "request" entry is highlighted with a red border.

Figure 1.43 – The line that caused the exception

In our case, we can see the exception was raised on line six of our `views.py` file, and by expanding its **Local vars** dropdown, we see that `name` has the value `world`, and the only other variable is the incoming `request` (*Figure 1.43*).

4. Go back to `views.py` and fix your `NameError` by renaming `invalid_name` back to `name`.
5. Save the file and refresh your browser. **Hello, world!** should display again (as shown in *Figure 1.41*).

In this exercise, we made our Django code raise an exception (`NameError`) by trying to use a variable that had not been set. We saw that Django automatically sent details of this exception and a stack trace to the browser to help us find the cause of the error. We then reversed our code change to make sure our view worked properly again.

Debugging

When you're trying to find problems in your code, it can help to use a debugger. This is a tool that lets you go through your code line by line, rather than executing it all at once. Each time the debugger is paused on a particular line of code, you can see the values of all the current variables. This is very useful for finding out errors in your code that don't raise exceptions.

For example, in Bookr, we talked about having a view that fetches a list of books from the database and renders them in an HTML template. If you view the page in the browser, you might see only one book when you expect several. You could have the execution pause inside your view function and see what values were fetched from the database. If your view is only receiving one book from the database, you know there is a problem with your database querying somewhere. If your view is

successfully fetching multiple books but only one is being rendered, then it's probably a problem with the template. Debugging helps you narrow down faults like this.

PyCharm has a built-in debugger to make it easy to step through your code and see what is happening on each line. To tell the debugger where to stop the execution of code, you need to set a *breakpoint* on one or more lines of code. They are named as such because the execution of code will *break* (stop) at that *point*.

For breakpoints to be activated, PyCharm needs to be set to run your project in its debugger. There is a small performance penalty, but it is usually not noticeable, so you might choose to always run your code inside the debugger so that you can quickly set a breakpoint without having to stop and restart the Django development server.

Running the Django development server inside the debugger is as simple as clicking the **Debug** icon instead of the **Play** icon (see *Figure 1.20*) to start it.

Exercise 1.09 – debugging your code

In this exercise, you will learn the basics of the PyCharm debugger. You will run the Django development server in the debugger and then set a breakpoint in your view function to pause execution so that you can examine the variables:

1. If the Django development server is running, stop it by clicking the **Stop** button in the top-right corner of the PyCharm window:

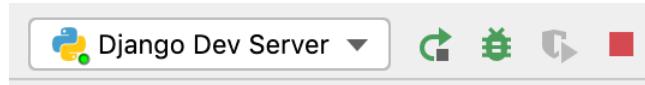
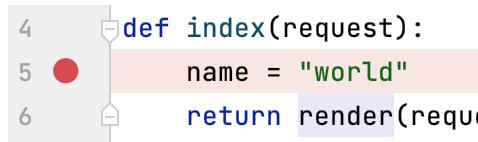


Figure 1.44 – The Stop button in the top-right corner of the PyCharm window

2. Start the Django development server again inside the debugger by clicking the **Debug** icon just to the left of the **Stop** button (*Figure 1.44*).

The server will take a few seconds to start, and then you should be able to refresh the page in your browser to make sure it's still loading. You shouldn't notice any changes; all the code is executed the same as before.

3. Now, we can set a breakpoint that will cause execution to stop so that we can see the state of the program. In PyCharm, click just to the right of the line numbers, on line 5, in the gutter on the left of the editor pane. A red circle will appear to indicate that the breakpoint is now active:



```
4     def index(request):  
5         name = "world"  
6         return render(request,
```

Figure 1.45 – A breakpoint on line 5

4. Go back to your browser and refresh the page. Your browser will not display any content; instead, it will just continue to try to load the page. Depending on your operating system, PyCharm should become active again; if not, bring it to the foreground. You should see that line 5 is highlighted, and at the bottom of the window, the debugger is shown. The stack frames (the chain of functions that were called to get to the current line) are on the left, and the current variables of the function are on the right.

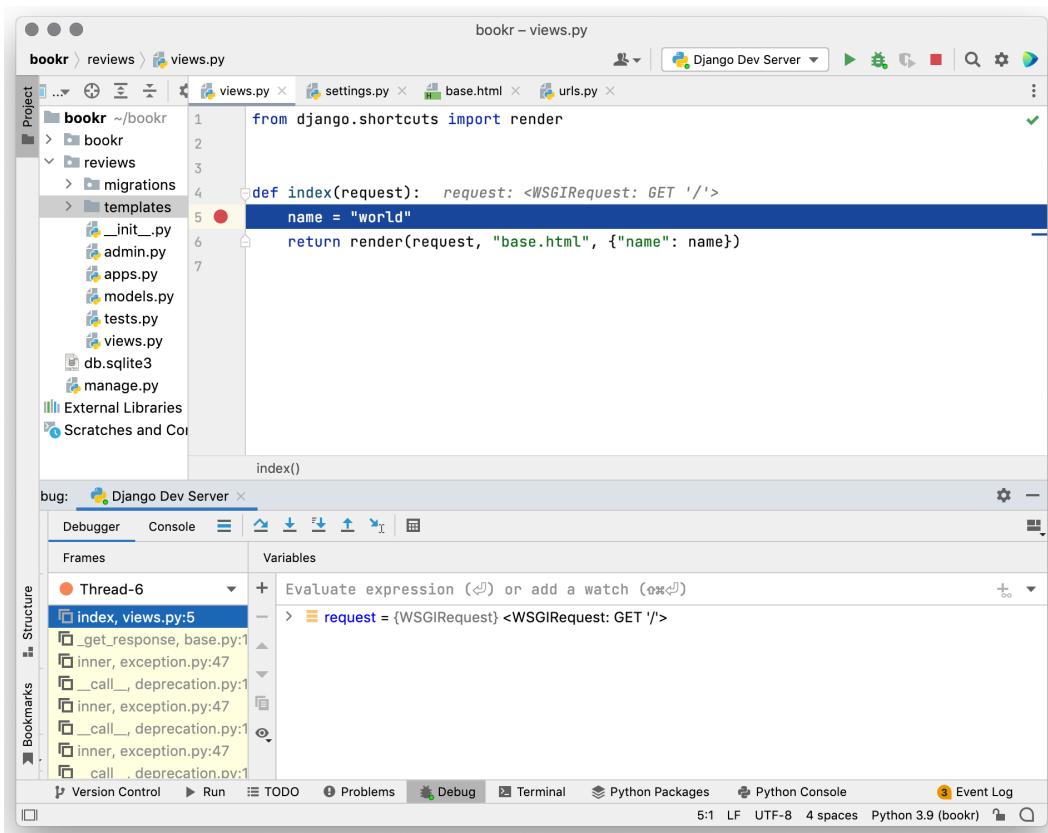


Figure 1.46 – The debugger paused with the current line (5) highlighted

There is currently one variable in scope, `request`. If you click the toggle triangle to the left of its name, you can show or hide the attributes it has set.

```

▼ └─ request = {WSGIRequest} <WSGIRequest: GET '/>
  > └─ COOKIES = {dict: 1} {'csrftoken': 'P86W4F9J4sMqGHeFma1h...'}
  > └─ FILES = {MultiValueDict: 0} <MultiValueDict: {}>
  > └─ GET = {QueryDict: 0} <QueryDict: {}>
  > └─ META = {dict: 79} {'PATH': '/Users/ben/.virtualenvs/bookr/bin:...'}
  > └─ POST = {QueryDict: 0} <QueryDict: {}>
  > └─ accepted_types = {list: 6} [<MediaType: text/html>, <MediaType: ...>]
  > └─ body = {bytes: 0} b""
  > └─ content_params = {dict: 0} {}
    ① content_type = {str} 'text/plain'
    ① csrf_processing_done = {bool} True
    ① encoding = {NoneType} None
  > └─ environ = {dict: 79} {'PATH': '/Users/ben/.virtualenvs/bookr/bin:...'}
  > └─ headers = {HttpHeaders: 15} {'Content-Length': '', 'Content-T...'}
    ① method = {str} 'GET'
    ① path = {str} '/'
    ① path_info = {str} '/'
  > └─ resolver_match = {ResolverMatch} ResolverMatch(func=revie...)
    ① scheme = {str} 'http'
  > └─ session = {SessionStore} <django.contrib.sessions.backends...
  > └─ upload_handlers = {list: 2} [<django.core.files.uploadhandler...
  > └─ user = {AnonymousUser} AnonymousUser
  > └─ Protected Attributes

```

Figure 1.47 – The attributes of the `request` variable

For example, if you scroll down through the list of attributes, you can see that `method` is `GET` and `path` is `/`.

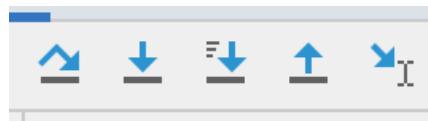


Figure 1.48 – The actions bar

5. The actions bar, shown in *Figure 1.48*, is above the stack frames and variables. Its buttons (from left to right) are as follows:
 - **Step Over:** This executes the current line of code and continues to the next line.
 - **Step Into:** This steps into the current line. For example, if the line contained a function, it would continue with the debugger inside this function.

- **Step into My Code:** This steps into the line being executed but continues until it finds the code you have written. For example, if you're stepping into third-party library code that later calls your code, it will not show you the third-party code, instead continuing through until it returns to the code that you have written.
- **Step Out:** This returns from the current code to the function or the method that called it. It is the opposite of the **Step In** action.
- **Run to Cursor:** If you have a line of code further along from where you currently are that you want to execute to, without having to click **Step Over** for all the lines in between, click to put your cursor on that line. Then, click **Run to Cursor**, and the execution will continue until that line.

Note that not all buttons are useful all the time. For example, it can be easy to step out of your view and end up confusing Django library code.

6. Click the **Step Over** button once to execute line 5.
7. You can see that the `name` variable has been added to the list of variables in the debugger view, and its value is `world`.

```
01 name = {str} 'world'  
>  02 request = {WSGIRequest} <WSGIRequest: GET '/'>
```

Figure 1.49 – The new variable `name` is now in scope, with the `world` value

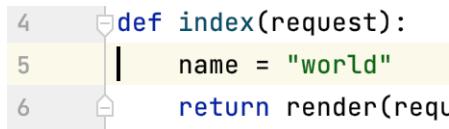
8. We are now at the end of our `index` view function, and if we were to step over this line of code, it would jump to Django library code, which we don't want to see. To continue executing and to send the response back to your browser, click the **Resume Program** button on the left of the window (Figure 1.50). You should see that your browser has now loaded the page again:



Figure 1.50 – Actions to control execution – the green play icon is the **Resume Program** button

There are more buttons in *Figure 1.50*; from the top, they are **Rerun** (stops the program and restarts it), **Modify Run Configuration...** (opens the run configuration window), **Resume Program** (continues running until the next breakpoint), **Pause Program** (breaks the program at its current execution point), **Stop** (stops the debugger), **View Breakpoints** (opens a window to see all the breakpoints you have set), and **Mute Breakpoints** (which will toggle all breakpoints on or off but not remove them).

9. For now, turn off the breakpoint in PyCharm by clicking on it (the red circle next to line 5).



A screenshot of the PyCharm code editor. The code is as follows:

```
4     def index(request):  
5         name = "world"  
6         return render(req
```

Line 5 is highlighted in yellow. On the far left of line 5, there is a small red circle with a white outline, representing a breakpoint. The line numbers 4, 5, and 6 are visible on the left.

Figure 1.51 – Clicking the breakpoint that was on line 5 disables it

This is just a quick introduction to how to set breakpoints in PyCharm. If you have used debugging features in other IDEs, then you should be familiar with the concepts – you can step through code, step in and out of functions, or evaluate expressions. Once you have set a breakpoint, you can right-click on it to change the options. For example, you can make the breakpoint conditional so that the execution only stops under certain circumstances. All these are beyond the scope of this book but are useful to know about when trying to solve problems in your code.

We're going to finish the chapter with two activities. First, you'll build a welcome screen for Bookr, and then you'll scaffold the book-searching page.

Activity 1.01 – creating a site welcome screen

The Bookr website that we are building needs to have a splash page that welcomes users and lets them know what site they are on. It will also contain links to other parts of the site, but these will be added in later chapters. For now, you will create a page with a welcome message.

These steps will help you complete the activity:

1. In your `index` view, render the `base.html` template.
2. Update the `base.html` template to contain the welcome message. It should be in both the `<title>` tag in `<head>` and in a new `<h1>` tag in the body.

After completing the activity, you should be able to see something like this:

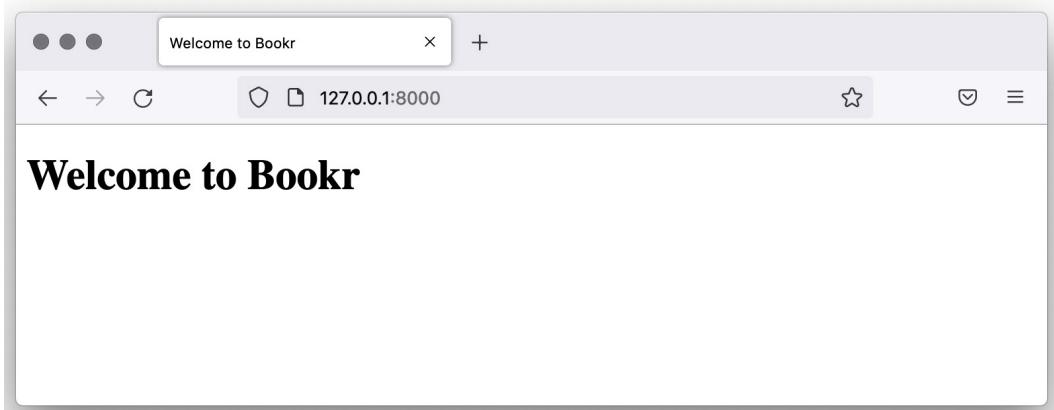


Figure 1.52 – The Bookr welcome page

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

In the next activity, you will build a basic page to show the results of a book search in Bookr.

Activity 1.02 – a book search scaffold

A useful feature of a site such as Bookr is the ability to search through data to find something on the site quickly. Bookr will implement book searching to allow users to find a particular book by part of its title. While we don't have any books to find yet, we can still implement a page that shows the text a user searched for. The user enters the search string as part of the URL parameters. We will implement searching and a form for easy text entry in *Chapter 6, Forms*.

These steps will help you complete the activity:

1. Create a search result HTML template. It should include a variable placeholder to show the search word(s) that were passed in through the render context. Show the passed-in variable in the `<title>` and `<h1>` tags. Use an `` tag around the search text in the body to make it italic.
2. Add a search view function in `views.py`. The view should read a search string from the URL parameters (in the request's `GET` attribute). It should then render the template you created in the previous step, passing in the search value to be substituted, using the context dictionary.

3. Add a URL mapping to your new view to `urls.py`. The URL can be something such as `/book-search`.

After completing this activity, you should be able to pass in a search value through the URL's parameters and see it rendered on the resulting page. It should look like this:



Figure 1.53 – Searching Django workshop

You should also be able to pass in special HTML characters such as `<` and `>` to see how Django automatically escapes them in the template.



Figure 1.54 – Note how HTML characters are escaped so that we are protected from tag injection

Note

The solution for this activity can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

This chapter was a quick introduction to Django. You first got up to speed on the HTTP protocol and the structure of HTTP requests and responses. We then saw how Django uses the MVT paradigm, and then how it parses a URL, generates an HTTP request, and sends it to a view to get an HTTP response. We scaffolded the Bookr project and then created the `reviews` app for it. We then built two example views to illustrate how to get data from a request and use it when rendering templates. You also experimented to see how Django escapes output in HTML when rendering a template.

We did all this with the PyCharm IDE, and you learned how to set it up to debug your application. The debugger will help you find out why things aren't working as they should.

In the next chapter, you will start to learn about Django's database integration and its model system, so you can start storing and retrieving real data for your application.

2

Models and Migrations

In the previous chapter, we had a brief introduction to Django and its use in developing web applications. Then, we learned about the **Model-View-Template (MVT)** concept. Later, we created a Django project and started the Django development server. We also briefly discussed Django's views, URLs, and templates.

In this chapter, you will be introduced to the concept of databases, the types of databases, and their importance in building web applications. You will start by creating a database using an open source database visualization tool called SQLite DB Browser. You will then perform some basic **create, read, update, and delete (CRUD)** database operations using SQL commands. Then, you will learn about Django's **object-relational mapping (ORM)**, through which your application can interact and seamlessly work with a relational database using simple Python code, eliminating the need to run complex SQL queries. You will learn about models and migrations, which are part of Django's ORM, used to propagate database schematic changes from the application to the database, and also perform database CRUD operations. Toward the end of the chapter, you will study the various types of database relationships and use that knowledge to perform queries across related records.

We will cover the following topics in this chapter:

- Understanding and using databases
- Understanding CRUD operations using SQL
- Exploring Django ORM
- Creating Django models and migrations
- Django's database CRUD operations
- Bulk create and bulk update operations
- Performing complex lookups using `Q` objects
- Activity 2.01 – create models for a project management application
- Populating the Bookr project's database

By the end of this chapter, you will have a good understanding of databases and how Django can be used to implement database operations for a web application.

Technical requirements

Throughout this book, you will be writing code. If you need to refer to the complete code for this chapter, you can find it in the GitHub repository here: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter02>.

Understanding and using databases

In most web applications, data forms the core part of the application. Unless we're talking about a very simple application such as a calculator, in most cases, we need to store data, process it, and display it to the user on a page. Since most operations in user-facing web applications involve data, there is a need to store data somewhere secure, easily accessible, and readily available. This is where databases come in handy. Imagine a library operation before the advent of computers. The librarian would have to maintain records of book inventories, book lending, returns from students, and so on. All of these would have been maintained in physical records. While carrying out their day-to-day activities, the librarian would modify these records for each operation, for example, when lending a book to someone or when the book was returned.

Today, we have databases to help us with such administrative tasks. A database looks like a spreadsheet or an Excel sheet containing records, with each table consisting of multiple rows and columns. An application can have many such tables. Here is an example table of a book inventory in a library:

Book Number	Author	Title	Number of Copies
Howto4563	Adam Chappel	How to Build a House	4
Travel5327	Charlie Hunt	How to Holiday in Switzerland	5
Fiction3453	Evan Stark	The Mystery Cat	2
Howto4453	Bruce Williams	Sailing Guide	7

Table 2.1: Table of a book inventory for a library

In the preceding table, we can see that there are columns with details about various attributes of the books in the library, while the rows contain entries for each book. To manage a library, there can be many such tables working together as a system. For example, along with an inventory, we may have other tables such as student information, book lending records, and so on. Databases are built with the same logic, where software applications can easily manage data.

A database is a structured collection of data that helps manage information easily. A software layer called the **database management system (DBMS)** is used to store, maintain, and perform operations on the data. There are two types of databases, relational databases, and non-relational databases. In the following sections, we will learn in brief about relational databases and how data is stored in such databases.

Relational databases

Relational databases or **structured query language (SQL)** databases store data in a pre-determined structure of rows and columns called **tables**. A database can be made up of more than one such table, and these tables have a fixed structure of attributes, data types, and relations with other tables. For example, as we just saw in *Figure 2.1*, the book inventory table has a fixed structure of columns comprising **Book Number**, **Author**, **Title**, and **Number of Copies**, and the entries form the rows in the table. There could be other tables as well, such as **Student Information** and **Lending Records**, which could be related to the inventory table. Also, whenever a book is lent to a student, the records will be stored per the relationships between multiple tables (say, the **Student Information** and the **Book Inventory** tables).

This pre-determined structure of rules defining the data types, tabular structures, and relationships across different tables acts like scaffolding or a blueprint for a database. This blueprint is collectively called a **database schema**. It will prepare the database to store application data when applied to a database. To manage and maintain these databases, there is a common language for relational databases called SQL. Some examples of relational databases are SQLite, PostgreSQL, MySQL, and OracleDB.

In the next section, you will be introduced to non-relational databases, some examples of non-relational databases, and the reasons for using them in web applications.

Non-relational databases

Non-relational databases or **not only SQL (NoSQL)** databases are designed to store unstructured data. They are well suited to large amounts of generated data that does not follow rigid rules, as is the case with relational databases. Some examples of non-relational databases are Cassandra, MongoDB, CouchDB, and Redis.

For example, imagine that you need to store the stock value of companies in a database using Redis. Here, the company name will be stored as the key and the stock value as the value. Using the key-value type NoSQL database in this use case is appropriate because it stores the desired value for a unique key and is faster to access.

For the scope of this book, we will be dealing only with relational databases, as Django does not officially support non-relational databases. However, if you wish to explore, there are many forked projects, such as **Django non-rel**, that support NoSQL databases.

In the next section, you will be introduced to database CRUD operations using SQL commands.

Database operations using SQL

SQL uses a set of commands to perform a variety of database operations, such as creating an entry, reading values, updating an entry, and deleting an entry. These operations are collectively called CRUD operations. To understand database operations in detail, let's first get some hands-on experience with SQL commands. Most relational databases share a similar SQL syntax; however, some operations will differ.

For the scope of this chapter, we will use SQLite as the database. SQLite is a lightweight relational database and is a part of Python's standard libraries. That's why Django uses SQLite as its default database configuration. However, we will also learn more about how to perform configuration changes to use other databases in *Chapter 17, Deploying a Django Application (Part 1 – Server Setup)*. This chapter can be downloaded from the GitHub repository of this book, here: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter17>.

In the next section, we will learn about the importance of data types in a database.

Data types in relational databases

Databases provide us with a way to restrict the type of data that can be stored in a given column. These are called data types. Some examples of data types for a relational database, such as SQLite3, are given here:

- **INTEGER:** This is used for storing integers
- **TEXT:** This can store text
- **REAL:** This is used for floating-point values

For example, you would want the title of a book to have TEXT as the data type. So, the database will enforce a rule that no type of data, other than text data, can be stored in that column. Similarly, the book's price can have a REAL data type, and so on.

Using some of the concepts learned before about databases, in the next section, we will get hands-on by creating a database and a database table.

Exercise 2.01 – creating a book database

In this exercise, you will create a book database for a book review application. For better visualization of the data in the SQLite database, you will install an open source tool called **DB Browser for SQLite**. This tool helps visualize the data and provides a shell to execute the SQL commands.

If you haven't done so already, visit <https://sqlitebrowser.org>, and from the Download section, install the application as per your operating system and launch it. Detailed instructions for DB Browser installation can be found in the *Preface*. To create a database, follow these steps:

Note

Database operations can be performed using a command-line shell as well.

1. After launching the application, create a new database by clicking **New Database** in the top-left corner of the application. Create a database named `bookr`, as you are working on a book review application:

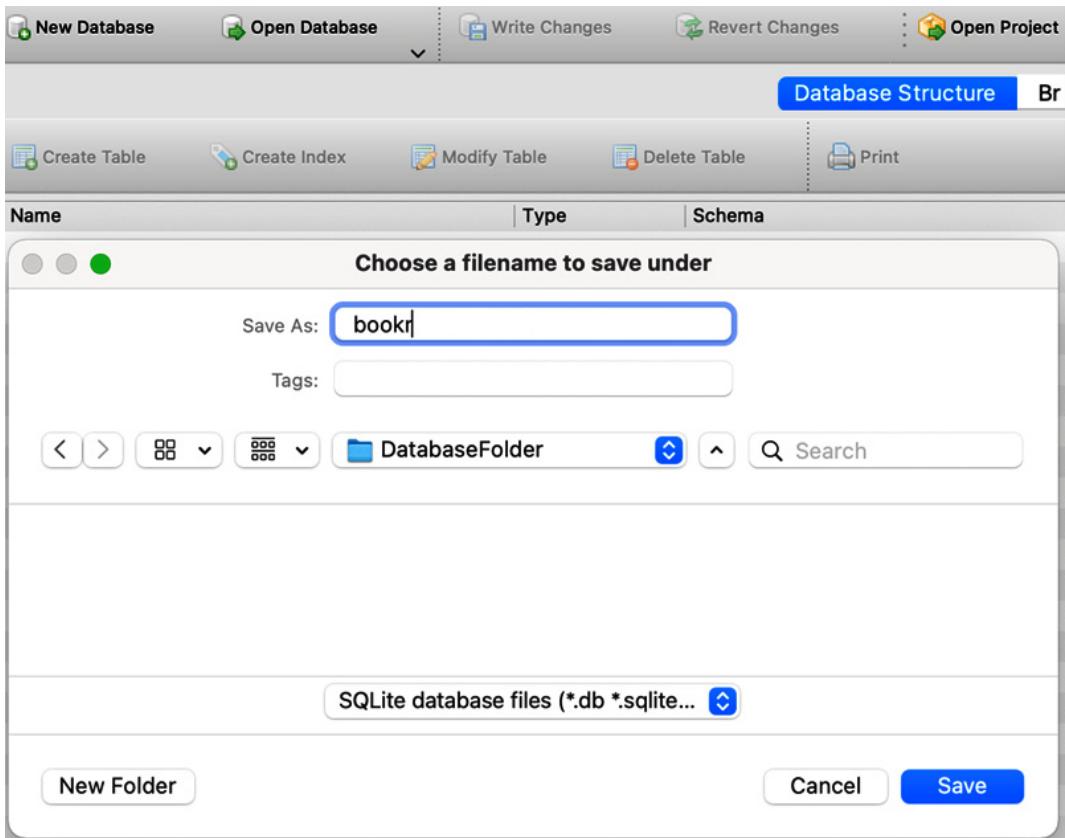


Figure 2.1 – Creating a database named `bookr`

2. Next, click on the **Create Table** button in the top-left corner and enter book as the table name.

Note

After clicking the **Save** button, you may find that the window for creating a table opens up automatically. In that case, you won't have to click the **Create Table** button; simply proceed with the creation of the book table as specified in the preceding step.

3. Now, click the **Add field** button, enter the field name as `title` and select the **TEXT** type from the drop-down menu. Here, **TEXT** is the data type for the `title` field in the database:

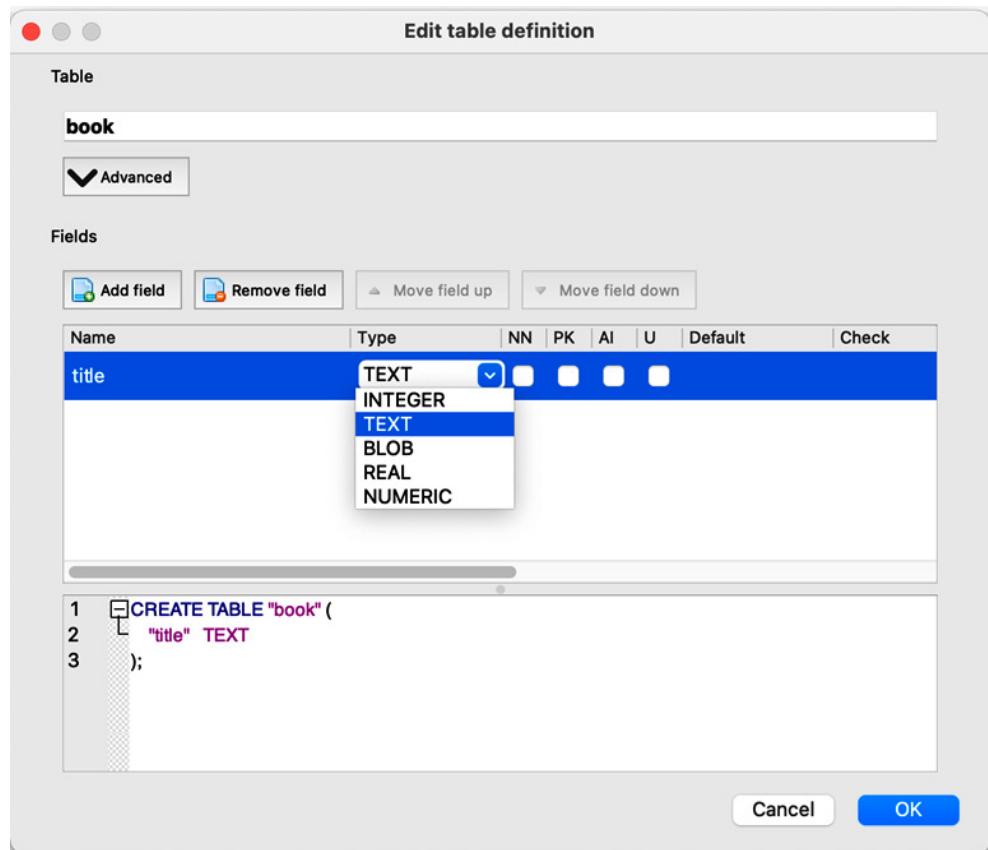


Figure 2.2: Adding a TEXT field named title

4. Similarly, add two more fields for the table named **publisher** and **author** and select the **TEXT** type for both fields. Then, click on the **OK** button:

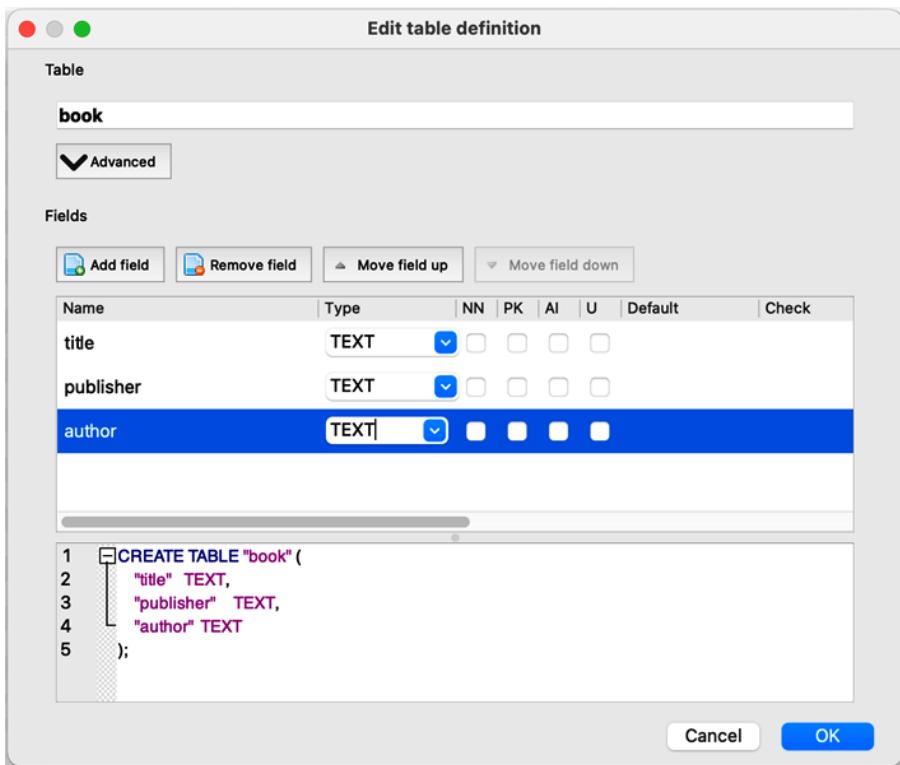


Figure 2.3: Creating TEXT fields named publisher and author

This creates a database table called book in the bookr database with the title, publisher, and author fields. This can be seen as follows:

Name	Type	Schema
Tables (1)		
book		CREATE TABLE "book" ("title" TEXT,"publisher" TEXT,"author" TEXT)
title	TEXT	"title" TEXT
publisher	TEXT	"publisher" TEXT
author	TEXT	"author" TEXT
Indices (0)		
Views (0)		
Triggers (0)		

Figure 2.4: Database with the title, publisher, and author fields

Overall, in this section about databases, we have learned about why we need databases and briefly about how they store data, the two most common types of databases. Additionally, we created a simple database and a table using an open source tool called DB Browser (SQLite). In the next section, we will learn more about how to do CRUD operations on the database using SQL.

Understanding CRUD operations using SQL

Let's assume that the editors or the users of our book review application want to make some modifications to the book inventory, such as adding a few books to the database, updating an entry in the database, and so on. SQL provides various ways to perform such CRUD operations. Before we dive into the world of Django models and migrations, let's explore these basic SQL operations first.

You will be running a few SQL commands for the CRUD operations that follow.

To run them, follow these steps:

1. Navigate to the **Execute SQL** tab in DB Browser. You can type in or paste the SQL commands we've listed in the sections that follow in the **SQL 1** window. You can spend some time modifying your queries, and understanding them before you execute them.
2. When you're ready, click the icon that looks like a **Play** button or press the **F5** key to execute the command. The results will show up in the window below the **SQL 1** window:

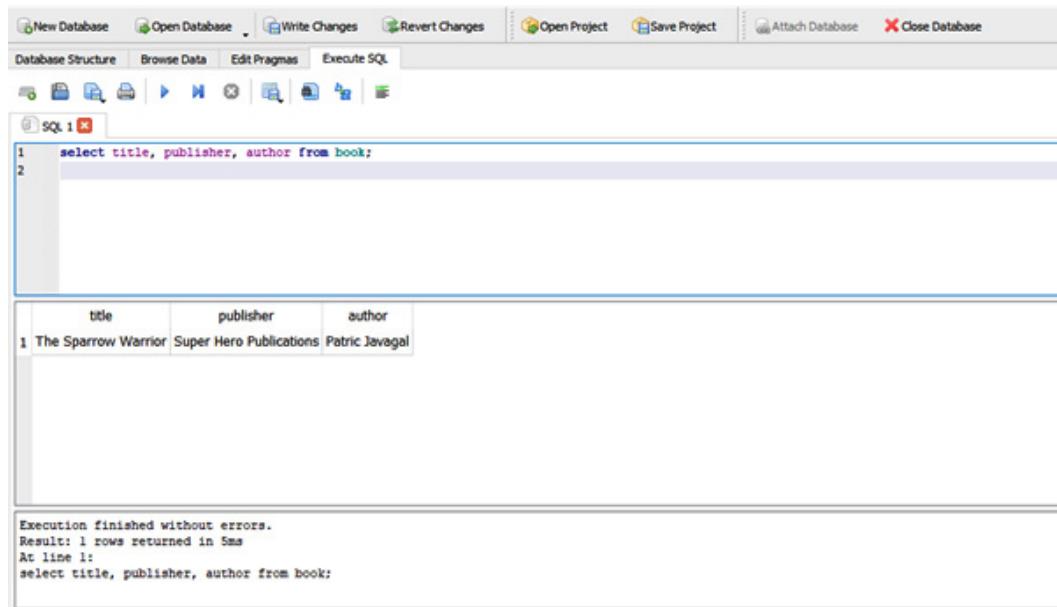


Figure 2.5: Executing SQL commands in DB Browser

Now we will try out some SQL create operations to create a few records in the database.

SQL create operations

A create operation in SQL is performed using the `insert` command, which, as the name implies, lets us insert data into the database. Let's go back to our `bookr` example. Since we have already created the database and the `book` table, we can now create or insert an entry in the database by executing the following command:

```
insert into book values ('The Sparrow Warrior', 'Super Hero Publications', 'Patric Javagal');
```

This inserts the values defined in the command into the `book` table. Here, `The Sparrow Warrior` is the title, `Super Hero Publications` is the publisher, and `Patric Javagal` is the author of the book. Note that the order of insertion corresponds with the way we have created our table; that is, the values are inserted into the columns representing title, publisher, and author, respectively. Similarly, let's execute two more inserts to populate the `book` table:

```
insert into book values ('Ninja Warrior', 'East Hill Publications', 'Edward Smith');
insert into book values ('The European History', 'Northside Publications', 'Eric Robbins');
```

The three inserts executed so far will insert three rows into the `book` table. But how do we verify that? How do we know whether those three entries we inserted are entered into the database correctly? Let's learn how to do that in the next section.

SQL read operations

We can read from the database using the `select` SQL operation. For example, the following SQL `select` command retrieves the selected entries created in the `book` table:

```
select title, publisher, author from book;
```

You should see the following output:

	title	publisher	author
1	The Sparrow Warrior	Super Hero Publications	Patric Javagal
2	Ninja Warrior	East Hill Publications	Edward Smith
3	The European History	Northside Publications	Eric Robbins

Figure 2.6: Output after using the `select` command

Here, `select` is the command that reads from the database, and the `title`, `publisher`, and `author` fields are the columns that we intend to select from the `book` table. Since these are all the columns the database has, the `select` statement has returned all the values present in the database. The `select` statement is also called an SQL query. An alternate way to get all the fields in the database is by using the `*` wildcard in the `select` query instead of specifying all the column names explicitly:

```
select * from book;
```

This will return the same output as shown in the preceding figure. Now, suppose we want to get the author's name for the book titled `The Sparrow Warrior`; in this case, the `select` query would be as follows:

```
select author from book where title="The Sparrow Warrior";
```

Here, we have added a special SQL keyword called `where` so that the `select` query returns only the entries that match the condition. The result of the query, of course, will be `Patric Javagal`. Now, what if we wanted to change the name of the book's publisher?

SQL update operations

In SQL, the way to update a record in the database is by using the `update` command:

```
update book set publisher = 'Northside Publications' where
title='The Sparrow Warrior';
```

Here, we set the `publisher` value to `Northside Publications` if the value of the `title` is `The Sparrow Warrior`. We can then run the `select` query we ran in the *SQL read operations* section to see how the updated table looks after running the `update` command:

	title	publisher	author
	Filter	Filter	Filter
1	The Sparrow Warrior	Northside Publications	Patric Javagal
2	Ninja Warrior	East Hill Publications	Edward Smith
3	The European History	Northside Publications	Eric Robbins

Figure 2.7: Updating the publisher value for The Sparrow Warrior

Next, what if we wanted to delete the title of the record we just updated? We will see just that in the next section.

SQL delete Operations

Here is an example of how to delete a record from the database using the `delete` command:

```
delete from book where title='The Sparrow Warrior';
```

The `delete` command is the SQL keyword for delete operations. Here, this operation will be performed only if the title is `The Sparrow Warrior`. Here is how the `book` table will look after the delete operation:

	title	publisher	author
1	Ninja Warrior	East Hill Publications	Edward Smith
2	The European History	Northside Publications	Eric Robbins

Figure 2.8 – Output after performing the delete operation

These are the basic operations of SQL. We will not go further into all the SQL commands and syntax, but feel free to explore more about database base operations using SQL.

Note

For further reading, you can start by exploring some advanced SQL `select` operations with `join` statements, which are used to query data across multiple tables. For a detailed course on SQL, you can refer to *The SQL Workshop* (<https://www.packtpub.com/product/the-sql-workshop/9781838642358>).

In this section, we learned how to interact with the database by performing CRUD operations on the database using SQL. In the next section, we will learn about how Django's ORM, which is an abstract layer, interacts with a database.

Exploring Django ORM

Web applications constantly interact with databases, and one of the ways to do so is by using SQL. If you decide to write a web application without a web framework such as Django and instead use Python alone, Python libraries such as `psycopg2` (for PostgreSQL) could be used to interact directly with the databases using SQL commands. But while developing a web application with multiple tables and fields, SQL commands can easily become overly complex and thus difficult to maintain. For this reason, popular web frameworks such as Django provide a level of abstraction that allows us to easily work with databases. The part of Django that helps us do this is called ORM.

Django ORM converts object-oriented Python code into actual database constructs, such as database tables with data type definitions, and facilitates all the database operations via simple Python code.

Because of this, we do not have to deal with SQL commands while performing database operations. This helps in faster application development and ease in maintaining the application source code.

Django supports relational databases such as SQLite, PostgreSQL, Oracle Database, and MySQL. Django's database abstraction layer ensures that the same Python/Django source code can be used across any of these relational databases with very little modification to the project settings. Since SQLite is part of the Python libraries and Django is configured by default to SQLite, for the scope of this chapter, we will use SQLite while we learn about Django models and migrations. Next, we will understand how database configuration is done with Django.

Database configuration and creating Django applications

As we have already seen in *Chapter 1, An Introduction to Django*, when we create a Django project and run the Django server, the default database configuration is SQLite3. The database configuration will be present in the project directory, in the `settings.py` file.

Note

Make sure you go through the `settings.py` file for the `bookr` app. Going through the entire file once will help you understand the concepts that follow. You can find the file here: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter02/final/bookr/bookr/settings.py>.

So, for our example project, the database configuration will be present at the following location: `bookr/settings.py`. The default database configuration present in this file when a Django project is created is as follows:

```
DATABASES = {
    'default': {
        'ENGINE': 'django.db.backends.sqlite3',
        'NAME': BASE_DIR / 'db.sqlite3',
    }
}
```

The `DATABASES` variable is assigned a dictionary containing the database details for the project. Inside the dictionary, there is a nested dictionary with a `default` key. This holds the configuration of a default database for the Django project. The reason we have a nested dictionary with `default` as a key is that a Django project could potentially interact with multiple databases, and the default database is the one used by Django for all operations unless explicitly specified. The `ENGINE` key represents which database engine is being used; in this case, `sqlite3`.

The NAME key defines the name of the database, which can have any value. But for SQLite3, since the database is created as a file, NAME can have the full path of the directory where the file needs to be created. The full path of the db file is processed by joining (or concatenating) the previously defined path in BASE_DIR with db.sqlite3. Note that BASE_DIR is the project directory as already defined in the `settings.py` file.

If you are using other databases, such as PostgreSQL, MySQL, and so on, changes will have to be made in the preceding database settings as shown here:

```
DATABASES = {
    'default': {
        'ENGINE': 'django.db.backends.postgresql',
        'NAME': 'bookr',
        'USER': <username>,
        'PASSWORD': <password>,
        'HOST': <host-IP-address>,
        'PORT': '5432', }}}
```

Here, changes have been made to ENGINE to use PostgreSQL. The host IP address and port number of the server need to be provided for HOST and PORT, respectively. As the names suggest, USER is the database username, and PASSWORD is the database password. In addition to changes in the configuration, we will have to install the database drivers or bindings along with the database host and credentials. This will be covered in detail in later chapters, but for now, since we are using SQLite3, the default configuration will be sufficient. Note that the preceding code is just an example to show the changes you'll need to make to use a different database such as PostgreSQL, but since we are using SQLite, we will use the database configuration that exists already, and there is no need to make any modifications to the database settings. Next, we will briefly read about Django apps and some of the default apps present in Django.

Django apps

A Django project can have multiple apps that often act as discrete entities. That's why, whenever required, an app can also be plugged into a different Django project. For example, if we develop an e-commerce web application, the web application can have multiple apps, such as a chatbot for customer support or a payment gateway to accept payments as users purchase goods from the application. These apps, if needed, can also be plugged into or reused in a different project.

Django comes with the following apps enabled by default. The following is a snippet from a project's `settings.py` file:

```
INSTALLED_APPS = [
    'django.contrib.admin',
```

```
'django.contrib.auth',
'django.contrib.contenttypes',
'django.contrib.sessions',
'django.contrib.messages',
'django.contrib.staticfiles',
]
```

These are a set of installed or default apps used for the admin site, authentication, content types, sessions, messaging, and an application to collect and manage static files. In the upcoming chapters, we shall study this in depth. For the scope of this chapter, though, we shall understand why Django migration is needed for these installed apps in the following section.

Django migration

As we learned before, Django's ORM helps make database operations simpler. A major part of the operation is transforming the Python code into database structures, such as database fields with stated data types and tables. In other words, the transformation of Python code into database structures is known as **migration**. For example, instead of creating database tables by running SQL commands, you would write models for them in Python; something you'll learn to do in the upcoming *Creating models and migrations* section. These models will have **fields** that form the blueprints of database tables. The fields, in turn, will have different field types giving us more information about the type of data stored there (recall how we specified the data type of our field as `TEXT` in *step 4 of Exercise 2.01 – creating a book database*).

Since we have a Django project set up, let's perform our first migration. Although we have not added any code yet to our project, we can migrate the applications listed in `INSTALLED_APPS`. This is necessary because Django's installed apps need to store the relevant data in the database for their operations, and migration will create the required database tables to store the data in the database. The following command should be entered in the terminal or shell to do this:

```
python manage.py migrate
```

Note

For macOS, you can use `python3` instead of `python` in the preceding command.

Here, `manage.py` is a script that was automatically created when the project was created. It is used for carrying out managerial or administrative tasks. By executing this command, we create all the database structures required by the installed apps.

As we are using DB Browser for SQLite to browse the database, let's take a look at the database for which changes have been made after executing the `migrate` command.

The database file will be created in the project directory under the `db.sqlite3` name. Follow these steps to verify that migration has been successful and database tables have been created:

1. Open DB Browser, click **Open Database**, navigate until you find the `db.sqlite3` file, and open it. You should see a set of newly created tables created by the Django migration. It will look as follows in DB Browser:

Name	Type	Schema
Tables (11)		
auth_group		CREATE TABLE "auth_group" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
auth_group_permissions		CREATE TABLE "auth_group_permissions" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
auth_permission		CREATE TABLE "auth_permission" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
auth_user		CREATE TABLE "auth_user" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
auth_user_groups		CREATE TABLE "auth_user_groups" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
auth_user_user_permissions		CREATE TABLE "auth_user_user_permissions" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
django_admin_log		CREATE TABLE "django_admin_log" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
django_content_type		CREATE TABLE "django_content_type" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
django_migrations		CREATE TABLE "django_migrations" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
django_session		CREATE TABLE "django_session" ("session_key" varchar(40) NOT NULL PRIMARY KEY)
sqlite_sequence		CREATE TABLE sqlite_sequence(name,seq)
Indices (15)		
Views (0)		
Triggers (0)		

Figure 2.9: Contents of the `db.sqlite3` file

2. Now, browse through the newly created database structure by clicking the database tables, and we should see the following:

Name	Type	Schema
Tables (11)		
auth_group		CREATE TABLE "auth_group" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
name	varchar(150)	"name" varchar(150) NOT NULL UNIQUE
auth_group_permissions		CREATE TABLE "auth_group_permissions" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
group_id	integer	"group_id" integer NOT NULL
permission_id	integer	"permission_id" integer NOT NULL
auth_permission		CREATE TABLE "auth_permission" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT)
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
content_type_id	integer	"content_type_id" integer NOT NULL
codename	varchar(100)	"codename" varchar(100) NOT NULL
name	varchar(255)	"name" varchar(255) NOT NULL

Figure 2.10: Browsing through the newly created database structure

3. Notice that the database tables created have different fields, each with their respective data types. Click on the **Browse data** tab in DB Browser and select a table from the drop-down menu. For instance, after clicking the `auth_group_permissions` table, you should see something like this:



The screenshot shows the DB Browser interface with the title bar 'Table: auth_group_permissions'. Below the title are several icons for database operations. The main area displays the table structure with three columns: 'id', 'group_id', and 'permission_id'. Each column has a 'Filter' button below it. The table is currently empty, showing only the column headers and a single row for the header.

id	group_id	permission_id
Filter	Filter	Filter

Figure 2.11: Viewing the auth_group_permissions table

You will see that there is no data available for these tables yet because Django migration only creates the database structure or the blueprint, and the actual data in the database is stored during the operation of the application. Now that we have migrated the built-in or default Django apps, let's try to create an app and perform a Django migration.

In this section, we learned about Django's ORM and how it can simplify database operations for a web application. Later, we learned about how Django's database configuration settings. We also briefly covered Django migrations, which we will go deeper into in the next section by creating models and migrations for our Bookr application.

Creating Django models and migrations

A Django model is essentially a Python class that holds the blueprint for creating a table in a database. The `models.py` file can have many such models, and each model is transformed into a database table. The attributes of the class form the fields and relationships of the database table as per the model definitions.

For our `reviews` application, we need to create the following models and their database tables consequently:

- **Book:** This should store information about books
- **Contributor:** This should store information about the person(s) who contributed to writing the book, such as the author, co-author, or editor
- **Publisher:** As the name implies, this refers to the book publisher
- **Review:** This should store all the book reviews written by the users of the application

Every book in our application will need to have a publisher, so let's create `Publisher` as our first model. Enter the following code in `reviews/models.py`:

```
from django.db import models

class Publisher(models.Model):
    """A company that publishes books."""
    name = models.CharField(
        max_length=50,
        help_text="The name of the Publisher.")
    website = models.URLField(
        help_text="The Publisher's website.")
    email = models.EmailField(
        help_text="The Publisher's email address.")
```

Note

You can take a look at the complete `models.py` file for the `bookr` app by clicking the following link: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter02/final/bookr/reviews/models.py>.

The first line of code imports the Django `models` module. While this line will be autogenerated at the time of the creation of the Django app, do make sure you add it if it is not present. Following the import, the rest of the code defines a class named `Publisher`, a subclass of Django's `models.Model`. Furthermore, this class will have attributes or fields such as `name`, `website`, and `email`. The following are the field types used while creating this model.

As we can see, each of these fields is defined to have the following types:

- `CharField`: This field type stores shorter string fields, such as `Packt Publishing`. For very large strings, we use `TextField`.
- `EmailField`: This is similar to `CharField` but validates whether the string represents a valid email address, for example, `customersupport@packtpub.com`.
- `URLField`: Again, this is similar to `CharField`, but validates whether the string represents a valid URL, for example, `https://www.packtpub.com`.

Next, we will look at the field options used when creating each of these fields.

Field options

Django provides a way to define field options for a model's field. These field options are used to set a value or a constraint, and so on. For example, we can set a default value for a field using `default=<value>` to ensure that every time a record is created in the database for the field, it is set to a default value specified by us. The following are the two field options that we used when defining the `Publisher` model:

- `help_text`: This field option helps us add descriptive text for a field that gets automatically included for Django forms
- `max_length`: This option is provided to `CharField` where it defines the maximum length of the field in terms of the number of characters

Django has many more field types and field options that can be explored from the extensive official Django documentation. As we go about developing our sample book review application, we will learn about the types and fields used for the project. Now, let's migrate the Django models into the database by following these steps:

1. Execute the following command in the shell or terminal to do that (run it from the folder where your `manage.py` file is stored):

```
python manage.py makemigrations reviews
```

The output of the command looks like this:

```
Migrations for 'reviews':  
  reviews/migrations/0001_initial.py  
    - Create model Publisher
```

The `makemigrations <appname>` command creates the migration scripts for the given app, in this case, for the `reviews` app. Notice that after running `makemigrations`, there is a new file created under the `migrations` folder:

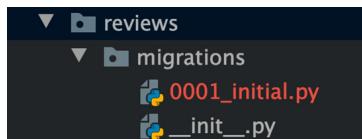


Figure 2.12: New file under the migrations folder

This is the migration script created by Django. When we run `makemigrations` without the app name, the migration scripts will be created for all the apps in the project.

2. Next, let's list the project migration status. Remember that earlier, we applied migrations to Django's installed apps, and now we have created a new app, `reviews`. Run the following command in the shell or terminal, and it will show the status of model migrations throughout the project (run it from the folder where your `manage.py` file is stored):

```
python manage.py showmigrations
```

The output for the preceding command is as follows:

```
admin
[X] 0001_initial
[X] 0002_logentry_remove_auto_add
[X] 0003_logentry_add_action_flag_choices
auth
[X] 0001_initial
[X] 0002_alter_permission_name_max_length
[X] 0003_alter_user_email_max_length
[X] 0004_alter_user_username_opts
[X] 0005_alter_user_last_login_null
[X] 0006_require_contenttypes_0002
[X] 0007_alter_validators_add_error_messages
[X] 0008_alter_user_username_max_length
[X] 0009_alter_user_last_name_max_length
[X] 0010_alter_group_name_max_length
[X] 0011_update_proxy_permissions
contenttypes
[X] 0001_initial
[X] 0002_remove_content_type_name
reviews
[ ] 0001_initial
sessions
[X] 0001_initial
```

Here, the `[X]` mark indicates that the migrations have been applied. Notice the difference that all the other apps' migrations have been applied except that of `reviews`. The `showmigrations` command can be executed to understand the migration status, but this is not a mandatory step while performing model migrations.

3. Next, let's understand how Django transforms a model into an actual database table. To do that, run the `sqlmigrate` command as follows:

```
python manage.py sqlmigrate reviews 0001_initial
```

We should see the following output:

```
BEGIN;
--
-- Create model Publisher
--
CREATE TABLE "reviews_publisher" ("id" integer
    NOT NULL PRIMARY KEY AUTOINCREMENT, "name"
    varchar(50) NOT NULL, "website" varchar(200)
    NOT NULL, "email" varchar(254) NOT NULL);
COMMIT;
```

The preceding snippet shows the SQL command equivalent used when Django migrates the database. In this case, we create the `reviews_publisher` table with the `name`, `website`, and `email` fields with defined field types. Furthermore, all these fields are defined to be `NOT NULL`, implying that the entries for these fields cannot be null and should have a value. The `sqlmigrate` command is not a mandatory step while doing the model migrations.

In the next section, we will learn about primary keys and their importance when storing data in a database.

Primary keys

Let's assume that a database table called `users`, as its name suggests, stores information about users. Let's say it has more than 1,000 records and there are at least 3 users with the same name, Joe Burns. How do we uniquely identify these users from the application? The solution is to have a way to uniquely identify each record in the database. This is done using **primary keys**. A primary key is unique for a database table, and as a rule, a table cannot have two rows with the same primary key. In Django, when the primary key is not explicitly mentioned in the database models, Django automatically creates `id` as the primary key (an integer type), which auto-increments as new records are created.

In the previous section, notice the output of the `python manage.py sqlmigrate` command. When creating the `Publisher` table, the SQL `CREATE TABLE` command added one more field called `id` to the table. An `id` value is defined as `PRIMARY KEY AUTOINCREMENT`. In relational databases, a primary key is used to uniquely identify an entry in the database. For example, the `book` table has `id` as the primary key, which has numbers starting from 1. This value increments by one as new records are created. The integer value of `id` is always unique across the `book` table. Since the migration script has already been created by executing `makemigrations`, let's now migrate the newly created model in the `reviews` app by executing the following command:

```
python manage.py migrate reviews
```

You should get the following output:

```
Operations to perform:
  Apply all migrations: reviews
Running migrations:
  Applying reviews.0001_initial... OK
```

This operation creates the database table for the `reviews` app. The following is a snippet from DB Browser indicating the new `reviews_publisher` table has been created in the database:

		CREATE TABLE "reviews_publisher" ("id" integer NOT N
	id	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
	name	"name" varchar(50) NOT NULL
	website	"website" varchar(200) NOT NULL
	email	"email" varchar(254) NOT NULL

Figure 2.13: The `reviews_publisher` table created after executing the migration command

So far, we have explored how to create a model and migrate it into the database. Let's now work on creating the rest of the models for our book review application. As we've already seen, the application will have the following database tables:

- **Book:** This is the database table that holds the information about the book itself. We have already created a `Book` model and have migrated this to the database.
- **Publisher:** This table holds information about the book publisher.
- **Contributor:** This table holds information about the contributor, that is, the author, co-author, or editor.
- **Review:** This table holds information about the review comments posted by the reviewers.

Let's add the `Book` and `Contributor` models, as shown in the following code snippet, into `reviews/models.py`:

```
class Book(models.Model):
    """A published book."""
    title = models.CharField(
        max_length=70, help_text="The title of the book.")
    publication_date = models.DateField(
        verbose_name="Date the book was published.")
    isbn = models.CharField(
        max_length=20, verbose_name="ISBN number of the
        book.")
class Contributor(models.Model):
    """A contributor to a Book, e.g. author, editor,
    co-author."""
```

```
first_names = models.CharField(
    max_length=50, help_text="The contributor's first
    name or names.")
last_names = models.CharField(
    max_length=50, help_text="The contributor's last
    name or names.")
email = models.EmailField(
    help_text="The contact email for the contributor.")
```

The code is self-explanatory. The `Book` model has the `title`, `publication_date`, and `isbn` fields. The `Contributor` model has the `first_names` and `last_names` fields and the `email` ID of the contributor. There are also some newly added models, apart from the ones we have seen in the `Publisher` model. They have `DateField` as a new field type, which, as the name suggests, is used to store a date. A new field option called `verbose_name` is also used. It provides a descriptive name for the field. Next, we will see how relationships work in a relational database.

Relationships

One of the powers of relational databases is the ability to establish relationships between data stored across database tables. Relationships help maintain data integrity by establishing the correct references across tables, which in turn helps maintain the database. Relationship rules, on the other hand, ensure data consistency and prevent duplicates.

In a relational database, there can be the following types of relations:

- Many-to-one
- Many-to-many
- One-to-one

Let's explore each relationship in detail in the following sections.

Many-to-one

In this relationship, many records (rows/entries) from one table can refer to one record (row/entry) in another table. For example, there can be many books produced by one publisher. This is an example of a **many-to-one relationship**. To establish this relationship, we need to use the database's foreign keys. A foreign key in a relational database establishes the relationship between a field from one table and a primary key from a different table.

For example, say you have data about employees belonging to different departments stored in a table called `employee_info` with their employee ID as the primary key alongside a column that stores their department name; this table also contains a column that stores that department's department ID. Now, there's another table called `departments_info`, which has the department ID as the primary key. In this case, then, the department ID is a foreign key in the `employee_info` table.

In our bookr app, the Book model can have a foreign key referring to the primary key of the Publisher table. Since we have already created the models for Book, Contributor, and Publisher, let's now establish a many-to-one relationship across the Book and Publisher models. For the Book model, add the last line:

```
class Book(models.Model):
    """A published book."""
    title = models.CharField(
        max_length=70, help_text="The title of the book.")
    publication_date = models.DateField(
        verbose_name="Date the book was published.")
    isbn = models.CharField(
        max_length=20, verbose_name="ISBN number of the
        book.")
    publisher = models.ForeignKey(
        Publisher, on_delete=models.CASCADE)
```

Now the newly added publisher field establishes a many-to-one relationship between Book and Publisher using a foreign key. This relationship ensures the nature of a many-to-one relationship, which is that many books can have one publisher:

- `models.ForeignKey`: This is the field option to establish a many-to-one relationship.
- `Publisher`: When we establish relationships with different tables in Django, we refer to the model that creates the table; in this case, the Publisher table is created by the Publisher model (or the Publisher Python class).
- `on_delete`: This is a field option that determines the action to be taken upon the deletion of the referenced object. In this case, the `on_delete` option is set to `CASCADE` (`models.CASCADE`), which deletes the referenced objects.

For example, assume a publisher has published a set of books. For some reason, if the publisher has to be deleted from the application, the next action is `CASCADE`, which means deleting all the referenced books from the application. There are many more `on_delete` actions, such as the following:

- `PROTECT`: This prevents the record deletion unless all the referenced objects are deleted
- `SET_NULL`: This sets a null value if the database field has been previously configured to store null values
- `SET_DEFAULT`: Sets to a default value on the deletion of the referenced object

We will only use the `CASCADE` option for our book review application.

Many-to-many

In this relationship, multiple records in a table can have a relationship with multiple records in a different table. For example, a book can have multiple co-authors, and each author (contributor) may have written multiple books. So, this forms a **many-to-many relationship** between the Book and Contributor tables:

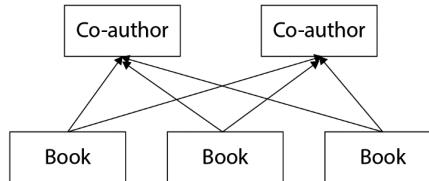


Figure 2.14: Many-to-many relationship between books and co-authors

In `models.py`, for the Book model, add the last line as shown here:

```

class Book(models.Model):
    """A published book."""
    title = models.CharField(
        max_length=70, help_text="The title of the book.")
    publication_date = models.DateField(
        verbose_name="Date the book was published.")
    isbn = models.CharField(
        max_length=20, verbose_name="ISBN number of the
        book.")
    publisher = models.ForeignKey(
        Publisher, on_delete=models.CASCADE)
    contributors = models.ManyToManyField(
        'Contributor', through="BookContributor")
  
```

The newly added `contributors` field establishes a many-to-many relationship with Book and Contributor using the `ManyToManyField` field type:

- `models.ManyToManyField`: This is the field type to establish a many-to-many relationship.
- `through`: This is a special field option for many-to-many relationships. When we have a many-to-many relationship across two tables, if we want to store some extra information about the relationship, then we can use this to establish the relationship via an intermediary table.

For example, we have two tables, namely `Book` and `Contributor`, where we need to store the information on the type of contributor for the book, such as `Author`, `Co-Author`, or `Editor`. Then the type of contributor is stored in an intermediary table called `BookContributor`. Here is how the `BookContributor` table/model looks. Make sure you include this model in `reviews/models.py`:

```
class BookContributor(models.Model):
    class ContributionRole(models.TextChoices):
        AUTHOR = "AUTHOR", "Author"
        CO_AUTHOR = "CO_AUTHOR", "Co-Author"
        EDITOR = "EDITOR", "Editor"

        book = models.ForeignKey(
            Book, on_delete=models.CASCADE)
        contributor = models.ForeignKey(
            Contributor, on_delete=models.CASCADE)
        role = models.CharField(
            verbose_name="The role this contributor had in the \
            book.", choices=ContributionRole.choices,
            max_length=20)
```

Note

The complete `models.py` file can be viewed here: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter02/final/bookr/reviews/models.py>.

An intermediary table such as `BookContributor` establishes relationships by using foreign keys to both the `Book` and `Contributor` tables. It can also have extra fields that can store information about the relationship the `BookContributor` model has with the following fields:

- `book`: This is a foreign key to the `Book` model. As we saw previously, `on_delete=models.CASCADE` will delete an entry from the relationship table when the relevant book is deleted from the application.
- `Contributor`: Again, this is a foreign key to the `Contributor` model/table. This is also defined as `CASCADE` upon deletion.
- `role`: This is the field of the intermediary model, which stores the extra information about the relationship between `Book` and `Contributor`.
- `class ContributionRole(models.TextChoices)`: This can be used to define a set of choices by creating a subclass of `models.TextChoices`. For example, `ContributionRole` is a subclass created out of `TextChoices`, which is used by the `roles` field to define `Author`, `Co-Author`, and `Editor` as a set of choices.

- **choices**: This refers to a set of choices defined in the models, and they are useful when creating Django forms using the models.

Note

When the `through` field option is not provided while establishing a many-to-many relationship, Django automatically creates an intermediary table to manage the relationship.

One-to-one relationships

In this relationship, one record in a table will have a reference to only one record in a different table. For example, a person can have only one driver's license, so a person with a driver's license could form a one-to-one relationship:



Figure 2.15: Example of a one-to-one relationship

`OneToOneField` can be used to establish a one-to-one relationship, as shown here:

```

class DriverLicence(models.Model):
    person = models.OneToOneField(
        Person, on_delete=models.CASCADE)
    licence_number = models.CharField(max_length=50)
  
```

Now that we have explored database relationships, let's come back to our `bookr` application and add one more model there.

Adding the Review model

We've already added the `Book` and `Publisher` models to the `reviews/models.py` file. The last model that we are going to add is the `Review` model. The following code snippet should help us do this:

```

from django.contrib import auth
class Review(models.Model):
    content = models.TextField(
        help_text="The Review text.")
    rating = models.IntegerField(
        help_text="The rating the reviewer has given.")
    date_created = models.DateTimeField(
        auto_now_add=True,
        help_text="The date and time the review was \
created.")
  
```

```
date_edited = models.DateTimeField(
    null=True, help_text="The date and time the review \
    was last edited.")
creator = models.ForeignKey(
    auth.get_user_model(), on_delete=models.CASCADE)
book = models.ForeignKey(
    Book, on_delete=models.CASCADE,
    help_text="The Book that this review is for.")
```

Note

The complete `models.py` file can be viewed here: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter02/final/bookr/reviews/models.py>.

The `review` model/table will be used to store user-provided review comments and ratings for books. It has the following fields:

- `content`: This field stores the text for a book review; hence, the field type used is `TextField`, as this can store a large amount of text.
- `rating`: This field stores the review rating of a book. Since the rating will be an integer, the field type used is `IntegerField`.
- `date_created`: This field stores the time and date when the review was written; hence the field type is `DateTimeField`.
- `date_edited`: This field stores the date and time whenever a review is edited. Again, the field type is `DateTimeField`.
- `Creator`: This field specifies the review creator or the person who writes the book review. Notice that this is a foreign key to `auth.get_user_model()`, which refers to the `User` model from Django's built-in authentication module. It has an `on_delete=models.CASCADE` field option. This explains that when a user is deleted from the database, all the reviews written by that user will be deleted.
- `Book`: Reviews have a field called `book`, a foreign key to the `Book` model. This is because reviews have to be written for a book review application, and a book can have many reviews, so this is a many-to-one relationship. This is also defined with an `on_delete=models.CASCADE` field option because once the book is deleted, there is no point in retaining the reviews in the application. So, when a book is deleted, all the reviews referring to the book will also get deleted.

In the next section, we will learn about and implement model methods.

Model methods

In Django, we can write methods inside a model class. These are called **model methods** and they can be custom or special methods that override the default methods of Django models. One such method is `__str__()`. This method returns the string representation of the Model instances and can be especially useful while using the Django shell. In the following example, where the `__str__()` method is added to the Publisher model, the string representation of the Publisher object will be the publisher's name; hence, `self.name` is returned, with `self` referring to the Publisher object:

```
class Publisher(models.Model):
    """A company that publishes books."""
    name = models.CharField(
        max_length=50,
        help_text="The name of the Publisher.")
    website = models.URLField(
        help_text="The Publisher's website.")
    email = models.EmailField(
        help_text="The Publisher's email address.")

    def str (self):
        return self.name
```

Add the `__str__()` methods to Contributor and Book as well, as follows:

```
class Book(models.Model):
    """A published book."""
    title = models.CharField(
        max_length=70, help_text="The title of the book.")
    publication_date = models.DateField(
        verbose_name="Date the book was published.")
    isbn = models.CharField(
        max_length=20,
        verbose_name="ISBN number of the book.")
    publisher = models.ForeignKey(
        Publisher,
        on_delete=models.CASCADE)
    contributors = models.ManyToManyField(
        'Contributor', through="BookContributor")

    def str (self):
        return self.title

class Contributor(models.Model):
    """A contributor to a Book, e.g. author, editor,
    co-author."""
```

```
first_names = models.CharField(  
    max_length=50,  
    help_text="The contributor's first name or names.")  
last_names = models.CharField(  
    max_length=50,  
    help_text="The contributor's last name or names.")  
email = models.EmailField(  
    help_text="The contact email for the contributor.")  
  
def str (self):  
    return self.first_names
```

Similarly, the string representation of book is title, so the returned value is `self.title`, with `self` referring to the Book object. The string representation of Contributor is the first name of the contributor; hence `self.first_names` is returned. Here, `self` refers to the Contributor object. Next, we will look at migrating the reviews app.

Migrating the reviews app

Since we have the entire model file ready, let's now migrate the models into the database, similar to what we did before with the installed apps. Since the reviews app has a set of models created by us, it is important to create the migration scripts before running the migration. Migration scripts help in identifying any changes to the models and will propagate these changes into the database while running the migration. Follow these steps to create migration scripts and then migrate the models into the database:

1. Execute the following command to create the migration scripts:

```
python manage.py makemigrations reviews
```

You should get an output similar to this:

```
reviews/migrations/0002_auto_20191007_0112.py  
- Create model Book  
- Create model Contributor  
- Create model Review  
- Create model BookContributor  
- Add field contributors to book  
- Add field publisher to book
```

Migration scripts will be created in a folder named `migrations` in the application folder.

2. Next, migrate all the models into the database using the `migrate` command:

```
python manage.py migrate reviews
```

You should see the following output:

```
Operations to perform:
  Apply all migrations: reviews
Running migrations:
  Applying reviews.0002_auto_20191007_0112... OK
```

After executing this command, we successfully created the database tables defined in the reviews app. You may use DB Browser for SQLite to explore the tables you have just created after the migration.

3. To do so, open DB Browser for SQLite, click on the **Open Database** button (*Figure 2.16*), and navigate to your project directory:



Figure 2.16: Click the Open Database button

4. Select the `db.sqlite3` database file to open it (*Figure 2.17*).

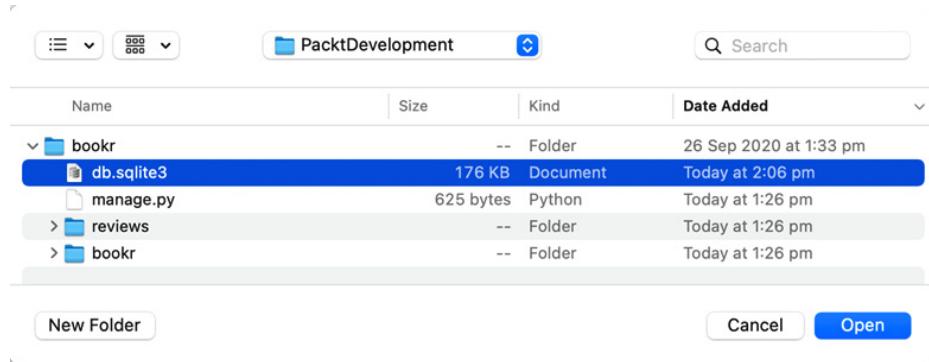


Figure 2.17: Locating `db.sqlite3` in the `bookr` directory

You should now be able to browse the new sets of tables created. The following screenshot shows the database tables defined in the reviews app:

reviews_book	CREATE TABLE "reviews_book" ("id" integer NOT NULL PRIMARY KEY	
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
title	varchar(70)	"title" varchar(70) NOT NULL
publication_date	date	"publication_date" date NOT NULL
isbn	varchar(20)	"isbn" varchar(20) NOT NULL
publisher_id	integer	"publisher_id" integer NOT NULL
reviews_bookcontributor	CREATE TABLE "reviews_bookcontributor" ("id" integer NOT NULL	
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
role	varchar(20)	"role" varchar(20) NOT NULL
book_id	integer	"book_id" integer NOT NULL
contributor_id	integer	"contributor_id" integer NOT NULL
reviews_contributor	CREATE TABLE "reviews_contributor" ("id" integer NOT NULL PRIM	
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
first_names	varchar(50)	"first_names" varchar(50) NOT NULL
last_names	varchar(50)	"last_names" varchar(50) NOT NULL
email	varchar(254)	"email" varchar(254) NOT NULL
reviews_publisher	CREATE TABLE "reviews_publisher" ("id" integer NOT NULL PRIMA	
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
name	varchar(50)	"name" varchar(50) NOT NULL
website	varchar(200)	"website" varchar(200) NOT NULL
email	varchar(254)	"email" varchar(254) NOT NULL
reviews_review	CREATE TABLE "reviews_review" ("id" integer NOT NULL PRIMARY	
id	integer	"id" integer NOT NULL PRIMARY KEY AUTOINCREMENT
content	text	"content" text NOT NULL
rating	integer	"rating" integer NOT NULL
date_created	datetime	"date_created" datetime NOT NULL
date_edited	datetime	"date_edited" datetime
book_id	integer	"book_id" integer NOT NULL
creator_id	integer	"creator_id" integer NOT NULL

Figure 2.18: Database tables as defined in the reviews app

In this section, we learned more about Django models and migrations and how Python's simple classes can transform themselves into database tables. We also learned about how various class attributes translate into appropriate database columns following the defined field types. Later, we learned about primary keys and different types of relationships that can exist in a database. We also created models for the book review application and migrated those models, translating them into database tables. In the next section, we will learn about how to perform database CRUD operations using Django's ORM.

Django's database CRUD operations

As we have created the necessary database tables for the book review application, let's work on understanding the basic database operations with Django.

We've already briefly touched on database operations using SQL statements in the *SQL CRUD operations* section. We tried creating an entry in the database using the `insert` statement, read from the database using the `select` statement, updated an entry using the `update` statement, and deleted an entry from the database using the `delete` statement.

Django's ORM provides the same functionality without having to deal with SQL statements. Django's database operations are simple Python code, hence we overcome the hassle of maintaining SQL statements among the Python code. Let's take a look at how these are performed.

To execute the CRUD operations, we will enter Django's command-line shell by executing the following command:

```
python manage.py shell
```

Note

For this chapter, we will designate Django shell commands using the `>>>` notation (highlighted) at the start of the code block. When pasting the query into DB Browser, make sure you exclude this notation every time.

When the interactive console starts, it looks as follows:

```
Type "help", "copyright", "credits" or "license" for more information.
(InteractiveConsole)
>>>
```

In the following exercise, we shall perform a create operation.

Exercise 2.02 – creating an entry in the bookr database

In this exercise, you will create a new entry in the database by saving a model instance. In other words, you will create an entry in a database table without explicitly running a SQL query. Follow these steps to do this:

1. First, import the Publisher class/model from `reviews.models`:

```
>>> from reviews.models import Publisher
```

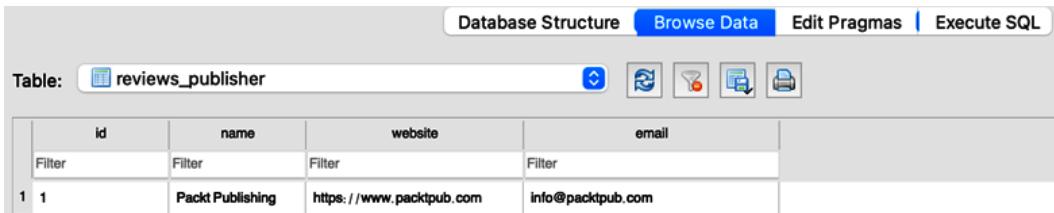
2. Create an object or an instance of the Publisher class by passing all the field values (`name`, `website`, and `email`) required by the Publisher model:

```
>>> publisher = Publisher(name='Packt Publishing',
    website='https://www.packtpub.com',
    email='info@packtpub.com')
```

3. Next, to write the object into the database, it is important to call the `save()` method, because until this is called, the entry will not be created in the database:

```
>>> publisher.save()
```

Now you can see a new entry created in the database using DB Browser:



	id	name	website	email
1	1	Packt Publishing	https://www.packtpub.com	info@packtpub.com

Figure 2.19: An entry created in the database

4. Use the object attributes to make any further changes to the object and save the changes to the database:

```
>>> publisher.email = 'info@packtpub.com'  
>>> publisher.email = 'customersupport@packtpub.com'  
>>> publisher.save()
```

You can see the changes using DB Browser as follows:



	id	name	website	email
1	1	Packt Publishing	https://www.packtpub.com	customersupport@packtpub.com

Figure 2.20: An entry with the updated email field

In this exercise, we created an entry in the database by creating an instance of the model object and used the `save()` method to write the model object into the database.

Note that by following the preceding method, the changes to the class instance are not saved until the `save()` method is called. However, if we use the `create()` method, Django saves the changes to the database in a single step. We'll use this method in the exercise that follows.

Exercise 2.03 – using the `create()` method to create an entry

Unlike the previous exercise where we executed two separate steps by first creating an object and later using the `save()` method to create an entry in the database, in this exercise, you will create a record in the `Contributor` table using the `create()` method in a single step:

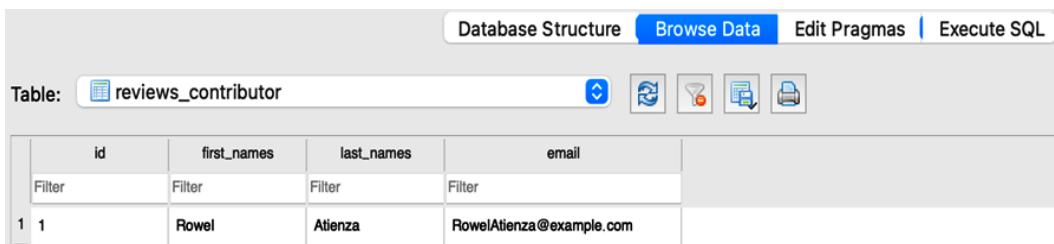
1. First, import the `Contributor` class as before:

```
>>> from reviews.models import Contributor
```

2. Invoke the `create()` method to create an object in the database in a single step. Ensure that you pass all the required parameters (`first_names`, `last_names`, and `email`):

```
>>> contributor =
    Contributor.objects.create(first_names="Rowel",
                                last_names="Atienza",
                                email="RowelAtienza@example.com")
```

3. Use DB Browser to verify that the contributor record has been created in the database. If DB Browser is not already open, open the `db.sqlite3` database file as we just did in the previous section. Click **Browse Data** and select the desired table – in this case, the `reviews_contributor` table from the **Table** drop-down menu, as shown in the screenshot – and verify the newly created database record:



The screenshot shows the DB Browser interface. The top navigation bar includes tabs for 'Database Structure', 'Browse Data' (which is selected and highlighted in blue), 'Edit Pragmas', and 'Execute SQL'. Below the navigation bar, a 'Table' dropdown is set to 'reviews_contributor'. The main area displays a table with four columns: 'id', 'first_names', 'last_names', and 'email'. A single row is present, with the values: '1', 'Rowel', 'Atienza', and 'RowelAtienza@example.com'. Above the table are several icons for filtering, sorting, and other database operations.

	id	first_names	last_names	email
Filter	Filter	Filter	Filter	
1	1	Rowel	Atienza	RowelAtienza@example.com

Figure 2.21: Verifying the creation of the record in DB Browser

In this exercise, we learned that using the `create()` method, we can create a record for a model in a database in a single step.

Creating an object with a foreign key

Similar to how we created a record in the `Publisher` and `Contributor` tables, let's now create one for the `Book` table. If you recall, the `Book` model has a foreign key to `Publisher` that cannot have a null value. So, a way to populate the publisher's foreign key is by providing the created `Publisher` object in the book's `publisher` field as shown in the following exercise.

Exercise 2.04 – creating records for a many-to-one relationship

In this exercise, you will create a record in the `Book` table, including a foreign key to the `Publisher` model. As you already know, the relationship between `Book` and `Publisher` is a many-to-one relationship, so you have to first fetch the `Publisher` object and then use it while creating the book record. To do that, follow these steps:

1. First, import the `Publisher` class:

```
>>> from reviews.models import Book, Publisher
```

2. Retrieve the publisher object from the database using the following command. The `get()` method is used to retrieve an object from the database. We still haven't explored database read operations. For now, use the following command; we will go deeper into reading/retrieving from databases in the next section:

```
>>> publisher =  
    Publisher.objects.get(name='Packt Publishing')
```

3. When creating a book, we need to supply a date object, as `publication_date` is a date field in the `Book` model. So, import `date` from `datetime` so that a `date` object can be supplied when creating the `book` object, as shown in the following code:

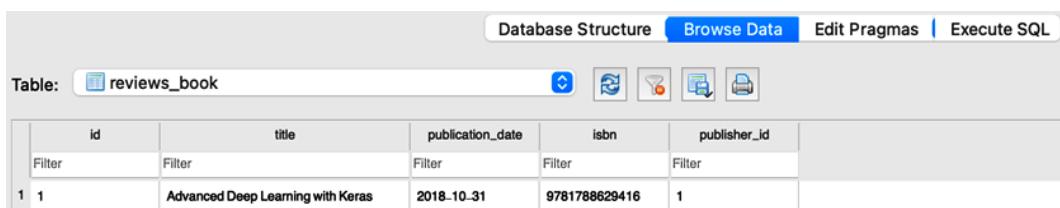
```
>>> from datetime import date
```

4. Use the `create()` method to create a record of the book in the database. Ensure that you pass all the fields, namely `title`, `publication_date`, `isbn`, and the `publisher` object:

```
>>> book = Book.objects.create(title="Advanced Deep  
    Learning with Keras", publication_date=date(2018,  
    10, 31), isbn="9781788629416",  
    publisher=publisher)
```

Note that since `publisher` is a foreign key and it is not nullable (cannot hold a `null` value), it is mandatory to pass a `publisher` object. When the mandatory foreign key object, `publisher`, is not provided, the database will throw an integrity error.

Figure 2.22 shows the `Book` table where the first entry is created. Notice that the foreign key field (`publisher_id`) points to the `id` value (primary key) of the `Publisher` table. The `publisher_id` entry in the book's record points to a `Publisher` record has a primary key of 1, as shown in the following two screenshots:



The screenshot shows a database browser interface with the following details:

- Header: Database Structure, Browse Data, Edit Pragmas, Execute SQL.
- Table: reviews_book
- Columns: id, title, publication_date, isbn, publisher_id.
- Row 1: 1, Advanced Deep Learning with Keras, 2018-10-31, 9781788629416, 1.
- Buttons: Filter, Sort, Refresh, New, Open, Save, Delete, Copy.

Figure 2.22: Foreign key pointing to the primary key for `reviews_book`

Table: reviews_publisher				
	id	name	website	email
Filter	Filter	Filter	Filter	Filter
1	1	Packt Publishing	https://www.packtpub.com	customersupport@packtpub.com

Figure 2.23: Foreign key pointing to the primary key for reviews_publisher

In this exercise, we learned that when creating a database record, an object can be assigned to a field if it is a foreign key. We know that the Book model also has a many-to-many relationship with the Contributor model. Let's now explore the ways to establish many-to-many relations as we create records in the database.

Exercise 2.05 – creating records with many-to-many relationships

In this exercise, you will create a many-to-many relationship between Book and Contributor using the BookContributor relationship model:

1. In case you have restarted the shell and lost the publisher and book objects, retrieve them from the database by using the following set of Python statements:

```
>>> from reviews.models import Book
>>> from reviews.models import Contributor
>>> contributor =
    Contributor.objects.get(first_names='Rowel')
>>> book = Book.objects.get(title="Advanced Deep
    Learning with Keras")
```

2. The way to establish a many-to-many relationship is by storing the information about the relationship in the intermediary model or the relationship model; in this case, BookContributor. Since we have already fetched the book and the contributor records from the database, let's use these objects when creating a record for the BookContributor relationship model. To do so, first, create an instance of the BookContributor relationship class and then save the object to the database. While doing so, ensure you pass the required fields, namely the book object, the contributor object, and role:

```
>>> from reviews.models import BookContributor
>>> book_contributor = BookContributor(book=book,
    contributor=contributor, role='AUTHOR')
>>> book_contributor.save()
```

Notice that we set role to AUTHOR while creating the `book_contributor` object. This is a classic example of storing relationship data while establishing a many-to-many relationship. The role can be AUTHOR, CO_AUTHOR, or EDITOR.

This established the relationship between the book *Advanced Deep Learning with Keras* and the contributor Rowel (Rowel being the author of the book).

In this exercise, we established a many-to-many relationship between Book and Contributor using the `BookContributor` relationship model. With regards to the verification of the many-to-many relationship that we just created, we will see this in detail in a few exercises later on in this chapter.

Exercise 2.06 – a many-to-many relationship using the `add()` method

In this exercise, you will establish a many-to-many relationship using the `add()` method. When we don't use a relationship to create the objects, we can use `through_defaults` to pass in a dictionary with the parameters defining the required fields. Continuing from the previous exercise, let's add one more contributor to the book titled *Advanced Deep Learning with Keras*. This time, the contributor is an editor of the book:

1. If you have restarted the shell, run the following two commands to import and fetch the desired book instance:

```
>>> from reviews.models import Book, Contributor
>>> book = Book.objects.get(title="Advanced Deep
Learning with Keras")
```

2. Use the `create()` method to create a contributor, as shown here:

```
>>> contributor =
    Contributor.objects.create(first_names='Packt',
    last_names='Example Editor',
    email='PacktEditor@example.com')
```

3. Add the newly created contributor to the book using the `add()` method. Ensure you provide the `role` relationship parameter as `dict`. Enter the following code:

```
>>> book.contributors.add(contributor,
    through_defaults={'role': 'EDITOR'})
```

Thus, we used the `add()` method to establish a many-to-many relationship between the book and the contributor while storing the relationship data role as EDITOR. Let's now take a look at other ways of doing this.

Using the `create()` and `set()` methods for many-to-many relationships

Assume the book *Advanced Deep Learning with Keras* has a total of two editors. Let's use the following method to add another editor to the book. If the contributor is not already present in the database, then we can use the `create()` method to simultaneously create an entry as well as to establish its relationship with the book:

```
>>> book.contributors.create(first_names='Packtp',
   last_names='Editor Example',
   email='PacktEditor2@example.com',
   through_defaults={'role': 'EDITOR'})
```

Similarly, we can also use the `set()` method to add a list of contributors for a book. Let's create a publisher, a set of two contributors who are the co-authors, and a `book` object. First, import the `Publisher` model, if not already imported, using the following code:

```
>>> from reviews.models import Publisher
```

The following code will help us do so:

```
>>> publisher =
    Publisher.objects.create(name='Pocket Books',
    website='https://pocketbookssampleurl.com',
    email='pocketbook@example.com')
>>> contributor1 =
    Contributor.objects.create(first_names='Stephen',
    last_names='Stephen', email='StephenKing@example.com')
>>> contributor2 =
    Contributor.objects.create(first_names='Peter',
    last_names='Straub', email='PeterStraub@example.com')

>>> book = Book.objects.create(title='The Talisman',
    publication_date=date(2012, 9, 25),
    isbn='9781451697216', publisher=publisher)
```

Since this is a many-to-many relationship, we can add a list of objects in just one go using the `set()` method. We can use `through_defaults` to specify the role of the contributors; in this case, they are co-authors:

```
>>> book.contributors.set([contributor1, contributor2],
   through_defaults={'role': 'CO_AUTHOR'})
```

Read operations

Django provides us with methods that allow us to read/retrieve from the database. We can retrieve a single object from the database using the `get()` method. We have already created a few records in the previous sections, so let's use the `get()` method to retrieve an object.

Exercise 2.07 – using the `get()` method to retrieve an object

In this exercise, you will retrieve an object from the database using the `get()` method. Follow these steps:

1. Fetch a `Publisher` object that has a `name` field of `Pocket Books`:

```
>>> from reviews.models import Publisher
>>> publisher =
    Publisher.objects.get(name='Pocket Books')
```

2. Re-enter the retrieved `publisher` object and press *Enter*:

```
>>> publisher
<Publisher: Pocket Books>
```

Notice that the output is displayed in the shell. This is called a string representation of an object. It is the result of adding the `__str__()` model method as we did in the *Model methods* section for the `Publisher` class.

3. Upon retrieving the object, you have access to all the object's attributes. Since this is a Python object, the attributes of the object can be accessed by using `.` followed by the attribute name. So, you can retrieve the publisher's name with the following command:

```
>>> publisher.name
'Pocket Books'
```

4. Similarly, to retrieve the publisher's website, use the following:

```
>>> publisher.website
'https://pocketbookssampleurl.com'
```

The publisher's email address can be retrieved as follows:

```
>>> publisher.email
'pocketbook@example.com'
```

In this exercise, we learned how to fetch a single object using the `get()` method. There are several disadvantages to using this method, though. Let's find out what they are next.

Returning an object using the `get()` method

It is important to note that the `get()` method can only fetch one object. If there is another object carrying the same value as the field mentioned, then we can expect a *returned more than one* error message. For example, if there are two entries in the `Publisher` table with the same value for the `name` field, we can expect an error. In such cases, there are alternate ways to retrieve those objects, which we will explore in the subsequent sections.

We can also get a *matching query does not exist* error message when there are no objects returned by the `get()` query. The `get()` method can be used with any of the object's fields to retrieve a record. In the following case, we use the `website` field:

```
>>> publisher =  
    Publisher.objects.get(  
        website='https://pocketbookssampleurl.com')
```

After retrieving the object, we can still get the publisher's name, as shown here:

```
>>> publisher.name  
'Pocket Books'
```

Another way to retrieve an object is by using its primary key, `pk`, as can be seen here:

```
>>> Publisher.objects.get(pk=2)  
<Publisher: Pocket Books>
```

Using `pk` for the primary key is a more generic way of using the primary key field. But for the `Publisher` table, since we know that `id` is the primary key, we can simply use the `id` field name to create our `get()` query:

```
>>> Publisher.objects.get(id=2)  
<Publisher: Pocket Books>
```

Note

For `Publisher` and all the other tables, the primary key is `id`, which Django automatically creates. This happens when a primary key field is not mentioned at the time of the creation of the table. But there can be instances where a field can be explicitly declared as a primary key.

In the next exercise, we will see how to query and retrieve all the objects for a given model.

Exercise 2.08 – using the `all()` method to retrieve a set of objects

We can use the `all()` method to retrieve a set of all objects. In this exercise, you will use this method to retrieve the names of all contributors. To do that, follow these steps:

1. Add the following code to retrieve all the objects from the `Contributor` table:

```
>>> from reviews.models import Contributor

>>> Contributor.objects.all()
<QuerySet [<Contributor: Rowel>, <Contributor: Packt>,
<Contributor: Packtp>, <Contributor: Stephen>, <Contributor:
Peter>] >
```

Upon execution, you will get `QuerySet` of all the objects.

2. We can use list indexing to look up a specific object or to iterate over the list using a loop to do any other operation:

```
>>> contributors = Contributor.objects.all()
```

3. Since `Contributor` is a list of objects, you can use indexing to access any element in the list, as shown in the following command:

```
>>> contributors[0]
<Contributor: Rowel>
```

In this case, the first element in the list is a contributor with a `first_names` value of 'Rowel' and a `last_names` value of 'Atienza', as you can see from the following code:

```
>>> contributors[0].first_names
'Rowel'
>>> contributors[0].last_names
'Atienza'
```

In this exercise, we learned how to retrieve all the objects using the `all()` method and how to use the retrieved set of objects as a list.

Retrieving objects by filtering

If we have more than one object for a field value, we cannot use the `get()` method since the `get()` method can return only one object. For such cases, we have the `filter()` method, which can retrieve all the objects that match a specified condition.

Exercise 2.09 – using the filter() method to retrieve objects

In this exercise, you will use the `filter()` method to get a specific set of objects for a certain condition. Specifically, you will retrieve all the contributors' names whose first name is Peter. To do that, follow these steps:

1. First, create two more contributors:

```
>>> from reviews.models import Contributor
>>> Contributor.objects.create(first_names='Peter',
    last_names='Wharton',
    email='PeterWharton@example.com')
>>> Contributor.objects.create(first_names='Peter',
    last_names='Tyrrell',
    email='PeterTyrrell@example.com')
```

2. To retrieve contributors who have the `first_names` value of Peter, add the following code:

```
>>> Contributor.objects.filter(first_names='Peter')
<QuerySet [<Contributor: Peter>, <Contributor: Peter>,
<Contributor: Peter>]>
```

3. The `filter()` method returns the object even if there is only one. You can see this here:

```
>>> Contributor.objects.filter(first_names='Rowel')
<QuerySet [<Contributor: Rowel>]>
```

4. Furthermore, the `filter()` method returns an empty `QuerySet` if nothing matches the query. This can be seen here:

```
>>> Contributor.objects.filter(first_names='Nobody')
<QuerySet []>
```

In this exercise, we saw the use of filters to retrieve a set of a few objects filtered by a certain condition. In the next section, we will learn about filtering by using field lookups.

Filtering by field lookups

Now, let's suppose we want to filter and query a set of objects using the object's fields by providing certain conditions. In such a case, we can use a **double-underscore lookup**. For example, the `Book` object has a `publication_date` field; let's say we want to filter and fetch all the books that were published after 01-01-2014. We can easily look these up by using the double-underscore method. To do this, we will first import the `Book` model:

```
>>> from reviews.models import Book
>>> book = Book.objects.filter(publication_date__gt=date(
    2014, 1, 1))
```

Here, `publication_date__gt` indicates the publication date, which is greater than (`gt`) a certain specified date, in this case, 01-01-2014. Similar to this, we have the following abbreviations:

- `lt`: Less than
- `lte`: Less than or equal to
- `gte`: Greater than or equal to

The result after filtering can be seen here:

```
>>> book
<QuerySet [<Book: Advanced Deep Learning with Keras>]>
```

Here is the publication date of the book that is part of the query set, which confirms that the publication date was after 01-01-2014:

```
>>> book[0].publication_date
datetime.date(2018, 10, 31)
```

Using pattern matching for filtering operations

For filtered results, we can also look up whether the parameter contains a part of the string we are looking for:

```
>>> book =
    Book.objects.filter(title__contains='Deep learning')
```

Here, `title__contains` looks for all objects with titles containing 'Deep learning' as a part of the string:

```
>>> book
<QuerySet [<Book: Advanced Deep Learning with Keras>]>

>>> book[0].title
'Advanced Deep Learning with Keras'
```

Similarly, we can use `icontains` if the string match needs to be case-insensitive. Using `startswith` matches any string starting with the specified string.

Retrieving objects by using the exclude() method

In the previous section, we learned about fetching a set of objects by matching a certain condition. Now, suppose we want to do the opposite; that is, we want to fetch all those objects that do not match a certain condition. In such cases, we can use the `exclude()` method to exclude a certain condition and fetch all the required objects. This will be clearer with an example. The following is a list of all contributors:

```
>>> Contributor.objects.all()
<QuerySet [<Contributor: Rowel>, <Contributor: Packt>,
<Contributor: Packtp>, <Contributor: Stephen>,
<Contributor: Peter>, <Contributor: Peter>,
<Contributor: Peter>] >
```

Now, from this list, we will exclude all those contributors who have a `first_names` value of Peter:

```
>>> Contributor.objects.exclude(first_names='Peter')
<QuerySet [<Contributor: Rowel>, <Contributor: Packt>,
<Contributor: Packtp>, <Contributor: Stephen>] >
```

We see here that the query returned all contributors whose first name is not Peter.

Retrieving objects using the order_by() method

We can retrieve a list of objects while ordering by a specified field using the `order_by()` method. For example, in the following code snippet, we order the books by their publication date:

```
>>> books = Book.objects.order_by("publication_date")
>>> books
<QuerySet [<Book: The Talisman>, <Book: Advanced Deep Learning with
Keras>] >
```

Let's examine the order of the query. Since the query set is a list, we can use indexing to check the publication date of each book:

```
>>> books[0].publication_date
datetime.date(2012, 9, 25)
>>> books[1].publication_date
datetime.date(2018, 10, 31)
```

Notice that the publication date of the first book with a 0 index is older than the publication date of the second book with a 1 index. So, this confirms that the queried list of books has been properly ordered as per their publication dates. We can also use a prefix with the negative sign for the field parameter to order results in descending order. This can be seen from the following code snippet:

```
>>> books = Book.objects.order_by("-publication_date")
>>> books

<QuerySet [<Book: Advanced Deep Learning with Keras>,
<Book: The Talisman>] >
```

Since we have prefixed a negative sign to the publication date, notice that the queried set of books has now been returned in the opposite order, where the first book object with a 0 index has a more recent date than the second book:

```
>>> books[0].publication_date
datetime.date(2018, 10, 31)

>>> books[1].publication_date
datetime.date(2012, 9, 25)
```

We can also order by using a string field or a numerical one. For example, the following code can be used to order books by their primary key or id:

```
>>> books = Book.objects.order_by('id')
>>> books
<QuerySet [<Book: Advanced Deep Learning with Keras>,
<Book: The Talisman>] >
```

The queried set of books has been ordered by book id in ascending order:

```
>>> books[0].id
1
>>> books[1].id
2
```

Again, to order it in descending order, the negative sign can be used as a prefix, as follows:

```
>>> books = Book.objects.order_by('-id')
>>> books
<QuerySet [<Book: The Talisman>, <Book: Advanced Deep Learning with
Keras>] >
```

Now, the queried set of books has been ordered by book `id` in descending order:

```
>>> books[0].id  
2  
>>> books[1].id  
1
```

To order by a string field in alphabetical order, we can do something like this:

```
>>> books = Book.objects.order_by('title')  
>>> books  
<QuerySet [<Book: Advanced Deep Learning with Keras>, <Book: The  
Talisman>] >
```

Since we have used the title of the book to order by, the query set has been ordered in alphabetical order. We can see this as follows:

```
>>> books[0]  
<Book: Advanced Deep Learning with Keras>  
>>> books[1]  
<Book: The Talisman>
```

Similar to what we've seen for the previous ordering types, the negative sign prefix can help us sort in reverse alphabetical order, as we can see here:

```
>>> books = Book.objects.order_by('-title')  
>>> books  
<QuerySet [<Book: The Talisman>, <Book: Advanced Deep Learning with  
Keras>] >
```

This will lead to the following output:

```
>>> books[0]  
<Book: The Talisman>  
>>> books[1]  
<Book: Advanced Deep Learning with Keras>
```

Yet another useful method offered by Django is `values()`. It helps us get a query set of dictionaries instead of objects. In the following code snippet, we're using this for a `Publisher` object:

```
>>> publishers = Publisher.objects.all().values()  
  
>>> publishers  
<QuerySet [{  
    'id': 1, 'name': 'Packt Publishing', 'website':  
    'https://www.packtpub.com', 'email':  
    'customersupport@packtpub.com'}, {  
    'id': 2, 'name':  
    'Pocket Books', 'website':  
    'https://pocketbookssampleurl.com',
```

```
'email': 'pocketbook@example.com'}] >
>>> publishers[0]
{'id': 1, 'name': 'Packt Publishing', 'website':
'https://www.packtpub.com', 'email':
'customersupport@packtpub.com'}

>>> publishers[0]
{'id': 1, 'name': 'Packt Publishing', 'website':
'https://www.packtpub.com', 'email':
'customersupport@packtpub.com'}
```

In the next few sections, we will explore how to query across relationships.

Querying across relationships

As we have studied in this chapter, the `reviews` app has two kinds of relationships: many-to-one and many-to-many. So far, we have learned various ways of making queries using `get()`, filters, field lookups, and so on. Now let's study how to perform queries across relationships. There are several ways to go about this – we could use foreign keys, object instances, and more. Let's explore these with the help of some examples.

Querying using foreign keys

When we have relationships across two models/tables, Django provides a way to perform a query using the relationship. The command shown in this section will retrieve all the books published by `Packt Publishing` by performing a query using model relationships. Similar to what we've seen previously, this is done using the double-underscore lookup. For example, the `Book` model has a foreign key of `publisher` pointing to the `Publisher` model. Using this foreign key, we can perform a query using double underscores and the `name` field in the `Publisher` model. This can be seen from the following code:

```
>>> Book.objects.filter(publisher__name='Packt Publishing')
<QuerySet [Book: Advanced Deep Learning with Keras]>
```

Querying using the model name

Another way of querying is using a relationship to do the query backward, using the model name in lowercase. For instance, let's say we want to query the publisher who published the book *Advanced Deep Learning with Keras* using model relationships in the query. For this, we can execute the following statement to retrieve the `Publisher` information object:

```
>>> Publisher.objects.get(book__title='Advanced Deep
Learning with Keras')
<Publisher: Packt Publishing>
```

Here, `book` is the model's name in lowercase. As we already know, the `Book` model has a `publisher` foreign key with a `name` value of `Packt Publishing`.

Querying across foreign key relationships using the object instance

We can also retrieve the information using the object's foreign key. Suppose we want to query the publisher's name for the title *The Talisman*:

```
>>> book = Book.objects.get(title='The Talisman')
>>> book.publisher
<Publisher: Pocket Books>
```

Using the object here is an example of where we use the reverse direction to get all the books published by a publisher by using the `set.all()` method:

```
>>> publisher = Publisher.objects.get(name='Pocket Books')

>>> publisher.book_set.all()
<QuerySet [<Book: The Talisman>]>
```

We can also create queries using chains of queries:

```
>>> Book.objects.filter(publisher__name='Pocket
      Books').filter(title='The Talisman')
<QuerySet [<Book: The Talisman>]>
```

Let's perform some more exercises to shore up our knowledge of the various kinds of queries we have learned about so far.

Exercise 2.10 – querying across a many-to-many relationship using the field lookup

We know that `Book` and `Contributor` have a many-to-many relationship. In this exercise, without creating an object, you will perform a query to retrieve all the contributors who contributed to writing the book titled *The Talisman*:

1. First, import the `Contributor` class:

```
>>> from reviews.models import Contributor
```

2. Now, add the following code to query for the set of contributors to *The Talisman*:

```
>>> Contributor.objects.filter(book__title='The
      Talisman')
```

You should see the following:

```
<QuerySet [<Contributor: Stephen>, <Contributor: Peter>]>
```

From the preceding output, we can see that Stephen and Peter are the contributors who contributed to writing the book *The Talisman*. The query uses the book model (written in lowercase) and does a field lookup for the `title` field using a double underscore, as shown in the command.

In this exercise, we learned how to perform queries across many-to-many relationships using a field lookup. Let's now look at using another method to carry out the same task.

Exercise 2.11 – a many-to-many query using objects

In this exercise, using a `Book` object, search for all the contributors who contributed to writing the book with the title *The Talisman*. The following steps will help you do that:

1. Import the `Book` model:

```
>>> from reviews.models import Book
```

2. Retrieve a `book` object with the title *The Talisman*, by adding the following line of code:

```
>>> book = Book.objects.get(title='The Talisman')
```

3. Then, retrieve all the contributors who worked on the book *The Talisman* using the `book` object. Add the following code to do so:

```
>>> book.contributors.all()  
<QuerySet [<Contributor: Stephen>, <Contributor: Peter>]>
```

Again, we can see that Stephen and Peter are the contributors who worked on the book *The Talisman*. Since the book has a many-to-many relationship with contributors, we have used the `contributors.all()` method to get a query set of all those contributors who worked on the book. Now, let's try using the `set()` method to perform a similar task.

Exercise 2.12 – a many-to-many query using the `set()` method

In this exercise, you will use a `contributor` object to fetch all the books written by the contributor named Rowel:

1. Import the `Contributor` model:

```
>>> from reviews.models import Contributor
```

2. Fetch a `contributor` object whose `first_names` is 'Rowel' using the `get()` method:

```
>>> contributor =  
Contributor.objects.get(first_names='Rowel')
```

3. Using the `contributor` object and the `book_set()` method, get all those books written by the contributor:

```
>>> contributor.book_set.all()
<QuerySet [<Book: Advanced Deep Learning with Keras>]>
```

Since `Book` and `Contributor` have a many-to-many relationship, we can use the `set()` method to query a set of objects associated with the model. In this case, `contributor.book_set.all()` returned all the books written by the contributor.

Exercise 2.13 – using the `update()` method

In this exercise, you will use the `update()` method to update an existing record:

1. Change `first_names` for a contributor who has the last name `Tyrrell`:

```
>>> from reviews.models import Contributor
>>> Contributor.objects.filter(last_names='Tyrrell').
    update(first_names='Mike')
1
```

The return value shows the number of records that have been updated. In this case, one record has been updated.

2. Fetch the contributor that was just modified using the `get()` method and verify that the first name has been changed to `Mike`:

```
>>> Contributor.objects.get(last_names='Tyrrell').
    first_names
'Mike'
```

Note

If the filter operation has more than one record, then the `update()` method will update the specified field in all the records returned by the filter.

In this exercise, we learned how to use the `update()` method to update a record in the database. Now, finally, let's try deleting a record from the database using the `delete()` method.

Exercise 2.14 – using the delete() method

An existing record in the database can be deleted using the `delete()` method. In this exercise, you will delete a record from the `contributors` table that has the `last_name` value of Wharton:

1. Fetch the object using the `get()` method and use the `delete()` method, as shown here:

```
>>> from reviews.models import Contributor
>>> Contributor.objects.get(last_name='Wharton').
    delete()
(1, {'reviews.Contributor': 1})
```

Notice that you called the `delete()` method without assigning the `contributor` object to a variable. Since the `get()` method returns a single object, you can access the object's method without actually creating a variable for it.

2. Verify the `contributor` object with `last_name` as 'Wharton' has been deleted:

```
>>> Contributor.objects.get(last_name='Wharton')
Traceback (most recent call last):
File "<console>", line 1, in <module>
File "/.../site-packages/django/db/models/manager.py",
line 82, in manager_method
    return getattr(self.get_queryset(), name)(*args,
**kwargs)
File "/.../site-packages/django/db/models/query.py",
line 417, in get
    self.model._meta.object_name
reviews.models.Contributor.DoesNotExist: Contributor
matching query does not exist.
```

As you can see, upon running the query, we got an *object does not exist* error. This is expected since the record has been deleted. In this exercise, we learned how to use the `delete()` method to delete a record from the database.

Bulk create and bulk update operations

When we have a large set of records that needs to be created or updated, we can perform bulk create and bulk update operations using the `bulk_create()` and `bulk_update()` methods, respectively. When we have a large number of records to be created or updated, using methods such as these can be efficient when performing database operations.

Typically, here is how a `bulk_create` method is called by supplying a list of objects:

```
Person.objects.bulk_create([
    Person(name='Robert', address='5, Byron bay, NSW'),
    Person(name='Mark', address="Unit 12, New town, NSW
2000"),])
```

Likewise, the bulk update is also performed on a list of similar objects. The object's attributes can be changed as shown in the following example and updated in the database in a single command using the `bulk_update()` method, as shown in the following snippet:

```
persons[0].address = "8, Byron bay, NSW"
persons[1].address = "15, New town, NSW"
Person.objects.bulk_update(persons, ["address"])
```

We will see further examples of both of these in the following exercises.

Exercise 2.15 – creating multiple records using `bulk_create`

In this exercise, we will use the `bulk_create()` method to create multiple entries into the `Publisher` table in one go. To do that, follow these steps:

1. Import the `Publisher` model if you have not already imported it into your Django shell:

```
>>> from reviews.models import Publisher
```

2. Create multiple records in the `Publisher` table by passing a list of the `Publisher` objects into the `bulk_create()` method:

```
>>> Publisher.objects.bulk_create([
    Publisher(name="New Town Publisher",
    website="www.newtownexample.com",
    email='newtow@email.com'), Publisher(name="Byron
    Bay Press", website="www.bryonbayexample.com",
    email='bryonbayexample@email.com'),
    Publisher(name="Katoomba Publisher",
    website="www.katoombaexample.com",
    email='katoombaexample@email.com')])
```



```
[<Publisher: New Town Publisher>, <Publisher: Byron Bay Press>,
<Publisher: Katoomba Publisher>]
```

The method returns a list of objects created as entries in the database. If a large number of entries need to be created, this might be one of the most efficient ways of carrying out the task.

Next, we shall try to bulk update records efficiently using the `bulk_update` method.

Exercise 2.16 – updating multiple records using bulk_update

Similar to the previous exercise, we can update multiple records in one go by using the `bulk_update` method:

1. First, import the `Publisher` model if you have not already imported it:

```
>>> from reviews.models import Publisher
```

2. In the next step, we will pick a couple of `Publisher` objects as a list.
3. Assuming both of these publishers were merged into a single company, and they need to share the same website, now let us update both the objects in a query set:

```
>>> publishers = [Publisher.objects.get(name='New Town
Publisher'), Publisher.objects.get(name='Byron Bay
Press')]
>>> publishers[0].website =
"www.newsouthwalespublisher.com"
>>> publishers[1].website =
"www.newsouthwalespublisher.com"

>>> Publisher.objects.bulk_update(publishers,
["website"])
2
```

Upon running the `bulk_update` method, it returns the number of objects updated; in this case, it is two objects. Again, this is an efficient way to update more than one object in one go.

Performing complex lookups using Q objects

`Q` objects are used to perform complex queries especially when a query involves the `AND` or `OR` operations in a `WHERE` clause. For instance, if we need to do a query similar to the SQL query shown here:

```
SELECT * FROM Person WHERE name LIKE "Rob%" OR name LIKE "Bob%";
```

The preceding SQL statement queries for any person whose name either starts with `Rob` or `Bob`.

In `LIKE "Rob%"`, here, the `LIKE` keyword pattern-matches a string to check whether the string starts with the specified `Rob` value.

`Q` objects use the `&` and `|` operators for the AND and OR operations when combining the WHERE clauses. The preceding query can be written as follows using `Q` objects:

```
Person.objects.get(Q(name__startswith='Rob') & Q(name__startswith='Bob'))
```

Exercise 2.17 – performing a complex query using a `Q` object

Next, using the concept learned before, `Q` objects let us perform a query for two of the publishers:

1. Import `Publisher` if you have not already done so and also import `Q` from `django.db.models`, as shown in the following command:

```
>>> from django.db.models import Q
>>> from reviews.models import Publisher
```

2. Execute the `Q` object query as shown next. Here, we check whether any of the `Publisher` objects has name starting with either `New` or `Idea`:

```
>>> Publisher.objects.filter(Q(name__startswith="New") | Q(name__startswith="Idea"))
<QuerySet [<Publisher: New Town Publisher>]>
```

Since there is no `Publisher` with name starting with `Idea` and there was an object with name starting with `New`, one object was returned.

This is an example `Q` object query with the AND operator. In this query, we query a `Publisher` object whose name starts with `New` and also whose name ends with `Publisher`:

```
>>> Publisher.objects.filter(Q(name__startswith="New") & Q(name__endswith="Publisher"))
<QuerySet [<Publisher: New Town Publisher>]>
```

In this case, only one object was returned, which is `New Town Publisher`, which satisfies the conditions stated in the query. This way, multiple such WHERE clauses can be combined using the AND or OR operators to get the desired results.

Exercise 2.18 – verifying whether a queryset contains a given object

In this exercise, we will use a `contains` method to check whether the query contains a specified object:

1. Import the `Publisher` model if you have not already imported it:

```
>>> from reviews.models import Publisher
```

- Run a query, as shown next. In this example, we are executing a Q object query, and the result has two publishers:

```
>>> publishers =
    Publisher.objects.filter(Q(name__startswith="New")
    | Q(name__endswith="Publisher"))
>>> publishers
<QuerySet [<Publisher: New Town Publisher>, <Publisher: Katoomba
Publisher>]>
```

- Fetch a new publisher object and verify whether it is a part of the previous query set:

```
>>> new_town_publisher =
    Publisher.objects.get(name='New Town Publisher')
>>> new_town_publisher
<Publisher: New Town Publisher>
>>> publishers.contains(new_town_publisher)
True
```

If the object is present, it returns a Boolean value such as `True` or `False`.

The following is an example when a certain query set does not have the specified object, in which case the result returned is `False`:

```
>>> publishers.contains(Publisher.objects.get(name='Byron
    Bay Press'))
False
```

In this exercise, we learned how a simple method, `contains()`, can be used to check whether an object is part of a query set.

Overall, in this entire section, we learned about and used Django's command-line interactive shell. Using the models created in the previous section for the book review application, we used the Django interactive shell to perform CRUD operations. We also explored various ways to filter and query (Read) from the database. All of these database operations will come in handy while developing any Django application. Using the concepts learned so far in this chapter, in the next section, you will create models for a sample application and perform some of the CRUD operations.

Activity 2.01 – creating models for a project management application

Imagine you are developing a project management application called Juggler. Juggler is an application that can track multiple projects, and each project can have multiple tasks associated with it. The following steps will help you complete this activity:

1. Using the techniques we have learned so far, create a Django project called `juggler`.
2. Create a Django app called `projectp`.
3. Add the app projects to the `juggler/settings.py` file.
4. Create two related model classes called `Project` and `Task` in `projectp/models.py`.
5. Create migration scripts and migrate the models' definitions to the database.
6. Open the Django shell and import the models.
7. Populate the database with an example and write a query displaying the list of tasks associated with a given project.

Note

The full solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Populating the Bookr project's database

Although we know how to create database records for the project, in the next few chapters, we will have to create a lot of records to work with the project. Therefore, we have created a script that can make things easy for us. This script populates the database by reading a **comma-separated values (CSV)** file consisting of many records. Follow the next few steps to populate the project's database:

1. Create the following folder structure inside the project directory:

```
bookr/reviews/management/commands/
```

2. Copy the `loadcsv.py` file from the following location and `WebDevWithDjangoData.csv` into the folder created. This can be found on the GitHub repository for this book at <http://packt.live/3pzbCLM>.

Because `loadcsv.py` is placed inside the `management/commands` folder, it now works like a Django custom management command. You can go through the `loadcsv.py` file and read more about writing Django custom management commands here: <https://docs.djangoproject.com/en/3.0/howto/custom-management-commands/>.

-
3. Now, let's recreate a fresh database. Delete the SQL database file present in the project folder:

```
rm db.sqlite3
```

4. To create a fresh database, execute the Django `migrate` command:

```
python manage.py migrate
```

Now you can see the newly created `db.sqlite3` file under the `bookr` folder.

5. Execute the `loadcsv` custom management command to populate the database:

```
python manage.py loadcsv --csv reviews/management/commands/  
WebDevWithDjangoData.csv
```

6. Using DB Browser for SQLite, verify that all the tables created by the `bookr` project have been populated.

With this, we have successfully populated our application's database, which will come in handy to work with in the upcoming chapters of this book.

Summary

In this chapter, we learned about some basic database concepts and their importance in application development. We used a free database visualization tool, DB Browser for SQLite, to understand what database tables and fields are and how records are stored in a database, and further performed some basic CRUD operations on the database using simple SQL queries.

We then learned how Django provides a valuable abstraction layer called ORM that helps us interact seamlessly with relational databases using simple Python code without having to compose SQL commands. As a part of ORM, we learned about Django models, migrations, and how they help propagate the changes to the Django models in the database.

We shored up our knowledge of databases by learning about database relationships and their key types in relational databases. We also worked with the Django shell, using Python code to perform the same CRUD queries we performed earlier using SQL. Later, we learned how to retrieve our data in a more refined manner using pattern matching and field lookups. As we learned these concepts, we also made considerable progress on our Bookr application. We created models for our `reviews` app and gained all the skills we needed to interact with the data stored inside the app's database. In the next chapter, we will learn how to create Django views, URL routing, and templates.

3

URL Mapping, Views, and Templates

In the previous chapter, we were introduced to databases, and we learned how to store, retrieve, update, and delete records from a database. We also learned how to create Django models and apply database migrations.

However, these database operations alone cannot display an application's data to a user. We need a way to display all the stored information in a meaningful way to the user – for example, displaying all the books present in our Bookr application's database, in a browser, and in a presentable format. This is where Django views, templates, and URL mapping come into play. Views are the part of a Django application that takes in a web request and provides a web response. For example, a web request could be a user trying to view a website by entering the website address, and a web response could be the website's home page loading in the user's browser. Views are one of the most important parts of a Django application, where the application logic is written. This application logic controls interactions with the database, such as creating, reading, updating, or deleting records from the database. It also controls how the data is displayed to the user. This is done with the help of Django HTML templates.

This chapter introduces you to three core concepts of Django: views, templates, and URL mapping. You will start by exploring the two main types of views in Django: function-based views and class-based views. Next, you will learn the basics of Django template language and template inheritance. Using these concepts, you will create a page to display the list of all the books in the Bookr application. You will also create another page to display the details, review comments, and ratings of books.

In this chapter, we will study the following topics:

- Understanding function-based views
- Understanding class-based views
- URL configuration
- Working with Django templates
- Django's template language

By using and learning about the previously mentioned topics, by the end of this chapter, we will have implemented a way to list all the books and display details of any given book in the Bookr application.

Django views can be broadly classified into two types: **function-based views** and **class-based views**. In this chapter, we will learn more about function-based views in Django.

Note

Class-based views, which is a more advanced topic, will be discussed in detail in *Chapter 11, Advanced Templating and Class-Based Views*.

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter03>.

Understanding function-based views

As the name implies, function-based views are implemented as Python functions. To understand how they work, consider the following snippet, which shows a simple view function named `home_page`:

```
from django.http import HttpResponse

def home_page(request):
    message = "<html><h1>Welcome to my Website</h1></html>"
    return HttpResponse(message)
```

The view function defined here, named `home_page`, takes a `request` object as an argument and returns an `HttpResponse` object with a `Welcome to my Website` message. The advantage of using function-based views is that, since they are implemented as simple Python functions, they are easier to learn and also easily readable for other programmers. The major disadvantage of function-based views is that the code cannot be reused and made as concisely as class-based views for generic use cases. The next section is a brief introduction to class-based views.

Understanding class-based views

As the name implies, class-based views are implemented as Python classes. Using the principles of class inheritance, these classes are implemented as subclasses of Django's generic view classes. Unlike function-based views, where all the view logic is expressed explicitly in a function, Django's generic view classes come with various pre-built properties and methods that can provide shortcuts to writing clean, reusable views. This property comes in handy quite often during web development; for example, developers often need to render an HTML page without needing any data inserted from the database,

or any customization specific to the user. In this case, it is possible to simply inherit from Django's `TemplateView`, and specify the path of the HTML file. The following is an example of a class-based view that can display the same message as in the function-based view example:

```
from django.views.generic import TemplateView

class HomePage(TemplateView):
    template_name = 'home_page.html'
```

In the preceding code snippet, `HomePage` is a class-based view inheriting Django's `TemplateView` from the `django.views.generic` module. The `template_name` class attribute defines the template to render when the view is invoked. For the template, we add an HTML file to our `templates` folder with the following content:

```
<html><h1>Welcome to my Website</h1></html>
```

This is a very basic example of class-based views, which will be explored further in *Chapter 11, Advanced Templating and Class-Based Views*. The major advantage of using class-based views is that fewer lines of code need to be used to implement the same functionality as compared to function-based views. Also, by inheriting Django's generic views, we can keep the code concise and avoid the duplication of code. However, a disadvantage of class-based views is that the code is often less readable for someone new to Django, which means that learning about it is usually a longer process as compared to function-based views.

In the next section, we will learn about URL configuration.

URL configuration

Django views cannot work on their own in a web application. When a web request is made to the application, Django's URL configuration takes care of routing the request to the appropriate view function to process the request. A typical URL configuration in the `urls.py` file in Django looks like this:

```
from . import views

urlpatterns = [
    path('url-path/', views.my_view, name='my-view'),
]
```

Here, `urlpatterns` is the variable defining the list of URL paths, and '`url-path/`' defines the path to match.

`views.my_view` is the view function to invoke when there is a URL match, and `name='my-view'` is the name of the view function used to refer to the view. There may be a situation wherein, elsewhere in the application, we want to get the URL of this view. We wouldn't want to hardcode the value, as it would then have to be specified twice in the code base. Instead, we can access the URL by using the name of the view, as follows:

```
from django.urls import reverse

url = reverse('my-view')
```

If needed, we can also use a regular expression in a URL path to match string patterns using `re_path()`:

```
urlpatterns = [
    re_path(r'^url-path/(?P<name>pattern)/$', 
            views.my_view, name='my-view')
]
```

Here, `name` refers to the pattern name, which can be any Python regular expression pattern, and this needs to be matched before calling the defined view function. You can also pass parameters from the URL into the view itself, as shown in this example:

```
urlpatterns = [
    path(r'^url-path/<int:id>/', views.my_view,
         name='my-view')
]
```

In the preceding example, `<int:id>` tells Django to look for URLs that contain an integer at this position in the string, and to assign the value of that integer to the `id` argument. This means that if the user navigates to `/url-path/14/`, the `id=14` keyword argument is passed to the view. This is often useful when a view needs to look up a specific object in the database and return corresponding data. For example, suppose we had a `User` model, and we wanted the view to display the user's name. The view could be written as follows:

```
def my_view(request, id):
    user = User.objects.get(id=id)
    return HttpResponseRedirect(f"This user's name is \
        {user.first_name} {user.last_name}")
```

When the user accesses `/url-path/14/`, the preceding view is called, and the `id=14` argument is passed into the function.

Here is the typical workflow when a URL such as `http://0.0.0.0:8000/url-path/` is invoked using a web browser:

1. An HTTP request would be made to the running application for the URL path. Upon receiving the request, it reaches for the `ROOT_URLCONF` setting present in the `settings.py` file:

```
ROOT_URLCONF = 'project_name.urls'
```

This determines the URL configuration file that will be used first. In this case, it is the URL file present in the project directory, `project_name/urls.py`.

2. Next, Django goes through the list named `urlpatterns`, and once it matches `url-path/` with the path present in the URL, `http://0.0.0.0:8000/url-path/`, it invokes the corresponding view function.

URL configuration is sometimes also referred to as **URL conf** or **URL mapping**, and these terms are often used interchangeably. Overall, we can say that URL configuration is present to route the URL request posted in a browser to an appropriate view function. To understand views and URL mapping better, let's start with a simple exercise.

Exercise 3.01 – implementing a simple function-based view

In this exercise, we will write a very basic function-based view and use the associated URL configuration to display the `Welcome to Bookr!` message in a web browser. We will also tell the user how many books we have in the database. Let's get started with the steps:

1. First, ensure that `ROOT_URLCONF` in `bookr/settings.py` points to the project's URL file by adding the following command:

```
ROOT_URLCONF = 'bookr.urls'
```

2. Open the `bookr/reviews/views.py` file and add the following code snippet:

```
from django.http import HttpResponse
from .models import Book

def welcome_view(request):
    message = f"<html><h1>Welcome to Bookr!</h1>
    <p>{Book.objects.count()} books and
    counting!</p></html>"
    return HttpResponse(message)
```

In the preceding snippet, we first import the `HttpResponse` class from the `django.http` module. Next, we define the `welcome_view` function, which can display the `Welcome to Bookr!` message in a web browser. The `request` object is a function parameter that carries the HTTP `request` object. The next line defines the `message` variable, which contains HTML that displays the header, followed by a line that counts the number of books available in the database.

In the last line, we return an `HttpResponse` object with the string associated with the `message` variable. When the `welcome_view` view function is called, it will display `Welcome to Bookr! 2 books and counting` in the web browser.

3. Now, create the URL mapping to call the newly created view function. Open the project URL file, `bookr/urls.py`, and add the list of `urlpatterns` as follows:

```
from django.contrib import admin
from django.urls import include, path

urlpatterns = [
    path('admin/', admin.site.urls),
    path('', include('reviews.urls'))
]
```

The first line in the list of `urlpatterns`, that is, `path('admin/', admin.site.urls)`, routes to the admin URLs if `admin/` is present in the URL path (for example, `http://0.0.0.0:8000/admin`).

Similarly, consider the second line, `path('', include('reviews.urls'))`. Here, the path mentioned is an empty string, `''`. If the URL does not have any specific path after `http://hostname:port-number/` (for example, `http://0.0.0.0:8000/`), it includes `urlpatterns`, present in `review.urls`.

The `include` function is a shortcut that allows you to combine URL configurations. It is common to keep one URL configuration per application in your Django project. Here, we've created a separate URL configuration for the `reviews` app and have added it to our project-level URL configuration.

4. Since we do not have the `reviews.urls` URL module yet, create a file called `bookr/reviews/urls.py`, and add the following lines of code:

```
from django.contrib import admin
from django.urls import path
from . import views

urlpatterns = [
    path('', views.welcome_view,
          name='welcome_view'),
]
```

Here, we have used an empty string again for the URL path. So, when `http://0.0.0.0:8000/` is invoked, after getting routed from `bookr/urls.py` into `bookr/reviews/urls.py`, this pattern invokes the `welcome_view` view function.

5. After making changes to the two files, we have the necessary URL configuration ready to call the `welcome_view` view. Now, start the Django server with `./manage.py runserver` and type `http://0.0.0.0:8000` or `http://127.0.0.1:8000` in your web browser. You should be able to see the **Welcome to Bookr!** message:



Welcome to Bookr!

18 books and counting!

Figure 3.1 – Displaying “Welcome to Bookr!” and the number of books on the home page

Note

If there is no URL match, Django invokes error handling, such as displaying a **404 Page not found** message or something similar.

In this exercise, we learned how to write a basic view function and do the associated URL mapping. We have created a web page that displays a simple message to the user and reports how many books are currently in our database.

However, astute readers will have noticed that it doesn’t look very nice to have HTML code sitting inside our Python function as in the preceding example. As our views get bigger, this will become even more unsustainable. Therefore, in the next section, we will turn our attention to where our HTML code is supposed to be – inside templates.

Working with Django templates

In *Exercise 3.01 – implementing a simple function-based view*, we saw how to create a view, do the URL mapping, and display a message in the browser. But if you recall, we hardcoded the `Welcome to Bookr!` HTML message in the view function itself and returned an `HttpResponse` object, as follows:

```
message = f"<html><h1>Welcome to Bookr!</h1> "
"<p>{Book.objects.count()} books and counting!</p></html>"
return HttpResponse(message)
```

Hardcoding HTML inside Python modules is not a good practice, because as the amount of content to be rendered on a web page increases, so does the amount of HTML code we need to write for it. Having a lot of HTML code among Python code can make the code hard to read and maintain in the long run.

For this reason, Django templates provide us with a better way to write and manage HTML templates. Django's templates not only work with static HTML content but also with dynamic HTML templates.

Django's template configuration is done in the `TEMPLATES` variable present in the `settings.py` file. This is how the default configuration looks:

```
TEMPLATES = [
{
    'BACKEND': 'django.template.backends.django.DjangoTemplates',
    'DIRS': [],
    'APP_DIRS': True,
    'OPTIONS': {
        'context_processors': [
            'django.template.context_processors.debug',
            'django.template.context_processors.request',
            'django.contrib.auth.context_processors.auth',
            'django.contrib.messages.context_processors.messages',
        ],
    },
},
```

Let's go through each keyword present in the preceding snippet:

- `'BACKEND' : 'django.template.backends.django.DjangoTemplates'`: This refers to the template engine to be used. A template engine is an API used by Django to work with HTML templates. Django is built with Jinja2 and the `DjangoTemplates` engine. The default configuration is the `DjangoTemplates` engine and Django template language. However, this can be changed for a different one if required, such as Jinja2 or any other third-party template engine. For our Bookr application though, we will leave this configuration as it is.
- `'DIRS' : []`: This refers to the list of directories where Django searches for the templates in the given order.
- `'APP_DIRS' : True`: This tells the Django template engine whether it should look for templates in the installed apps defined under `INSTALLED_APPS` in the `settings.py` file. The default option for this is `True`.
- `'OPTIONS'`: This is a dictionary containing template engine-specific settings. Inside this dictionary, there is a default list of context processors, which helps the Python code to interact with templates to create and render dynamic HTML templates.

The current default settings are mostly fine for our purposes. However, in the next exercise, we will create a new directory for our templates, and we will need to specify the location of this folder. For example, if we have a directory called `my_templates`, we need to specify its location by adding it to the `TEMPLATES` settings as follows:

```
TEMPLATES = [
{
    'BACKEND' :
        'django.template.backends.django.DjangoTemplates',
    'DIRS': [os.path.join(BASE_DIR, 'my_templates')],
    'APP_DIRS': True,
    'OPTIONS': {
        'context_processors': [
            'django.template.context_processors.debug',
            'django.template.context_processors.request',
            'django.contrib.auth.context_processors.auth',
            'django.contrib.messages.context_processors.messages',
        ],
    },
},
```

BASE_DIR is the directory path to the project folder. This is defined in the `settings.py` file. The `os.path.join()` method joins the project directory with the `templates` directory, returning the full path for the `templates` directory. In the following exercise, we shall use a template to display a message to the user.

Exercise 3.02 – using templates to display a greeting message

In this exercise, we will create our first Django template, and, just as we did in the previous exercise, we will display the `Welcome to Bookr!` message using the templates. Let's get started with the steps:

1. Create a directory called `templates` in the `bookr` project directory and inside it, create a file called `base.html`. The directory structure should look like the following figure:

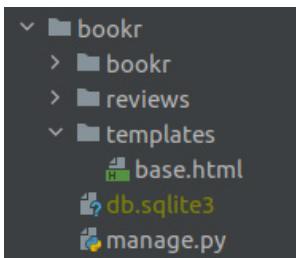


Figure 3.2 – Directory structure for bookr

Note

When the default configuration is used (that is, when `DIRS` is an empty list), Django searches for templates present only in the app folders' `templates` directory (the `reviews/templates` folder in the case of a book review application). Since we have included the new template directory in the main project directory, Django's template engine will not be able to find the directory unless the directory is included in the '`DIRS`' list.

2. Add the folder to the `TEMPLATES` settings:

```
TEMPLATES = [
{
    'BACKEND': 'django.template.backends.django.DjangoTemplates',
    'DIRS': [os.path.join(BASE_DIR, 'my_templates')],
    'APP_DIRS': True,
    'OPTIONS': {
        'context_processors': [
            'django.template.context_processors.debug',
```

```
        'django.template.context_processors.request',
        'django.contrib.auth.context_processors.auth',
        'django.contrib.messages.context_processors
            .messages',
    ],
},
},
]
```

3. Add the following lines of code into the `base.html` file:

```
<!doctype html>
<html lang="en">
<head>
    <meta charset="utf-8">
    <title>Home Page</title>
</head>
<body>
    <h1>Welcome to Bookr!</h1>
</body>
</html>
```

This is simple HTML that displays the `Welcome to Bookr!` message in the header.

4. Modify the code inside `bookr/reviews/views.py` so that it looks as follows:

```
from django.shortcuts import render

def welcome_view (request):
    return render(request, 'base.html')
```

Since we have already configured the `'templates'` directory in the `TEMPLATES` configuration, `base.html` is available for use for the template engine. The code renders the `base.html` file using the imported `render` method from the `django.shortcuts` module.

5. Save the files, run `./manage.py runserver`, and open the `http://0.0.0.0:8000/` or `http://127.0.0.1:8000/` URL to check the newly added template loading in the browser:



Figure 3.3 – Displaying “Welcome to Bookr!” on the home page

In this exercise, we created an HTML template and used Django templates and views to return the `Welcome to Bookr!` message.

Overall, we can say that Django templates help us to work with HTML templates and present information to the user in a desired format. Next, we will learn about the Django template language, which can be used to render the application’s data along with HTML templates.

Django’s template language

Django templates not only return static HTML templates but can also add dynamic application data while generating the templates. Along with data, we can also include some programmatic elements in the templates. All of these put together form the basics of **Django’s template language**. The following few sections look at some of the basic parts of the Django template language, such as template variables, template tags, comments, and filters.

Template variables

A template variable is represented between two curly braces, as shown here:

```
{ { variable } }
```

When this is present in the template, the value carried by the variables will be replaced in the template. Template variables help add the application’s data into the templates:

```
template_variable = "I am a template variable."  
  
<body>  
    { { template_variable } }  
</body>
```

Template tags

A tag is similar to a programmatic control flow, such as an `if` condition or a `for` loop. A tag is represented between two curly braces and percentage signs, as shown. Here is an example of a `for` loop iterating over a list using template tags:

```
{% for element in element_list %}  
  
{% endfor %}
```

Unlike Python programming, we also add the end of the control flow by adding the `end` tag, such as `{% endfor %}`. This can be used along with template variables to display the elements in the list, as shown here:

```
<ul>  
    {% for element in element_list %}  
        <li>{{ element.title }}</li>  
    {% endfor %}  
</ul>
```

Comments

Comments in the Django template language can be written as shown here; anything in between `{% comment %}` and `{% endcomment %}` will be commented out:

```
{% comment %}  
    <p>This text has been commented out</p>  
{% endcomment %}
```

Filters

Filters can be used to modify a variable to represent it in a different format. The syntax for a filter is a variable separated from the filter name using a pipe (|) symbol:

```
{{ variable|filter }}
```

Here are some examples of built-in filters:

- `{{ variable|lower }}`: This converts the variable string into lowercase
- `{{ variable|title }}`: This converts the first letter of every word into uppercase

Let's use the concepts we have learned to develop the book review application.

Exercise 3.03 – displaying a list of books and reviews

In this exercise, we will create a web page that can display a list of all books, their ratings, and the number of reviews present in the book review application. For this, we will be using some features of the Django template language, such as variables and template tags, to pass the book review application data into the templates to display meaningful data on the web page:

1. Create a file called `utils.py` under `bookr/reviews/utils.py` and add the following code:

```
def average_rating(rating_list):  
    if not rating_list:  
        return 0  
  
    return round(sum(rating_list) / len(rating_list))
```

This is a helper method that will be used to calculate the average rating of a book.

2. Remove all the code present inside `bookr/reviews/views.py` and add the following code to it:

```
from django.shortcuts import render, get_object_or_404  
  
from .models import Book  
from .utils import average_rating  
  
  
def index(request):  
    return render(request, "base.html")  
  
  
def book_search(request):  
    search_text = request.GET.get("search", "")  
    return render(request,  
                 "reviews/search-results.html",  
                 {"search_text": search_text})
```

This is a view to display the list of books for the book review application. The first three lines import Django modules, model classes, and the helper method we just added.

Here, `books_list` is the view method. In this method, we start by querying the list of all books. Next, for every book, we calculate the average rating and the number of reviews posted. All this information for each book is appended to a list called `book_list` as a list of dictionaries. This list is then added to a dictionary named `context` and is passed to the `render` function.

The `render` function takes three parameters: the first one is the `request` object that was passed into the view, the second is the `books_list.html` HTML template, which will display the list of books, and the third is `context`, which we pass to the template.

Since we have passed `book_list` as a part of the context, the template will be using this to render the list of books using template tags and template variables.

3. Create the `book_list.html` file in the `bookr/reviews/templates/reviews/books_list.html` path and add the following HTML code to the file:

```
<!doctype html>
<html lang="en">
<head>
    <meta charset="utf-8">
    <title>Bookr</title>
</head>
<body>
    <h1>Book Review application</h1>
    <hr>
```

You can find the complete code at https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter03/Exercise3.03/bookr/reviews/templates/reviews/book_list.html.

This is a simple HTML template with template tags and variables iterating over `book_list` to display the list of books.

4. In `bookr/reviews/urls.py`, add the following URL pattern to invoke the `books_list` view:

```
from django.urls import path
from . import views

urlpatterns = [
    path('books/', views.book_list,
         name='book_list'),
]
```

This does the URL mapping for the `books_list` view function.

5. Save all the modified files and wait for the Django service to restart. Open `http://0.0.0.0:8000/books/` in the browser, and you should see something similar to the following figure:



Book Review application

- Title: Advanced Deep Learning with Keras
Publisher: Packt Publishing
Publication Date: Oct. 31, 2018
Rating: 4
Number of reviews: 2
- Title: Hands-On Machine Learning for Algorithmic Trading
Publisher: Packt Publishing
Publication Date: Dec. 31, 2018
Rating: None
Number of reviews: 0
Provide a rating and write the first review for this book.
- Title: Architects of Intelligence
Publisher: Packt Publishing
Publication Date: Nov. 23, 2018
Rating: None
Number of reviews: 0
Provide a rating and write the first review for this book.
- Title: Deep Reinforcement Learning Hands-On
Publisher: Packt Publishing
Publication Date: June 20, 2018
Rating: None
Number of reviews: 0
Provide a rating and write the first review for this book.
- Title: Natural Language Processing with TensorFlow
Publisher: Packt Publishing
Publication Date: May 30, 2018
Rating: None
Number of reviews: 0
Provide a rating and write the first review for this book.

Figure 3.4 – List of books present in the book review application

In this exercise, we created a view function, created templates, and also did the URL mapping, which can display a list of all books present in the application. Although we were able to display a list of books using a single template, next, let's explore a bit about how to work with multiple templates in an application that has common or similar code.

Template inheritance

As we build the project, the number of templates will increase. It is highly probable that when we design the application, some of the pages will look similar and have common HTML code for certain features. Using template inheritance, other HTML files can inherit the common HTML code. This is similar to class inheritance in Python, where the parent class has all the common code, and the child class has those extras that are unique to the child's requirements.

For example, let's consider the following to be a parent template that is named `base.html`:

```
<!doctype html>
<html lang="en">
<head>
    <meta charset="utf-8">
    <title>Hello World</title>
</head>
<body>
    <h1>Hello World using Django templates!</h1>
    {%
        block content %
    %}
    {%
        endblock %
    </body>
</html>
```

The following is an example of a child template:

```
{% extends 'base.html' %}
{% block content %}
<h1>How are you doing?</h1>
{% endblock %}
```

In the preceding snippet, the `{% extends 'base.html' %}` line extends the template from `base.html`, which is the parent template. After extending from the parent template, any HTML code in between the block content will be displayed along with the parent template. Once the child template has been rendered, here is how it looks in the browser:



Figure 3.5 – Greeting message after extending the `base.html` template

In the next section, we will do some template styling using Bootstrap.

Template styling with Bootstrap

We have seen how to display all the books using views, templates, and URL mapping. Although we were able to display all the information in the browser, it would be even better if we could add some styling and make the web page look better. For this, we can add a few elements of **Bootstrap**. Bootstrap is an open source **Cascading Style Sheets (CSS)** framework that is particularly good for designing responsive pages that work across desktop and mobile browsers.

Using Bootstrap is simple. First, you need to add Bootstrap CSS to your HTML. You can experiment yourself by creating a new file called `example.html`. Populate it with the following code and open it in a browser:

```
<!doctype html>
<html lang="en">
  <head>
    <!-- Required meta tags -->
    <meta charset="utf-8">
    <meta name="viewport" content="width=device-width,
    initial-scale=1, shrink-to-fit=no">

    <!-- Bootstrap CSS -->
    <link rel="stylesheet" href=
      "https://stackpath.bootstrapcdn.com/bootstrap/4.4.1/
      css/bootstrap.min.css" integrity="sha384-
      Vko08x4CGsO3+Hhxv8T/Q5PaXtkKtu6ug5TOeNV6gBiFeWPGFN9M
      uhOf23Q9Ifjh" crossorigin="anonymous">

  </head>
  <body>
    Content goes here
  </body>
</html>
```

The Bootstrap CSS link in the preceding code adds the Bootstrap CSS library to your page. This means that certain HTML element types and classes will inherit their styles from Bootstrap. For example, if you add the `btn-primary` class to the class of a button, the button will be rendered in blue with white text.

Try adding the following between <body> and </body>:

```
<h1>Welcome to my Site</h1>
<button type="button" class="btn btn-primary">
    Checkout my Blog!</button>
```

You will see that the title and button are both styled nicely, using Bootstrap's default styles:

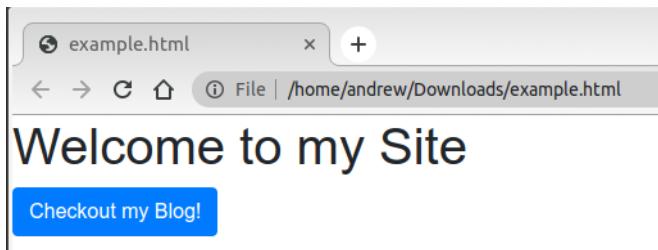


Figure 3.6 – Display after applying Bootstrap

This is because in the Bootstrap CSS code, it specifies the color of the `btn-primary` class with the following code:

```
.btn-primary {
    color: #fff;
    background-color: #007bff;
    border-color: #007bff
}
```

You can see that using third-party CSS libraries such as Bootstrap allows you to quickly create nicely styled components without needing to write too much CSS. Next, we shall use template inheritance from a base template and also add some styling elements such as a navigation bar.

Note

We recommend that you explore Bootstrap further with its tutorial here:
<https://getbootstrap.com/docs/4.4/getting-started/introduction/>.

Exercise 3.04 – adding template inheritance and a Bootstrap navigation bar

In this exercise, we will use template inheritance to inherit the template elements from a base template and reuse them in the `book_list` template to display the list of books. We will also use certain elements of Bootstrap in the base HTML file to add a navigation bar to the top of our page. The bootstrap code for `base.html` was taken from <https://getbootstrap.com/docs/4.4/getting-started/introduction/> and <https://getbootstrap.com/docs/4.4/components/navbar/>. Let's get started with the steps:

1. Open the `base.html` file from the `bookr/templates/base.html` location. Remove any existing code and replace it with the following code:

```
<!doctype html>
{%
  load static %
}
<html lang="en">
<head>
  <!-- Required meta tags -->
  <meta charset="utf-8">
  <meta name="viewport" content="width=device-width,
    initial-scale=1, shrink-to-fit=no">
  <!-- Bootstrap CSS -->
```

You can view the entire code for this file at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter03/Exercise3.04/bookr/reviews/templates/reviews/base.html>.

This is a `base.html` file with all the Bootstrap elements for styling and the navigation bar.

2. Next, open the template at `bookr/reviews/templates/reviews/books_list.html`, remove all the existing code, and replace it with the following code:

```
{% extends 'base.html' %}

{% block content %}
<ul class="list-group">
  {% for item in book_list %}
  <li class="list-group-item">
    <span class="text-info">Title: </span> <span>{ {
      item.book.title } }</span>
    <br>
    <span class="text-info">Publisher:
      </span><span>{ {
        item.book.publisher } }</span>
```

You can view the complete code for this file at https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter03/Exercise3.04/bookr/reviews/templates/reviews/book_list.html.

This template has been configured to inherit the `base.html` file, and it has also been added with a few styling elements to display the list of books. The part of the template that helps in inheriting the `base.html` file is as follows:

```
{% extends 'base.html' %}

{% block content %}
{% endblock %}
```

3. After adding the two new templates, open either of the URLs (`http://0.0.0.0:8000/books/` or `http://127.0.0.1:8000/books/`) in your web browser to see the book list page, which should now look neatly formatted:

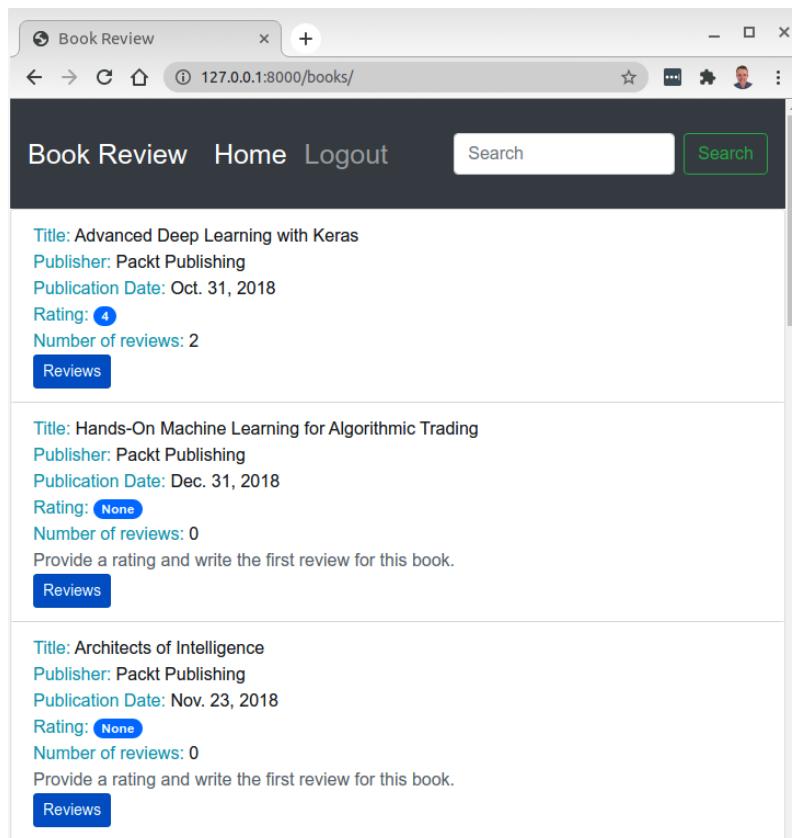


Figure 3.7 – Neatly formatted book list page

In this exercise, we added some styling to the application using Bootstrap, and we also used template inheritance while we displayed the list of books from the book review application. So far, we have worked extensively on displaying all the books present in the application. In the next activity, you will display the details and reviews of an individual book.

Activity 3.01 – implementing the book details view

In this activity, you will implement a new view, template, and URL mapping to display these details of a book: the title, publisher, publication date, and overall rating. In addition to these details, the page should also display all the review comments, specifying the name of the commenter and the dates on which the comments were written and (if applicable) modified. The following steps will help you complete this activity:

1. Create a book details endpoint that extends the base template.
2. Create a book details view that takes a specific book's primary key as the argument and returns an HTML page listing the book's details and any associated reviews.
3. Do the required URL mapping in `urls.py`. The book details view URL should be `http://0.0.0.0:8000/books/1/` (where 1 will represent the ID of the book being accessed). You can use the `get_object_or_404` method to retrieve the book with the given primary key.

Note

The `get_object_or_404` function is a useful shortcut for retrieving an instance based on its primary key. You could also do this using the `.get()` method described in *Chapter 2, Models and Migrations*: `Book.objects.get(pk=pk)`. However, `get_object_or_404` has the added advantage of returning an `HTTP 404 Not Found` response if the object does not exist. If we simply use `get()` and someone attempts to access an object that does not exist, our Python code will hit an exception and return an `HTTP 500 Server Error` response. This is undesirable because it looks as though our server has failed to handle the request correctly.

4. At the end of the activity, you should be able to click the **Reviews** button on the book list page and get a detailed view of a book. The detailed view should have all the details displayed, as shown in the following screenshot:



Book Details

Title: Advanced Deep Learning with Keras
Publisher: Packt Publishing
Publication Date: Oct. 31, 2018
Overall Rating: 4

Review Comments

Review comment: A must read for all
Created on: Nov. 22, 2020, 10:47 p.m.
Modified on: Jan. 4, 2020, 4:31 p.m.
Rating: 5
Creator: peterjones@test.com

Review comment: An ok read
Created on: Nov. 22, 2020, 10:47 p.m.
Modified on: Jan. 4, 2020, 4:31 p.m.
Rating: 3
Creator: marksandler@test.com

© 2020 Copyright: [Packt Publications](#)
Website by: Packt Publications

Contact information: emailexample@email.com

Figure 3.8 – Page displaying the book details

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

In the templates section, we learned how to use Django templates to present static information to the user, template language to display the application's dynamic data, and Bootstrap to add styling to the information and make it presentable to users.

Summary

This chapter covered the core infrastructure required to handle an HTTP request to our website. The request is first mapped via URL patterns to an appropriate view. Parameters from the URL are also passed into the view to specify the object displayed on the page. The view is responsible for compiling any necessary information to display on the website, and then passes this dictionary through to a template, which renders the information as HTML code that can be returned as a response to the user. We covered both class- and function-based views and learned about the Django template language and template inheritance. We created two new pages for the book review application, one displaying all the books present and the other being the book details view page. In the next chapter, we will learn about Django admins and superusers, registering models, and performing **create, read, update, and delete (CRUD)** operations using the admin site.

4

An Introduction to Django Admin

This chapter introduces you to the basic functionality of the Django admin app. You will start by creating superuser accounts for the Bookr app, before moving on to executing **Create, Read, Update, and Delete (CRUD)** operations with the admin app. You will learn how to integrate your Django app with the admin app, and you'll also look at the behavior of `ForeignKeys` in the admin app. At the end of this chapter, you will see how you can customize the admin app according to a unique set of preferences by sub-classing the `AdminSite` and `ModelAdmin` classes, making its interface more intuitive and user-friendly.

In this chapter, we will cover the following topics:

- Creating a superuser account
- CRUD operations using the Django admin app
- Managing Django users and groups
- Registering models with the admin app
- Customizing the admin interface

By the end of this chapter, you will be able to do the following:

- Create an admin account
- Create user accounts and groups, and assign permissions
- Register your models with the admin app
- Customize global properties of the admin app
- Customize the functionality of change list pages and forms for administering your models

When developing an app, there is often a need to populate it with data and then alter that data. We have already seen in *Chapter 2, Models and Migrations*, how this can be done in the command line using the Python `manage.py` shell. In *Chapter 3, URL Mapping, Views, and Templates*, we learned how to develop a web form interface for our model using Django's views and templates. But neither of these approaches is ideal for administering the data from the classes in `reviews/models.py`. Using the shell to manage data is too technical for non-programmers, and building individual web pages would be a laborious process as it would see us repeating the same view logic and very similar template features for each table in the model. Fortunately, a solution to this problem was devised in the early days of Django when it was still being developed.

The Django admin site is actually written as a Django app. It offers an intuitively rendered web interface to give administrative access to the model data. The admin interface is designed to be used by the website administrators. It is not intended to be used by non-privileged users who interact with the site. In our case of a book review system, the general population of book reviewers will never encounter the admin app. They will see the app pages, like those we built with views and templates in *Chapter 3, URL Mapping, Views, and Templates*, and will write their reviews on the pages.

Also, while developers put a lot of effort into creating a simple and inviting web interface for general users, the admin interface, aimed at administrative users, maintains a utilitarian feel that typically displays the intricacies of the model. It may have escaped your attention, but you already have an admin app in your Bookr project. Look at the list of installed apps in `bookr/settings.py`:

```
INSTALLED_APPS = [
    'django.contrib.admin',
    ...
]
```

Now, look at the URL patterns in `bookr/urls.py`:

```
urlpatterns = [
    path('admin/', admin.site.urls),
    ...
]
```

If we put this path into our browser, we can see that the link to the admin app on the development server is `http://127.0.0.1:8000/admin/`. Before we make use of it, though, we need to create a superuser through the command line.

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter04>.

Creating a superuser account

Our Bookr application has just found a new user. Her name is Alice, and she wants to start adding her reviews right away. Bob, who is already using Bookr, has just informed us that his profile seems incomplete and needs to be updated. David no longer wants to use the application and wants his account to be deleted. For security reasons, we do not want just any user performing these tasks for us. That's why we need to create a **superuser** with elevated privileges. Let's start by doing just that.

In Django's authorization model, a superuser is one with the `Staff` attribute set. We will examine this later in the chapter and learn more about this authorization model in *Chapter 9, Sessions and Authentication*.

We can create a superuser by using the `manage.py` script that we have explored in earlier chapters. Again, we need to be in the project directory when we enter the command. We will use the `createsuperuser` subcommand by entering the following command in the command line (if you're using Windows you will need to type either `python` or `py` instead of `python3`):

```
python3 manage.py createsuperuser
```

Note

In this chapter, we will use email addresses that fall under the *example.com* domain. This follows an established convention to use this reserved domain for testing and documentation. You could use your own email addresses if you prefer.

Let's go ahead and create our superuser.

Exercise 4.01 – creating a superuser account

In this exercise, you will create a superuser account that lets the user log in to the admin site. This functionality will also be used in the upcoming exercises to implement changes that only a superuser can. The following steps will help you complete this exercise:

1. Enter the following command to create a superuser:

```
python manage.py createsuperuser
```

On executing the preceding command, you will be prompted to create a superuser. This command will prompt you for a superuser name, an optional email address, and a password.

2. Add the username and email for the superuser as follows. Here, we enter `bookradmin` (highlighted) at the prompt and press the *Enter* key. Similarly, at the next prompt, which asks you to enter your email address, you can add `bookradmin@example.com` (highlighted). Press the *Enter* key to continue:

```
Username (leave blank to use 'django'): bookradmin
Email address: bookradmin@example.com
Password:
```

This will assign the name `bookradmin` to the superuser. Note that you won't see any output immediately.

3. The next prompt in the shell is for your password. Add a strong password and press the *Enter* key to confirm it once again:

```
Password:
Password (again):
```

You should see the following message on your screen:

```
Superuser created successfully.
```

Note that the password is validated as per the following criteria:

- It cannot be among the 20,000 most common passwords
- It should have a minimum of eight characters
- It cannot be only numerical characters
- It cannot be derived from the username, first name, last name, or email address of the user

With this, you have created a superuser named `bookradmin` who can log in to the admin app. The following screenshot shows how this looks in the shell:

```
> python manage.py createsuperuser
Username (leave blank to use 'django'): bookradmin
Email address: bookradmin@example.com
Password:
Password (again):
Superuser created successfully.
> █
```

Figure 4.1 – Creating a superuser

4. Visit the admin app at `http://127.0.0.1:8000/admin` and log in with the superuser account that you have created:



The image shows the Django administration login interface. It features a dark blue header bar with the text "Django administration". Below this is a light blue form area. The form has two text input fields: the first is labeled "Username:" and contains the text "bookadmin"; the second is labeled "Password:" and contains a series of masked dots. At the bottom of the form is a blue "Log in" button.

Figure 4.2 – The Django administration login form

In this exercise, you created a superuser account that we will be using for the rest of this chapter to assign or remove privileges as needed.

With a superuser account created, we are now able to log into the Django admin app and learn about the CRUD facilities that it provides.

CRUD operations using the Django admin app

Let's get back to the requests we got from Bob, Alice, and David. As a superuser, your tasks will involve creating, updating, retrieving, and deleting various user accounts, reviews, and title names. This set of activities is collectively termed CRUD. CRUD operations are central to the behavior of the admin app. It turns out that the admin app is already aware of the models from another Django app, `Authentication and Authorization` – referenced in `INSTALLED_APPS` as `'django.contrib.auth'`. When logging into `http://127.0.0.1:8000/admin/`, we are presented with the models from the authorization application, as shown in *Figure 4.3*:



The image shows the Django administration interface. At the top, there is a dark blue header bar with the text "Django administration", the name "BOOKADMIN" (in green), and links for "VIEW SITE", "CHANGE PASSWORD", and "LOG OUT". Below this is a light blue "Site administration" bar. The main content area has a dark blue header titled "AUTHENTICATION AND AUTHORIZATION". Under this header, there are two entries: "Groups" and "Users", each with a green "Add" button and a yellow "Change" button. To the right of the main content area, there is a sidebar with two sections: "Recent actions" and "My actions". The "Recent actions" section is empty, and the "My actions" section also says "None available".

Figure 4.3 – The Django administration window

When the admin app is initialized, it calls its `autodiscover()` method to detect whether any other installed apps contain an admin module. If so, these admin models are imported. In our case, it has discovered `'django.contrib.auth.admin'`. Now that the modules are imported, and our superuser account is ready, let's start by working on the requests from Bob, Alice, and David.

Create

Before Alice starts writing her reviews, we need to create an account for her through the admin app. Once that is done, we can look at the administration access levels we can assign to her. To **Create** a user, we need to click the **+ Add** option next to **Users** (refer to *Figure 4.3*), and fill out the form, as shown in *Figure 4.4*.

Note

We don't want any random user to have access to the Bookr users' accounts. Therefore, it is imperative that we choose strong, secure passwords.

The screenshot shows the 'Add user' page in the Django Admin. The URL in the browser is `Home > Authentication and Authorization > Users > Add user`. On the left, there's a sidebar with a search bar and two main categories: 'AUTHENTICATION AND AUTHORIZATION' (Groups and Users) and 'CONTENT' (Books and Authors). The 'Users' category is highlighted with a yellow background and has a green '+ Add' button. The main content area is titled 'Add user' and contains instructions: 'First, enter a username and password. Then, you'll be able to edit more user options.' Below this, there are two input fields: 'Username' (containing 'alice') and 'Password' (containing '*****'). To the right of the 'Password' field are four validation messages: 'Your password can't be too similar to your other personal information.', 'Your password must contain at least 8 characters.', 'Your password can't be a commonly used password.', and 'Your password can't be entirely numeric.' Below the password fields is a 'Password confirmation' field with '*****' and a note: 'Enter the same password as before, for verification.' At the bottom of the form are three buttons: 'Save and add another', 'Save and continue editing', and a dark blue 'SAVE' button.

Figure 4.4 – The Add user page

There are three buttons at the bottom of the form:

- **Save and add another:** This creates the user and renders the same **Add user** page again, with blank fields.
- **Save and continue editing:** This creates the user and loads the **Change user** page. The **Change user** page lets you add additional information that wasn't present on the **Add user** page,

such as **First name**, **Last name**, and more (see *Figure 4.5*). Note that **Password** does not have an editable field on the form. Instead, it shows information about the hashing technique that the password is stored with, in addition to a link to a separate *change password* form:

 The user "alice" was added successfully. You may edit it again below.

Change user

alice

Username: HISTORY

Required: 150 characters or fewer. Letters, digits and @/./+/-/_ only.

Password: **algorithm:** pbkdf2_sha256 **iterations:** 320000 **salt:** cF2Ldy***** **hash:** a84crt*****
Raw passwords are not stored, so there is no way to see this user's password, but you can change the password using [this form](#).

Personal info

First name:

Last name:

Email address:

Permissions

Active
Designates whether this user should be treated as active. Unselect this instead of deleting accounts.

Staff status
Designates whether the user can log into this admin site.

Superuser status
Designates that this user has all permissions without explicitly assigning them.

Figure 4.5 – The Change user page presented after clicking Save and continuing editing

- **SAVE:** This creates the user and lets the user navigate to the **Select user to change** list page, as depicted in *Figure 4.6*.

Activity 4.01 – creating a user account

We have now seen how to create user accounts through the Django admin interface.

It is time to put your knowledge into practice to create an account for a new user David Green, using the email address `david.green@example.com`.

As was the case with Alice White's account, just set this one as an active user without ticking the **Staff Status** or **Superuser Status** boxes.

Note

The solution for this activity can be found <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Now that you have created a user through the Django admin interface, we can examine the other three CRUD operations – retrieve, update, and delete.

Retrieve

The administrative tasks need to be divided among some users, and for this, the admin (the person with the superuser account) would like to view those users whose email addresses end with `n@example.com` and assign the tasks to these users. This is where the **retrieve** functionality can come in handy. After we have clicked the **SAVE** button on the **Add user** page (refer to *Figure 4.4*), we are taken to the **Select user to change** list page (as shown in *Figure 4.6*), which carries out the retrieve operation. Note that the **Create** form is also reachable by clicking on the **ADD USER +** button on the **Select user to change** list page. So, after we have added a few more users, the change list will look something like this:

USERNAME	EMAIL ADDRESS	FIRST NAME	LAST NAME	STAFF STATUS
alice	alice.white@example.com	Alice	White	✗
bob	bob.black@example.com	Bob	Black	✗
bookadmin	bookadmin@example.com			✓
carol	carol.brown@example.com	Carol	Brown	✗
david	david.green@example.com	David	Green	✗

5 users

Figure 4.6 – The Select user to change page

At the top of the form is a **Search** bar that searches the content such as username, email address, and first and last names of users. On the right-hand side is a **FILTER** panel that narrows down the selection based on the values of **By staff status**, **By superuser status**, and **By active**. In *Figure 4.7*, you can see what happens when we search the string `n@example.com`. This will return only the names of the users whose email addresses consist of a username ending with an `n` and a domain starting with `example.com`. We can only see three users with email addresses matching this requirement – `bookradmin@example.com`, `carol.brown@example.com`, and `david.green@example.com`:

Select user to change ADD USER 

Search

3 results (5 total)

Action:	----	Go	0 of 3 selected		
	USERNAME	EMAIL ADDRESS	FIRST NAME	LAST NAME	STAF
<input type="checkbox"/>	bookradmin	bookradmin@example.com			✓
<input type="checkbox"/>	carol	carol.brown@example.com	Carol	Brown	✗
<input type="checkbox"/>	david	david.green@example.com	David	Green	✗

3 users

FILTER

By staff status

All Yes No

By superuser status

All Yes No

By active

All Yes No

Figure 4.7 – Searching for users by a portion of their email address

Retrieve functionality, such as the list page and search bar, allows us to obtain records for update and delete operations.

Update

Remember that Bob wanted his profile to be updated. Let's **update** Bob's unfinished profile by clicking the **bob** username link in the **Select user to change** list:

Select user to change ADD USER +

Search

Action: Go 0 of 5 selected

<input type="checkbox"/>	USERNAME	EMAIL ADDRESS	FIRST NAME	LAST NAME	STAFF
<input type="checkbox"/>	alice	alice.white@example.com	Alice	White	✖
<input type="checkbox"/>	bob	bob.black@example.com	Bob	Black	✖
<input type="checkbox"/>	bookradmin	bookradmin@example.com			✓
<input type="checkbox"/>	carol	carol.brown@example.com	Carol	Brown	✖
<input type="checkbox"/>	david	david.green@example.com	David	Green	✖

5 users

FILTER

By staff status

All

Yes

No

By superuser status

All

Yes

No

By active

All

Yes

No

Figure 4.8 – Selecting bob from the Select user to change list

This will take us back to the **Change user** form where the values for **First name**, **Last name**, and **Email address** can be entered:

The screenshot shows the Django admin interface for a user named 'bob'. The top navigation bar includes 'WELCOME, BOOKADMIN. VIEW SITE / CHANGE PASSWORD / LOG OUT'. Below the navigation, the breadcrumb navigation shows 'Home > Authentication and Authorization > Users > bob'. The main title is 'Change user'. On the right, there is a 'HISTORY' button. The 'Username' field is set to 'bob'. A note below it states: 'Required. 150 characters or fewer. Letters, digits and @/./+/-/_ only.' The 'Password' field displays a hashed password: 'algorithm: pbkdf2_sha256 iterations: 180000 salt: uC2Obk***** hash: 4vvITw*****'. A note below it says: 'Raw passwords are not stored, so there is no way to see this user's password, but you can change the password using this form.' Below this, a 'Personal info' section is expanded, showing fields for 'First name' (Bob), 'Last name' (Black), and 'Email address' (bob.black@example.com).

Figure 4.9 – Adding personal info

As can be seen from *Figure 4.9*, we are adding personal information about Bob here – his name, surname, and email address, specifically.

Another type of update operation is “soft deleting.” The **Active** Boolean property allows us to deactivate a user rather than deleting the entire record and losing all the data that has dependencies on the account. This practice of using a Boolean flag to denote a record as inactive or removed (and subsequently filtering these flagged records out of queries) is referred to as a **soft delete**. Similarly, we can upgrade the user to **Staff status** or **Superuser status** by ticking the respective checkboxes for those:

Permissions

Active
Designates whether this user should be treated as active. Unselect this instead of deleting accounts.

Staff status
Designates whether the user can log into this admin site.

Superuser status
Designates that this user has all permissions without explicitly assigning them.

Figure 4.10 – The Active, Staff status, and Superuser status Booleans

Delete

David no longer wants to use the Bookr application and has requested that we delete his account. The auth admin caters to this too. Select a user or user records on the **Select user to change** list page and choose the **Delete selected users** option from the **Action** drop-down menu. Then, hit the **Go** button (Figure 4.11):

Select user to change ADD USER +

Action: **Delete selected users** Go 1 of 5 selected

<input type="checkbox"/> USERNAME	EMAIL ADDRESS	FIRST NAME	LAST NAME	STAFF
<input type="checkbox"/> alice	alice.white@example.com	Alice	White	✖
<input type="checkbox"/> bob	bob.black@example.com	Bob	Black	✖
<input type="checkbox"/> bookadmin	bookadmin@example.com			✓
<input type="checkbox"/> carol	carol.brown@example.com	Carol	Brown	✖
<input checked="" type="checkbox"/> david	david.green@example.com	David	Green	✖

5 users

FILTER

By staff status

All
 Yes
 No

By superuser status

All
 Yes
 No

By active

All
 Yes
 No

Figure 4.11 – Deleting from the Select user to change list page

You will be presented with a confirmation screen and taken back to the **Select user to change** list once you have deleted the object:

Are you sure?

Are you sure you want to delete the selected user? All of the following objects and their related items will be deleted:

Summary

- Users: 1

Objects

- User: [david](#)

[Yes, I'm sure](#)

[No, take me back](#)

Figure 4.12 – User deletion confirmation

You will see the following message once the user is deleted:

 Successfully deleted 1 user.

Figure 4.13 – User deletion notification

After that confirmation, you will find that David's account no longer exists.

So far, we have learned how we can add a new user, get the details of another user, make changes to the data for a user, and delete a user. These skills helped us cater to Alice, Bob, and David's requests. As the number of users of our app grows, managing requests from hundreds of users will eventually become quite difficult. One way around this problem would be to delegate some of the administrative responsibilities to a selected set of users. We'll learn how to do that in the section that follows.

Managing Django users and groups

Django's authentication model consists of users, groups, and permissions. Users can belong to many groups, which is a way of categorizing users. It also streamlines the implementation of permissions by allowing permissions to be assigned to collections of users as well as individuals.

In *Exercise 4.01, Creating a superuser account*, we saw how we could cater to Alice, David, and Bob's requests to make modifications to their profiles. It was quite easy to do, and our application seems well equipped to handle their requests.

What will happen when the number of users grows? Will the admin user be able to manage 100 or 150 users at once? As you can imagine, this can be quite a complicated task. To overcome this, we can give elevated permissions to a certain set of users, and they can help ease the admin's tasks. And that's where groups come in handy. Though we'll learn more about users, groups, and permissions in *Chapter 9, Sessions and Authentication*, we can start understanding groups and their functionality by creating a **Help Desk user group** that contains accounts that have access to the admin interface but lack many powerful features, such as the ability to add, edit, or delete groups or to add or delete users.

Exercise 4.02 – adding and modifying users and groups through the admin app

In this exercise, we will grant a certain level of administrative access to one of our Bookr users, Carol. First, we will define the level of access for a group, and then we will add Carol to the group. This will allow Carol to update user profiles and check user logs. The following steps will help you implement this exercise:

1. Visit the admin interface at `http://127.0.0.1:8000/admin/` and log in as `bookradmin` using the account set up with the `superuser` command.
2. In the admin interface, follow the links to **Home** | **AUTHENTICATION AND AUTHORIZATION** | **Groups**:

Site administration

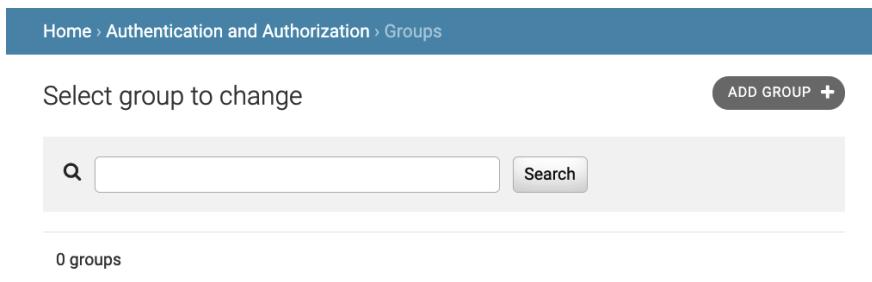


The screenshot shows the 'AUTHENTICATION AND AUTHORIZATION' section of the Django Admin. It has two main sections: 'Groups' and 'Users'. Each section has a 'Add' button (with a green plus sign) and a 'Change' button (with a yellow pencil icon).

AUTHENTICATION AND AUTHORIZATION		
Groups	Add	Change
Users	Add	Change

Figure 4.14 – The Groups and Users options on the AUTHENTICATION AND AUTHORIZATION page

3. Use **ADD GROUP +** in the top right-hand corner to add a new group:



Home > Authentication and Authorization > Groups

Select group to change

ADD GROUP +

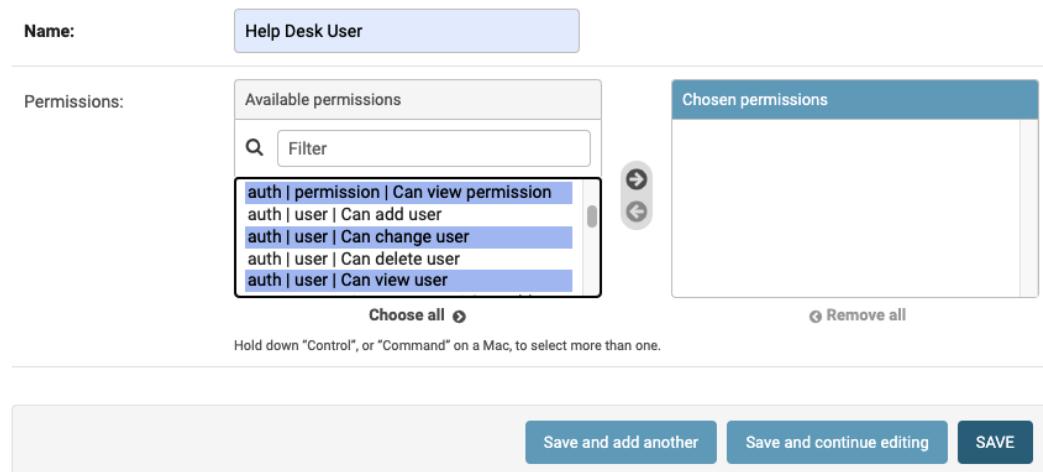
0 groups

Figure 4.15 – Adding a new group

4. Name the group `Help Desk User` and give it the following permissions, as shown in *Figure 4.16*:

- **Can view log entry**
- **Can view permission**
- **Can change user**
- **Can view user**

Add group



Name: Help Desk User

Permissions:

Available permissions

Filter

auth | permission | Can view permission
auth | user | Can add user
auth | user | Can change user
auth | user | Can delete user
auth | user | Can view user

Choose all Remove all

Hold down "Control", or "Command" on a Mac, to select more than one.

Save and add another Save and continue editing SAVE

Figure 4.16 – Selecting permissions

This can be done by selecting the permissions from **Available permissions** and clicking the right arrow in the middle so that they appear under **Chosen permissions**. Note that to add multiple permissions at a time, you can hold down the *Ctrl* key (or *Command* for Macintosh) to select more than one:

Add group

Name:	Help Desk User
Permissions:	<div style="display: flex; align-items: center;"> <div style="flex: 1;"> <p>Available permissions</p> <input type="text" value="Filter"/> <ul style="list-style-type: none"> auth user Can add user auth user Can delete user contenttypes content type Can add content contenttypes content type Can change content contenttypes content type Can delete content </div> <div style="flex: 1; border-left: 1px solid #ccc; padding-left: 10px;"> <p>Chosen permissions</p> <ul style="list-style-type: none"> admin log entry Can view log entry auth permission Can view permission auth user Can change user auth user Can view user </div> </div>
<input style="margin-right: 10px;" type="button" value="Choose all"/> <input type="button" value="Remove all"/>	
<p>Hold down "Control", or "Command" on a Mac, to select more than one.</p>	
<input style="margin-right: 10px;" type="button" value="Save and add another"/> <input type="button" value="Save and continue editing"/> <input style="background-color: #0070C0; color: white; border: 1px solid #0070C0; padding: 2px 10px; border-radius: 5px; font-weight: bold; margin-left: 10px;" type="button" value="SAVE"/>	

Figure 4.17 – Adding selected permissions into Chosen permissions

Once you click the **SAVE** button, you will see a confirmation message stating **The group “Help Desk User” was added successfully**:

Home › Authentication and Authorization › Groups

✓ The group “Help Desk User” was added successfully.

Select group to change	ADD GROUP +
<input type="text" value="Q"/> <input type="button" value="Search"/>	
Action: <input type="button" value="—"/> <input type="button" value="Go"/> 0 of 1 selected	
<input type="checkbox"/> GROUP	
<input type="checkbox"/> Help Desk User	
1 group	

Figure 4.18 – The group Help Desk User was added successfully

5. Now, navigate to **Home | AUTHENTICATION AND AUTHORIZATION | Users** and click the link of the user with the first name **carol**:

Select user to change ADD USER +

Action: 0 of 4 selected

<input type="checkbox"/>	USERNAME	EMAIL ADDRESS	FIRST NAME	LAST NAME	STAFF STATUS
<input type="checkbox"/>	alice	alice.white@example.com	Alice	White	✖
<input type="checkbox"/>	bob	bob.black@example.com	Bob	Black	✖
<input type="checkbox"/>	bookadmin	bookadmin@example.com			✓
<input type="checkbox"/>	carol	carol.brown@example.com	Carol	Brown	✖

4 users

FILTER

By staff status

- All
- Yes
- No

By superuser status

- All
- Yes
- No

By active

- All
- Yes
- No

Figure 4.19 – Clicking on the username carol

6. Scroll down to the **Permissions** fields set and select the **Staff status** checkbox. This is required for Carol to be able to log in to the admin app:

Permissions

Active
Designates whether this user should be treated as active. Unselect this instead of deleting accounts.

Staff status
Designates whether the user can log into this admin site.

Superuser status
Designates that this user has all permissions without explicitly assigning them.

Groups:

Available groups

- Help Desk User

Chosen groups

+

Choose all ↻

Remove all ✖

The groups this user belongs to. A user will get all permissions granted to each of their groups. Hold down "Control", or "Command" on a Mac, to select more than one.

Figure 4.20 – Clicking the Staff status checkbox

7. Add Carol to the **Help Desk User** group that we created in the previous steps by selecting it from the **Available groups** selection box (refer to *Figure 4.20*) and clicking the right arrow to shift it into her list of **Chosen groups** (as shown in *Figure 4.21*). Note that unless you do this, Carol won't be able to log in to the admin interface using her credentials:

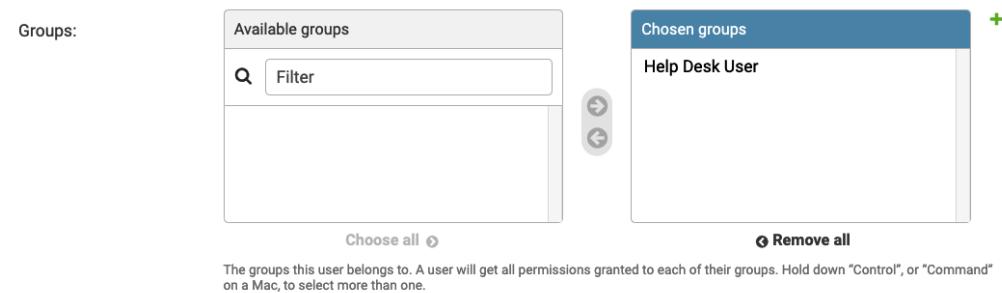


Figure 4.21 – Shifting the Help Desk User group into the list of Chosen groups for Carol

8. Let's test whether what we've done up till now has yielded the right outcome. To do this, log out of the admin site and log in again as Carol. Upon logging out, you should see the following on your screen:

Logged out

Thanks for spending some quality time with the Web site today.

[Log in again](#)

Figure 4.22 – The Logged out screen

Note

If you don't recall the password that you initially gave her, you can change the password in the command line by typing `python3 manage.py changepassword carol`.

Upon successful login, on the admin dashboard, you can see that there is no link to **Groups**:

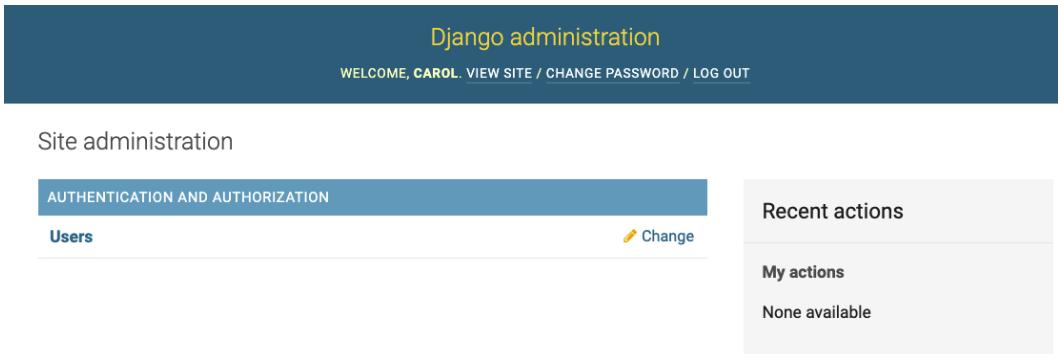


Figure 4.23 – The admin dashboard

As we did not assign any group permissions, not even **auth | group | Can view group**, to the **Help Desk User** group, when Carol logs in, the **Groups** admin interface is not available to her. Similarly, when you navigate to **Home | AUTHENTICATION AND AUTHORIZATION | Users** and click a user link, you will see that there are no options to edit or delete the user. This is because of the permissions that were granted to the **Help Desk User** group, of which Carol is a member. The members of the group can view and edit users but cannot add or delete any user.

In this exercise, we learned how to grant a certain amount of administrative privileges to users of our Django app.

Now that we have an understanding of how the admin app works with the **AUTHENTICATION AND AUTHORIZATION** models for **Users** and **Groups**, we can turn our attention to registering the models that we have developed in the reviews app.

Registering models with the admin app

The Django admin app produces workable CRUD interfaces to Django objects based on the characteristics of the models, with a minimal amount of coding. In this section, we will first look at how to register the models with the admin app and how elements of the interface, such as change lists and change pages, are derived from the model properties. We will conclude with an important exercise that demonstrates how foreign key settings in the model determine the behavior of the object deletion process.

Registering the reviews model

Let's say that Carol is tasked with improving the **Reviews** section in Bookr; that is, only the most relevant and comprehensive reviews should be shown, and duplicate or spam entries should be removed. For this, she will need access to the `reviews` model. As we have seen previously with our investigation of groups and users, the admin app already contains admin pages for the models from the Authentication and Authorization app, but it does not yet reference the models in our `Reviews` app. To make the admin app aware of the models, we need to explicitly register them with the admin app. Fortunately, we don't need to modify the admin app's code to do so, as we can instead import the admin app into our project and use its API to register our models. This has already been done in the Authentication and Authorization app, so let's try it with our `Reviews` app. Our aim is to be able to use the admin app to edit the data in our `reviews` model.

Take a look at the `reviews/admin.py` file. It is a placeholder file that was generated with the `startapp` subcommand that we used in *Chapter 1, An Introduction to Django*, and currently contains these lines:

```
from django.contrib import admin

# Register your models here.
```

Now we can try to expand this. To make the admin app aware of our models, we can modify the `reviews/admin.py` file and import the models. Then we can register the models with the `AdminSite` object, `admin.site`. The `AdminSite` object contains the instance of the Django admin application (later, we will learn how to subclass this `AdminSite` and override many of its properties). Then, `reviews/admin.py` will look as follows:

```
from django.contrib import admin
from reviews.models import (Publisher, Contributor,
                            Book, BookContributor, Review)

# Register your models here.
admin.site.register(Publisher)
admin.site.register(Contributor)
admin.site.register(Book)
admin.site.register(BookContributor)
admin.site.register(Review)
```

The `admin.site.register` method makes the models available to the admin app by adding it to a registry of classes contained in `admin.site._registry`. If we chose not to make a model accessible through the admin interface, we would simply not register it. When you reload `http://127.0.0.1:8000/admin/` in your browser, you will see the following on the admin app landing page:

Site administration

AUTHENTICATION AND AUTHORIZATION		
Groups	 Add	 Change
Users	 Add	 Change
REVIEWS		
Book contributors	 Add	 Change
Books	 Add	 Change
Contributors	 Add	 Change
Publishers	 Add	 Change
Reviews	 Add	 Change

Figure 4.24 – Admin app landing page

Note the change in the appearance of the admin page after the `reviews` model has been imported.

With the five models of the Reviews app registered with the admin app, we can examine some of the default interfaces that it has created for us.

Change lists

We now have change lists populated for our models. If we click the **Publishers** link, we will be taken to `http://127.0.0.1:8000/admin/reviews/publisher` and see a change list containing links to the publishers. These links are designated by the `id` field of the `Publisher` objects. If your database has been populated with the script in *Chapter 3, URL Mapping, Views, and Templates*, you will see a list with seven publishers, as shown in *Figure 4.25*.

Note

Depending on the state of your database and based on the activities you have completed, the object IDs, URLs, and links in these examples may be numbered differently from those listed here.

Select publisher to change

Action: 0 of 9 selected

PUBLISHER

Pocket Books

CreateSpace Independent Publishing Platform

Houghton Mifflin Harcourt

Simon and Schuster

Bay Back Books

Scribner

Penguin Classics

Harper Collins

Packt Publishing

9 publishers

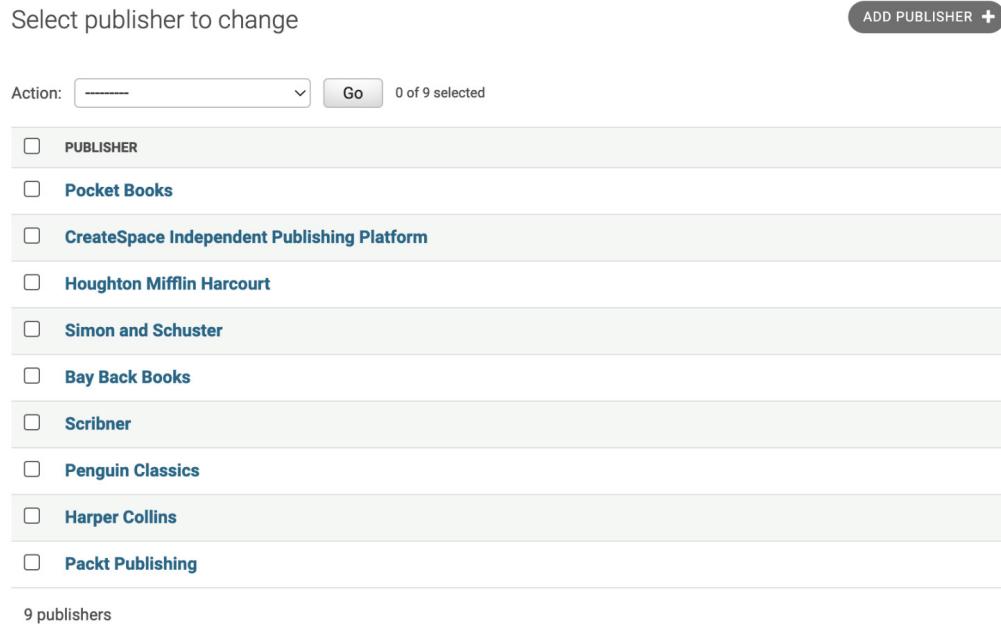


Figure 4.25 – The Select publisher to change list

We can follow one of the links in the change list and examine the **Change publisher** page.

The Change publisher page

The **Change publisher** page at `http://127.0.0.1:8000/admin/reviews/publisher/1` contains what we might expect (see *Figure 4.26*). There is a form for editing the publisher's details. These details have been derived from the `reviews.models.Publisher` class:

Change publisher

HISTORY

Name:	<input type="text" value="Packt Publishing"/> The name of the Publisher.
Website:	Currently: https://www.packtpub.com/ change: <input type="text" value="https://www.packtpub.com/"/> The Publisher's website.
Email:	<input type="text" value="info@packtpub.com"/> The Publisher's email address.

Delete **Save and add another** **Save and continue editing** **SAVE**

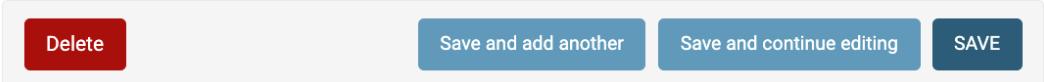


Figure 4.26 – The Change publisher page

If we had clicked the **ADD PUBLISHER +** button, the admin app would have returned a similar form for adding a publisher. The beauty of the admin app is that it gives us all of this CRUD functionality with just one line of coding – `admin.site.register(Publisher)` – using the definition of the `reviews.models.Publisher` attributes as a schema for the page content:

```
class Publisher(models.Model):  
    """A company that publishes books."""  
    name = models.CharField(max_length=50,  
                           help_text="The name of the Publisher.")  
    website = models.URLField(  
                           help_text="The Publisher's website.")  
    email = models.EmailField(  
                           help_text="The Publisher's email address.")
```

The publisher **Name** field is constrained to 50 characters as specified in the model. The help text that appears in gray below each field is derived from the `help_text` attributes specified on the model. We can see that `models.CharField`, `models.URLField`, and `models.EmailField` are rendered in HTML as input elements of the `text`, `url`, and `email` types, respectively.

The fields in the form come with validation where appropriate. Unless model fields are set to `blank=True` or `null=True`, the form will throw an error if the field is left blank, as is the case for the `Publisher.name` field. Similarly, as `Publisher.website` and `Publisher.email` are respectively defined as instances of `models.URLField` and `models.EmailField`, they are validated accordingly. In *Figure 4.27*, we can see the validation of **Name** as a required field, validation of **Website** as a URL, and validation of **Email** as an email address:

Add publisher

Please correct the errors below.

Name:	<div style="border: 1px solid #ccc; padding: 5px; border-radius: 3px; position: relative;"> <div style="color: red; font-size: 0.8em; margin-bottom: 5px;">This field is required.</div> <input name="name" style="width: 100%; height: 30px; border: none;" type="text" value=""/> <div style="font-size: 0.8em; margin-top: 2px;">The name of the Publisher.</div> </div>
Website:	<div style="border: 1px solid #ccc; padding: 5px; border-radius: 3px; position: relative;"> <div style="color: red; font-size: 0.8em; margin-bottom: 5px;">Enter a valid URL.</div> <input name="website" style="width: 100%; height: 30px; border: none;" type="text" value="packtcom"/> <div style="font-size: 0.8em; margin-top: 2px;">The Publisher's website.</div> </div>
Email:	<div style="border: 1px solid #ccc; padding: 5px; border-radius: 3px; position: relative;"> <div style="color: red; font-size: 0.8em; margin-bottom: 5px;">Enter a valid email address.</div> <input name="email" style="width: 100%; height: 30px; border: none;" type="text" value="info@packt.com"/> <div style="font-size: 0.8em; margin-top: 2px;">The Publisher's email address.</div> </div>

Save and add another
Save and continue editing
SAVE

Figure 4.27 – Field validation

It is useful to examine how the admin app renders elements of the models to understand how it functions. In your browser, right-click **View Page Source** and examine the HTML that has been rendered for this form. You will see a browser tab displaying something like this:

```
<fieldset class="module aligned ">
<div class="form-row errors field-name">
  <ul class="errorlist"><li>This field is
    required.</li></ul>
<div>
  <label class="required" for="id_name">Name:</label>
  <input type="text" name="name" class="vTextField"
    maxlength="50" required="" id="id_name">
```

```
    <div class="help">The name of the Publisher.</div>
  </div>
</div>
<div class="form-row errors field-website">
  <ul class="errorlist"><li>Enter a valid URL.</li></ul>
<div>
  <label class="required"
    for="id_website">Website:</label>
  <input type="url" name="website" value="packtcom"
    class="vURLField" maxlength="200" required=""
    id="id_website">
  <div class="help">The Publisher's website.</div>
</div>
</div>
<div class="form-row errors field-email">
  <ul class="errorlist"><li>Enter a valid email
    address.</li></ul>
<div>
  <label class="required" for="id_email">Email:</label>
  <input type="email" name="email"
    value="info@packt.com"
    class="vTextField" maxlength="254"
    required="" id="id_email">
  <div class="help">The Publisher's email
    address.</div>
</div>
</div>
</fieldset>
```

The form has a `publisher_form` ID and it contains a fieldset with HTML elements corresponding to the data structure of the `Publisher` model in `reviews/models.py`, shown as follows:

```
class Publisher(models.Model):
    """A company that publishes books."""
    name = models.CharField(max_length=50,
                           help_text="The name of the Publisher.")
    website = models.URLField(
        help_text="The Publisher's website.")
    email = models.EmailField(
        help_text="The Publisher's email address.")
```

Note that for the `name`, the input field is rendered like this:

```
<input type="text" name="name" value="Packt Publishing"
  class="vTextField" maxlength="50"
  required="" id="id_name">
```

It is a required field, and it has a `text` type and `maxlength` of 50, as defined by the `max_length` parameter in the model definition:

```
name = models.CharField(max_length=50,
    help_text="The name of the Publisher.")
```

Similarly, we can see the website and email being defined in the model as `URLField` and `EmailField` are rendered in HTML as input elements of the `url` and `email` types, respectively:

```
<input type="url" name="website"
    value="https://www.packtpub.com/"
    class="vURLField" maxlength="200" required=""
    id="id_website">
<input type="email" name="email" value="info@packtpub.com"
    class="vTextField" maxlength="254"
    required="" id="id_email">
```

We have learned that the Django admin app derives sensible HTML representations of Django models based on the model definitions that we have provided.

Now let's examine the Book change page.

The Book change page

Similarly, there is a change page that can be reached by selecting **Books** from the **Site administration** page and then selecting a specific book from the change list:

The screenshot shows the Django Site administration interface. The top navigation bar says "Site administration". Below it is a "Recent actions" sidebar with a list of recent user actions:

- carol User
- Help Desk User Group
- david User
- carol User
- bob User
- bob User

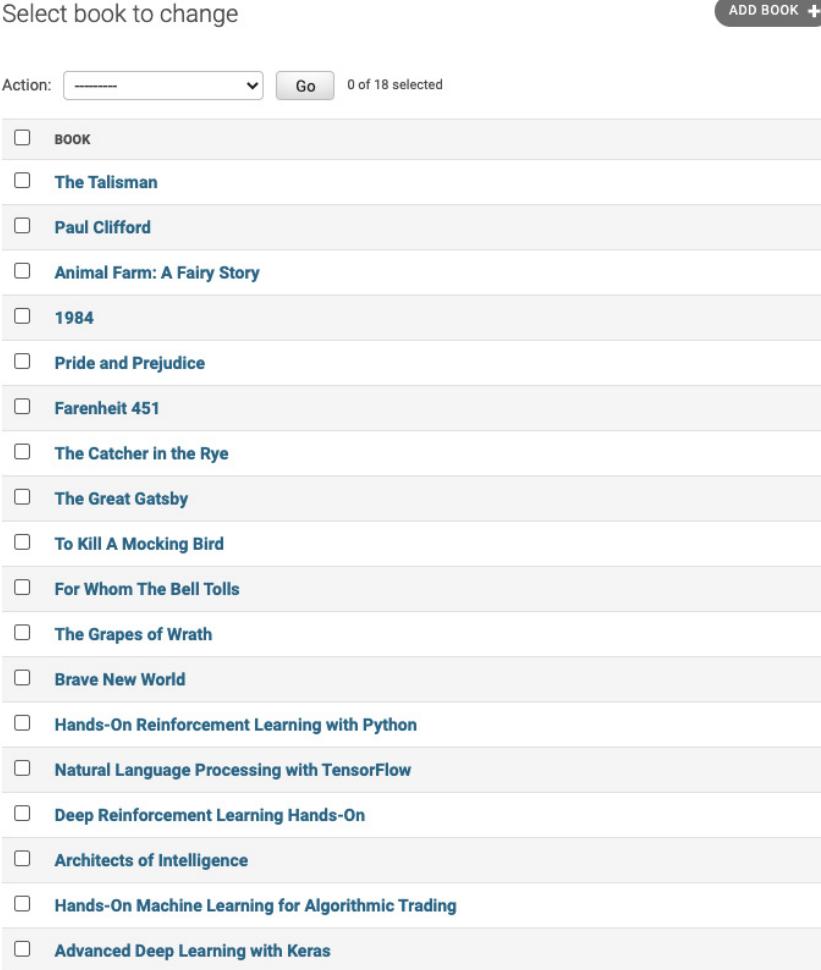
The main content area is divided into sections:

- AUTHENTICATION AND AUTHORIZATION** (Groups, Users)
- REVIEWS** (Book contributors, Books, Contributors, Publishers, Reviews)

Under the "Books" section, there is a link to "Books".

Figure 4.28 – Selecting Books from the Site administration page

After clicking **Books**, as shown in the preceding screenshot, you will see the following on your screen:



The screenshot shows a list of 18 books in the Django Admin interface. The list is titled "Select book to change" and includes the following titles:

- BOOK
- [The Talisman](#)
- [Paul Clifford](#)
- [Animal Farm: A Fairy Story](#)
- [1984](#)
- [Pride and Prejudice](#)
- [Fahrenheit 451](#)
- [The Catcher in the Rye](#)
- [The Great Gatsby](#)
- [To Kill A Mocking Bird](#)
- [For Whom The Bell Tolls](#)
- [The Grapes of Wrath](#)
- [Brave New World](#)
- [Hands-On Reinforcement Learning with Python](#)
- [Natural Language Processing with TensorFlow](#)
- [Deep Reinforcement Learning Hands-On](#)
- [Architects of Intelligence](#)
- [Hands-On Machine Learning for Algorithmic Trading](#)
- [Advanced Deep Learning with Keras](#)

At the bottom of the list, it says "18 books".

Figure 4.29 – The Book change page

In this instance, selecting the book *Architects of Intelligence* will take us to the URL `http://127.0.0.1:8000/admin/reviews/book/3/change/`. In the previous example, all the model fields were rendered as simple HTML text widgets. The rendering of some other subclasses of `django.db.models.Field` used in `models.Book` are worthy of closer examination:

Django administration

WELCOME, BOOKADMIN. [VIEW SITE](#) / [CHANGE PASSWORD](#) / [LOG OUT](#)

Home > Reviews > Books > Architects of Intelligence

Change book

Title: Architects of Intelligence

The title of the book.

Date the book was published: 2018-11-23 [Today](#)

Note: You are 11 hours ahead of server time.

ISBN number of the book: 9781789954531

Publisher: Packt Publishing

[Delete](#) [Save and add another](#) [Save and continue editing](#) **SAVE**

Figure 4.30 – The Change book page

Here, `publication_date` is defined using `models.DateField`. It is rendered using a date selection widget. The visual representation of the widgets will vary between different operating systems and choices of browser:

Title: Architects of Intelligence

The title of the book.

Date the book was published: 2018-11-23 [Today](#)

Note: You are 11 hours ahead of server time.

ISBN number of the book: 9781789954531

Publisher: Packt Publishing

NOVEMBER 2018						
S	M	T	W	T	F	S
				1	2	3
4	5	6	7	8	9	10
11	12	13	14	15	16	17
18	19	20	21	22	23	24
25	26	27	28	29	30	

[Yesterday](#) | [Today](#) | [Tomorrow](#)

Cancel

Figure 4.31 – Date selection widget

As Publisher is defined as a foreign key relation, it is rendered as a **Publisher** drop-down menu with a list of the Publisher objects:

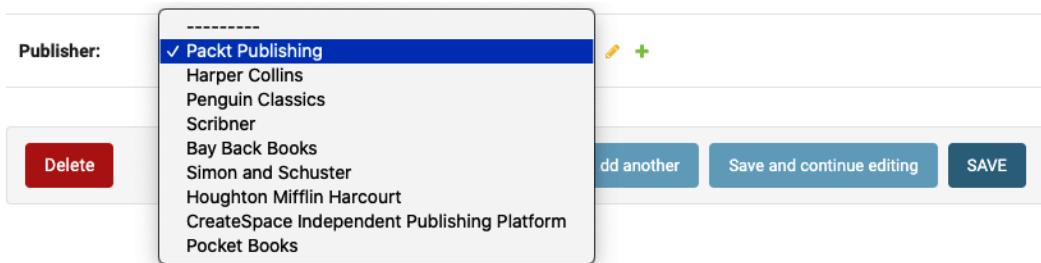


Figure 4.32 – Publisher drop-down menu

This brings us to how the admin app handles deletion. The admin app takes its cue from the models' foreign key constraints when determining how to implement deletion functionality. In the `BookContributor` model, `Contributor` is defined as a foreign key. The code in `reviews/models.py` looks as follows:

```
contributor = models.ForeignKey(Contributor,  
    on_delete=models.CASCADE)
```

By setting `on_delete=CASCADE` on a foreign key, the model specifies the database behavior required when a record is deleted; the deletion is cascaded to other objects that are referenced by the foreign key.

It is important to understand the impact of foreign key settings on the flow of the deletion process, and the following exercise examines this.

Exercise 4.03 – foreign keys and deletion behavior in the admin app

At present, all the `ForeignKey` relations in the `reviews` models are defined with an `on_delete=CASCADE` behavior. For instance, think of a case wherein an admin deletes one of the publishers. This would delete all the books that are associated with the publisher. We do not want that to happen, and that is precisely the behavior that we will be changing in this exercise:

1. Visit the **Contributors** change list at `http://127.0.0.1:8000/admin/reviews/contributor/` and select a contributor to delete. Make sure that the contributor is the author of a book.
2. Click the **Delete** button, but don't click **Yes, I'm sure** on the confirmation dialog. You will see a message like the one in *Figure 4.33*:

Home > Reviews > Contributors > Rowel > Delete

Are you sure?

Are you sure you want to delete the contributor "Rowel"? All of the following related items will be deleted:

Summary

- Contributors: 1
- Book contributors: 1

Objects

- Contributor: [Rowel](#)
- Book contributor: [BookContributor object \(19\)](#)

[Yes, I'm sure](#)

[No, take me back](#)

Figure 4.33 – Cascading delete confirmation dialog

In accordance with the `on_delete=CASCADE` argument to the foreign key, we are warned that deleting this `Contributor` object will have a cascading effect on a `BookContributor` object.

3. In the `reviews/models.py` file, modify the `Contributor` attribute of `BookContributor` to the following and save the file:

```
contributor = models.ForeignKey(Contributor,
    on_delete=models.CASCADE)
```

4. Now, try deleting the `Contributor` object again. You will see a message similar to the one in *Figure 4.34*:

Home > Reviews > Contributors > Rowel > Delete

Cannot delete contributor

Deleting the contributor 'Rowel' would require deleting the following protected related objects:

- Book contributor: [BookContributor object \(19\)](#)

Figure 4.34 – Foreign key protection error

Because the `on_delete` argument is `PROTECT`, our attempt to delete the object with dependencies will throw an error. If we used this approach in our model, we would need to delete objects in the `ForeignKey` relation before we deleted the original object. In this case, it would mean deleting the `BookContributor` object before deleting the `Contributor` object.

-
5. Now that we have learned about how the admin app handles the `ForeignKey` relations, let's revert the `ForeignKey` definition in the `BookContributor` class to the following:

```
contributor = models.ForeignKey(Contributor,  
    on_delete=models.CASCADE)
```

We have examined how the admin app's behavior adapts to the `ForeignKey` constraints that are expressed in model definitions. If the `on_delete` behavior is set to `models.PROTECT`, the admin app returns an error explaining why a protected object is blocking the deletion. This functionality can come in handy while building real-world apps, as there is often a chance of a manual error inadvertently leading to the deletion of important records. In the next section, we will look at how we can customize our admin app interface for a smoother user experience.

Customizing the admin interface

When first developing an application, the convenience of the default admin interface is excellent for building a rapid prototype of the app. Indeed, for many simpler applications or projects that require minimal data maintenance, this default admin interface may be entirely adequate. However, as the application matures to the point of release, the admin interface will generally need to be customized to facilitate more intuitive use and to robustly control data, subject to user permissions. You might want to retain certain aspects of the default admin interface and, at the same time, make some tweaks to certain features to better suit your purposes. For example, you would want the publisher list to show the complete names of the publishing houses instead of `Publisher (1)`, `Publisher (2)`, and so on. In addition to the aesthetic appeal, this makes it easier to use and navigate through the app.

In this section, we will be examining the `AdminSite` object and learning about some of its important properties that can be used to modify the global appearance of the admin app. Then we will examine approaches to subclassing `AdminSite` within our project so that we can customize it instead of using its default appearance. We will end by testing our knowledge of these techniques by creating a new project and customizing its admin interface.

Site-wide Django admin customizations

We have seen a page titled **Log in | Django site admin** containing a **Django Administration** form. However, an administrative user of the Bookr application may be somewhat perplexed by all this Django jargon, and it would be very confusing and a recipe for error if they had to deal with multiple Django apps that all had identical admin apps. As a developer of an intuitive and user-friendly application, you should customize this. Global properties such as these are specified as attributes of the `AdminSite` object. The following table details some of the simplest customizations to improve the usability of your app's admin interface:

AdminSite Attribute	Base Value	Description
site_title	"Django site admin"	Populates the <title> tag on each page of the admin interface.
site_header	"Django administration"	Sets the header on the login form.
index_title	"Site administration"	Sets the heading on the admin index page (where the models are listed).
index_template	None	Provides the path to find the admin index template. If unset, the admin/index.html template is used.
app_index_template	None	Provides the path to find the app admin index template. If unset, the admin/app_index.html template is used.
login_template	None	Provides the path to find the login template. If unset, the admin/login.html template is used.
logout_template	None	Provides the path to find the logout template. If unset, the registration/logout.html template is used.
password_change_template	None	Provides the path to find the password change template. If unset, the registration/password_change_form.html template is used.
password_change_done_template	None	Provides the path to find the password change done template. If unset, the registration/password_change_done.html template is used.

Figure 4.35 – Important AdminSite attributes

If this seems a little abstract, it might help if we examine the properties in this table using the Django shell.

Examining the AdminSite object from the Python shell

Let's take a deeper look at the `AdminSite` class. We have already encountered an object of the `AdminSite` class. It is the `admin.site` object that we used in the previous *Registering the reviews model* section. If the development server is not running, start it now with the `runserver` subcommand as follows (use `python` instead of `python3` for Windows):

```
python3 manage.py runserver
```

We can examine the `admin.site` object by importing the `admin` app in the Django shell, using the `manage.py` script again:

```
python3 manage.py shell
>>> from django.contrib import admin
```

We can interactively examine the default values of `site_title`, `site_header`, and `index_title` and see that they match the expected values of '`Django site admin`', '`Django administration`', and '`Site administration`' that we have already observed on the rendered web pages of the Django admin app:

```
>>> admin.site.site_title
'Django site admin'
>>> admin.site.site_header
'Django administration'
>>> admin.site.index_title
'Site administration'
```

The `AdminSite` class also specifies which forms and views are used to render the admin interface and determine its global behavior.

Now that we have examined the object's properties, we will attempt to subclass it so that we can customize them in our Django project.

Subclassing AdminSite – a first approach

We can make some modifications to the reviews/admin.py file. Instead of importing the django.contrib.admin module and using its site object, we will import AdminSite, subclass it, and instantiate our customized admin_site object. Consider the following code snippet. Here, BookrAdminSite is a subclass of AdminSite that contains custom values for site_title, site_header, and index_title; admin_site is an instance of BookrAdminSite, and we can use this instead of the default admin.site object to register our models. The reviews/admin.py file will look as follows:

```
from django.contrib.admin import AdminSite
from reviews.models import (Publisher, Contributor, Book,
    BookContributor, Review)

class BookrAdminSite(AdminSite):
    title_header = 'Bookr Admin'
    site_header = 'Bookr administration'
    index_title = 'Bookr site admin'

admin_site = BookrAdminSite(name='bookr')

# Register your models here.
admin_site.register(Publisher)
admin_site.register(Contributor)
admin_site.register(Book)
admin_site.register(BookContributor)
admin_site.register(Review)
```

Now we have created our own admin_site object that overrides the behavior of the admin.site object, we need to remove the existing references in our code to the admin.site object. In bookr/urls.py, we need to point admin to the new admin_site object and update our URL patterns. Otherwise, we would still be using the default admin site, and our customizations would be ignored. The change will look as follows:

```
from reviews.admin import admin_site
import reviews.views
from django.urls import include, path

urlpatterns = [
    path("admin/", admin_site.urls),
    path("book-search", reviews.views.book_search),
    path("", include("reviews.urls")),
]
```

This produces the expected results on the login screen:

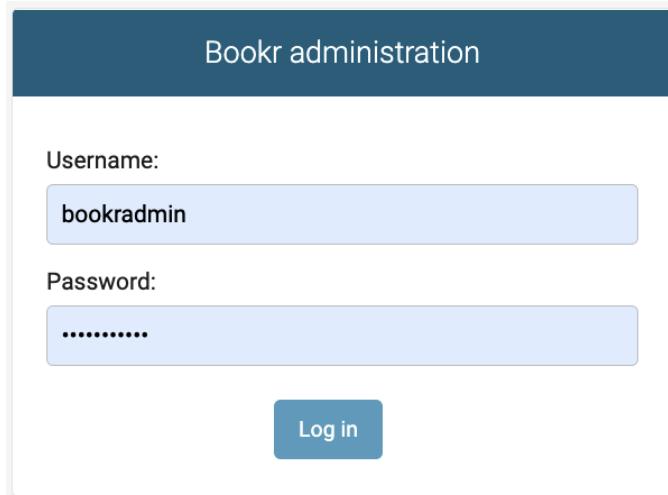


Figure 4.36 – Customizing the login screen

However, now there is a problem; we have lost the interface for auth objects. Previously, the admin app discovered the models registered in `reviews/admin.py` and `django.contrib.auth.admin` through the auto-discovery process, but now we have overridden this behavior by creating a new `AdminSite`:

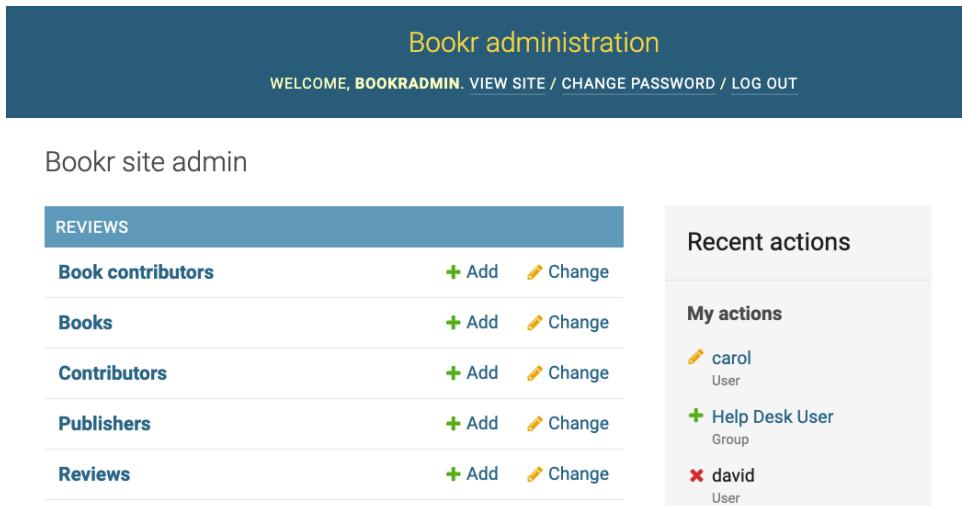


Figure 4.37 – Customized AdminSite is missing Authentication and Authorization

We could go down the path of referencing both the `AdminSite` objects to URL patterns in `bookr/urls.py`, but this approach would mean that we would end up with two separate admin apps for authentication and reviews. So, the URL `http://127.0.0.1:8000/admin` will take you to the original admin app derived from the `admin.site` object, while `http://127.0.0.1:8000/bookradmin` will take you to our `BookrAdminSite`, `admin_site`. This is not what we want to do, as we are still left with the admin app without the customizations that we added when we subclassed `BookrAdminSite`:

```
from django.contrib import admin
from reviews.admin import admin_site
from django.urls import path
urlpatterns = [path('admin/', admin.site.urls),
               path('bookradmin/', admin_site.urls),]
```

To avoid this problem, let's see what we can do in the next subsection.

Subclassing AdminSite – a better way

This has been a clumsy problem with the Django admin interface that has led to a lot of ad hoc solutions in earlier versions. However, since Django 2.1 came out, there has been a simple way of integrating a customized interface for the admin app without breaking auto-discovery or any of its other default features. Let's see how to do this with the following steps:

1. As `BookrAdminSite` is project-specific, the code does not really belong under our `reviews` folder. So, move `BookrAdminSite` to a new file called `admin.py` at the top level of the `Bookr` project directory:

```
from django.contrib import admin

class BookrAdminSite(admin.AdminSite):
    title_header = 'Bookr Admin'
    site_header = 'Bookr administration'
    index_title = 'Bookr site admin'
```

The URL settings path in `bookr/urls.py` changes to `path('admin/', admin.site.urls)` and we define our `ReviewsAdminConfig`.

2. Add a file called `reviews/adminconfig.py` containing these lines:

```
from django.contrib.admin.apps import AdminConfig
class ReviewsAdminConfig(AdminConfig):
    default_site = 'admin.BookrAdminSite'
```

3. Replace `django.contrib.admin` with `reviews.apps.ReviewsAdminConfig`, so that `INSTALLED_APPS` in the `bookr/settings.py` file will look as follows:

```
INSTALLED_APPS = [
    'reviews.adminconfig.ReviewsAdminConfig',
    'django.contrib.auth',
    'django.contrib.contenttypes',
    'django.contrib.sessions',
    'django.contrib.messages',
    'django.contrib.staticfiles',
    'reviews']
```

With the `ReviewsAdminConfig` specification of `default_site`, we no longer need to replace references to `admin.site` with a custom `AdminSite` object, `admin_site`. We can replace those `admin_site` calls with the `admin.site` calls that we had originally. Now, `reviews/admin.py` reverts to the following:

```
from django.contrib import admin
from reviews.models import (Publisher, Contributor,
    Book, BookContributor, Review)

# Register your models here.
admin.site.register(Publisher)
admin.site.register(Contributor)
admin.site.register(Book, BookAdmin)
admin.site.register(BookContributor)
admin.site.register(Review)
```

There are other aspects of `AdminSite` that we can customize, but we will revisit these in *Chapter 9, Sessions and Authentication*, once we have a fuller understanding of Django's templates and forms.

The following activity is aimed at testing the skills that you have learned to configure and customize the Django admin app for a new Django project.

Activity 4.02 – customizing the AdminSite object

You have learned how to modify attributes of the `AdminSite` object in a Django project. This activity will challenge you to use these skills to customize a new project and override its site title, site header, and index header. Also, you will replace the logout message by creating a project-specific template and setting it in our custom `AdminSite` object. You are developing a Django project that implements a message board called `Comment8or`. `Comment8or` is geared toward a technical demographic, so you need to make the phraseology succinct and abbreviated:

- The `Comment8or` admin site will be referred to as `c8admin`. This will appear on the site header and index title.
- For the title header, it will say `c8 site admin`.

- The default Django admin logout message is Thanks for spending some quality time with the Web site today. In Comment8or, it will say Bye from c8admin.

These are the steps that you need to follow to complete this activity:

1. Following the process that you learned in *Chapter 1, An Introduction to Django*, create a new Django project called `comment8or`, an app called `messageboard`, and run the migrations. Create a superuser called `c8admin`.
2. In the Django source code, there is a template for the logout page located in `djangotoolbox/contrib/admin/templates/registration/logged_out.html`. Make a copy of it in your project's directory under `comment8or/templates/comment8or` and modify the message in the template following the requirements.
3. Inside the project, create an `admin.py` file that implements a custom `AdminSite` object. Set the appropriate values for the `index_title`, `title_header`, `site_header`, and `logout_template` attributes based on the requirements.
4. Add a custom `AdminConfig` subclass to `messageboard/adminconfig.py`.
5. Replace the `admin` app with the custom `AdminConfig` subclass in `comment8or/settings.py`.
6. Configure the `TEMPLATES` setting so that the project's template is discoverable.

When the project was first created, the login page looked like this:

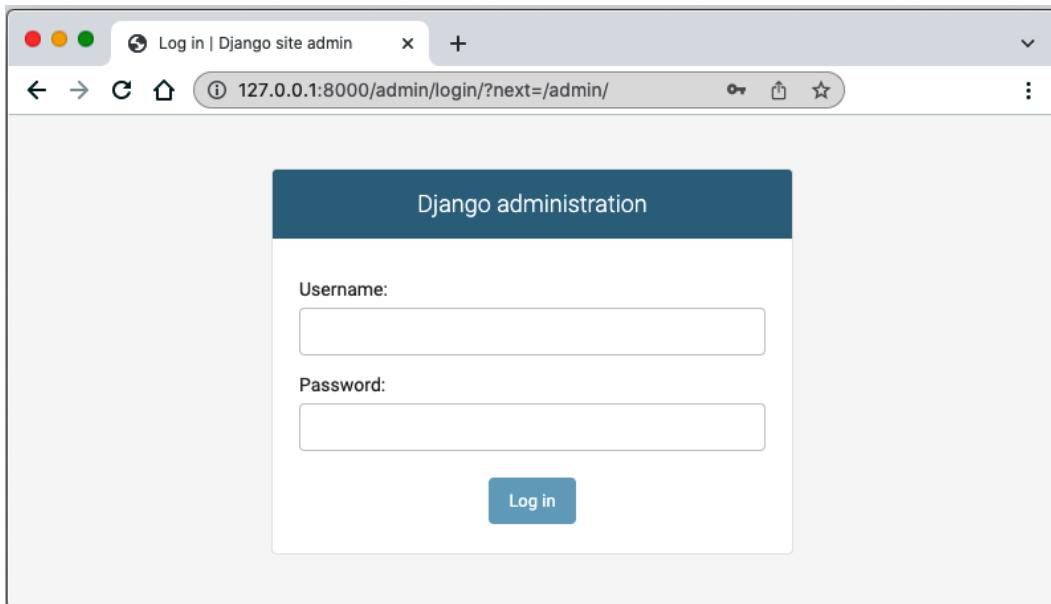


Figure 4.38 – Login page for the project

The app index page looked like this:

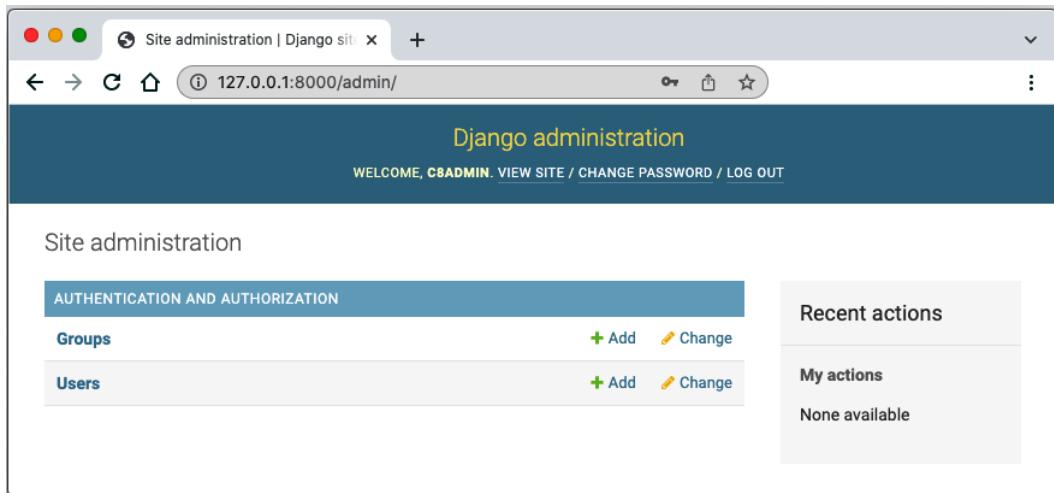


Figure 4.39 – App index page for the project

Finally, the logout page looked like this:

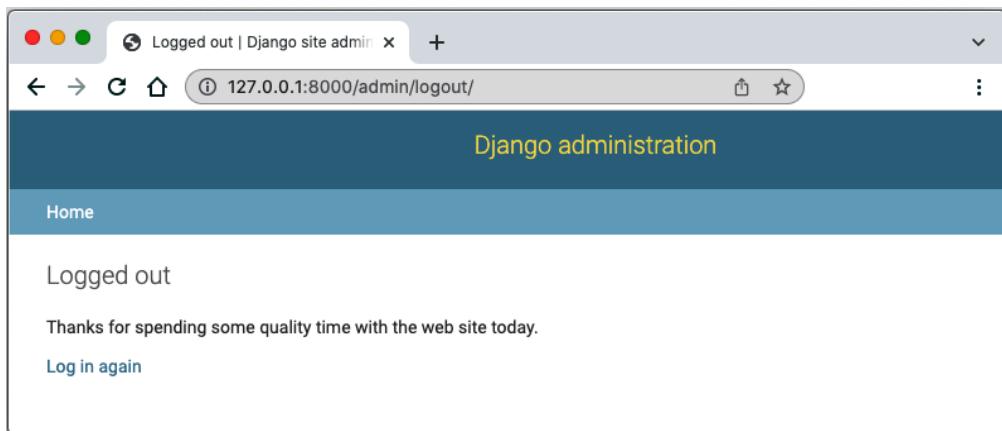


Figure 4.40 – Logout page for the project

After you have completed this activity with all the customizations, the login page will look like this:

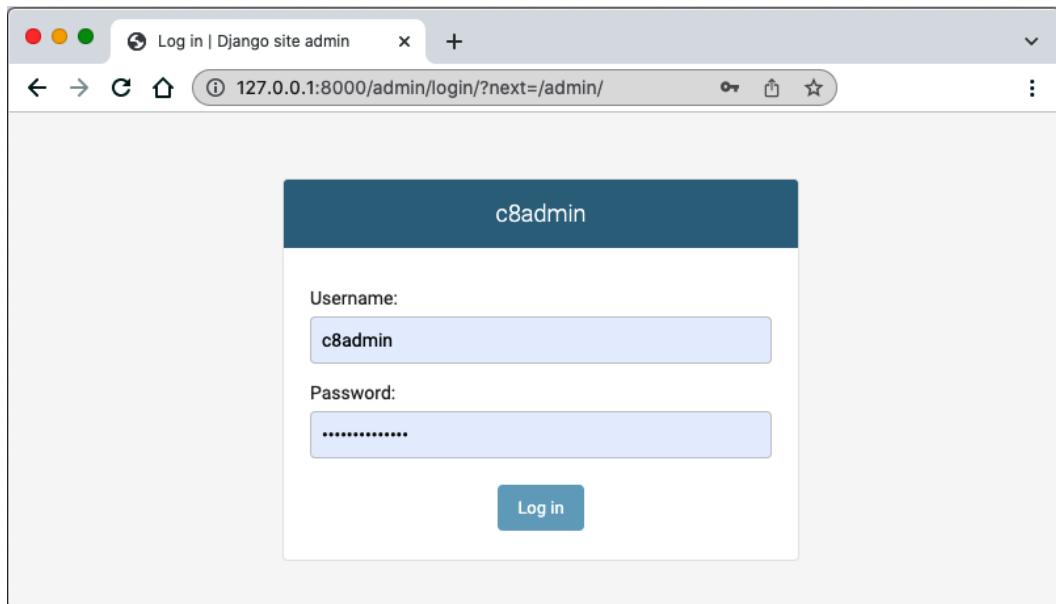


Figure 4.41 – Login page after customization

The app index page will look like this:

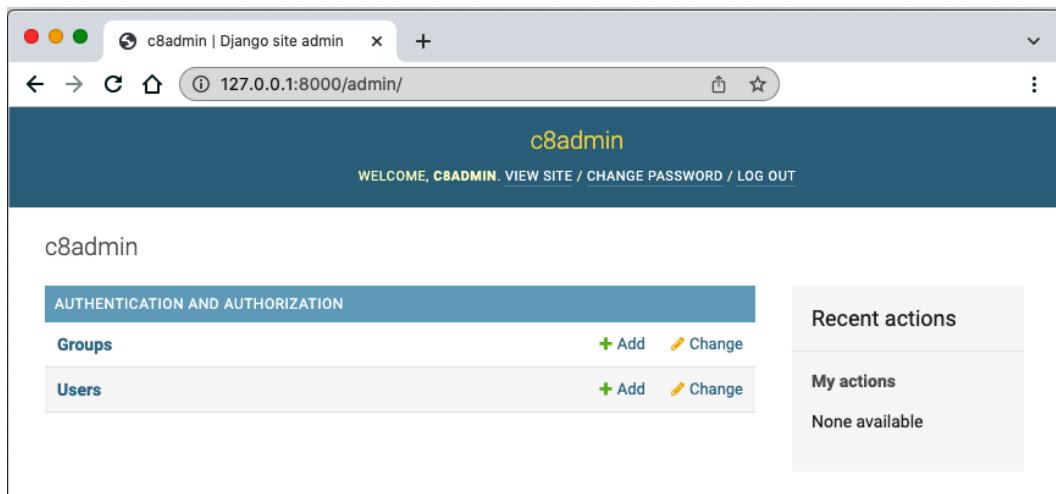


Figure 4.42 – App index page after customization

Finally, the logout page will look like this:

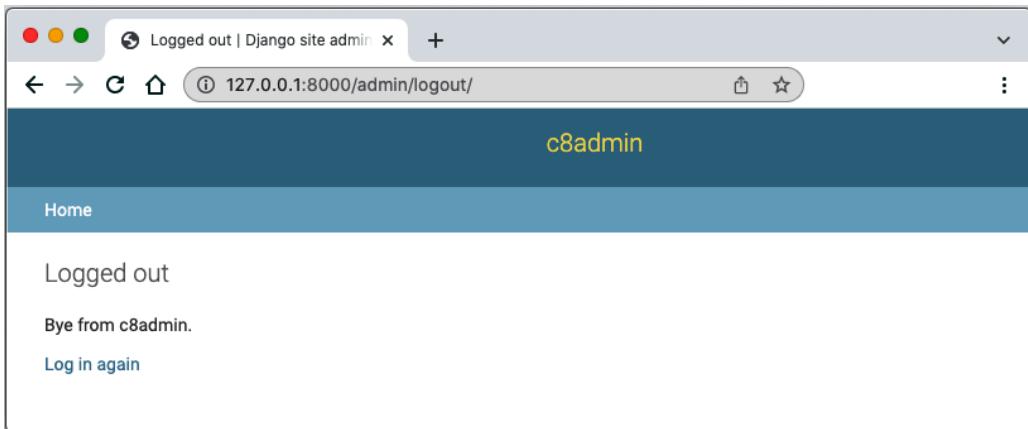


Figure 4.43 – Logout page after customization

You have successfully customized the admin app by subclassing `AdminSite`.

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

In the next section, we will examine the `AdminSite` object and learn about some of its important properties that can be used to modify the global appearance of the admin app. Then we will examine approaches to subclassing `AdminSite` within our project so that we can customize it away from its default appearance. We will end by testing our knowledge of these techniques by creating a new project and customizing its admin interface.

Customizing the ModelAdmin classes

Now that we've learned how a subclassed `AdminSite` can be used to customize the global appearance of the admin app, we will look at how to customize the admin app's interface to individual models. Owing to the admin interface being generated automatically from the models' structure, it has an overly generic appearance and needs to be customized for the sake of aesthetics and usability. Click one of the **Books** links in the admin app and compare it to the **Users** link. Both links take you to change list pages. These are the pages that a Bookr administrator visits when they want to add new books or add or alter the privileges of a user. As explained previously, a change list page presents a list of model objects with the option of selecting a group of them for bulk deletion (or other bulk activity), examining an individual object to edit it, or adding a new object. Notice the difference between the

two change list pages with a view to making our vanilla **Books** page as fully featured as the **Users** page. The following screenshot from the **Authentication and Authorization** app contains useful features such as a search bar, sortable column headers for important user fields, and a result filter:

	USERNAME	EMAIL ADDRESS	FIRST NAME	LAST NAME	STAFF STATUS
<input type="checkbox"/>	alice	alice.white@example.com	Alice	White	X
<input type="checkbox"/>	bob	bob.black@example.com	Bob	Black	X
<input type="checkbox"/>	bookadmin	bookadmin@example.com			✓
<input type="checkbox"/>	carol	carol.brown@example.com	Carol	Brown	✓

4 users

Figure 4.44 – The Users change list contains the customized ModelAdmin features

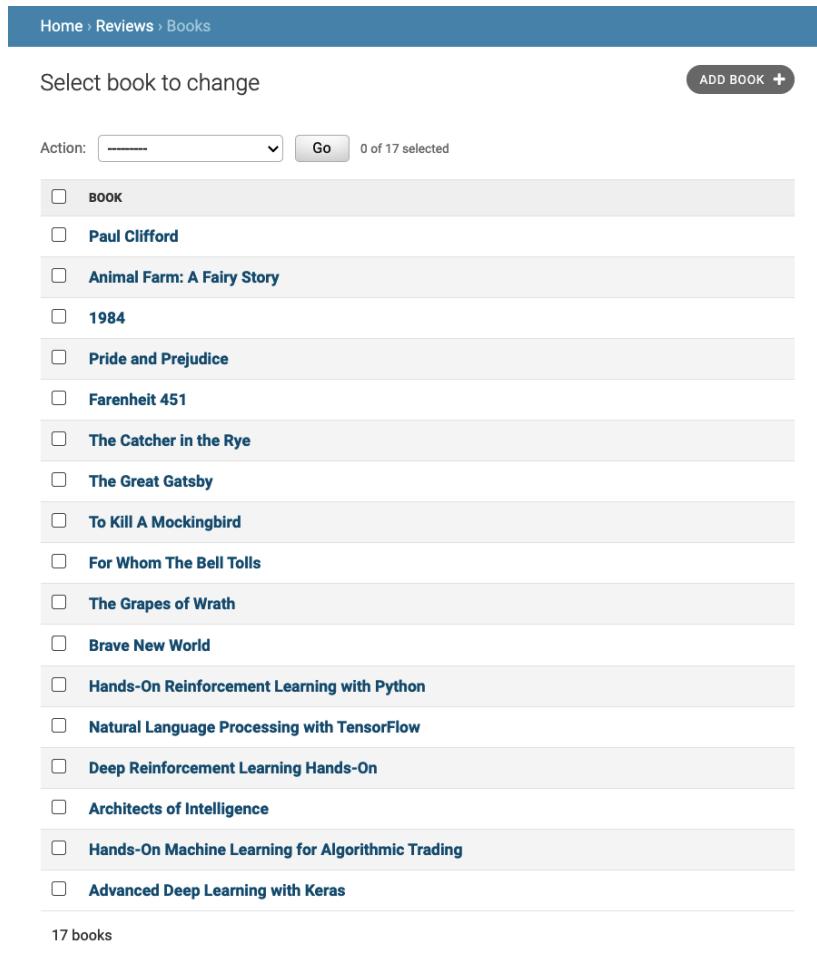
In this section, we will learn how the `ModelAdmin` classes can be developed to customize the fields on the change list page and to implement filters and search bars. We will look at how model fields can be grouped in admin CRUD forms or be excluded from the form if the fields should retain a default or system-generated value.

The list display fields

On the **Users** change list page, you will see the following:

- There is a list of user objects presented, summarized by their **USERNAME**, **EMAIL ADDRESS**, **FIRST NAME**, **LAST NAME**, and **STAFF STATUS** attributes.
- These individual attributes are sortable. The sorting order can be changed by clicking the headers.
- There is a search bar at the top of the page.
- In the right-hand column, there is a selection filter that allows the selection of several user fields, including some not appearing in the list display.

However, the behavior for the **Books** change list page is a lot less helpful. The books are listed by their titles but not in alphabetical order. The title column is not sortable and there are no filter or search options present:



Home : Reviews : Books

Select book to change

Action: 0 of 17 selected

BOOK

Paul Clifford

Animal Farm: A Fairy Story

1984

Pride and Prejudice

Fahrenheit 451

The Catcher in the Rye

The Great Gatsby

To Kill A Mockingbird

For Whom The Bell Tolls

The Grapes of Wrath

Brave New World

Hands-On Reinforcement Learning with Python

Natural Language Processing with TensorFlow

Deep Reinforcement Learning Hands-On

Architects of Intelligence

Hands-On Machine Learning for Algorithmic Trading

Advanced Deep Learning with Keras

17 books

Figure 4.45 – The Books change list

Recall from *Chapter 2, Models and Migrations*, that we defined the `__str__` methods on the Publisher, Book, and Contributor classes. In the case of the Book class, it had a `__str__()` representation that returns the book object's title:

```
class Book(models.Model):  
    ...  
    def __str__(self):  
        return self.title
```

If we had not defined the `__str__()` method on the Book class, it would have inherited it from the base Model class, `django.db.models.Model`.

This base class provides an abstract way to give a string representation of an object. For example, when we have Book with a primary key, in this case, the `id` field, with a value of 17, then we will end up with a string representation of `Book object (17)`:

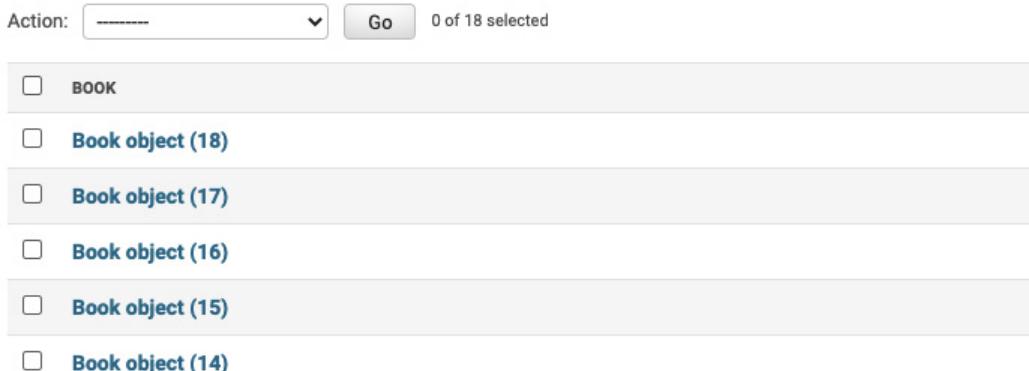
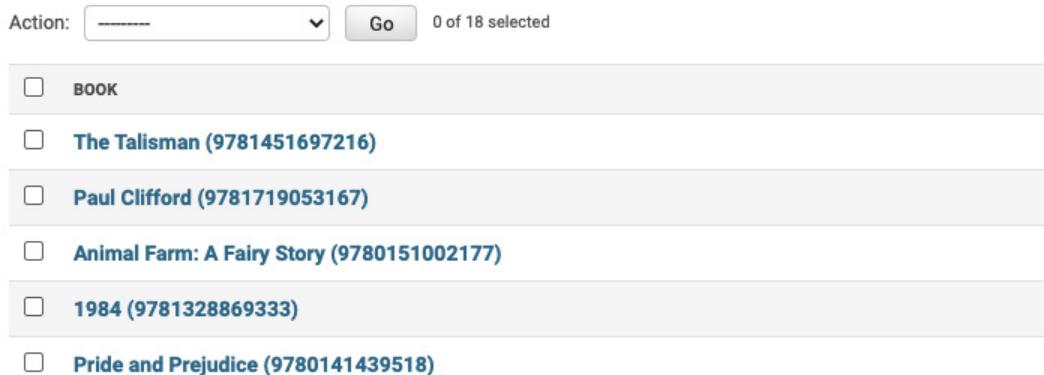


Figure 4.46 – The Books change list using the Model `__str__` representation

It could be useful in our application to represent a Book object as a composite of several fields. For example, if we wanted the books to be represented as *Title (ISBN)*, the following code snippet would produce the desired results:

```
class Book(models.Model):  
    ...  
    def __str__(self):  
        return "{} ({})".format(self.title, self.isbn)
```

This is a useful change in and of itself as it makes the representation of the object more intuitive in the app:



The screenshot shows a portion of the Django admin 'Books' change list. At the top, there is a 'Action' dropdown menu, a 'Go' button, and a status message '0 of 18 selected'. Below this, there is a list of books, each with a checkbox to its left. The list includes:

- BOOK
- The Talisman (9781451697216)
- Paul Clifford (9781719053167)
- Animal Farm: A Fairy Story (9780151002177)
- 1984 (9781328869333)
- Pride and Prejudice (9780141439518)

Figure 4.47 – A portion of the Books change list with the custom string representation

We are not limited to using the `__str__` representation of the object in the `list_display` field. The columns that appear in the list display are determined by the `ModelAdmin` class of the Django admin app. In the Django shell, we can import the `ModelAdmin` class and examine its `list_display` attribute:

```
python manage.py shell
>>> from django.contrib.admin import ModelAdmin
>>> ModelAdmin.list_display
('__str__',)
```

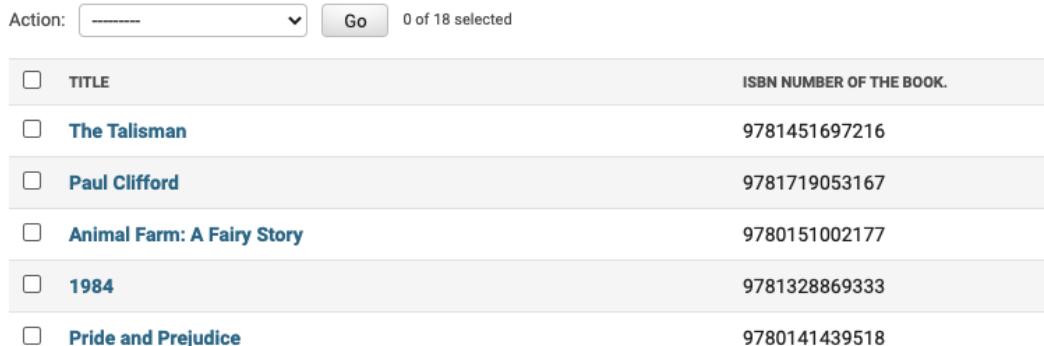
This explains why the default behavior of `list_display` is to display a single-columned table of the objects' `__str__` representations so that we can customize the list display by overriding this value. The best practice is to subclass `ModelAdmin` for each object. If we wanted the book list display to contain two separate columns for **Title** and **ISBN**, rather than having a single column containing both values, as in *Figure 4.47*, we would subclass `ModelAdmin` as `BookAdmin` and specify the custom `list_display`. The benefit of doing this is that we are now able to sort books by **Title** and by **ISBN**. We can add this class to `reviews/admin.py`:

```
class BookAdmin(admin.ModelAdmin):
    list_display = ('title', 'isbn')
```

Now that we've created a `BookAdmin` class, we should reference it when we register our `reviews.models.Book` class with the admin site. In the same file, we also need to modify the model registration to use `BookAdmin` instead of the default value of `admin.ModelAdmin`, so the `admin.site.register` call now becomes the following:

```
admin.site.register(Book, BookAdmin)
```

Once these two changes have been made to the `reviews/admin.py` file, we will get a **Books** change list page that looks like this:



The screenshot shows a portion of the Django admin interface for the 'Books' model. At the top, there is a search bar with placeholder text 'Action: -----', a dropdown menu, a 'Go' button, and a status message '0 of 18 selected'. Below this is a table with two columns. The first column contains checkboxes and the book titles: 'The Talisman', 'Paul Clifford', 'Animal Farm: A Fairy Story', '1984', and 'Pride and Prejudice'. The second column contains the ISBN numbers: '9781451697216', '9781719053167', '9780151002177', '9781328869333', and '9780141439518' respectively.

<input type="checkbox"/> TITLE	ISBN NUMBER OF THE BOOK.
<input type="checkbox"/> The Talisman	9781451697216
<input type="checkbox"/> Paul Clifford	9781719053167
<input type="checkbox"/> Animal Farm: A Fairy Story	9780151002177
<input type="checkbox"/> 1984	9781328869333
<input type="checkbox"/> Pride and Prejudice	9780141439518

Figure 4.48 – A portion of the Books change list with a two-column list display

This gives us a hint as to how flexible `list_display` is. It can take four types of values:

- It takes field names from the model, such as `title` or `isbn`.
- It takes a function that takes the model instance as an argument, such as this function that gives an initialized version of a person's name:

```
def initialled_name(obj):
    """ obj.first_names='Jerome David',
        obj.last_names='Salinger'=> 'Salinger, JD' """
    initials = ''.join([name[0] for name in
                       obj.first_names.split(' ')])
    return "({}), ({})".format(obj.last_names, initials)

class ContributorAdmin(admin.ModelAdmin):
    list_display = (initialled_name,)
```

- It takes a method from the `ModelAdmin` subclass that takes the model object as a single argument. Note that this needs to be specified as a string argument as it would be out of scope and undefined within the class:

```
class BookAdmin(admin.ModelAdmin):

    list_display = ('title', 'isbn13')
    def isbn13(self, obj):
        """ '9780316769174' => '978-0-31-676917-4' """
        return "{}-{}-{}-{}-{}".format(obj.isbn[0:3],
                                       obj.isbn[3:4], obj.isbn[4:6],
                                       obj.isbn[6:12], obj.isbn[12:13])
```

- It takes a method (or a non-field attribute) of the model class, such as `__str__`, as long as it accepts the model object as an argument. For example, we could convert `isbn13` to a method on the `Book` model class:

```
class Book(models.Model):

    def isbn13(self):
        """ '9780316769174' => '978-0-31-676917-4' """
        return "{}-{}-{}-{}-{}".format(self.isbn[0:3],
                                       self.isbn[3:4], self.isbn[4:6],
                                       self.isbn[6:12], self.isbn[12:13])
```

Now when viewing the **Books** change list at `http://127.0.0.1:8000/admin/reviews/book`, we can see the hyphenated ISBN13 field:

<input type="checkbox"/> TITLE	ISBN13
<input type="checkbox"/> The Talisman	978-1-45-169721-6
<input type="checkbox"/> Paul Clifford	978-1-71-905316-7
<input type="checkbox"/> Animal Farm: A Fairy Story	978-0-15-100217-7
<input type="checkbox"/> 1984	978-1-32-886933-3
<input type="checkbox"/> Pride and Prejudice	978-0-14-143951-8

Figure 4.49 – A portion of the Books change list with the hyphenated ISBN13

It is worth noting that computed fields such as `__str__` or our `isbn13` methods do not make for sortable fields on the summary page. Also, we cannot include fields of the `ManyToManyField` type in `list_display`.

While columns that are derived from model attributes derive their column headers and properties from the field attributes, we can specify properties of display columns that are computed fields using the `display` decorator.

The `display` decorator

When using a callable in the `list_display`, as in the cases of `initialled_name` and `isbn13`, we can use the `admin.display` decorator to specify the column name that will appear in the header of the change list using the `description` argument. We can also use it to get around the limitation of calculated fields not being sortable by specifying `ordering` on the callable. The `empty_value` argument can be used to specify how a `None` value or empty string is displayed. The default `empty_value` display is a single dash character:

```
@admin.display(
    ordering='isbn',
    description='ISBN-13',
    empty_value='--'
)
def isbn13(self, obj):
    """ '9780316769174' => '978-0-31-676917-4' """
    return "{}-{}-{}-{}-{}-{}".format(obj.isbn[0:3],
        obj.isbn[3:4], obj.isbn[4:6], obj.isbn[6:12],
        obj.isbn[12:13])
```

The `boolean` argument to `admin.display` can be used to flag a value to be represented in Boolean form:

```
@admin.display(
    boolean=True,
    description='Has ISBN',
)
def has_isbn(self, obj):
    """ '9780316769174' => True """
    return bool(obj.isbn)
```

Together these display decorator settings will give us display columns that look like this:

<input type="checkbox"/> TITLE	ISBN-13	HAS ISBN
To Kill A Mocking Bird	—	✗
Brave New World	978-0-06-085052-4	✓
Pride and Prejudice	978-0-14-143951-8	✓
The Grapes of Wrath	978-0-14-303943-3	✓
Animal Farm: A Fairy Story	978-0-15-100217-7	✓

Figure 4.50 – The Books change list with the admin display settings

The filter

Once the admin interface needs to deal with a significant number of records, it is convenient to narrow down the results that appear on change list pages. The simplest filters select individual values. For example, the user filter depicted in *Figure 4.6* allows the selection of users by choosing **By staff status**, **By superuser status**, and **By active**. We've seen on the user filter that `BooleanField` can be used as a filter. We can also implement filters on `CharField`, `TextField`, `DateTimeField`, `IntegerField`, `ForeignKey`, and `ManyToManyField`. In this case, by adding `publisher` as a `ForeignKey` of `Book`, it is defined on the `Book` class as follows:

```
publisher = models.ForeignKey(Publisher,  
                             on_delete=models.CASCADE)
```

Filters are implemented using the `list_filter` attribute of a `ModelAdmin` subclass. In our `Book` app, filtering by book title or ISBN would be impractical as it would produce a large list of filter options that return only one record. The filter that would occupy the right-hand side of the page would take up more space than the actual change list. A practical option would be to filter books by publisher. We defined a custom `__str__` method for the `Publisher` model that returns the publisher's name attribute, so our filter options will be listed as publisher names.

We can specify our change list filter in `reviews/admin.py` in the `BookAdmin` class like so:

```
list_filter = ('publisher',)
```

Here is how the **Books** change list page should look now:

Select book to change ADD BOOK +

Action: 0 of 6 selected

<input type="checkbox"/> TITLE	ISBN NUMBER OF THE BOOK.
<input type="checkbox"/> Hands-On Reinforcement Learning with Python	9781788836524
<input type="checkbox"/> Natural Language Processing with TensorFlow	9781788478311
<input type="checkbox"/> Deep Reinforcement Learning Hands-On	9781788834247
<input type="checkbox"/> Architects of Intelligence	9781789954531
<input type="checkbox"/> Hands-On Machine Learning for Algorithmic Trading	9781789346411
<input type="checkbox"/> Advanced Deep Learning with Keras	9781788629416

6 books

FILTER

By publisher

- All
- [Packt Publishing](#)
- [Harper Collins](#)
- [Penguin Classics](#)
- [Scribner](#)
- [Bay Back Books](#)
- [Simon and Schuster](#)
- [Houghton Mifflin Harcourt](#)
- [HardPress](#)

Figure 4.50 – The Books change list page with the publisher filter

With that line of code, we have implemented a useful publisher filter on the **Books** change list page.

In order to consolidate our knowledge of filters, we will add further filters to the **Books** change list page.

Exercise 4.04 – adding the date list_filter and date_hierarchy filters

We have seen that the `admin.ModelAdmin` class provides useful attributes to customize filters on change list pages. For example, filtering by date is a crucial functionality for many applications and can also help us make our app more user-friendly. In this exercise, we will examine how date filtering can be implemented by including a date field in the filter and look at the `date_hierarchy` filter:

1. Edit the `reviews/admin.py` file and modify the `list_filter` attribute in the `BookAdmin` class to include `'publication_date'`:

```
class BookAdmin(admin.ModelAdmin):
    list_display = ('title', 'isbn13')
    list_filter = ('publisher', 'publication_date')
```

-
2. Reload the **Books** change list page and confirm that the filter now includes date settings:

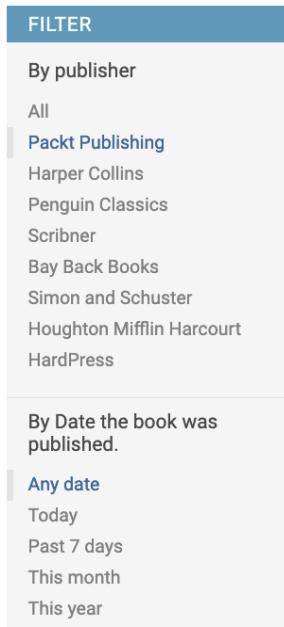


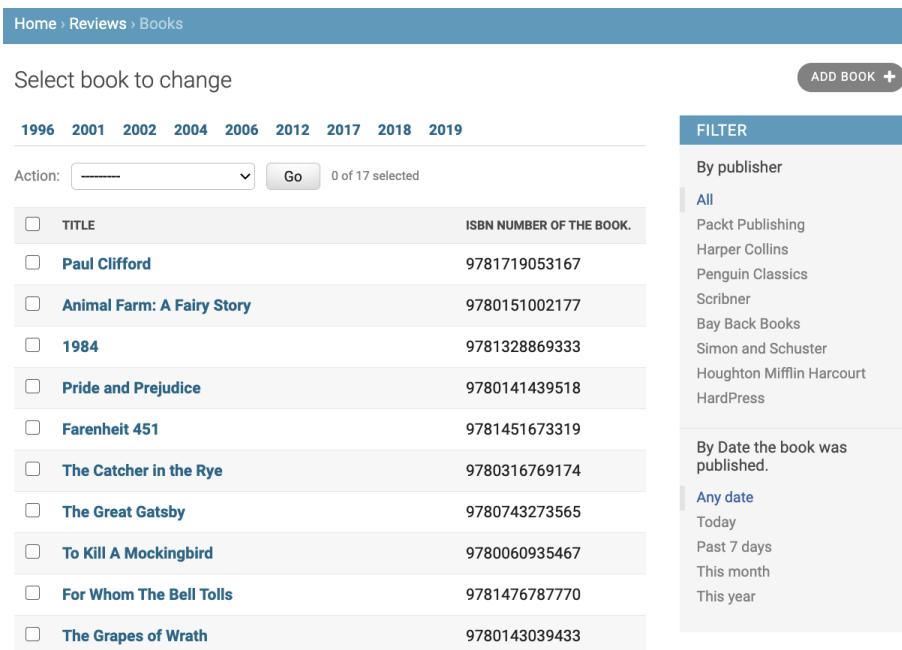
Figure 4.51 – Confirming that the Books change list page includes date settings

This publication date filter would be convenient if the Bookr project was receiving a lot of new releases, and we wanted to filter books by what was published in the last seven days or a month. Sometimes though, we might like to filter by a specific year or a specific month in a specific year. Fortunately, the `admin.ModelAdmin` class comes with a custom filter attribute geared towards navigating hierarchies of temporal information. It is called `date_hierarchy`.

3. Add a `date_hierarchy` attribute to `BookAdmin` and set its value to `publication_date`:

```
class BookAdmin(admin.ModelAdmin):  
    date_hierarchy = 'publication_date'  
    list_display = ('title', 'isbn13')  
    list_filter = ('publisher', 'publication_date')
```

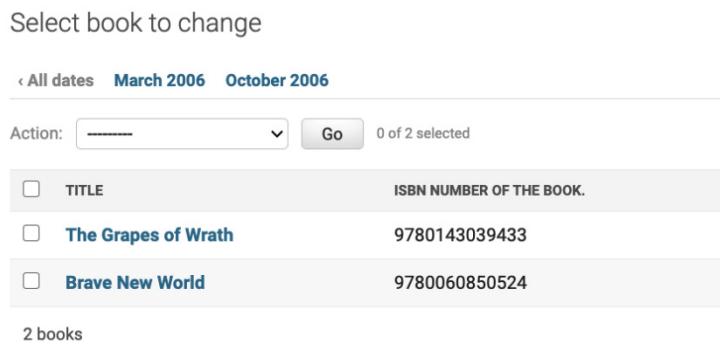
4. Reload the **Books** change list page and confirm that the date hierarchy appears above the **Action** drop-down menu:



The screenshot shows the Django Admin interface for the 'Books' model. At the top, there is a breadcrumb navigation: Home > Reviews > Books. Below the navigation, a title 'Select book to change' is displayed. On the right, there is a 'ADD BOOK +' button. A 'FILTER' sidebar is on the right, containing sections for 'By publisher' (with 'All' selected, showing Packt Publishing, Harper Collins, Penguin Classics, Scribner, Bay Back Books, Simon and Schuster, Houghton Mifflin Harcourt, and HardPress) and 'By Date the book was published' (with 'Any date' selected, showing Today, Past 7 days, This month, and This year). The main content area shows a table of books. The table has two columns: 'TITLE' and 'ISBN NUMBER OF THE BOOK.'. The books listed are: Paul Clifford (ISBN 9781719053167), Animal Farm: A Fairy Story (ISBN 9780151002177), 1984 (ISBN 9781328869333), Pride and Prejudice (ISBN 9780141439518), Farenheit 451 (ISBN 9781451673319), The Catcher in the Rye (ISBN 9780316769174), The Great Gatsby (ISBN 9780743273565), To Kill A Mockingbird (ISBN 9780060935467), For Whom The Bell Tolls (ISBN 9781476787770), and The Grapes of Wrath (ISBN 9780143039433). At the top of the table, there is a dropdown menu labeled 'Action' with a 'Go' button and a message '0 of 17 selected'.

Figure 4.52 – Confirming that the date hierarchy appears above the Action drop-down menu

5. Select a year from the date hierarchy and confirm that it contains a list of months in that year containing book titles and a total list of books:



The screenshot shows the Django Admin interface for the 'Books' model, similar to Figure 4.52 but with a different date selection. The 'FILTER' sidebar shows 'By Date the book was published' with 'March 2006' selected. The main content area shows a table of books for March 2006. The table has two columns: 'TITLE' and 'ISBN NUMBER OF THE BOOK.'. The books listed are: The Grapes of Wrath (ISBN 9780143039433) and Brave New World (ISBN 9780060850524). At the top of the table, there is a dropdown menu labeled 'Action' with a 'Go' button and a message '0 of 2 selected'. Below the table, a message '2 books' is displayed.

Figure 4.53 – Confirming that the selection of a year from the date hierarchy shows the books published that year

-
6. Confirm that selecting one of these months further filters down to days in the month:

Select book to change

« 2006 March 28

Action: 0 of 1 selected

<input type="checkbox"/> TITLE	ISBN NUMBER OF THE BOOK.
<input type="checkbox"/> The Grapes of Wrath	9780143039433

1 book

Figure 4.54 – Filtering months down to days in the month

The `date_hierarchy` filter is a convenient way of customizing a change list that contains a large set of time-sortable data in order to facilitate faster record selection, as we saw in this exercise.

Let's now look at the implementation of a search bar in our app.

The search bar

This brings us to the remaining piece of functionality that we wanted to implement – the search bar. Like filters, a basic search bar is quite simple to implement. We only need to add the `search_fields` attribute to the `ModelAdmin` class. The obvious character fields in our `Book` class to search on are `title` and `isbn`. At present, the `Books` change list appears with a date hierarchy across the top of the change list. The search bar will appear above this:

Select book to change ADD BOOK +

[1996](#) [2001](#) [2002](#) [2004](#) [2006](#) [2012](#) [2017](#) [2018](#) [2019](#)

Action: 0 of 17 selected

<input type="checkbox"/> TITLE	ISBN NUMBER OF THE BOOK.
<input type="checkbox"/> Paul Clifford	9781719053167
<input type="checkbox"/> Animal Farm: A Fairy Story	9780151002177
<input type="checkbox"/> 1984	9781328869333
<input type="checkbox"/> Pride and Prejudice	9780141439518
<input type="checkbox"/> Fahrenheit 451	9781451673319
<input type="checkbox"/> The Catcher in the Rye	9780316769174
<input type="checkbox"/> The Great Gatsby	9780743273565
<input type="checkbox"/> To Kill A Mockingbird	9780060935467
<input type="checkbox"/> For Whom The Bell Tolls	9781476787770
<input type="checkbox"/> The Grapes of Wrath	9780143039433

FILTER

By publisher

[All](#)
[Packt Publishing](#)
[Harper Collins](#)
[Penguin Classics](#)
[Scribner](#)
[Bay Back Books](#)
[Simon and Schuster](#)
[Houghton Mifflin Harcourt](#)
[HardPress](#)

By Date the book was published.

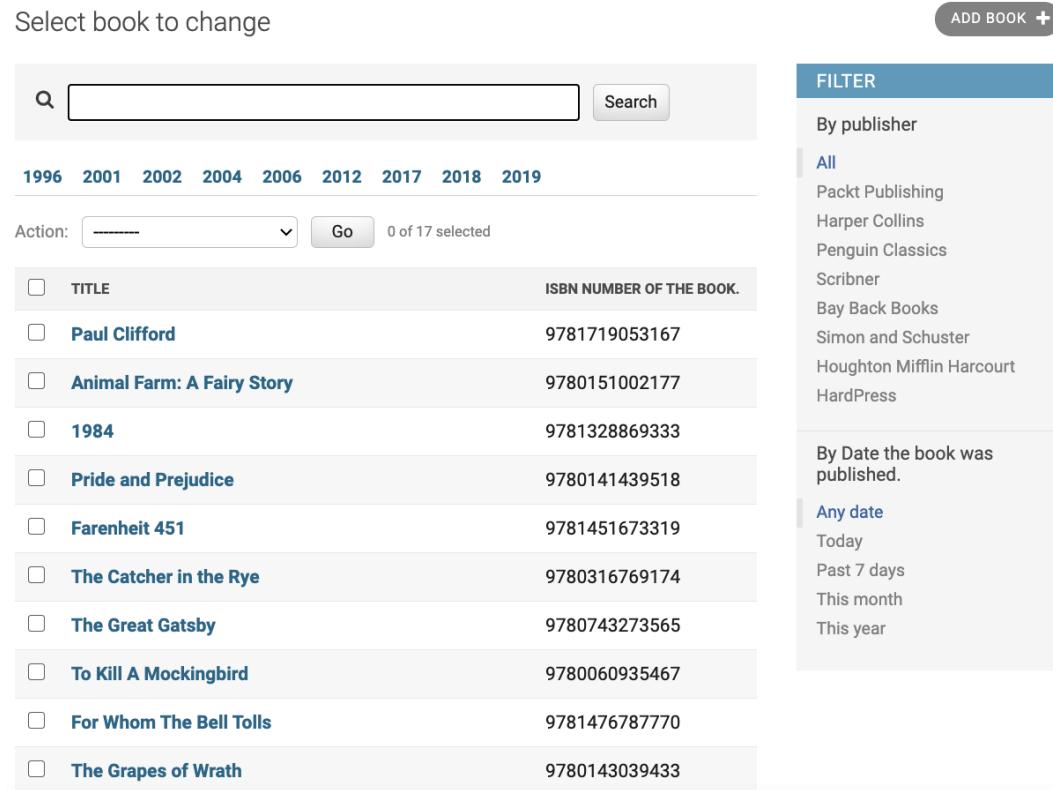
[Any date](#)
[Today](#)
[Past 7 days](#)
[This month](#)
[This year](#)

Figure 4.55 – The Books change list before the search bar is added

We can start by adding this attribute to BookAdmin in `reviews/admin.py` and examine the result:

```
search_fields = ('title', 'isbn')
```

The result would look like this:



The screenshot shows a Django ModelAdmin change list for the 'Books' model. At the top, there is a search bar with a magnifying glass icon and a 'Search' button. To the right of the search bar is a blue 'ADD BOOK +' button. On the far right, there is a 'FILTER' section with a 'By publisher' dropdown set to 'All'. The dropdown menu lists various publishers: Packt Publishing, Harper Collins, Penguin Classics, Scribner, Bay Back Books, Simon and Schuster, Houghton Mifflin Harcourt, and HardPress. Below the publisher filter is another dropdown for 'By Date the book was published', with options: 'Any date', 'Today', 'Past 7 days', 'This month', and 'This year'. The main content area displays a table of book data. The table has two columns: 'TITLE' and 'ISBN NUMBER OF THE BOOK.'. The data includes the following entries:

TITLE	ISBN NUMBER OF THE BOOK.
Paul Clifford	9781719053167
Animal Farm: A Fairy Story	9780151002177
1984	9781328869333
Pride and Prejudice	9780141439518
Fahrenheit 451	9781451673319
The Catcher in the Rye	9780316769174
The Great Gatsby	9780743273565
To Kill A Mockingbird	9780060935467
For Whom The Bell Tolls	9781476787770
The Grapes of Wrath	9780143039433

Figure 4.56 – The Books change list with the search bar

Now we can perform a simple text search on fields that match the title field or ISBN. This search requires precise string matches, so “color” won’t match “colour.” It also lacks the deep semantic processing that we expect from more sophisticated search facilities such as **Elasticsearch**. ISBN lookup is a very good feature if you happen to have a barcode scanner. Limiting our search to fields on the Books model is quite restrictive. We might want to search by publisher name too. Fortunately, `search_fields` is flexible enough to accomplish this. To search on `ForeignKeyField` or `ManyToManyField`, we just need to specify the field name on the current model and the field on the related model separated by two underscores. In this case, `Book` has a foreign key, `publisher`, and we want to search on the `Publisher.name` field so it can be specified as `'publisher__name'` on `BookAdmin`. `search_fields`:

```
search_fields = ('title', 'isbn', 'publisher__name')
```

If we wanted to restrict a search field to an exact match rather than return results that contain the search string, then the field can be suffixed with ' `exact`' . So, replacing ' `isbn`' with ' `isbn__exact`' will require the complete ISBN to be matched, and we won't be able to get a match using a portion of the ISBN.

Similarly, we constrain the search field to only return results that start with the search string by using the ' `startswith`' suffix. Qualifying the publisher name search field as ' `publisher__name__startswith`' means that we will get results searching for "pack" but not for "ackt."

This concludes our examination of common customizations of the change list pages. We may also want to customize the behavior of the create and update forms in our app.

Excluding and grouping fields

There are occasions when it is appropriate to restrict the visibility of some of the fields in the model in the admin interface. This can be achieved with the `exclude` attribute.

This is the review form screen with the **Date edited** field visible. Note that the **Date created** field does not appear – it is already a hidden view because `date_created` is defined on the model with the `auto_now_add` parameter:

Add review

Content:	An excellent introduction to the topic.	
The Review text.		
Rating:	5	
The rating the reviewer has given.		
Date edited:	Date: 2022-02-17	Today
	Time: 01:22:33	Now
Note: You are 11 hours ahead of server time.		
The date and time the review was last edited.		
Creator:	bookadmin	
Book:	Hands-On Machine Learning for Algorithmic Trading	
The Book that this review is for.		
<input type="button" value="Save and add another"/> <input type="button" value="Save and continue editing"/> <input type="button" value="SAVE"/>		

Figure 4.57 – The review form

If we wanted to exclude the **Date edited** field from the review form, we would do this in the `ReviewAdmin` class:

```
exclude = ('date_edited',)
```

Then the review form would appear without **Date edited**:

Add review

The screenshot shows the 'Add review' form in the Django admin. The fields are:

- Content:** A text area containing "An excellent introduction to the topic." with a placeholder "The Review text."
- Rating:** A dropdown menu with a pencil icon.
- Creator:** A dropdown menu with a pencil and plus icon.
- Book:** A dropdown menu with a pencil and plus icon.

At the bottom, there are three buttons: "Save and add another", "Save and continue editing", and a large "SAVE" button.

Figure 4.58 – The review form with the Date edited field excluded

Conversely, it might be more prudent to restrict the admin fields to those that have been explicitly permitted. This is achieved with the `fields` attribute. The advantage of this approach is that if new fields are added in the model, they won't be available in the admin form unless they have been added to the `fields` tuple in the `ModelAdmin` subclass:

```
fields = ('content', 'rating', 'creator', 'book')
```

This will give us the same result that we saw earlier.

Another option is to use the `fieldsets` attribute of the `ModelAdmin` subclass to specify the form layout as a series of grouped fields. Each grouping in `fieldsets` consists of a title followed by a dictionary containing a 'fields' key pointing to a list of field name strings:

```
fieldsets = (
    ("Linkage", {"fields": ("creator", "book")}),
    ("Review content", {"fields": ("content",
        "rating")}),
)
```

The review form should look as follows:

Add review

The screenshot shows the Django Admin 'Add review' form. It is divided into two main sections: 'Linkage' and 'Review content', each with its own title bar.

Linkage Section:

- Creator:** A dropdown menu showing 'bookadmin' with a pencil and plus icon for editing or adding.
- Book:** A dropdown menu showing 'Hands-On Machine Learning for Algorithmic Trading' with a pencil and plus icon for editing or adding. Below it, a note says 'The Book that this review is for.'

Review content Section:

- Content:** A text area containing the text 'An excellent introduction to the topic.' Below it, a note says 'The Review text.'
- Rating:** A text input field containing the number '5'. Below it, a note says 'The rating the reviewer has given.'

Bottom Action Bar:

- Save and add another
- Save and continue editing
- SAVE (in a dark blue button)

Figure 4.59 – The review form with fieldsets

If we want to omit the title on a fieldset, we can do so by assigning the `None` value to it:

```
fieldsets = (
    (None, {'fields': ('creator', 'book')}),
    ('Review content', {'fields': ('content',
        'rating')}))
```

Now, the review form should appear as shown in the following screenshot:

Add review

The screenshot shows a Django ModelAdmin 'Add review' form. At the top, it says 'Add review'. The first fieldset, titled 'Review content', contains fields for 'Creator' (a dropdown menu showing 'bookadmin' with edit and add icons) and 'Book' (a dropdown menu showing 'Hands-On Machine Learning for Algorithmic Trading' with edit and add icons). Below this, the 'Review content' fieldset (titled 'Review content' in a blue header) contains a 'Content' text area with the placeholder 'An excellent introduction to the topic.' and a 'Rating' dropdown menu. At the bottom, there are three buttons: 'Save and add another', 'Save and continue editing', and a large 'SAVE' button.

Figure 4.60 – The review form with the first fieldset untitled

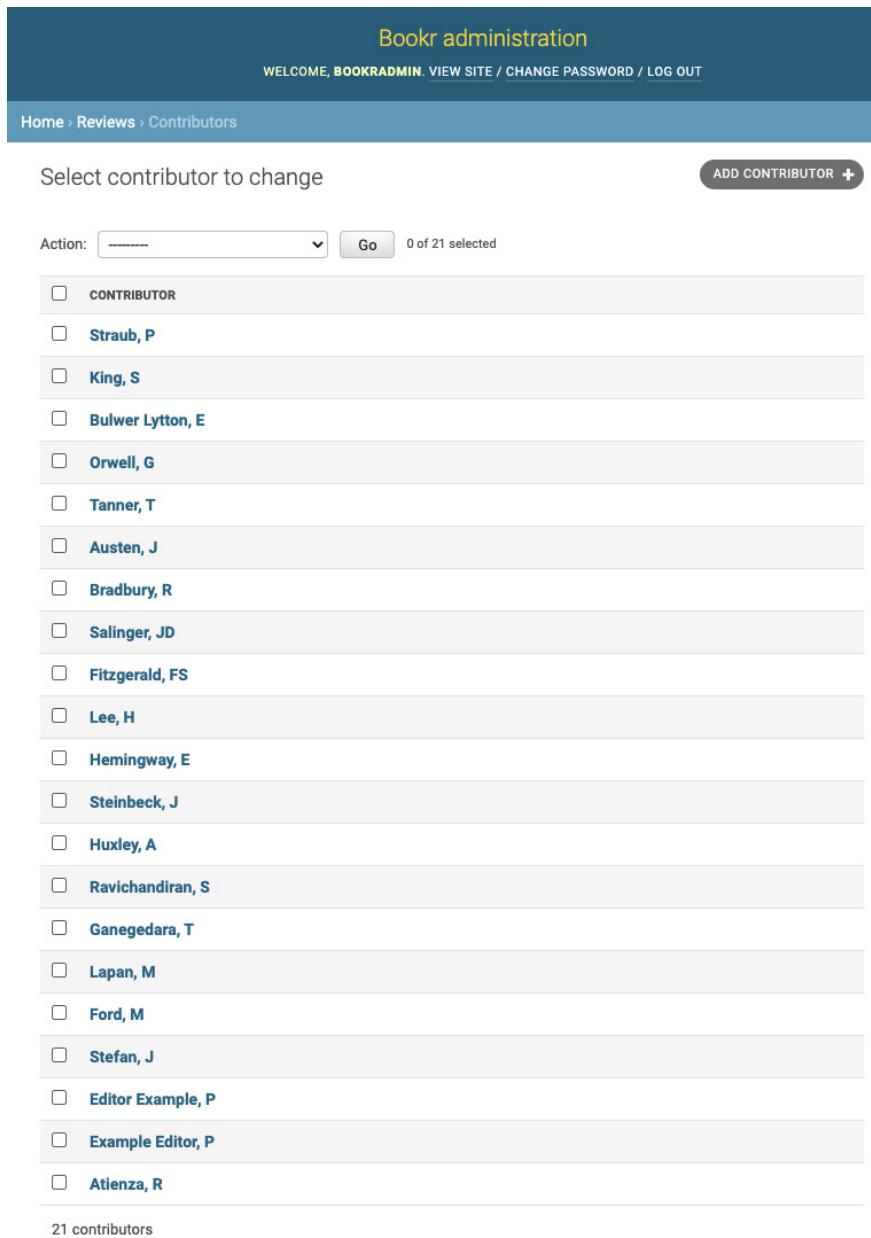
Now that we have learned about change lists and form customizations, we can put our knowledge into practice with a comprehensive activity.

Activity 4.03 – customizing the Model admins

In our data model, the `Contributor` class is used to store data for book contributors; they can be authors, contributors, or editors. This activity focuses on modifying the `Contributor` class and adding a `ContributorAdmin` class to improve the user-friendliness of the admin app. At present, the `Contributor` change list defaults to a single column, `FirstNames`, based on the `__str__` method created in *Chapter 2, Models and Migrations*. We will investigate some alternative ways of representing this. These steps will help you complete the activity:

1. Edit `reviews/models.py` to add additional functionality to the `Contributor` model.
2. Add an `initialled_name` method to `Contributor` that takes no arguments (like the `Book.isbn13` method).
3. The `initialled_name` method will return a string containing `Contributor.last_names` followed by a comma and the initials of the given names. For example, for a `Contributor` object with `first_names` of `Jerome David` and `last_names` of `Salinger`, `initialled_name` will return `Salinger, JD`.
4. Replace the `__str__` method for `Contributor` with one that calls `initialled_name()`.

At this point, the **Contributors** display list will look like this:



The screenshot shows the 'Contributors' display list in the Bookr administration interface. The title bar reads 'Bookr administration' with links for 'WELCOME, BOOKADMIN', 'VIEW SITE / CHANGE PASSWORD / LOG OUT'. The breadcrumb navigation shows 'Home > Reviews > Contributors'. The main content area is titled 'Select contributor to change' and includes a 'Action:' dropdown, a 'Go' button, and a status message '0 of 21 selected'. A large list of contributors is displayed in a table format, each with a selection checkbox. The contributors listed are: Straub, P; King, S; Bulwer Lytton, E; Orwell, G; Tanner, T; Austen, J; Bradbury, R; Salinger, JD; Fitzgerald, FS; Lee, H; Hemingway, E; Steinbeck, J; Huxley, A; Ravichandiran, S; Ganegedara, T; Lapan, M; Ford, M; Stefan, J; Editor Example, P; Example Editor, P; and Atienza, R. At the bottom of the list, it says '21 contributors'.

Action:	CONTRIBUTOR
<input type="checkbox"/>	Straub, P
<input type="checkbox"/>	King, S
<input type="checkbox"/>	Bulwer Lytton, E
<input type="checkbox"/>	Orwell, G
<input type="checkbox"/>	Tanner, T
<input type="checkbox"/>	Austen, J
<input type="checkbox"/>	Bradbury, R
<input type="checkbox"/>	Salinger, JD
<input type="checkbox"/>	Fitzgerald, FS
<input type="checkbox"/>	Lee, H
<input type="checkbox"/>	Hemingway, E
<input type="checkbox"/>	Steinbeck, J
<input type="checkbox"/>	Huxley, A
<input type="checkbox"/>	Ravichandiran, S
<input type="checkbox"/>	Ganegedara, T
<input type="checkbox"/>	Lapan, M
<input type="checkbox"/>	Ford, M
<input type="checkbox"/>	Stefan, J
<input type="checkbox"/>	Editor Example, P
<input type="checkbox"/>	Example Editor, P
<input type="checkbox"/>	Atienza, R

Figure 4.61 – The Contributors display list

5. Edit the `reviews/admin.py` file. We will modify `ContributorAdmin` that was subclassed from `admin.ModelAdmin` in the *The list display fields* section.
6. Modify it so that on the `Contributors` change list, records are displayed with two sortable columns: `Last Names` and `First Names`. We can remove the existing `list_display` attribute from `ContributorAdmin` along with the `initialised_name` function.
7. Add a search bar that searches on `Last Names` and `First Names`. Modify it so that it only matches the start of `Last Names`.
8. Add a filter on `Last Names`.

By completing the activity, you should see something like this:

The screenshot shows the Django admin interface for the 'Contributors' model. The top navigation bar includes 'WELCOME, BOOKADMIN', 'VIEW SITE / CHANGE PASSWORD / LOG OUT', and a 'Home' link. The main title is 'Bookr administration' and the current page is 'Home > Reviews > Contributors'. The main content area is titled 'Select contributor to change' and contains a table with 21 rows. The table has two columns: 'LAST NAMES' and 'FIRST NAMES'. The first few rows are: Straub (Peter), King (Stephen), Bulwer Lytton (Edward), Orwell (George), Tanner (Tony), Austen (Jane), Bradbury (Ray), Salinger (Jerome David), Fitzgerald (Francis Scott), Lee (Harper), Hemingway (Ernest), Steinbeck (John), Huxley (Aldous), Ravichandiran (Sudharsan), Ganegedara (Thushan), Lapan (Maxim), Ford (Martin), Stefan (Jansen), Editor Example (Packtp), Example Editor (Packt), and Atienza (Rowel). To the right of the table is a 'FILTER' sidebar with a dropdown for 'By last names' and a list of names: All, Atienza, Austen, Bradbury, Bulwer Lytton, Editor Example, Example Editor, Fitzgerald, Ford, Ganegedara, Hemingway, Huxley, King, Lapan, Lee, Orwell, Ravichandiran, Salinger, Stefan, Steinbeck, Straub, and Tanner. At the bottom of the table, it says '21 contributors'.

Figure 4.62 – Expected output

Changes such as these can be made to improve the functionality of the admin user interface. By implementing the `First Names` and `Last Names` columns as separate columns in the `Contributors` change list, we give the user an option to sort on either of the fields. By considering what columns are most useful in search retrieval and filter selections, we can improve the efficient retrieval of records.

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we saw how to create superusers through the Django command line and how to use them to access the admin app. Then, after a brief tour of the admin app's basic functionality, we examined how to register our models with it to produce a CRUD interface for our data.

Then we learned how to refine this interface by modifying site-wide features. We altered how the admin app presents model data to the user by registering custom model admin classes with the admin site. This allowed us to make fine-grained changes to the representation of our models' interfaces. These modifications included customizing change list pages by adding additional columns, filters, date hierarchies, and search bars. We also modified the layout of the model admin pages by grouping and excluding fields.

This was only a very shallow dive into the functionality of the admin app. We will revisit the rich functionality of `AdminSite` and `ModelAdmin` in *Chapter 10, Advanced Django Admin and Customization*. But first, we need to learn some more intermediate features of Django. In the next chapter, we will learn how to organize and serve static content, such as CSS, JavaScript, and images, from a Django app.

5

Serving Static Files

A web application with just plain **HyperText Markup Language (HTML)** is quite limiting. We can enhance the look of web pages with **Cascading Style Sheets (CSS)** and images, and add interaction with **JavaScript**. We call all these kinds of files “static files.” They are developed and then deployed as part of the application. We can compare this to dynamic responses, which are generated in real time when a request is made. All the views you have written generate a dynamic response by rendering a template. Note that we will not consider templates to be static files as they are not sent verbatim to a client; instead, they are rendered first and sent as part of a dynamic response.

During development, the static files are created on the developer’s machine, and then they must be moved to the production web server. If you must move to production in a short timeframe (say, a few hours) then it can be time-consuming to collect all the static assets, move them to the correct directory, and upload them to the server. When developing web applications using other frameworks or languages, you might need to manually put all of your static files into a specific directory that your web server hosts. Making changes to the URL from which static files are served might mean updating values throughout your code.

Django can manage static assets for us to make this process easier. It provides tools for serving them with its development server during development. When your application goes to production, it can also collect all your assets and copy them to a folder for a dedicated web server to host. This allows you to keep your static files segregated in a meaningful way during development, and automatically bundle them for deployment.

This functionality is provided by Django’s built-in `staticfiles` app. It adds several useful features for working with and serving static files:

- The `static` template tag automatically builds the static URL for an asset and includes it in your HTML.
- A view (called `static`) serves static files in development.
- Static file finders can be used to customize where assets are found on your filesystem.

- The `collectstatic` management command finds all static files and moves them into a single directory for deployment.
- The `findstatic` management command shows which static file on the disk has been loaded for a particular request. This also helps us debug if a particular file is not being loaded.

In the exercises and activities of this chapter, we will be adding static files (images and CSS) to the Bookr application. Each file will be stored inside the Bookr project directory during development. We will need to generate a URL for each so that the templates can reference them, and the browser can download them. Once the URL has been generated, Django will need to serve these files. When we deploy the Bookr application to production, all the static files need to be found and moved to a directory where they can be served by the production web server. If there are static files that are not loading as expected, we need some method of determining what the cause is.

For the sake of simplicity, let's take a single static file as an example: `logo.png`. We will briefly introduce the role of each feature we mentioned in the previous paragraph and explain it in depth throughout this chapter. There are five parts of the static files app that we will look at:

- The `static` template tag is used to convert a filename into a URL or path that can be used in a template – for example, from `logo.png` to `/static/logo.png`.
- The `static` view receives a request to load the static file at the `/static/logo.png` path. It reads the file and sends it to the browser.
- A static file finder (or just `finder`) is used by the `static` view to locate the static file on the disk. There are different finders, but in this example, one is just converting from the `/static/logo.png` URL path to the path on the disk, `bookr/static/logo.png`.
- When deploying to production, the `collectstatic` management command must be used. This will copy the `logo.png` file from the Bookr project directory to a web server directory, such as `/var/www/bookr/static/logo.png`.
- If a static file is not working (for example, the request for it returns a `404 Not Found` response, or the wrong file is being served), then we can use the `findstatic` management command to try to determine the reason. This command takes the filename as a parameter and will output which directories were looked through and where it was able to locate that requested file.

These are the most common features that are used day to day, but there are others that we will also discuss.

In this chapter, you will start by learning the difference between static and dynamic responses. You will then see how the Django `staticfiles` app helps manage static files. As you continue to work on the Bookr app, you will enhance it with images and CSS. You'll see the different ways you can lay out your static files for your project and examine how Django consolidates them for production deployment. Django includes tools for referencing static files in templates; you'll see how these tools help reduce the amount of work you must do when deploying the application to production. After this,

you'll explore the `findstatic` command, which can be used to debug issues with your static files. Later, you'll get an overview of how to write code for storing static files on a remote service. Finally, you'll look at caching web assets and how Django can help with cache invalidation.

We will cover the following main topics in this chapter:

- Static file serving
- Introduction to Static Files Finder
- Generating static URLs with the `static` template tag
- `FileSystem` Finder
- Static file finders – use during `collectstatic`
- `STATICFILES_DIRS` prefixed mode
- The `findstatic` command
- Serving the latest files (for cache validation)
- Custom storage engines

Technical requirements

All the code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter05>.

Static file serving

In the introduction, we mentioned that Django includes a view function called `static` that serves static files. The first important point to make regarding serving static files is that Django does not intend to serve them in production. It is not Django's role, and in production, Django will refuse to serve static files. This is normal and intended behavior. If Django is just reading from the filesystem and sending out a file, then it has no advantage over a normal web server, which will probably be more performant at this task. Further, if you serve static files with Django, you will keep the Python process busy for the duration of this request and it will be unable to serve the dynamic requests for which it is more suited.

For these reasons, the Django `static` view is designed only for use during development and will not work if your `DEBUG` setting is `False`. Since during development we usually only have one person accessing the site at a time (the developer), then Django is fine to serve static files. We will discuss how the `staticfiles` app supports production deployment more shortly. The entire production deployment process will be covered in *Chapter 17, Deploying a Django Application (Part 1 – Server Setup)*, which is hosted on the GitHub repository of this book.

A URL mapping to the `static` view is automatically set up when running the Django development server, provided that your `settings.py` file consists of the following elements:

- Has `DEBUG` set to `True`
- Contains '`django.contrib.staticfiles`' in its `INSTALLED_APPS`

Both settings exist by default.

The URL mapping that is created is roughly equivalent to having the following map in your `urlpatterns`:

```
path(settings.STATIC_URL, django.conf.urls.static)
```

Any URL starting with `settings.STATIC_URL` (which is `/static/` by default) gets mapped to the `static` view.

Note

You could use the `static` view without having `staticfiles` in `INSTALLED_APPS`, but you must set up an equivalent URL mapping manually.

Next, we'll talk about how Django locates static files using Static Files Finder.

Introduction to Static Files Finder

There are three times when Django needs to locate static files on disk, and for this, it uses **static file finder**. It can be thought of as a plugin. It is a class that implements methods for converting URL paths into disks and iterates through the project directory to find static files.

The first time Django does this is when the Django `static` view receives a request to load a particular static file; here, it needs to convert the path in the URL into a location on disk. For example, let's say the URL's path is `/static/logo.png`, and it is converted into the `bookr/static/logo.png` path on the disk. As we noted in the previous section, this only happens during development. On a production server, Django should not receive this request as it will be handled directly by the web server.

The second time is when using the `collectstatic` management command. This gathers all the static files in the project directory and copies them to a single directory to be served by the production web server. `bookr/static/logo.png` will get copied to the web server root – for example, `/var/www/bookr/static/logo.png`. Static File Finder contains code for locating all the static files inside your project directory.

The last time Static File Finder is used is while executing the `findstatic` management command. This is similar to the first usage, where it accepts a static filename (such as `logo.png`), but it outputs the full path (`bookr/static/logo.png`) to the terminal instead of loading the file's content.

Django comes with some built-in finders, but you can also write your own if you want to store static files in a custom directory layout. The list of finders Django uses is defined by the `STATICFILES_FINDERS` setting in `settings.py`. In this chapter, we will cover the behavior of the default Static File Finders, `AppDirectoriesFinder` and `FileSystemFinder`, in the `AppDirectoriesFinder` and `FileSystemFinder` sections, respectively.

Note

If you look in `settings.py`, you won't see the `STATICFILES_FINDERS` setting defined by default. This is because Django will use its built-in default setting, which is defined as `['django.contrib.staticfiles.finders.FileSystemFinder', 'django.contrib.staticfiles.finders.AppDirectoriesFinder']`. If you add the `STATICFILES_FINDERS` setting to your `settings.py` file to include a custom finder, be sure to include these defaults if you're using them.

In this section, we will first discuss static file finders and their use in the first case – responding to a request. Then, we will introduce some more concepts and return to the behavior of `collectstatic` and how it uses static file finders.

Static file finders – use during a request

When Django receives a request for a static file (remember, Django will only serve static files during development), each static file finder that has been defined will be queried until a file on the disk has been found. Or, if none of the finders can locate a file, the `static` view will return an HTTP 404 Not Found response.

For example, the URL of the request will be something like `/static/main.css` or `/static/reviews/logo.png`. Each finder will be queried in turn with the path from the URL and will return a path such as `bookr/static/main.css` for the first file and `bookr/reviews/static/reviews/logo.png` for the second. Each finder will use its logic to convert from a URL path into a filesystem path – we will discuss this logic in the `AppDirectoriesFinder` and `FileSystemFinder` sections.

AppDirectoriesFinder

The `AppDirectoriesFinder` class is used to find static files inside each app directory, in a directory called `static`. The application must be listed in the `INSTALLED_APPS` setting in your `settings.py` file (we did this in *Chapter 1, An Introduction to Django*). As we also mentioned in *Chapter 1, An Introduction to Django*, it is good for apps to be self-contained. By letting each application have a `static` directory, we can continue the self-contained design by also storing app-specific static files inside the app directory.

Before we use `AppDirectoriesFinder`, we will explain a problem that can occur if multiple static files have the same name, and also how to solve this problem.

Static file namespacing

In the prior section, *Static file finders – use during a request*, we discussed serving a file named `logo.png`. This would provide a logo for the `reviews` application. The filename (`logo.png`) could be quite common – you could imagine that if we added a `store` app (for purchasing books), it would also have a logo. Not to mention that third-party Django apps might also want to use a common name such as `logo.png`. The problem we are about to describe could apply to any static file that has a common name, such as `styles.css` or `main.js`.

Let's consider the `reviews` and `stores` examples. We can add a `static` directory in each of these apps. Then, each `static` directory would have a `logo.png` file (although it would be a different logo). The directory structure is shown in *Figure 5.1*:

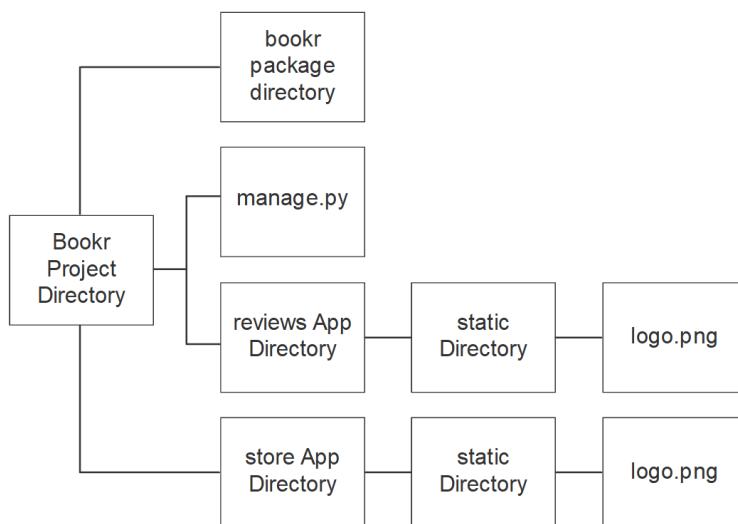


Figure 5.1 – Directory layout with static directories inside app directories

The URL path that we use to download the static file is relative to the `static` directory. Therefore, it is unclear which `logo.png` is being referenced if we make an HTTP request for `/static/logo.png`. Django will check the `static` directory for each application in turn (in the order they are specified in the `INSTALLED_APPS` setting). The first `logo.png` it locates, it will serve. There is no way, in this directory layout, to specify which `logo.png` you want to load.

We can solve this problem by **namespacing** our static files. This is the process of using another directory inside the `static` directory, named the same as the app. The `reviews` app has a `reviews` directory inside its `static` directory, and the `store` app has a `store` directory inside its `static` directory. The respective `logo.png` files are then moved inside these sub-directories. The new directory layout is shown in *Figure 5.2*:



Figure 5.2 – Directory layout with namespaced directories

To load a specific file, we must include the namespaced directory too. For the reviews logo, the URL path is `/static/reviews/logo.png`, which maps to `bookr/reviews/static/reviews/logo.png` on the disk. Similarly, for the store logo, the path is `/static/store/logo.png`, which maps to `bookr/store/static/store/logo.png`. You might have noticed that the example paths for the `logo.png` file were already namespaced in the *Static file finders – use during a request* section.

Note

If you are considering writing a Django app that might be released as a standalone plugin, you could use an even more explicit sub-directory name. For example, you could choose one that contains the entire dotted project path: `bookr/reviews/static/bookr.reviews`. In most cases, though, it's fine for the sub-directory name to be unique for just your project.

Now that we have introduced `AppDirectoriesFinder` and static file namespacing, we can use them to serve our first static file. In the first exercise of this chapter, we will create a new Django project for a basic business site. We will then serve a logo file from an app called `landing` that we will create in this project. The `AppDirectoriesFinder` class is used to find static files inside each app directory, in a directory called `static`. The application must be listed in the `INSTALLED_APPS` setting in your `settings.py` file. As we mentioned in *Chapter 1, An Introduction to Django*, it is good for apps to be self-contained. By letting each application have a `static` directory, we can continue the self-contained design by also storing app-specific static files inside the app directory.

The easiest way to serve a static file is from an app directory. This is because we don't need to make any settings changes. Instead, we just need to create the files in the correct directory, and they will be served using the default Django configuration.

The business site project

For the exercises in this chapter, we'll create a new Django project, and use it to demonstrate the static file concepts. The project will be a basic business site with a simple landing page that has a logo. The project will have one app, called *landing*.

You can refer to *Exercise 1.01 – creating a project and app, and starting the dev server* from *Chapter 1, An Introduction to Django*, to refresh your memory on creating a Django project.

Exercise 5.01 – serving a file from an app directory

In this exercise, you will add a logo file for the *landing* app. This will be done by putting a `logo.png` file in a `static` directory inside the *landing* app directory. Once you've done this, you can test that the static file is being served correctly and confirm the URL that will serve it:

1. Start by creating the new Django project. You can reuse the Bookr virtual environment that already has Django installed.
2. Open a new terminal and activate the virtual environment (refer to the *Preface* for instructions on how to create and activate a virtual environment).
3. Run the `django-admin` command in the terminal (or Command Prompt) to start a Django project named `business_site`:

```
django-admin startproject business_site
```

There won't be any output. This command will scaffold the Django project in a new directory named `business_site`.

4. Create a new Django app in this project by using the `startapp` management command. The app should be called *landing*. To do this, `cd` into the `business_site` directory, then run the following:

```
python3 manage.py startapp landing
```

Note that there won't be any output again. The command will create the `landing` app directory inside the `business_site` directory.

Note

Remember that on Windows, the command is `python manage.py startapp landing`.

5. Launch PyCharm and open the `business_site` directory. If you already have a project open, you can do this by choosing **File | Open**; otherwise, just click **Open** in the **Welcome to PyCharm** window. Navigate to the `business_site` directory, select it, then click **Open**. If prompted, choose **Trust Project**. The `business_site` project window should look similar to what's shown in *Figure 5.3*:

Note

For detailed instructions on how to set up and configure PyCharm to work with your Django project, refer to *Exercise 1.02 – project setup in PyCharm*, in *Chapter 1, An Introduction to Django*.

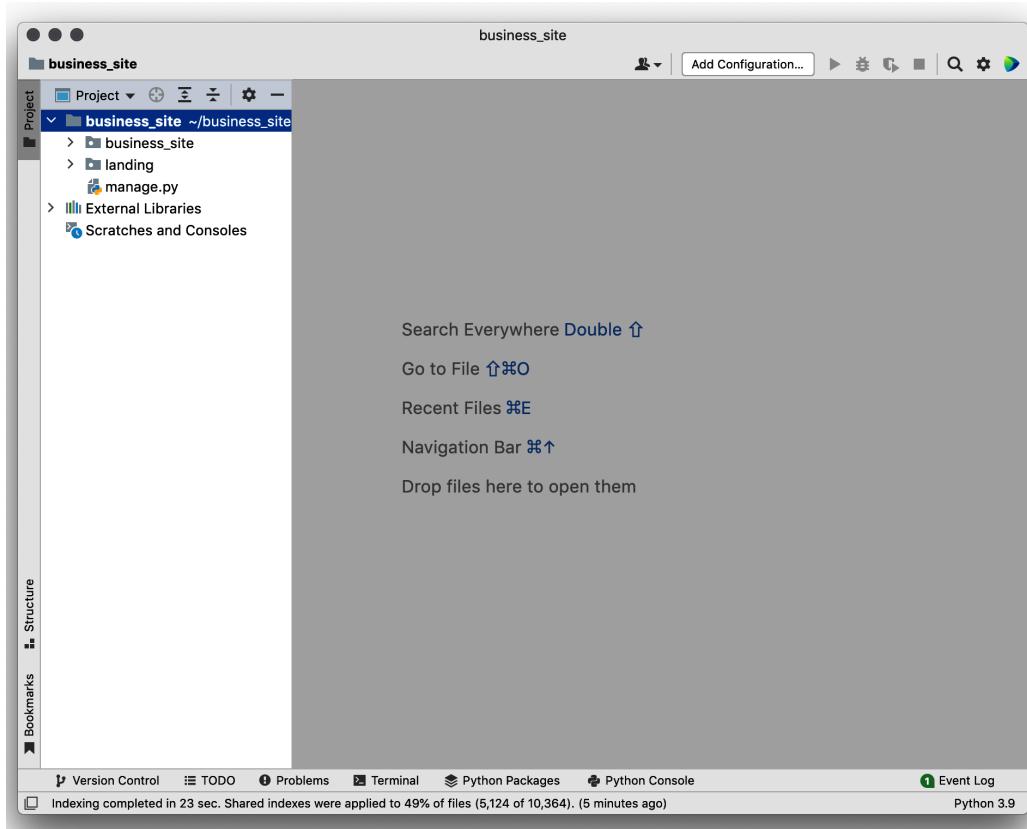


Figure 5.3 – The `business_site` project

6. Create a new run configuration to execute `manage.py runserver` for the project. You can reuse the Bookr virtual environment again. The **Run/Debug Configurations** window should look similar to what's shown in *Figure 5.4* when you're done:

Note

If you are unsure of how to configure these settings in PyCharm, refer to *Exercise 1.02 – project setup in PyCharm*, in *Chapter 1, An Introduction to Django*.

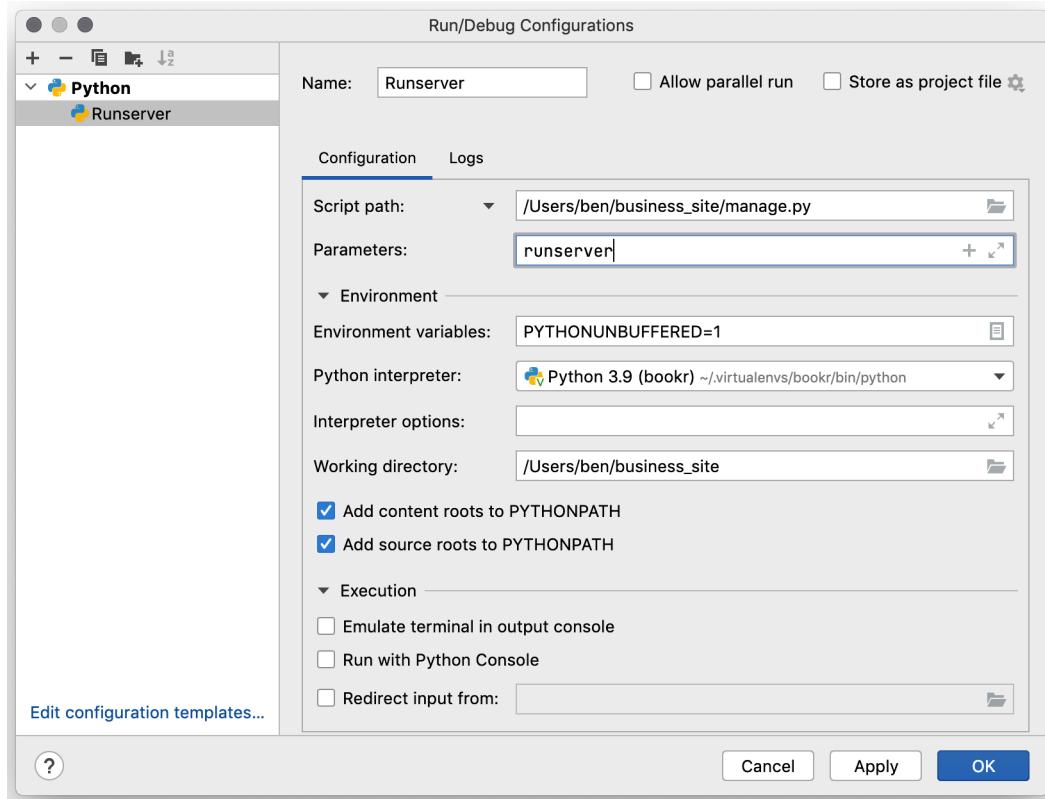


Figure 5.4 – Run/Debug Configurations for Runserver

You can test that the configuration has been set up correctly by clicking the **Run** button, and then visiting `http://127.0.0.1:8000/` in your browser. You should see the Django welcome screen. If the debug server fails to start or you see the Bookr main page, then you probably still have the Bookr project running. Try stopping the Bookr `runserver` process (press `Ctrl + C` in the terminal that's running it) and then starting the new one you just set up.

7. Open `settings.py` in the `business_site` directory and add '`landing`' to the `INSTALLED_APPS` setting. We learned how to do this in *Step 1 of Exercise 1.05 – creating a templates directory and base template*, in *Chapter 1, An Introduction to Django*.
8. In PyCharm, right-click on the `landing` directory in the project pane and select **New | Directory**.

9. Enter `static` and press the *Enter* key:

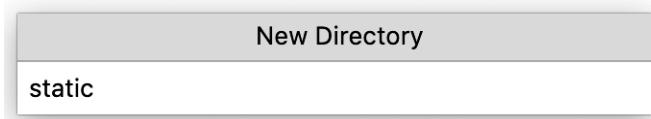


Figure 5.5 – Naming the directory static

10. Right-click on the `static` directory you just created and select **New | Directory** again.
11. Enter `landing` and press *Enter*. This will implement namespacing of the static files directory, as we discussed earlier:

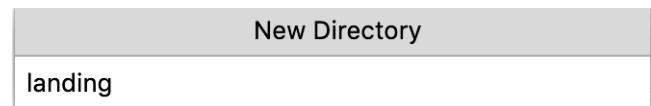


Figure 5.6 – Naming the new directory reviews to implement namespacing

12. Download the `logo.png` file from https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Exercise5.01/business_site/landing/static/landing/logo.png and move it into the `landing/static/landing` directory.
13. Start the Django dev server. If it is not already running, then navigate to `http://127.0.0.1:8000/static/landing/logo.png`. You should see the image being served in your browser:

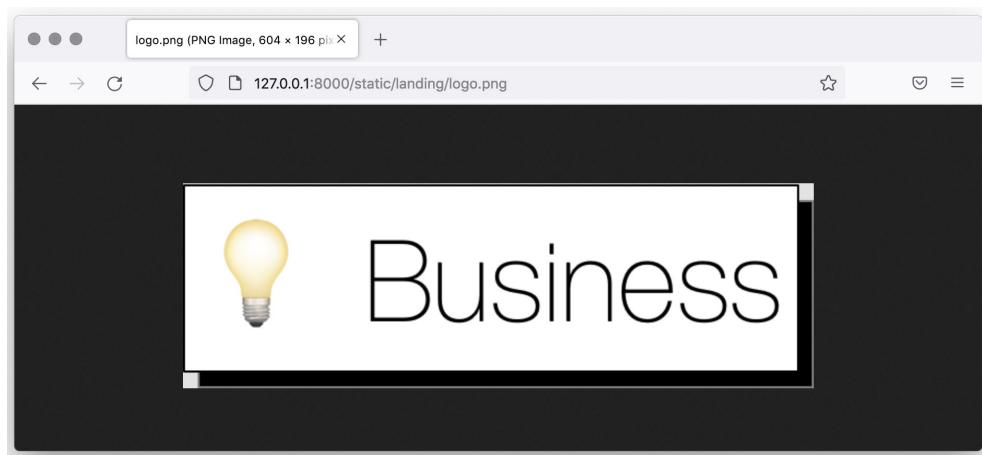


Figure 5.7 – Image served by Django

If you see the image shown in *Figure 5.7*, you have set up static file serving correctly.

Now that you've seen that loading static images works correctly, we'll learn more about how to prevent hardcoding these URLs. In the next section, you'll see how Django can automatically generate correct static URLs for use in templates.

Generating static URLs with the static template tag

In *Exercise 5.01 – serving a file from an app directory*, you set up an image file to be served by Django. We saw that the URL of the image was `http://127.0.0.1:8000/static/landing/logo.png`, which you could use inside an HTML template. For example, to display the image with an `img` tag, you could use this code in your template:

```

```

Or, since Django is also serving the media and has the same host as the dynamic template response, you can simplify this by just including the path, as follows:

```

```

Both addresses (URLs and paths) have been hardcoded into the template – that is, we include the full path to the static file and make assumptions about where the file is being hosted. This works fine with the Django dev server or if you host your static files and Django website on the same domain. For more performance as your site becomes more popular, you might consider serving static files from their own domain or **Content Delivery Network (CDN)**.

Note

A CDN is a service that can host parts or all of your website for you. It provides several web servers and can seamlessly speed up the loading of your website. For example, it might serve files to a user from the server that is geographically closest to them. There are several CDN providers and depending on how they are set up, they might require you to specify a certain domain from which to serve your static files.

Take, for instance, a common separation approach: using a different domain for static file serving. You host your main website at `https://www.example.com` but want to serve static files from `https://static.example.com`. During development, we could use just the path to the logo file, as in the example we just saw. But when we deploy to the production server, our URLs would need to change so that they include the domain, like so:

```

```

Since all the links are hardcoded, this would need to be done for every URL throughout our templates, every time we deploy to production. Once they have been changed, though, the URL will no longer work in the Django dev server. Luckily, Django provides a solution to this problem.

The `staticfiles` app provides a template tag, `static`, for dynamically generating the URL to a static file inside a template. Since the URLs are all being dynamically generated, we can change the URL for all of them by changing just one setting (`STATIC_URL` in `settings.py` – more on this soon). Furthermore, later, we will introduce a method for invalidating browser caches for static files that relies on the use of the `static` template tag.

The `static` tag is very simple – it takes a single argument, which is the project-relative path to a static asset. It will then output this path, prepended by the `STATIC_URL` setting. But first, it must be loaded into the template with the `{% load static %}` template tag.

Django has a set of default template tags and filters (or tag sets) that it automatically makes available to every template. Django (and third-party libraries) also provides tag sets that are not automatically loaded. In these cases, we need to load these extra template tags and filters into a template before we can use them. This can be done with the `load` template tag, which should come near the start of a template (although it must be after the `extends` template tag if one is used). The `load` template tag takes one or more packages/libraries to load; take the following example:

```
{% load package_one package_two package_three %}
```

This would load the template tag and filters set provided by the (made up) `package_one`, `package_two`, and `package_three` packages.

The `load` template tag must be used in the actual template that requires the loaded package. For example, if your template extends another template, and that base template has loaded a certain package, your dependent template does not automatically have access to that package. Your template must still load the package to access the new tag set. The `static` template tag is not part of the default set, which is why we need to load it.

Then, it can be used to interpolate anywhere inside the template file. For example, by default, Django uses `/static/` as `STATIC_URL`. If we wanted to generate the static URL for our `logo.png` file, we would use the tag in a template like this:

```
{% static landing/logo.png %}
```

The output inside the template would be this:

```
/static/landing/logo.png
```

This will become clearer with an example, so let's look at how the `static` tag could be used to generate a URL for several different assets.

We can include the logo as an image on the page with an `img` tag, as follows:

```

```

This is rendered in the template like so:

```

```

Alternatively, we could use the `static` tag to generate the URL for a linked CSS file, as follows:

```
<link href="{% static 'path/to/file.css' %}" rel="stylesheet">
```

This will be rendered like so:

```
<link href="/static/path/to/file.css" rel="stylesheet">
```

It can be used in a `script` tag to include a JavaScript file, using the following line of code:

```
in a script tag:<script src="{% static 'path/to/file.js' %}"></script>
```

This is rendered as follows:

```
<script src="/static/path/to/file.js"></script>
```

You can even use it to generate a link to a static file for download, as we've done here:

```
<a href="{% static 'path/to/document.pdf' %}">Download PDF</a>
```

Note

Note that this won't generate the actual PDF content – it will just create a link to an already existing file.

This is rendered as follows:

```
<a href="/static/path/to/document.pdf">Download PDF</a>
```

While referring to these examples, we can now demonstrate the advantage of using the `static` tag instead of hardcoding. When we are ready to deploy to production, we can just change the `STATIC_URL` value in `settings.py`. None of the values in the templates need to be changed.

For example, we can change `STATIC_URL` to `https://static.example.com/`, and then when the page gets rendered next, the examples we've seen will automatically update.

The following line shows this for the image:

```

```

The following is for the CSS link:

```
<link href="https://static.example.com/path/to/files.css" rel="stylesheet">
```

For the script, it's as follows:

```
<script src="https://static.example.com/path/to/file.js"></script>
```

And finally, the following is for the link:

```
<a href="https://static.example.com/path/to/document.pdf">Download PDF</a>
```

Note that in all these examples, a literal string is being passed as an argument (it is quoted). You can also use a variable as an argument – for example, if you were rendering a template with a context, such as in this example code:

```
def view_function(request):
    context = {"image_file": "logofile.png"}
    return render(request, "example.html", context)
```

We are rendering the `example.html` template with a variable called `image_file`. This variable has a value of `logo.png`.

You would pass this variable to the `static` tag without quotes:

```

```

It would render like this (assuming we changed `STATIC_URL` back to `/static/`):

```

```

The template tag can also be used with the `as [variable]` suffix to assign the result to a variable for use later in the template. This can be useful if the static file lookup takes a long time and you want to refer to the same static file multiple times (such as by including an image in multiple places).

The first time you refer to the static URL, give it a variable name to assign to. In this case, we are creating the `logo_path` variable:

```

```

This renders the same as the examples we've seen before:

```

```

However, we can then use the assigned variable (`logo_path`) again later in the template:

```

```

This variable is now just a normal context variable in the template scope and can be used anywhere in the template. Be careful, though, as you might override a variable that has already been defined – although this a general warning when using any of the template tags that assign variables (for example, `{% with %}`).

In the next exercise, we will put the `static` template into practice by adding a template to the `business_site` project, then including the example image.

Exercise 5.02 – using the static template tag

In *Exercise 5.01 – serving a file from an app directory*, you tested serving `logo.png` from the `static` directory. In this exercise, you will continue with the business site project and create an `index.html` file as the template for our landing page. Then, you'll include the logo inside this page using the `{% static %}` template tag:

1. In PyCharm (make sure you're in the `business_site` project), right-click on the landing directory and create a new folder called `templates`.
2. Right-click on the new `templates` directory and select **New | HTML File**. Select **HTML 5 file** and name it `index.html`:

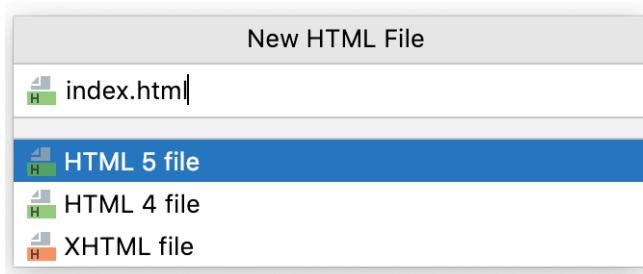


Figure 5.8 – New index.html

3. Open the `index.html` file and first, load the static tag library to make the `static` tag available in the template. Do this with the `load` template tag. On the second line of the file (just after `<!DOCTYPE html>`), add this line to load the static library:

```
{% load static %}
```

4. You can also make the template a bit nicer with some extra content. Enter the Business Site text inside the `<title>` tags:

```
<title>Business Site</title>
```

Then, inside the body, add an `<h1>` element with the Welcome to my Business Site text:

```
<h1>Welcome to my Business Site</h1>
```

5. Underneath the heading text, use the `{% static %}` template tag to set the source of an ``. You will use it to refer to the logo from *Exercise 5.01 – serving a file from an app directory*:

```

```

6. Finally, to flesh out the site a bit, add a `<p>` element under ``. Give it some text about the business:

```
<p>Welcome to the site for my Business. For all your Business needs!</p>
```

Although the extra text and title are not too important, they give us an idea of how to use the `{% static %}` template tag around the rest of the content. Save the file. It should look like this once complete: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Exercise5.02/business_site/landing/templates/index.html.

7. Next, set up a URL to use to render the template. You will also use the built-in `TemplateView` to render the template without having to create a view. Open `urls.py` in the `business_site` package directory. At the start of the file, import `TemplateView` as follows:

```
from django.views.generic import TemplateView
```

You can also remove this Django admin import line since we're not using it in this project:

```
from django.contrib import admin
```

8. Add a URL map from `/` to `TemplateView`. The `as_view` method of `TemplateView` takes `template_name` as an argument, which is used in the same way as a path that you might pass to the `render` function. Your `urlpatterns` should look like this:

```
urlpatterns = [
    path("", TemplateView.as_view(template_name="index.html")),
]
```

Save the `urls.py` file. Once complete, it should look like this: `https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Exercise5.02/business_site/business_site/urls.py`.

9. Start the Django dev server, if it's not already running. Navigate to `http://127.0.0.1:8000` in your browser. You should see your new landing page, as shown in *Figure 5.9*:



Figure 5.9 – My business site with a logo

In this exercise, we added a base template for `landing` and loaded the static library into the template. Once the static library was loaded, we were able to use the `static` template tag to load an image. Then, we were able to see our business logo rendered in the browser.

So far, the static file loading has used `AppDirectoriesFinder` because it required no extra configuration to use. In the next section, we will look at `FileSystemFinder`, which is more flexible but requires a small amount of configuration to use.

FileSystemFinder

So far, we've learned about `AppDirectoriesFinder`, which loads static files inside Django app directories. However, we expect well-designed apps to be self-contained, and therefore they should only contain static files that they rely on. If we have other static files that are used throughout the website or across different apps, we should store them outside the app directory.

Note

As a general rule, your CSS is probably consistent throughout your site and could be kept in a global directory. Some images and JavaScript code could be specific to apps, so these would be stored in the `static` directory for that application. This is just general advice, though: you can store static files anywhere that makes the most sense for your project.

In our business site application, we will be storing a CSS file in a site `static` directory as it will be used not only in the `landing` app but also throughout the site as we add more apps.

Django provides support for serving static files from arbitrary directories using its `FileSystemFinder` static file finder. The directories can be anywhere on the disk. Usually, you will have a `static` directory inside your project directory, but if your company has a global `static` directory that is used in many different projects (including non-Django web applications), then you could use this as well.

`FileSystemFinder` uses the `STATICFILES_DIRS` setting in the `settings.py` file to determine which directories to search for static files in. This is not present when the project is created and must be set by the developer. We will add it in the next exercise. There are two options for building this list:

- Setting a list of directories
- Setting a list of tuples in the form of `(prefix, directory)`

The second use case will be easier to understand once we have covered some more of the fundamentals, so we will return to it after explaining and demonstrating the first case. This will be covered in the `STATICFILES_DIRS` *prefixed mode* section. For now, we will just explain the first use case, which is just a list of one or more directories.

In `business_site`, we will add a `static` directory inside the project directory (that is, in the same directory that contains the `landing` app and the `manage.py` file). We can use the `BASE_DIR` setting when building the list to assign to `STATICFILES_DIRS`:

```
STATICFILES_DIRS = [BASE_DIR / "static"]
```

We also mentioned earlier in this section that you might want to set multiple directory paths in this list, for example, if you had some company-wide static data shared by multiple web projects. Simply add extra directories to the `STATICFILES_DIRS` list:

```
STATICFILES_DIRS = [BASE_DIR / "static",
    "/Users/username/projects/company-static/"
]
```

Each of these directories would be checked, in the order specified, to find a matching file. If a file existed in both directories, the first one found would be served. For example, if the `static/main.css` (inside the `business_site` project directory) and `/Users/username/projects/company-static/bar/main.css` files both existed, a request for `/static/main.css` would serve the `business_site` project's `main.css` as it is first in the list. Consider this when deciding the order in which you add directories to `STATICFILES_DIRS`; you may choose to prioritize your project static files over the global ones or vice versa.

In our `business_site` (and later with `Bookr`), we will only use one `static` directory in this list, so we won't have to worry about this problem.

In the next exercise, we will add a `static` directory with a CSS file inside. Then, we will configure the `STATICFILES_DIRS` setting so that it can be served from the `static` directory.

Exercise 5.03 – serving from a project static directory

We showed an example of serving an application-specific image file in *Exercise 5.01 – serving a file from an app directory*. Now, we want to serve a CSS file that is to be used throughout our project to set styles, so we will serve this from a `static` directory right inside the project folder.

In this exercise, you'll set up your project to serve static files from a specific directory, and then use the `{% static %}` template tag again to include it in the template. This will be done using the `business_site` example project:

1. Open the `business_site` project in PyCharm if it's not already open. Then, right-click on the `business_site` project directory (the top-level `business_site` directory, not the `business_site` package directory) and select **New | Directory**.
2. In the **New Directory** dialog, enter `static` and click **OK**.
3. Right-click on the `static` directory you just created and select **New | File**.
4. In the **Name New File** dialog, enter `main.css` and click **OK**.

5. The blank `main.css` file should open automatically. Enter a couple of simple CSS rules to center the text and set a font and background color, like so:

```
body {  
    font-family: Arial, sans-serif;  
    text-align: center;  
    background-color: #f0f0f0;  
}
```

You can now save and class `main.css`. You can take a look at the complete file for reference at https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Exercise5.03/business_site/static/main.css.

6. Open `business_site/settings.py`. Here, set a list of directories to the `STATICFILES_DIRS` settings. In this case, the list will have just one item. Define a new variable called `STATICFILES_DIRS` at the bottom of the `settings.py` file using the following code:

```
STATICFILES_DIRS = [BASE_DIR / "static"]
```

In the `settings.py` file, `BASE_DIR` is a variable that contains the path to the project directory. You can build the full path to the `static` directory you created in *Step 2* by joining `static` to `BASE_DIR`. Then, you can put this inside a list. The complete `settings.py` file should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Exercise5.03/business_site/business_site/settings.py.

7. Start the Django dev server if it is not running. You can verify that the settings are correct by checking whether you can load the `main.css` file. Note that this is not namespaced, so the URL is `http://127.0.0.1:8000/static/main.css`. Open this URL in your browser and check that the content matches what you just entered and saved:



Figure 5.10 – CSS served by Django

If the file does not load, check your `STATICFILES_DIRS` settings. You may need to restart the Django dev server if it was running while you made changes to `settings.py`.

8. Now, you need to include `main.css` in your `index.html` template. Open `index.html` in the `templates` folder. Before the closing `</head>` tag, add this `<link>` tag to load the CSS:

```
<link rel="stylesheet" href="{% static 'main.css' %}">
```

This links in the `main.css` file, using the `{% static %}` template tag. As mentioned earlier, since `main.css` is not namespaced, you can just include its name. Save the file; it should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Exercise5.03/business_site/landing/templates/index.html.

9. Load `http://127.0.0.1:8000/` in your browser; you should see the background color, fonts, and alignment all change:



Figure 5.11 – CSS applied with custom fonts visible

Your business landing page should look like what's shown in *Figure 5.11*. Since you included the CSS in the `index.html` template, it will be available in all the templates that extend this template (although none do at the moment; it's good planning for the future).

In this exercise, we put some CSS rules into a file and served them using Django's `FileSystemFinder`. This was accomplished by creating a `static` directory inside the `business_site` project directory and specifying it in the Django settings (the `settings.py` file) using the `STATICFILES_DIRS` setting. We linked the `main.css` file to the `index.html` template using the `static` template tag. We loaded the main page in our browser and saw that the font and color changes applied.

We've now covered how static file finders are used during a request (to load a specific static file when given a URL). We'll now look at their other use case: finding and copying static files for production deployment when running the `collectstatic` management command.

Static file finders – use during collectstatic

Once we have finished working on our static files, they need to be moved into a specific directory that can be served by our production web server. We can then deploy our website by copying our Django code and static files to our production web server. In the case of `business_site`, we will want to move `logo.png` and `main.css` (along with other static files that Django itself includes) into a single directory that can be copied to the production web server. This is the role of the `collectstatic` management command.

We have already discussed how Django uses static file finders during request handling. Now, we will cover the other use case: collecting static files for deployment. Upon running the `collectstatic` management command, Django uses each finder to list static files on the disk. Every static file that is found is then copied into the `STATIC_ROOT` directory (also defined in `settings.py`). This is a little bit like the reverse of handling a request. Instead of getting a URL path and mapping it to a filesystem path, the filesystem path is copied to a location that can be predicted by the frontend web server. This allows the frontend web server to handle a request for a static file independently of Django.

Note

A frontend web server is a web server that's designed to route requests to applications (such as Django) or read static files from disk. They can handle requests faster but aren't able to generate dynamic content in the same way as something such as Django. Frontend web servers are software such as Apache HTTPD, Nginx, and lighttpd.

For some specific examples of how `collectstatic` works, we'll use the two files from *Exercise 5.01 – serving a file from an app directory* and *Exercise 5.03 – serving from a project status directory*, respectively: `landing/static/landing/logo.png` and `static/main.css`.

Let's assume that `STATIC_ROOT` has been set to a directory being served by a normal web server – this would be something such as `/var/www/business_site/static`. The destination for these files would be `/var/www/business_site/static/reviews/logo.png` and `/var/www/business_site/static/main.css`, respectively.

Now, when a request for a static file comes in, the web server will be able to serve it easily because the paths are mapped consistently:

- `/static/main.css` is served from the `/var/www/business_site/static/main.css` file
- `/static/reviews/logo.png` is served from the `/var/www/business_site/static/reviews/logo.png` file

This means the web server root is `/var/www/business_site/` and static paths are loaded directly from disk in the usual manner that a web server would load files.

With that, we have demonstrated how Django locates static files during development and can serve them itself. In production, we need the frontend web server to be able to serve static files without involving Django, for both safety and speed.

Without having to run `collectstatic`, a web server would not be able to map a URL back to a path. For example, it would not know that `main.css` must be loaded from the project static directory while `logo.png` is to be loaded from the landing app directory – it has no concept of the Django directory layout.

You might be tempted to serve files directly from the Django project directory by setting your web server root to this directory – **don't do this**. There is a **security risk** in sharing your entire Django project directory as it would make it possible to download `settings.py` or other sensitive files. Running `collectstatic` will copy the files to a directory that can be moved outside the Django project directory to the web server root for security.

So far, we have talked about using `collectstatic` to copy static files directly to the web server root. You could also have Django copy them to an intermediary directory and have your deployment process move to a CDN or another server afterward. We won't go into detail on specific deployment processes; how you choose to copy static files to the web server will depend on your or your company's existing setup (for example, a continuous delivery pipeline).

Note

The `collectstatic` command does not take into consideration the use of `static` template tags. It will collect all the static files inside static directories, even those that your project does not include inside a template.

In the next exercise, we will see the `collectstatic` command in action. We'll use it to copy all the `business_site` static files that we have so far into a temporary directory.

Exercise 5.04 – collecting static files for production

While we won't be covering deployment to a web server in this chapter, we can still use the `collectstatic` management command and see its result. In this exercise, we will create a temporary holding location for the static files to be copied into. This directory will be called `static_production_test` and will be located inside the `business_site` project directory. As part of the deployment process, you could copy this directory to your production web server. However, since we won't be setting up a web server until *Chapter 17, Deploying a Django Application (Part 1 – Server Setup)*, we will just examine its contents to understand how files are copied and organized:

1. In PyCharm, create a temporary directory to put the collected files in. Right-click on the `business_site` project directory (this is the top-level folder, not the `business_site` module) and select **New | Directory**.
2. In the **New Directory** dialog, enter `static_production_test` and click **OK**.
3. Open `settings.py` and, at the bottom of the file, define a new setting for `STATIC_ROOT`. Set it to the path of the directory you just created:

```
STATIC_ROOT = BASE_DIR / "static_production_test"
```

This will join `static_dir` to `BASE_DIR` (the `business_site` project path) to generate the full path. Save the `settings.py` file. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Exercise5.04/business_site/business_site/settings.py.

4. In a terminal, run the `collectstatic` manage command:

```
python3 manage.py collectstatic
```

You should see an output like the following:

```
130 static files copied to '/Users/ben/business_site/static_production_test'.
```

This might seem like a lot if you were expecting it to copy just two files, but remember that it will copy all the files for all installed apps. In this case, as you have the Django admin app installed, most out of the 132 files support that.

5. Let's look through the `static_production_test` directory to check what has been created. An expanded view of this directory (from the PyCharm project page) is shown in *Figure 5.12*, for reference. Yours should be similar:



Figure 5.12 – Destination directory of the `collectstatic` command

You should notice three items inside it:

- The `admin` directory: This contains files from the Django admin app. If you look inside this, you'll see it has been organized into subfolders for `css`, `fonts`, `img`, and `js`.
- The `landing` directory: This is the static directory from your landing app. Inside it is the `logo.png` file. This directory has been created to match the namespacing of the directory that we created.
- The `main.css` file: This is from your project static directory. Since you didn't place it inside a namespacing directory, this has been placed directly inside `STATIC_ROOT`.

If you want, you can open any of these files and verify that their content matches the files you have just been working on – they should do as they are simply copies of the original files.

In this exercise, we collected all the static files from `business_site` (including the `admin` static files that Django includes). They were copied into the directory defined by the `STATIC_ROOT` setting (`static_production_test` inside the `business_site` project directory). We saw that `main.css` was directly inside this folder but that other static files were namespaced inside their app directories (`admin` and `reviews`). This folder could have been copied to a production web server to deploy our project.

In the next section, we'll look at how a prefix can be added to each static file directory source to customize how static files are organized after being collected.

STATICFILES_DIRS prefixed mode

As mentioned earlier, the `STATICFILES_DIRS` setting also accepts items as tuples in the form of `(prefix, directory)`. These modes of operation are not mutually exclusive; `STATICFILES_DIRS` may contain both non-prefixed (string) or prefixed (tuple) items. Essentially, this allows you to map a certain URL prefix to a directory. In Bookr, we do not have enough static assets to warrant setting this up, but it can be useful if you want to organize your static assets differently. For example, you can keep all your images in a certain directory, and all your CSS in another directory. You might need to do this if you use a third-party CSS generation tool such as Node.js with LESS.

Note

LESS is a CSS pre-processor that uses Node.js. It allows you to write CSS using variables and other programming-like concepts that don't exist natively. Node.js will then compile this to CSS. A more in-depth explanation is outside the scope of this book; suffice it to say that if you use it (or a similar tool), then you might want to serve directly from the directory to which it saves its compiled output.

The easiest way to explain how prefixed mode works is with a short example. This will expand on the STATICFILES_DIRS setting we created in *Exercise 5.03 – serving from a project static directory*; however, this is not an exercise. In this example, two prefixed directories are being added to this setting – one for serving images and one for serving CSS:

```
STATICFILES_DIRS = [
    BASE_DIR / "static",
    ("images", BASE_DIR / "static_images"),
    ("css", BASE_DIR / "static_css"),
]
```

As well as the `static` directory that was already being served with no prefix, we have added serving of the `static_images` directory inside the `business_site` project directory. This has a prefix of `images`. We have also added serving for the `static_css` directory inside the Bookr project directory, with a prefix of `css`.

Then, we can serve three files, `main.js`, `main.css`, and `main.jpg`, from the `static`, `static_css`, and `static_images` directories, respectively. The directory layout will be as shown in *Figure 5.13*:

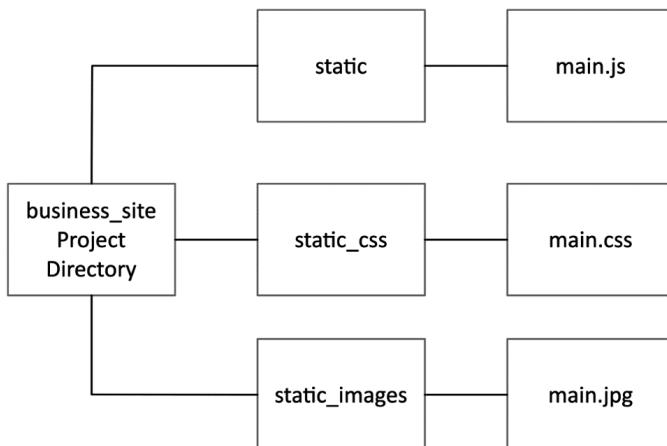


Figure 5.13 – Directory layout for use with prefixed URLs

In terms of accessing these via a URL, the mapping is as shown in *Figure 5.14*:

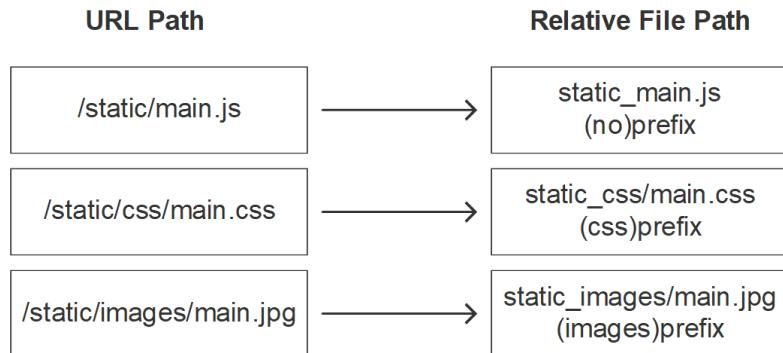


Figure 5.14 – Mapping a URL to a file based on the prefix

Django routes any static URL with a prefix to the directory that matches that prefix.

When used with the `static` template tag, use the prefix and filename, not the directory name; take the following example:

```
{% static 'images/main.jpg' %}
```

When the static files are gathered using the `collectstatic` command, they are moved into a directory with the prefix name, inside `STATIC_ROOT`. The source paths and the target paths inside the `STATIC_ROOT` directory are shown in *Figure 5.15*:

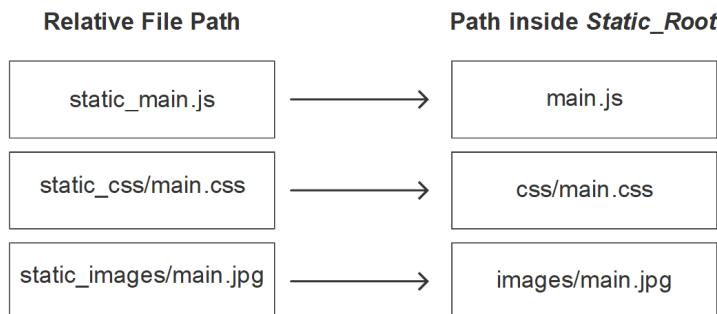


Figure 5.15 – Mapping from the path in the project directory to the path in `STATIC_ROOT`

Django creates the prefix directories inside `STATIC_ROOT`. Because of this, the paths can be kept consistent, even when using a web server and not routing the URL lookup through Django.

Next, we'll look at a management command that's useful for troubleshooting when static files aren't being loaded as you expect: `findstatic`.

The `findstatic` command

The `staticfiles` application also provides one more management command: `findstatic`. This command allows you to enter the relative path to a static file (the same as what would be used inside a `static` template tag) and Django will tell you where that file was located. It can also be used in verbose mode to output the directories it is searching through.

Note

You may not be familiar with the concept of verbosity, or verbose mode. Having a higher verbosity (or simply turning on verbose mode) will cause a command to generate more output. Many command-line applications can be executed with more or less verbosity. This can be helpful when you're trying to debug programs you are using. To see an example of verbose mode in action, you can try running the Python shell in verbose mode. Enter `python -v` (instead of just `python`) and press *Enter*. Python will start in verbose mode, which prints out the path of every file it imports.

This command is mostly useful for debugging/troubleshooting purposes. If the wrong file is loading, or a particular file can't be found, you can use this command to try to find out why. The command will display which file on disk is being loaded for a specific path, or let you know that the file can't be found and what directories were searched.

This can help solve issues where multiple files have the same name, and the precedence is not how you expect. You may also see that Django is not searching in a directory you expect for the file, in which case the static directory might need to be added to the `STATICFILES_DIRS` setting.

In the next exercise, you will execute the `findstatic` management command so that you are familiar with what the outputs are for good (file found correctly) and bad (file missing) scenarios.

Exercise 5.05 – finding files using `findstatic`

You will now run the `findstatic` command with a variety of options and understand what its output means. First, we will use it to find a file that exists and see that it displays the path to the file. Then, we will try to find a file that doesn't exist and check the error that is output. We will then repeat this process with multiple levels of verbosity and different ways of interacting with the command. While this exercise will not make changes to or progress the Bookr project, it is good to be familiar with the command in case you need to use it when working on your own Django applications. Let's get started:

1. Open a terminal and navigate to the `business_site` project directory.
2. Execute the `findstatic` command with no options. It will output some help explaining how it is used:

```
python3 manage.py findstatic
```

The following help output will be displayed:

```
usage: manage.py findstatic [-h] [--first] [--version] [-v
{0,1,2,3}] [--settings SETTINGS] [--pythonpath PYTHONPATH]
[--traceback] [--no-color] [--force-color]
[--skip-checks] staticfile [staticfile ...]
manage.py findstatic: error: Enter at least one label.
```

3. You can find one or more files at a time. Let's start with one that we know exists: `main.css`. We'll use the `findstatic` management command to locate it on disk by entering the following command:

```
python3 manage.py findstatic main.css
```

The preceding command outputs the path at which `main.css` was found:

```
Found 'main.css' here:
/Users/ben/business_site/static/main.css
```

Your full path will be different (unless you are also called Ben), but you can see that when Django locates `main.css` in a request, it will load the `main.css` file from the project's `static` directory.

This can be useful if a third-party application you have installed has not namespaced its static files correctly and is conflicting with one of your files.

4. Let's try finding a file that does not exist, `logo.png`:

```
python3 manage.py findstatic logo.png
```

Django will display an error stating that the file could not be found:

```
No matching file found for 'logo.png'.
```

Django is unable to locate this file because we have namespaced it – we must include the full relative path, the same way we used it in the `static` template tag.

5. Try finding the `logo.png` again, but this time using the full path:

```
python3 manage.py findstatic landing/logo.png
```

Django can find the file now:

```
Found 'landing/logo.png' here:  
/Users/ben/business_site/landing/static/landing/logo  
.png
```

6. You can find multiple files at once by adding each file as an argument:

```
python3 manage.py findstatic landing/logo.png  
missing-file.js main.css
```

The location status for each file is shown as follows:

```
No matching file found for 'missing-file.js'.  
Found 'landing/logo.png' here:  
/Users/ben/business_site/landing/static/landing/logo  
.png  
Found 'main.css' here:  
/Users/ben/business_site/static/main.css
```

7. The command can be executed with a verbosity of 0, 1, or 2. By default, it executes at verbosity 1. To set the verbosity, use the `--verbosity` or `-v` flag. Decreasing the verbosity to 0 only outputs the paths it locates without any extra information. No errors are displayed for missing paths:

```
python3 manage.py findstatic -v0 landing/logo.png  
missing-file.js main.css
```

The output shows only found paths – notice that no error is shown for the missing file, `missing-file.js`:

```
/Users/ben/business_site/landing/static/landing/  
logo.png  
/Users/ben/business_site/static/main.css
```

This level of verbosity can be useful if you are piping the output to another file or command.

8. To get more information about which directories Django is searching in for the file you have requested, increase the verbosity to 2:

```
python3 manage.py findstatic -v2 landing/logo.png
missing-file.js main.css
```

The output contains much more information, including the directories that have been searched for the requested file. You can see that as the `admin` application is installed, Django is also searching in the Django admin application directory for static files:

```
(bookr) → business_site python3 manage.py findstatic -v2 landing/logo.png missing-file.js main.css
No matching file found for 'missing-file.js'.
```

```
Looking in the following locations:
/Users/ben/business_site/static
/Users/ben/.virtualenvs/bookr/lib/python3.9/site-packages/django/contrib/admin/static
/Users/ben/business_site/landing/static
Found 'landing/logo.png' here:
/Users/ben/business_site/landing/static/landing/logo.png
Looking in the following locations:
/Users/ben/business_site/static
/Users/ben/.virtualenvs/bookr/lib/python3.9/site-packages/django/contrib/admin/static
/Users/ben/business_site/landing/static
Found 'main.css' here:
/Users/ben/business_site/static/main.css
Looking in the following locations:
/Users/ben/business_site/static
/Users/ben/.virtualenvs/bookr/lib/python3.9/site-packages/django/contrib/admin/static
/Users/ben/business_site/landing/static
```

Figure 5.16 – `findstatic` executed with verbosity 2, showing exactly which directories were searched

The `findstatic` command is not something that you will use day to day when working with Django, but it is useful to know about when you're trying to troubleshoot problems with static files. We saw the command's output, the full path to a file that existed, as well as the error messages when files did not exist. We also ran the command and supplied multiple files at once and saw that information about all the files was outputted. Finally, we ran the command with different levels of verbosity. The `-v0` flag suppressed errors about missing files. `-v1` was the default and displayed found paths and errors. Increasing the verbosity using the `-v2` flag also printed out the directories that were being searched through for a particular static file.

In the next section, we'll take a look at caching static files and how Django helps us ensure our users are served the latest versions.

Serving the latest files (for cache invalidation)

If you are not familiar with caching, the basic idea is that some operations can take a long time to perform. We can speed up a system by storing the results of the operation in a place that is faster to access so that the next time we need them, they can be retrieved quickly. The “operation” that takes a long time can be anything – a function that takes a long time to run, an image that takes a long time to render, or a large asset that takes a long time to download over the internet. We are interested in this last scenario.

You might have noticed that the first time you ever visit a particular website, it can be slow to load, but then the next time, it loads much faster. This is because your browser has cached some (or all) of the static files the site needs to load.

To use our business site as an example, we have a page that includes the `logo.png` file. The first time we visit the business site, we must download the dynamic HTML, which is small and quick to transfer. Our browser parses the HTML and sees that the `logo.png` file should be included. It can then download this file too, which is much larger and can take longer to download. Note that this scenario assumes that our business site is now hosted on a remote server and not on our local machine – which is very fast for us to access.

If the web server has been set up correctly, your browser will store `logo.png` on your computer. The next time we visit the *landing* page (or indeed any page that includes the `logo.png` file), our browser will recognize the URL and can load the file from disk instead of having to download it again, thus speeding up the browsing experience.

Note

We said that the browser will cache “if the web server is set up correctly.” What does this mean? The frontend web server should be configured to send special HTTP headers as part of a static file response. It can send a `Cache-Control` header, which can have values such as `no-cache` (the file should never be cached – that is, the latest version should be requested every time) or `max-age=seconds` (the file should only be downloaded again if it was last retrieved more than `seconds` ago). The response could also contain the `Expires` header, with the value being a date. The file is considered to be “stale” once this date has been reached, and at that point, the new version should be requested.

One of the hardest problems in computer science is cache invalidation. For instance, if we change `logo.png`, how does our browser know it should download the new version? The only surefire way of knowing it had changed would be to download the file again and compare it with the version we had already saved every time. Of course, this defeats the purpose of caching since we would still be downloading every time the file changed (or not). We can cache for an arbitrary or server-specified amount of time, but if the static file changed before that time was up, we wouldn’t know. We would use the old version until we considered it expired, at which time we’d download the new version.

If we had a one-week expiry and the static file changed the next day, we'd still be using the old one for six days. Of course, the browser can be made to reload the page without using the cache (how this is done depends on the browser; for example, *Shift + F5* or *Cmd + Shift + R*) if you want to forcefully download all static assets again.

There is no need to try to cache our dynamic responses (rendered templates). Since they are designed to be dynamic, we would want to make sure that the user gets the latest version on every page load, so they should not be cached. They are also quite a small size (compared to assets such as images), so there is not much speed advantage when caching them.

Django provides a built-in solution. During the `collectstatic` phase, when the files are copied, Django can append a hash of their content to the filename. For example, the source file, `logo.png`, will be copied to `static_production_test/landing/logo.f30ba08c60ba.png`. This is only done when using the `ManifestFilesStorage` storage engine.

Since the filename only changes when the content changes, the browser will always download the new content.

Using `ManifestFilesStorage` is just one way of invalidating caches. There may be other options that are more suitable for your particular application.

Note

A hash is a one-way function that generates a string of a fixed length, regardless of the length of the input. There are several different hash functions available, and Django uses **MD5** for content hashing. While no longer cryptographically secure, it is adequate for this purpose. To illustrate the fixed-length property, the MD5 hash of the `a` string is `0cc175b9c0f1b6a831c399e269772661`. The MD5 hash of the `a much longer string` string is `69fc4316c18cdd594a58ec2d59462b97`. They are both 32 characters long.

You can choose the storage engine by changing the `STATICFILES_STORAGE` value in `settings.py`. This is a string with a dotted path to the module and class to use. The class that implements the hash-addition functionality is `django.contrib.staticfiles.storage.ManifestStaticFilesStorage`.

Using this storage engine doesn't require you to make any changes to your HTML templates, provided you are including static assets with the `static` template tag. Django generates a manifest file (`staticfiles.json`, in JSON format) that contains a mapping between the original filename and the hashed filename. It will automatically insert the hashed filename when using the `static` template tag. If you are including your static files without using the `static` tag and instead just manually inserting the static URL, then your browser will attempt to load the non-hashed path and the URL will not automatically be updated when the cache should be invalidated.

For example, you can include `logo.png` with the `static` tag:

```

```

When the page is rendered, the latest hash will be retrieved from `staticfiles.json` and the output will be like this:

```

```

Whereas, if we had not used the `static` tag and instead hardcoded the path, it would always appear as written:

```

```

Since this does not contain a hash, our browser will not see the path changing and thus never attempt to download the new file.

Django retains the previous version of files with the old hash when running `collectstatic`, so older versions of your application can still refer to it if they need to. The latest version of the file is also copied with no hash so that non-Django applications can refer to it without needing to look up the hash.

In the next exercise, we will change our project settings to use the `ManifestFilesStorage` engine, then run the `collectstatic` management command. This will copy all the static assets as in *Exercise 5.04 – collecting static files for production*; however, they will now have their hash included in the filename.

Exercise 5.06 – exploring the `ManifestFilesStorage` storage engine

In this exercise, you will temporarily update `settings.py` to use `ManifestFilesStorage`, then run `collectstatic` to see how the files are generated with a hash:

1. In PyCharm (still in the `business_site` project), open `settings.py` and add a `STATICFILES_STORAGE` setting at the bottom of the file:

```
STATICFILES_STORAGE = "django.contrib.staticfiles.storage.  
ManifestStaticFilesStorage"
```

The completed file should look like this: https://github.com/PacktWorkshops/The-Django-Workshop/blob/master/Chapter05/Exercise5.06/business_site/business_site/settings.py.

2. Open a terminal, navigate to the `business_site` project directory, and run the `collectstatic` command like before:

```
python3 manage.py collectstatic
```

If your `static_production_test` directory is not empty (which will probably be the case as files were moved there in *Exercise 5.04 – collecting static files for production*), then you will be prompted to allow the overwriting of the existing files:

```
(bookr) ➔ business_site python3 manage.py collectstatic
```

```
You have requested to collect static files at the destination
location as specified in your settings:
```

```
/Users/ben/business_site/static_production_test
```

```
This will overwrite existing files!
Are you sure you want to do this?
```

```
Type 'yes' to continue, or 'no' to cancel:
```

Figure 5.17 – Prompt to allow overwriting during static collection

Just type `yes` and then press `Enter` to allow the overwrite.

The output from this command will tell you the number of files that were copied, as well as the number that were processed and had the hash added to the filename:

```
0 static files copied to '/Users/ben/business_site/static_
production_test', 130 unmodified, 98 post-processed.
```

Since you haven't changed any files since we last ran `collectstatic`, no files have been copied. Instead, Django post-processes the files (28 of them) – that is, it generates their hashes and appends them to the filename.

3. The static files were copied into the `static_production_test` directory as they were before; however, there are now two copies of each file: one named with the hash and one without. For example, `static/main.css` has been copied to `static_production_test/main.856c74fb7029.css` (this filename might be different if your CSS file contents differ; for example, if it has extra spaces or newlines):

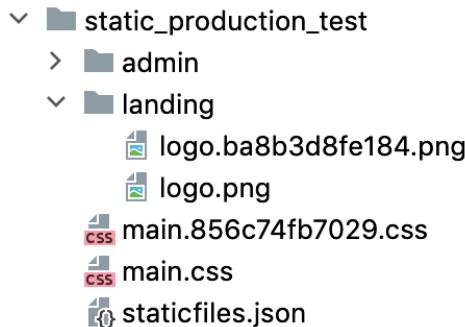


Figure 5.18 – Expanded static_production_test directory with hashed filenames

Figure 5.18 shows the expanded static_production_test directory layout. You can see two copies of each file's static file, as well as the staticfiles.json manifest file. To take logo.png as an example, you can see that landing/static/landing/logo.png has been copied to static_production_test/landing/logo.ba8d3d8fe184.png.

Let's make a change to the main.css file and see how the hash changes.

4. Add some blank lines to the end of the file and save it (make sure that you edit the source file and not the one that was already copied to the static_production_test directory). This won't change the effect of the CSS but the change in the file will affect its hash. Re-run the collectstatic command in a terminal:

```
python3 manage.py collectstatic
```

Once again, you may have to enter yes to confirm the overwrite, which will give you the following output:

```
You have requested to collect static files at the destination
location as specified in your settings:
```

```
/Users/ben/business_site/static_production_test
```

```
This will overwrite existing files!
```

```
Are you sure you want to do this?
```

```
Type 'yes' to continue, or 'no' to cancel: yes
```

```
1 static file copied to '/Users/ben/business_site/static_
production_test', 129 unmodified, 98 post-processed.
```

Since only one file was changed, only one static file was copied (`main.css`).

5. Look inside the `static_production_test` directory again. You should see that the old file with the old hash was retained and that a new file with a new hash has been added:



Figure 5.19 – Another `main.css` file with the latest hash was added

In this case, we have `main.856c74fb7029.css` (existing), `main.02b59bcc5cd9.css` (new), and `main.css`. Your hashes may differ.

The `main.css` file (no hash) always contains the newest content – that is, the contents of the `main.02b59bcc5cd9.css` and `main.css` files are identical. During the execution of `collectstatic`, Django will copy the file with a hash, as well as without a hash.

6. Now, examine the `staticfiles.json` file that Django generates. This is the mapping that allows Django to look up the hashed path from the normal path. Open `static_production_test/staticfiles.json`. All the content may appear in one line. If it does, enable text soft wrapping from the **View** menu | **Active Editor** | **Soft Wrap**. Scroll to the end of the file; you should see an entry for the `main.css` file; take the following example:

```
"main.css": "main.02b59bcc5cd9.css"
```

This is how Django can populate the correct URL in a template when using the `static` template tag: by looking up the hashed path in this mapping file.

7. We've finished using `business_site`, which we were just using for testing. You can delete the project or keep it around for reference during the activities.

Note

Unfortunately, we can't examine how the hashed URL is interpolated in the template because, when running in debug mode, Django does not look up the hashed version of the file. As we know, the Django dev server only runs in debug mode, so if we were to turn debug mode off to try to view the hashed interpolation, the Django dev server would not start. You will need to examine this interpolation yourself when going to production and when using a frontend web server.

In this exercise, we configured Django to use `ManifestFilesStorage` for its static file storage by adding the `STATICFILES_STORAGE` setting to `settings.py`. Then, we executed the `collectstatic` command to see how the hashes are generated and added to the filename of the copied files. We saw the `staticfiles.json` manifest file, which stored a lookup from the original path to the hashed path. Finally, we cleaned up the settings and directories that we added in this exercise and *Exercise 5.04 – collecting static files for production*. These were the `STATIC_ROOT` setting, the `STATICFILES_STORAGE` setting, and the `static_production_test` directory.

So far, we've only seen how to locate static files stored on our disk. Django can be customized to locate static files stored in other locations (such as in the cloud). In the next section, we'll briefly look at how to use custom storage engines.

Custom storage engines

In the previous section, we set the storage engine to `ManifestFilesStorage`. This class is provided by Django, but it is also possible to write a custom storage engine. For example, you could write a storage engine that uploads your static files to a CDN, Amazon S3, or Google Cloud bucket when you run `collectstatic`.

Writing a custom storage engine is beyond the scope of this book. Third-party libraries already exist that support uploading to a variety of cloud services. One such library is `Django Storages`, which can be found at <https://django-storages.readthedocs.io/>.

The following code is a short skeleton indicating which methods you should implement to create a custom file storage engine:

```
from django.conf import settings
from django.contrib.staticfiles import storage

class CustomFilesStorage(storage.StaticFilesStorage):
    def __init__(self):
        """The class must be able to be instantiated
        without any arguments.
        Create custom settings in settings.py and read them
        instead."""

```

```
self.setting = settings.CUSTOM_STORAGE_SETTING
```

The class you define (in this example, it's called `CustomFilesStorage`) must be able to be instantiated without any arguments. The `__init__` function must be able to load any settings from global identifiers (in this case, from our Django settings):

```
def delete(self, name):
    """Implement delete of the file from the remote
    service."""
```

The `delete` method should be able to delete the file, specified by the `name` argument, from the remote service:

```
def exists(self, name):
    """Return True if a file with name exists in the
    remote service."""
```

The `exists` method should query the remote service to check whether the file specified by `name` exists. It should return `True` if the file exists, or `False` if it doesn't:

```
def listdir(self, path):
    """List a directory in the remote service. Return
    should be a 2-tuple of lists, the first a list
    of directories, the second a list of files."""
```

The `listdir` method should query the remote service to list the directory at `path`. It should then return a two-tuple list. The first element of this tuple is a list of directories inside `path`, while the second element is a list of files; take the following example:

```
return (["directory1", "directory2"], ["main.css", "document.txt",
"image.jpg"])
```

If `path` contains no directories or no files, then an empty list should be returned for that element. You would return two empty lists if the directory was empty:

```
def size(self, name):
    """Return the size in bytes of the file with
    name."""
```

The `size` method should query the remote service and get the size of the file specified by name:

```
def url(self, name):  
    """Return the URL where the file of with name can  
    be access on the remote service. For example,  
    this might be URL of the file after it has been  
    uploaded to a specific remote host with a  
    specific domain."""
```

The `url` method should determine the URL to access the file specified by name. This can be built by appending name to a specific static hosting URL:

```
def _open(self, name, mode="rb"):  
    """Return a File-like object pointing to file with  
    name. For example, this could be a URL handle  
    for a remote file."""
```

The `_open` method will provide a handle remote file, specified by name. How you implement this will depend on the type of remote service. You might have to download the file and then use a memory buffer (such as an `io.BytesIO` object) to simulate opening the file:

```
def _save(self, name, content):  
    """Write the content for a file with name. In this  
    method you might upload the content to a remote  
    service."""
```

The `_save` method should save content to the remote file at name. The method of implementing this will depend on your remote service. It might transfer the file over SFTP, or upload it to a CDN.

While this example does not implement any transferring to or from a remote service, you can refer to it to get an idea of how to implement a custom storage engine.

After implementing your custom storage engine, you can make it active by setting its dotted module path in the `STATICFILES_STORAGE` setting in `settings.py`.

You'll now work on the first activity in this chapter. In this activity, you'll add a static logo image to the Bookr application. The logo will be displayed in the **Reviews** section of the website.

Activity 5.01 – adding a Reviews logo

The Bookr app should have a logo that is specific to pages in the Reviews app. This will involve adding a base template just for the Reviews app and updating our current Reviews templates to inherit from it. Then, you'll include the Bookr Reviews logo on this base template.

These steps will help you complete this activity:

1. Add a CSS rule to position the logo. Put this rule into the existing `base.html` file, after the `.navbar-brand` rule:

```
.navbar-brand > img {  
    height: 60px;  
}
```

2. Add a `brand block` template tag that inheriting templates can override. Put this inside the `<a>` element with the `navbar-brand` class. The default contents of `block` should be left as *Book Review*.
3. Add a static directory inside the Reviews app, containing a namespaced directory. Download the Reviews logo `.png` file from <https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Activity5.01/bookr/reviews/static/reviews/logo.png> and put it inside this directory.
4. Create a `templates` directory for the Bookr project (inside the Bookr project directory). Then, move the Reviews app's current `base.html` into this directory so that it becomes a base template for the whole project.
5. Add the new `templates` directory's path to the `TEMPLATES ["DIRS"]` setting in `settings.py`.
6. Create another `base.html` template specifically for the Reviews app. Put it inside the Reviews app's `templates` directory. The new template should extend the existing (now global) `base.html` file.
7. The new `base.html` file should override the content of the `brand block`. This block should contain just a `` whose `src` attribute is set using the `{% static %}` template tag. The image source should be the logo we added in *Step 2*.

Refer to the following screenshots to see what your pages should look like after these changes. Note that although you are making changes to the base template, it will not change the layout of the main page:

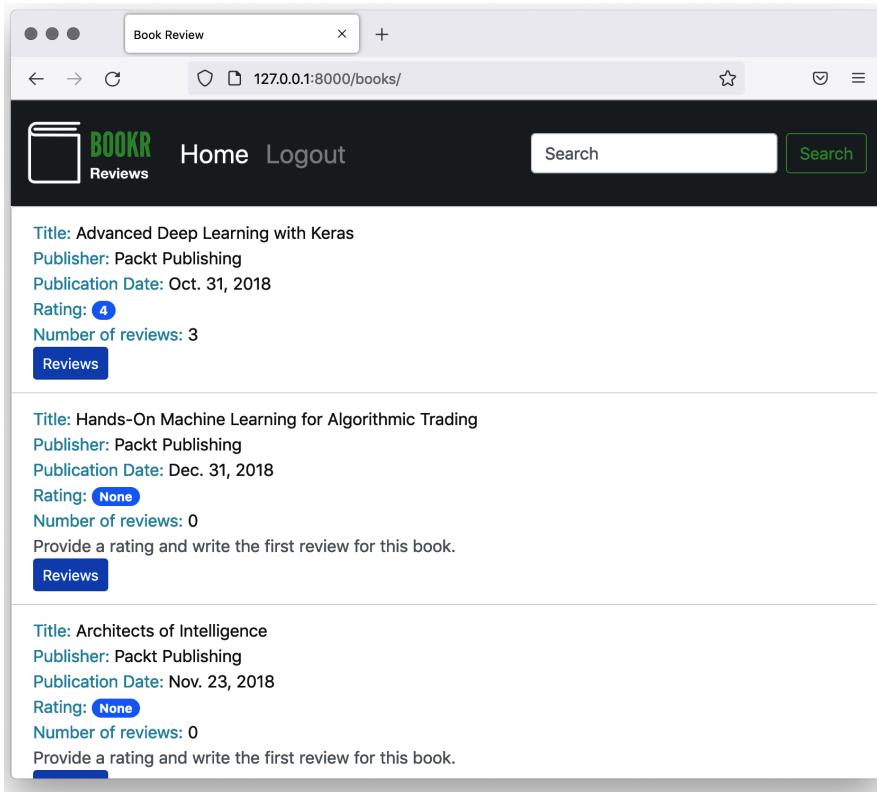


Figure 5.20 – The Book List page after adding the Reviews logo

On the **Book Details** page, you will also see the Reviews logo:

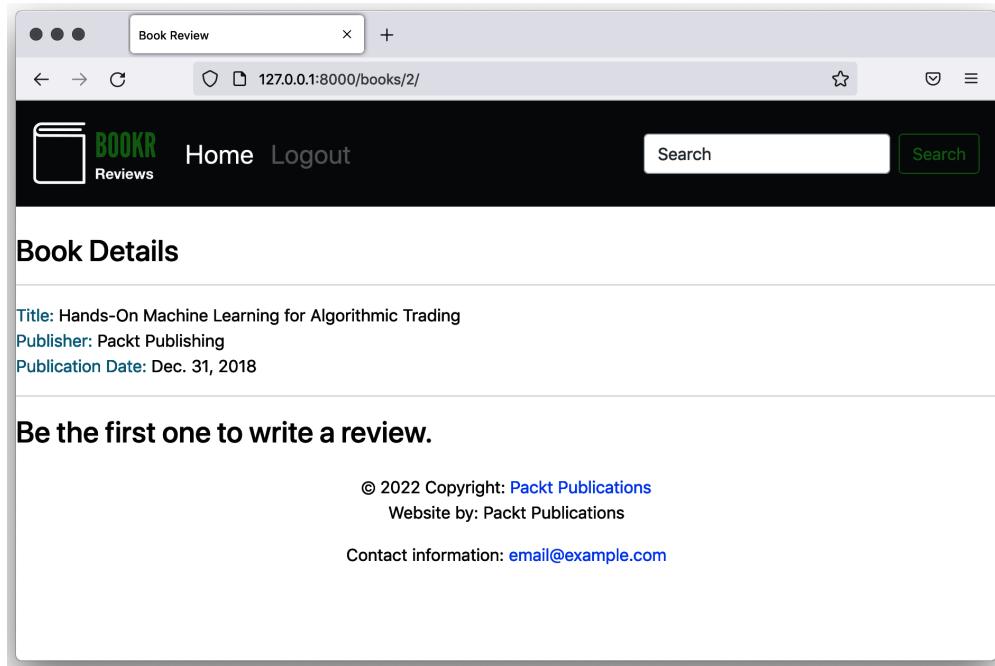


Figure 5.21 – The Book Details page after adding the Reviews logo

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Activity 5.02 – CSS enhancements

Currently, the CSS is kept inline in the `base.html` template. As a best practice, it should be moved into its own file so that it can be cached separately and decrease the size of the HTML downloads. As part of this, you'll also add some CSS enhancements such as fonts and colors, and link in Google Fonts CSS to support these changes.

These steps will help you complete this activity:

1. Create a directory named `static` in the Bookr project directory. Then, create a new file inside it named `main.css`.
2. Copy the contents of the `<style>` element from the main `base.html` template into the new `main.css` file, then remove the `<style>` element from the template. Add these extra rules to the end of the CSS file:

```
body {  
    font-family: 'Source Sans Pro', sans-serif;  
    background-color: #e6efe8  
    color: #393939;  
}  
  
h1, h2, h3, h4, h5, h6 {  
    font-family: 'Libre Baskerville', serif;  
}
```

3. Link to the new `main.css` file with a `<link rel="stylesheet" href="...">` tag. Use the `{% static %}` template tag to generate the URL for the `href` attribute, and don't forget to load the `static` library.
4. Link in the Google Fonts CSS by adding this code to the base template:

```
<link rel="stylesheet" href="https://fonts.googleapis.com/  
css?family=Libre+Baskerville|Source+Sans+Pro&display  
=swap">
```

(You will need to have an active internet connection so that your browser can include this remote CSS file.)

5. Update your Django settings to add `STATICFILES_DIRS`, which is set to the `static` directory you created in *Step 1*. When you're finished, your Bookr application should look like *Figure 5.22*:

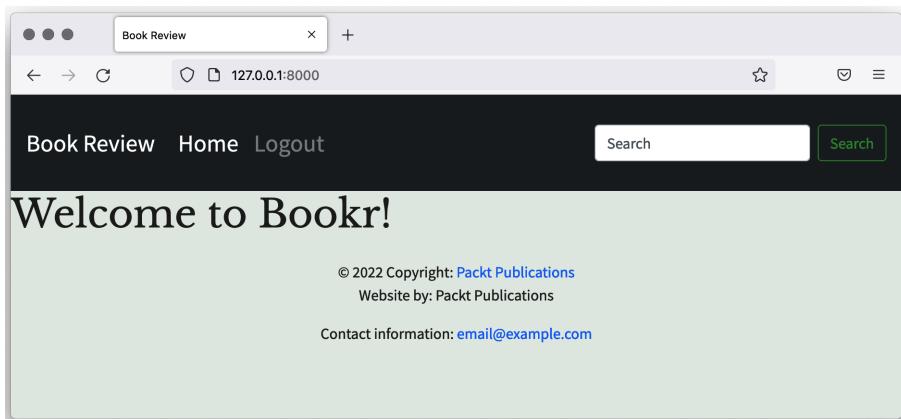


Figure 5.22 – Main page with a new font and background color

Notice the new font and background color. These should be displayed on all the Bookr pages.

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Activity 5.03 – adding a global logo

You have already added a logo that is served on pages for the Reviews app. We have another logo to be used globally by default, but other apps will be able to override it:

1. Download the Bookr logo (`logo.png`) from <https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter05/Activity5.03/bookr/static/logo.png>.
2. Save it in the main `static` directory for the project.
3. Edit the main `base.html` file. We already have a block for the logo (`brand`), so a `` can be placed inside here. Use the `static` template tag to refer to the logo you just downloaded.
4. Check that your pages work. On the main URL, you should see the Bookr logo, but on the **Book List** and **Book Details** pages, you should see the Bookr Reviews logo.

When you're finished, you should see the Bookr logo on the main page:

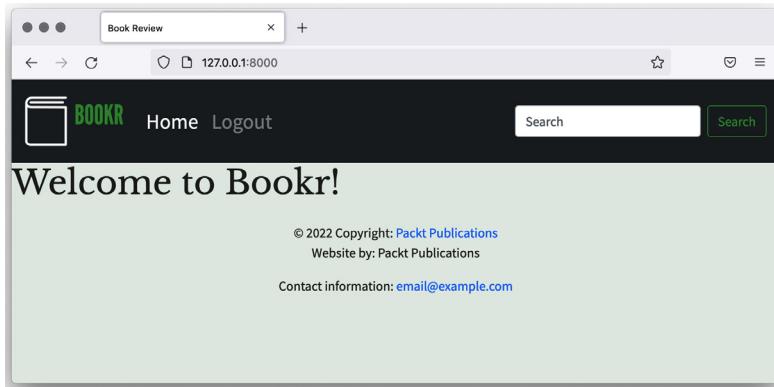


Figure 5.23 – Bookr logo on the main page

When you visit a page that had the Bookr Reviews logo before, such as the **Book List** page, it should still show the Bookr Reviews logo:

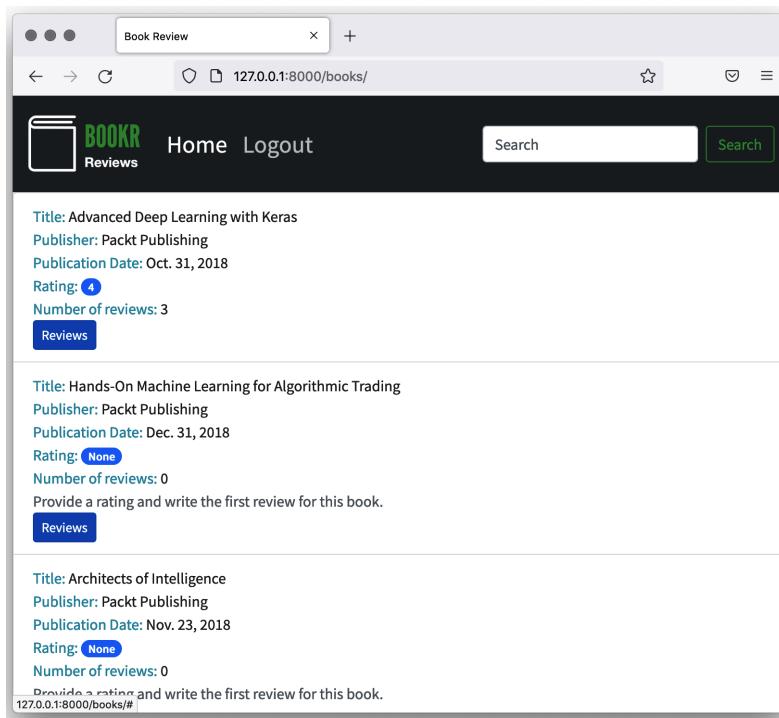


Figure 5.24 – Bookr Reviews logo still appears on the Reviews pages

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we showed you how to use Django's `staticfiles` app to find and serve static files. We used the built-in `static` view to serve these files with the Django dev server in `DEBUG` mode. We showed different places to store static files by using a directory that is global to the project or in a specific directory for the application; global resources should be stored in the former while application-specific resources should be stored in the latter. We showed the importance of namespacing static file directories to prevent conflicts. After serving the assets, we used the `static` tag to include them in our template. We then demoed how the `collectstatic` command copies all the assets into the `STATIC_ROOT` directory, for production deployment. We showed how to use the `findstatic` command to debug the loading of static files. To invalidate caches automatically, we looked at using `ManifestFilesStorage` to add a hash of the file's content to the static file URL. Finally, we briefly talked about using a custom file storage engine.

So far, we have only fetched web pages using content that already exists. In the next chapter, we will start adding forms so that we can interact with web pages by sending data to them over HTTP.

6

Forms

This chapter introduces web forms, a method of sending information from the browser to a web server. We will start by introducing forms in general and discussing how data is encoded to be sent to the server.

So far, the views we have been building for Django have been one-way only. Our browser is retrieving data from the views we have written, but it doesn't send any data back to them. In *Chapter 4, An Introduction to Django Admin*, we created model instances by using the Django admin and submitting forms, but those were using views built into Django, not created by us. In this chapter, we will use the Django Forms library to start accepting user-submitted data. The data will be provided through GET requests in the URL parameters and/or POST requests in the body of the request. But before we get into the details, first, let's understand what forms are in Django. You'll learn about the differences between sending form data in a GET HTTP request and sending it in a POST HTTP request and how to choose which one to use.

In this chapter, we will be covering the following topics:

- What is a form?
- The Django Forms library
- Validating forms and retrieving Python values
- Activity 1 – Book Search

By the end of this chapter, you will have learned how Django's Forms library is used to build and validate forms automatically and how it cuts down on the amount of manual HTML to write.

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter06>.

What is a form?

When working with an interactive web app, we not only want to provide data to users but also accept data from them to either customize the responses we're generating or to let them submit data to the site. When browsing the web, you most definitely will have used forms. Whether you're logged into your internet banking account, surfing the web with a browser, posting a message on social media, or writing an email to an online email client, in all these cases, you're entering data in a form. A form is made up of inputs that define key-value pairs of data to submit to the server. For example, when logging in to a website, the data being sent would have the keys *username* and *password*, with values of your username and your password, respectively. We will go into the different types of inputs in more detail in an upcoming section. Each input in the form has a *name*, and this is how its data is identified on the server side (in a Django view). There can be multiple inputs with the same *name*, whose data is available in a list containing all the posted values with this name – for example, a list of checkboxes with permissions to apply to users. Each checkbox would have the same name but a different value. The form has attributes that specify which URL the browser should submit the data to and what method it should use to submit the data (browsers only support *GET* or *POST*).

The GitHub Login page shown in the following figure is an example of a form:

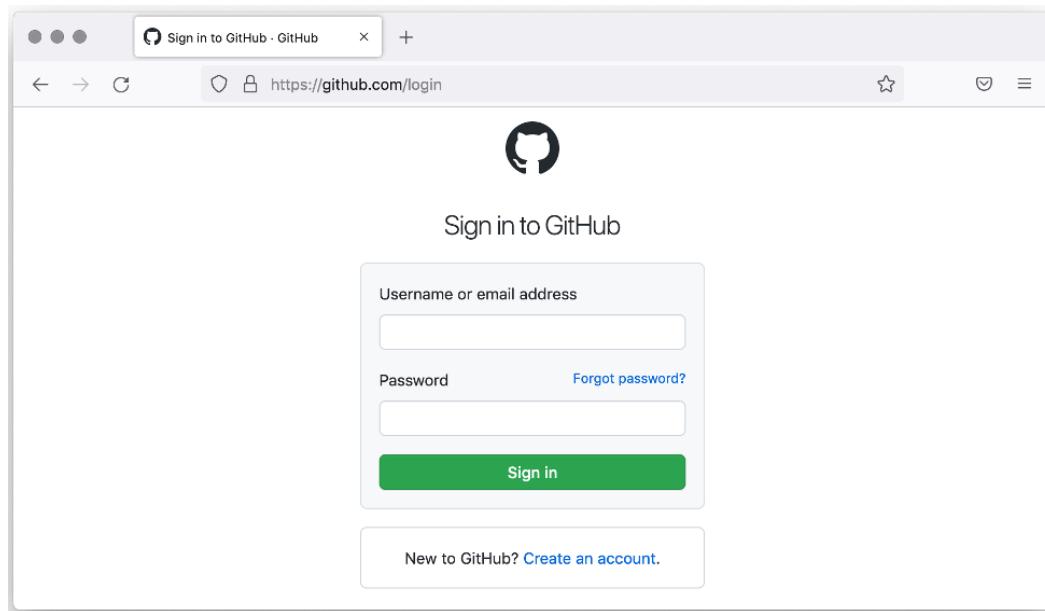
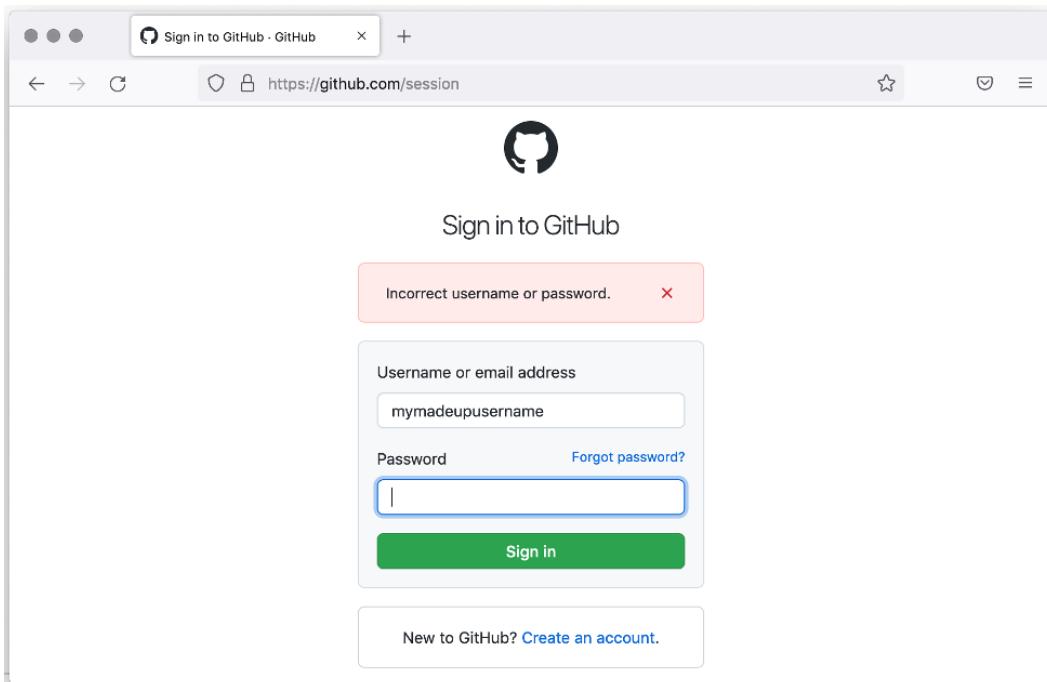


Figure 6.1: The GitHub Login page is an example of a form

It has three visible inputs: a text field (**Username or email address**), a password field (**Password**), and a submit button (**Sign in**). It also has a field that is not visible – its type is `hidden`, and it contains a special token for security called a CSRF token. We will discuss this later in this chapter. When you click the **Sign In** button, the form data is submitted with a *POST* request. If you entered a valid username and password, you will be logged in; otherwise, the form will display an error, as follows:



The screenshot shows a web browser window for GitHub. The address bar indicates the URL is <https://github.com/session>. The main content is a 'Sign in to GitHub' form. At the top, there is a red error message box containing the text 'Incorrect username or password.' with a close button. Below this, the form fields are visible: a text input for 'Username or email address' containing 'mymadeupusername', a password input field with a placeholder 'Password', and a 'Sign in' button. To the right of the password field is a 'Forgot password?' link. At the bottom of the form, there is a link 'New to GitHub? [Create an account](#)'.

Figure 6.2: Form submitted with incorrect username or password

There are two states a form can have: pre-submit and post-submit. The first is the initial state when the page is first loaded. All the fields will have a default value (usually empty), and no errors will be displayed. If all the information that has been entered into a form is valid, then usually, when it is submitted, you will be taken to a page showing the results of submitting the form. This might be a search results page or a page showing you the new object that you created. In this case, you will not see the form in its post-submit state.

If you did not enter valid information into the form, then it will be rendered again in its post-submit state. In this state, you will be shown the information that you entered, as well as any errors to help you resolve the problems with the form. These errors may be field errors or non-field errors. Field errors apply to a specific field – for example, leaving a required field blank or entering a value that is too large, too small, too long, or too short. If a form required you to enter your name and you left it blank, this would be displayed as a field error next to that field.

Non-field errors either do not apply to a field or apply to multiple fields and are displayed at the top of the form. In the preceding figure, we can see a message that either the username or password may be incorrect when logging in. For security, GitHub does not reveal if a username is valid, so this is displayed as a non-field error rather than a field error for a username or password (Django also follows this convention). Non-field errors also apply to fields that depend on each other. For example, on a credit card form, if the payment is rejected, we might not know if the credit card number or security code is incorrect; therefore, we can't show that error on a specific field. It applies to the form as a whole.

Let's look at the different parts of an HTML form, starting with the `form` element itself.

The `form` element

All inputs used during form submission must be contained inside an `<form>` element. There are three HTML attributes that you must use to modify the behavior of the form:

- `method`

This is the HTTP method for submitting the form; there's either `get` or `post`. If omitted, this defaults to `get` (because this is the default method when typing in a URL in the browser and hitting *Enter*).

- `action`

This refers to the URL (or path) to send the form data to. If omitted, the data gets sent back to the current page.

- `enctype`

This sets the encoding type of the form. You only need to change this if you are using the form to upload files. The most common values are `application/x-www-form-urlencoded` (the default if this value is omitted) or `multipart/form-data` (set this if you're uploading files). Note that you don't have to worry about the encoding type in your view – Django handles the different types automatically.

Here is an example of a form without any of its attributes set:

```
<form>
  <!-- Input elements go here -->
</form>
```

It will submit its data using a `GET` request to the current URL that the form is being displayed on, using the `application/x-www-form-urlencoded` encoding type.

In the following example, we'll set all three attributes on a form:

```
<form method="post" action="/form-submit"
enctype="multipart/form-data">
    <!-- Input elements go here -->
</form>
```

This form will submit its data with a POST request to the /form-submit path, encoding the data as multipart/form-data.

How do GET and POST requests differ in how the data is sent? Recall in *Chapter 1, An Introduction to Django*, that we discussed what the underlying HTTP request and response data that your browser sends look like. In the following two examples, we will submit the same form twice, the first time using GET and the second time using POST. The form will have two inputs: a first name and a last name.

A form submitted using GET sends its data in the URL like this:

```
GET /form-submit?first_name=Joe&last_name=Bloggs HTTP/1.1
Host: www.example.com
```

Whereas a form submitted using POST sends its data in the body of the request, like this:

```
POST /form-submit HTTP/1.1
Host: www.example.com
Content-Length: 31
Content-Type: application/x-www-form-urlencoded

first_name=Joe&last_name=Bloggs
```

You'll notice that the form data is encoded the same way in both cases; it is just placed differently for the GET and POST requests. In the *Choosing between GET or POST* section, we'll discuss how to choose between these two types of requests.

Types of inputs

We've seen four examples of inputs so far (*text*, *password*, *submit*, and *hidden*). Most inputs are created with a `<input>` tag, and their type is specified with its *type* attribute. Each input has a *name* attribute that defines the key for the key-value pairs that are sent to the server in the HTTP request.

In the next exercise, we'll look at how we can build a form in HTML. This will allow you to get up to speed on many different form fields.

Exercise 6.01 – building a form in HTML

For the first few exercises of this chapter, we will need an HTML form to test with. We will manually code one in this exercise. This will also allow you to experiment with how different fields are validated and submitted. This will be done in a new Django project so that we don't interfere with Bookr. You can refer to *Chapter 1* to refresh your memory on creating a Django project:

1. We'll start by creating the new Django project. You can reuse the Bookr virtual environment that already has Django installed. Open a new terminal and activate the virtual environment. Then, use `django-admin` to start a Django project named `form_project`. To do this, run the following command:

```
django-admin startproject form_project
```

This will scaffold the Django project in a directory named `form_example`.

2. Create a new Django app in this project by using the `startapp` management command. The app should be called `form_example`. To do this, `cd` into the `form_project` directory, then run the following:

```
python3 manage.py startapp form_example
```

This will create the `form_example` app directory inside the `form_project` directory.

3. Launch PyCharm, then open the `form_project` directory. If you already have a project open, you can do this by choosing **File** -> **Open**; otherwise, just click **Open** in the **Welcome to PyCharm** window. Navigate to the `form_project` directory, select it, then click **Open**.

If prompted, be sure to choose **Trust Project**.

The `form_project` project window should be shown like this:

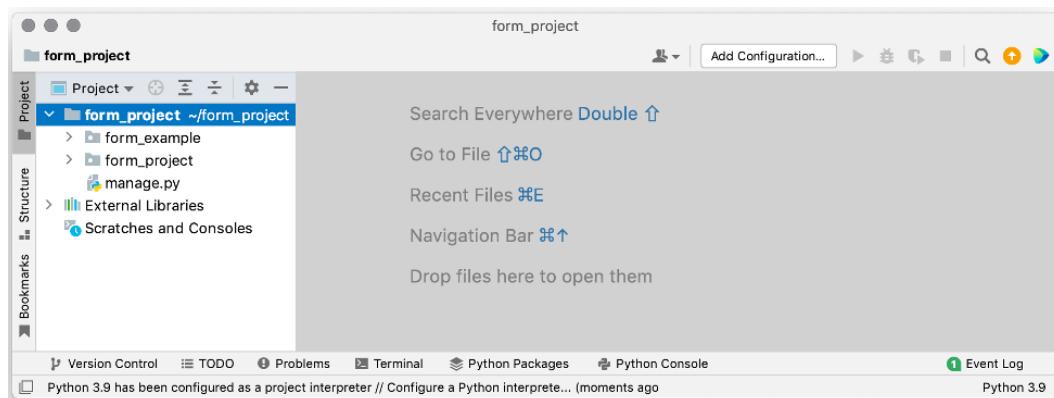


Figure 6.3: `form_project` project open

4. Create a new run configuration to execute `manage.py runserver` for the project. You can reuse the Bookr virtual environment again. The **Run/Debug Configurations** window should look similar to the following figure when you're done:

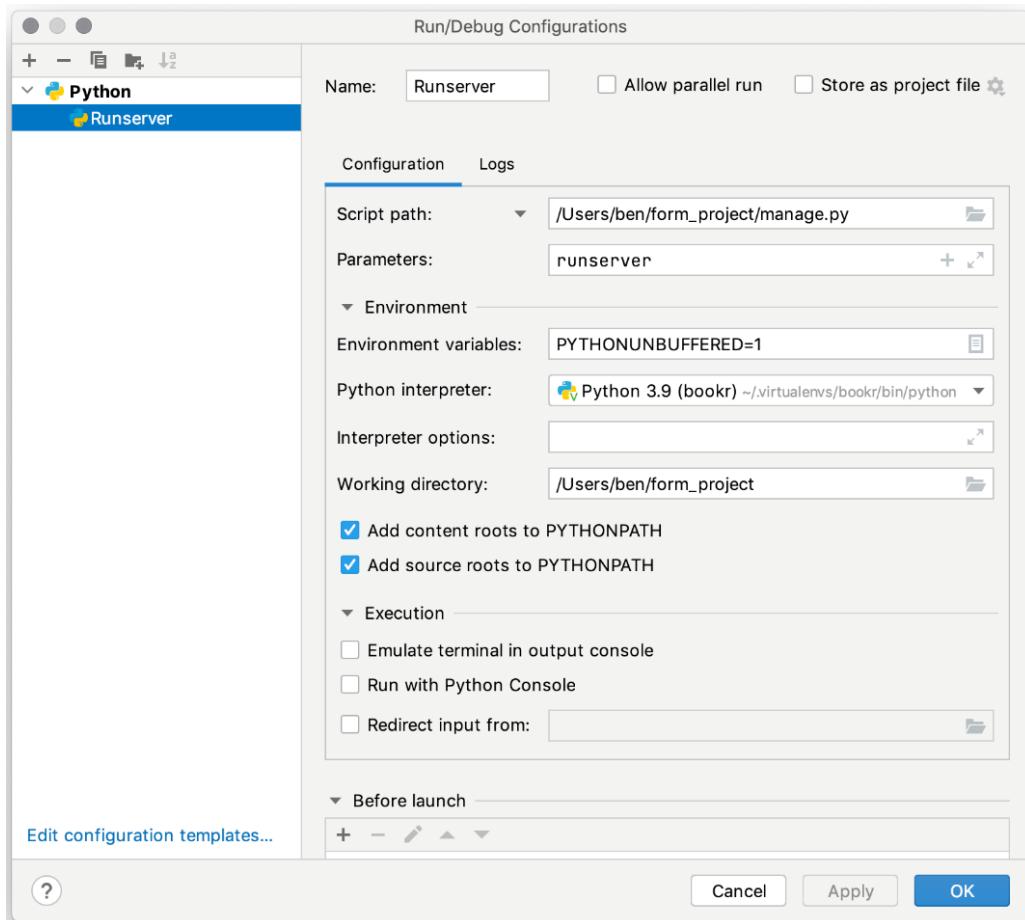


Figure 6.4: Run/Debug Configurations for Runserver

You can test that the configuration has been set up correctly by clicking the **Run** button, then visiting `http://127.0.0.1:8000/` in your browser. You should see the Django welcome screen. If the debug server fails to start or you see the Bookr main page, then you probably still have the Bookr project running. Try stopping the Bookr `runserver` process and then starting the new one you just set up.

5. Open `settings.py` in the `form_project` directory and add `"form_example"` to the `INSTALLED_APPS` setting.

6. The last step in setting up this new project is to create a templates directory for the `form_example` app. Right-click on the `form_example` directory then select **New | Directory**. Name it `templates`.
7. We need an HTML template to display our form. Create one by right-clicking on the `templates` directory you just created, then choosing **New | HTML File**.

In the dialog box that appears, enter the name `form-example.html` and press *Enter* to create it.

8. The `form-example.html` file should now be open in the editor pane of PyCharm. Start by creating the `form` element. We'll set its `method` attribute to `post`. The `action` attribute will be omitted, which means the form will submit back to the same URL on which it was loaded.

Insert this code between the `<body>` and `</body>` tags:

```
<form method="post">  
</form>
```

9. Now, let's add a few inputs. To add a little bit of spacing between each input, we'll wrap them inside `<p>` tags. We will start with a text field and a password field. This code should be inserted between the `<form>` tags you just created:

```
<p>  
    <label for="id_text_input">Text Input</label><br>  
    <input id="id_text_input" type="text"  
        name="text_input" value="" placeholder=  
        "Enter some text">  
</p>  
<p>  
    <label for="id_password_input">Password Input  
        </label><br>  
    <input id="id_password_input" type="password"  
        name="password_input" value="" placeholder="Your  
        password">  
</p>
```

10. Next, we will add two checkboxes and three radio buttons. Insert this code after the HTML you added in the previous step; it should come before the `</form>` tag:

```
<p>  
    <input id="id_checkbox_input" type="checkbox"  
        name="checkbox_on" value="Checkbox Checked"  
        checked>  
    <label for="id_checkbox_input">Checkbox</label>  
</p>  
<p>  
    <input id="id_radio_one_input" type="radio"
```

```
    name="radio_input" value="Value One">
<label for="id_radio_one_input">Value One</label>
<input id="id_radio_two_input" type="radio"
    name="radio_input" value="Value Two" checked>
<label for="id_radio_two_input">Value Two</label>
<input id="id_radio_three_input" type="radio"
    name="radio_input" value="Value Three">
<label for="id_radio_three_input">Value Three
</label>
</p>
```

11. Next is a drop-down select menu, for choosing a favorite book. Add this code after that of the previous step but before the `</form>` tag:

```
<p>
    <label for="id_favorite_book">Favorite Book
    </label><br>
    <select id="id_favorite_book" name=
        "favorite_book">
        <optgroup label="Non-Fiction">
            <option value="1">Deep Learning with
                Keras</option>
            <option value="2">Web Development with
                Django</option>
        </optgroup>
        <optgroup label="Fiction">
            <option value="3">Brave New World</option>
            <option value="4">The Great
                Gatsby</option>
        </optgroup>
    </select>
</p>
```

It will display four options that are split into two groups. The user will only be able to select one option.

12. The next is a multiple select (achieved by using the `multiple` attribute). Add this code after that of the previous step but before the `</form>` tag:

```
<p>
    <label for="id_books_you_own">Books You Own
    </label><br>
    <select id="id_books_you_own" name="books_you_own"
            multiple>
        <optgroup label="Non-Fiction">
            <option value="1">Deep Learning with
                Keras</option>
            <option value="2">Web Development with
                Django</option>
        </optgroup>
        <optgroup label="Fiction">
            <option value="3">Brave New World</option>
            <option value="4">The Great Gatsby
            </option>
        </optgroup>
    </select>
</p>
```

The user can select zero or more options from the four. They are displayed in two groups.

13. Next is a `textarea`. It is like a text field but has multiple lines. This code should be added, like in the previous steps, before the closing `</form>` tag:

```
<p>
    <label for="id_text_area">Text Area</label><br>
    <textarea name="text_area" id="id_text_area"
              placeholder="Enter multiple lines of text">
    </textarea>
</p>
```

-
14. Next, we'll add some fields for specific data types: number, email, and date inputs. Add this all before the `</form>` tag:

```
<p>
  <label for="id_number_input">Number Input
  </label><br>
  <input id="id_number_input" type="number"
    name="number_input" value="" step="any"
    placeholder="A number">
</p>
<p>
  <label for="id_email_input">Email Input
  </label><br>
  <input id="id_email_input" type="email"
    name="email_input" value="" placeholder="Your
    email address">
</p>
<p>
  <label for="id_date_input">Date Input</label><br>
  <input id="id_date_input" type="date"
    name="date_input" value="2019-11-23">
</p>
```

15. We'll now add some buttons to submit the form. Once again, insert this before the closing `</form>` tag:

```
<p>
  <input type="submit" name="submit_input"
    value="Submit Input">
</p>
<p>
  <button type="submit" name="button_element"
    value="Button Element">Button With<strong>
    Styled</strong> Text
  </button>
</p>
```

This will demonstrate two ways of creating submit buttons, either as an `<input>` or `<button>`.

16. Finally, we'll add a hidden field. Insert this before the closing `</form>` tag:

```
<input type="hidden" name="hidden_input" value="Hidden  
Value">
```

This field can't be seen or edited, so it has a fixed value.

You can now save and close `form-example.html`.

17. As with any template, we can't see it unless we have a view to render it. Open the `form_example` app's `views.py` file and add a new view called `form_example`. It should render and return the template you just created, like so:

```
def form_example(request):  
    return render(request, "form-example.html")
```

You can now save and close `views.py`.

18. You should be familiar with the next step now, which is to add a URL mapping to the view. Open the `urls.py` file in the `form_project` package directory. Add a mapping for the `form-example` path to your `form_example` view, to the `urlpatterns` variable. It should look like this:

```
path('form-example/', form_example.views.form_example)
```

Make sure you also add an import of `form_example.views`. Save and close `urls.py`.

19. Start the Django Dev Server (if it is not already running), then load your new view in your web browser; the address is `http://127.0.0.1:8000/form-example/`. Your page should look like this:

The screenshot shows a web browser window with the title 'Title' and the URL '127.0.0.1:8000/form-example/'. The page displays a form with the following fields:

- Text Input:** A text input field with placeholder text 'Enter some text'.
- Password Input:** A password input field with placeholder text 'Your password'.
- Checkbox:** A checkbox labeled 'Checkbox' with the checked attribute.
- Radio buttons:** Three radio buttons labeled 'Value One', 'Value Two', and 'Value Three', with 'Value Two' being the selected option.
- Favorite Book:** A dropdown menu showing 'Deep Learning with Keras'.
- Books You Own:** A list of books categorized into 'Non-Fiction' and 'Fiction'.
 - Non-Fiction:** Deep Learning with Keras, Web Development with Django.
 - Fiction:** Dune, Now World.
- Text Area:** A text area with placeholder text 'Enter multiple lines of text'.
- Number Input:** A number input field with placeholder text 'A number'.
- Email Input:** An email input field with placeholder text 'Your email address'.
- Date Input:** A date input field showing '23 / 11 / 2019' with a clear button.
- Submit Input:** A 'Submit Input' button at the bottom of the form.

Figure 6.5: Example inputs page

You can now familiarize yourself with the behavior of the web forms and see how they are generated from the HTML you specified. One activity to try is to enter invalid data into the number, date, or email inputs and click the submit button – the built-in HTML validation should prevent the form from being submitted:

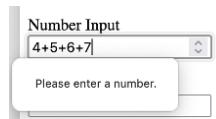


Figure 6.6: Browser error due to an invalid number

We have not yet set up everything for form submission, so if you correct all the errors in the form and try to submit it (by clicking either of the submit buttons), you will receive an error stating **CSRF verification failed**, as we can see in the following figure. We will talk about what this means and how to fix it later in this chapter:

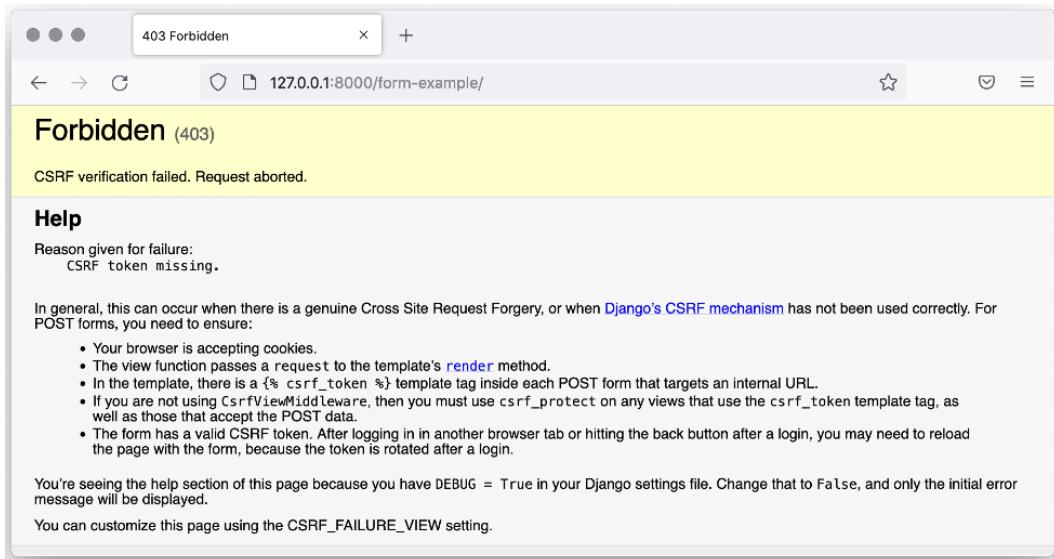


Figure 6.7: CSRF verification error

20. If you do receive an error, just go **Back** in your browser to return to the input example page.

In this exercise, you created an example page showcasing many HTML inputs, then created a view to render it and a URL to map to it. You loaded the page in your browser and experimented with changing data and trying to submit the form when it contained errors.

We just saw an example of Django's CSRF protection. In the next section, we'll learn what CSRF is and how Django's protection works.

Form security with Cross-Site Request Forgery Protection

Throughout this book, we have mentioned features that Django includes to prevent certain types of security exploits. One of these features is protection against **Cross-Site Request Forgery (CSRF)**.

The CSRF attack exploits the fact that a form on a website can be submitted to any other website. The `action` attribute of this `form` just needs to be set appropriately. Let's look at an example for Bookr. We don't have this set up yet, but we will be adding a view and a URL that allows us to post a review for a book. To do this, we'll have a form for posting the review content and selecting the rating. Its HTML is like this:

```
<form method="post" action="http://127.0.0.1:8000/books/4/reviews/">
  <p>
    <label for="id_review_text">Your Review
```

```
</label><br/>
<textarea id="id_review_text" name="review_text"
placeholder="Enter your review"></textarea>
</p>
<p>
    <label for="id_rating">Rating</label><br/>
    <input id="id_rating" type="number" name="rating"
placeholder="Rating 1-5">
</p>
<p>
    <button type="submit">Create Review</button>
</p>
</form>
```

And on a web page, it would look like this:

Your Review

Rating

Create Review

Figure 6.8: Example review create form

Someone could take this form, make a few changes, and host it on their website. For example, they could make the inputs hidden and hardcode a good review and rating for a book, and then make it look like some other kind of form, like this:

```
<form method="post"
action="http://127.0.0.1:8000/books/4/reviews/">
    <input type="hidden" name="review_text" value="This
book is great!">
    <input type="hidden" name="rating" value="5">
    <p>
        <button type="submit">Enter My Website</button>
    </p>
</form>
```

Of course, the hidden fields don't display, so the form looks like this on the malicious website:

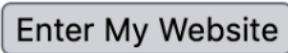


Figure 6.9: Hidden inputs are not visible

The user would think they were clicking a button to enter a website, but while clicking it, they would submit the hidden values to the original view on Bookr. Of course, a user could check the source code of the page they were on to check what data is being sent and where, but most users are unlikely to inspect every form they come across. The attacker could even have the form with no submit button and just use JavaScript to submit it, which means the user would be submitting the form without even realizing it.

You may think that requiring the user to log in to Bookr will prevent this type of attack, and it does limit its effectiveness somewhat, as the attack would then only work for logged-in users. But because of the way authentication works, once a user is logged in, they have a cookie set in their browser that identifies them to the Django application. This cookie is sent on every request so that the user does not have to provide their login credentials on every page. Because of the way web browsers work, they will include the server's authentication cookie in *all* requests they send to that particular server. Even though our form is hosted on a malicious site, ultimately, it is sending a request to our application, so it will send through our server's cookies.

How can we prevent CSRF attacks? Django uses something called a CSRF token, which is a small random string that is unique to each site visitor – in general, you can consider a visitor to be one browser session. Different browsers on the same computer would be different visitors, and the same Django user logged in at two different browsers would also be different visitors. When the form is read, Django puts the token into the form as a hidden input. The CSRF token must be included in all POST requests being sent to Django, and it must match the token Django has stored on the server side for the visitor; otherwise, a 403 status HTTP response is returned. This protection can be disabled – either for the whole site or for an individual view – but it is not advisable to do so unless you need to. The CSRF token must be added into the HTML for every form being sent and is done with the `{% csrf_token %}` template tag. We'll add it to our example review form now; the code in the template will look like this:

```
<form method="post" action="http://127.0.0.1:8000/books/4/reviews/">
    {% csrf_token %}
    <p>
        <label for="id_review_text">Your Review
        </label><br/>
        <textarea id="id_review_text" name="review_text">
```

```
    placeholder="Enter your review"></textarea>
</p>
<p>
    <label for="id_rating">Rating</label><br/>
    <input id="id_rating" type="number" name="rating"
    placeholder="Rating 1-5">
</p>
<p>
    <button type="submit">Enter My Website</button>
</p>
</form>
```

When the template gets rendered, the template tag is interpolated, so the output HTML ends up like this (note that the inputs are still in the output – they have just been removed here for brevity):

```
<form method="post" action=
"http://127.0.0.1:8000/books/4/reviews/">
    <input type="hidden" name="csrfmiddlewaretoken"
    value="tETZjLDUXev1tiYqGCSbMQkhWiesHCnutxpt6mutHI6YH6
    4F0nin5k2JW3B68IeJ">
    ...
</form>
```

Since this is a hidden field, the form on the page does not look any different from how it did before.

The CSRF token is unique to every visitor on the site. If an attacker were to copy the HTML from our site, they would get their own CSRF token, which would not match that of any other user, so Django would reject the form when it was posted by someone else.

CSRF tokens also change periodically. This limits how long the attacker would have to take advantage of a particular user and token combination. Even if they were able to get the CSRF token of a user that they were trying to exploit, they would have a short window of time to use it.

Now that you know what a CSRF token is used for, we will be adding it to our example form in the next exercise. Before we get to that, let's have a quick refresh about accessing data in a request using `QueryDict` objects.

Accessing data in the View

As we discussed in *Chapter 1*, Django provides two `QueryDict` objects on the `HttpRequest` instances that are passed to the view function. These are `request.GET`, which contains parameters passed in the URL, and `request.POST`, which contains parameters in the HTTP request body. Even though `request.GET` has `GET` in its name, this variable is populated even for non-`GET` HTTP requests. This is because the data it contains is parsed from the URL. Since all HTTP requests have a URL, all HTTP requests may contain `GET` data, even if they are `POST` or `PUT`, and so on.

In the next exercise, we will add code to our view to read and display the `POST` data.

Exercise 6.02 – working with POST data in a view

We will now add some code to our example view to print out the received `POST` data to the console. We will also insert the HTTP method that was used to generate the page, into the HTML output. This will allow us to be sure of what method was used to generate the page (`GET` or `POST`) and see how the form differs for each type:

1. First, in PyCharm, open the `form_example` app's `views.py` file. Alter the `form_example` view so that it prints each value in the `POST` to the console by adding this code inside the function:

```
for name in request.POST:  
    print("{}: {}".format(name,  
                          request.POST.getlist(name)))
```

This code iterates over each key in the `request.POST` data's `QueryDict` and prints the key and list of values to the console. We already know that each `QueryDict` can have multiple values for a key, so we use the `getlist` function to get them all.

2. Pass the template's `request.method` into a context variable named `method`. Do this by updating the call to `render` in the view so that it's like this:

```
return render(request, "form-example.html", {"method":  
    request.method})
```

3. We will now display the `method` variable in the template. Open the `form-example.html` template and use a `<h4>` tag to show the `method` variable. Put this just after the opening `<body>` tag, like so:

```
<body>  
  <h4>Method: {{ method }}</h4>
```

Note that we could access the method directly inside the template without passing it in a context dictionary, by using the `request.method` variable. We know from *Chapter 3, URL Mapping, Views, and Templates*, that by using the `render` shortcut function, the request is always available in the template. We just demonstrated how to access the method in the view here because later on, we will change the behavior of the page based on the method.

4. We also need to add the CSRF token to the form HTML. We do this by putting the `{% csrf_token %}` template tag after the opening `<form>` tag. The start of the form should look like this:

```
<form method="post">
    {% csrf_token %}
```

Now, save the file.

5. Start the Django Dev Server if it's not already running. Load the example page (`http://127.0.0.1:8000/form-example/`) in your browser; you should see it now displays the method at the top of the page (**GET**):



Figure 6.10: Method at the top of the page

6. Enter some text or data in each of the inputs and submit the form by clicking the **Submit Input** button:



The screenshot shows a web browser window with the following details:

- Title Bar:** Title, x, +
- Address Bar:** 127.0.0.1:8000/form-example/
- Form Fields:**
 - Text Area:** Text Area
 - Number Input:** 145
 - Email Input:** user@example.com
 - Date Input:** 22 / 03 / 2022
 - Submit Input:** A button labeled "Submit Input".
 - Button With Styled Text:** A button labeled "Button With **Styled** Text".
- Buttons:** Back, Forward, Stop, Refresh, Home, and a search bar.

Figure 6.11: Clicking the Submit Input to submit the form

You should see the page reload and the method display change to **POST**:



The screenshot shows a web browser window with the following details:

- Title Bar:** Title, x, +
- Address Bar:** 127.0.0.1:8000/form-example/
- Form Fields:**
 - Method:** POST
 - Text Input:** Enter some text
 - Password Input:** Your password
 - Checkbox:** A checked checkbox.
- Buttons:** Back, Forward, Stop, Refresh, Home, and a search bar.

Figure 6.12: Method updated to POST after form submitted

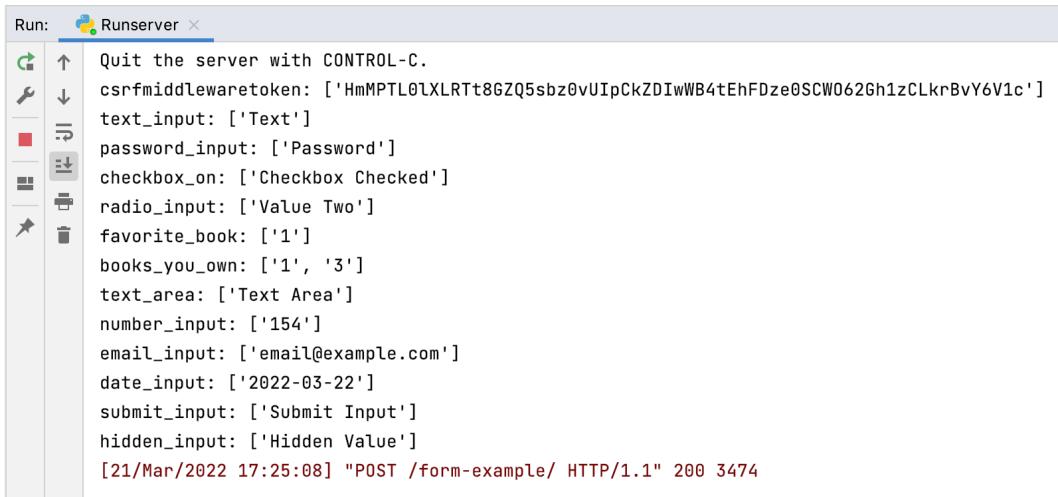
7. Switch back to PyCharm and look in the **Run** console at the bottom of the window. If the console is not visible, click the **Run** button at the bottom of the window to show it:



Figure 6.13: Clicking the Run button at the bottom of the window to display the console

If you're running the dev server in **Debug** mode, click **Debug** instead.

Inside the **Run** console, the list of the values that were posted to the server should be displayed:



The screenshot shows a terminal window titled "Run: Runserver". The content of the terminal is a list of posted form values. The values are listed in key-value pairs, where the key is the input type and the value is a list of values. The list includes:

- csrfmiddlewaretoken: ['HmMPTL0lXLRTt8GZQ5sbz0vUIpCkZDIwWB4tEhFDze0SCW062Gh1zCLkrBvY6V1c']
- text_input: ['Text']
- password_input: ['Password']
- checkbox_on: ['Checkbox Checked']
- radio_input: ['Value Two']
- favorite_book: ['1']
- books_you_own: ['1', '3']
- text_area: ['Text Area']
- number_input: ['154']
- email_input: ['email@example.com']
- date_input: ['2022-03-22']
- submit_input: ['Submit Input']
- hidden_input: ['Hidden Value']

At the bottom of the terminal, the timestamp [21/Mar/2022 17:25:08] and the response details [POST /form-example/ HTTP/1.1 200 3474] are displayed.

Figure 6.14: Input values shown in the Run console

Here are some things you should note:

- All values are sent as text, even number and date inputs
 - For the `select` inputs, the `selected value` attributes of the selected options are sent, not the text content of the `option` tag.
 - If you select multiple options for `books_you_own`, then you will see multiple values in the request. This is why we use the `getlist` method since multiple values are sent for the same input name.
 - If the checkbox was checked, you will have a `checkbox_on` input in the debug output. If it was not checked, then the key will not exist at all (that is, there is no key instead of having the key existing with an empty string or `None` value).
 - We have a value for the `submit_input` name, which is the text `Submit Input`. You submitted the form by clicking the **Submit Input** button, so we receive its value. Notice that no value is set for the `button_element` input since that button was not clicked.
8. We will experiment with two other ways of submitting the form, first by pressing *Enter* on the keyboard when your cursor is in a text-like input (*text*, *password*, *date*, *email*, but not *text area* as pressing *Enter* will add a new line).

If you submit a form in this way, the form will act as though you had clicked the first submit button on the form, so the `submit_input` input value will be included. The output you see should match that of the previous figure.

The other way to submit the form is by clicking the `button_element` submit input. We will try clicking this button to submit the form. You should see that `submit_button` is no longer in the list of posted values, while `button_element` is now present:

```
csrfmiddlewaretoken: ['HmMPTL01XLRT']
text_input: ['Text']
password_input: ['']
checkbox_on: ['Checkbox Checked']
radio_input: ['Value Two']
favorite_book: ['1']
books_you_own: ['1', '3']
text_area: ['Text Area']
number_input: ['154']
email_input: ['email@example.com']
date_input: ['2022-03-22']
button_element: ['Button Element']
hidden_input: ['Hidden Value']
```

Figure 6.15: `submit_button` is now gone from the inputs, and `button_element` has been added

You can use this multiple-submit technique to alter how your view behaves, depending on which button was clicked. You can even have multiple submit buttons with the same `name` attribute to make the logic easier to write.

In this exercise, you added a CSRF token to your form element by using the `{% csrf_token %}` template tag. This means that your form can be submitted to Django successfully without it generating an **HTTP Permission Denied** response. We then added some code to output the values that our form contained when it was submitted. We tried submitting the form with various values to see how they are parsed into Python variables on the `request.POST` part of `QueryDict`. We will now discuss some more of the theory around the difference between GET and POST requests, then move on to the Django Forms library, which makes designing and validating forms easier.

Choosing between GET or POST

Choosing when to use a GET or POST request requires you to consider several factors. The most important is deciding whether the request should be idempotent. A request can be said to be idempotent if it can be repeated and produce the same result each time. Let's look at some examples.

If you type in any web address into your browser (such as any of the Bookr pages we have built so far), it will do a GET request to fetch the information. You can refresh the page, and no matter how many times you click to refresh, you will get the same data back. The request you make does not affect the content on the server. We call these requests idempotent.

Now, remember when you added data through the Django admin interface (in *Chapter 4*). You typed the information for the new book into a form, then clicked **Save**. Your browser made a POST request to create a new book on the server. If you repeated that POST request, the server would create *another* Book and would do so each time you repeated the request. Since the request is updating information, it is not idempotent. Your browser will warn you about this. If you've ever tried to refresh a page that you were sent to after submitting a form, you may have received a message asking if you want to "*Repost form data?*" (or something more verbose, as shown in the following figure). This is a warning that you're sending the form data again, which might cause the action you just undertook to be repeated:

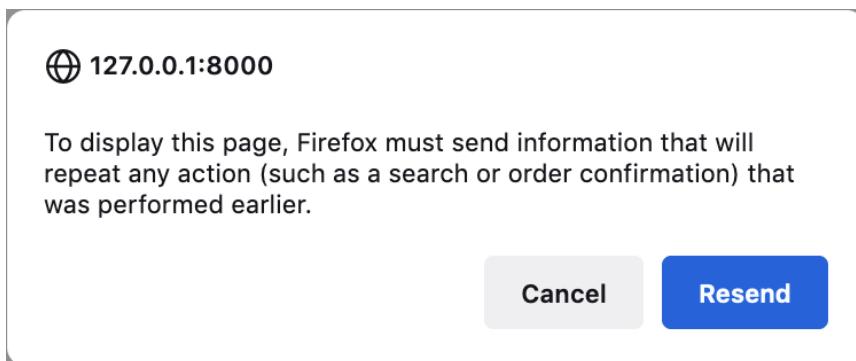


Figure 6.16: Firefox confirming if information should be resent

This is not to suggest that all GET requests are idempotent and all POST requests are not – your backend application can be designed in any way you want. Although it is not best practice, a developer might have decided to make data get updated during a GET request in their web application. When you are building your applications, you should try to make sure GET requests are idempotent and leave data-altering to POST requests only. Stick to these principles unless you have a good reason not to.

Another point to consider is that Django only applies CSRF protection to POST requests. Any GET request, including one that alters data, can be accessed without a CSRF token.

Sometimes, it can be hard to decide if a request is idempotent or not – for example, a login form. Before you submitted your username and password, you were not logged in, and afterward, the server considered you to be logged in, so could we consider that non-idempotent as it changed your authentication status with the server? On the other hand, once logged in, if you were able to send your credentials again, you would remain logged in. This implies that the request is idempotent and repeatable. So, should the request be a GET or a POST? This brings us to the second point to consider when choosing what method to use: form data visibility.

If sending form data with a GET request, the form parameters will be visible in the URL. For example, if we made a login form using a GET request, the login URL might be `https://www.example.com/login?username=user&password=password1`. The username, and worse, the password, are visible in the web browser's address bar. It would also be stored in the browser history, so anyone who used the browser after the real user could log in to the site. The URL is often stored in web server log files as well, meaning the credentials would be visible there too. In short, regardless of the idempotency of a request, don't pass sensitive data through URL parameters.

Sometimes, knowing that the parameter will be visible in the URL might be something you desire. For example, when searching with a search engine, usually, the search parameter will be visible in the URL. To see this in action, try visiting `https://www.google.com` and searching for something. You'll notice that the page with the results has your search term as the `q` parameter in the URL. For example, a search for "Django" will take you to the URL `https://www.google.com/search?q=Django`. This allows you to share search results with someone else by sending them this URL. In *Activity 1*, you will add a search form that passes a parameter similarly.

Another consideration is that the maximum length of a URL allowed by a browser can be short compared to the size of a POST body – sometimes, only around 2,000 characters (or about 2 KB) compared to many megabytes or gigabytes that a POST body can be (assuming your server is set up to allow these size requests).

As we mentioned earlier, URL parameters are available in `request.GET`, regardless of the type of request being made (GET, POST, PUT, and so on). You might find it useful to send some data in URL parameters and others in the request body (available in `request.POST`). For example, you could specify a `format` argument in the URL that sets what format some output data will be transformed to, but the input data is provided in the POST body.

Setting parameters in the GET query string might seem redundant when values can be passed as part of the URL path and then parsed by Django to route the URL. Next, we'll look at some reasons why you might prefer using the query string.

Why use GET when we can put parameters in the URL mappings?

Django allows us to easily define URL maps that contain variables. We could, for example, set up a URL mapping for a search view like this:

```
path("/search/<str:search>", reviews.views.search)
```

This probably looks like a good approach at first, but when we start wanting to customize the results view with arguments, it can get complicated quickly. For example, we might want to be able to move from one results page to the next, so we'll add a `page` argument:

```
path("/search/<str:search>/<int:page>", reviews.views.search)
```

And then, we might also want to order the search results by a specific category, such as author name or date of publishing, so we'll add another argument for that:

```
path("/search/<str:search>/<int:page>/<str:order >", reviews.views.  
    search)
```

You might be able to see the problem with this approach – we can't order the results without providing a page. If we wanted to add a `results_per_page` argument too, we wouldn't be able to use that without setting a `page` and `order` key.

Contrast this to using query parameters – all of them are optional, so you could search like this:

```
?search=search+term:
```

Or you could set a page like this:

```
?search=search+term&page=2
```

Or just set the results ordering like this:

```
?search=search+term&order=author
```

Or you could combine them all:

```
?search=search+term&page=2&order=author
```

Another reason to use URL query parameters is that when submitting a form, the browser always sends the input values in this manner – it cannot be changed so that parameters are submitted as path components in the URL. Therefore, when submitting a form using GET, the URL query parameters must be used as the input data.

Now that we've introduced the fundamentals of forms in HTTP requests and HTML, let's look at the Django Forms library, which makes dealing with forms and requests much easier.

The Django Forms library

We've looked at how to manually write forms in HTML and how to access the data on the request object using `QueryDict`. We saw that the browser provides some validation for us for certain field types (such as email or numbers), but we have not tried validating the data in the Python view. We should validate the form in the Python view for two reasons:

- It is not safe to rely solely on browser-based validation of input data. A browser may not implement certain validation features, meaning the user could post any type of data. For example, older browsers don't validate number fields, so a user can type in a number outside the range we are expecting. Further, a malicious user could try to send harmful data without using a browser at all. The browser validation should be considered a nicety for the user and that's all.

- The browser does not allow us to do cross-field validation. For example, we can use the `required` attribute for inputs that are mandatory to be filled in. Often, though, we want to set the `required` attribute based on the value of another input. For example, the email address input should only be set as `required` if the user has checked the `Register My Email` checkbox.

The Django Forms library allows you to quickly define a form using a Python class. This is done by creating a subclass of the base Django Form class. You can then use an instance of this class to render the form in your template and validate the input data. We refer to our classes as forms, similar to how we subclass Django Models to create Model classes. Forms contain one or more fields of a certain type (such as text fields, number fields, or email fields). You'll notice this sounds like Django Models, and forms *are* similar to Models but use different field classes. You can even automatically create a form from a Model – we will cover this in *Chapter 7*.

Defining a form

Creating a Django form is similar to creating a Django model: you define a class that inherits from the `django.forms.Form` class. The class has attributes that are instances of different `django.forms.Field` subclasses. When rendered, the attribute name in the class corresponds to its input name in HTML. To give you a quick idea of what fields there are, some examples are `CharField`, `IntegerField`, `BooleanField`, `ChoiceField`, and `DateField`. Each field generally corresponds to one input when rendered in HTML, but there's not always a one-to-one mapping between a form field class and an input type. Form fields are more coupled to the type of data they collect rather than how they are displayed.

To illustrate this, consider a `text` input and a `password` input. They both accept some typed-in text data, but the main difference between them is that the text is displayed in a `text` input whereas, with a `password` input, the text is obscured. In a Django form, both of these fields are represented using a `CharField`. Their difference in how they are displayed is set by changing the `widget` the field is using.

Note

If you're not familiar with the word `widget`, it's a term to describe the actual input that is being interacted with and how it is displayed. Text inputs, password inputs, select menus, checkboxes, and buttons are all examples of different widgets. The inputs we have seen in HTML correspond one-to-one with widgets. In Django, this is not the case, and the same type of `Field` class can be rendered in multiple ways, depending on the `widget` that is specified.

Django defines several widget classes that define how a field should be rendered as HTML. They inherit from `django.forms.widgets.Widget`. A widget can be passed to the `Field` constructor to change how it is rendered. For example, a `CharField` renders as a `text` `<input>` by default. If we use the `PasswordInput` widget, it will instead render as a `password` `<input>`. The other widgets we will use are as follows:

- `RadioSelect`, which renders a `ChoiceField` as radio buttons instead of a `<select>` menu
- `Textarea`, which renders a `CharField` as a `<textarea>`
- `HiddenInput`, which renders a field as a hidden `<input>`

We'll look at an example form and add fields and features one by one. First, let's just create a form with a text input and a password input:

```
from django import forms

class ExampleForm(forms.Form):
    text_input = forms.CharField()
    password_input =
        forms.CharField(widget=forms.PasswordInput)
```

The `widget` argument can be just a widget subclass, which can be fine a lot of the time. If you want to further customize the display of the input and its attributes, you can set the `widget` argument to an instance of the widget class instead. We will look at customizing the widget displays further soon. In this case, we're using just the `PasswordInput` class, since we are not customizing beyond changing the type of input being displayed.

When the form is rendered in a template, it looks like this:

Text input:

Password input:

Figure 6.17: Django form rendered in the browser

Note that the inputs do not contain any content when the page loads; the text has been entered to illustrate the different input types.

If we examine the page source, we will see the HTML that Django generates. For the first two fields, it looks like this (some spacing has been added for readability):

```
<p>
    <label for="id_text_input">Text input:</label>
    <input type="text" name="text_input" required
           id="id_text_input">
</p>
<p>
    <label for="id_password_input">Password input:</label>
    <input type="password" name="password_input" required
           id="id_password_input">
</p>
```

Notice that Django has automatically generated a `label` with text derived from the field name. The `name` and `id` attributes have been set automatically. Django also automatically adds the `required` attribute to the input. Similar to model fields, form field constructors also accept a `required` argument – this defaults to `True`. Setting this to `False` removes the `required` attribute from the generated HTML.

Next, we'll look at how a checkbox is added to the form.

A checkbox is represented with a `BooleanField` as it can have only two values: checked or unchecked. It's added to the form in the same way as the other field:

```
class ExampleForm(forms.Form):
    ...
    checkbox_on = forms.BooleanField()
```

The HTML Django generates for this new field is similar to the previous two fields:

```
<label for="id_checkbox_on">Checkbox on:</label>
<input type="checkbox" name="checkbox_on" required
       id="id_checkbox_on">
```

Next are select inputs:

- We need to provide a list of choices to display in the `<select>` dropdown.

- The field class constructor takes a `choices` argument. The choices are provided as a tuple of two-element tuples. The first element in each sub-tuple is the value of the choice, while the second element is the text or description of the choice. For example, choices could be defined like this:

```
BOOK_CHOICES = (
    ("1", "Deep Learning with Keras"),
    ("2", "Web Development with Django"),
    ("3", "Brave New World"),
    ("4", "The Great Gatsby")
)
```

Note that you can use lists instead of tuples if you want (or a combination of the two). This can be useful if you want your choices to be mutable:

```
BOOK_CHOICES = [
    ["1", "Deep Learning with Keras"],
    ["2", "Web Development with Django"],
    ["3", "Brave New World"],
    ["4", "The Great Gatsby"]
]
```

- To implement an `optgroup`, we can nest the choices. To implement the choices as we did in our previous examples, we can use a structure like this:

```
BOOK_CHOICES = (
(
    "Non-Fiction", (
        ("1", "Deep Learning with Keras"),
        ("2", "Web Development with Django")
    )
),
(
    "Fiction", (
        ("3", "Brave New World"),
        ("4", "The Great Gatsby")
    )
)
)
```

The select is added to the form by using a `ChoiceField`. The widget defaults to a `select` input so that no configuration is necessary apart from setting `choices`:

```
class ExampleForm(forms.Form):
    ...
    favorite_book =
        forms.ChoiceField(choices=BOOK_CHOICES)
```

This is the HTML that is generated:

```
<label for="id_favorite_book">Favorite book:</label>
<select name="favorite_book" id="id_favorite_book">
    <optgroup label="Non-Fiction">
        <option value="1">Deep Learning with Keras
        </option>
        <option value="2">Web Development with Django
        </option>
    </optgroup>
    <optgroup label="Fiction">
        <option value="3">Brave New World</option>
        <option value="4">The Great Gatsby</option>
    </optgroup>
</select>
```

- Making a multiple select requires the use of `MultipleChoiceField`. It takes a `choices` argument in the same format as the regular `ChoiceField` for single selects:

```
class ExampleForm(forms.Form):
    ...
    books_you_own =
        forms.MultipleChoiceField(choices=BOOK_CHOICES)
```

And its HTML is similar to that of the single select, except it has the `multiple` attribute added:

```
<label for="id_books_you_own">Books you own:</label>
<select name="books_you_own" required
        id="id_books_you_own" multiple>
    <optgroup label="Non-Fiction">
        <option value="1">Deep Learning with Keras
        </option>
        <option value="2">Web Development with Django
        </option>
    </optgroup>
    <optgroup label="Fiction">
        <option value="3">Brave New World</option>
        <option value="4">The Great Gatsby</option>
    </optgroup>
</select>
```

- Choices can also be set after the form has been instantiated. You may want to generate the list/tuple choice inside your view dynamically and then assign it to the field's `choices` attribute, like so:

```
form = ExampleForm()
form.fields["books_you_own"].choices = [ ("1", "Deep
Learning with Keras"), ...]
```

Next are radio inputs, which are similar to selects:

- Like selects, radio inputs use `ChoiceField`, as they provide a single choice between multiple options.
- The options to choose between are passed into the field constructor with the `choices` argument.
- The choices are provided as a tuple of two-element tuples, also like selects:

```
choices = (
    ("1", "Option One"),
    ("2", "Option Two"),
    ("3", "Option Three")
)
```

- `ChoiceField` defaults to displaying as a select input, so the widget must be set to `RadioSelect` to have it render as radio buttons. Putting the choice setting together with this, we can add radio buttons to the form like this:

```
RADIO_CHOICES = (
    ("Value One", "Value One"),
    ("Value Two", "Value Two"),
    ("Value Three", "Value Three")
)

class ExampleForm(forms.Form):
    ...
    radio_input =
        forms.ChoiceField(choices=RADIO_CHOICES,
                           widget=forms.RadioSelect)
```

Here is the HTML that is generated:

```
<label for="id_radio_input_0">Radio input:</label>
<ul id="id_radio_input">
<li>
    <label for="id_radio_input_0">
        <input type="radio" name="radio_input"
               value="Value One" required
               id="id_radio_input_0">
        Value One
    </label>
</li>
<li>
    <label for="id_radio_input_1">
        <input type="radio" name="radio_input"
               value="Value Two" required
               id="id_radio_input_1">
        Value Two
    </label>
</li>
<li>
    <label for="id_radio_input_2">
        <input type="radio" name="radio_input"
               value="Value Three" required
               id="id_radio_input_2">
        Value Three
    </label>
</li>
</ul>
```

Django automatically generates a unique label and ID for each of the three radio buttons.

- To create a `textarea`, use a `CharField` with a `Textarea` widget:

```
class ExampleForm(forms.Form):
    ...
    text_area = forms.CharField(widget=forms.Textarea)
```

You might notice that this `textarea` is much larger than the previous ones we have seen (see the following figure):



Figure 6.18: Normal `textarea` (top) versus Django default `textarea` (bottom)

This is because Django automatically adds `cols` and `rows` attributes. These set the number of columns and rows that the text field displays, respectively:

```
<label for="id_text_area">Text area:</label>
<textarea name="text_area" cols="40" rows="10"
required id="id_text_area"></textarea>
```

- Note that the `cols` and `rows` settings do not affect the amount of text that can be entered into a field, only the amount that is displayed at a time. Also, note that the size of `textarea` can be set using CSS (for example, the `height` and `width` properties). This will override the `cols` and `rows` settings.

To create number inputs, you might expect Django to have a `NumberField` type, but it does not.

- Remember that the Django form fields are data-centric rather than display-centric, so instead, Django provides different `Field` classes, depending on what type of numeric data you want to store.
- For integers, use an `IntegerField`.
- For floating point numbers, use `FloatField` or `DecimalField`. The latter two differ in how they convert their data into a Python value.
- `FloatField` will convert into a float, while a `DecimalField` is a decimal.
- Decimal values offer better accuracy in representing numbers than float values but may not integrate well into your existing Python code.

We'll add all three fields to the form at once:

```
class ExampleForm(forms.Form):  
    ...  
    integer_input = forms.IntegerField()  
    float_input = forms.FloatField()  
    decimal_input = forms.DecimalField()
```

Here's the HTML for all three:

```
<p>  
    <label for="id_integer_input">Integer input:</label>  
    <input type="number" name="integer_input" required  
          id="id_integer_input">  
</p>  
<p>  
    <label for="id_float_input">Float input:</label>  
    <input type="number" name="float_input" step="any"  
          required id="id_float_input">  
</p>  
<p>  
    <label for="id_decimal_input">Decimal input:</label>  
    <input type="number" name="decimal_input" step="any"  
          required id="id_decimal_input">  
</p>
```

- The `IntegerField`'s generated HTML is missing the `step` attribute that the other two have, which means the widget will only accept integer values.
- The other two fields (`FloatField` and `DecimalField`) generate very similar HTML, their behavior is the same in the browser, and they differ only when their values are used in Django code.

As you might have guessed, an `email` input can be created with an `EmailField`:

```
class ExampleForm(forms.Form):  
    ...  
    email_input = forms.EmailField()
```

Its HTML is similar to the `email` input we created manually:

```
<label for="id_email_input">Email input:</label>  
<input type="email" name="email_input" required  
id="id_email_input">
```

Continuing with our manually created form, the next field we will look at is `DateField`:

- By default, Django will render a `DateField` as a `text` input, and the browser will not show a calendar popup when the field is clicked on

We can add `DateField` to the form with no arguments, like this:

```
class ExampleForm(forms.Form):  
    ...  
    date_input = forms.DateField()
```

When rendered, it just looks like a normal `text` input:

Date input:

Figure 6.19: Default `DateField` display in the form

Here is the HTML generated by default:

```
<label for="id_date_input">Date input:</label>
<input type="text" name="date_input" required
id="id_date_input">
```

The reason for using a `text` input is that it allows the user to enter the date in several different formats. For example, by default, the user can type in the date in *Year-Month-Day* (dash-separated) or *Month/Day/Year* (slash-separated) formats.

The accepted formats can be specified by passing a list of formats to the `DateField` constructor using the `input_formats` argument. For example, we could accept dates in the *Day/Month/Year* or *Day/Month/Year-with-century* format, like this:

```
DateField(input_formats = ["%d/%m/%y", "%d/%m/%Y"])
```

We can override any attributes on a field's widget by passing the `attrs` argument to the widget constructor. This accepts a dictionary of attribute keys/values that will be rendered into the input's HTML.

We have not used this yet, but we will see it again in the next chapter when we customize the field rendering further. For now, we'll just set one attribute, `type`, that will overwrite the default input type:

```
class ExampleForm(forms.Form):
    ...
    date_input =
        forms.DateField(widget=forms.DateInput(attrs={"type": "date"}))
```

When rendered, it now looks like the date field we had before, and clicking on it brings up the calendar date picker:



Figure 6.20: DateField with date input

Examining the generated HTML now, we can see it uses the `date` type:

```
<label for="id_date_input">Date input:</label>
<input type="date" name="date_input" required
id="id_date_input">
```

The final input that we are missing is the hidden input:

- Once again, due to the data-centric nature of Django forms, there is no `HiddenField`.
- Instead, we choose the type of field that needs to be hidden and set its `widget` to a `HiddenInput`. We can then set the value of the field using the field constructor's `initial` argument:

```
class ExampleForm(forms.Form):  
    ...  
    hidden_input =  
        forms.CharField(widget=forms.HiddenInput,  
        initial="Hidden Value")
```

Here is the generated HTML:

```
<input type="hidden" name="hidden_input" value="Hidden  
Value" id="id_hidden_input">
```

- Note that as this is a hidden input, Django does not generate a `label` or surrounding `p` elements.

There are other form fields that Django provides that work in similar ways. These range from `DateTimeField` (for capturing a date and a time), to `GenericIPAddressField` (for either IPv4 or IPv6 addresses) and `URLField` (for URLs). A full list of fields is available at <https://docs.djangoproject.com/en/4.0/ref/forms/fields/>.

Rendering a form in a template

We've now seen how to create a form and add fields, and we've seen what the form looks like and what HTML is generated. But how is the form rendered in the template? We simply instantiate the form class and pass it to the `render` function in a view, using the context, just like any other variable.

For example, here's how we can pass our `ExampleForm` to a template:

```
def view_function(request):  
    form = ExampleForm()  
    return render(request, "template.html", {"form": form})
```

Django does not add the `<form>` element or submit button(s) for you when rendering the template; you should add these around where your form is placed in the template. The form can be rendered like any other variable.

We mentioned briefly earlier that the form was being rendered in the template using the `as_p` method. This layout method was chosen as it most closely matches the example form we built manually. Django offers three layout methods that can be used:

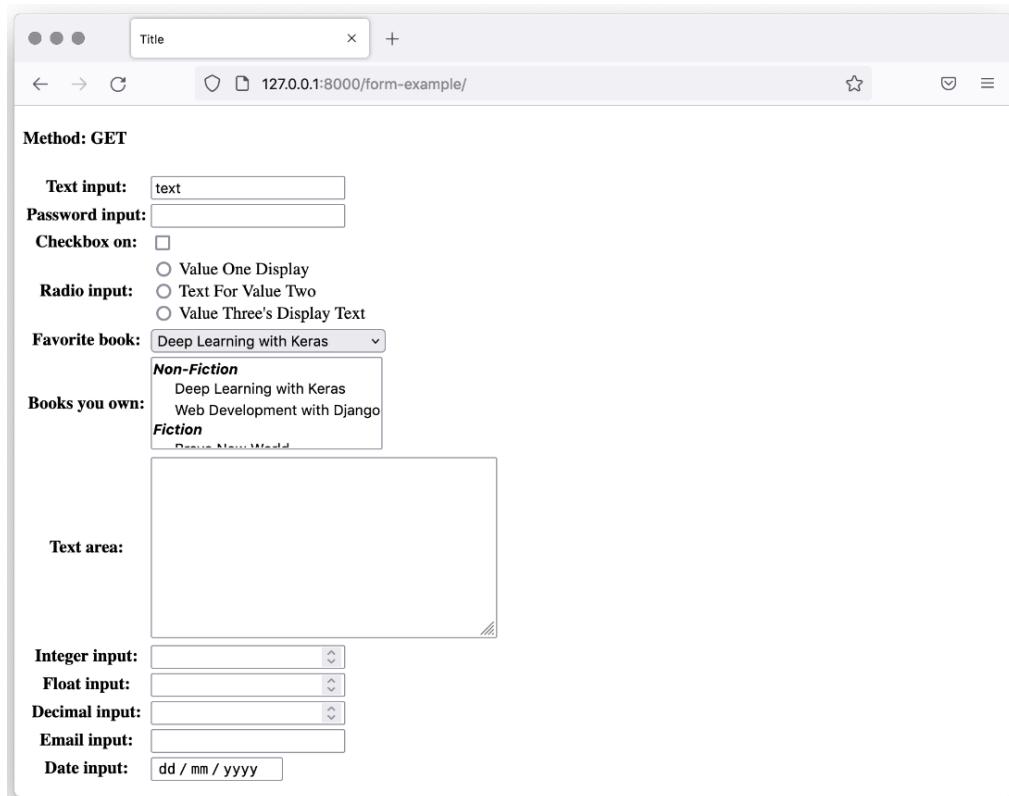
- `as_table`

The form is rendered as table rows, with each input on its own row. Django does not generate the surrounding `table` element, so you should wrap the form yourself; for example:

```
<form method="post">
    <table>
        {{ form.as_table }}
    </table>
</form>
```

`as_table` is the default rendering method, so `{{ form.as_table }}` and `{{ form }}` are equivalent.

When rendered, the form looks like this:



The screenshot shows a web browser window with a title bar 'Title' and a URL bar '127.0.0.1:8000/form-example/'. The content area displays a form rendered as a table. The form fields include:

- Text input:** A text input field with the value 'text'.
- Password input:** A password input field.
- Checkbox on:** A checkbox input field.
- Radio input:** Three radio buttons labeled 'Value One Display', 'Text For Value Two', and 'Value Three's Display Text'. The first one is selected.
- Favorite book:** A dropdown menu currently showing 'Deep Learning with Keras'.
- Books you own:** A dropdown menu with two sections:
 - Non-Fiction:** 'Deep Learning with Keras', 'Web Development with Django'.
 - Fiction:** 'Dune'.
- Text area:** A text area input field.
- Integer input:** An integer input field with a dropdown arrow.
- Float input:** A float input field with a dropdown arrow.
- Decimal input:** A decimal input field with a dropdown arrow.
- Email input:** An email input field.
- Date input:** A date input field with the placeholder 'dd / mm / yyyy'.

Figure 6.21: Form rendered as a table

Here is a small sample of HTML that is generated:

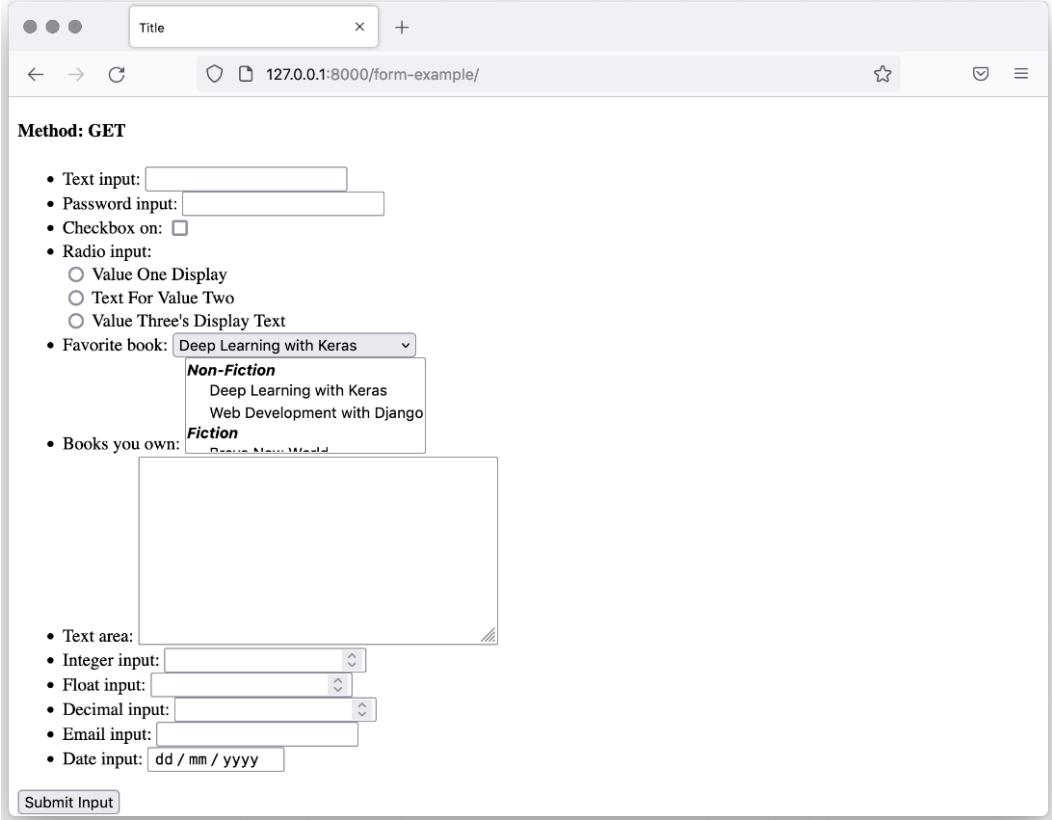
```
<tr>
    <th>
        <label for="id_text_input">Text input:</label>
    </th>
    <td>
        <input type="text" name="text_input" required
               id="id_text_input">
    </td>
</tr>
<tr>
    <th>
        <label for="id_password_input">Password input:
        </label>
    </th>
    <td>
        <input type="password" name="password_input"
               required id="id_password_input">
    </td>
</tr>
```

- **as_ul**

This renders the form fields as list items (`li`) inside either a `ul` or `ol` element. Like with `as_table`, the containing element (`` or ``) is not created by Django and must be added by you:

```
<form method="post">
    <ul>
        {{ form.as_ul }}
    </ul>
</form>
```

Here's how the form renders using `as_ul`:



Method: GET

- Text input:
- Password input:
- Checkbox on:
- Radio input:
 - Value One Display
 - Text For Value Two
 - Value Three's Display Text
- Favorite book:
 - Non-Fiction**
 - Deep Learning with Keras
 - Web Development with Django
 - Fiction**
 - Game of Thrones
- Books you own:
- Text area:
- Integer input:
- Float input:
- Decimal input:
- Email input:
- Date input:

Figure 6.22: Form rendered using `as_ul`

And here's a sample of the generated HTML:

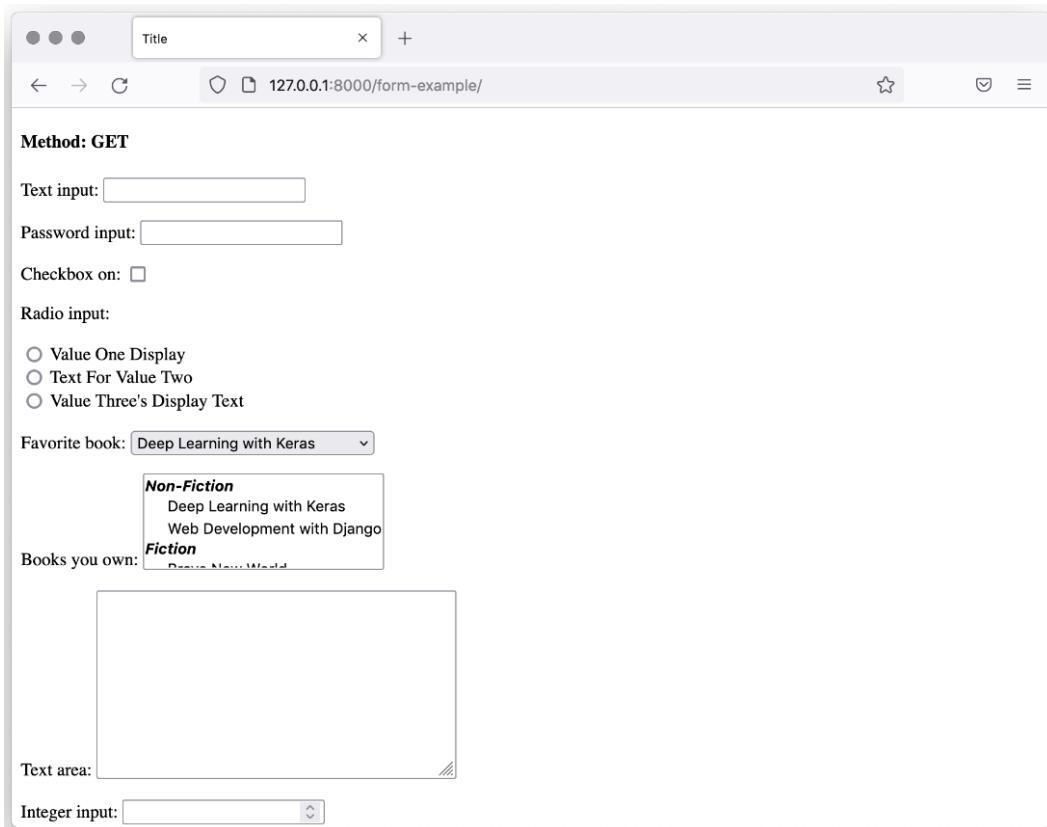
```
<li>
    <label for="id_text_input">Text input:</label>
    <input type="text" name="text_input" required
           id="id_text_input">
</li>
<li>
    <label for="id_password_input">Password input:
    </label>
    <input type="password" name="password_input"
           required id="id_password_input">
</li>
```

- `as_p`

Finally, there's the `as_p` method, which we used in our previous examples. Each input is wrapped within `p` tags, which means that you don't have to wrap the form manually (in a `<table>` or ``) as you did with the previous methods:

```
<form method="post">
    {{ form.as_p }}
</form>
```

Here's what the rendered form looks like:



The screenshot shows a web browser window with a Django form rendered using the `as_p` method. The browser title is "Title". The address bar shows "127.0.0.1:8000/form-example/". The form content is as follows:

Method: GET

Text input:

Password input:

Checkbox on:

Radio input:

Value One Display
 Text For Value Two
 Value Three's Display Text

Favorite book:

Books you own:

Non-Fiction

Deep Learning with Keras
Web Development with Django

Fiction

... (truncated)

Text area:

Integer input:

Figure 6.23: Form rendered using `as_p`

You've seen this before, but once again, here's a sample of the HTML generated:

```
<p>
    <label for="id_text_input">Text input:</label>
    <input type="text" name="text_input" required
           id="id_text_input">
</p>
<p>
    <label for="id_password_input">Password input:
        </label>
    <input type="password" name="password_input"
           required id="id_password_input">
</p>
```

It is up to you to decide which method you want to use to render your form, depending on which suits your application best. In terms of their behavior and use in your view, they are all identical. In *Chapter 15*, we will also introduce a method of rendering forms that will make use of the Bootstrap CSS classes.

Now that you have been introduced to Django forms, we can update our example form page to use a Django form instead of manually writing all HTML. We'll do this in the next exercise by replacing hand-written HTML with a Django form.

Exercise 6.03 – building and rendering a Django form

In this exercise, you will build a Django form using all the fields we have seen. The form and view will behave similarly to the form built manually; however, you will be able to see how much less code is required when writing forms using Django. Your form will also automatically get field validation, and if we make changes to the form, we don't have to then make changes to the HTML as it will update dynamically based on the form definition:

1. In PyCharm, create a new file called `forms.py` inside the `form_example` app directory.
2. Import the Django forms library at the top of your `forms.py` file:

```
from django import forms
```

3. Define the choices for the radio buttons by creating a `RADIO_CHOICES` variable. Populate it as follows:

```
RADIO_CHOICES = (
    ("Value One", "Value One Display"),
    ("Value Two", "Text For Value Two"),
    ("Value Three", "Value Three's Display Text")
)
```

You will use this soon when you create a `ChoiceField` called `radio_input`.

4. Define the nested choices for the book and select inputs by creating a BOOK_CHOICES variable. Populate it as follows:

```
BOOK_CHOICES = (
    (
        "Non-Fiction", (
            ("1", "Deep Learning with Keras"),
            ("2", "Web Development with Django")
        )
    ),
    (
        "Fiction", (
            ("3", "Brave New World"),
            ("4", "The Great Gatsby")
        )
    )
)
```

5. Create a class called ExampleForm, which inherits from the forms.Form class:

```
class ExampleForm(forms.Form):
```

Add the following fields as attributes to the class:

```
text_input = forms.CharField()
password_input =
    forms.CharField(widget=forms.PasswordInput)
checkbox_on = forms.BooleanField()
radio_input =
    forms.ChoiceField(choices=RADIO_CHOICES,
                      widget=forms.RadioSelect)
favorite_book =
    forms.ChoiceField(choices=BOOK_CHOICES)
books_you_own =
    forms.MultipleChoiceField(choices=BOOK_CHOICES)
text_area = forms.CharField(widget=forms.Textarea)
integer_input = forms.IntegerField()
float_input = forms.FloatField()
decimal_input = forms.DecimalField()
email_input = forms.EmailField()
date_input =
    forms.DateField(widget=forms.DateInput(attrs=
        {"type": "date"}))
hidden_input =
    forms.CharField(widget=forms.HiddenInput,
                   initial="Hidden Value")
```

Save the file.

6. Open your `form_example` app's `views.py` file. At the top of the file, add a line to import `ExampleForm` from your `forms.py` file:

```
from .forms import ExampleForm
```

7. Inside the `form_example` view, instantiate the `ExampleForm` class and assign it to the `form` variable:

```
form = ExampleForm()
```

8. Add the `form` variable to the context dictionary using the `form` key. The `return` line should look like this:

```
return render(request, "form-example.html",
{"method": request.method, "form": form})
```

Save the file.

Make sure you haven't removed the code that prints out the data the form has sent, as we will use it again later in this exercise.

9. Open the `form-example.html` file inside the `form_example` app's `templates` directory. You can remove nearly all of the contents of the `form` element, except the `{% csrf_token %}` template tag and the submit buttons. When you're done, it should look like this:

```
<form method="post">
{% csrf_token %}

<p>
    <input type="submit" name="submit_input"
           value="Submit Input">
</p>
<p>
    <button type="submit" name="button_element"
           value="Button Element">
        Button With <strong>Styled</strong> Text
    </button>
</p>
</form>
```

10. Add a rendering of the `form` variable using the `as_p` method. Put this on the line after the `{% csrf_token %}` template tag. The whole `form` element should now look like this:

```
<form method="post">
{% csrf_token %}
{{ form.as_p }}
<p>
```

```
<input type="submit" name="submit_input"
       value="Submit Input">
</p>
<p>
    <button type="submit" name="button_element"
            value="Button Element">
        Button With <strong>Styled</strong> Text
    </button>
</p>
</form>
```

11. Start the Django Dev Server if it is not already running, then visit the form example page in your browser, at <http://127.0.0.1:8000/form-example/>. It should look as follows:

Method: GET

Text input:

Password input:

Checkbox on:

Radio input:

Value One Display
 Text For Value Two
 Value Three's Display Text

Favorite book:

Non-Fiction
Deep Learning with Keras
Web Development with Django

Fiction
None listed

Books you own:

Text area:

Integer input:

Figure 6.24: Django ExampleForm rendered in the browser

12. Enter some data into the form – since Django marks all fields as required, you will need to enter some text or select values for all fields, including ensuring that the checkbox is checked. Submit the form.
13. Switch back to PyCharm and look in the Debug Console at the bottom of the window. You should see that all the values being submitted by the form are printed out to the console, similar to in *Exercise 6.02*:

```
csrfmiddlewaretoken: ['aozp1KTS...']  
text_input: ['Text']  
password_input: ['password']  
checkbox_on: ['on']  
radio_input: ['Value Two']  
favorite_book: ['1']  
books_you_own: ['1', '2']  
text_area: ['Some text in the text area.']  
integer_input: ['10']  
float_input: ['10.5']  
decimal_input: ['11.5']  
email_input: ['user@example.com']  
date_input: ['2022-04-20']  
hidden_input: ['Hidden Value']  
submit_input: ['Submit Input']
```

Figure 6.25: Values as submitted by the Django form

You can see that the values are still strings, and the names match those of the attributes of the `ExampleForm` class. Notice that the submit button that you clicked is included, as well as the CSRF token. The form you submit can be a mix of Django form fields and arbitrary fields you add; both will be contained in the `request . POST QueryDict` object.

In this exercise, you created a Django form with many different types of form fields. You instantiated it into a variable in your view, then passed it to `form-example.html`, where it was rendered as HTML. Finally, you submitted the form and looked at the values it posted. Notice that the amount of code we had to write to generate the same form was greatly reduced. We did not have to manually code any HTML, and we now have one place that defines how the form will display and how it will validate.

In the next section, we will examine how Django forms can automatically validate the submitted data, as well as how the data is converted from strings into Python objects.

Validating forms and retrieving Python values

So far, we have seen how Django forms make it much simpler to define a form using Python code and have it automatically rendered. We will now look at the other part of what makes Django forms useful: their ability to automatically validate the form and then retrieve native Python objects and values from them.

In Django, a form can either be *unbound* or *bound*. These terms describe whether or not the form has had the submitted POST data sent to it for validation. So far, we have only seen unbound forms – they are instantiated without arguments, like this:

```
form = ExampleForm()
```

A form is bound if it is called with some data to be used for validation, such as the POST data. A bound form can be created like this:

```
form = ExampleForm(request.POST)
```

A bound form allows us to start using built-in validation-related tools on the form instance. First, there's the `is_valid` method, which checks the form validity, then the `cleaned_data` attribute, which contains the values converted from strings into Python objects. The `cleaned_data` attribute is only available after the form has been *cleaned*, which is the process of cleaning up the data and converting it from strings into Python objects. The cleaning process runs during the `is_valid` call. An `AttributeError` will be raised if you try to access `cleaned_data` before calling `is_valid`.

A short example of how to access the cleaned data of `ExampleForm` looks like this:

```
form = ExampleForm(request.POST)

if form.is_valid(): # cleaned_data is only populated if the
    form is valid
    if form.cleaned_data["integer_input"] > 5:
        do_something()
```

In this example, `form.cleaned_data["integer_input"]` is the integer value 10, so it can be compared to the number 5. Compare this to the value that was posted, which is the string 10. The cleaning process performs this conversion for us. Other fields such as dates or Booleans are converted accordingly.

The cleaning process also sets any errors on the form and fields, which will be displayed when the form is rendered again. Let's see all this in action. Modern browsers provide a large amount of client-side validation, so they prevent forms from being submitted unless their basic validation rules are met. You might have already seen this if you tried to submit the form in the previous exercise with empty fields:

Form submission prevented by the browser

Address: 127.0.0.1:8000/form-example/

Errors:

- Checkbox on: Please fill out this field.

Fields:

- Text area: (empty)
- Integer input: (empty)
- Float input: (empty)
- Decimal input: (empty)

Figure 6.26: Form submission prevented by the browser

Figure 6.26 shows the browser preventing form submission. Since the browser is preventing the submission, Django never gets the opportunity to validate the form itself. To allow the form to be submitted, we need to add some more advanced validation that the browser is unable to validate itself. We will discuss the different types of validations that can be applied to form fields in the next section, but for now, we will just add a `max_digits` setting of 3 to `decimal_input` of our `ExampleForm`. This means the user should not enter more than three digits in the form.

Note

Why should Django validate the form if the browser is already doing this and preventing submission? A server-side application should never trust input from the user: the user might be using an older browser or other HTTP clients to send the request, thus not receiving any errors from their browser. Also, as we have just mentioned, there are types of validation that the browser does not understand, so Django must validate these on its end.

ExampleForm can be updated like this:

```
class ExampleForm(forms.Form):  
    ...  
    decimal_input = forms.DecimalField(max_digits=3)  
    ...
```

Now, the view should be updated to pass `request.POST` to the form class when the method is `POST`, for example, like this:

```
if request.method == "POST":  
    form = ExampleForm(request.POST)  
else:  
    form = ExampleForm()
```

If you pass `request.POST` into the form constructor when the method is not `POST`, then the form will always contain errors when first rendered as `request.POST` will be empty.

Now, the browser will let us submit the form, but we will get an error displayed if `decimal_input` contains more than three digits:

- Ensure that there are no more than 3 digits in total.



Figure 6.27: An error displayed when a field is not valid

Django automatically renders the form differently in the template when it has errors. But how can we make the view behave differently depending on the validity of the form? As we mentioned earlier, we should use the form's `is_valid` method. A view using this check might contain code like this:

```
form = ExampleForm(request.POST)  
  
if form.is_valid():  
    # perform operations from with data from  
    # form.cleaned_data  
    return redirect("/success-page")  # redirect to a  
    # success page
```

In this example, we are redirecting to a success page if the form is valid. Otherwise, we assume the execution flow continues as before and pass the invalid form back to the `render` function to be displayed to the user with errors.

Note

Why do we return a redirect on success? For two reasons: first, an early return prevents the execution of the rest of the view (that is, the failure branch); second, it prevents the message about resending the form data if the user then reloads the page.

In the next exercise, we will see the form validation in action and change the view's execution flow based on the validity of the form.

Exercise 6.04 – validating forms in a view

In this exercise, we will update the example view to instantiate the form differently depending on the HTTP method. We will also change the form so that it prints out the cleaned data instead of the raw POST data, but only if the form is valid:

1. In PyCharm, open the `forms.py` file inside the `form_example` app directory. Add a `max_digits=3` argument to the `ExampleForm`'s `decimal_input`:

```
class ExampleForm(forms.Form):
    ...
    decimal_input = forms.DecimalField(max_digits=3)
```

Once this argument has been added, we can submit the form since the browser does not know how to validate this rule, but Django does.

2. Open the `reviews` app's `views.py` file. We need to update the `form_example` view so that if the request's method is POST, `ExampleForm` is instantiated with the POST data; otherwise, it's instantiated without arguments. Replace the current form initialization with this code:

```
def form_example(request):
    if request.method == "POST":
        form = ExampleForm(request.POST)
    else:
        form = ExampleForm()
```

3. Next, also for the POST request method, we will check if the form is valid using the `is_valid` method. If the form is valid, we will print out all the cleaned data. Add a condition after the `ExampleForm` instantiation to check `form.is_valid()`, then move the debug print loop inside this condition. Your POST branch should look like this:

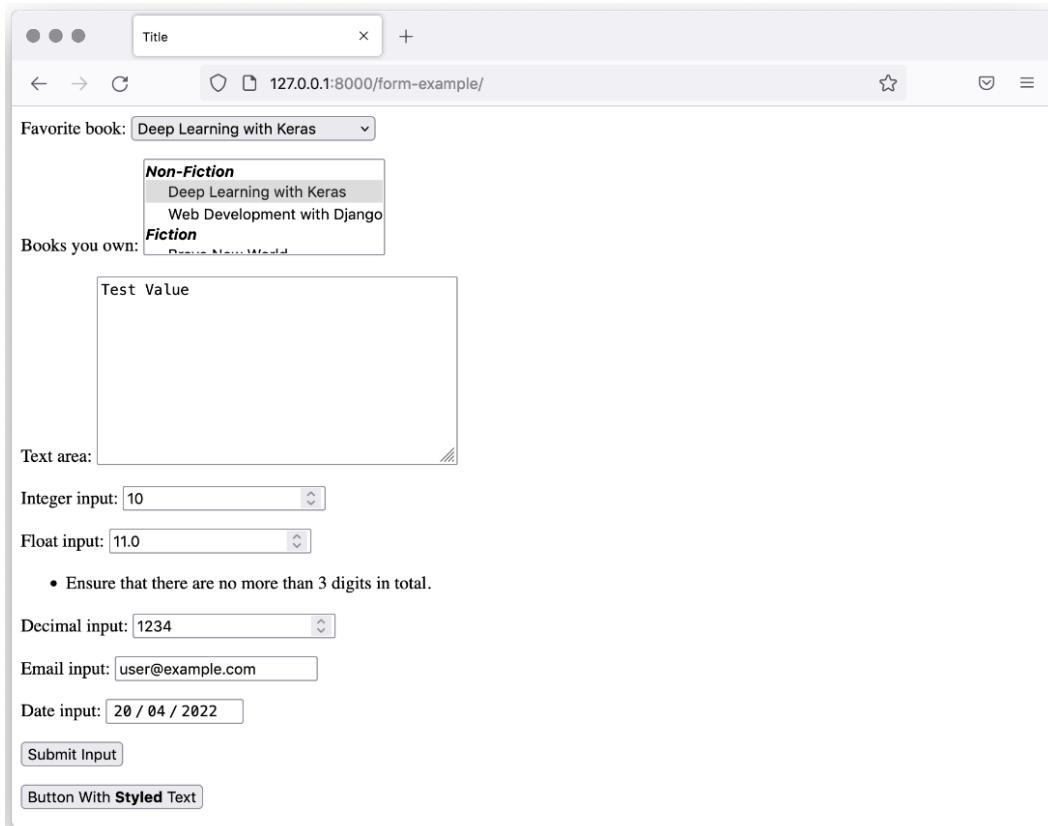
```
if request.method == "POST":
    form = ExampleForm(request.POST)
    if form.is_valid():
        for name in request.POST:
            print("{}: {}".format(name,
            request.POST.getlist(name)))
```

- Instead of iterating over the raw `request.POST` QueryDict (in which all the data is strings), we will iterate over the form's `cleaned_data`. This is a normal dictionary and contains the values converted into Python objects. Replace the `for` line and `print` line with these two:

```
for name, value in
    form.cleaned_data.items():
        print("{}: ({}) {}".format(name,
        type(value), value))
```

We don't need to use `getlist()` anymore since `cleaned_data` has already converted multi-value fields into lists.

- Start the Django Dev Server, if it's not already running. Switch to your browser and browse to the example form page at `http://127.0.0.1:8000/form-example/`. The form should look as it did before. Fill in all the fields, but make sure that you enter four or more numbers into the **Decimal input** field to make the form invalid. Submit the form; you should see the error message for the **Decimal input** field show up when the page refreshes:



The screenshot shows a web browser window with the title 'Title'. The address bar displays '127.0.0.1:8000/form-example/'. The page content is a Django form for 'form-example'. The form includes the following fields:

- Favorite book:** A dropdown menu showing 'Deep Learning with Keras' and 'Web Development with Django'. The 'Deep Learning with Keras' option is selected.
- Books you own:** A dropdown menu showing 'Non-Fiction' and 'Fiction'. The 'Non-Fiction' option is selected. A tooltip for 'Non-Fiction' lists 'Deep Learning with Keras' and 'Web Development with Django'. A tooltip for 'Fiction' lists 'Django Web Development' and 'Django REST Framework'.
- Text area:** A text area containing 'Test Value'.
- Integer input:** A text input field containing '10'.
- Float input:** A text input field containing '11.0'.
- Decimal input:** A text input field containing '1234'. An error message below it states: '• Ensure that there are no more than 3 digits in total.'
- Email input:** A text input field containing 'user@example.com'.
- Date input:** A text input field containing '20 / 04 / 2022'.
- Submit Input:** A button labeled 'Submit Input'.
- Button With Styled Text:** A button labeled 'Button With Styled Text'.

Figure 6.28: Decimal Input error displayed after the form is submitted

6. Fix the form errors by making sure only three digits are in the **Decimal input** field, then submit the form again. Switch back to PyCharm and check the Debug Console. You should see that all the cleaned data has been printed out:

```
text_input: (<class 'str'>) Text
password_input: (<class 'str'>) password
checkbox_on: (<class 'bool'>) True
radio_input: (<class 'str'>) Value One
favorite_book: (<class 'str'>) 1
books_you_own: (<class 'list'>) ['1']
text_area: (<class 'str'>) Test Value
integer_input: (<class 'int'>) 10
float_input: (<class 'float'>) 11.0
decimal_input: (<class 'decimal.Decimal'>) 123
email_input: (<class 'str'>) user@example.com
date_input: (<class 'datetime.date'>) 2022-04-20
hidden_input: (<class 'str'>) Hidden Value
```

Figure 6.29: The cleaned data from the form has been printed out

Notice the conversions that have taken place. The CharFields have been converted into `str`, `BooleanField` into `bool`, and `IntegerField`, `FloatField`, and `DecimalField` to `int`, `float`, and `Decimal`, respectively. `DateField` becomes a `datetime.date` and the choice fields will retain the string values of their initial choice values. Notice that `books_you_own` is automatically converted into a list of `str`.

Also, note that unlike when we iterated over all of the POST data, `cleaned_data` only contains form fields. The other data (such as the CSRF token and the submit button that was clicked) is present in the POST's `QueryDict` but is not included as they are not form fields.

In this exercise, you updated `ExampleForm` so that the browser allowed it to be submitted even though Django would consider it to be invalid. This allowed Django to perform its validation on the form. You then updated the `form_example` view to instantiate the `ExampleForm` class differently, depending on the HTTP method, passing in the request's POST data for a POST request. The view also had its debug output code updated to print out the `cleaned_data` dictionary. Finally, you tested submitting valid and invalid form data to see the different execution paths and the types of data that the form generated. We saw that Django automatically converted the POST data from strings into Python types based on the field class.

Next, we will look at how to add more validation options to fields, which will allow us to more tightly control the values that can be entered.

Built-in field validation

We have not yet discussed the standard validation arguments that can be used on fields. Although we already mentioned the `required` argument (which is `True` by default), many others can be used to control the data being entered into a field more tightly. Here are a few useful ones:

- `max_length`

Sets the maximum number of characters that can be entered into the field, available on a `CharField` (and `FileField` – which we will cover in *Chapter 8*).

- `min_length`

Sets the minimum number of characters that must be entered into the field, available on `CharField` (and `FileField` – again, more about it in *Chapter 8*).

- `max_value`

Sets the maximum value that can be entered into a numeric field, available on `IntegerField`, `FloatField`, and `DecimalField`.

- `min_value`

Sets the minimum value that can be entered into a numeric field, available on `IntegerField`, `FloatField`, and `DecimalField`.

- `max_digits`

This sets the maximum number of digits that can be entered. This includes before and after a decimal point (if one exists). For example, the number `12.34` has four digits, while the number `56.7` has three. This is used in a `DecimalField`.

- `decimal_places`

This sets the maximum number of digits that can be after the decimal point. This is used in conjunction with `max_digits`, and the number of decimal places will always count toward the number of digits, even if that number of decimals has not been entered after the decimal place. For example, using `max_digits` of four and `decimal_places` of three: if the number `12.34` was entered, it would be interpreted as the value `12.340` – that is, zeroes are appended until the number of digits after the decimal point is equal to the `decimal_places` setting. Since we set three as the value for `decimal_places`, the total number of digits ends up being five, which exceeds the `max_digits` setting of four. The number `1.2` would be valid since even after expanding to `1.200`, the total number of digits is only four.

You can mix and match the validation rules (provided the fields support them). `CharField` can have a `max_length` and a `min_length`; numeric fields can have both a `min_value` and `max_value`, and so on.

If you need more validation options, you can write custom validators, which we will cover in the next section. Right now, we will add some validators to our `ExampleForm` to see them in action.

Exercise 6.05 – adding extra field validation

In this exercise, we will add and modify the validation rules on the fields of `ExampleForm`. We will then see how these changes affect how the form behaves, both in the browser and when Django validates the form:

1. In PyCharm, open the `forms.py` file inside the `form_example` app directory.
2. We will make `text_input` require at most three characters. Add a `max_length=3` argument to the `CharField` constructor:

```
text_input = forms.CharField(max_length=3)
```

3. Make `password_input` more secure by requiring a minimum of eight characters. Add a `min_length=8` argument to the `CharField` constructor:

```
password_input = forms.CharField(min_length=8,
                                 widget=forms.PasswordInput)
```

4. The user may have no books, so the `books_you_own` field should not be required. Add a `required=False` argument to the `MultipleChoiceField` constructor:

```
books_you_own =
forms.MultipleChoiceField(required=False,
                           choices=BOOK_CHOICES)
```

5. The user should only be able to enter a value between 1 and 10 in `integer_input`. Add `min_value=1` and `max_value=10` arguments to the `IntegerField` constructor:

```
integer_input = forms.IntegerField(min_value=1,
                                    max_value=10)
```

6. Finally, add `max_digits=5` and `decimal_places=3` to the `DecimalField` constructor:

```
decimal_input = forms.DecimalField(max_digits=5,
                                    decimal_places=3)
```

Save the file.

7. Start the Django Dev Server if it's not already running. We do not have to make any changes to any other files to get these new validation rules since Django automatically updates the HTML generation and validation logic. This is a great benefit you get from using Django Forms. Just visit or refresh `http://127.0.0.1:8000/form-example/` in your browser and the new validation will be automatically added. The form should not look any different until you try to submit it with incorrect values, in which case your browser can automatically show errors. Some things to try are as follows:

- Enter more than three characters into the **Text input** field is something you will not be able to do.
- Type less than eight characters into the password field and then clicking away from it might cause your browser to show an error indicating that this is not valid (only some browsers support this).
- Do not select any values for the **Books you own** field. This will not prevent you from submitting the form anymore.
- Use the stepper buttons on **Integer input**. You will only be able to enter a value between 1 and 10. If you type in a value outside this range, your browser should show an error.
- **Decimal input** is the only field that does not validate the Django rules in the browser. You will need to type in an invalid value (such as `123.456`) and submit the form before an error (generated by Django) is displayed.

The following figure shows some of the fields that the browser can use to validate itself:

Integer input: (up/down arrows)

Float input: (up/down arrows)

Decimal input: (up/down arrows)

Figure 6.30: Some browsers highlight invalid fields

Figure 6.31 shows an alternate way that validation errors might be presented by the browser – by placing an error message next to a particular field on form submit:

Password input:

Checkbox on: Please use at least 8 characters (you are currently using 4 characters).

Radio input:

Figure 6.31: Browsers might show messages next to fields on submit

Figure 6.32 shows an error that can only be generated by Django as the browser does not understand the `DecimalField` validation rules:

- Ensure that there are no more than 3 digits in total.

Decimal input: 

Figure 6.32: The browser considers the form valid, but Django does not

In this exercise, we implemented some basic validation rules on our form fields. We then loaded the form example page in the browser without having to make any changes to our template or view. We tried to submit the form with different values to check how the browser can validate the form compared to Django.

In the activity for this chapter, we will implement the Book Search view using a Django Form.

Activity 1 – Book Search

In this activity, you will finish the Book Search view that was started in *Chapter 1*. You will build a `SearchForm` that submits and accepts a search string from `request.GET`. It will have a `select` field to choose to search for a `title` or `contributor`. It will then search for all books containing the given text in their `title` or contributor's `first_names` or `last_names`. You will then render this list of books in the `search-results.html` template. The search term should not be required, but if it exists, it should have a length of more than three characters. Since the view will search even when using the `GET` method, the form will always have its validation checked. If we made the field required, it would always show an error whenever the page loads.

There will be two ways of performing the search. The first is by submitting the search form that is in the `base.html` template and thus in the top-right corner of every page. This will only search through Book titles. The other method is by submitting a `SearchForm` that is rendered on the `search-results.html` page. This form will display a `ChoiceField` for selecting between `title` or `contributor` search.

These steps will help you complete this activity:

1. Create a `forms.py` file, import the Django forms library, then create a `SearchForm` class.
2. `SearchForm` should have two fields. The first is a `CharField` with the name `search`. This field should not be required but should have a minimum length of 3.
3. The second field on `SearchForm` is a `ChoiceField` named `search_in`. This will allow you to select between `title` and `contributor` (with labels of **Title** and **Contributor**, respectively). It should not be required.

4. Update the `book_search` view so that it instantiates a `SearchForm` using data from `request`. GET.
5. Add code to search for `Book` models using `title__icontains` (for case-insensitive searching). This should be done if searching by *title*. The search should only be performed if the form is valid and contains some search text. The `search_in` value should be retrieved from `cleaned_data` using the GET method since it might not exist as it's not required. Default it to `title`.
6. When searching for contributors, use `first_names__icontains` or `last_names__icontains`, then iterate the contributors and retrieve the books for each contributor. This should be done if you're searching by *contributor*. The search should only be performed if the form is valid and contains some search text. There are many ways to combine the search results for first or last names. The easiest method using the techniques that you have been introduced to so far is to perform two queries, one for matching first names and then for last names, and iterating them separately.
7. Update the `render` call so that it includes the `form` variable and books that were retrieved in the context (as well as `search_text`, which was already being passed). The location of the template was changed in *Chapter 3*, so update the second argument to `render` accordingly.
8. The `search-results.html` template we created in *Chapter 1* is essentially redundant now, so you can clear its content. Update `search-results.html` so that it extends from `base.html` instead of being a standalone template file.
9. Add a *title* block that will display `Search Results for <search_text>` if the form is valid and `search_text` was set; otherwise, just display `Book Search`. This block will also be added to `base.html` later in this activity.
10. Add a *content* block, which should show a `<h2>` heading with the text **Search for Books** under the `<h2>` render the form. The `<form>` element can have no attributes and it will default to making a GET request to the same URL it's on. Add a submit button, as we have used in previous activities, with the `btn btn-primary` class.
11. Under the form, show a message stating `Search results for <search_text>` if the form is valid and search text was entered; otherwise, show no message. This should be displayed in `<h3>` tags, and the search text should be wrapped in `` tags.
12. Iterate over the search results and render each one. Show the book title and any contributors' first and last names. The book title should link to the `book_detail` page. If the books list is empty, show the text **No results found**. You should wrap the results in `` tags with a class called `list-group`; each result should be a `` item with a class of `list-group-item`. This will be like the `book_list` page, but we won't show as much information (just the title and contributors).

13. Update `base.html` so that it includes an `action` attribute in the search `<form>` tag. Use the `url` template tag to generate the URL for this attribute.
14. Set the `name` attribute of the search field to `search`, and the `value` attribute to the entered `search_text`. Also, ensure the minimum length of the field is 3.
15. While in `base.html`, add a title block to the `title` tag that was overridden by other templates (as in *step 9*). Add a `block` template tag inside the `<title>` HTML element. It should contain the content **Bookr**.

After completing this activity, you should be able to open the Book Search page at `http://127.0.0.1:8000/book-search/`; it will look like *Figure 6.33*:

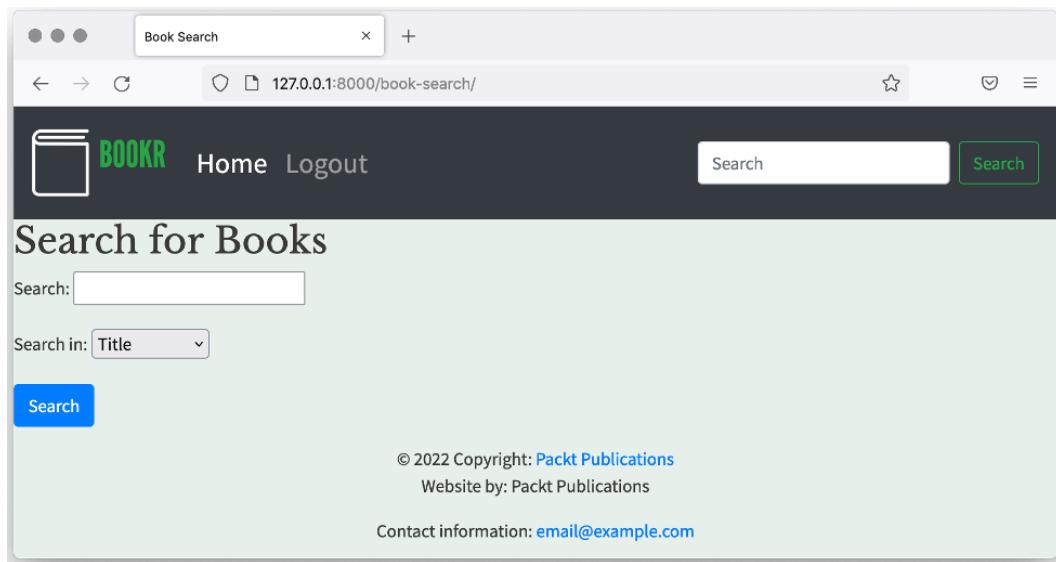


Figure 6.33: The Book Search page without a search

When searching for something using just two characters, your browser should prevent you from submitting either of the search fields.

If you search for something that returns no results, you will see a message that there were no results. Searching by title (this can be done with either field) will show matching results:

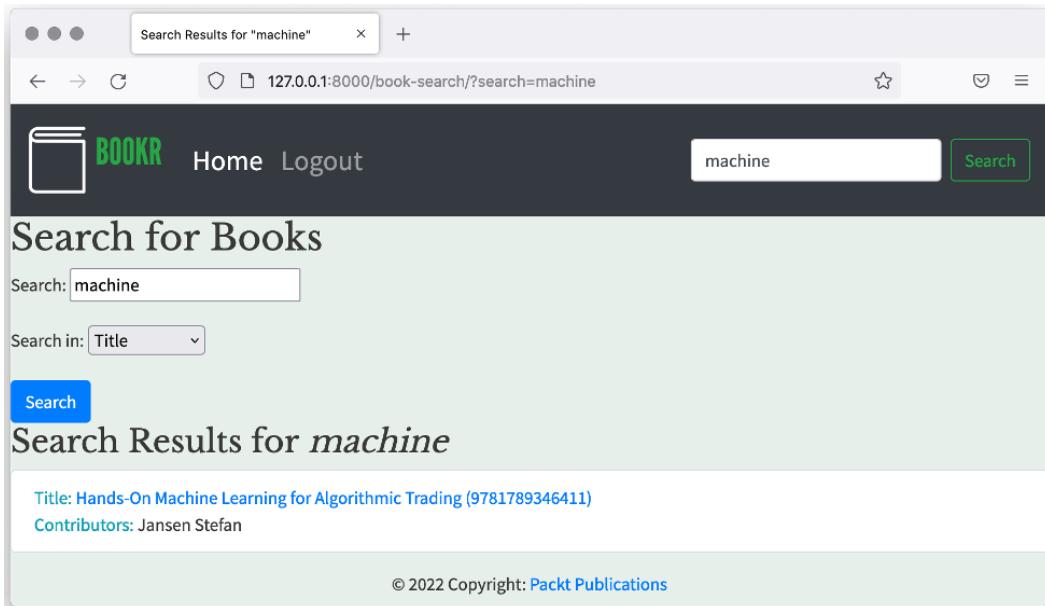


Figure 6.34: Title search

Similarly, when searching by contributor (although this can only be done in the lower form), you'll see matching results:

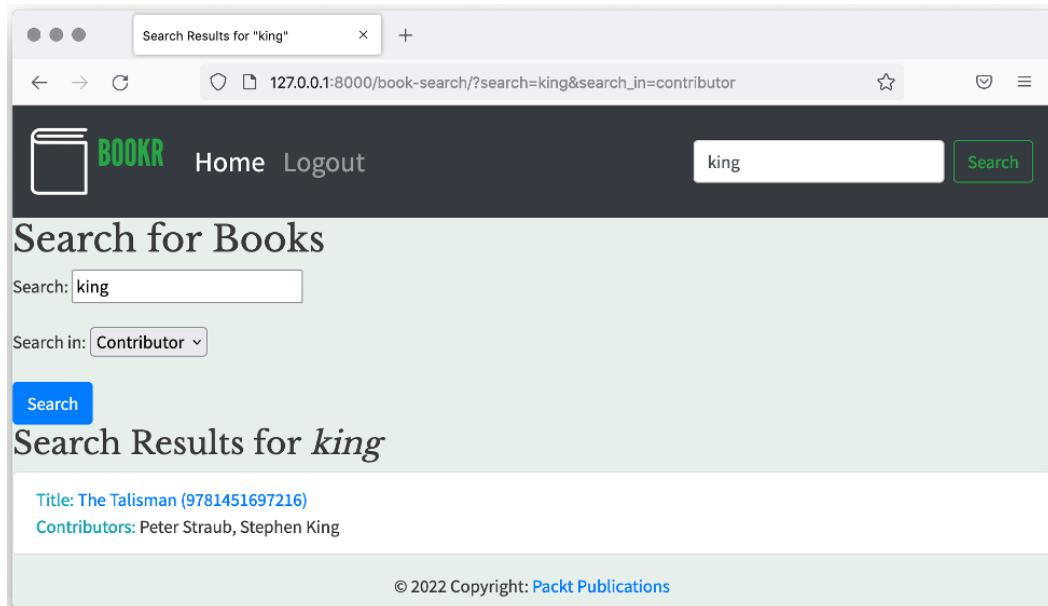


Figure 6.35: A contributor search

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

This chapter provided an introduction to forms in Django. We introduced several HTML inputs for entering data onto a web page. We talked about how data is submitted to a web application and when to use GET and POST requests. We then looked at how Django's form classes can make generating the form HTML simpler, as well as allow us to automatically build forms using models. Finally, we enhanced Bookr some more by building the Book Search functionality.

In the next chapter, we will dive deeper into forms and learn how to customize the display of form fields, how to add more advanced validation to a form, and how to automatically save model instances by using the `ModelForm` class.

Advanced Form Validation and Model Forms

Continuing your journey with the `form_project` application you started in the previous chapter, you'll begin this chapter by adding a new form to your app with custom multi-field validation and form cleaning. You'll learn how to set the initial values on your form and customize the widgets (the HTML input elements that are being generated). Then, you'll be introduced to the `ModelForm` class, which allows a form to be automatically created from a model. You'll use it in a view to automatically save the new or changed `Model` instance.

In this chapter, we will take our knowledge of Django form validation further by introducing concepts that form the fundamentals of most production websites.

For instance, a certain field might only be required if another field is set. Let's say we want to add a checkbox to allow users to sign up for our monthly newsletter. It has a text box below it that lets them enter their email address. With some basic validation, we can check the following:

- Whether the user has checked the checkbox
- Whether the user has entered their email address

When the user clicks the **Submit** button, we'll be able to validate whether both fields are actioned. But what if the user doesn't want to sign up for our newsletter? If they click the **Submit** button, ideally, both fields should be blank. That's where validating each field might not work.

Another example could be a case where we have two fields where each has a maximum value of, say, 50. But the total values added to each one must be less than 75. We'll start this chapter by looking at how to write custom validation rules to solve such problems.

Later, as we progress in this chapter, we will look at how to set initial values on a form. This can be useful when automatically filling out information that is already known to the user. For example, we can automatically put a user's contact information into a form if that user is logged in.

We'll finish this chapter by looking at model forms, which will let us automatically create a form from a Django `Model` class. This cuts down the amount of code that needs to be written to create a new `Model` instance.

In this chapter, we will cover the following topics:

- Custom field validation and cleaning
- Adding placeholders and initial values
- Creating or editing Django models

By the end of this chapter, you will know how to add extra multi-field validation to Django forms, how to customize and set form widgets for fields, how to use `ModelForms` to automatically create a form from a Django model, and how to automatically create `Model` instances from `ModelForms`.

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter07>

Custom field validation and cleaning

We have seen how a Django form converts values from an HTTP request, which are strings, into Python objects. In a non-custom Django form, the target type is dependent on the field class. For example, the Python type derived from `IntegerField` is `int`, and string values are given to us verbatim, as the user entered them. But, we can also implement methods on our `Form` class to alter the output values from our fields in any way we choose. This allows us to clean or filter the user's input data to fit what we expect better. We could round an integer to the nearest multiple of 10 so that it fits into a batch size for ordering specific items. Or, we could transform an email address into lowercase so that the data is consistent for searching.

We can also implement some custom validators. We'll look at a couple of different ways of validating fields: by writing a custom validator, and by writing a custom `clean` method for the field. Each method has its pros and cons: a custom validator can be applied to different fields and forms, so you don't have to write the validation logic for each field; a custom `clean` method must be implemented on each form you want to clean, but is more powerful and allows validation since you can use other fields in the form or change the cleaned value that the field returns.

Custom validators

A validator is simply a function that accepts a value and raises `django.core.exceptions.ValidationError` if the value is invalid – the validity is determined by the code you write. The value is a Python object (that is, `cleaned_data`, which has already been converted from the POST request string).

Here's a simple example that validates whether a value is in lowercase:

```
from django.core.exceptions import ValidationError

def validate_lowercase(value):
    if value.lower() != value:
        raise ValidationError("{} is not lowercase."
                             .format(value))
```

Notice the function does not return anything, for either success or failure. It will just raise `ValidationError` if the value is not valid.

Note

Note that the behavior and handling of `ValidationError` are different than how we've seen other exceptions behave in Django. Normally, if you raise an exception in your view, you'll end up with an HTTP response with a 500 status code from Django (if you don't handle the exception in your code).

When raising `ValidationError` in your validation/cleaning code, the Django `Form` class will catch the error for you; then, the `is_valid` method of `form` will return `False`. You don't have to write `try/except` handlers around the code that might raise `ValidationError`.

The validator can be passed to the `validators` argument of a field constructor on a form, inside a list; for example, to our `text_input` field from our `ExampleForm`:

```
class ExampleForm(forms.Form):
    text_input =
        forms.CharField(validators=[validate_lowercase])
```

Now, if we submit the form and the fields contain uppercase values, we'll get an error, as shown in the following figure:

- Text is not lowercase.

Text input:

Figure 7.1 – Lowercase text validator in action

The validator function can be used on any number of fields. In our example, if we wanted lots of fields to have lowercase enforced, `validate_lowercase` could be passed to all of them. Now, let's look at how we could implement this another way, with a custom `clean` method.

Cleaning methods

A clean method is created on the `Form` class and is named in the `clean_field_name` format. For example, the clean method for `text_input` would be called `clean_text_input`, the clean method for `books_you_own` would be called `clean_books_you_own`, and so on.

Cleaning methods take no arguments; instead, they should use the `cleaned_data` attribute on `self` to access the field data. This dictionary will contain the data after being cleaned in the standard Django way, as we saw in the previous example. The clean method must return the cleaned value, which will replace the original value in the `cleaned_data` dictionary. Even if the method does not change the value, a value must be returned. You can also use the clean method to raise `ValidationError`, and the error will be attached to the field (the same as with a validator).

Let's re-implement the lowercase validator as a clean method, like this:

```
class ExampleForm(forms.Form):
    text_input = forms.CharField()
    ...

    def clean_text_input(self):
        value = self.cleaned_data["text_input"]
        if value.lower() != value:
            raise ValidationError("{} is not lowercase."
                                  .format(value))
        return value
```

You can see that the logic is essentially the same, except we must return the validated value at the end. If we submit the form, we will get the same result as the previous time we tried (*Figure 7.1*).

Let's look at one more cleaning example. Instead of raising an exception when the value is invalid, we could just convert the value into lowercase. We would implement that with this code:

```
class ExampleForm(forms.Form):
    text_input = forms.CharField()
    ...

    def clean_text_input(self):
        value = self.cleaned_data['text_input']
        return value.lower()
```

Now, let's say we enter text into the input as uppercase:

Text input:

Figure 7.2 – ALL UPPERCASE text entered

If we were to examine the cleaned data using our debug output from the view, we would see that it is lowercase:

```
text_input: (<class 'str'>) all uppercase
```

Figure 7.3 – The cleaned data has been transformed into lowercase

These were just a couple of simple examples of how to validate fields using both validators and clean methods. You can, of course, make each type of validation much more complex if you wish and transform the data in more complex ways using a clean method.

So far, you've only learned simple methods for form validation, where you've treated each field independently. A field is valid (or not) based only on the information it contains and nothing else. What if the validity of one field depends on what the user entered into another field? An example of this might be that you have an `email` field to collect someone's email address if they want to be signed up for a mailing list. The field is only required if they check a checkbox that indicates they wanted to be signed up. Neither of these fields is required on their own – we do not want the checkbox to be required to be checked, but if it is checked, then the `email` field should be required too.

In the next section, we will show how you can validate a form whose fields depend on each other by overriding the `clean` method in your form.

Multi-field validation

We've just looked at the `clean_field_name` methods that can be added to a Django form, to clean a specific field. Django also allows us to override the `clean` method, in which we can access all `cleaned_data` from all fields, and we know that all custom field methods have been called. This allows the fields to be validated based on another field's data.

Referring to our previous example with a form that has an email address that is only required if a checkbox is checked, we'll see how we can implement this using the `clean` method.

First, create a `Form` class and add two fields – make them both optional with the `required=False` argument:

```
class NewsletterSignupForm(forms.Form):  
    signup = forms.BooleanField(label="Sign up to  
    newsletter?", required=False)  
    email = forms.EmailField(help_text="Enter your email  
    address to subscribe", required=False)
```

We've also introduced two new arguments that can be used for any field:

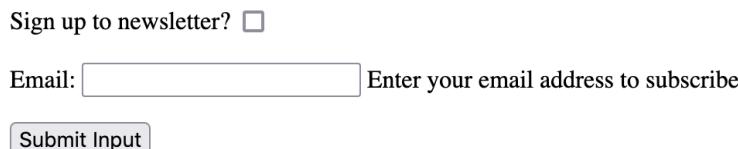
- `label`

This allows you to set the label text for a field. As we've seen, Django will automatically generate label text from the field name. If you set the `label` argument, you can override this default. Use this argument if you want to have a more descriptive label.

- `help_text`

If you need more information to be displayed regarding what input a field requires, you can use this argument. By default, it's displayed after the field.

When rendered, the form looks like this:



Sign up to newsletter?

Email: Enter your email address to subscribe

Figure 7.4 – Email signup form with a custom label and help text

If we were to submit the form now, without entering any data, nothing would happen. Neither field is required, so the form validates fine.

Now, we can add the multi-field validation to the `clean` method. We will check whether the `Sign up to newsletter` checkbox is checked, and then check that the `Email` field has a value. The built-in Django methods have already validated that the email address is valid at this point, so we just need to check that a value exists for it. We will then use the `add_error` method to set an error for the `email` field. This is a method you haven't seen before but it's very simple; it takes two arguments – the name of the field to set the error on, and the text of the error.

Here's the code for the `clean` method:

```
class NewsletterSignupForm(forms.Form):  
    ...  
  
    def clean(self):  
        cleaned_data = super().clean()  
  
        if cleaned_data["signup"] and not  
        cleaned_data.get("email"):  
            self.add_error("email", "Your email address is  
            required if signing up for the newsletter.")
```

Your `clean` method must always call the `super().clean()` method to retrieve the cleaned data. When `add_error` is called to add errors to the form, the form will no longer validate (the `is_valid` method returns `False`).

Now, if we submit the form without the checkbox checked, still, no error is generated, but if you check the checkbox without an email address, you will receive the error we just wrote the code for:

Sign up to newsletter?

- Enter your email address to subscribe

Email: Enter your email address to subscribe

Figure 7.5 – Error displayed when attempting to sign up with no email address

You might notice that we are retrieving the email from the `cleaned_data` dictionary using the `get` method. The reason for doing this is that if the `email` value in the form is invalid, then the `email` key will not exist in the dictionary. The browser should prevent the user from submitting the form if an invalid email has been entered, but a user might be using an older browser that does not support this client-side validation, so, for safety, we use the `get` method. Since the `signup` field is `BooleanField`, and not required, it will only be invalid if a custom validation function is used. We're not using one here, so it's safe to access its value using square bracket notation.

There's one more validation scenario to consider before moving on to our first exercise, and that's adding errors that are not specific to any field. Django calls these *non-field errors*. There are many scenarios where you might want to use these when multiple fields are dependent on each other.

Take, for example, a shopping website. Your order form could have two numeric fields whose totals can't exceed a certain value. If the total were exceeded, the value of either field could be decreased to bring the total below the maximum value, so the error is not specific to either of the fields. To add a non-field error, call the `add_error` method with `None` as the first argument.

Let's look at how to implement this. In this example, we'll have a form where the user can specify a certain number of items to order, for item A or item B. The user cannot order more than 100 items in total. The fields will have a `max_value` of 100, and a `min_value` of 0, but custom validation in the `clean` method will need to be written to handle the validation of the total amount:

```
class OrderForm(forms.Form):
    item_a = forms.IntegerField(min_value=0, max_value=100)
    item_b = forms.IntegerField(min_value=0, max_value=100)

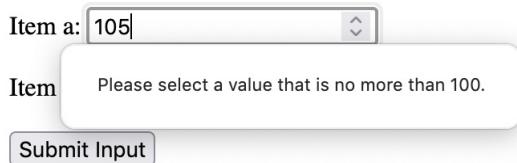
    def clean(self):
        cleaned_data = super().clean()
        if cleaned_data.get("item_a", 0) +
            cleaned_data.get("item_b", 0) > 100:
            self.add_error(None, "The total number of items
            must be 100 or less.")
```

The fields (`item_a` and `item_b`) are added in the normal way, with standard validation rules. You can see that we have used the `clean` method the same way we used it before. Moreover, we have implemented the maximum item logic inside this method. The following line is what registers the non-field error if the maximum items are exceeded:

```
    self.add_error(None, "The total number of items must be 100
            or less.")
```

Once again, we access the values of `item_a` and `item_b` using the `get` method, with a default value of 0. This is in case the user has an older browser (from 2011 or earlier) and was able to submit the form with invalid values.

In a browser, the field-level validation ensures values between 0 and 100 have been entered in each field, and prevents the form from being submitted otherwise:



Item a: 105

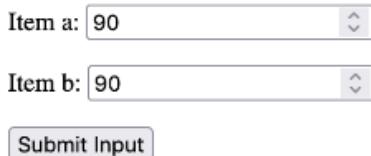
Item a: Please select a value that is no more than 100.

Submit Input

Figure 7.6 – The form cannot be submitted if one field exceeds the maximum value

However, if we put in two values that sum to more than 100, we can see how Django displays the non-field error:

- The total number of items must be 100 or less.



Item a: 90

Item b: 90

Submit Input

Figure 7.7 – Django non-field error displayed at the start of the form

Django non-field errors are always displayed at the start of a form, before other fields or errors.

In the following exercise, we will build a form that implements a validation function, a field clean method, and a form clean method.

Exercise 7.01 – custom clean and validation methods

In this exercise, you will build a new form that allows the user to create an order for books or magazines. It must have the following validation criteria:

- The user may order up to 80 magazines and/or 50 books, but the total number of items must not be more than 100
- The user can choose to receive an order confirmation, and if they do, they must enter an email address
- The user should not enter an email address if they have not chosen to receive an order confirmation

- To ensure they are part of our company, the email address must be part of our company domain (in our case, we'll just use `example.com`)
- For consistency with other email addresses in our fictional company, the address should be converted into lowercase

This sounds like a lot of rules, but with Django, it's simple if we tackle them one by one. We'll carry on with the `form_project` app you started in *Chapter 6, Forms*. If you haven't completed *Chapter 6, Forms*, you can download the code from https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter06/final/form_project:

1. In PyCharm, open the `form_example` app's `forms.py` file.

Note

Make sure the Django dev server is not running; otherwise, it may crash as you make changes to this file, causing PyCharm to jump into the debugger.

Since our work with `ExampleForm` is done, you can remove it from this file.

2. Create a new class called `OrderForm` that inherits from `forms.Form`:

```
class OrderForm(forms.Form):
```

3. Add four fields to the class, as follows:

- `magazine_count`, `IntegerField` with a `min_value` of 0 and a `max_value` of 80
- `book_count`, `IntegerField` with a `min_value` of 0 and a `max_value` of 50
- `send_confirmation`, `BooleanField`, which is not required
- `email`, `EmailField`, which is also not required

The class should look like this:

```
class OrderForm(forms.Form):
    magazine_count = forms.IntegerField(min_value=0,
                                         max_value=80)
    book_count = forms.IntegerField(min_value=0,
                                    max_value=50)
    send_confirmation =
        forms.BooleanField(required=False)
    email = forms.EmailField(required=False)
```

4. Add a validation function to check that the user's email address is on the right domain. First, `ValidationError` needs to be imported; add this line to the top of the file:

```
from django.core.exceptions import ValidationError
```

Then, write this function after the `import` line (before the `OrderForm` class implementation):

```
def validate_email_domain(value):  
    if value.split("@")[-1].lower() != "example.com":  
        raise ValidationError("The email address must  
        be on the domain example.com.")
```

The function splits the email address on the `@` symbol, then checks whether the part after it is equal to `example.com`. This function alone would validate non-email addresses. For example, the `not-valid@some other domain@example.com` string would not cause a `ValidationError` to be raised in this function. This is acceptable in our case because, since we're using `EmailField`, the other standard field validators will check the email address's validity.

5. Add the `validate_email_domain` function as a validator to the `email` field on `OrderForm`. Update the `EmailField` constructor call to add a `validators` argument, passing in a list containing the validation function:

```
class OrderForm(forms.Form):  
    ...  
    email = forms.EmailField(required=False,  
                            validators=[validate_email_domain])
```

6. Add a `clean_email` method to the form to make sure the email address is lowercase:

```
class OrderForm(forms.Form):  
    # truncated for brevity  
  
    def clean_email(self):  
        return self.cleaned_data['email'].lower()
```

7. Now, add the `clean` method to perform all the cross-field validation. First, we will just add the logic for making sure that an email address is only entered if an order confirmation is requested:

```
class OrderForm(forms.Form):  
    # truncated for brevity  
    def clean(self):  
        cleaned_data = super().clean()  
        if cleaned_data["send_confirmation"] and not  
        cleaned_data.get("email"):
```

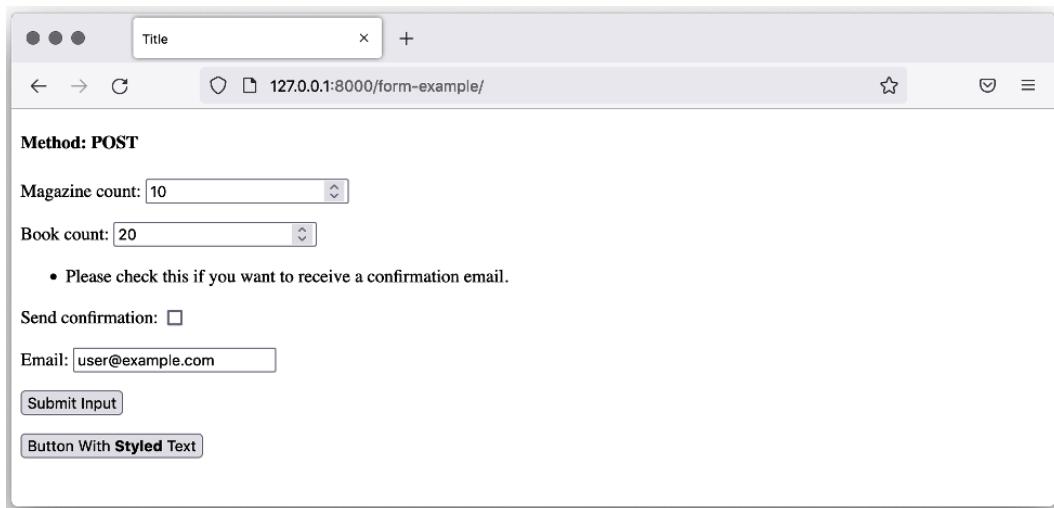
```
        self.add_error("email", "Please enter an
                      email address to receive the confirmation
                      message.")
    elif cleaned_data.get("email") and not
        cleaned_data["send_confirmation"]:
        self.add_error("send_confirmation",
                      "Please check this if you want to receive
                      a confirmation email.")
```

This will add an error to the **Email** field if **Send confirmation** is checked but no email address is added:



Figure 7.8 – Error if Send confirmation is checked but no email address is added

Similarly, an error will be added to **Email** if an email address is entered but **Send confirmation** is not checked:



Method: POST

Magazine count:

Book count:

- Please check this if you want to receive a confirmation email.

Send confirmation:

Email:

Figure 7.9 – Error because an email has been entered but the user has not chosen to receive confirmation

8. Add the final check, also inside the `clean` method. The total number of items should not be more than 100. We will add a non-field error if the sum of `magazine_count` and `book_count` is greater than 100:

```
class OrderForm(forms.Form):  
    ...  
    def clean(self):  
        ...  
        item_total =  
            cleaned_data.get("magazine_count", 0) +  
            cleaned_data.get("book_count", 0)  
  
        if item_total > 100:  
            self.add_error(None, "The total number of  
            items must be 100 or less.")
```

This will add a non-field error by passing `None` as the first argument to the `add_error` call.

Note

Refer to https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter07/Exercise7.01/form_project/form_example/forms.py for the complete code.

Save `forms.py`.

9. Open the `reviews` app's `views.py` file. We will change the form's import so that `OrderForm` is being imported instead of `ExampleForm`. Consider the following import line:

```
from .forms import ExampleForm
```

Change it as follows:

```
from .forms import OrderForm
```

10. In the `form_example` view, change the two lines that use `ExampleForm` to use `OrderForm` instead. Consider the following line of code:

```
form = ExampleForm(request.POST)
```

Change this as follows:

```
form = OrderForm(request.POST)
```

Similarly, consider the following line of code:

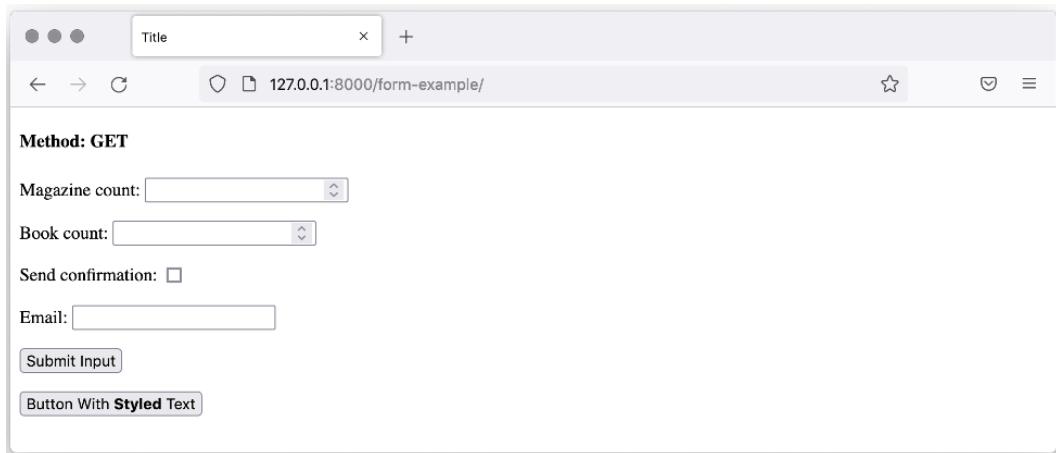
```
form = ExampleForm()
```

Change this as follows:

```
form = OrderForm()
```

The rest of the function can stay as-is.

We don't have to make changes to the template. Start the Django dev server and navigate to `http://127.0.0.1:8000/form-example/` in your browser. You should see the form rendered similarly to what's shown in *Figure 7.10*:



A screenshot of a web browser window titled "Title". The address bar shows "127.0.0.1:8000/form-example/". The page content is labeled "Method: GET". It contains the following form fields:

- Magazine count:
- Book count:
- Send confirmation:
- Email:
- Submit Input
- Button With **Styled Text**

Figure 7.10 – OrderForm in the browser

11. Try submitting the form with a **Magazine count** of 80 and a **Book count** of 50. The browser will allow this but since they sum to more than 100, an error will be triggered by the `clean` method in the form and displayed on the page:



A screenshot of a web browser window titled "Title". The address bar shows "127.0.0.1:8000/form-example/". The page content is labeled "Method: POST". It contains the following form fields:

- Magazine count:
- Book count:
- Send confirmation:
- Email:
- Submit Input
- Button With **Styled Text**

A validation error message is displayed above the form fields:

- The total number of items must be 100 or less.

Figure 7.11 – A non-field error displayed on the form when the maximum number of allowed items is exceeded

12. Try submitting the form with **Send confirmation** checked but the **Email** field blank. Then, fill in the **Email** text box but uncheck **Send confirmation**. Either combination will give an error stating that both must be present. The error will differ based on which field is missing:



The screenshot shows a web browser window with a title bar 'Title' and a URL bar '127.0.0.1:8000/form-example/'. The page content is a form with the following fields:

- Method: POST**
- Magazine count:** 20
- Book count:** 20
- Send confirmation:**
- Email:**

Below the 'Email' field, there is an error message: **• Please enter an email address to receive the confirmation message.**

At the bottom of the form are two buttons: **Submit Input** and **Button With Styled Text**.

Figure 7.12 – Error message if no email address is present

13. Now, try submitting the form with **Send confirmation** checked, and an email address that is on the `example.com` domain. You should receive a message that your email address must have the `example.com` domain. You should also receive a message that `email` must be set – this is because the email does not end up in the `cleaned_data` dictionary. After all, it is not valid:



The screenshot shows a web browser window with a title bar 'Title' and a URL bar '127.0.0.1:8000/form-example/'. The page content is a form with the following fields:

- Method: POST**
- Magazine count:** 20
- Book count:** 20
- Send confirmation:**
- Email:**

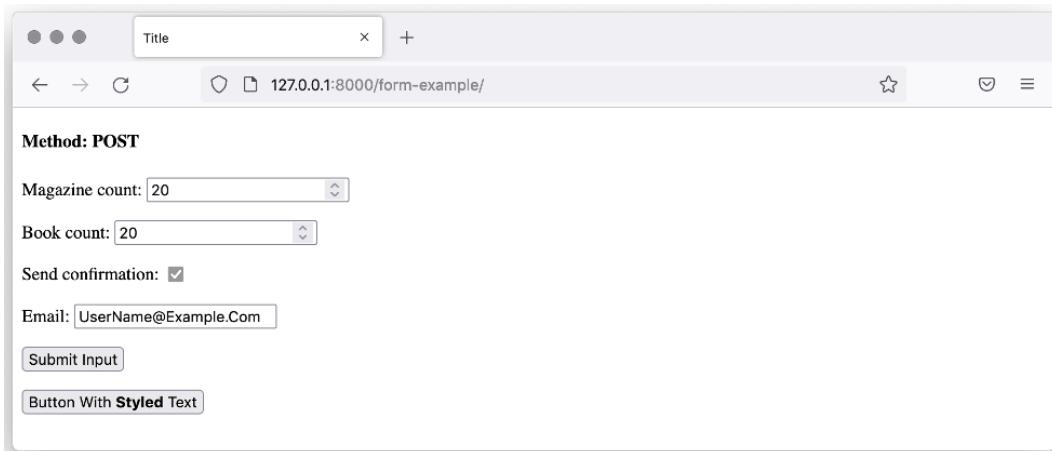
Below the 'Email' field, there are two error messages:

- The email address must be on the domain example.com.**
- Please enter an email address to receive the confirmation message.**

At the bottom of the form are three buttons: **Submit Input**, **Button With Styled Text**, and a third button with a different style.

Figure 7.13 – The error message that's shown when the email domain is not example.com

- Finally, enter valid values for **Magazine count** and **Book count** (such as 20 and 20). Check **Send confirmation** and enter `UserName@Example.Com` as the email (make sure you match the letter case, including the mixed uppercase and lowercase characters):



Method: POST

Magazine count:

Book count:

Send confirmation:

Email:

Figure 7.14 – The form after being submitted with valid values

- Switch to PyCharm and look in the debug console. You'll see that the email is converted into lowercase when it is printed by our debug code:

```
magazine_count: (<class 'int'>) 20
book_count: (<class 'int'>) 20
send_confirmation: (<class 'bool'>) True
email: (<class 'str'>) username@example.com
```

Figure 7.15 – Email in lowercase, as well as other fields

This is our `clean_email` method in action – even though we entered data in both uppercase and lowercase, it has been converted into all lowercase.

In this exercise, we created a new `OrderForm` that implemented form and field clean methods. We used a custom validator to ensure that the **Email** field met our specific validation rules – only a specific domain was allowed. We used a custom field cleaning method (`clean_email`) to convert the email address into lowercase. We then implemented a `clean` method to validate the forms that were dependent on each other. In this method, we added both field and non-field errors.

So far, we have learned about multi-field validations, and how to implement them to make sure that we can validate the form in one go rather than highlighting errors one by one.

In the next section, we will cover how to add placeholders and initial values to the form.

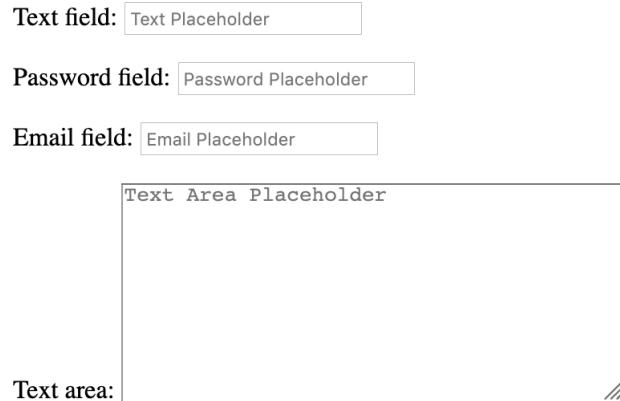
Adding placeholders and initial values

There are two things our first manually built form had that our current Django form still does not have – placeholders and initial values. Adding placeholders is simple; they are just added as attributes to the widget constructor for the form field. This is similar to what we've already seen for setting the type of `DateField` in our previous examples.

Here's an example:

```
class ExampleForm(forms.Form):
    text_field = forms.CharField(
        widget=forms.TextInput(attrs={"placeholder": "Text
Placeholder"})
    )
    password_field = forms.CharField(
        widget=forms.PasswordInput(attrs={"placeholder":
"Password Placeholder"})
    )
    email_field = forms.EmailField(
        widget=forms.EmailInput(attrs={"placeholder":
"Email Placeholder"})
    )
    text_area = forms.CharField(
        widget=forms.Textarea(attrs={"placeholder": "Text
Area Placeholder"})
    )
```

Here is what the preceding form looks like when rendered in the browser:



The image shows a screenshot of a Django form rendered in a browser. It consists of five input fields: a text input, a password input, an email input, a text area, and a text input. Each field has a placeholder text inside it. The text input fields have the placeholder 'Text Placeholder'. The password input field has the placeholder 'Password Placeholder'. The email input field has the placeholder 'Email Placeholder'. The text area has the placeholder 'Text Area Placeholder'.

Text field: Text Placeholder

Password field: Password Placeholder

Email field: Email Placeholder

Text area: Text Area Placeholder

Figure 7.16 – Django form with placeholders

Of course, if we're manually setting `Widget` for each field, we need to know which `Widget` class to use. The ones that support placeholders are `TextInput`, `NumberInput`, `EmailInput`, `URLInput`, `PasswordInput`, and `Textarea`.

While we're examining the `Form` class itself, we'll look at the first of two ways of setting an initial value for a field. We can do this by using the `initial` argument on a `Field` constructor, like this:

```
text_field = forms.CharField(initial="Initial Value", ...)
```

The other method is to pass in a dictionary of data when instantiating the form in our view. The keys are the field names. The dictionary should have zero or more items (that is, an empty dictionary is valid). Any extra keys are ignored. This dictionary should be supplied as the `initial` argument in our view, as follows:

```
initial = {"text_field": "Text Value", "email_field": "user@example.com"}  
form = ExampleForm(initial=initial)
```

Or, for a POST request, pass in `request.POST` as the first argument, as usual:

```
initial = {"text_field": "Text Value", "email_field": "user@example.com"}  
form = ExampleForm(request.POST, initial=initial)
```

Values in `request.POST` will override values in `initial`. This means that even if we have an initial value for a required field, if it's left blank when submitted, then it will not validate. The field will not fall back to the value in `initial`.

Whether you decide to set initial values in the `Form` class itself or the view is up to you and depends on your use case. If you had a form that was used in multiple views but usually had the same value, it would be better to set the `initial` value in the form. Otherwise, it can be more flexible to use `setting` in the view.

In the next exercise, we will add placeholders and initial values to the `OrderForm` class from the previous exercise.

Exercise 7.02 – placeholders and initial values

In this exercise, you will enhance the `OrderForm` class by adding placeholder text. You will simulate passing an initial email address to the form. It will be a hardcoded address but once the user can log in, it could be an email address associated with their account – you will learn about sessions and authentication in *Chapter 9, Sessions and Authentication*:

1. In PyCharm, open the `reviews` app's `forms.py` file. You will add placeholders to the `magazine_count`, `book_count`, and `email` fields on the `OrderForm`, which means also setting `widget`.

To the `magazine_count` field, add a `NumberInput` widget with `placeholder` in the `attrs` dictionary. This `placeholder` should be set to **Number of Magazines**. Write the following code:

```
magazine_count = forms.IntegerField(  
    min_value=0,  
    max_value=80,  
    widget=forms.NumberInput(attrs={"placeholder":  
        "Number of Magazines"}),  
)
```

2. Add a placeholder to the `book_count` field in the same manner. The placeholder text should be **Number of Books**:

```
book_count = forms.IntegerField(  
    min_value=0,  
    max_value=50,  
    widget=forms.NumberInput(attrs={"placeholder":  
        "Number of Books"}),  
)
```

3. The final change to `OrderForm` is to add a placeholder to the email field. This time, the widget is `EmailInput`. The placeholder text should be **Your company email address**:

```
email = forms.EmailField(  
    required=False,  
    validators=[validate_email_domain],  
    widget=forms.EmailInput(attrs={"placeholder":  
        "Your company email address"}),  
)
```

Note that the `clean_email` and `clean` methods should remain as they were in *Exercise 7.01 – custom clean and validation methods*. Save the file.

4. Open the `form_project` app's `views.py` file. In the `form_example` view function, create a new dictionary variable called `initial` with one key, `email`, like this:

```
initial = {"email": "user@example.com"}
```

5. In the two places that you are instantiating `OrderForm`, also pass in the `initial` variable using the `initial` kwarg. The first instance is as follows:

```
form = OrderForm(request.POST, initial=initial)
```

The second instance is as follows:

```
form = OrderForm(initial=initial)
```

The complete code for `views.py` can be found at https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter07/Exercise7.02/form_project/form_example/views.py.

Save the `views.py` file.

6. Start the Django dev server if it is not already running. Browse to `http://127.0.0.1:8000/form-example/` in your browser. You should see that your form now has placeholders and an initial value set:



The screenshot shows a web browser window with a title bar 'Title' and a tab '127.0.0.1:8000/form-example/'. The content area displays an 'Order' form. The 'Method: GET' is indicated. The form fields are:

- Magazine count:
- Book count:
- Send confirmation:
- Email:
-
-

Figure 7.17 – Order form with initial values and placeholders

In this exercise, we added placeholders to form fields. This was done by setting a `form` widget when defining the `form` field on the form class and setting a `placeholder` value in the `attrs` dictionary. We also set an initial value for the form using a dictionary and passed it to the `form` instance using the `initial` kwarg.

In the next section, we will talk about how to work with Django models using data from forms, and how `ModelForm` makes this easier.

Creating or editing Django models

We've seen how to define a form, and in *Chapter 2, Models and Migrations*, you learned how to create Django model instances. By using these things together, you can build a view that displays a form and also saves a model instance to the database. This gives you an easy way to save data without having to write a lot of boilerplate code or create custom forms. In Bookr, we will use this method to allow users to add reviews without requiring access to the Django admin site. Without using `ModelForm`, you can do something like this:

- You can create a form based on an existing model; for example, `Publisher`. The form would be called `PublisherForm`.

- You can manually define the fields on `PublisherForm`, using the same rules defined on the `Publisher` model, as shown here:

```
class PublisherForm(forms.Form):  
    name = forms.CharField(max_length=50)  
    website = forms.URLField()  
    ...
```

- In the view, the `initial` values would be retrieved from the model queried from the database, then passed to the form using the `initial` argument. If you were creating a new instance, the `initial` value would be blank – something like this:

```
if create:  
    initial = {}  
else:  
    publisher = Publisher.objects.get(pk=pk)  
    initial = { "name": publisher.name, "website":  
                publisher.website, ...}  
  
form = PublisherForm(initial=initial)
```

- Then, in the POST flow of the view, you can either create or update the model based on `cleaned_data`:

```
form = PublisherForm(request.POST, initial=initial)  
if create:  
    publisher = Publisher()  
else:  
    publisher = Publisher.objects.get(pk=pk)  
  
publisher.name = form.cleaned_data['name']  
publisher.website = forms.cleaned_data['website']  
...  
publisher.save()
```

This is a lot of work, and we have to consider how much duplicated logic we have. For example, we're defining the length of the name in the name form field. If we made a mistake here, we could allow a longer name in the field than the model allows. We also have to remember to set all the fields in the `initial` dictionary, as well as set the values on the new or updated model with `cleaned_data` from the form. There are many opportunities to make mistakes here, as well as remembering to add or remove field setting data for each of these steps if the model changes. All this code would have to be duplicated for each Django model you work with as well, expounding the duplication problem.

The ModelForm class

Luckily, Django provides a method of building `Model` instances from forms much more simply, with the `ModelForm` class. A `ModelForm` is a form that is built automatically from a particular model. It will inherit the validation rules from the model (such as whether fields are required or the maximum length of `CharField` instances, and so on). It provides an extra `__init__` argument (called `instance`) for automatically populating the initial values from an existing model. It also adds a `save` method to automatically persist the form data to the database. All you need to do to set up `ModelForm` is specify its model and what fields should be used: this is done on the `class Meta` attribute of the `form` class. Let's see how to build a form from `Publisher`.

Inside the file that contains the form (for example, the `forms.py` file we have been working with), the only change is that the model must be imported:

```
from .models import Publisher
```

Then, the `Form` class can be defined. The class requires a `class Meta` attribute, which, in turn, must define a `model` attribute and either the `fields` or `excludes` attributes:

```
class PublisherForm(forms.ModelForm):
    class Meta:
        model = Publisher
        fields = ("name", "website", "email")
```

`fields` is a list or tuple of the fields to include in the form. When manually setting the list of fields, if you add extra fields to the model, you must also add their name here to have them displayed on the form.

You can also use the special `__all__` value instead of a list or tuple to automatically include all the fields, like this:

```
class PublisherForm(forms.ModelForm):
    class Meta:
        model = Publisher
        fields = "__all__"
```

If the `model` field has its `editable` attribute set to `False`, then it will not be automatically included.

On the contrary, the `exclude` attribute sets the fields to not display in the form. Any fields added to the model will automatically be added to the form. We could define the preceding form using `exclude` with any empty tuple since we want all the fields. The code is like this:

```
class PublisherForm(forms.ModelForm):  
    class Meta:  
        model = Publisher  
        exclude = ()
```

This saves you some work because you don't need to add a field to both the model and the `fields` list; however, it is not as safe since you might automatically expose fields to the end user that you don't want to. For example, if you had a `User` model with `UserForm`, you might add an `is_admin` field to the `User` model to give admin users extra privileges. If this field was not in `exclude`, it would be displayed to the user. A user would then be able to make themselves an administrator, which is something you probably wouldn't want.

Whichever of these three approaches to choosing the forms to display that we decide to use, in our case, they will display the same in the browser. This is because we are choosing to display *all* the fields. They all look like this when rendered in the browser:

Name: The name of the Publisher.
Website: The Publisher's website.
Email: The Publisher's email address.

Figure 7.18 – PublisherForm

Note that `help_text` from the `Publisher` model is automatically rendered as well.

Usage in a view is similar to the other forms we have seen. Also, as mentioned, there is an extra argument that can be provided, called `instance`. This can be set to `None`, which will render an empty form.

Assuming, in your view function, you have some method of determining whether you are creating or editing a model instance (we will discuss how to do this later), this will determine a variable called `is_create` (True if creating an instance, or False if editing an existing one). Your view function to create the form could then be written like this:

```
if is_create:  
    instance = None  
else:  
    instance = get_object_or_404(Publisher, pk=pk)  
  
if request.method == "POST":  
    form = PublisherForm(request.POST, instance=instance)  
    if form.is_valid():
```

```

        # we'll cover this branch soon
    else:
        form = PublisherForm(instance=instance)

```

As you can see, in either branch, the instance is passed to the PublisherForm constructor, although it is `None` if we're in create mode.

If the form is valid, we can save the `model` instance. This can be done simply by calling the `save` method on the form. This will automatically create the instance, or simply save changes to the old one:

```

if form.is_valid():
    form.save()
    return redirect(success_url)

```

The `save` method returns the `model` instance that was saved. It takes one optional argument, `commit`, which determines whether the changes should be written to the database. You can pass `False` instead, which allows you to make additional changes to the instance before manually saving the changes. This would be required to set attributes that have not been included in the form. As we mentioned, maybe you would set the `is_admin` flag to `False` on a `User` instance:

```

if form.is_valid():
    new_user = form.save(False)
    new_user.is_admin = False
    new_user.save()
    return redirect(success_url)

```

In *Activity 7.02 – Review Creation UI*, at the end of this chapter, we will be using this feature as well.

If your model uses `ManyToMany` fields, and you also call `form.save(False)`, you should also call `form.save_m2m()` to save any many-to-many relationships that have been set. It's not necessary to call this method if you call the `form.save` method with `commit` set to `True` (that is, the default).

Model forms can be customized by making changes to their `Meta` attributes. The `widgets` attribute can be set. It can contain a dictionary keyed on the field names, with widget classes or instances as the values. For example, this is how to set up `PublisherForm` to have placeholders:

```

class PublisherForm(forms.ModelForm):
    class Meta:
        model = Publisher
        fields = "__all__"
        widgets = {
            "name": forms.TextInput(attrs={
                "placeholder": "The publisher's name."
            })
        }

```

The values behave the same as setting the `kwarg` `widget` in the field definition; they can be a class or an instance. For example, to display `CharField` as a password input, the `PasswordInput` class can be used; it does not need to be instantiated:

```
widgets = {
    "password": forms.PasswordInput
}
```

Model forms can also be augmented with extra fields added in the same way as they are added to a normal form. For example, suppose we wanted to provide the option of sending a notification email after saving a `Publisher`. We can add an `email_on_save` field to `PublisherForm` like this:

```
class PublisherForm(forms.ModelForm):
    email_on_save = forms.BooleanField(required=False,
        help_text="Send notification email on save")

    class Meta:
        model = Publisher
        fields = "__all__"
```

When rendered, the form looks like this:

Name: The name of the Publisher.

Website: The Publisher's website.

Email: The Publisher's email address.

Email on save: Send notification email on save

Figure 7.19 – PublisherForm with an additional field

Additional fields are placed after the `Model` fields. The extra fields are not handled automatically – they do not exist on the model, so Django won't attempt to save them on the `model` instance. Instead, you should handle how their values are saved by examining the `cleaned_data` values of the form, as you would with a standard form, for example (inside your view function):

```
if form.is_valid():
    if form.cleaned_data.get("email_on_save"):
        send_email()  # assume this function is defined
        elsewhere
    form.save()  # save the instance regardless of sending
    the email or not
    return redirect(success_url)
```

In the following exercise, you will write a new view function in Bookr that will allow you to create and edit a Publisher.

Exercise 7.03 – creating and editing a publisher

In this exercise, we will return to Bookr. We want to add the ability to create and edit a Publisher without using the Django admin. To do this, we'll add a `ModelForm` for the `Publisher` model. It will be used in a new view function. This view function will take an optional argument, `pk`, which will either be the ID of the `Publisher` object being edited or `None` to create a new Publisher. We will add two new URL maps to facilitate this. When this is complete, we will be able to see and update any publisher using their ID. For example, information for *Publisher 1* will be viewable/editable at the `/publishers/1` URL path:

1. In PyCharm, open the `reviews` app's `forms.py` file. After the `forms` import, import the `Publisher` model:

```
from .models import Publisher
```

2. Create a `PublisherForm` class, inheriting from `forms.ModelForm`:

```
class PublisherForm(forms.ModelForm):
```

3. Define the `class Meta` attribute on `PublisherForm`. The attributes that `Meta` requires are the `model` (`Publisher`) and `fields` ("`__all__`"):

```
class PublisherForm(forms.ModelForm):
    class Meta:
        model = Publisher
        fields = "__all__"
```

Save `forms.py`.

Note

The full file can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter07/Exercise7.03/bookr/reviews/forms.py>.

4. Open the reviews app's `views.py` file. At the top of the file, import `PublisherForm`:

```
from .forms import PublisherForm, SearchForm
```

5. Make sure you import the `get_object_or_404` and `redirect` functions from `django.shortcuts`, if you aren't already:

```
from django.shortcuts import render, get_object_or_404, redirect
```

6. Also, make sure you're importing the `Publisher` model if you aren't already. You may already be importing this and other models:

```
from .models import Book, Contributor, Publisher
```

7. The final import you will need is the `messages` module. This will allow us to register a message letting the user know that a `Publisher` object was edited or created:

```
from django.contrib import messages
```

Once again, add this import if you do not already have it.

8. Create a new view function called `publisher_edit`. It takes two arguments: `request` and `pk` (the ID of the `Publisher` object to edit). This is optional, and if it is `None`, then a `Publisher` object will be created instead:

```
def publisher_edit(request, pk=None):
```

9. Inside the view function, we need to try to load the existing `Publisher` instance if `pk` is not `None`. Otherwise, `publisher` should be `None`:

```
def publisher_edit(request, pk=None):  
    if pk is not None:  
        publisher = get_object_or_404(Publisher,  
                                      pk=pk)  
    else:  
        publisher = None
```

10. After getting a `Publisher` instance or `None`, complete the branch for a `POST` request. Instantiate the form in the same way as you saw earlier in this chapter, but now, make sure that it takes `instance` as a `kwarg`. Then, if the form is valid, save it using the `form.save()` method. This method will return the updated `Publisher` instance, which is stored in the `updated_publisher` variable. Then, register a different success message depending on whether the `Publisher` was created or updated. Finally, redirect back to this `publisher_edit` view, since `updated_publisher` will always have an ID at this point:

```
def publisher_edit(request, pk=None):  
    ...  
    if request.method == "POST":  
        form = PublisherForm(request.POST,  
                             instance=publisher)  
        if form.is_valid():  
            updated_publisher = form.save()  
            if publisher is None:  
                messages.success(request, "Publisher  
                "{}" was created.  
                .format(updated_publisher))  
            else:  
                messages.success(request, "Publisher  
                "{}" was updated.  
                .format(updated_publisher))  
  
            return redirect("publisher_edit",  
                           updated_publisher.pk)
```

If the form is not valid, the execution will fall through and just return the `render` function call with the invalid form (this will be implemented in *step 12*). The redirect uses a named URL map, which will be added later in this exercise.

11. Next, fill in the non-`POST` branch of the code. In this case, just instantiate the form with your `instance`:

```
def publisher_edit(request, pk=None):  
    ...  
    if request.method == "POST":  
        ...  
    else:  
        form = PublisherForm(instance=publisher)
```

12. In *step 14*, you'll create a `form-example.html` file to render, but we can add the call to the `render` function now. Render it in the view with the `render` function by passing in the HTTP method and `form` as the context:

```
def publisher_edit(request, pk=None):  
    ...  
    return render(request, "form-example.html",  
    {"method": request.method, "form": form})
```

Save this file. You can refer to this at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter07/Exercise7.03/bookr/reviews/views.py>.

13. Open `urls.py` in the `reviews` directory. Add two new URL maps; they will both go to the `publisher_edit` view. One will capture the ID of Publisher, which we will want to edit and pass into the view as the `pk` argument. The other will use `new` instead, and will not pass the `pk`, which will indicate we want to create a new Publisher.

To your `urlpatterns` variable, add the `"publishers/<int:pk>/"` path mapping to the `reviews.views.publisher_edit` view with the name `"publisher_edit"`.

Also, add the `"publishers/new/"` path mapping to the `reviews.views.publisher_edit` view with the name `"publisher_create"`:

```
urlpatterns = [  
    ...  
    path("publishers/<int:pk>/", views.publisher_edit,  
    name="publisher_edit"),  
    path("publishers/new/", views.publisher_edit,  
    name="publisher_create")  
]
```

Since the second mapping does not capture anything, whichever `pk` that is passed to the `publisher_detail` view function is `None`.

Save the `urls.py` file. The completed version for reference is at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter07/Exercise7.03/bookr/reviews/urls.py>.

14. Create a `form-example.html` file inside the `reviews` app's `templates` directory by using the **New HTML File** function in PyCharm. Since this is a standalone template (it does not extend any other templates), we need to render the messages inside it. Add this code just after the opening `<body>` tag to iterate through all the messages and display them:

```
{% for message in messages %}  
<p><em>{{ message.level_tag|title }} :{{ message }}</em>  
</p>  
{% endfor %}
```

This will loop over the messages we have added and display the tag (in our case, `Success`) and then the message.

15. Then, add the normal form rendering and submission code:

```
<form method="post">  
    {% csrf_token %}  
    {{ form.as_p }}  
    <p>  
        <input type="submit" value="Submit">  
    </p>  
</form>
```

Save and close this file.

You can refer to the full version of this file at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter07/Exercise7.03/bookr/reviews/templates/form-example.html>.

16. Start the Django dev server, then navigate to `http://127.0.0.1:8000/publishers/new/`. You should see a blank `PublisherForm` being displayed:

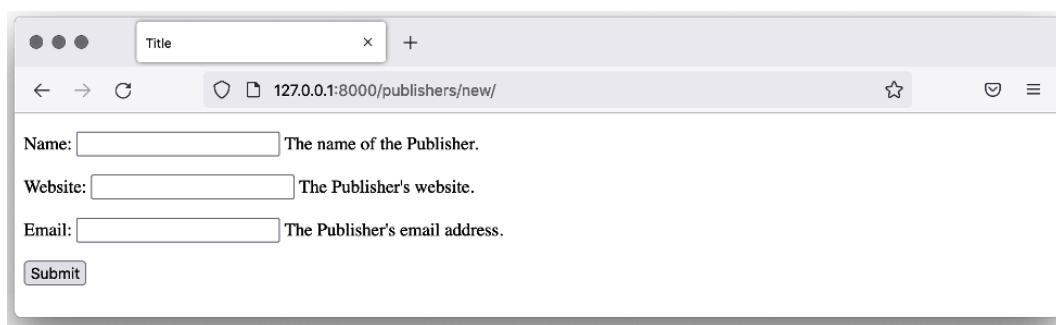


Figure 7.20 – Blank publisher form

17. The form has inherited the model's validation rules, so you cannot submit the form with too many characters for Name or with an invalid Website or Email. Put in some valid information, then submit the form. After submission, you should see the success message; the form will be populated with information that was saved to the database:



Success: Publisher "Test Publisher" was created.

Name: The name of the Publisher.

Website: The Publisher's website.

Email: The Publisher's email address.

Figure 7.21 – Form after submission

Notice that the URL has also been updated and now includes the ID of the publisher that was created. In this case, it is `http://127.0.0.1:8000/publishers/10/`; the ID in your setup will depend on how many Publisher instances were already in your database.

Notice that if you refresh the page, you will not receive a message confirming whether you want to re-send the form data. This is because we redirected after saving, so it is safe to refresh this page as many times as you want and no new Publisher instances will be created. If you had not redirected it, then every time the page was refreshed, a new Publisher instance would be created.

If you have other Publisher instances in your database, you can change the ID in the URL to edit other ones. Since the ID in this instance is 3, we can assume that Publisher 1 and Publisher 2 already exist and can substitute their IDs to see the existing data. Here is the view of the existing Publisher 1 information (at `http://127.0.0.1:8000/publishers/1/`) – your information may be different:

Name: The name of the Publisher.

Website: The Publisher's website.

Email: The Publisher's email address.

Figure 7.22 – Existing Publisher 1 information

Try making changes to the existing **Publisher** instance. Notice that after you save, the message is different – it is telling the user that the **Publisher** instance was *updated* rather than *created*:

Success: Publisher "Packt Publishing" was updated.

Name: The name of the Publisher.

Website: The Publisher's website.

Email: The Publisher's email address.

Figure 7.23 – Publisher after updating instead of creating

In this exercise, we implemented a `ModelForm` from a model (`PublisherForm` was created from `Publisher`) and saw how Django automatically generated the form fields with the correct validation rules. We then used the form's built-in `save` method to save changes to the `Publisher` instance (or automatically create it) inside the `publisher_edit` view. We mapped two URLs to the view. The first was for editing an existing `Publisher` and passing a `pk` to the view. The other did not pass this `pk` to the view, indicating that the `Publisher` instance should be created. Finally, we used the browser to experiment with creating a new `Publisher` instance and then editing an existing one.

Activity 7.01 – styling and integrating the publisher form

In *Exercise 7.03 – creating and editing a publisher*, you added `PublisherForm` to create and edit `Publisher` instances. You built this with a standalone template that did not extend any other templates, so it lacked the global styles. In this activity, you will build a generic form detail page that will display a Django form, similar to `form-example.html` but extending from a base template.

The template will accept a variable to display the type of model being edited. You will also update the main base.html template to render the Django messages, using Bootstrap styling.

These steps will help you complete this activity:

1. Start by editing the project's base.html file. Wrap the content block in a div container for a nicer layout with some spacing. Surround the existing content block with an `<div>` element with `class="container-fluid"`.
2. Render each message in messages (similar to step 14 of *Exercise 7.03 – creating and editing a publisher*). Add the `{% for %}` block after the `<div>` container you just created but before the content block. You should use the Bootstrap framework classes – this snippet will help you:

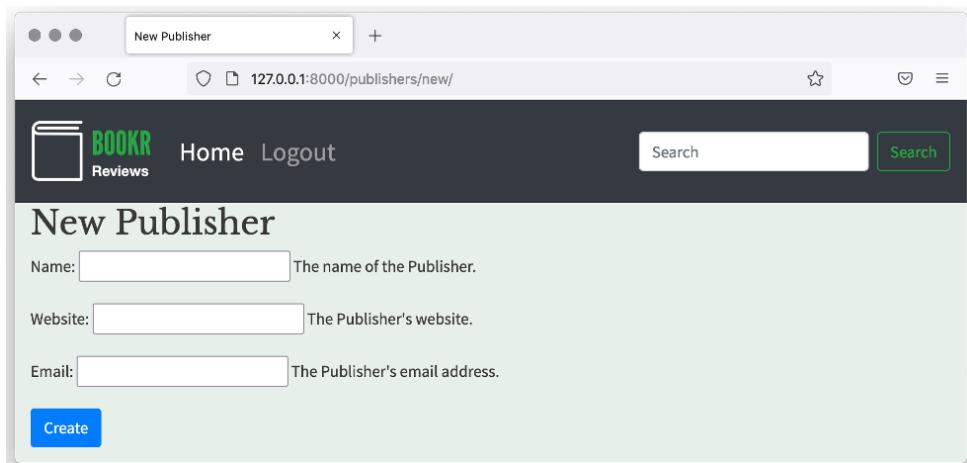
```
<div class="alert alert-{{ if message.level_tag == 'error' %}}  
danger{{% else %}}{{ message.level_tag }}{{% endif %}}"  
role="alert">  
{{ message }}  
</div>
```

The Bootstrap class and Django message tags have corresponding names for the most part (for example, success and alert-success). The exception is Django's error tag. The corresponding Bootstrap class is alert-danger. See more information about Bootstrap alerts at <https://getbootstrap.com/docs/4.0/components/alerts/>. This is why you need to use the if template tag in this snippet.

3. Create a new template called `instance-form.html`, inside the reviews app's namespaced `templates` directory.
4. `instance-form.html` should extend from the reviews app's `base.html`.
5. The context that's being passed to this template will contain a variable called `instance`. This will be the `Publisher` instance being edited, or `None` if we are creating a new `Publisher` instance. The context will also contain a `model_type` variable, which is a string indicating the model type (in this case, `Publisher`). Use these two variables to populate the `title` block template tag:
 - If the `instance` is `None`, the title should be *New Publisher*
 - Otherwise, the title should be *Editing Publisher <Publisher Name>*
6. `instance-form.html` should contain a `content` block template tag, to override the `base.html` `content` block.
7. Add an `<h2>` element inside the `content` block and populate it using the same logic as the title. For better styling, wrap the publisher name in an `` element.
8. Add an `<form>` element to the template with a `method` of `post`. Since we are posting back to the same URL, an `action` does not need to be specified.
9. Include the CSRF token template tag in the `<form>` body.

10. Render the Django form (its context variable will be `form`) inside `<form>`, using the `as_p` method.
11. Add a `submit` `<button>` to the form. Its text should depend on whether you're editing or creating. Use the `Save` text for editing or `Create` for creating. You can use the Bootstrap classes for the button styling here. It should have the `class="btn btn-primary"` attribute.
12. In `reviews/views.py`, the `publisher_edit` view does not need many changes. Update the `render` call to render `instance-form.html` instead of `form-example.html`.
13. Update the context dictionary being passed to the `render` call. It should include the Publisher instance (the `publisher` variable that was already defined) and the `model_type` string. The context dictionary already includes `form` (a `PublisherForm` instance). You can remove the `method` key.
14. Since we're finished with the `form-example.html` template, it can be deleted.

When you've finished, the Publisher creation page (at `http://127.0.0.1:8000/publishers/new/`) should look like what's shown in *Figure 7.24*:



The screenshot shows a web browser window with the title "New Publisher". The address bar displays "127.0.0.1:8000/publishers/new/". The page content is titled "New Publisher". It contains three input fields: "Name" (placeholder: "The name of the Publisher."), "Website" (placeholder: "The Publisher's website."), and "Email" (placeholder: "The Publisher's email address."). Below these fields is a blue "Create" button. The top of the page has a dark header with a logo (a book icon and the word "BOOKR"), "Home", and "Logout" links. There is also a search bar with a "Search" button.

Figure 7.24 – The Publisher creation page

When editing a Publisher (for example, at `http://127.0.0.1:8000/publishers/1/`), your page should look like what's shown in *Figure 7.25*:

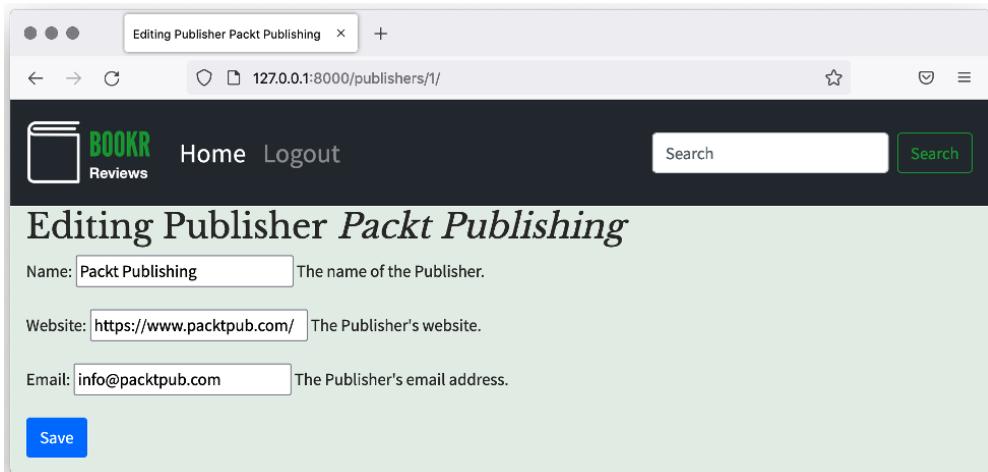


Figure 7.25 – The Editing Publisher page

After saving a Publisher instance, whether creating or editing, you should see a success message at the top of the page (*Figure 7.26*):

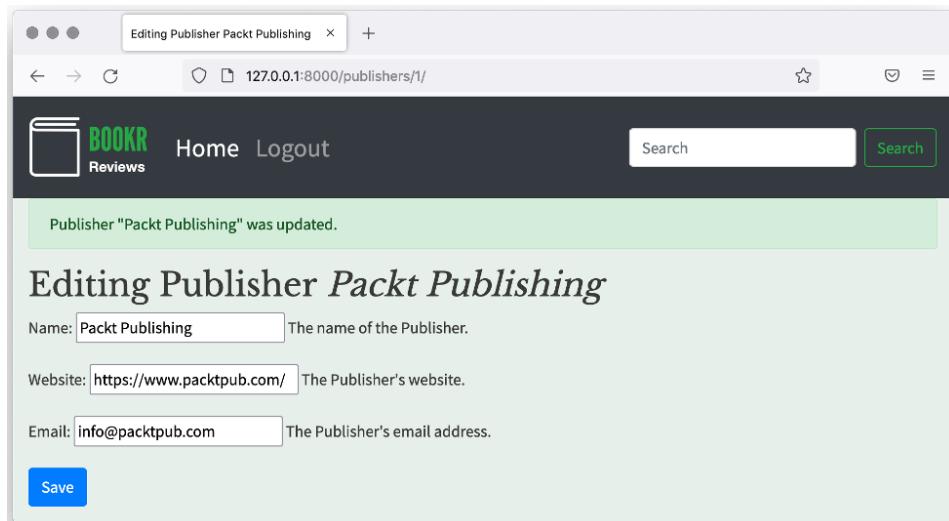


Figure 7.26 – Success message rendered as a Bootstrap alert

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Activity 7.02 – Review Creation UI

Activity 7.01 – styling and integrating the Publisher form, was quite extensive; however, by completing it, you have created a foundation that makes it easier to add other *edit* and *create* views. You will experience this first-hand in this activity when you will build forms for creating and editing reviews. Because the `instance-form.html` template was made generically, you can reuse it in other views.

In this activity, you will create a review `ModelForm`, then add a `review_edit` view to create or edit a `Review` instance. You can reuse `instance-form.html` from *Activity 7.01 – styling and integrating the Publisher form*, and pass in different context variables to make it work with the `Review` model. When working with reviews, you'll operate within the context of a book – that is, the `review_edit` view must accept a Book PK as an argument. You'll fetch this Book separately and assign it to the `Review` instance that you create.

These steps will help you complete this activity:

1. In `forms.py`, add a `ReviewForm` subclass of `ModelForm`; its model should be `Review` (make sure you import the `Review` model).

`ReviewForm` should exclude the `date_edited` and `book` fields since the user should not be setting these in the form. The database allows any rating, but we can override the `rating` field with an `IntegerField` that requires a minimum value of `0` and a maximum value of `5`.

Create a new view called `review_edit`. It should accept two arguments after `request`: `book_pk`, which is required, and `review_pk`, which is optional (defaults to `None`). Fetch the `Book` instance and `Review` instance using the `get_object_or_404` shortcut (call it once for each type). When fetching the review, make sure the review belongs to the book. If `review_pk` is `None`, then the `Review` instance should be `None` too.

2. If the `request` method is `POST`, then instantiate a `ReviewForm` using `request.POST` and the `review` instance. Make sure you import this `ReviewForm`.

If the form is valid, save the form but set the `commit` argument from `save` to `False`. Then, set the `book` attribute on the returned `Review` instance to the book that was fetched in *step 2*.

If the `Review` instance is being updated instead of created, then you should also set the `date_edited` attribute to the current date and time. Use the `from django.utils.timezone.now()` function. Then, save the `Review` instance.

3. Finish the valid form branch by registering a success message and redirecting back to the `book_detail` view. Since the `Review` model doesn't contain a meaningful text description, use the `Book` title in the message – for example, `Review for "<book title>" created.`
4. If the `request` method is not `POST`, instantiate a `ReviewForm` and just pass in the `Review` instance.
5. Render the `instance-form.html` template. In the context dictionary, include the same items as were used in `publisher_view: form, instance, and model_type (Review)`. Include two extra items – `related_model_type`, which should be `Book`, and `related_instance`, which will be the `Book` instance.
6. Edit `instance-form.html` to add a place to display the related instance information added in *step 6*. Under the `<h2>` element, add an `<p>` element that is only displayed if both `related_model_type` and `related_instance` are set. It should show the text `For <related_model_type> <related_instance>` – for example, `For Book Advanced Deep Learning with Keras`. Put the `related_instance` output in an `` element for better readability.
7. In the `reviews` app's `urls.py`, add URL maps to the `review_edit` view. The `/books/` and `/books/<pk>/` URLs have already been configured. Add `/books/<book_pk>/reviews/new/` to create a review and `/books/<book_pk>/reviews/<review_pk>/` to edit a review. Make sure you give these names such as `review_create` and `review_edit`.
8. Inside the `book_detail.html` template, add links that a user can click on to create or edit a review. Add a link inside the content block, just before the `endblock` closing template tag. It should use the `url` template tag to link to the `review_edit` view when in creation mode. Also, use the `class="btn btn-primary"` attribute to make the link display like a Bootstrap button. The link text should be `Add Review`.
9. Finally, add a link to edit a review, inside the `for` loop that iterates over `Reviews` for `Book`. After all the instances of `text-info `, add a link to the `review_edit` view using the `url` template tag. You will need to provide `book.pk` and `review.pk` as arguments. The text of the link should be `Edit Review`.

When you're finished, the **Review Comments** page should look like what's shown in *Figure 7.27*:

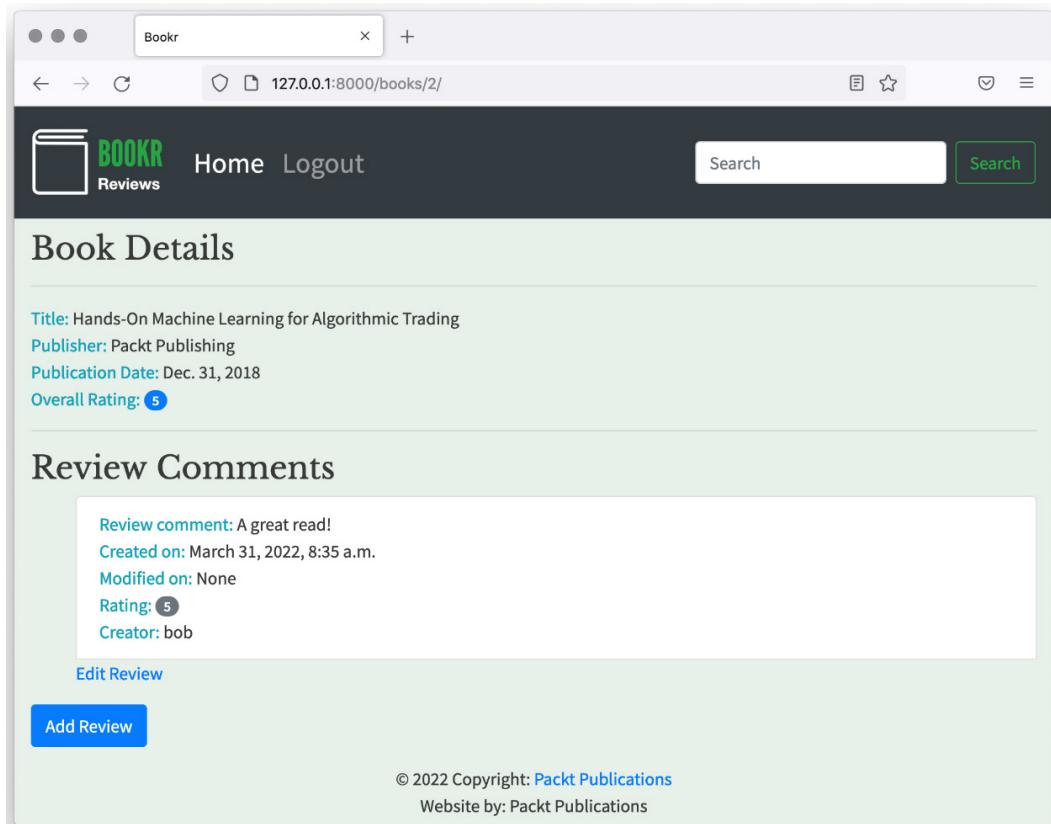
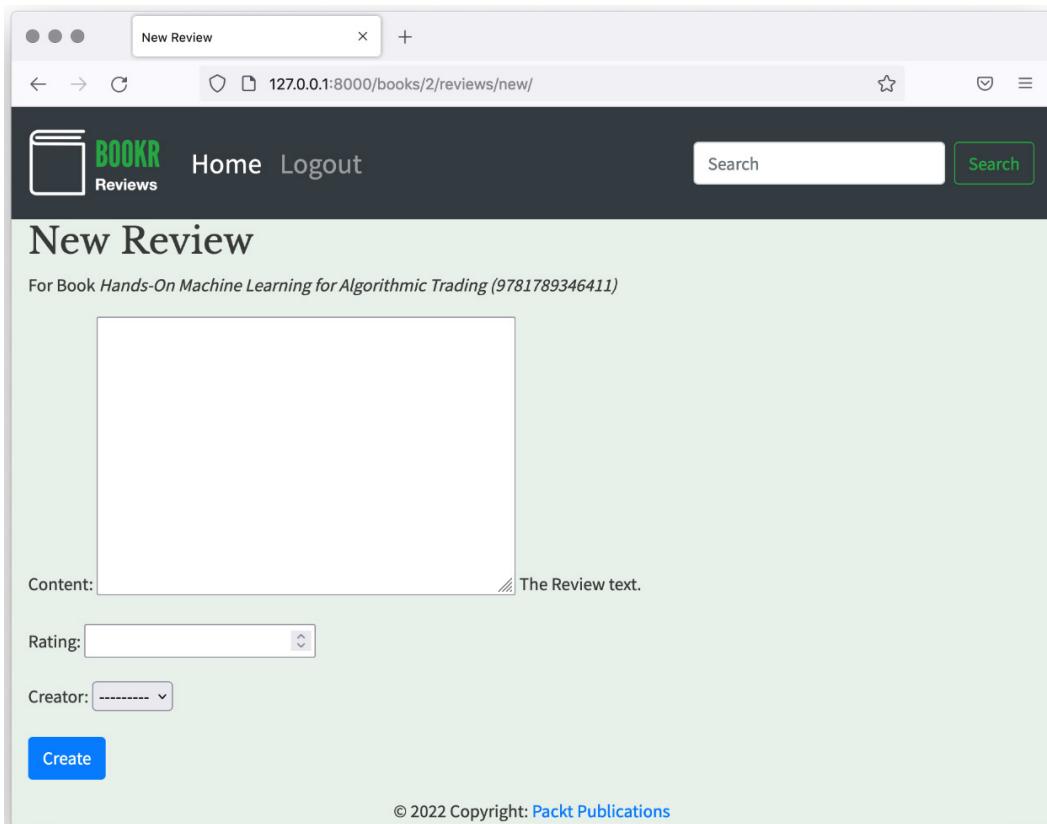


Figure 7.27 – The Book Details page with the added Add Review button

You will see the **Add Review** button. Clicking it will take you to the **Create Book Review** page, which should look like what's shown in *Figure 7.28*:



The screenshot shows a web browser window with the title 'New Review'. The URL in the address bar is '127.0.0.1:8000/books/2/reviews/new/'. The page has a dark header with a logo 'BOOKR Reviews', 'Home', 'Logout', a 'Search' bar, and a 'Search' button. The main content area is titled 'New Review' and specifies 'For Book *Hands-On Machine Learning for Algorithmic Trading* (9781789346411)'. It contains a large text input field labeled 'Content:' with the placeholder 'The Review text.' Below it are dropdown menus for 'Rating:' and 'Creator:', both currently set to '-----'. A blue 'Create' button is at the bottom left. The footer contains the copyright notice '© 2022 Copyright: Packt Publications'.

Figure 7.28 – The Create Book Review page

Enter some details in the form and click **Create**. You'll be redirected to the **Book Details** page. You should see the success message and your review, as shown in *Figure 7.29*:

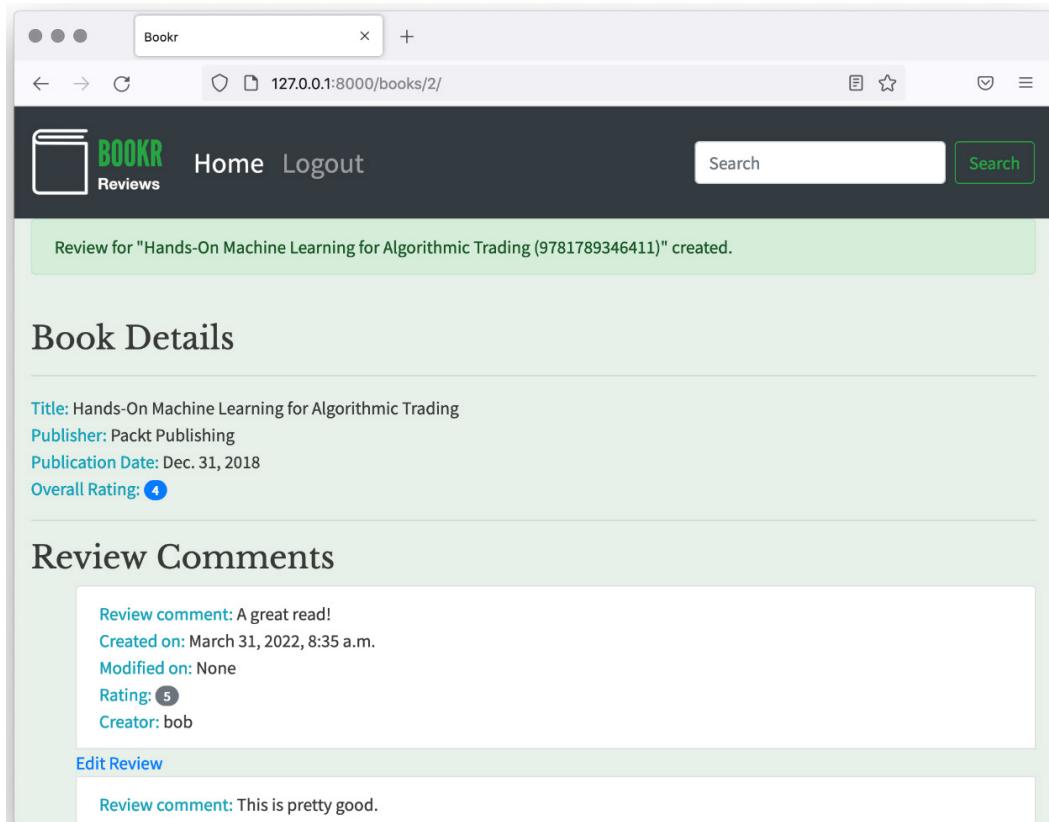
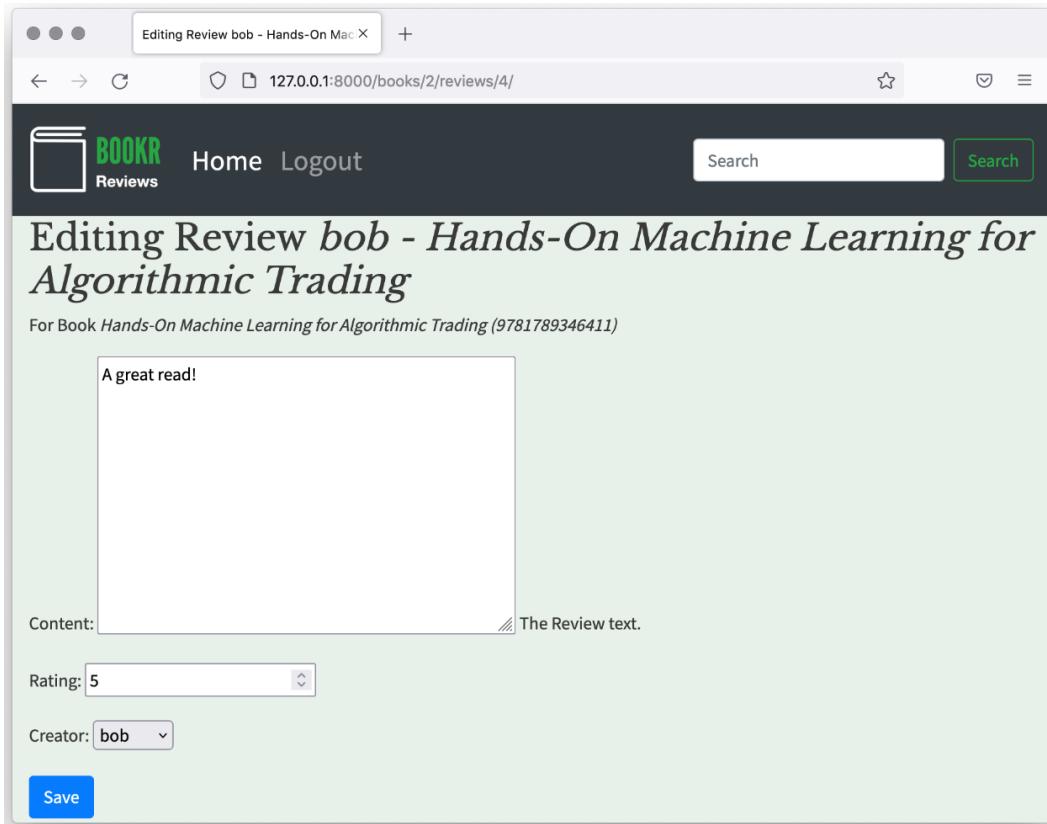


Figure 7.29 – The Book Details page with a review added

You will also see the **Edit Review** link. If you click it, you will be taken to a form that has been pre-populated with your review data (see *Figure 7.30*):



Editing Review bob - Hands-On Machine Learning for Algorithmic Trading

For Book *Hands-On Machine Learning for Algorithmic Trading* (9781789346411)

A great read!

Content: _____

The Review text.

Rating: 5

Creator: bob

Save

Figure 7.30 – Review form when editing a review

After saving an existing review, you should see that the **Modified on** date has been updated on the **Book Details** page (*Figure 7.31*):

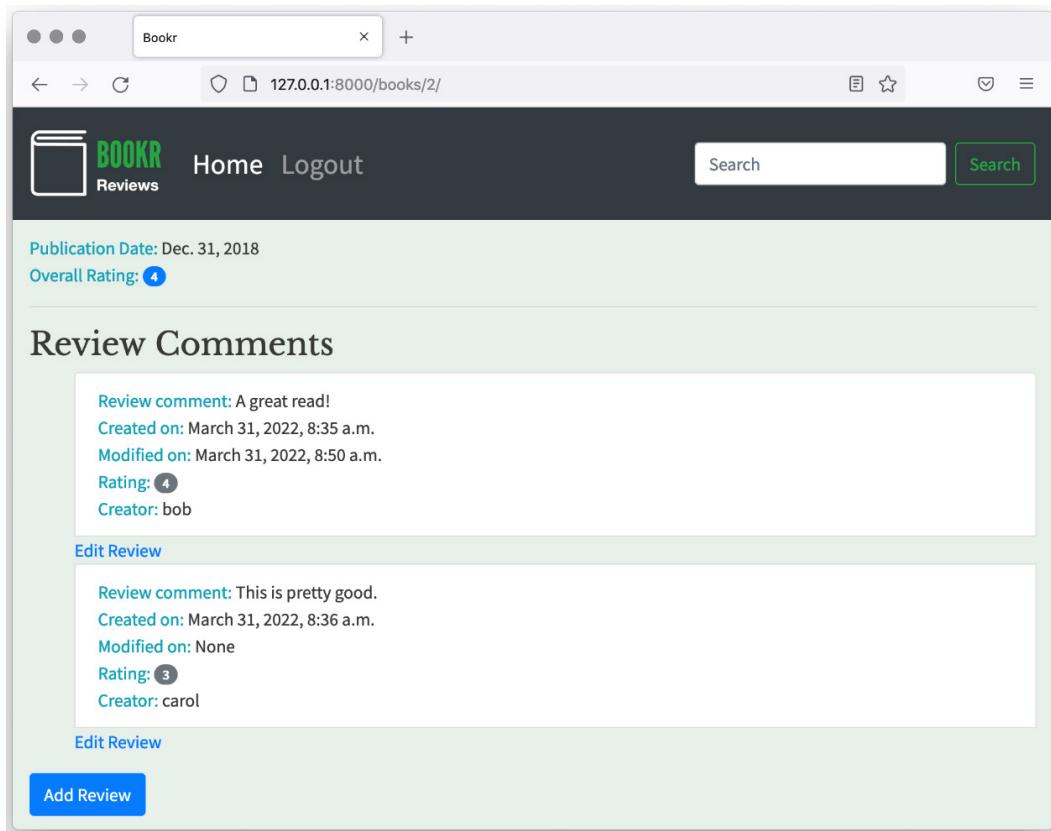


Figure 7.31 – The Modified on: date field has now been populated

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

This chapter provided a deep dive into forms. We saw how to enhance Django forms with custom validation advanced rules for cleaning data and validating fields. We also saw how custom cleaning methods can transform the data that we get out of forms. A nice feature we saw that can be added to forms is the ability to set initial and placeholder values on fields so that the user does not have to fill them out.

We then looked at how to use the `ModelForm` class to automatically create a form from a Django model. We saw how to only show some fields to the user and how to apply custom form validation rules to the `ModelForm`. We also saw how Django can automatically save the new or updated model instance to the database inside the view. In the activities for this chapter, we enhanced Bookr some more by adding forms for creating and editing publishers and submitting reviews.

In the next chapter, we will continue with the theme of submitting user input, and we'll discuss how Django handles file uploads and downloads.

8

Media Serving and File Uploads

Media files refer to extra files that can be added after deployment to enrich your Django application. Usually, they are extra images that you would use in your site, but any type of file (including video, audio, PDF, text, documents, or even HTML) can be served as media.

You can think of them as somewhere between dynamic data and static assets. They aren't dynamic data that Django generates on the fly, such as when rendering a template. They also aren't the static files that are included by the site developer when the site is deployed. Instead, they are extra files that can be uploaded by users or generated by your application for later retrieval.

Some common examples of media files (that you'll see in *Activity 8.01 – image and PDF upload of books*, later in this chapter) are book covers and preview PDFs that can be attached to a Book object. You can also use media files to allow users to upload images for a blog post or avatars for a social media site. If you wanted to use Django to build your own video-sharing platform, you would store the uploaded videos as media. Your website will not function well if all these files are static files, as users won't be able to upload their own book covers, videos, and so on and will be stuck with the ones you deployed.

In this chapter, we will be covering the following topics:

- Settings for media uploads and serving
- File uploads using HTML forms
- File uploads with Django forms
- Image uploads with Django forms
- Serving uploaded (and other) files using Django
- ModelForms and file uploads
- Activity 8.01 – image and PDF upload for books
- Activity 8.02 – displaying cover and sample link

Technical requirements

The code for the exercises and activities in this chapter can be found in this book's GitHub repository at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter08>.

Settings for media uploads and serving

In *Chapter 5, Serving Static Files*, we looked at how Django can be used to serve static files. Serving media files is quite similar. Two settings must be configured in `settings.py`: `MEDIA_ROOT` and `MEDIA_URL`. These are analogous to `STATIC_ROOT` and `STATIC_URL` for serving static files.

- `MEDIA_ROOT`: This is the path on the disk where the media (such as uploaded files) will be stored. As with static files, your web server should be configured to serve directly from this directory to take the load off Django.
- `MEDIA_URL`: This is similar to `STATIC_URL`, but as you might guess, it's the URL that should be used to serve media. It must end in `/`. Generally, you will use something such as `/media/`.

Note

For security reasons, the path for `MEDIA_ROOT` must not be the same as the path for `STATIC_ROOT`, and `MEDIA_URL` must not be the same as `STATIC_URL`. If they were the same, a user might replace your static files (such as JavaScript or CSS files) with malicious code and exploit your users.

`MEDIA_URL` is designed to be used in templates so that you don't hardcode URLs, and can change them easily. For example, you might want to set it to a specific host or **content delivery network (CDN)** when you deploy to production. We will discuss the use of `MEDIA_URL` in templates in an upcoming section.

Serving media files in development

As with static files, when serving media in production, your web server should be configured to serve directly from the `MEDIA_ROOT` directory to prevent Django from being tied up in servicing the request. The Django dev server can serve media files in development. However, unlike static files, the URL mapping and view are not set up automatically for media files.

Django provides the `static` URL mapping that can be added to your existing URL maps to serve media files. It is added to your `urls.py` file like this:

```
from django.conf import settings
from django.conf.urls.static import static

urlpatterns = [
    # your existing URL maps
]
if settings.DEBUG:
    urlpatterns += static(settings.MEDIA_URL,
        document_root=settings.MEDIA_ROOT)
```

This will serve the `MEDIA_ROOT` setting defined in `settings.py` to the `MEDIA_URL` setting that is also defined there. The reason we check for `settings.DEBUG` before appending the map, so we don't add this map in production.

For example, if your `MEDIA_ROOT` was set to `/var/www/bookr/media`, and your `MEDIA_URL` was set to `/media/`, then the `/var/www/bookr/media/image.jpg` file would be available at `http://127.0.0.1:8000/media/image.jpg`.

The `static` URL map does not work when the Django `DEBUG` setting is `False`, so it can't be used in production. However, as has been mentioned, in production, your web server should be serving these requests so Django will not need to handle them.

In the first exercise, you will create a new add `MEDIA_ROOT` and `MEDIA_URL` to your `settings.py`. You will then add the `static` media serving URL map and add a test file to ensure media serving is configured correctly.

Exercise 8.01 – configuring media storage and serving

In this exercise, you will set up a new Django project as an example project to use throughout this chapter. Then, you'll configure it to be able to serve media files. You'll do this by creating a `media` directory and adding the `MEDIA_ROOT` and `MEDIA_URL` settings. Then you'll set up the URL mapping for `MEDIA_URL`.

To check that everything is configured and being served correctly, you will put a test file inside the `media` directory:

- As with the previous example Django projects you've set up, you can reuse the existing `bookr` virtual environment. In a terminal, activate the `bookr` virtual environment. Then, start a new project named `media_project` using `django-admin.py`:

Note

To learn how to create and activate a virtual environment, refer to the *Preface* section of this book.

```
django-admin startproject media_project
```

Change (or `cd` into) into the `media_project` directory that was created, then use the `startapp` management command to start an app called `media_example`:

```
python3 manage.py startapp media_example
```

- Open the `media_project` directory in PyCharm. Set up a run configuration for the `runserver` command in the same manner as for the other Django projects you've opened:

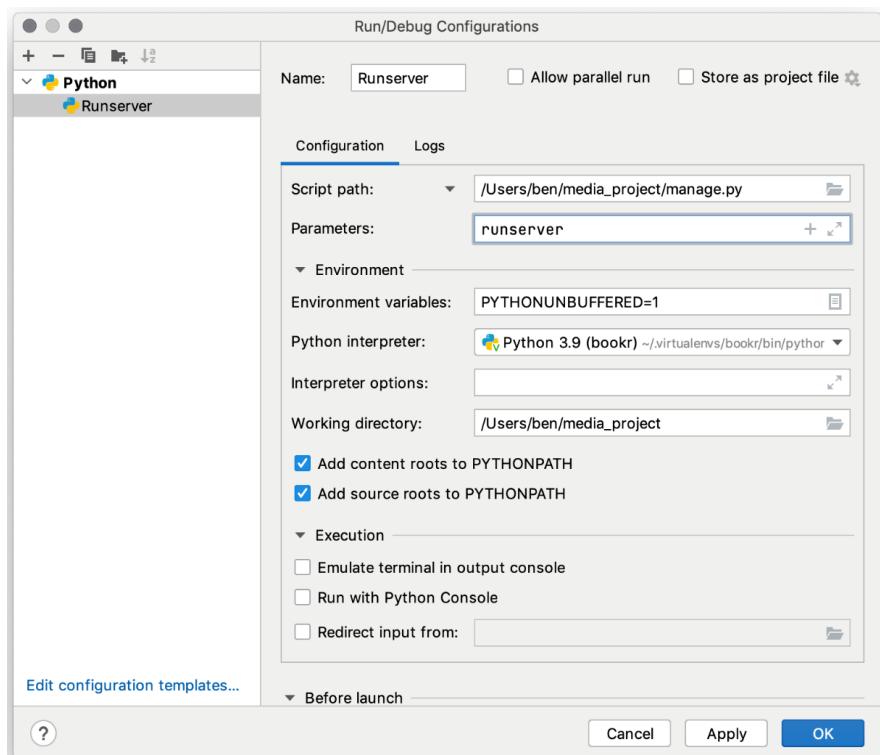


Figure 8.1 – Runserver configuration

Figure 8.1 shows the **Runserver** configuration of the project in PyCharm.

3. Create a new directory named `media` inside the `media_project` project directory. Then, create a new file in this directory named `test.txt`. The directory structure of this will look like *Figure 8.2*:

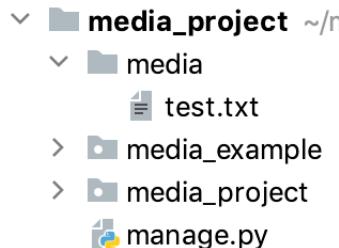


Figure 8.2: The media directory and the `test.txt` layout

4. The `test.txt` file will also open automatically. Enter `Hello, world!` into it, then you can save and close the file.
5. Open `settings.py` inside the `media_project` package directory. At the end of the file, add a setting for `MEDIA_ROOT` using the path to the media directory you just created. Join the name of the directory (`media`) onto `BASE_DIR`:

```
MEDIA_ROOT = BASE_DIR / "media"
```

6. Directly below the line added in step 5, add another setting for `MEDIA_URL`; this should just be `"/media/"`:

```
MEDIA_URL = "/media/"
```

With these changes made, `settings.py` should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.01/media_project/media_project/settings.py.

7. Open the `media_project` package's `urls.py` file. After the `urlpatterns` definition, add the following code to add the media serving URL if running in the `DEBUG` mode. First, you will need to import the Django settings and static serving view by adding the highlighted import lines above the `urlpatterns` definition:

```
from django.contrib import admin
from django.urls import path
from django.conf import settings
from django.conf.urls.static import static

urlpatterns = [
```

```
        path("admin/", admin.site.urls),  
    ]
```

8. Then, add the following code right after your `urlpatterns` definition (refer to the code block in the previous step) to conditionally add a mapping from the `MEDIA_URL` setting to the `static` view, which will serve from `MEDIA_ROOT`:

```
if settings.DEBUG:  
    urlpatterns += static(settings.MEDIA_URL,  
        document_root=settings.MEDIA_ROOT)
```

You can now save this file. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.01/media_project/media_project/urls.py.

9. Start the Django dev server, if it is not already running, then visit `http://127.0.0.1:8000/media/test.txt`. If you did everything correctly, you should see the **Hello, world!** text in your browser:



Figure 8.3: Serving a media file

If your browser looks like *Figure 8.3*, it means that the media files are being served from the `MEDIA_ROOT` directory.

The `test.txt` file we created was just for testing, but we will use it in *Exercise 8.02 – template settings, and using MEDIA_URL in templates*, so don't delete it yet.

In this exercise, we configured Django to serve media files. We served a test file just to make sure everything worked as expected, and it did. We'll now look at how we can automatically generate media URLs in templates.

Context processors and using MEDIA_URL in templates

To use MEDIA_URL in a template, we could pass it in through the rendering context dictionary in our view. For example, as shown in the following code block:

```
from django.conf import settings

def my_view(request):
    return render(request, "template.html", {"MEDIA_URL": settings.MEDIA_URL, "username": "jbloggs"})
```

This will work, but the problem is that MEDIA_URL is a common variable that we might want to use in many places, so we'd have to pass it through in practically every view.

Instead, we can use a **context processor**, which is a way of adding one or more variables automatically to the context dictionary on every `render` call.

A context processor is a function that accepts one argument, the current request. It returns a dictionary of context information that will be merged with the dictionary that was passed to the `render` call.

We can look at the source code of the `media` context processor, which illustrates how they work:

```
def media(request):
    """
    Add media-related context variables to the context.
    """
    return {"MEDIA_URL": settings.MEDIA_URL}
```

With the `media` context processor activated, `MEDIA_URL` will be added to your context dictionaries. We could change our `render` call, seen previously, to this:

```
return render(request, "template.html", {"username": "jbloggs"})
```

The same data would be sent to the template, as the context processor would add `MEDIA_URL`.

The full module path to the `media` context processor is `django.template.context_processors.media`. Some examples of other context processors that Django provides are the following:

- `django.template.context_processors.debug`: This returns the `{"DEBUG": settings.DEBUG}` dictionary.
- `django.template.context_processors.request`: This returns the `{"request": request}` dictionary; that is, it just adds the current HTTP request to the context.

To enable a context processor, its module path must be added to the `context_processors` option of your `TEMPLATES` setting. For example, to enable the media context processor, add `django.template.context_processors.media`. We will cover how to do this in detail in *Exercise 8.02 – Template settings and using MEDIA_URL in templates*.

Once the media context processor is enabled, the `MEDIA_URL` variable can be accessed inside a template just like a normal variable:

```
{{ MEDIA_URL }}
```

You could use it, for example, to source an image:

```

```

Note that, unlike with static files, there is no template tag for loading media files (i.e., there is no equivalent to the `{% static %}` template tag).

Custom context processors can also be written. For example, referring back to the Bookr application that we built, we might want to show a list of the five latest reviews in a sidebar that's on every page. A context processor like this would perform the following code:

```
from reviews.models import Review

def latest_reviews(request):
    return {"latest_reviews": Review.objects.order_by("-date_created")[:5]}.
```

This would be saved in a file named `context_processors.py` in the Bookr project directory, then referred to in the `context_processors` setting by its module path, `context_processors.latest_reviews`. Or we could save it inside the `reviews` app and refer to it as `reviews.context_processors.latest_reviews`. It is up to you to decide whether a context processor should be considered project-wide or app specific. However, bear in mind that regardless of where it is stored, once activated, it applies to all the `render` calls for all apps.

A context processor can return a dictionary with multiple items or even zero items. It would do this if it had conditions to only add items if certain criteria were met. For example, show the latest reviews only if the user is logged in.

Let us explore this in detail in the next exercise.

Exercise 8.02 – template settings and using MEDIA_URL in templates

In this exercise, you will continue with `media_project` and configure Django to automatically add the `MEDIA_URL` setting to every template. You do this by adding `django.template.context_processors.media` to the `TEMPLATES context_processors` setting. You'll then add a template that uses this new variable and an example view to render it. You will make changes to the view and template throughout the exercises in this chapter:

1. In PyCharm, open `settings.py`. First, you'll need to add `media_example` to the `INSTALLED_APPS` setting since it wasn't done when the project was set up:

```
INSTALLED_APPS = [
    # other apps truncated for brevity
    "media_example",
]
```

2. About halfway down the file, you'll find the `TEMPLATES` setting, which is a dictionary. Inside it is the `OPTIONS` item (another dictionary). Inside `OPTIONS` is the `context_processors` setting.

Add the following to the end of this list:

```
"django.template.context_processors.media"
```

The full list should look like this:

```
TEMPLATES = [
    {
        "BACKEND": "django.template.backends.djangoproject.DjangoTemplates",
        "DIRS": [],
        "APP_DIRS": True,
        "OPTIONS": {
            "context_processors": [
                "django.template.context_processors.debug",
                "django.template.context_processors.request",
                "django.contrib.auth.context_processors.auth",
                "django.contrib.messages.context_processors.messages",
                "django.template.context_processors.media",
            ],
        },
    },
]
```

```

        },
    },
]
```

The complete file should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.02/media_project/media_project/settings.py.

3. Open the `media_example` app's `views.py` file and create a new view called `media_example`. For now, it can just render a template named `media-example.html` (you will create this in *step 5*). The entire code of the view function is like this:

```

def media_example(request):
    return render(request, "media-example.html")
```

Save `views.py`. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.02/media_project/media_example/views.py.

4. You need a URL mapping to the `media_example` view. Open the `media_project` package's `urls.py` file.

First, import `media_example.views` with the other imports in the file:

```
import media_example.views
```

Then add `path` into `urlpatterns` to map `media-example/` to the `media_example` view:

```

path('media-example/',
      media_example.views.media_example)
```

Your full `urlpatterns` should look like this code block:

```

from django.conf.urls.static import static

import media_example.views

urlpatterns = [
    path("admin/", admin.site.urls),
    path("media-example/",
          media_example.views.media_example)
]

if settings.DEBUG:
    urlpatterns += static(settings.MEDIA_URL,
                      document_root=settings.MEDIA_ROOT)
```

You can save and close the file.

5. Create a `templates` directory inside the `media_example` app directory. Then, create a new HTML file inside the `media_project` project `templates` directory. Select **HTML 5 file** and name the file `media-example.html`:

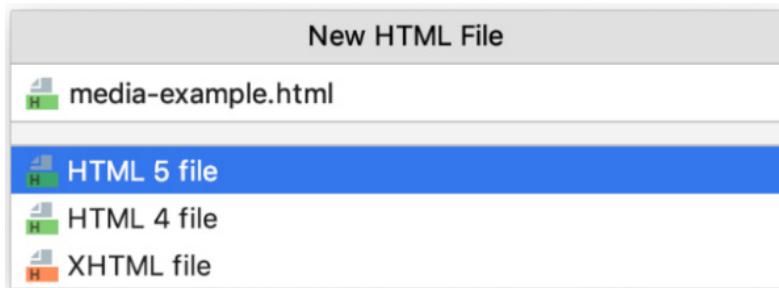


Figure 8.4: Creating `media-example.html`

6. The `media-example.html` file should open automatically. You are just going to add a link inside the file to the `test.txt` file you created in *Exercise 8.01 – configuring media storage and serving*. Inside the `<body>` element, add the highlighted code:

```
<body>
  <a href="{{ MEDIA_URL }}test.txt">Test Text
  File</a>
</body>
```

Note that there is no `/` between `MEDIA_URL` and the file name – this is because we already added a trailing slash when we defined it in `settings.py`. You can save the file. The complete file will look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.02/media_project/media_example/templates/media-example.html.

7. Start the Django dev server if it's not already running, then visit `http://127.0.0.1:8000/media-example/`. You should see a simple page like in *Figure 8.5*:



Figure 8.5: Basic media link page

If you click the link, you will be taken to the `test.txt` display and see the *Hello, world!* text you created in *Exercise 8.01 – configuring media storage and serving* (*Figure 8.5*). This means you have configured the Django `context_processors` settings correctly.

We have finished with `test.txt`, so you can delete the file now. We will use the `media_example` view and template in the other exercises, so retain them. In the next section, we will talk about how to upload files using a web browser and how Django accesses them in a view.

File uploads using HTML forms

In *Chapter 6, Forms*, we learned about HTML forms. We discussed how to use the `method` attribute of `<form>` for the GET or POST requests. And so far, we have only submitted text data using a form, but it is also possible to submit one or more files using a form.

When submitting files, we must ensure that there are at least two attributes on the form: `method` and `enctype`. You may still also need other attributes, such as `action`. A form that supports file uploads might look like this:

```
<form method="post" enctype="multipart/form-data">
```

File uploads are only available for POST requests. They are not possible with GET requests, as it would be impossible to send all the data for a file through a URL. The `enctype` attribute must be set to let the browser know it should send the form data as multiple parts, one part for the text data of the form and separate parts for each of the files that have been attached to the form. This encoding is seamless to the user, they don't know how the browser is encoding the form, nor do they need to do anything different.

To attach files to a form, you need to create an input of the `file` type. You can manually write the HTML code like this:

```
<input type="file" name="file-upload-name">
```

When the input is rendered in the browser, it looks like this when empty:

Browse... No file selected.

Figure 8.6: Empty file input

The title of the button might be different depending on your browser.

Clicking the **Browse...** button will display a *file open* dialog box:

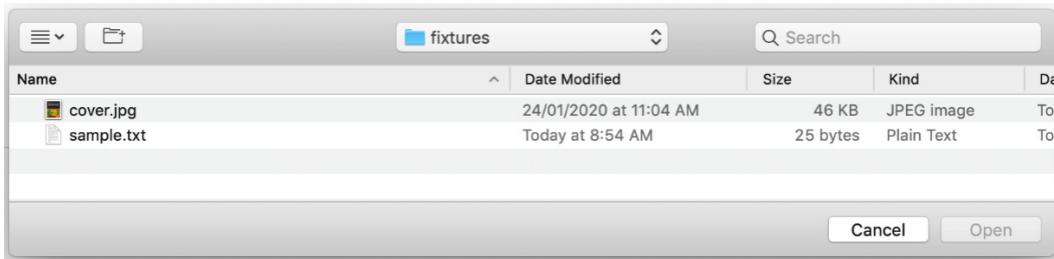


Figure 8.7: File browser on macOS

After selecting a file, the name of the file is shown in the field:

Browse... cover.jpg

Figure 8.8: File input with cover.jpg selected

Figure 8.8 shows a file input with a file named `cover.jpg` having been selected.

We briefly looked at how file upload fields look in the browser, let's start looking at the Django side of the file upload process.

Working with uploaded files in a view

In addition to text data, if a form also contains file uploads, Django will populate the `request.FILES` attribute with uploaded files. The `request.FILES` attribute is a dictionary-like object that is keyed on the name attribute given to the `file` input.

In the form example in the previous section, the file input had the name `file-upload-name`. So the file would be accessible in the view using `request.FILES["file-upload-name"]`.

The objects that `request.FILES` contains are file-like objects (specifically a `django.core.files.uploadedfile.UploadedFile` instance), so to use them, you must read their data. For example, to get the content of an uploaded file in your view, you can write the following:

```
content = request.FILES["file-upload-name"].read()
```

A more common action is to write the file contents to disk. When files are uploaded, they are stored in a temporary location (in memory if they are under 2.5 MB, otherwise in a temporary file on disk). To store the file data in a known location, the contents must be read and then written to disk at that location. `UploadedFile` has a `chunks` method that will read the file data one chunk at a time to prevent too much memory from being used by reading the entirety of the file at once.

So, instead of simply using the `read` and `write` functions, use the `chunks` method to only read small chunks of the file into memory at a time:

```
with open("/path/to/output.jpg", "wb+") as output_file:
    uploaded_file = request.FILES["file-upload-name"]
    for chunk in uploaded_file.chunks():
        output_file.write(chunk)
```

Note that in some of the upcoming examples, we will refer to this code as the `save_file_upload` function. Assume the function is defined like this:

```
def save_file_upload(upload, save_path):
    with open(save_path, "wb+") as output_file:
        for chunk in upload.chunks():
            output_file.write(chunk)
```

The previous example code could then be refactored to call the function:

```
uploaded_file = request.FILES["file-upload-name"]
save_file_upload(uploaded_file, "/path/to/output.jpg")
```

Each `UploadedFile` object (the `uploaded_file` variable in the previous example code snippets) also contains extra metadata about the uploaded file, such as the file's name, size, and content type. The attributes you'll find most useful are the following:

- `size`: As the name suggests, this is the size of the uploaded file in bytes.
- `name`: This refers to the name of the uploaded file, for example, `image.jpg`, `file.txt`, `document.pdf`, and so on. This value is sent by the browser.
- `content_type`: The content type (MIME type) of the uploaded file. For example, `image/jpeg`, `text/plain`, `application/pdf`, and so on. Like `name`, this value is sent by the browser.
- `charset`: This refers to the charset or text encoding of the uploaded file for text files. This will be something like `UTF-8` or `ASCII`. Once again, this value is also determined and sent by the browser.

Here is a quick example of accessing these attributes (such as inside a view):

```
upload = request.FILES["file-upload-name"]
size = upload.size
name = upload.name
content_type = upload.content_type
charset = upload.charset
```

Security and trust of browsers' sent values

As we just described, `UploadedFile` values for `name`, `content_type`, and `charset` are determined by the browser. This is important to consider because a malicious user could send fake values in place of real ones to disguise the actual files being uploaded. Django does not automatically try to determine the content type or charset of the uploaded file, so it relies on the client to be accurate when it sends this information.

If we manually handle the saving of file uploads without suitable checks, then a scenario like this could happen:

1. A user of the site uploads a malicious `malware.exe` executable but sends the `image/jpeg` content type.
2. Our code checks the content type and considers it to be safe, so it saves `malware.exe` to the `MEDIA_ROOT` file.
3. Another user of the site downloads what they think is a book cover image but is the `malware.exe` executable. They open the file, and their computer is infected with malware.

This scenario has been simplified, the malicious file would probably have a name that wasn't so obvious (maybe something like `cover.jpg.exe`), but the general process has been illustrated.

How you choose to handle the security of your uploads will depend on the specific use case, but for most cases, these tips will help:

- When you save the file to disk, generate a name instead of using the one provided by the uploader. You should replace the file extension with what you expect. For example, if a file is named `cover.exe`, but the content type is `image/jpeg`, save the file as `cover.jpg`. You could also generate a completely random filename for extra security.
- Check that the filename extension matches the content type. This method is not foolproof, as there are so many MIME types that if you are handling uncommon files you might not get a match. The built-in Python `mimetypes` module can help you here. Its `guess_type` function takes a filename and returns a `mimetype` tuple (content type) and encoding. Here's a short snippet showing its use in a Python console:

```
>>> import mimetypes
>>> mimetypes.guess_type('file.jpg')
('image/jpeg', None)
>>> mimetypes.guess_type('text.html')
('text/html', None)
>>> mimetypes.guess_type('unknownfile.abc')
(None, None)
>>> mimetypes.guess_type('archive.tar.gz')
('application/x-tar', 'gzip')
```

Either element of the tuple might be `None` if the type or encoding can't be guessed. Once imported into your file by doing `import mimetypes`, you can use it like this in your view function:

```
upload = request.FILES["file-upload-name"]
mimetype, encoding = mimetypes.guess_type(upload.name)
if mimetype != upload.content_type:
    raise TypeError("Mimetype doesn't match file
extension.")
```

This method will work for common file types such as images but as mentioned, many uncommon types may return `None` for `mimetype`.

- If you are expecting image uploads, use the `Pillow` library to try to open the uploaded file as an image. If it is not a valid image, then `Pillow` will be unable to open it. This is what Django does when using `ImageField` to upload images. We will show how to use this technique to open and manipulate an image in *Exercise 8.05 – image uploads using Django forms*.

- You can also consider the `python-magic` Python package, which examines the actual content of files to try to determine their type. It is installed using PIP, and its GitHub project is found here: <https://github.com/ahupp/python-magic>. Once installed and imported into your file with `import magic`, you can use it like this in your view function:

```
upload = request.FILES["field_name"]
mimetype = magic.from_buffer(upload.read(2048),
    mime=True)
```

You can then verify that `mimetype` is in a list of allowed types.

This is not a definitive list of all the ways of protecting against malicious file uploads. The best approach will depend on what type of application you are building. You might build a site for hosting arbitrary files, in which case you wouldn't need any kind of content checking at all.

Let us now see how we can build an HTML form and view that allows files to be uploaded. We will then store them inside the `media` directory and retrieve the downloaded files in our browser.

Exercise 8.03 – file upload and download

In this exercise, you will add a form with a file field to the `media-example.html` template. This will allow you to upload a file to the `media_example` view using your browser. You'll also update the `media_example` view to save the file to the `MEDIA_ROOT` directory so that it's available for download. You will then test that this all works by downloading the file again:

1. In PyCharm, open the `media-example.html` template located inside the `templates` folder. Inside the `<body>` element, remove the `<a>` link that was added in step 6 of *Exercise 8.02 – template settings and using MEDIA_URL in templates*. Replace it with a `<form>` element (highlighted in the following code snippet). Make sure the opening tag has `method="post"` and `enctype="multipart/form-data"`:

```
</head>
<body>
    <form method="post" enctype="multipart/form-data">

        </form>
    </body>
```

2. Insert the `{% csrf_token %}` template tag inside the `<form>` body.
3. After `{% csrf_token %}`, add an `<input>` element, with `type="file"` and `name="file_upload"`:

```
<input type="file" name="file_upload">
```

4. Finally, before the closing `</form>` tag, add a `<button>` element with `type="submit"` and the `Submit` text content:

```
<button type="submit">Submit</button>
```

Your HTML body should now look like this:

```
<body>
  <form method="post" enctype="multipart/form-data">
    {% csrf_token %}
    <input type="file" name="file_upload">
    <button type="submit">Submit</button>
  </form>
</body>
```

Now, save and close the file. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.03/media_project/media_example/templates/media-example.html.

5. Open the `media_example` app's `views.py` file. Inside the `media_example` view, add code to save the uploaded file to the `MEDIA_ROOT` directory. For this, you need access to `MEDIA_ROOT` from settings, so import the Django settings at the top of the file:

```
from django.conf import settings
```

6. The uploaded file should only be saved if the request method is `POST`. Inside the `media_example` view, add an `if` statement to validate that `request.method` is `POST`:

```
def media_example(request):
    if request.method == "POST":
        ...
```

7. Inside the `if` statement added in the previous step, generate the output path by joining the uploaded filename to the `MEDIA_ROOT` directory. Then, open this path in the `wb` mode and iterate over the uploaded file using the `chunks` method. Finally, write each chunk to the saved file:

```
def media_example(request):
    if request.method == 'POST':
        save_path = settings.MEDIA_ROOT /
        request.FILES["file_upload"].name

        with open(save_path, "wb") as output_file:
            for chunk in
            request.FILES["file_upload"].chunks():
                output_file.write(chunk)

    return render(request, "media-example.html")
```

Note that the uploaded file and its metadata are being accessed from the `request.FILES` dictionary using the key that matches the name given to the file input (in our case, this is `file_upload`). You can save and close `views.py`. It should now look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.03/media_project/media_example/views.py.

8. Start the Django dev server if it's not already running, then navigate to `http://127.0.0.1:8000/media-example/`. You should see the file upload field and the **Submit** button, as can be seen here:



Figure 8.9: The file upload form

Click **Browse...** (or the equivalent on your browser) and select a file to upload. The name of the file will appear in the file input. Then, click **Submit**.

The page will reload, and the form will be empty again. This is normal – in the background, the file should have been saved.

9. Try to download the file you uploaded using `MEDIA_URL`. In this example, a file named `cover.jpg` was uploaded. It will be downloadable at `http://127.0.0.1:8000/media/cover.jpg`. Your URL will depend on the name of the file you uploaded:

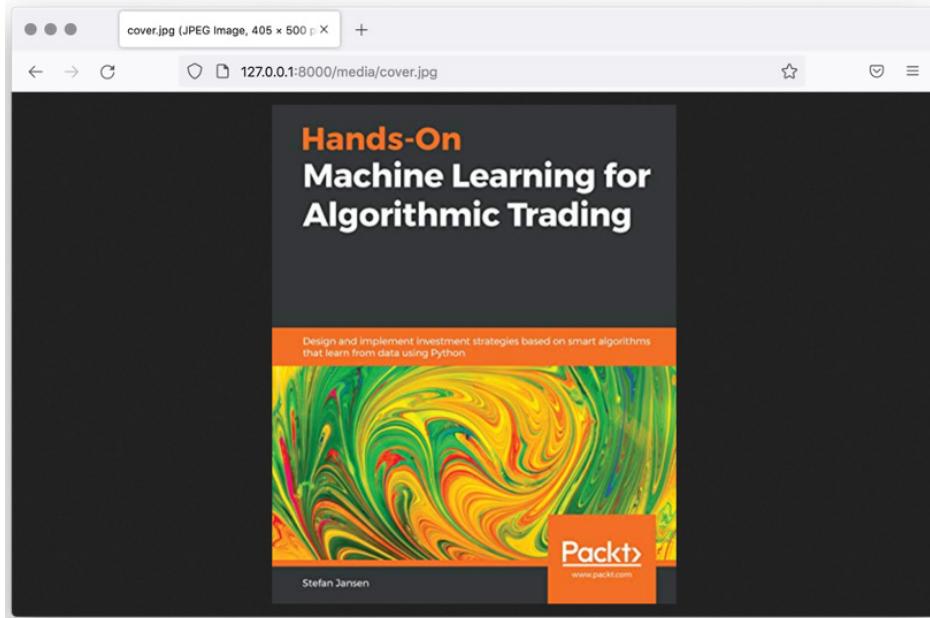


Figure 8.10: Uploaded file visible inside MEDIA_URL

If you upload an image file, HTML file, or another type of file your browser can display, you will be able to view it inside the browser. Otherwise, your browser will just download it to disk again. In both cases, it means the upload was successful.

You can also confirm the upload was successful by looking inside the `media` directory in the `media_project` project directory:



Figure 8.11: cover.jpg inside the media directory

Figure 8.11 shows the `cover.jpg` file inside the `media` directory in PyCharm.

In this exercise, you added an HTML form with `enctype` set to `multipart/form-data` so that it allowed file uploads. It contained a `file` input to select a file to upload. You then added a saving functionality to the `media_example` view to save the uploaded file to disk.

In the next section, we will look at how to simplify form generation and add validation using Django forms for file upload handling.

File uploads with Django forms

In *Chapter 6, Forms*, we saw how Django makes it easy to define forms and automatically render them to HTML. In the previous example, we defined our form manually and wrote the HTML. We can replace this with a Django form and implement the file input with a `FileField` constructor.

Here's how `FileField` is defined on a form:

```
from django import forms

class ExampleForm(forms.Form):
    file_upload = forms.FileField()
```

The `FileField` constructor can take the following keyword arguments:

- `Required`: This should be `True` for required fields and `False` if the field is optional
- `max_length`: This refers to the maximum length of the filename of the file being uploaded
- `allow_empty_file`: A field with this argument is considered to be valid even if the uploaded file is empty (has a size of 0)

Apart from these three keyword arguments, the constructor can also accept the standard `Field` arguments, such as `widget`. The default widget class for `FileField` is `ClearableFileInput`; this is a file input that can display a checkbox that can be checked to send a null value and clear the saved file on a Model field.

Using a form with a `FileField` constructor in a view is similar to other forms, but when the form has been submitted (i.e., `request.METHOD` is `POST`), then `request.FILES` should be passed into the form constructor as well. This is because Django needs to access `request.FILES` to find information about uploaded files when validating the form.

The basic flow in a view function is, therefore, similar to this:

```
def view(request):
    if request.method == "POST":
        # instantiate the form with POST data and files
        form = ExampleForm(request.POST, request.FILES)
        if form.is_valid():
            # process the form and save files
            return redirect("success-url")
    else:
        # instantiate an empty form as we've seen before
        form = ExampleForm()
```

```
# render a template, the same as for other forms
return render(request, "template.html", {"form": form})
```

When working with uploaded files and forms, you can interact with the uploaded files by accessing them through `request.FILES`, or through `form.cleaned_data`; the values will return to the same object. In the preceding example, we could process the uploaded file like this:

```
if form.is_valid():
    save_file_upload("/path/to/save.jpg",
                    request.FILES["file_upload"])
    return redirect("/success-url/")
```

Or, since they contain the same object, you can use `form.cleaned_data`:

```
if form.is_valid():
    save_file_upload("/path/to/save.jpg",
                    form.cleaned_data["file_upload"])
    return redirect("/success-url/")
```

The data that is saved will be the same.

Note

In *Chapter 6, Forms*, you experimented with forms and submitted them with invalid values. When the page refreshed to show the form errors, the data that you had previously entered was populated when the page reloaded. This does not occur with file fields; instead, the user will have to navigate and select the file again if the form is invalid.

In the next exercise, we will put what we have seen with `FileFields` into practice by building an example form and then modifying our view to save the file only if the form is valid.

Exercise 8.04 – file uploads with a Django form

In the previous exercise, you created a form in HTML and used it to upload a file to a Django view. If you tried submitting the form without selecting a file, you would get a Django exception screen. You did not do any validation on the form, so this method is quite fragile.

In this exercise, you will create a Django form with `FileField`, which will allow you to use form validation functions to make the view more robust as well as reduce the amount of code:

1. In PyCharm, inside the `media_example` app, create a new file named `forms.py`. It will open automatically. At the start of the file, import the Django `forms` library:

```
from django import forms
```

Then, create a `forms.Form` subclass, and name it `UploadForm`. Add to it a `FileField` field named `file_upload`. Your class should have this code:

```
class UploadForm(forms.Form):  
    file_upload = forms.FileField()
```

You can save and close this file. The complete file should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.04/media_project/media_example/forms.py.

2. Open the `form_example` app's `views.py` file. At the start of the file, right below the existing `import` statements, you will need to import your new class like this:

```
from .forms import UploadForm
```

3. If you are in the `POST` branch of the view, `UploadForm` needs to be instantiated with both `request.POST` and `request.FILES`. If you don't pass in `request.FILES`, then the form instance will not be able to access the uploaded files. Under the `if request.method == "POST"` check, instantiate `UploadForm` with these two arguments:

```
form = UploadForm(request.POST, request.FILES)
```

4. The existing lines that define `save_path` and store the file contents can be retained, but they should be indented by one block and put inside a form validity check, so they are only executed if the form is valid. Add the `if form.is_valid():` line, and then indent the other lines, so the code looks like this:

```
if form.is_valid():  
    save_path = os.path.join(settings.MEDIA_ROOT,  
    request.FILES["file_upload"].name)  
  
    with open(save_path, "wb") as output_file:  
        for chunk in  
            request.FILES["file_upload"].chunks():  
                output_file.write(chunk)
```

5. Since you are using a form now, you can access the file upload through the form. Replace usages of `request.FILES["file_upload"]` with `form.cleaned_data["file_upload"]`:

```
if form.is_valid():
    save_path = settings.MEDIA_ROOT /
    form.cleaned_data["file_upload"].name

    with open(save_path, "wb") as output_file:
        for chunk in
            form.cleaned_data["file_upload"].chunks():
                output_file.write(chunk)
```

6. Finally, add an `else` branch to handle non-POST requests, which simply instantiates a form without any arguments:

```
if request.method == "POST":
    ...
else:
    form = UploadForm()
```

7. Add a context dictionary argument to the `render` call and set the `form` variable in the `form` key:

```
return render(request, "media-example.html", {"form": form})
```

You can now save and close this file. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.04/media_project/media_example/views.py.

8. Finally, open the `media-example.html` template and remove your manually defined `<input>` file. Replace it with `form`, rendered using the `as_p` method (highlighted):

```
<body>
    <form method="post" enctype="multipart/form-data">
        {% csrf_token %}
        {{ form.as_p }}
        <button type="submit">Submit</button>
    </form>
</body>
```

You should not change any other parts of the file. You can save and close this file. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.04/media_project/media_example/templates/media-example.html.

9. Start the Django dev server if it's not already running, then navigate to `http://127.0.0.1:8000/media-example/`. You should see the file upload field and the **Submit** button as follows:



File upload: No file selected.

Figure 8.12: The file upload Django form rendered in the browser

10. Since we're using a Django form, we get its built-in validation automatically. If you try to submit the form without selecting a file, your browser should prevent it and show an error, as can be seen here:



File upload: No file selected.

Please select a file.

Figure 8.13: Form submission prevented by the browser

11. Finally, repeat the upload test that you performed in *Exercise 8.03 – file upload and download*, by selecting a file and submitting the form. You should then be able to retrieve the file using `MEDIA_URL`. In this case, a file named `cover.jpg` is being uploaded again (see the following screenshot):



Figure 8.14: Uploading a file named `cover.jpg`

You can then retrieve the file at `http://127.0.0.1:8000/media/cover.jpg`, and you can see it in the browser as follows:

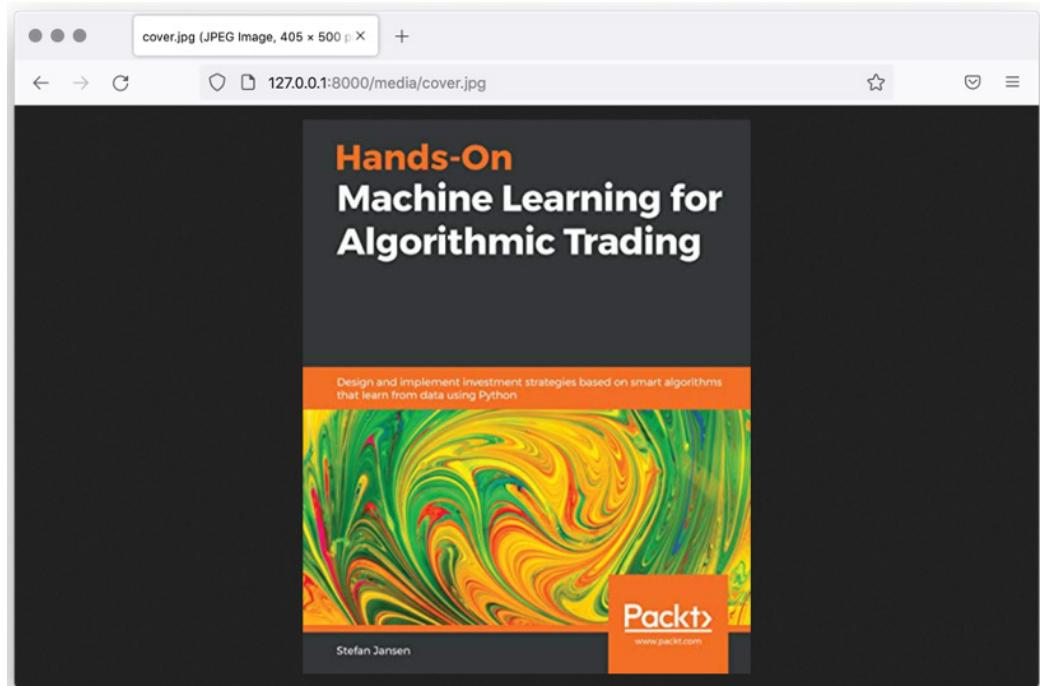


Figure 8.15: The file uploaded using the Django form is also visible in the browser

In this exercise, we replaced a manually built form with a Django form containing `FileField`. We instantiated the form in the view by passing in both `request.POST` and `request.FILES`. We then used the standard `is_valid` method to check the validity of the form and only saved the file upload if the form was valid. We tested the file uploading and saw we were able to retrieve uploaded files using `MEDIA_URL`.

In the next section, we will look at `ImageField`, which is like `FileField`, but specifically for images.

Image uploads with Django forms

If you want to work with images in Python, the most common library that you'll use is called `Pillow`, and this is the library Django uses to validate images. Originally there was a library called **Python Imaging Library (PIL)**. It was not kept up to date and, eventually, a fork of the library was created and is still maintained, `Pillow`. To maintain backward compatibility, the package is still called `PIL` when installed. For example, the `Image` object is imported from `PIL`:

```
from PIL import Image
```

The terms Python Imaging Library, `PIL`, and `Pillow` are often used interchangeably. You can assume that if someone refers to `PIL`, they mean the latest `Pillow` library.

`Pillow` provides various methods of retrieving data about images or manipulating them. You can find out the width and height of images, scale, crop, and apply transformations to them. There are too many operations available to cover in this chapter, so we will just introduce a simple example (scaling an image), which you will use in the next exercise.

Since images are one of the most common types of files that a user may want to upload, Django also includes `ImageField`. This behaves similarly to `FileField` but also automatically validates that the data is an image file. This helps mitigate against security issues where we expect an image, but the user uploads a malicious file.

An `UploadedFile` from an `ImageField` has all the same attributes and methods as that of `FileField` (`size`, `content_type`, `name`, `chunks()`, etc.) but adds an extra attribute: `image`. This is an instance of the `PIL Image` object that is used to verify that the file being uploaded is a valid image.

After checking that the form is valid, the underlying `PIL Image` object is closed. This is done to free up memory and prevent the Python process from holding too many files open, which could cause performance issues. What this means for the developer is that you can access some of the metadata about the image (such as `width`, `height`, and `format`), but you can't access the actual image data without re-opening the image.

To illustrate, we'll have a form with `ImageField`, named `picture`:

```
class ExampleForm(forms.Form):
    picture = forms.ImageField()
```

Inside the `view` function, the `picture` field can be accessed in the form's `cleaned_data`:

```
if form.is_valid():
    picture_field = form.cleaned_data["picture"]
```

Then, the `picture` field's `Image` object can be retrieved:

```
image = picture_field.image
```

Now that we have a reference to the image in the view, we can get some metadata:

```
w = image.width  # an integer, e.g. 600
h = image.height # also an integer, e.g. 420
f = image.format # the format of the image as a string, e.g. "PNG"
```

Django will also automatically update the `content_type` attribute of `UploadedFile` to the correct type for the `picture` field. This overwrites the value that the browser has sent when uploading the file.

Attempting to use a method that accesses the actual image data (rather than just the metadata) will cause an exception to be raised. This is because Django has already closed the underlying image file.

For example, the following code snippet will raise `AttributeError`:

```
image.getdata()
```

Instead, we need to re-open the image. The image data can be opened with the `ImageField` reference after importing the `Image` class:

```
from PIL import Image

image = Image.open(picture_field)
```

Now that the image has been opened, you can perform operations on it. In the next section, we will look at a simple example – resizing the uploaded image.

Resizing an image with Pillow

Pillow supports many operations that you might want to perform on an image before saving it. We can't explain them all in this book, so we will just use a common operation: resizing an image to a specific size before saving it. This will help us save storage space and improve download speed. For example, a user may upload large cover images in Bookr that are bigger than are needed for our purposes. When saving the file (writing it back to disk), we must specify the format to use. We could determine the type of image being uploaded with a number of methods (such as checking `content_type` of the uploaded file or `format` from the `Image` object), but in our example, we'll always just save it as a JPEG file.

The PIL `Image` class has a `thumbnail` method that will resize an image to a maximum size while retaining the aspect ratio. For example, we could set a maximum size of 50 **pixels (px)** by 50 px. A 200 px by 100 px image would be resized to 50 px by 25 px: the aspect ratio is retained by setting the maximum dimension to 50 px. Each dimension is scaled by a factor of 0.25:

```
from PIL import Image

size = 50, 50 # a tuple of width, height to resize to
image = Image.open(image_field) # open the image as before
image.thumbnail(size) # perform the resize
```

At this point, the resize has been done in memory only. The change is not saved to disk until the `save` method is called, like so:

```
image.save("path/to/file.jpg")
```

The output format is automatically determined from the file extension used, in this case, JPEG. The `save` method can also take a `format` argument to override it, for example:

```
image.save("path/to/file.png", "JPEG")
```

Despite having the `.png` extension, the format is specified as JPEG, and so the output will be in JPEG format. As you might imagine, this can be very confusing, so you might decide to stick with specifying the extension only.

In the next exercise, we will change `UploadForm` we have been working with to use `ImageField` instead of `FileField`, then implement resizing of an uploaded image before saving it to the media directory.

Exercise 8.05 – image uploads using Django forms

In this exercise, you will update the `UploadForm` class you created in *Exercise 8.04 – file uploads with a Django form*, to use `ImageField` instead of `FileField` (this will involve simply changing the field's class). You will then see that the form renders the same in the browser. Next, you will try uploading some non-image files and see how Django validates the form to disallow them. Finally, you will update your view to use PIL to resize the image before saving it, and then test it in action:

1. Open the `media_example` app's `forms.py` file. In the `UploadForm` class, change `file_upload`, so that it's an instance of `ImageField` instead of `FileField`. After updating, `UploadForm` should look like this:

```
class UploadForm(forms.Form):  
    file_upload = forms.ImageField()
```

Save and close the file. Your `forms.py` file should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.05/media_project/media_example/forms.py.

2. Start the Django dev server if it's not already running, then navigate to `http://127.0.0.1:8000/media-example/`. You should see the form rendered, and it will look identical to when we used `FileField` (see the following screenshot):



Figure 8.16: `ImageField` looks the same as `FileField`

3. You will notice the difference when you try to upload a non-image file. Click the **Browse...** button and try to select a non-image file. Depending on your browser or operating system, you might not be able to select anything other than an image file, as in *Figure 8.17*:

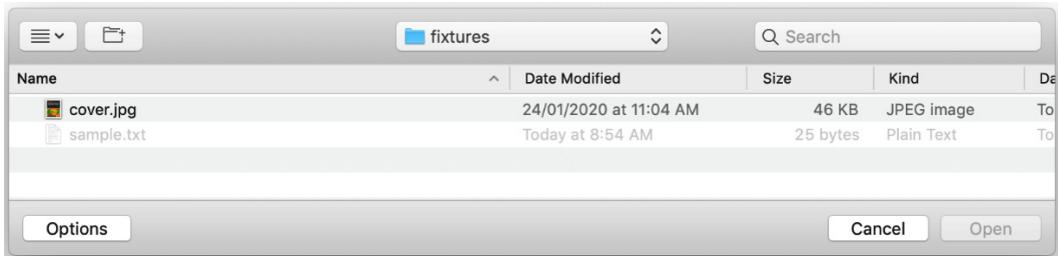


Figure 8.17: Only image files are selectable

Your browser may allow you to select an image but show an error in the form after selection. Or your browser may allow you to select a file and submit the form, and Django will raise `ValidationError`. Regardless, you can be sure that in your view, the form's `is_valid` view will only return `True` if an image has been uploaded.

Note

You don't need to test uploading a file at this point, as the result would be the same as in *Exercise 8.04 – file uploads with a Django form*.

4. The first thing you will need to do is make sure the Pillow library is installed. In a terminal (making sure your virtual environment has been activated), run the following command:

```
pip3 install pillow
```

(In Windows, this is `pip install pillow`). You'll get output like *Figure 8.18*:

```
media_project — ben@BensMBP — ~/media_project — -zsh — 80x24
(bookr) → media_project pip3 install pillow
Looking in indexes: http://localhost:3141/root/pypi/+simple/
Collecting pillow
  Downloading http://localhost:3141/root/pypi/%2Bf/80c/a33961ced9c63/Pillow-9.0.1-cp39-cp39-macosx_10_10_x86_64.whl (3.0 MB)
  3.0/3.0 MB 241.7 MB/s eta 0:00:00
Installing collected packages: pillow
Successfully installed pillow-9.0.1
(bookr) → media_project
```

Figure 8.18: pip3 installing pillow

If Pillow is already installed, you'll see the Requirement already satisfied output message.

5. Now, we can update the `media_example` view to resize the image before saving it. Switch back to PyCharm and open the `media_example` app's `views.py` file, then import PIL's `Image` class by adding this import line near the top of the file:

```
from PIL import Image
```

6. Go to the `media_example` view. Under the line that generates `save_path`, take out the three lines that open the output file, iterate over the uploaded file, and write out its chunks. Replace this with the code that opens the uploaded file with PIL, resizes it, then saves it:

```
image = Image.open(form.cleaned_data["file_upload"])
image.thumbnail((50, 50))
image.save(save_path)
```

The first line creates an `Image` instance by opening the uploaded file, the next performs the thumbnail conversion (to a maximum size of 50 px by 50 px), and the third line saves the file to the same save path that we have been generating in previous exercises. You can save the file. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.05/media_project/media_example/views.py.

7. The Django dev server should still be running from *step 2*, but you should start it if it is not. Then, navigate to `http://127.0.0.1:8000/media-example/`. You'll see the familiar `UploadForm`. Select an image and submit the form. If the upload and resize was successful, the form will refresh and be empty again.
8. View the uploaded image using `MEDIA_URL`. For example, a file named `cover.jpg` will be downloadable from `http://127.0.0.1:8000/media/cover.jpg`. You should see the image has been resized to have a maximum dimension of just 50 px:

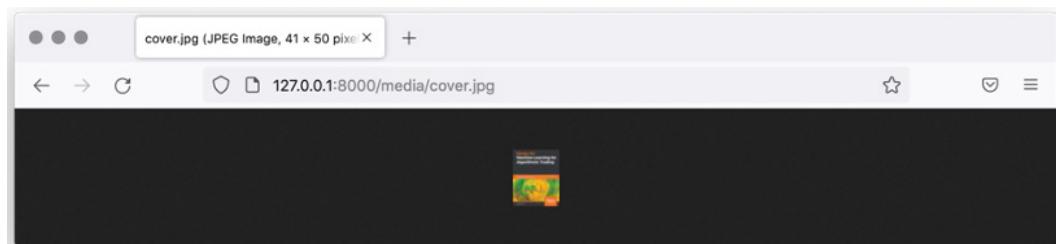


Figure 8.19: The resized cover

While a thumbnail this size might not be that useful, it at least lets us be sure that the image resize has worked correctly.

In this exercise, we changed `FileField` on `UploadForm` to `ImageField`. We saw that the browser wouldn't let us upload anything other than images. We then added code to the `media_example` view to resize the uploaded image using PIL.

We have encouraged the use of a separate web server to serve static and media files for performance reasons. However, in some cases, you might want to use Django to serve files, for example, to provide authentication before allowing access. In the next section, we will discuss how to use Django to serve media files.

Serving uploaded (and other) files using Django

Throughout this chapter and *Chapter 5, Serving Static Files*, we have discouraged serving files using Django. This is because it would needlessly tie up a Python process just serving a file – something that the webserver is capable of handling. Unfortunately, web servers do not usually provide dynamic access control, that is, allowing only authenticated users to download a file. Depending on your web server used in production, you might be able to have it authenticate against Django and then serve the file itself; however, the specific configuration of specific web servers is outside the scope of this book.

One approach you can take is to specify a subdirectory of your `MEDIA_ROOT` directory and have your web server prevent access to just this specific folder. Any protected media should be stored inside it. If you do this, only Django will be able to read the files inside. For example, your web server could serve everything in the `MEDIA_ROOT` directory except for a `MEDIA_ROOT/protected` directory.

Another approach would be to configure a Django view to serve a specific file from disk. The view will determine the path of the file on disk to send, then send it using the `FileResponse` class. The `FileResponse` class takes an open file handle as an argument and tries to determine the correct content type from the file's content. Django will close the file handle after the request completes.

The `view` function will accept the request and a relative path to the file to be downloaded as parameters. This relative path is the path inside the `MEDIA_ROOT/protected` folder.

In our case, we will just check whether the user is anonymous (not logged in). We'll do this by checking the `request.user.is_anonymous` property. If they aren't logged in, then we will raise a `django.core.exceptions.PermissionDenied` exception, which returns an **HTTP 403 Forbidden** response to the browser. This will stop the execution of the view and not return any file:

```
import os.path
from django.conf import settings
from django.http import FileResponse
from django.core.exceptions import PermissionDenied
```

```
def download_view(request, relative_path):
    if request.user.is_anonymous:
        raise PermissionDenied
    full_path = os.path.join(settings.MEDIA_ROOT,
        "protected", relative_path)
    file_handle = open(full_path, "rb")
    return FileResponse(file_handle)  # Django sends the
                                    # file then closes
                                    # the handle
```

The URL mapping to this view could be like this, using the `path` path converter inside your `urls.py` file:

```
urlpatterns = [
    ...
    path("downloads/<path:relative_path>",
        views.download_view)
]
```

There are many ways that you could choose to implement a view that sends files. The important thing is that you use the `FileResponse` class, which is designed to stream the file to the client in chunks instead of loading it all into memory. This will reduce the load on the server and lessen the impact on resource usage if you must resort to sending files with Django.

Now that we know how to upload and process images, let's look at how to store them on models.

Storing files on model instances

So far, we have manually managed the uploading and saving of files. You can also associate a file with a model instance by assigning the path to which it was saved to `CharField`. However, as with much of Django, this capability (and more) is already provided with the `models.FileField` class. A `FileField` instance does not actually store the file data; instead, they store the path where the file is stored (like `CharField` would), but it also provides helper methods. These methods assist with loading files (so you don't have to open them manually) and generating disk paths for you based on the ID of the instance (or other attributes).

`FileField` can accept two specific optional arguments in its constructor (as well as the base `Field` arguments, such as `required`, `unique`, `help_text`, etc.):

- `max_length`: Like `max_length` in the form's `ImageField`, this is the maximum length of the filename that is allowed.
- `upload_to`: The `upload_to` argument has three different behaviors depending on what type of variable is passed to it. Its simplest use is with a string or a `pathlib.Path` object. The path is simply appended to `MEDIA_ROOT`.

In this example, `upload_to` is just defined as a string:

```
class ExampleModel(models.Model):  
    file_field = models.FileField(upload_to="files/")
```

Files saved to `FileField` will be stored in the `MEDIA_ROOT/files` directory.

You can achieve the same result using a `pathlib.Path` instance, too:

```
import pathlib  
  
class ExampleModel(models.Model):  
    file_field =  
        models.FileField(upload_to=pathlib.Path("files/"))
```

The next way of using `upload_to` is with a string that contains the `strftime` formatting directives (for example, `%Y` to substitute the current year, `%m` for the current month, and `%d` for the current day of the month). The full list of these directives is extensive and can be found at <https://docs.python.org/3/library/time.html#time.strftime>. Django will automatically interpolate these values when saving the file.

For example, if you defined the model and `FileField` like this:

```
class ExampleModel(models.Model):  
    file_field =  
        models.FileField(upload_to="files/%Y/%m/%d/")
```

For the first file uploaded on a specific day, Django would create the directory structure for that day. For example, for the first file uploaded on January 1st, 2020, Django would create the `MEDIA_ROOT/2020/01/01` directory and then store the uploaded file there. The next file (and all subsequent ones) uploaded on the same day would also be stored in that directory. Similarly, on January 2nd, 2020, Django will create the `MEDIA_ROOT/2020/01/02` directory, and files will be stored there.

If you have many thousands of files being uploaded every day, you could even have the files split up further by including the hour and minute in the `upload_to` argument (`upload_to="files/%Y/%m/%d/%H/%M/"`). This may not be necessary if you only have a small volume of uploads, though.

By utilizing this method of the `upload_to` argument, you can have Django automatically segregate uploads and prevent too many files from being stored within a single directory (which can be hard to manage).

The final method of using `upload_to` is by passing a function that will be called to generate the storage path. Note that this is different from the other uses of `upload_to`, as it should generate the full path, including the file name, rather than just the directory. The function takes two arguments: `instance` and `filename`. The `instance` argument is the model instance that `FileField` is attached to, and `filename` is the name of the uploaded file.

Here is an example function that takes the first two characters of a filename to generate the save directory. This means that each uploaded file will be grouped into parent directories, which can help organize files and prevent there from being too many in one directory:

```
def user_grouped_file_path(instance, filename):
    return "{}/{}/{}/{}".format(instance.username,
                                filename[0].lower(), filename[1].lower(), filename)
```

If this function is called with the `Test.jpg` filename, it will return `<username>/t/e/test.jpg`. If called with `example.txt`, it will return `<username>e/x/example.txt`, and so on. `username` is retrieved from the instance that is being saved. To illustrate, here is a model with `FileField` that uses this function. It also has a `username`, which is `CharField`:

```
class ExampleModel(models.Model):
    file_field =
        models.FileField(upload_to=user_grouped_file_path)
    username = models.CharField(unique=True)
```

You can use any attribute of the instance in the `upload_to` function, but be aware that if this instance is in the process of being created, then the file save function will be called before it is saved to the database. Therefore, some of the automatically generated attributes on the instance (such as `id`/`pk`) will not yet be populated and shouldn't be used to generate a path.

Whatever path is returned from the `upload_to` function is appended to `MEDIA_ROOT`, so the uploaded files would be saved at `MEDIA_ROOT/<username>/t/e/test.jpg` and `MEDIA_ROOT/<username>/e/x/example.txt`, respectively.

Note that `user_grouped_file_path` is just an illustrative function that has intentionally been kept short, so it will not work correctly with single-character file names or if the username has invalid characters. For example, if the username has a / character in it, then this would act as a directory separator in the generated path.

Now, we've done a deep dive into setting up `FileField` on a model, but how do we actually save an uploaded file to it? It's as easy as assigning the uploaded file to the attribute of the model, as you would with any type of value. Here's a quick example with a view and the simple `ExampleModel` we were using as an example earlier in this section:

```
class ExampleModel(models.Model):
    file_field = models.FileField(upload_to="files/")

    def view(request):
        if request.method == "POST":
            m = ExampleModel()  # Create a new ExampleModel
                                instance
            m.file_field = request.FILES["uploaded_file"]
            m.save()
        return render(request, "template.html")
```

In this example, we create a new `ExampleModel` and assign the uploaded file (called `uploaded_file` on the form) to its `file_field` attribute. When we saved the model instance, Django automatically wrote the file with its name to the `upload_to` directory path. If the uploaded file had the name `image.jpg`, the save path would be `MEDIA_ROOT/upload_to/image.jpg`.

We could just have easily updated the file field on an existing model or used a form (validating it before saving). Here is another simple example demonstrating this:

```
class ExampleForm(forms.Form):
    uploaded_file = forms.FileField()

    def view(request, model_pk):
        form = ExampleForm(request.POST, request.FILES)
        if form.is_valid():
            # Get an existing model instance
            m = ExampleModel.object.get(pk=model_pk)

            # store the uploaded file on the instance
            m.file_field = form.cleaned_data["uploaded_file"]
            m.save()
        return render(request, "template.html")
```

You can see that updating `FileField` on an existing model instance is the same process as setting it on a new instance and if you choose to use a Django form or just access `request.FILES` directly, the process is just as simple.

Now we'll look at how to store images on models using `ImageField`.

Storing images on model instances

While `FileField` can store any type of file, including images, there also exists `ImageField`. As you would expect, this is only for storing images. The relationship between the models' forms. `FileField` and forms. `ImageField` is similar to that between `models.FileField` and `models.ImageField`, meaning `ImageField` extends `FileField` and adds extra methods for working with images.

The `ImageField` constructor takes the same arguments as `FileField` and adds two extra optional arguments:

- `height_field`: This is the name of the field on the model that will be updated with the height of the image every time the model instance is saved
- `width_field`: This is the width counterpart to `height_field`: the field that stores the width of the image that is updated every time the model instance is saved

Both of these arguments are optional, but the fields they name must exist if used. That is, it is valid to have `height_field` or `width_field` unset, but if they are set to the name of a field that does not exist, then an error will occur. The purpose of this is to assist with searching the database for files of a particular dimension.

Here is an example model using `ImageField`, which updates the image dimension fields:

```
class ExampleModel(models.Model):  
    image = models.ImageField(upload_to="images/%Y/%m/%d/",  
                             height_field="image_height",  
                             width_field="image_width")  
    image_height = models.IntegerField()  
    image_width = models.IntegerField()
```

Notice that `ImageField` is using the `upload_to` parameter with date formatting directives that are updated on save. The behavior of `upload_to` is identical to that of `FileField`.

Upon saving an `ExampleModel` instance, its `image_height` field would be updated with the height of the image and `image_width` with the width of the image.

We won't show examples for setting the `ImageField` values in a view, as the process is the same as for a plain `FileField`.

Working with `FieldFile`

When you access the `FileField` or `ImageField` attributes of a model instance, you will not get a native Python `file` object. Instead, you will be working with a `FieldFile` object. The `FieldFile` class is a wrapper around `file` that adds extra methods. Yes, it can be confusing to have classes called `FileField` and `FieldFile`.

The reason that Django uses `FieldFile` instead of just a `file` object is twofold. First, it adds extra methods to open, read, delete, and generate the URL of the file. Second, it provides an abstraction to allow alternative storage engines to be used.

Custom storage engines

We looked at custom storage engines in *Chapter 5, Serving Static Files*, regarding storing static files. We won't examine custom storage engines in detail about media files since the code outlined in *Chapter 5, Serving Static Files*, for static files also applies to media files. The important thing to note is that which storage engine you're using can be changed without updating your other code. This means that you can have your media files stored on your local drive during development and then save them to a CDN when your application is deployed to production.

The default `storage_engine` class can be set with `DEFAULT_FILE_STORAGE` in `settings.py`. The storage engine can also be specified on a per-field basis (for `FileField` or `ImageField`) with the `storage` argument, for example:

```
storage_engine = CustomStorageEngine()

class ExampleModel(models.Model):
    image_field = ImageField(storage=storage_engine)
```

This demonstrates what actually happens when you upload or retrieve a file. Django delegates to the storage engine to write or read it respectively. This happens even while saving to disk; however, this is fundamental and is invisible to the user.

Reading a stored FieldFile

Now that we've learned about custom storage engines, let's look at reading from `FieldFile`. In the previous sections, we saw how to set the file on the model instance. Reading the data back again is just as easy – we have a couple of different methods that can help us depending on our use case.

In the following few code snippets, assume we are inside a view and have retrieved our model instance in some manner, and it's stored in an `m` variable:

```
m = ExampleModel.objects.get(pk=model_pk)
```

We can read all the data from the file with the `read` method:

```
data = m.file_field.read()
```

Or, we can manually open the file with the `open` method. This might be useful if we want to write our own generated data to the file:

```
with m.file_field.open("wb") as f:
    chunk = f.write(b"test") # write bytes to the file
```

If we want to read the file in chunks, we can use the `chunks` method. This works the same as reading chunks from the uploaded file as we saw earlier:

```
for chunk in m.file_field.chunks():
    write_chunk(open_file, chunk) # assume this method is
        defined somewhere
```

We can also manually open the file ourselves by using its `path` attribute:

```
open(m.file_field.path)
```

If we want to stream `FileField` for download, the best way is by using the `FileResponse` class, as we saw earlier. Combine this with the `open` method on `FileField`. Note that if we're just trying to serve a media file, we should only implement a view to do this if we're trying to restrict access to the file. Otherwise, we should just serve the file using `MEDIA_URL` and allow the web server to handle the request. Here's how we'd write `download_view` to use `FileField` instead of the manually specified path:

```
def download_view(request, model_pk):
    if request.user.is_anonymous:
        raise PermissionDenied
    m = ExampleModel.objects.get(pk=model_pk)
    return FileResponse(m.file_field.open()) # Django
        sends the file then closes the handle
```

Django opens the correct path and closes it after the response. Django will also attempt to determine the correct MIME type for the file. We assume that `FileField` has its `upload_to` attribute set to a protected directory that the web server is preventing direct access to.

Storing existing files or content on FileField

We've seen how to store an uploaded file on an image field: simply assign it to the field like so:

```
m.file_field = request.FILES["file_upload"]
```

But how can we set the field value to that of an existing file that we might already have on disk? You might think you can use a standard Python `file` object, but this won't work:

```
m.file_field = open("/path/to/file.txt", "rb") # Don't do  
this
```

You might also try setting the file using some content:

```
m.file_field = "new file content" # Don't do this
```

This won't work either.

You instead need to use the `FileField` `save` method, which accepts an instance of a Django `File` object or a `ContentFile` object (these classes' full paths are `django.core.files.File` and `django.core.files.base.ContentFile`, respectively). We'll briefly discuss the `save` method and its arguments and then return to these classes.

The `FileField` `save` method takes three arguments:

- `name`: This is the name of the file you are saving and is the name the file will have when saved to the storage engine (in our case, to disk, inside `MEDIA_ROOT`).
- `content`: This is an instance of `File` or `ContentFile` that we just saw; again, we will discuss these soon.
- `save`: This argument is optional and defaults to `True`. This indicates whether or not to save the model instance to the database after saving the file. If set to `False` (i.e., the model is not saved), then the file will still be written to the storage engine (to disk), but the association is not stored on the model. The previous file path (or no file if one was not set) will still be stored in the database until the model instance's `save` method is called manually. You should only set this argument to `False` if you intend to make other changes to the model instance and then save it manually.

Back to `File` and `ContentFile`: the one to use depends on what you want to store in `FileField`.

`File` is used as a wrapper around a Python `file` object, and you should use it if you have an existing `file` or file-like object that you want to save. File-like objects include `io.BytesIO` or `io.StringIO` instances. To instantiate a `File` instance, just pass the native file object to the constructor, for example:

```
f = open("/path/to/file.txt", "rb")
file_wrapper = File(f)
```

Use `ContentFile` when you already have some data loaded, either a `str` object or a `bytes` object. Pass the data to the `ContentFile` constructor:

```
string_content = ContentFile("A string value")
bytes_content = ContentFile(b"A bytes value")
```

Now that you have either a `File` instance or a `ContentFile` instance, saving the data to `FileField` is easy, using the `save` method:

```
m = ExampleModel.objects.first()
with open("/path/to/file.txt") as f:
    file_wrapper = File(f)
    m.file_field.save("file.txt", f)
```

Since we did not pass a value for `save` to the `save` method, it will default to `True`, so the model instance is automatically persisted to the database.

Next, we will look at how to store an image that has been manipulated with a PIL back to an image field.

Writing PIL images to ImageField

In *Exercise 8.05 – image uploads using Django forms*, you used PIL to resize an image and save it to disk. When working with a model, you might want to perform a similar operation but have Django handle the file storage using `ImageField` so that you don't have to do it manually. As in the exercise, you could save the image to disk and then use the `File` class to wrap the stored path – something like this:

```
image = Image.open(request.FILES["image_field"])
image.thumbnail((150, 150))
image.save("/tmp/thumbnaill.jpg") # save thumbnail to temp
                                location

with open("/tmp/thumbnaill.jpg", "rb") as f:
    image_wrapper = File(f)
    m.image_field.save("thumbnaill.jpg", image_wrapper)

os.unlink("/tmp/thumbnaill.jpg") # clean up temp file
```

In this example, we're having PIL store to a temporary location with the `Image.save()` method, and then re-opening the file.

This method works but is not ideal, as it involves writing the file to disk and then reading it out again, which can sometimes be slow. Instead, we can perform this whole process in memory.

Note

`io.BytesIO` and `io.StringIO` are useful objects. They behave like files but exist in memory only. `BytesIO` is used for storing raw bytes, and `StringIO` accepts Python 3's native Unicode strings. You can use `read`, `write`, and `seek`, just like a normal file. Unlike a normal file, though, they don't get written to disk and, instead, will disappear when your program terminates, or they go out of scope and are garbage collected. They are very useful if a function wants to write to something like a file, but you want to access the data immediately.

First, we will save the image data to an `io.BytesIO` object. Then, we'll wrap the `BytesIO` object in a `django.core.files.images.ImageField` instance (a subclass of `File` that is specifically for images and provides the `width` and `height` attributes). Once we have this `ImageFile` instance, we can use it in the `save` method of `ImageField`.

Note

`ImageFile` is a file or file-like wrapper just like `File`. It provides two extra attributes: `width` and `height`. `ImageFile` does not generate any errors if you use it to wrap a non-image. For example, you could use `open()` to open a text file and pass the file handle to the `ImageFile` constructor without issue. You can check whether the image file you passed in was valid by trying to access the `width` or `height` attributes: if these are `None`, then PIL was unable to decode the image data. You can check for the validity of these values yourself and throw an exception if they are `None`.

Let's have a look at this in practice in a view:

```
from io import BytesIO
from PIL import Image
from django.core.files.images import ImageFile

def index(request, pk):
    # trim out logic for checking if method is POST

    # get a model instance, or create a new one
    m = ExampleModel.objects.get(pk=pk)

    # store the uploaded image in a variable for shorter
```

```
code
uploaded_image = request.FILES["image_field"]

# load a PIL image instance from the uploaded file
image = Image.open(uploaded)

# perform the image resize
image.thumbnail((150, 150))

# Create a BytesIO file-like object to store
image_data = BytesIO()

# Write the Image data back out to the BytesIO object
# Retain the existing format from the uploaded image
image.save(fp=image_data, uploaded_image.format)

# Wrap the BytesIO containing the image data
image_file = ImageFile(image_data)

# Save the wrapped image file data with the original
# name
m.image_field.save(uploaded_image.name, image_file)
# this also saves the model instance
return redirect("/success-url/")
```

You can see this is a little bit more code, but it saves on writing the data to disk. You can choose to use either method (or another one that you come up with) depending on your needs.

Referring to media in templates

Once we have uploaded a file, we want to be able to refer to it in a template. For an uploaded image, such as a book cover, we will want to display the image on the page. We saw in *Exercise 8.02 – template settings, and using MEDIA_URL in templates*, how to build URL using MEDIA_URL in a template. When working with FileField or ImageField on a model instance, it is not necessary to do this, as Django provides this functionality for you.

The `url` attribute of FileField will automatically generate the full URL to the media file based on the `MEDIA_URL` in your settings.

Note

Note that references we make to FileField in this section also apply to ImageField, as it is a subclass of FileField.

This can be used anywhere that you have access to the instance and field, such as in a view or a template. The following example is in a view:

```
instance = ExampleModel.objects.first()
url = instance.file_field.url # Get the URL
```

Or in a template (assuming `instance` has been passed to the template context):

```

```

In the next exercise, we will create a new model with `FileField` and `ImageField` and then show how Django can automatically save these. We'll also demonstrate how to retrieve the URL for an uploaded file.

Exercise 8.06 – FileField and ImageField on models

In this exercise, we will create a model with `FileField` and `ImageField`. After doing this, we will have to generate a migration and apply it. We'll then change `UploadForm` we've been using, so it has both `FileField` and `ImageField`. The `media_example` view will be updated to store the uploaded files onto the model instance. Finally, we'll add `` into the example template to show the previously uploaded image:

1. In PyCharm, open the `media_example` app's `models.py` file. Create a new model called `ExampleModel`, with two fields: `ImageField` named `image_field`, and `FileField` called `file_field`. `ImageField` should have `upload_to` set to `images/`, and `FileField` should have `upload_to` set to `files/`. The finished model should look like this:

```
class ExampleModel(models.Model):
    image_field =
        models.ImageField(upload_to="images/")
    file_field = models.FileField(upload_to="files/")
```

Your `models.py` file should now look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.06/media_project/media_example/models.py.

2. Open a terminal and navigate to the `media_project` project directory. Make sure your `bookr` virtual environment is active. Run the `makemigrations` management command to generate the migrations for this new model (for Windows, you can use `python` instead of `python3` in the following code):

```
python3 manage.py makemigrations
```

Note

To learn how to create and activate a virtual environment, refer to the *Preface* section of the book.

The output should be like the following:

```
(bookr)$ python3 manage.py makemigrations
Migrations for 'media_example':
  media_example/migrations/0001_initial.py
    - Create model ExampleModel
```

3. Apply the migration by running the `migrate` management command:

```
python3 manage.py migrate
```

The output will be as follows:

```
(bookr)$ python3 manage.py migrate
Operations to perform:
  Apply all migrations: admin, auth, contenttypes,
  reviews, sessions
Running migrations:
  # output trimmed for brevity
  Applying media_example.0001_initial... OK
```

Note that all the initial Django migrations will also be applied since we did not apply those after creating the project.

4. Switch back to PyCharm and open the `media_example` app's `forms.py` file. Rename the existing `ImageField` from `file_upload` to `image_upload`. Then, add a new `FileField` named `file_upload`. After making these changes, your `UploadForm` code should look like this:

```
class UploadForm(forms.Form):
    image_upload = forms.ImageField()
    file_upload = forms.FileField()
```

You can save and close the file. It should look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.06/media_project/media_example/forms.py.

5. Open the `media_example` app's `views.py` file. First, import `ExampleModel` into the file. To do this, add this line at the top of the file after the existing `import` statements:

```
from .models import ExampleModel
```

Some imports will no longer be required, so you can remove these lines:

```
from PIL import Image
from django.conf import settings
```

6. In the `media_example` view, set a default for the instance you will render in case one is not created. After the function definition, define a variable called `instance` and set it to `None`:

```
def media_example(request):  
    instance = None
```

7. You can completely remove the contents of the `form.is_valid()` branch as you no longer need to manually save the file. Instead, it will automatically be saved when the `ExampleModel` instance is saved. You will instantiate an `ExampleModel` instance and set the file and image fields from the uploaded form.

Add this code under the `if form.is_valid():` line:

```
instance = ExampleModel()  
instance.image_field =  
    form.cleaned_data["image_upload"]  
instance.file_field = form.cleaned_data["file_upload"]  
instance.save()
```

8. Pass the instance through to the template in the context dictionary that is passed to `render`. Use the `instance` key:

```
return render(request, "media-example.html", {"form":  
    form, "instance": instance})
```

Now, your completed `media_example` view should look like this https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.06/media_project/media_example/views.py.

You can now save and close this file.

9. Open the `media-example.html` template. Add an `` element that displays the last uploaded image. Under the closing `</form>` tag, add an `if` template tag that checks whether `instance` has been provided. If so, display `` with a `src` attribute of `instance.image_field.url`:

```
{% if instance %}  
      
{% endif %}
```

You can save and close this file. It should now look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.06/media_project/media_example/templates/media-example.html.

10. Start the Django dev server if it's not already running, then navigate to `http://127.0.0.1:8000/media-example/`. You should see the form rendered with two fields:



A screenshot of a web browser window. The title bar says "Title". The address bar shows "127.0.0.1:8000/media-example/". The page content is a form with two fields: "Image upload" and "File upload", each with a "Browse..." button and the message "No file selected.". A "Submit" button is at the bottom.

Figure 8.20: UploadForm with two fields

11. Select a file for each field – for `ImageField`, you must select an image, but any type of file is allowed for `FileField`. See *Figure 8.20*, which shows the fields with files selected:



A screenshot of a web browser window, identical to Figure 8.20, but with files selected. The "Image upload" field now shows "cover.jpg" and the "File upload" field shows "sample.txt". The "Submit" button is at the bottom.

Figure 8.21: ImageField and FileField with files selected

Then, submit the form. If the submission is successful, the page will reload, and the last image we uploaded will be displayed (*Figure 8.22*):

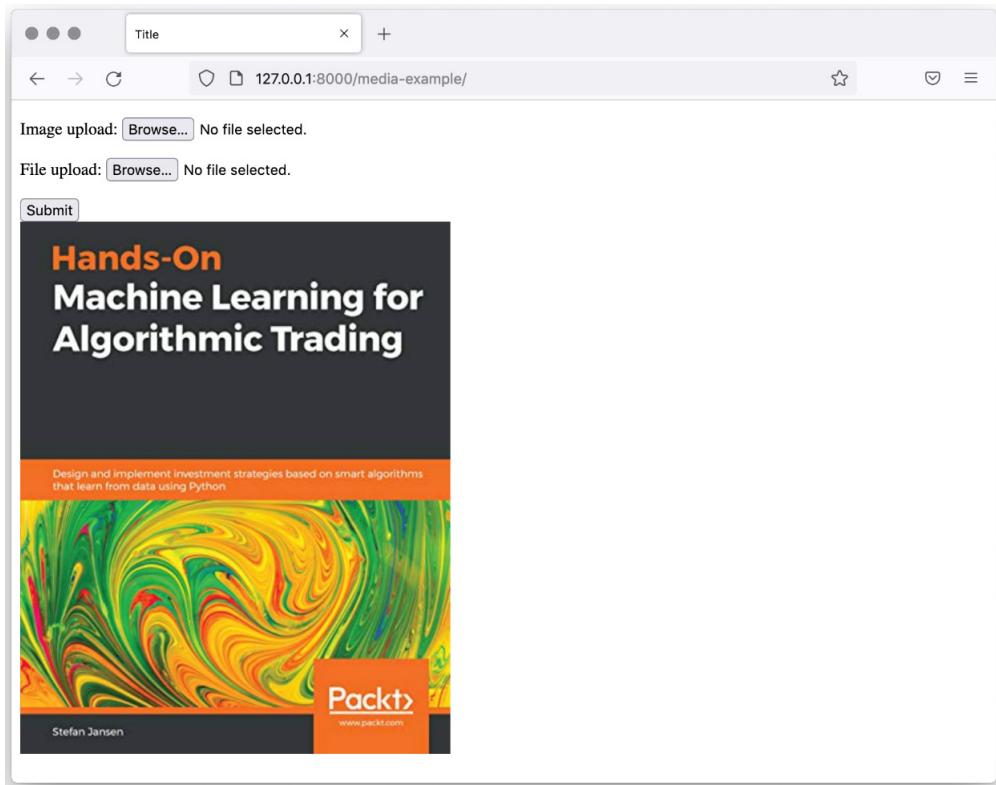


Figure 8.22: The last image that was uploaded is displayed

12. You can see how Django stores the files by looking in the `MEDIA_ROOT` directory. *Figure 8.23* shows the directory layout in PyCharm:

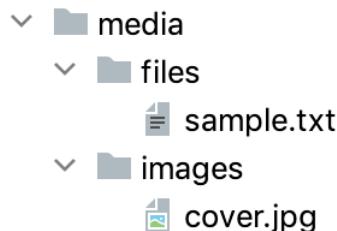


Figure 8.23: Uploaded files that Django has created

You can see that Django has created the `files` and `images` directories. These were what you set in the `upload_to` arguments on `ImageField` and `FileField` of the model.

We could also verify these uploads by attempting to download them, for example, at `http://127.0.0.1:8000/media/files/sample.txt` or `http://127.0.0.1:8000/media/images/cover.jpg`.

In this exercise, we created `ExampleModel` with `FileField` and `ImageField` and saw how to store uploaded files on it. We saw how to generate a URL to an uploaded file for use in a template. We tried uploading some files and saw that Django automatically created the `upload_to` directories (`media/files` and `media/images`) and then stored the files inside.

In the next section, we will look at how we can simplify the process even further by using `ModelForm` to generate the form and save the model without having to manually set the files in the view.

ModelForms and file uploads

We've seen how using `form.ImageField` on a form can prevent non-images from being uploaded. We've also seen how `models.ImageField` makes it easy to store an image for a model. But we need to be aware that Django does not stop you from setting a non-image file to `ImageField`. For example, consider a form that has both `FileField` and `ImageField`:

```
class ExampleForm(forms.Form):
    uploaded_file = forms.FileField()
    uploaded_image = forms.ImageField()
```

In the following view, the form would not validate if the `uploaded_image` field on the form was not an image, so some data validity is ensured for uploaded data. The following is an example of this:

```
def view(request):
    form = ExampleForm(request.POST, request.FILES)
    if form.is_valid():
        m = ExampleModel()
        m.file_field = form.cleaned_data["uploaded_file"]
        m.image_field =
            forms.cleaned_data["uploaded_image"]
        m.save()
    return render(request, "template.html")
```

Since we're sure the form is valid, we know that `forms.cleaned_data["uploaded_image"]` must contain an image. Therefore, we would never assign a non-image to the model instance's `image_field`.

However, what if we made a mistake in our code and wrote something like the following:

```
m.image field = forms.cleaned data["uploaded file"]
```

That is, if we accidentally reference `FileField` by mistake, Django does not validate that a (potential) non-image is being assigned to `ImageField`, and so it does not throw an exception or generate any kind of error. We can mitigate the potential for issues like this by using `ModelForm`.

We introduced `ModelForm` in *Chapter 7, Advanced Form Validation and Model Forms*. These are forms whose fields are automatically defined from a model. We saw that `ModelForm` has a `save` method that automatically creates or updates the model data in the database. When used with a model that has `FileField` or `ImageField`, then the `save` method of `ModelForm` will also save uploaded files.

Here's an example of using `ModelForm` to save a new model instance in a view. Here, we are just making sure to pass `request.FILES` to the `ModelForm` constructor:

```
class ExampleModelForm(forms.ModelForm):
    class Meta:
        model = ExampleModel # The same ExampleModel class
                            # we've seen previously
        fields = "__all__"

def view(request):
    if request.method == "POST":
        form = ExampleModelForm(request.POST,
                               request.FILES)
        form.save()
        return redirect("/success-page")
    else:
        form = ExampleModelForm()
    return (request, "template.html", {"form": form})
```

As with any `ModelForm`, the `save` method can be called with the `commit` argument set to `False`. Then the model instance will not be saved to the database, and the `FileField` and `ImageField` files will not be saved to disk. The `save` method should be called on the model instance itself – this will commit changes to the database and save the files. In this next short example, we set a value on the model instance before saving it:

```

        m.save()  # save the model instance, also write the
                  files to disk
        return redirect("/success-page/")
    else:
        form = ExampleModelForm()
    return (request, "template.html", {"form": form})

```

Calling the `save` method on the model instance both saves the model data to the database and the uploaded files to disk. In the next exercise, we will build `ModelForm` from `ExampleModel` we created in *Exercise 8.06 – FileField and ImageField on models*, then test uploading files with it.

Exercise 8.07 – file and image uploads using ModelForm

In this exercise, you will update `UploadForm` to be a subclass of `ModelForm` and have it built automatically from `ExampleModel`. You will then change the `media_example` view to save the instance automatically from the form, so you can see how the amount of code can be reduced:

1. In PyCharm, open the `media_example` apps' `forms.py` file. You need to use `ExampleModel` in this chapter, so use `import` to import it at the top of the file after `from django import forms` statement. Insert this line:

```
from .models import ExampleModel
```

2. Change `UploadForm` to be a subclass of `forms.ModelForm`. Remove the class body and replace it with a `class Meta` definition, its `model` should be `ExampleModel`. Set the `fields` attribute to `__all__`. After completing this step, `UploadForm` should look like this:

```

class UploadForm(forms.ModelForm):
    class Meta:
        model = ExampleModel
        fields = "__all__"

```

Save and close the file. It should now look like this: https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.07/media_project/media_example/forms.py.

3. Open the `media_example` app's `views.py` file. Since you no longer need to reference `ExampleModel` directly, you can remove its import at the top of the file. Remove the following line:

```
from .models import ExampleModel
```

4. In the `media_example` view, remove the entirety of the `form.is_valid()` branch and replace it with a single line:

```
instance = form.save()
```

The form's `save` method will handle persisting the instance to the database and saving the files. It will return an instance of `ExampleModel`, the same as the other `ModelForm` we have worked with in *Chapter 7, Advanced Form Validation and Model Forms*.

After completing this step, your `media_example` function should look like https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Exercise8.07/media_project/media_example/views.py. Save and close `views.py`.

5. Start the Django dev server if it's not already running, then navigate to `http://127.0.0.1:8000/media-example/`. You should see the form rendered with two fields, `Image field` and `File field` (*Figure 8.24*):

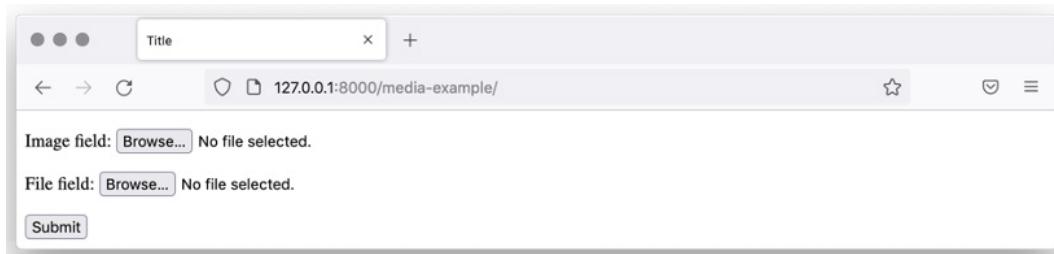


Figure 8.24: UploadForm as a ModelForm rendered in the browser

Note that the names of these fields now match those of the model rather than the form, since the form just uses the model's fields.

6. Browse and select an image and file (*Figure 8.25*), then submit the form:



Figure 8.25: Image and file selected

7. The page will reload, and as in *Exercise 8.06 – FileField and ImageField on models*, you will see the previously uploaded image (*Figure 8.26*):

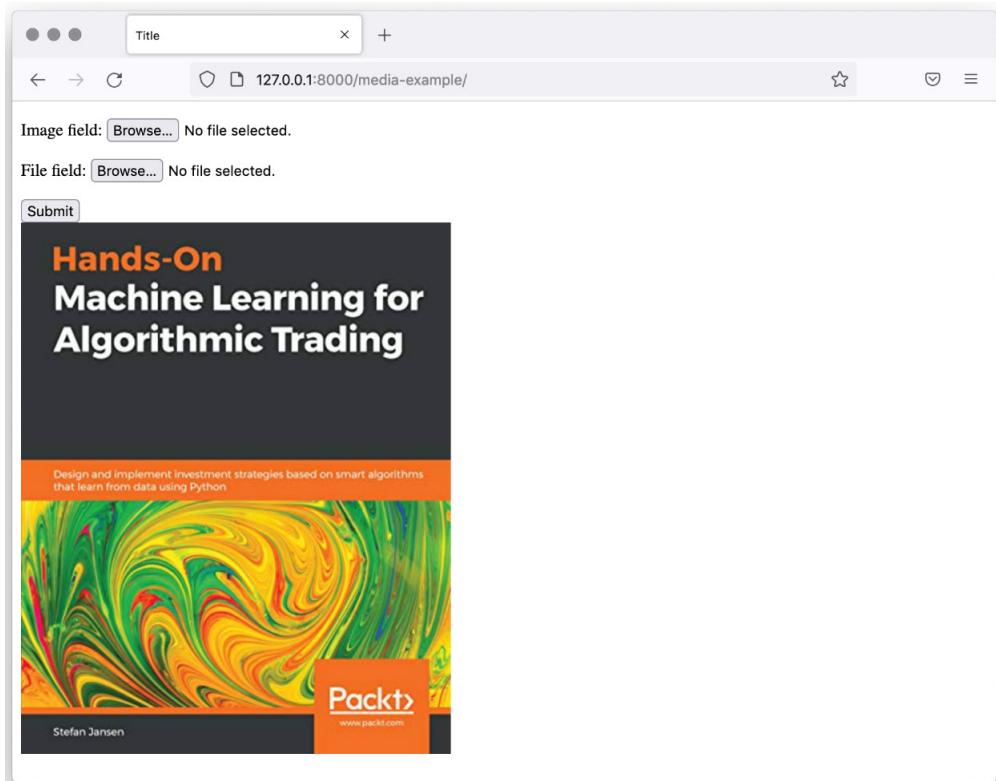


Figure 8.26: Image being displayed after upload

- Finally, examine the contents of the `media` directory. You should see the directory layout matches that of *Exercise 8.06 – FileField and ImageField on models*, with images inside the `images` directory and files inside the `files` directory:

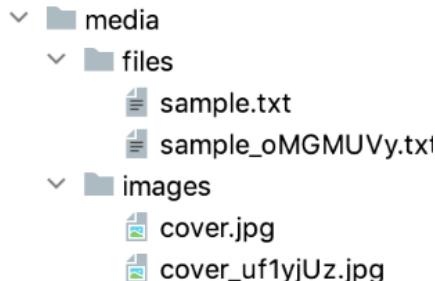


Figure 8.27: The uploaded files directory matches Exercise 6

In this exercise, we changed `UploadForm` to a `ModelForm` subclass, which allowed us to automatically generate the upload fields. We could replace the code that stored the uploaded files on the models with a call to the form's `save` method.

Before moving on to the activities for this chapter, we need to briefly discuss how to handle file saving when using an instance.

Handling file saving

As you know from working with `ModelForms` in *Chapter 7*, you can pass an `instance` argument to `ModelForm` to cause it to update an existing model instead of creating a new one when saving.

For example, do the following to update an example model in a view:

```
def update_view(request, pk):
    instance = ExampleModel.objects.get(pk=pk)
    form = UploadForm(request.POST, request.FILES,
                      instance=instance)
    if form.is_valid():
        form.save()
```

When accessing the `cleaned_data` attribute of the form, we will get back the data that was contained in `request.POST`. If there was no value for that key, then we'll get the value from the `instance` instead.

This normally doesn't cause any problems, but when working with `ImageField` or `FileField`, the type of data being returned will differ from the `request` or the `instance`.

If a file was uploaded, then the cleaned data will be `InMemoryUploadedFile` or `TemporaryUploadedFile` (depending on its size; only small files will fit in memory). If no file was uploaded, then the type will be `FieldFile` for `FileField` or `ImageFieldFile` for `ImageField`.

To explain, consider this code:

```
def update_view(request, pk):
    instance = ExampleModel.objects.get(pk=pk)
    form = UploadForm(request.POST, request.FILES,
                      instance=instance)
    if form.is_valid():
        form_file = form.cleaned_data["file_field"]
        form_image = form.cleaned_data["image_field"]
```

In the case where data was uploaded (that is, the user selected an image and file in the form), then `form_file` and `form_image` will be `InMemoryUploadedFile` or `TemporaryUploadedFile`. However, if the user did not upload anything for the fields (assuming the fields are not required in the form), then `form_file` will be a `FieldFile` instance, and `form_image` will be an `ImageFieldFile` instance.

To differentiate between them, there are a few different checks you could do. One of the simplest methods is to check for the existence of a `path` attribute. Since uploaded files haven't been saved to disk yet, they have no path:

```
if hasattr(form_file, "path"):
    print("This is a FieldFile")
else:
    print("This is an uploaded file")
```

Sometimes you will need to check what kind of file you are working with before performing some action, and perhaps skip that action if the user hadn't uploaded a new file. For example, we wouldn't want to resize an image that was already assigned to a model. We'd only do the image resize on newly uploaded images:

```
if hasattr(image_file, "path"):
    print("This image was on the instances so no need to
          resize")
else:
    perform_resize(image_file)
```

You will need to make a check like this in the first activity for this chapter.

We've now covered everything you need to start enhancing Bookr with file uploads. In the activity for this chapter, we'll add support for uploading a cover image and sample document (PDF, text file, and more) for a book. The book cover will be resized using PIL before it is saved.

Activity 8.01 – image and PDF upload of books

In this activity, you will start by cleaning up (deleting) the example views, templates, forms, models, and URL maps that we were using throughout the exercises in this chapter. You'll then need to generate and apply a migration to delete `ExampleModel` from the database.

You can then start adding the Bookr enhancements, first by adding `ImageField` and `FileField` to the `Book` model to store the Book cover and sample. Then you will create a migration and apply it to add these fields to the database. You can then build a form that will display just these new fields. You'll add a view that uses this form to save the model instance with the uploaded files after first resizing the image to a thumbnail size. You will be able to reuse the `instance-form.html` template from *Chapter 7, Advanced Form Validation and Model Forms*, with a minor change to allow file uploads.

These steps will help you complete the activity:

1. Update the Django settings to add the `MEDIA_ROOT` and `MEDIA_URL` settings.
2. The `/media/` URL mapping should be added to the main `urls.py` file. Use the `static` view and utilize `MEDIA_ROOT` and `MEDIA_URL` from Django settings. Remember, this mapping should only be added if `DEBUG` is `True`.
3. Add `ImageField` (named `cover`) and `FileField` (named `sample`) to the `Book` model. The fields should upload to `book_covers/` and `book_samples/`, respectively. They should both allow the `null` and `blank` values.
4. Run `makemigrations` and `migrate` again to apply the `Book` model changes to the database.
5. Create `BookMediaForm` as a subclass of `ModelForm`. Its model should be `Book`, and the fields should only be the fields you added in *step 3*.
6. Add a `book_media` view. This will not allow you to create `Book`, instead, it will only allow you to add media to an existing `Book` (so it must take `pk` as a required argument).
7. The `book_media` view should validate the form and use `save` to save it but not commit the instance. The uploaded cover should first be resized using the `thumbnail` method, as demonstrated in the *Writing PIL images to ImageField* section. The maximum size should be 300 px by 300 px. It should then be stored on the instance, and the instance should be saved. Remember that the `cover` field is not required, so you should check this before trying to manipulate the image, and you won't need to resize the image if one is already on the instance.

On a successful POST, register a success message that Book was updated, then redirect to the `book_detail` view.

8. Render `instance-form.html`, passing a context dictionary containing `form`, `model_type`, and `instance`, as you did in *Chapter 6, Forms*. Also, pass another item, `is_file_upload`, set to True. This variable will be used in the next step.
9. In the `instance-form.html` template, use the `is_file_upload` variable to add the correct `enctype` attribute to the form. This will allow you to switch the modes for the form to enable file uploads when required.
10. Finally, add a URL map that maps `/books/<pk>/media/` to the `book_media` view.

When you're finished, you should be able to start the Django dev server and load the `book_media` view at `http://127.0.0.1:8000/books/<pk>/media/`, for example, `http://127.0.0.1:8000/books/2/media/`. You should see `BookMediaForm` rendered in the browser, like in *Figure 8.28*:

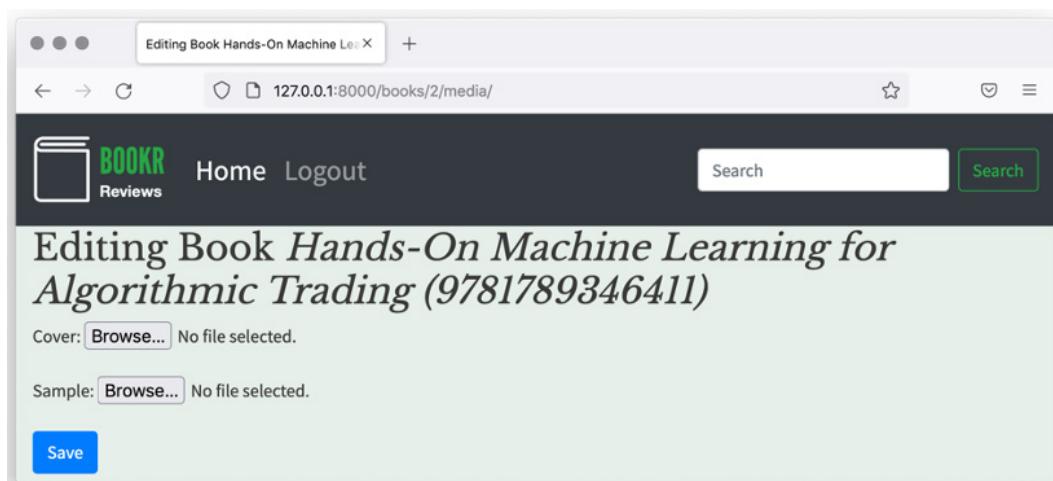


Figure 8.28: `BookMediaForm` in the browser

Select a cover image and a sample file for the book. You can use the image at <https://raw.githubusercontent.com/PacktPublishing/Web-Development-with-Django-Second-Edition/main/Chapter08/Activity8.01/bookr/fixtures/machine-learning-for-algorithmic-trading.png> and the PDF at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter08/Activity8.01/bookr/fixtures/machine-learning-for-trading.pdf> (or you can use any other image/PDF of your choosing).

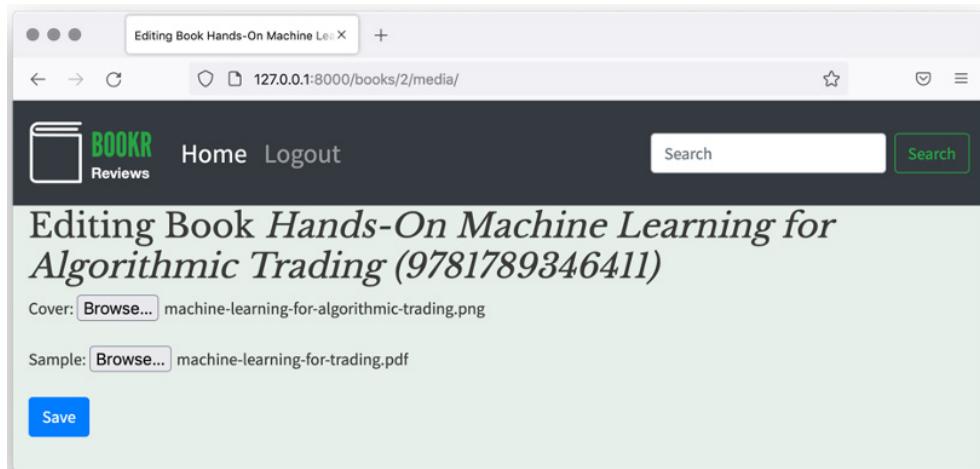


Figure 8.29: Book cover and sample selected

After submitting the form, you will be redirected to the Book Details view and see the success message (Figure 8.30):

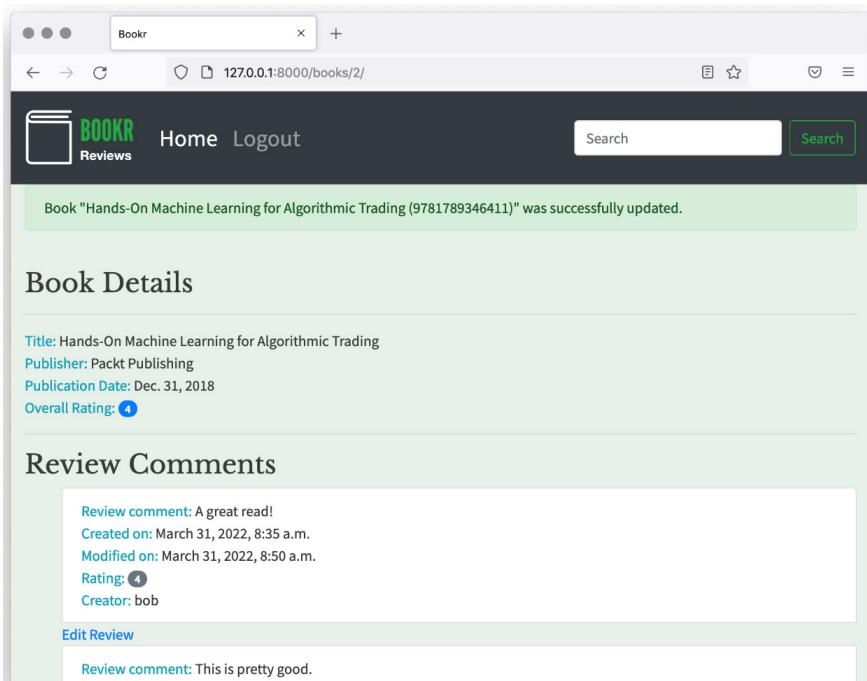


Figure 8.30: Success message on the Book Details page

If you go back to the same book's media page, you should see the fields are now filled in with an option to clear the data from them:

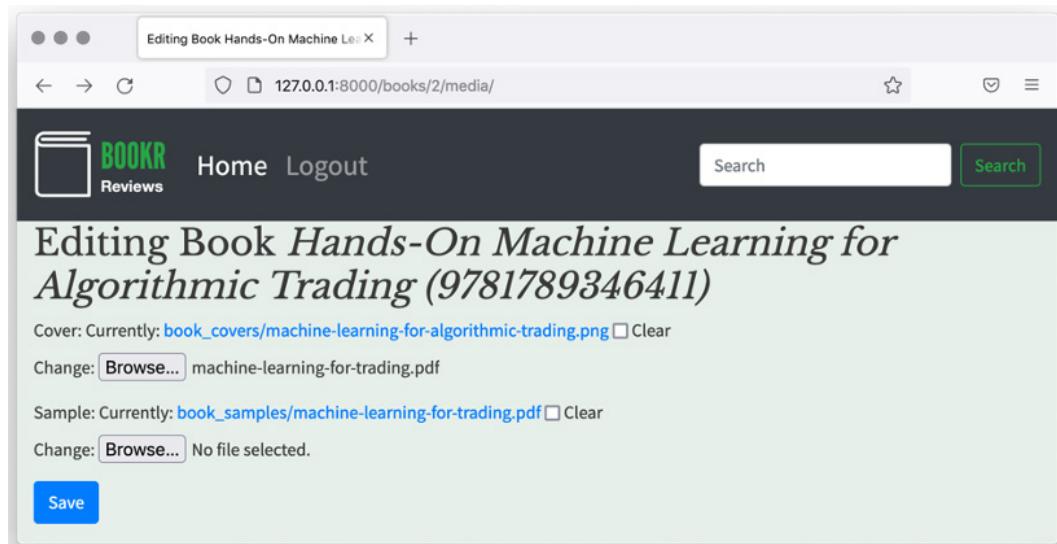


Figure 8.31: BookMediaForm with existing values

In *Activity 8.02 – displaying Cover and Sample Link*, you will add these uploaded files to the Book Detail view, but for now, if you want to check that uploads have worked, you can look inside the `media` directory in the Bookr project:

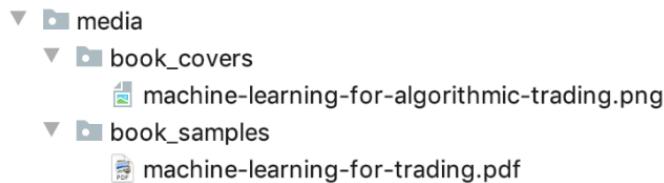


Figure 8.32: Book media

You should see the directories that were created and the uploaded files, as per *Figure 8.32*. Open an uploaded image, and you should see its maximum dimension is 300 px.

Note

The solution to this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Activity 8.02 – displaying the cover and sample link

In this activity, you will update the `book_detail.html` template to show the cover for Book (if one is set). You will also add a link to download the sample again, only if one is set. You will use the `FileField` and `ImageField` `url` attributes to generate the URLs to the media files.

These steps will help you complete this activity:

1. Inside the Book Details display in the `book_detail.html` view, add an `` element if the book has a `cover` image. Then, display the cover of the book inside it. Use `
` after the `` tag, so the image is on its own line.
2. After the *Publication Date* display, add a link to the sample file. It should only be displayed if a `sample` file has been uploaded. Make sure you add another `
` tag, so it displays correctly.
3. In the section that has a link to add a review, add another link that goes to the media page for the book. Follow the same styling as the **Add Review** link.

When you've completed these steps, you should be able to load a book's detail page. If the book has no `cover` or `sample`, then the page should look very similar to what it did before, except you should see the new link to the **Media** page at the bottom (*Figure 8.33*):

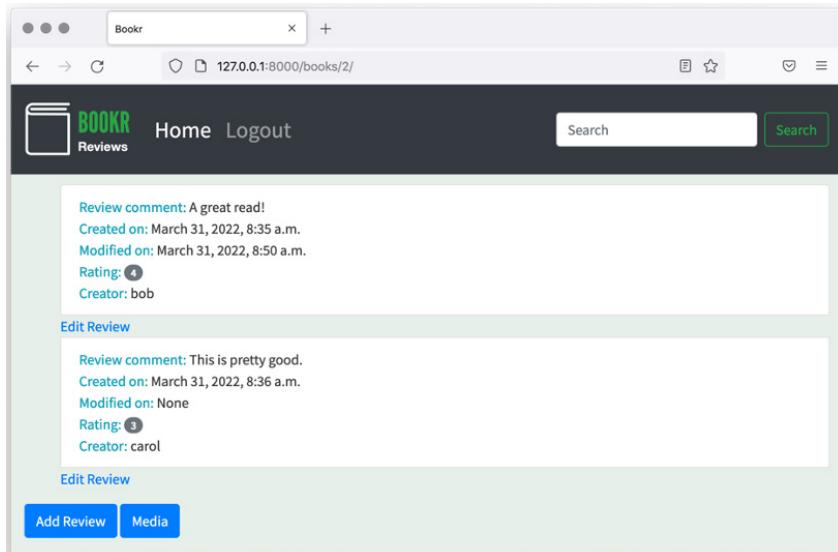


Figure 8.33: The new Media button visible on the Book Detail page

Once you have uploaded cover and/or sample for a book, the cover image and the sample link should be displayed (Figure 8.34):

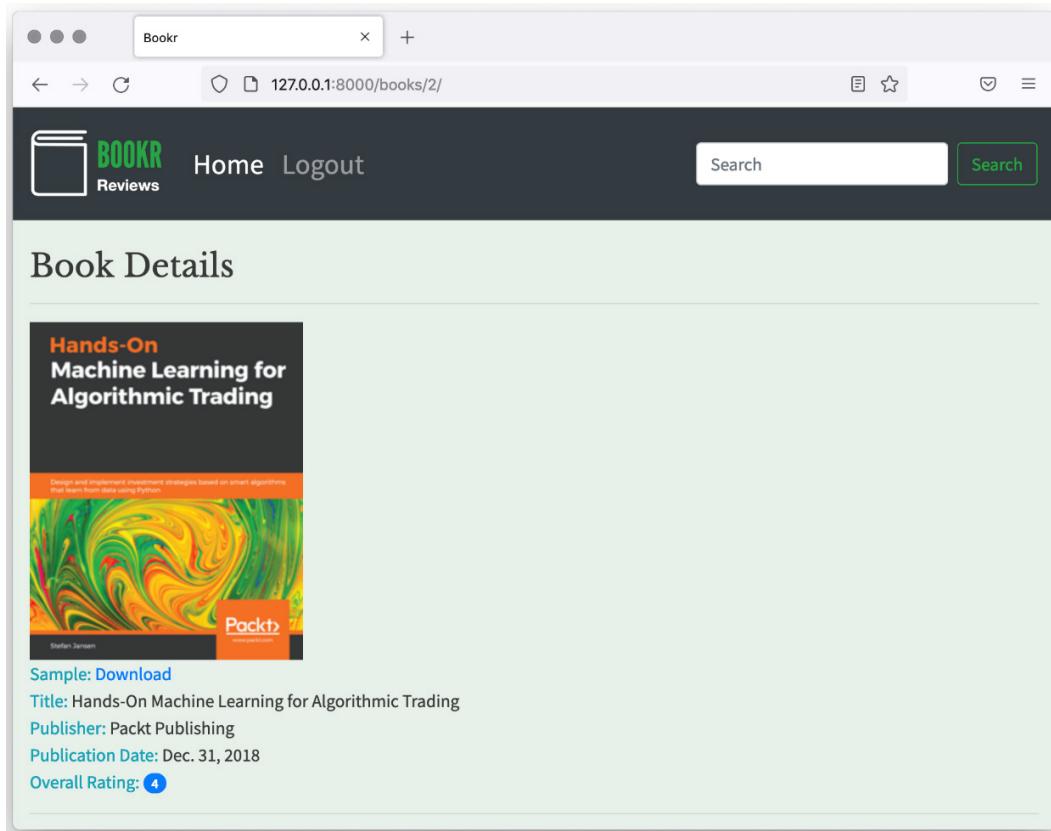


Figure 8.34: The book cover and the sample link displayed

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we added the `MEDIA_ROOT` and `MEDIA_URL` settings and a special URL map to serve media files. We then created a form and a view to upload files and save them to the `media` directory. We saw how to add the media context processor to automatically have access to the `MEDIA_URL` setting in all our templates. We then enhanced and simplified our form code by using a Django form with `FileField` or `ImageField`, instead of manually defining one in HTML.

We looked at some of the enhancements Django provides for images with `ImageField`, and how to interact with an image using Pillow. We showed an example view that would be able to serve files that required authentication using the `FileResponse` class. Then, we saw how to store files on models using `FileField` and `ImageField` and refer to them in a template using the `FileField.url` attribute. We were able to reduce the amount of code we had to write by automatically building `ModelForm` from a model instance. Finally, we did two activities to enhance Bookr by adding a cover image and a sample file to the `Book` model.

In *Chapter 9, Sessions and Authentication*, we will learn how to add authentication to a Django application to protect it from unauthorized users.

9

Sessions and Authentication

So far, we have used Django to develop dynamic applications that allow users to interact with application models, but we have not attempted to secure these applications from unwanted use. For example, our Bookr app allows unauthenticated users to add reviews and upload media. This is a critical security issue for any online web app as it leaves the site open to the posting of spam or other inappropriate material and the vandalism of existing content. We want the creation and modification of content to be strictly limited to authenticated users who have registered with the site.

The **authentication** app supplies Django with the models for representing users, groups, and permissions. It also provides middleware, utility functions, decorators, and mixins that help us integrate user authentication into our apps. Furthermore, the authentication app allows us to group and name certain sets of users.

In *Chapter 4, Introduction to Django Admin*, we used the Admin app to create a help desk user group with the Can view log entry, Can view permission, Can change user, and Can view user permissions. Those permissions could be referenced in our code using their corresponding codenames: `view_logentry`, `view_permissions`, `change_user`, and `view_user`. In this chapter, we will learn how to customize Django behavior based on specific user permissions.

Permissions are directives that delineate what is permissible by classes of users. Permissions can be assigned either to groups or directly to individual users. From an administrative point of view, it is cleaner to assign permissions to groups. Groups make it easier to model roles and organizational structures. If a new permission is created, it is less time-consuming to modify a few groups than to remember to assign it to a subset of users.

We are already familiar with creating users and groups and assigning permissions using several methods, such as the option of instantiating users and groups through the model using scripts and the convenience of creating them through the Django Admin app. The authentication app also offers us programmatic ways of creating and deleting users, groups, and permissions and assigning relationships between them.

As we go through this chapter, we'll learn how to use authentication and permissions to implement application security and how to store user-specific data to customize the user's experience. This will help us secure the `bookr` project from unauthorized content changes and make it contextually relevant for different types of users. Adding this basic security to our `bookr` project is crucial before we consider deploying it on the internet.

This chapter begins with a brief introduction to **middleware** before delving into the concepts of **authentication models** and **session engines**. You will implement Django's authentication model to restrict permissions to only a specific set of users. Then, you will learn how to leverage Django authentication to provide a flexible approach to application security. After that, you will learn how Django supports multiple session engines to retain user data. By the end of this chapter, you will be proficient at using sessions to retain information on past user interactions and to maintain user preferences for when pages are revisited.

We will cover the following topics in this chapter:

- Middleware
- Password storage in Django
- Authentication decorators and redirection
- Enhancing templates with authentication data
- Sessions

Authentication, as well as session management (which we'll learn about in the *Sessions* section), is handled by something known as a **middleware stack**. Before we implement authentication in our `bookr` project, let's learn a bit about this middleware stack and its modules.

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter09/>.

Middleware

In *Chapter 3, URL Mapping, Views, and Templates*, we discussed Django's implementation of the request/response process, along with its view and rendering functionality. In addition to these, another feature that plays an extremely important role when it comes to Django's core web processing is **middleware**. Django's middleware refers to a variety of software components that intervene in this request/response process to integrate important functionalities such as security, session management, and authentication.

So, when we write a view in Django, we don't have to explicitly set a series of important security features in the response header. These additions to the response object are automatically made by the `SecurityMiddleware` instance after the view returns its response. As middleware components wrap the view and perform a series of pre-processes on the request and post-processes on the response, the view is not cluttered with a lot of repetitive code and we can concentrate on coding application logic rather than worrying about low-level server behavior. Rather than building these functionalities into the Django core, Django's implementation of a middleware stack allows these components to be both optional and replaceable.

Middleware modules

When we run the `startproject` subcommand, a default list of middleware modules is added to the `MIDDLEWARE` variable in the `<project>/settings.py` file, as follows:

```
MIDDLEWARE = [  
    "django.middleware.security.SecurityMiddleware",  
    "django.contrib.sessions.middleware.SessionMiddleware",  
    "django.middleware.common.CommonMiddleware",  
    "django.middleware.csrf.CsrfViewMiddleware",  
    "django.contrib.auth.middleware.AuthenticationMiddleware",  
    "django.contrib.messages.middleware.MessageMiddleware",  
    "django.middleware.clickjacking.XframeOptionsMiddleware",  
]
```

This is a minimal middleware stack that is suitable for most Django applications. The following list elaborates on the general purpose of each module:

- `SecurityMiddleware` provides common security enhancements such as handling SSL redirects and adding response headers to prevent common hacks
- `SessionMiddleware` enables session support and seamlessly associates a stored session with the current request
- `CommonMiddleware` implements a lot of miscellaneous features, such as rejecting requests from the `DISALLOWED_USER_AGENTS` list, implementing URL rewrite rules, and setting the `Content-Length` header
- `CsrfViewMiddleware` adds protection against **Cross-Site Request Forgery (CSRF)**
- `AuthenticationMiddleware` adds the `user` attribute to the `request` object
- `MessageMiddleware` adds “flash” message support

- `XFrameOptionsMiddleware` protects against `X-Frame-Options` header clickjacking attacks

The middleware modules are loaded in the order that they appear in the `MIDDLEWARE` list. This makes sense because we want to call the middleware that deals with initial security issues first so that dangerous requests are rejected before further processing occurs. Django also comes with several other middleware modules that perform important functions, such as using `gzip` file compression, redirect configuration, and web cache configuration.

This chapter is devoted to discussing two important aspects of stateful application development that are implemented as middleware components – `SessionMiddleware` and `AuthenticationMiddleware`.

The `process_request` method of `SessionMiddleware` adds a `session` object as an attribute of the `request` object. The `process_request` method of `AuthenticationMiddleware` adds a `user` object as an attribute of the `request` object.

It is possible to write a Django project without these layers of the middleware stack if a project does not require user authentication or a means of preserving the state of individual interactions. However, most of the default middleware plays an important role in application security. If you don't have a good reason for changing the middleware components, it is best to maintain these initial settings. The Admin app requires `SessionMiddleware`, `AuthenticationMiddleware`, and `MessageMiddleware` to run, and the Django server will throw errors such as these if the Admin app is installed without them:

```
django.core.management.base.SystemCheckError: SystemCheckError: System
check identified some issues:

ERRORS:
?: (admin.E408) 'django.contrib.auth.middleware.
AuthenticationMiddleware' must be in MIDDLEWARE in order to use the
admin application.
?: (admin.E409) 'django.contrib.messages.middleware.MessageMiddleware'
must be in MIDDLEWARE in order to use the admin application.
?: (admin.E410) 'django.contrib.sessions.middleware.SessionMiddleware'
must be in MIDDLEWARE in order to use the admin application.
```

Now that we know about the middleware modules, let's look at one way we can enable authentication in our project using the authentication app's views and templates.

Implementing authentication views and templates

We encountered the login form on the Admin app in *Chapter 4, Introduction to Django Admin*. This is the authentication entry point for staff users who have access to the Admin app. We also need to create a login capability for ordinary users who want to give book reviews. Fortunately, the authentication app comes with the tools to make this possible.

As we work through the forms and views of the authentication app, we will encounter a lot of flexibility in its implementation. We are free to implement our own login pages, define either very simple or fine-grained security policies at the view level, and authenticate against external authorities.

The authentication app exists to accommodate a lot of different approaches to authentication so that Django doesn't rigidly enforce a single mechanism. For a first-time user encountering the documentation, this can be quite bewildering. For the most part in this chapter, we will follow Django's defaults, but some of the important configuration options will be noted.

A Django project's `settings` object contains attributes for login behavior. `LOGIN_URL` specifies the URL of the login page. `'/accounts/login/'` is the default value. `LOGIN_REDIRECT_URL` specifies the path to where a successful login is redirected. The default path is `'/accounts/profile/'`.

The authentication app supplies standard forms and views for carrying out typical authentication tasks. These forms are located in `django.contrib.auth.forms`, while the views are located in `django.contrib.auth.views`.

The views are referenced by these URL patterns present in `django.contrib.auth.urls`:

```
urlpatterns = [
    path('login/', views.LoginView.as_view(),
         name='login'),
    path('logout/', views.LogoutView.as_view(),
         name='logout'),
    path('password_change/',
         views.PasswordChangeView.as_view(),
         name='password_change'),
    path('password_change/done/',
         views.PasswordChangeDoneView.as_view(),
         name='password_change_done'),
    path('password_reset/',
         views.PasswordResetView.as_view(),
         name='password_reset'),
    path('password_reset/done/',
         views.PasswordResetDoneView.as_view(),
         name='password_reset_done'),
    path('reset/<uidb64>/<token>/',
         views.PasswordResetConfirmView.as_view(),
```

```
        name='password_reset_confirm'),
    path('reset/done/',
        views.PasswordResetCompleteView.as_view(),
        name='password_reset_complete'),
]
```

If this style of views looks unfamiliar, it is because they are class-based views rather than the function-based views that we have previously encountered. We will learn more about class-based views in *Chapter 11, Advanced Templating and Class-Based Views*. For now, note that the authentication app makes use of class inheritance to group the functionality of views and prevent a lot of repetitive coding.

If we want to maintain the default URLs and views that are presupposed by the authentication app and Django settings, we can include the authentication app's URLs in our project's `urlpatterns`.

By taking this approach, we have saved a lot of work. We only need to include the authentication app's URLs in our `<project>/urls.py` file and assign it to the 'accounts' namespace. Designating this namespace ensures that our reverse URLs correspond to the default template values of the views. The authlib views contain code references to templates named `password_reset_done` and `password_reset_complete`, so their paths need to be explicitly included in our `urlpatterns` too:

```
from django.contrib import admin, auth
...
urlpatterns = [
    path("admin/", admin.site.urls),
    path("accounts/", include(("django.contrib.auth.urls",
        "auth"),
        namespace="accounts")),
    path("accounts/password_reset/done/",
        auth.views.PasswordResetDoneView.as_view(),
        name="password_reset_done",),
    path("accounts/reset/done/",
        auth.views.PasswordResetCompleteView.as_view(),
        name="password_reset_complete",),
    path("", reviews.views.index),
    path("book-search/", reviews.views.book_search,
        name="book_search"),
    path("", include("reviews.urls")),
]
```

Though the authentication app comes with forms and views, it lacks the templates needed to render these components as HTML. *Figure 9.1* lists the templates that we require to implement the authentication functionality in our project. Fortunately, the Admin app does implement a set of templates that we can utilize for our purposes.

We could just copy the template files from the Django source code in the `django/contrib/admin/templates/registration` directory and `django/contrib/admin/templates/admin/login.html` into our project's `templates/registration` directory:

Note

When we say Django source code, it's the directory where your Django installation resides. If you installed Django in a virtual environment (as detailed in the *Preface*), you can find these template files at `<name of your virtual environment>/lib/python3.X/site-packages/django/contrib/admin/templates/registration/`. Provided your virtual environment is activated and Django is installed in it, you can also retrieve the complete path to the `site-packages` directory by running the `python -c "import sys; print(sys.path)"` command in a terminal.

Template Path

<code>templates/registration/login.html</code>
<code>templates/registration/password_reset_email.html</code>
<code>templates/registration/password_change_form.html</code>
<code>templates/registration/password_change_done.html</code>
<code>templates/registration/password_reset_form.html</code>
<code>templates/registration/password_reset_done.html</code>
<code>templates/registration/password_reset_confirm.html</code>
<code>templates/registration/password_reset_complete.html</code>
<code>templates/registration/logged_out.html</code>

Figure 9.1 – Default paths for authentication templates

Note

We only need to copy the templates that are dependencies for the views and should avoid copying the `base.html` or `base_site.html` files.

This gives a promising result at first, but as they stand, the admin templates do not meet our precise needs, as we can see from the login page (*Figure 9.2*):

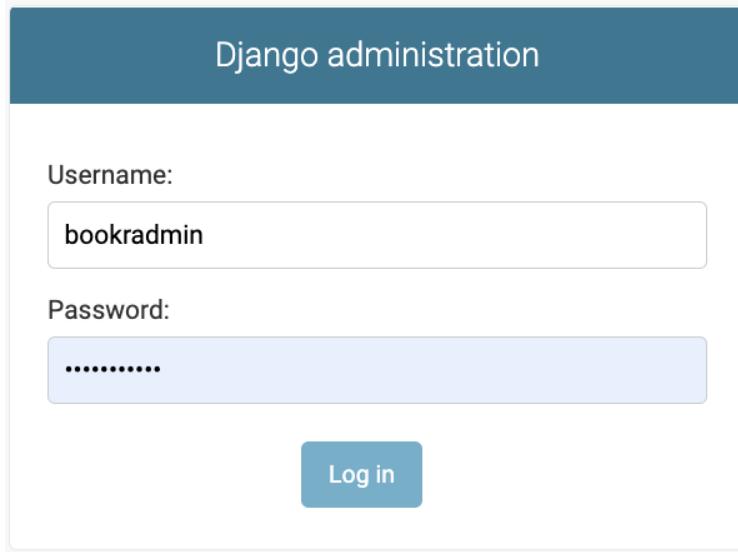


Figure 9.2 – A first attempt at a user login screen

As these authentication pages inherit from the Admin app's `admin/base_site.html` template, they follow the style of the Admin app. We would prefer for these pages to follow the style of the `bookr` project that we have developed. We can do this by following these three steps on each Django template that we have copied from the Admin app to our project:

1. The first change that needs to be made is to replace the `{% extends "admin/base_site.html" %}` tag with `{% extends "base.html" %}`.
2. Given that `template/base.html` only contains the `title`, `brand`, and `content` block definitions, we should examine the additional block substitutions from our templates in the `bookr` folder. We are not using the content from the `userlinks` and `breadcrumbs` blocks in our app's template layouts, so these blocks can be removed entirely.

Some of these blocks, such as `content_title` and `reset_link`, contain HTML content that is relevant to our application, so they should be retained.

For example, the `password_change_done.html` template contains the `userlinks` and `breadcrumbs` blocks:

```
{% extends "admin/base_site.html" %}
{% load i18n %}
{% block userlinks %}{% url 'django-admindocs-docroot' as
docsroot %}{% if docsroot %}<a href="{{ docsroot }}">{%
translate 'Documentation' %}</a> / {% endif %}{%
translate 'Change password' %} / <a href="{% url 'admin:logout' %}">{%
translate 'Log out' %}</a>{% endblock %}
{% block breadcrumbs %}
<div class="breadcrumbs">
<a href="{% url 'admin:index' %}">{%
translate 'Home' %}</a>
&rsaquo; {%
translate 'Password change' %}
</div>
{% endblock %}

{% block content %}
<p>{%
translate 'Your password was changed.' %}</p>
{% endblock %}
```

It will be simplified to this template in the `bookr` project:

```
{% extends "base.html" %}
{% load i18n %}

{% block content %}
<p>{%
translate 'Your password was changed.' %}</p>
{% endblock %}
```

3. Likewise, there are reverse URL patterns that need to change to reflect the current path, so `{% url 'login' %}` gets replaced with `{% url 'accounts:login' %}`.

Given these considerations, the next exercise will focus on transforming the Admin app's login template into a login template for the `bookr` project.

Note

The `i18n` module is used for creating multilingual content. If you intend to develop multilingual content for your website, leave the `i18n` import, `translate` tags, and `translateblock` statements in the templates. For brevity, we will not be covering those in detail in this chapter.

Exercise 9.01 – repurposing the Admin app login template

We started this chapter without a login page for our project. By adding the URL patterns for authentication and copying the templates from the Admin app to our project, we can implement the functionality of a login page. But this login page is not satisfactory as it is directly copied from the Admin app and is disconnected from the Bookr design. In this exercise, we will follow the steps needed to repurpose the Admin app's login template for our project. The new login template will need to inherit its style and format directly from the bookr project's `templates/base.html`:

1. Create a directory inside your project for `templates/registration`.

The Admin login template is located in the Django source directory at the `django/contrib/admin/templates/admin/login.html` path. It begins with an `extends` tag, a `load` tag, importing the `i18n` and `static` modules, and a series of block extensions that override the blocks defined in the child template, `django/contrib/admin/templates/admin/base.html`. A truncated snippet of the `login.html` file is shown in the following code block:

```
{% extends "admin/base_site.html" %}
{% load i18n static %}

{% block extrastyle %}{{ block.super }}...
{% endblock %}

{% block bodyclass %}{{ block.super }} login{% endblock %}
{% block user-tools %}{% endblock %}
{% block nav-global %}{% endblock %}
{% block content_title %}{% endblock %}
{% block breadcrumbs %}{% endblock %}
```

2. Copy this Admin login template, `django/contrib/admin/templates/admin/login.html`, into `templates/registration` and begin editing the file using PyCharm.
3. As the login template you are editing is located at `templates/registration/login.html` and extends the base template (`templates/base.html`), replace the argument of the `extends` tag at the top of `templates/registration/login.html` with the following:

```
{% extends "base.html" %}
```

4. We don't need most of the contents of this file. Just retain the `content` block, which contains the login form. The remainder of the template will consist of loading the `i18n` and `static` tag libraries:

```
{% load i18n static %}

{% block content %}
...
{% endblock %}
```

5. Now, you must replace the paths and reverse URL patterns in `templates/registration/login.html` with ones that are appropriate to your project. As you don't have an `app_path` variable defined in your template, it needs to be replaced with the reverse URL for the login, `'accounts:login'`. So, consider the following line:

```
<form action="{{ app_path }}" method="post" id="login-form">
```

This line changes as follows:

```
<form action="{% url 'accounts:login' %}" method="post" id="login-form">
```

No `'admin_password_reset'` has been defined in your project paths, so it will be replaced with `'accounts:password_reset'`.

Consider the following line:

```
{% url 'admin_password_reset' as password_reset_url %}
```

This changes to the following line:

```
{% url 'accounts:password_reset' as password_reset_url %}
```

Your login template will look as follows:

templates/registration/login.html

```
{% extends "base.html" %}  
{% load i18n static %}  
  
{% block content %}  
{% if form.errors and not form.non_field_errors %}  
<p class="errornote">  
{% if form.errors.items|length == 1 %}{% translate  
"Please correct the error  
below." %}{% else %}{% translate "Please correct  
the errors below." %}{% endif %}  
</p>  
{% endif %}
```

You can find the complete code for this file at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter09/Exercise9.01/bookr>.

6. As mentioned previously, to use the standard Django authentication views, we must add URL mappings to them. Open the `urls.py` file in the `bookr` project directory, ensure that `auth` is imported from `django.contrib`, and add the three paths to the URL patterns:

```
from django.contrib import admin, auth
...
urlpatterns = [
    ...
    path("accounts/", include(("django.contrib.auth.urls", "auth"),
        namespace="accounts")),
    path("accounts/password_reset/done/", auth.views.PasswordResetDoneView.as_view(),
        name="password_reset_done",),
    path("accounts/reset/done/", auth.views.PasswordResetCompleteView.as_view(),
        name="password_reset_complete",),
    ...
]
```

7. Now, when you visit the login link at `http://127.0.0.1:8000/accounts/login/`, you will see this page:

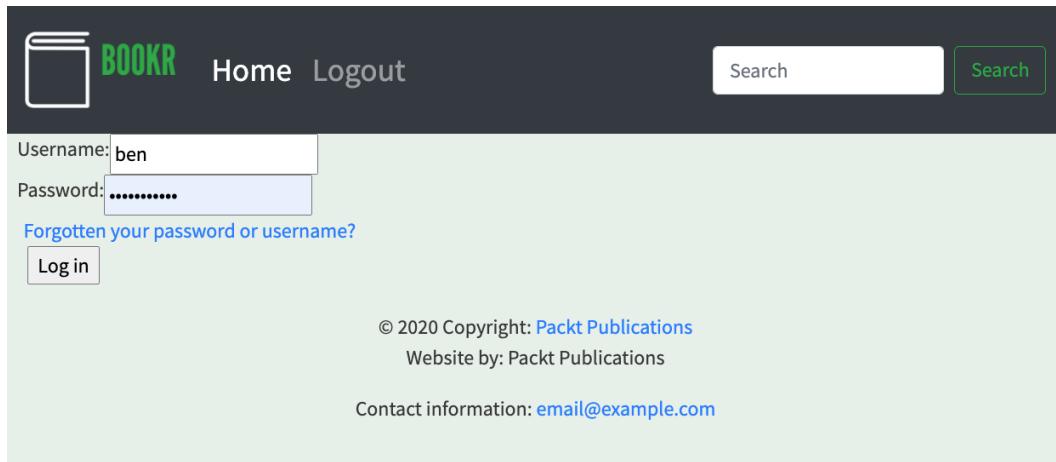


Figure 9.3 – The Bookr login page

By completing this exercise, you have created the template required for non-admin authentication in your project.

Note

Before you proceed, you'll need to make sure the rest of the templates in the `registration` directory follow the `bookr` project's style; that is, they inherit from the Admin app's `admin/base_site.html` template. You've already seen this done with `password_change_done.html` and the `login.html` templates. Go ahead and apply what you've learned in this exercise (and the section before it) to the rest of the files in the `registration` directory. Alternatively, you may download the modified files from this book's GitHub repository: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter09/Activity9.01/bookr>.

Now that we have modified our project to enable authentication, we will learn about password and session storage in Django.

Password storage in Django

Django does not store passwords in plain text form in the database. Instead, passwords are digested with a hashing algorithm, such as **PBKDF2/SHA256**, **BCrypt/SHA256**, or **Argon2**. As hashing algorithms are a one-way transformation, this prevents a user's password from being decrypted from the hash stored in the database. This often comes as a surprise to users who expect a system administrator to retrieve their forgotten password, but it is best practice in security design. So, if we query the database for the password, we will see something like this:

```
sqlite> select password from auth_user;
pbkdf2_sha256$320000$auohDtoK58a6LAoaG79Tvl$Mxzssu6gfd5uZvI1lADFI4faAjP0g3puP2JxxFdWZI0=
pbkdf2_sha256$320000$qoVFBnhgo4zgw4juRUNtAK$TATQEeqXtbRfc4Ccuuwr9nRHVuShZo5W6N90ygpvVbA=
pbkdf2_sha256$320000$s3eprgniu7e4q0ux9Xj5nh$/SAFKgXuYIWSm08PWEwBUsMfmPCYd60rAaz0JZnLKGs=
```

Figure 9.4 – Password hashes in a Django SQLite3 database

The components of this string are `<algorithm>$<iterations>$<salt>$<hash>`. As several hashing algorithms have been compromised over time and we sometimes need to work with mandated security requirements, Django is flexible enough to accommodate new algorithms and can maintain data encrypted in multiple algorithms.

The profile page and the `request.user` object

When a login is successful, the login view redirects to `/accounts/profile`. However, this path is not included in the existing `auth.url`, nor does the authentication app provide a template for it. To avoid a `Page not Found` error, a view and an appropriate URL pattern are required.

Each Django request has a `request.user` object. If the request is made by an unauthenticated user, `request.user` will be an `AnonymousUser` object. If the request is made by an authenticated user, then `request.user` will be a `User` object. This makes it easy to retrieve personalized user information in a Django view and render it in a template.

In the next exercise, we will add a profile page to our `bookr` project.

Exercise 9.02 – adding a profile page

In this exercise, we will add a profile page to our project. To do so, we need to include the path to it in our URL patterns and also include it in our views and templates. The profile page will simply display the following attributes from the `request.user` object:

- `username`
- `first_name` and `last_name`
- `date_joined`
- `email`
- `last_login`

Perform the following steps to complete this exercise:

1. Add `bookr/views.py` to the project. It needs a trivial profile function to define our view:

```
from django.shortcuts import render

def profile(request):
    return render(request, 'profile.html')
```

2. In the `templates` folder of your main `bookr` project, create a new file called `profile.html`. In this template, the attributes of the `request.user` object can easily be referenced by using a notation such as `{{ request.user.username }}`:

```
{% extends "base.html" %}

{% block title %}Bookr{% endblock %}

{% block content %}
<h2>Profile</h2>
<div>
    <p>
        Username: {{ request.user.username }} <br>
        Name: {{ request.user.first_name }} {{
```

```
    request.user.last_name }}<br>
    Date Joined: {{ request.user.date_joined }} <br>
    Email: {{ request.user.email }}<br>
    Last Login: {{ request.user.last_login }}<br>
  </p>
</div>
{%- endblock %}
```

We also added a block containing the profile details of the user. More importantly, we made sure that `profile.html` extends `base.html`.

3. Finally, this path needs to be added to the top of the `urlpatterns` list in `bookr/urls.py`. First, import the new views and then add a path linking the `accounts/profile/` URL to `bookr.views.profile`:

```
from bookr.views import profile

urlpatterns = [...
    path('accounts/profile/', profile,
         name='profile'),
    path("", reviews.views.index),
    path("book-search/",
         reviews.views.book_search,
         name="book_search"),
    path('', include('reviews.urls'))]
```

This is a good start on a user profile page. When Alice logs in and visits `http://localhost:8000/accounts/profile/`, it is rendered as shown in *Figure 9.5*. Remember, if the server needs to be started, use the `python manage.py runserver` command:

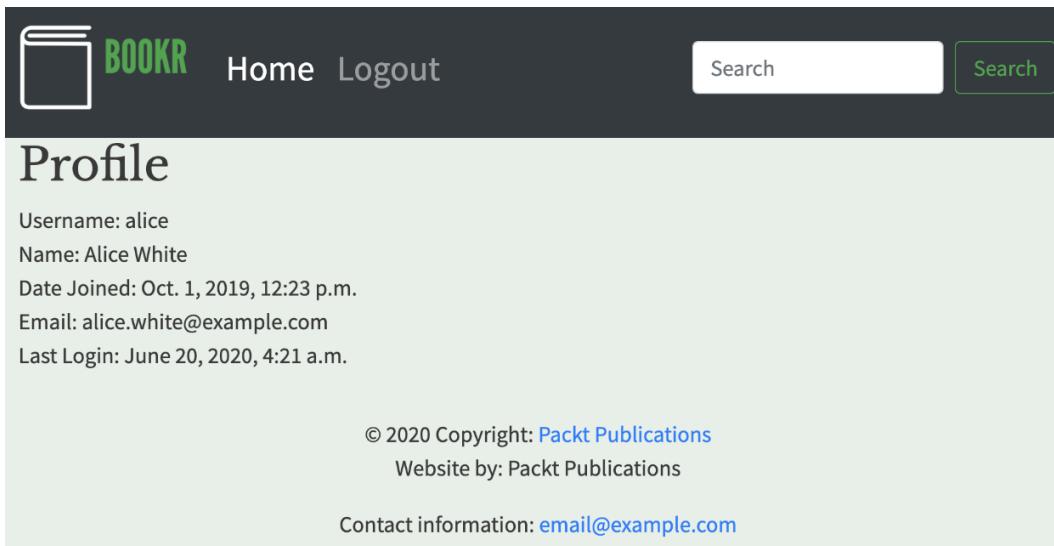


Figure 9:5 – Alice visits her user profile

With that, we've seen how we can redirect a user to their profile page, once they've successfully logged in. Now, let's discuss how we can give content access to specific users only.

Authentication decorators and redirection

Now that we have learned how to allow ordinary users to log in to our project, we can discover how to restrict content to authenticated users. The authentication module comes with some useful decorators that can be used to secure views according to the current user's authentication or access.

Unfortunately, if, say, a user named Alice was to log out of Bookr, the profile page would still render and display empty details. Instead of this happening, it would be preferable for any unauthenticated visitor to be directed to the login screen:

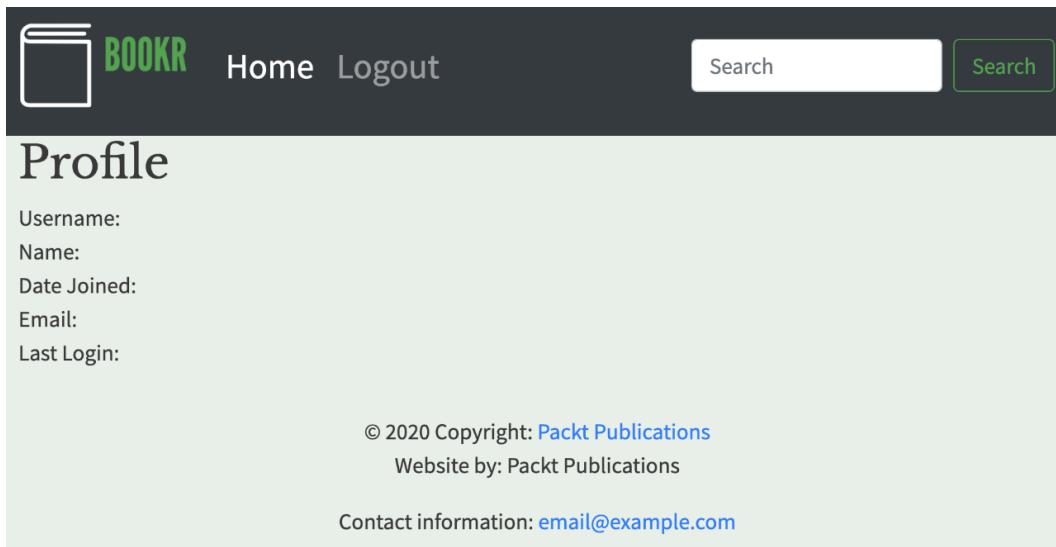


Figure 9.6 – An unauthenticated user visits a user profile

The authentication app comes with useful decorators for adding authentication behavior to Django views. In this situation of securing our profile view, we can use the `login_required` decorator:

```
from django.contrib.auth.decorators import login_required

@login_required
def profile(request):
    ...
```

Now, if an unauthenticated user visits the `/accounts/profile` URL, they will be redirected to `http://localhost:8000/accounts/login/?next=/accounts/profile/`.

This URL takes the user to the login URL. The `next` parameter in the GET variables tells the login view where to redirect to after a successful login. The default behavior is to redirect back to the current view, but this can be overridden by specifying the `login_url` argument to the `login_required` decorator. For example, if we had some need to redirect to a different page after login, we could have explicitly stated it in the decorator call like this:

```
@login_required(login_url='/accounts/profile2')
```

If we had rewritten our login view to expect the redirection URL to be specified in a different URL argument to 'next', we could explicate this in the decorator call with the `redirect_field_name` argument:

```
@login_required(redirect_field_argument='redirect_to')
```

There are often situations where a URL should be restricted to users or groups holding a specific condition. Consider the case where we have a page for staff users to view any user profile. We don't want this URL to be accessible to all users, so we want to limit this URL to users or groups with the 'view_user' permission and forward the unauthorized requests to the login URL:

```
from django.contrib.auth.decorators import (login_required,
                                             permission_required)
...
@permission_required('view_group')
def user_profile(request, uid):
    user = get_object_or_404(User, id=uid)
    permissions = user.get_all_permissions()
    return render(request, 'user_profile.html',
                  {'user': user, 'permissions': permissions})
```

Som with this decorator applied to our `user_profile` view, an unauthorized user visiting `http://localhost:8000/accounts/users/123/profile/` would be redirected to `http://localhost:8000/accounts/login/?next=/accounts/users/123/profile/`.

Sometimes, though, we need to structure more subtle conditional permissions that don't fall into the scope of these two directors. For this purpose, Django provides a custom decorator that takes an arbitrary function as an argument. The `user_passes_test` decorator requires a `test_func` argument:

```
user_passes_test(test_func, login_url=None,
                  redirect_field_name='next')
```

Here's an example where we have a view, `veteran_features`, that is only available to users who have been registered on the site for more than a year:

```
from django.contrib.auth.decorators import (login_required,
                                             permission_required, user_passes_test)
...
def veteran_user(user):
    now = datetime.datetime.now()
    if user.date_joined is None:
        return False
    return now - user.date_joined > datetime.timedelta(days=365)
```

```
@user_passes_test(veteran_user)
def veteran_features(request):
    user = request.user
    permissions = user.get_all_permissions()
    return render(request, 'veteran_profile.html',
                  {'user': user, 'permissions': permissions})
```

Sometimes, the logic in our views cannot be handled with one of these decorators and we need to apply the redirect within the control flow of the view. We can do this using the `redirect_to_login` helper function. It takes the same arguments as the decorators, as shown in the following snippet:

```
redirect_to_login(next, login_url=None,
                  redirect_field_name='next')
```

We will consolidate our knowledge of authentication decorators by applying them to the Bookr project in the following exercise.

Exercise 9.03 – adding authentication decorators to the views

Having learned about the flexibility of the authentication app's permission and authentication decorators, we will now set about putting them to use in the Reviews app. We need to ensure that only authenticated users can edit reviews and that only staff users can edit publishers. There are several ways of doing this, so we will attempt a few approaches. All the code in these steps is in the `reviews/views.py` file:

1. Your first instinct to solve this problem would be to think that the `publisher_edit` method needs an appropriate decorator to enforce that the user has `edit_publisher` permission. For this, you could easily do something like this:

```
from django.contrib.auth.decorators import
    permission_required
...
@permission_required('edit_publisher')
def publisher_edit(request, pk=None):
    ...
```

- Using this method is fine and it's one way to add permissions checking to a view. You can also use a slightly more complicated but more flexible method. Instead of using a permission decorator to enforce permission rights on the `publisher_edit` method, you can create a test function that requires a staff user and apply this test function to `publisher_edit` with the `user_passes_test` decorator. Writing a test function allows more customization on how you validate users' access rights or permissions. If you made changes to your `views.py` file in *step 1*, feel free to comment the decorator out (or delete it) and write the following test function instead:

```
from django.contrib.auth.decorators import
user_passes_test
...
def is_staff_user(user):
    return user.is_staff
...
@user_passes_test(is_staff_user)
...
```

- Ensure that login is required for the `review_edit` and `book_media` functions by adding the appropriate decorator:

```
...
from django.contrib.auth.decorators import
(login_required, user_passes_test)
...
@login_required
def review_edit(request, book_pk, review_pk=None):
    ...
@login_required
def book_media(request, pk):
    ...
```

- In the `review_edit` method, add logic to the view that requires that the user be either a staff user or the owner of the review. The `review_edit` view controls the behavior of both review creation and review updates. The constraint that we are developing only applies to the case where an existing review is being updated. So, the place to add code is after a `Review` object has been successfully retrieved. If the user is not a staff account or the review's creator doesn't match the current user, we need to raise a `PermissionDenied` error:

```
...
from django.core.exceptions import PermissionDenied
from PIL import Image
from django.contrib import messages
...
@login_required
def review_edit(request, book_pk, review_pk=None):
    book = get_object_or_404(Book, pk=book_pk)

    if review_pk is not None:
        review = get_object_or_404(Review,
                                   book_id=book_pk,
                                   pk=review_pk)
        user = request.user
        if not user.is_staff and review.creator.id != user.id:
            raise PermissionDenied
    else:
        review = None
...

```

Now, when a non-staff user attempts to edit another user's review, a `Forbidden` error will be thrown, as in *Figure 9.7*. In the next section, we will look at applying conditional logic in templates so that users aren't taken to pages that they don't have sufficient permission to access:



403 Forbidden

Figure 9.7 – Access is forbidden to non-staff users

In this exercise, we used authentication decorators to secure views in a Django app. The authentication decorators that were applied provided a simple mechanism to restrict views from users lacking necessary permissions, non-staff users, and unauthenticated users. Django's authentication decorators provide a robust mechanism that follows Django's role and permission framework, while the `user_passes_test` decorator provides an option to develop custom authentication.

Now that we have used decorators to control the authentication flow of the application, we can customize templates with user authentication data.

Enhancing templates with authentication data

In *Exercise 9.02 – adding a profile page*, we saw that we can pass the `request.user` object to the template to render the current user's attributes in the HTML. We can also take the approach of giving different template renderings according to the user type or permissions held by a user. Consider that we want to add an edit link that only appears to staff users. We might apply an `if` condition to achieve this:

```
{% if user.is_staff %}  
  <p><a href="{% url 'review:edit' %}">Edit this Review</a>  
  </p>  
{% endif %}
```

If we didn't take the time to conditionally render links based on permissions, users would have a frustrating experience navigating the application as many of the links that they click on would lead to 403 Forbidden pages. The following exercise will show how we can use templates and authentication to present contextually appropriate links in our project.

Exercise 9.04 – toggling login and logout links in the base template

In the bookr project's base template, located in `templates/base.html`, we have a placeholder `logout` link in the header. It is coded in HTML as follows:

```
<li class="nav-item">  
  <a class="nav-link" href="#">Logout</a>  
</li>
```

We don't want the `logout` link to appear after a user has logged out. So, this exercise aims to apply conditional logic to the template so that the `Login` and `Logout` links are toggled, depending on whether the user is authenticated:

1. Edit the `templates/base.html` file. Copy the structure of the `Logout` list element and create a `Login` list element. Then, replace the placeholder links with the correct URLs for the `Logout` and `Login` pages – `/accounts/logout` and `/accounts/login`, respectively – as follows:

```
<li class="nav-item">
    <a class="nav-link" href="/accounts/logout">Logout
    </a>
</li>
<li class="nav-item">
    <a class="nav-link" href="/accounts/login">Login</a>
</li>
```

2. Now, put our two li elements inside an `if ... else ... endif` conditional block. The logic condition that we are applying is `if user.is_authenticated`:

```
{% if user.is_authenticated %}
<li class="nav-item">
    <a class="nav-link" href="/accounts/logout">Logout
    </a>
</li>
{% else %}
<li class="nav-item">
    <a class="nav-link" href="/accounts/login">Login
    </a>
</li>
{% endif %}
```

3. Now, visit the user profile page at `http://localhost:8000/accounts/profile/`. When authenticated, you will see the Logout link:

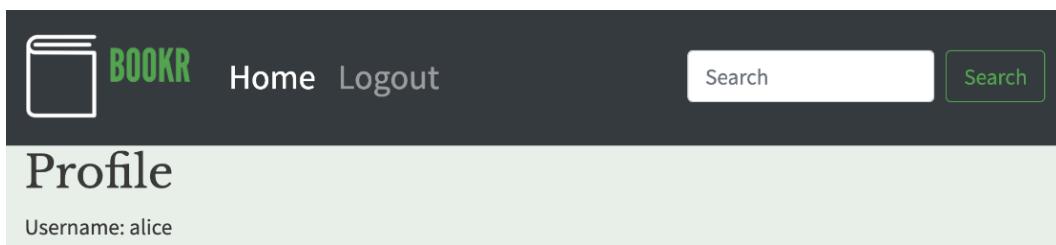


Figure 9.8 – An authenticated user sees the Logout link

4. Now, click the Logout link; you will be taken to the /accounts/logout page. The Login link appears in the menu, confirming that the link is contextually dependent on the authentication state of the user:

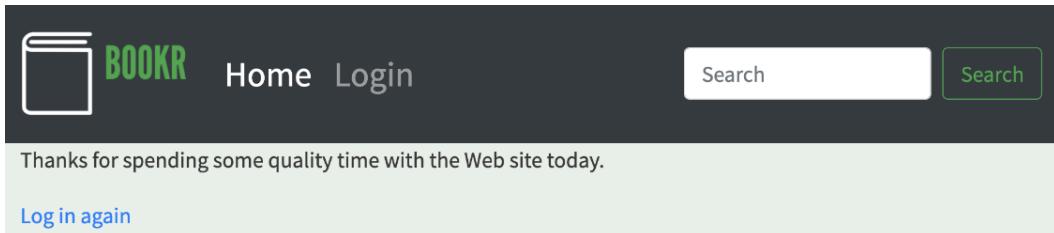


Figure 9.9 – An unauthenticated user sees the Login link

This exercise was a simple example of how Django templates can be used with authentication information to create a stateful and contextual user experience. We also do not want to provide links that a user does not have access to or actions that are not permissible for the user's permission level. The following activity will use this templating technique to fix some of these problems in Bookr.

Activity 9.01 – authentication-based content using conditional blocks in templates

In this activity, you will apply conditional blocks to templates that modify content based on user authentication and user status. Users should not be presented with links that they are not permitted to visit or actions that they are not authorized to carry out. The following steps will help you complete this activity:

1. In the `book_detail` template, in the file at `reviews/templates/reviews/book_detail.html`, hide the `Add Review` and `Media` buttons from non-authenticated users.
2. Also, hide the heading that says **Be the first one to write a review** as that is not an option for non-authenticated users.
3. In the same template, make the **Edit Review** link only appear for the staff or the user that wrote the review. The conditional logic for the template block is very similar to the conditional logic that we used in the `review_edit` view in the previous section:

Review Comments

Review comment: I learnt a lot from reading this book.

Created on: June 20, 2020, 5:05 a.m.

Modified on: June 20, 2020, 5:20 a.m.

Rating: 5

Creator: alice

[Edit Review](#)

[Add Review](#)

[Media](#)

Figure 9.10 – The Edit Review link appears on Alice’s review when Alice is logged in

When a different user is logged in, the **Edit Review** link will not be present:

Review Comments

Review comment: I learnt a lot from reading this book.

Created on: June 20, 2020, 5:05 a.m.

Modified on: June 20, 2020, 5:20 a.m.

Rating: 5

Creator: alice

[Add Review](#)

[Media](#)

Figure 9.11 – There is no Edit Review link on Alice’s review when Bob is logged in

4. Modify `template/base.html` so that it displays the currently authenticated user’s username to the right of the search form in the header, linking to the user profile page.

By completing this activity, you will have added dynamic content to the template that reflects the authentication state and identity of the current user, as can be seen from the following screenshot:

The screenshot shows a dark-themed user interface. At the top, there is a navigation bar with a logo icon (a book), the word 'BOOKR', and links for 'Home' and 'Logout'. To the right of the navigation bar is a search bar with a 'Search' button. To the right of the search bar, the text 'User: alice' is displayed. Below the navigation bar, the word 'Profile' is centered. At the bottom of the screenshot, the text 'Username: alice' is visible.

Figure 9.12 – An authenticated user’s name appears after the search form

Note

You can find the solution for this activity on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

With that, we have reviewed the available authentication mechanisms in Django and can examine how Django implements sessions.

Sessions

It is worth looking at some theory to understand why sessions are a common solution in web applications for managing user content. The HTTP protocol defines the interactions between a client and a server. It is said to be a “stateless” protocol as no stateful information is retained by the server between requests. This protocol design worked well for delivering hypertextual information in the early days of the World Wide Web, but it did not suit the needs of secured web applications delivering customized information to specific users.

We are now acquainted with seeing websites adapt to our viewing habits. Shopping sites recommend similar products to the ones that we have recently viewed and tell us about products that are popular in our region. These features all required a stateful approach to website development. One of the most common ways to implement a stateful web experience is through **sessions**. A session refers to a user’s current interaction with a web server or application and requires that data be persisted for the duration of the interaction. This may include information about the links that the user has visited, the actions that they have performed, and the preferences that they have made in their interactions.

If a user sets a blogging site to a dark theme on one page, there is an expectation that the next page will use the same theme as well. We describe this behavior as **maintaining state**. A session key is stored client-side as a browser cookie, which can be identified with server-side information that persists while the user is logged in.

In Django, sessions are implemented as a form of middleware. When we initially created the app in *Chapter 4, Introduction to Django Admin*, session support was activated by default.

The session engine

Information about current and expired sessions needs to be stored somewhere. In the early days of the World Wide Web, this was done through saving session information in files on the server, but as web server architectures have become more elaborate and their performance demands have increased, other more efficient strategies such as a database or in-memory storage have become the norm. By default, in Django, session information is stored in a project’s database.

This is a reasonable default for most small projects. However, Django’s middleware implementation of sessions gives us the flexibility to store our project’s session information in a variety of ways to

suit our system architecture and performance requirements. Each of these different implementations is called a session engine. If we want to change the session configuration, we need to specify the `SESSION_ENGINE` setting in the project's `settings.py` file:

- **Cached sessions:** In some environments, caching session information in memory or a database is an approach that is suited to high performance. Django provides the `django.contrib.sessions.backends.cache` and `django.contrib.sessions.backends.cached_db` session engines for this purpose.
- **File-based sessions:** As stated earlier, this is a somewhat antiquated way of maintaining session information but may suit some sites where performance is not an issue and there are reasons not to store dynamic information in a database.
- **Cookie-based sessions:** Rather than keeping session information on the server side, you can keep them entirely in the web browser client by serializing the content of the session as JSON and storing it in a browser-based cookie.

Do you need to flag cookie content?

All of Django's implementations of sessions require storing a session ID in a cookie on the user's web browser.

Regardless of the session engine employed, all these middleware implementations involve storing a site-specific cookie in the web browser. In the early days of web development, it was not uncommon to pass session IDs as URL arguments, but this approach has been eschewed in Django for reasons of security.

In many jurisdictions, including the European Union, websites are legally required to warn users if the site sets cookies in their browsers. If there are such legislative requirements in the region where you intend to operate your site, it is your responsibility to ensure that the code meets these obligations. Be sure to use up-to-date implementations and avoid using abandoned projects that have not kept pace with legislative changes.

Note

To cater to these changes and legislative requirements, there are many useful apps, such as **Django Cookie Consent** and **Django Cookie Law**, that are designed to work with several legislative frameworks. You can find out more by going to the following links:

- <https://github.com/jazzband/django-cookie-consent>
- <https://github.com/TyMaszWeb/django-cookie-law>

Many JavaScript modules exist that implement similar cookie consent mechanisms.

With this in mind, we can examine how session data is implemented in Django.

Pickle or JSON storage?

Python provides the **pickle** module in its standard library for serializing Python objects into a byte stream representation. A pickle is a binary structure that has the benefit of being interoperable between different architectures and different versions of Python so that a Python object can be serialized to a pickle on a Windows PC and deserialized to a Python object on a Linux Raspberry Pi.

This flexibility comes with security vulnerabilities and it is not recommended that it is used to represent untrusted data. Consider the following Python object, which contains several types of data. It can be serialized using `pickle`:

```
import datetime
data = dict(viewed_books=[17, 18, 3, 2, 1],
            search_history=['1981', 'Machine Learning',
                            'Bronte'],
            background_rgb=(96, 91, 92),
            foreground_rgb=(17, 17, 17),
            last_login_login=datetime.datetime(2019, 12, 3,
                                              15, 30, 30),
            password_change=datetime.datetime(2019, 9, 2,
                                              8, 41, 25),
            user_class='Veteran',
            average_rating=4.75,
            reviewed_books={18, 3, 7})
```

Using the `dumps` (dump string) method of the `pickle` module, we can serialize the `data` object to produce a byte representation:

```
import pickle
data_pickle = pickle.dumps(data)
```

JSON stands for **JavaScript Object Notation**. The syntax of JSON is a small subset of the JavaScript language. It is a widespread standard for messaging and data exchange, commonly used for transferring data between web browsers and servers. We can serialize JSON with a similar approach to the one that we outlined with the `pickle` format:

```
import json
data_json = json.dumps(data)
```

Because data contains Python `datetime` and `set` objects, which aren't serializable with JSON, when we attempt to serialize the structure, a type error will be thrown:

```
TypeError: Object of type datetime is not JSON serializable
```

For serializing to JSON, we could convert the `datetime` objects into the `string` type and `set` them in a list:

```
data['last_login_login'] =
    data['last_login_login'].strftime("%Y%d%m%H%M%S")
data['password_change'] =
    data['password_change'].strftime("%Y%d%m%H%M%S")
data['reviewed_books'] = list(data['reviewed_books'])
```

As JSON data is human-readable, it is easy to examine:

```
{"viewed_books": [17, 18, 3, 2, 1], "search_history": ["1981",
"Machine Learning", "Bronte"], "background_rgb": [96, 91, 92],
"foreground_rgb": [17, 17, 17], "last_login_login": "20190312153030",
"password_change": "20190209084125", "user_class": "Veteran",
"average_rating": 4.75, "reviewed_books": [18, 3, 7]}
```

Note that we had to explicitly convert the `datetime` and `set` objects, but the tuple is automatically converted into a list by the JSON. Django ships with `PickleSerializer` and `JSONSerializer`. If the serializer needs to be altered, it can be changed by setting the `SESSION_SERIALIZER` variable in the project's `settings.py` file:

```
SESSION_SERIALIZER =
'django.contrib.sessions.serializers.JSONSerializer'
```

Exercise 9.05 – examining the session key

The purpose of this exercise is to query the project's SQLite database and perform queries on the session table to become familiar with how session data is stored. You will then create a Django management command for examining session data that is stored using `JSONSerializer`:

1. Start the *DB Browser for SQLite* application and open your `db.sqlite3` database:

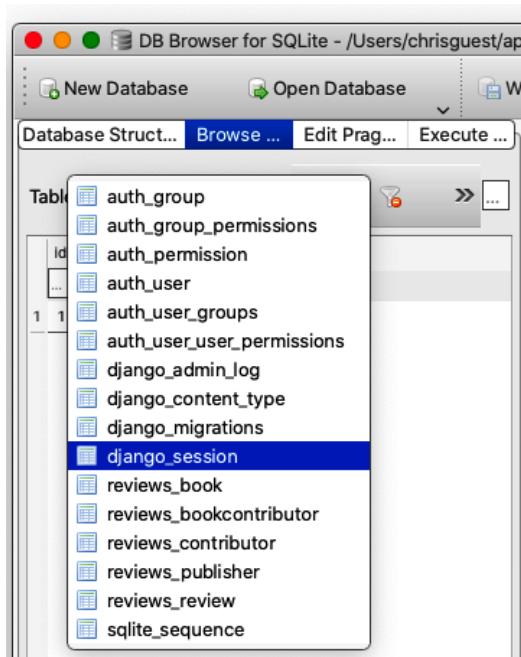


Figure 9.13 – Table selection in DB Browser for SQLite

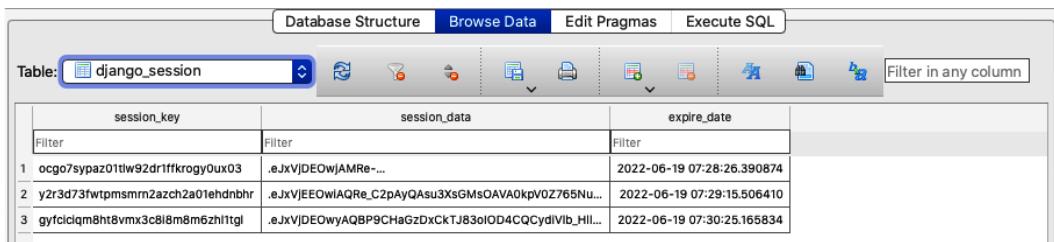
2. Under the **Database Structure** tab, expand the `django_session` table:

Database Structure			Browse Data	Edit Pragmas	Execute SQL
Name		Type	Schema		
▼	Tables (16)				
►	auth_group		CREATE TABLE "auth_group" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, "name" va		
►	auth_group_permissions		CREATE TABLE "auth_group_permissions" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT,		
►	auth_permission		CREATE TABLE "auth_permission" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, "cont		
►	auth_user		CREATE TABLE "auth_user" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, "password"		
►	auth_user_groups		CREATE TABLE "auth_user_groups" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, "use		
►	auth_user_user_permissions		CREATE TABLE "auth_user_user_permissions" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT,		
►	django_admin_log		CREATE TABLE "django_admin_log" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, "act		
►	django_content_type		CREATE TABLE "django_content_type" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, ";		
►	django_migrations		CREATE TABLE "django_migrations" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, "ap		
►	django_session		CREATE TABLE "django_session" ("session_key" varchar(40) NOT NULL PRIMARY KEY, "session_d		
	session_key	varchar(40)	"session_key" varchar(40) NOT NULL		
	session_data	text	"session_data" text NOT NULL		
	expire_date	datetime	"expire_date" datetime NOT NULL		
►	reviews_book		CREATE TABLE "reviews_book" ("id" integer NOT NULL PRIMARY KEY AUTOINCREMENT, "title" va		

Figure 9.14 – Table definitions in DB Browser for SQLite

This reveals that the `django_session` table in the database stores session information in the `session_key`, `session_data`, and `expire_date` fields.

3. Click on the **Browse Data** tab and select the `django_session` table:



	session_key	session_data	expire_date
1	ocgo7sypaz01tw92dr1ffkrogy0ux03	.eJxVjDEOwjlAMRe-...	2022-06-19 07:28:26.390874
2	y2r3d73fwtpmsmr2azch2a01ehdnbhr	.eJxVjEE0wlAQRe_C2pAyQAsu3XsGMsOAVA0kpV0Z765Nu...	2022-06-19 07:29:15.506410
3	gyfclclqm8ht8vmx3c8l8m6m6zh11tg1	.eJxVjDEOwyAQBP9ChagDxCkTJ83oI0D4CQCydlVlb_HII...	2022-06-19 07:30:25.165834

Figure 9.15 – Querying data in the `django_session` table

4. Django stores the session data using cryptographical signed values. Because by default we are using the database to store our sessions, we will import the `Session` model, as well as the `SessionStore` class of `django.contrib.sessions.backends.db.SessionStore` has a `decode` method that we can use to read the encoded `session_data` from the `Session` object.
5. This code snippet can be executed using the Django shell by entering `python manage.py shell` at the command line to get an interactive prompt:

```
>>> from django.contrib.sessions.backends.db import SessionStore
>>> from django.contrib.sessions.models import Session
>>> session_store = SessionStore()
>>> sessions = Session.objects.all()
>>> session_store.decode(sessions[0].session_data)
{'_auth_user_id': '2', '_auth_user_backend': 'django.contrib.auth.backends.ModelBackend', '_auth_user_hash': 'eb5de6bf54460bafc9e0f555c65d16b17b3fdf754617d1d73529d9605fea2f0c'}
```

6. Using this approach, we can create a management command for examining session data. Sensitive data may be stored in sessions, so creating a management command ensures that the data is only accessible by a user with console access to the server that is hosting Django.
7. Using **PyCharm**, in the `reviews/management/commands` directory of the Bookr project, create a Python file called `sessioninfo.py`.

8. Let's start by importing the necessary modules. We will use the `pformat` command from the `pprint` module from Python's Standard Library to format session data. We will need the `User` model from `django.contrib.auth.models` and the `Session` model from `django.contrib.sessions.models` to query `User` and `Session` objects. The `BaseCommand` class from `django.core.management.base` provides the scaffolding for our Django Admin command:

```
from pprint import pformat

from django.contrib.auth.models import User
from django.contrib.sessions.models import Session
from django.contrib.sessions.backends.db import SessionStore
from django.core.management.base import BaseCommand
```

9. At a minimum, a user command needs to be defined by subclassing `BaseCommand` as the `Command` class and defining a `handle` method. Typically, additional command-line parameters would be defined in an `add_arguments` method, but we are aiming to create a very simple example. A `help` attribute of the `Command` class is used to provide information when the user enters `python manage.py --help` or `python manage.py sessioninfo --help`:

```
class Command(BaseCommand):
    help = "List all user sessions and the data that
they contain."
```

10. Write a `handle` method of the `Command` class. First, instantiate a `SessionStore` object, then iterate through the `Session` objects and decode the data for each using the `SessionStore.decode` method. The `_auth_user_id` key of the decrypted data references a `User.id` from the database so that we can retrieve the user using the `User` model.

Now, for each session, write the session key, user ID, username, and email address to the console. We will develop a simple Python utility for decrypting this session information. Note that it is best practice to call `self.stdout.write` in a management command instead of `print` when writing to the console:

```
def handle(self, *args, **options):
    session_store = SessionStore()
    for session in Session.objects.all():
        data = session_store.decode(
            session.session_data)
        user = User.objects.get(
            id=data['_auth_user_id'])
        self.stdout.write(
            f"Session Key: {session.session_key}"
            f" User: {user.id} {user.username}"
            f" {user.email}")
        self.stdout.write(pformat(data))
```

```
{user.email} ") self.stdout.write(pformat(data))
```

11. Now, you can use this management command to examine session data that is stored in the database. You can call it on the command line as follows:

```
python manage.py sessioninfo
```

This will be useful for debugging session behavior when you attempt the final activity:

```
> python manage.py sessioninfo
Session Key: ocgo7sypaz01tlw92dr1ffkrogy0ux03 User: 2 carol carol.brown@example.com
{'_auth_user_backend': 'django.contrib.auth.backends.ModelBackend',
 '_auth_user_hash': 'eb5de6bf54460bafc9e0f555c65d16b17b3fdf754617d1d73529d9605fea2f0c',
 '_auth_user_id': '2'}
Session Key: y2r3d73fwtpmsmrn2azch2a01ehdnbhr User: 1 bob bob.black@example.com
{'_auth_user_backend': 'django.contrib.auth.backends.ModelBackend',
 '_auth_user_hash': '322625efba7788f598d02fcfa5ad86bb14db4893f3c92bc4099822ea3bf4d1e4f',
 '_auth_user_id': '1'}
Session Key: gyfciciqm8ht8vmx3c8i8m8m6zh1tgc User: 3 bookadmin bookadmin@example.com
{'_auth_user_backend': 'django.contrib.auth.backends.ModelBackend',
 '_auth_user_hash': 'a30d9315c9b5dda38802da5e8e64f94b0b063f0136fad8456205ed8cced8ff30',
 '_auth_user_id': '3'}
```

Figure 9.16 – Python script

This script outputs the decoded session information. At present, the session only contains three keys:

- `_auth_user_backend` is a string representation of the class of the user backend. As our project stores user credentials in the model, `ModelBackend` is used.
- `_auth_user_hash` is a hash of the user's password.
- `_auth_user_id` is the user ID obtained from the model's `User.id` attribute.

This exercise helped you become familiar with how session data is stored in Django. We will now turn our attention to adding additional information to Django sessions.

Storing data in sessions

We've already covered the way sessions are implemented in Django. Now, we are going to briefly examine some of the ways that we can make use of sessions to enrich our user experience. In Django, the session is an attribute of the `request` object. It is implemented as a dictionary-like object. In our views, we can assign keys to the `session` object like a typical dictionary, as shown here:

```
request.session['books_reviewed_count'] = 39
```

But there are some restrictions. First, the keys in the session must be strings, so integers and timestamps are not allowed. Secondly, keys starting with an underscore are reserved for internal system use. Data is limited to values that can be encoded as JSON, so some byte sequences that can't be decoded as UTF-8, such as `binary_key` listed previously, can't be stored as JSON data. The other warning is to avoid reassigning `request.session` to a different value. We should only assign or delete keys. So, don't do this:

```
request.session = {'books_read_count':30, 'books_reviewed_count': 39}
```

Instead, do this:

```
request.session['books_read_count'] = 30
request.session['books_reviewed_count'] = 39
```

With those restrictions in mind, we will investigate how we can make use of session data in our Reviews application.

Exercise 9.06 – storing recently viewed books in sessions

The purpose of this exercise is to use the session to keep information about the 10 books that have been most recently browsed by the authenticated user. This information will be displayed on the profile page of the bookr project. When a book is browsed, the `book_detail` view is called. In this exercise, we will edit `reviews/views.py` and add some additional logic to the `book_detail` method. We will add a key to the session called `viewed_books`. Using basic knowledge of HTML and CSS, the page can be created to show the profile details and viewed books stored in separate divisions of the page, as follows:

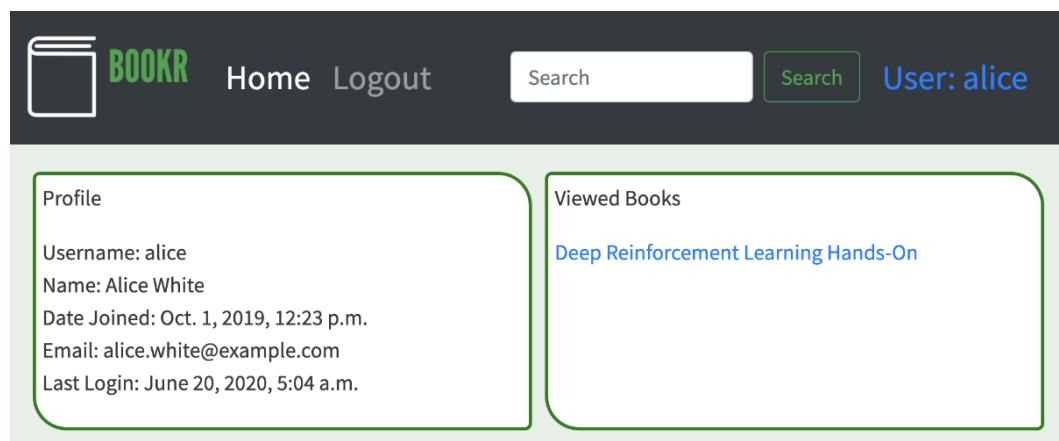


Figure 9.17 – The Profile page incorporating Viewed Books

Let's get started with the steps:

1. Edit `reviews/views.py` and the `book_detail` method. We are only interested in adding session information for authenticated users, so add a conditional statement to check whether the user is authenticated and set `max_viewed_books_length`, the maximum length of the viewed books list, to 10:

```
def book_detail(request, pk):  
    ...  
    if request.user.is_authenticated:  
        max_viewed_books_length = 10
```

2. Within the same conditional block, add code to retrieve the current value of `request.session['viewed_books']`. If this key isn't present in the session, start with an empty list:

```
viewed_books = request.session.get(  
    'viewed_books', [])
```

3. If the current book's primary key is already present in `viewed_books`, the following code will remove it:

```
viewed_book = [book.id, book.title]  
if viewed_book in viewed_books:  
    viewed_books.pop(  
        viewed_books.index(viewed_book))
```

4. The following code inserts the current book's primary key at the start of the `viewed_books` list:

```
viewed_books.insert(0, viewed_book)
```

5. Add the following key to only keep the first 10 elements of the list:

```
viewed_books = viewed_books[  
    :max_viewed_books_length]
```

6. The following code will add our `viewed_books` back to `session['viewed_books']` so that it is available in subsequent requests:

```
request.session['viewed_books'] = viewed_books
```

7. As before, at the end of the `book_detail` function, render the `reviews/book_detail.html` template given the request and context data:

```
return render(request, "reviews/book_detail.html", context)
```

Once complete, the `book_detail` view will have this conditional block:

```
def book_detail(request, pk):
    ...
    if request.user.is_authenticated:
        max_viewed_books_length = 10
        viewed_books = request.session.get("viewed_books", [])
        viewed_book = [book.id, book.title]
        if viewed_book in viewed_books:
            viewed_books.pop(viewed_books.index(viewed_book))
        viewed_books.insert(0, viewed_book)
        viewed_books = viewed_books[:max_viewed_books_length]
        request.session["viewed_books"] = viewed_books
    return render(request, "reviews/book_detail.html", context)
```

8. Modify the page layout and CSS of `templates/profile.html` to accommodate the viewed book division. As we may add more divisions to this page in the future, one convenient layout concept is the **flexbox**. We will add this CSS and separate the content into nested `div` instances that will be arranged horizontally on the page. We will refer to the internal `div` instances as `infocell` instances and style them with green borders and rounded edges:

```
<style>
.flexrow { display: flex;
            border: 2px black;
}
.flexrow > div { flex: 1; }

.infocell {
    border: 2px solid green;
    border-radius: 5px 25px;
    background-color: white;
    padding: 5px;
    margin: 20px 5px 5px 5px;
}

</style>

<div class="flexrow" >
    <div class="infocell" >
        <p>Profile</p>
        ...
    </div>

    <div class="infocell" >
        <p>Viewed Books</p>
        ...
    </div>
</div>
```

```
</div>
</div>
```

9. Modify the Viewed Books div element in `templates/profile.html` so that if there are books present, their titles are displayed and linked to the individual book detail pages. This will be rendered as follows:

```
<a href="/books/1">Advanced Deep Learning with Keras
</a><br>
```

A message should be displayed if the list is empty. The entire div, including the iteration through `request.session.viewed_books`, will look like this:

```
<div class="infocell" >
  <p>Viewed Books</p>
  <p>
    {%
      for book_id, book_title in
      request.session.viewed_books %
    }
    <a href="/books/{{ book_id }}"/>{{ book_title }}
    </a><br>
    {%
      empty %
    }
    No recently viewed books found.
    {%
      endfor %
    }
  </p>
</div>
```

The complete profile template will look like this once all these changes have been incorporated:

templates/profile.html

```
{% extends "base.html" %}

{% block title %}Bookr{% endblock %}

{% block heading %}Profile{% endblock %}

{% block content %}

<style>
```

You can find the complete code for this file at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter09/Exercise06/bookr/templates/profile.html>.

This exercise has enhanced the profile page by adding a list of recently viewed books. Now, when you visit the login link at `http://127.0.0.1:8000/accounts/profile/`, you will see this page:

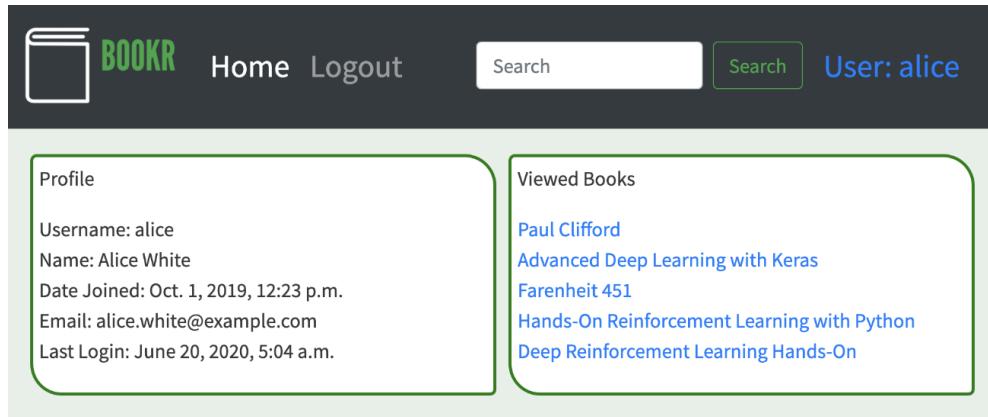

 A screenshot of a web application interface. At the top, there is a dark header bar with a logo icon (a book) and the text 'BOOKR'. To the right of the logo are links for 'Home' and 'Logout'. On the far right of the header, it says 'User: alice'. Below the header, there are two main sections. The left section is titled 'Profile' and contains user information: 'Username: alice', 'Name: Alice White', 'Date Joined: Oct. 1, 2019, 12:23 p.m.', 'Email: alice.white@example.com', and 'Last Login: June 20, 2020, 5:04 a.m.'. The right section is titled 'Viewed Books' and lists five recently viewed books: 'Paul Clifford', 'Advanced Deep Learning with Keras', 'Farenheit 451', 'Hands-On Reinforcement Learning with Python', and 'Deep Reinforcement Learning Hands-On'.

Figure 9.18 – Recently viewed books

We can use the `sessioninfo` management command that we developed in *Exercise 9.05 – examining the session key*, to examine the user's session once this feature is implemented. It can be called on the command line by passing the session data as an argument:

```
python manage.py sessioninfo
```

We can see that the book IDs and titles are listed in the `viewed_books` key. Remember that the encoded data is obtained by querying the `django_session` table in the SQLite database:

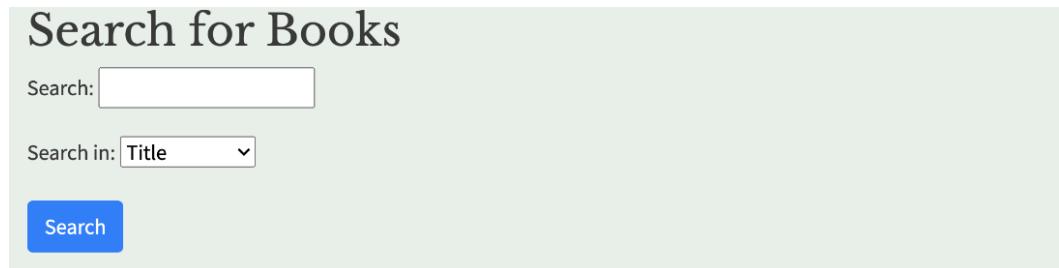
```
> python manage.py sessioninfo
Session Key: mjbppj6e6y0qayh1mj6sgw27kfugk1qc User: 4 alice alice.white@example.com
{'_auth_user_backend': 'django.contrib.auth.backends.ModelBackend',
 '_auth_user_hash': '384340e0724b154062dcad2c94a9fc94380cd323b9e1e433aff3aef799b99a3e',
 '_auth_user_id': '4',
 'viewed_books': [[17, 'Paul Clifford'],
                  [1, 'Advanced Deep Learning with Keras'],
                  [13, 'Farenheit 451'],
                  [6, 'Hands-On Reinforcement Learning with Python'],
                  [4, 'Deep Reinforcement Learning Hands-On']]}
```

Figure 9.19 – The viewed books are stored in the session data

In this exercise, we used Django's session mechanism to store ephemeral information about user interactions with the Django project. We have learned how this information can be retrieved from the user session and displayed in a view that informs users about their recent activity.

Activity 9.02 – using session storage for the Book Search page

Sessions are a useful way to store short-lived information that assists in maintaining a stateful experience on a site. Users frequently revisit pages such as search forms, and it would be convenient to store their most recently used form settings when they return to those pages. In *Chapter 3, URL Mapping, Views, and Templates*, we developed a book search feature for the bookr project. The Book Search page has two options for Search in – Title and Contributor. Currently, each time the page is visited, it defaults to Title:



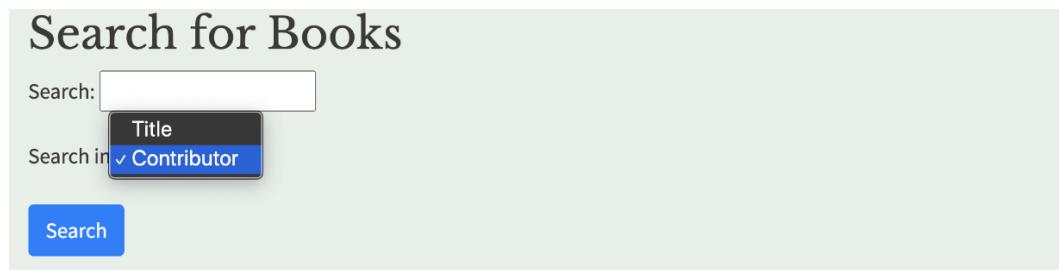
The screenshot shows a 'Search for Books' form. It has a 'Search:' input field, a 'Search in:' dropdown menu with 'Title' selected, and a 'Search' button.

Figure 9.20 – The Search and Search in fields of the Book Search form

In this activity, you will use session storage so that when the Book Search page, /book-search, is visited, it will default to the most recently used search option. You will also add a third infocell to the profile page that contains a list of links to the most recently used search terms. Follow these steps to complete this activity:

1. Edit the book_search view and retrieve search_history from the session.
2. When the form has received valid input and a user is logged in, append the search option and search text to the session's search history list.

If the form hasn't been filled (for example, when the page is first visited), render the form with the previously used Search in option selected – that is, either Title or Contributor (Figure 9.21):



The screenshot shows a 'Search for Books' form. It has a 'Search:' input field, a 'Search in:' dropdown menu with 'Contributor' selected (indicated by a checked checkbox), and a 'Search' button.

Figure 9.21 – Selecting Contributor in the Search field

3. In the profile template, include an additional `infocell` division for `Search History`.
4. List the search history as a series of links to the book search page. The links will take this form: `/book-search?search=Python&search_in=title`.

This activity will challenge you to apply session data to solve a usability issue in a web form. This approach will have applicability in many real-world situations and will give you some idea of the use of sessions in creating a stateful web experience. After completing this activity, the profile page will contain the third `infocell`, as shown in *Figure 9.22*:

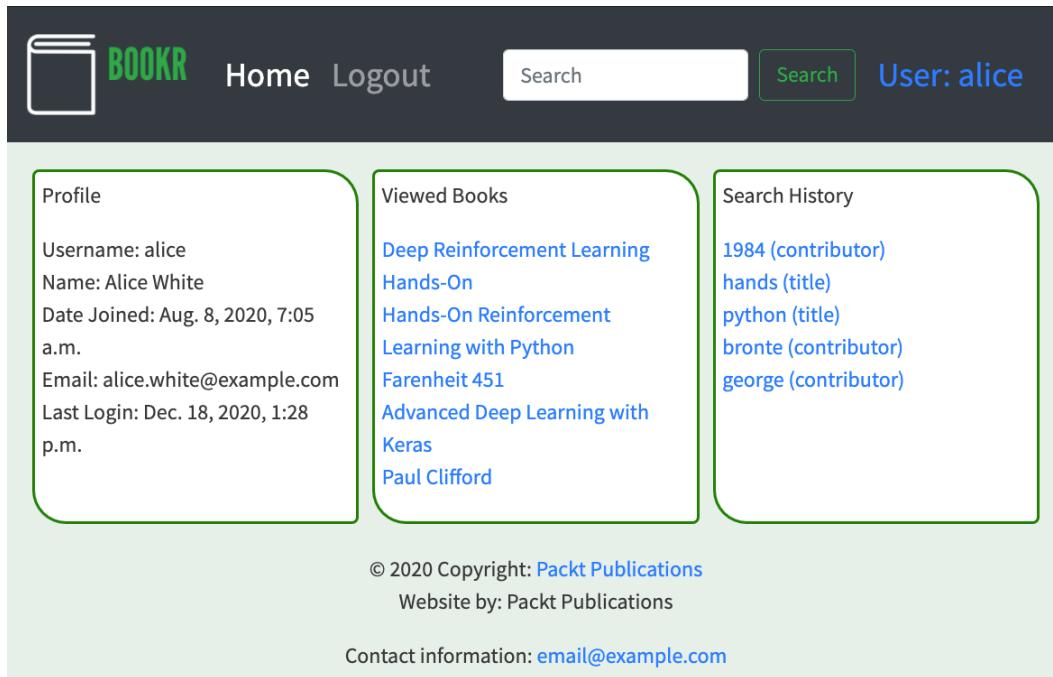


Figure 9.22 – The profile page with the Search History infocell

Note

You can find the solution for this activity on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we examined Django's middleware implementation of authentication and sessions. We also learned how to incorporate authentication and permission logic into views and templates. Now, we can set permissions on specific pages and limit their access to authenticated users. We also examined how to store data in a user's session and render it on subsequent pages.

With that, you have the skills to customize a Django project to deliver a personalized web experience. You can limit the content to authenticated or privileged users and you can personalize a user's experience based on their prior interactions. In the next chapter, we will revisit the Admin app and learn about some advanced techniques we can use to customize our user model and apply fine-grained changes to the admin interface for our models.

10

Advanced Django Admin and Customizations

Let's say we want to customize the front page of a large organization's admin site. We want to show the health of the different systems in the organization and see any high-priority alerts that are active. If this were an internal website built on top of Django, we would need to customize it. Adding these kinds of functionalities will require the developers in the IT team to customize the default admin panel and create their own custom `AdminSite` module, which will render a different index page in comparison to what is provided by the default admin site. Fortunately, Django makes these kinds of customizations easy.

In this chapter, we will look at how we can leverage Django's framework and its extensibility to customize Django's default admin interface (as shown in *Figure 10.1*). We'll not just learn how to make the interface more personal; we will also learn how we can control the different aspects of the admin site to make Django load a custom admin site, instead of the one that ships with the default framework. Such customization can come in handy when we want to introduce features into the admin site that are not present by default:

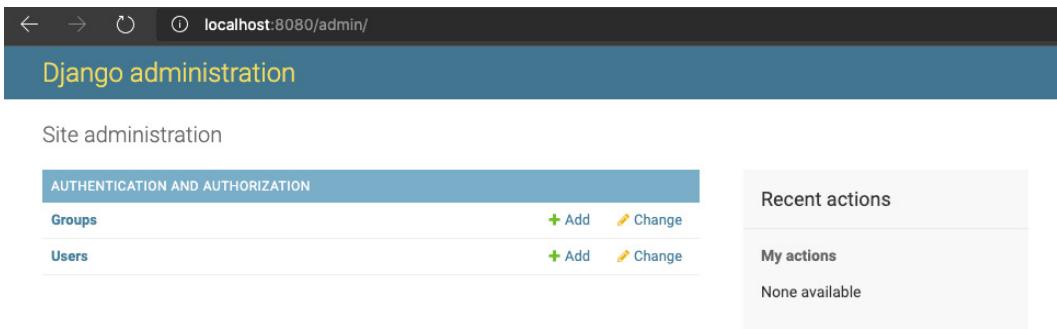


Figure 10.1 – Default Django administration panel interface

This chapter builds upon the skills we practiced in *Chapter 4, Introduction to Django Admin*. Just to recap, we learned how to use the Django admin site to take control of the administration and authorization for our Bookr app. We also learned how to register models to read and edit their contents and also to customize Django's admin interface using the `admin.site` properties.

The following topics are covered in this chapter:

- Customizing the admin site
- Adding views to the admin site

Now, let's expand our knowledge further by taking a look at how we can start customizing the admin site by utilizing Django's `AdminSite` module to add powerful new functionalities to the admin portal of our web application.

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter10>

Customizing the admin site

Django, as a web framework, provides a lot of customization options for building web applications. We will be using this same freedom provided by Django when we work on building the admin application for our project.

In *Chapter 4, Introduction to Django Admin*, we looked at how we can use the `admin.site` properties to customize the elements of Django's admin interface. However, what if we require more control over how our admin site behaves? For example, what if we wanted to show a custom template for the login page or the logout page when the user comes to the Bookr admin panel? In this case, the `admin.site` properties provided might not be enough, and we will need to build customizations that can extend the default admin site's behavior. Luckily, this can be easily achieved by extending the `AdminSite` class from Django's `admin` module. However, before we jump into building our admin site, let's first understand how Django discovers admin files and how we can use this admin file discovery mechanism to build a new app inside Django that will act as our admin site app.

Discovering admin files in Django

When we build applications in our Django project, we use the `admin.py` file frequently to register our models or create `ModelAdmin` classes that customize our interactions with the models inside the admin interface. These `admin.py` files store and provide this information to our project's admin interface. The discovery of these files is affected automatically by Django once we add `django.contrib.admin` to the `INSTALLED_APPS` section inside our `settings.py` file:

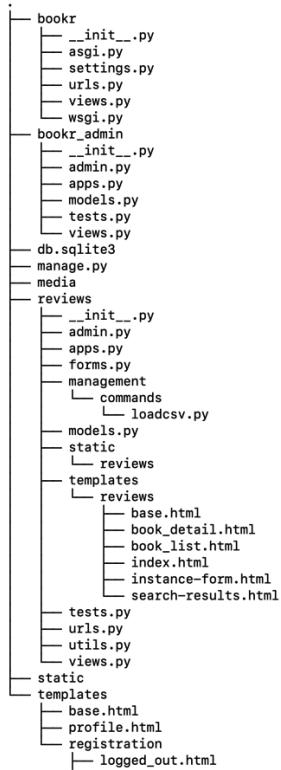


Figure 10.2 – The Bookr application structure

As we can see in the preceding figure, we have an `admin.py` file under the `reviews` application directory that is used by Django to customize the admin site for Bookr.

When the admin application is added, it tries to find the `admin` module inside every app of the Django project we are working on, and if a module is found, the admin application loads the contents from it.

Django's AdminSite class

Before we start customizing Django's admin site, we must understand how the default admin site is generated and handled by Django.

To provide us with the default admin site, Django packages a module known as the `admin` module, which holds a class known as `AdminSite`. This class implements a lot of useful functionalities and intelligent defaults that the Django community considers important for implementing a useful administration panel for most Django websites. The default `AdminSite` class provides a lot of built-in properties that not only control the look and feel of how the default admin site is rendered in the web browser but also control the way we can interact with it, and how a particular interaction will result in an action.

Some of these defaults include the site template properties, such as text to be shown in the site header, text to be shown in the title bar of the web browser, integration with Django's `auth` module for authenticating to the admin site, and a host of other properties.

As we progress on our path to building a custom admin site for our Django web project, it is more than desirable to retain a lot of the useful functionalities that are already built into Django's `AdminSite` class. This is where the concepts of Python object-oriented programming come to our rescue.

As we start moving forward to create our custom admin site, we will try to leverage the existing useful set of functionalities that are provided by Django's default `AdminSite` class. To do this, instead of building everything from scratch, we will work on creating a new child class that inherits from Django's `AdminSite` class to leverage the existing set of functionalities and useful integration that Django already provides us with. This kind of approach allows us to focus on adding a new and useful set of functionalities to our custom admin site, rather than spending time implementing the basic set of functionalities from scratch. For example, the following code snippet shows how we can create a child class of Django's `AdminSite` class:

```
class MyAdminSite(admin.AdminSite):  
    ...
```

To start working on our custom admin site for our web application, let's start by overriding some of the basic properties of Django's admin panel through the use of the custom `AdminSite` class we are going to work on.

Some of the properties that can be overridden include `site_header`, `site_title`, and others.

Note

When creating a custom admin site, we will have to register once again any `Model` and `ModelAdmin` classes that we might have registered using the default `admin.site` variable earlier. This happens because a custom admin site doesn't inherit the instance details from the default admin site provided by Django, and so unless we re-register our `Model` and `ModelAdmin` interfaces, our custom admin site will not show them.

Now, with the knowledge of how Django discovers what to load into the admin interface and how we can start building our custom admin site, let's go ahead and try to create our custom admin app for Bookr, which extends the existing `admin` module provided by Django. In the exercise that follows, we are going to create a custom admin site interface for our Bookr application using Django's `AdminSite` class.

Exercise 10.01 – creating a custom admin site for Bookr

In this exercise, you will create a new application that extends the default Django admin site and allows you to customize the components of the interface. Consequently, you will customize the default title of Django's admin panel. Once that is done, you will override the default value of Django's `admin.site` property to point to your custom admin site:

1. Before you can start working on your custom admin site, you first need to make sure that you are in the correct directory in your project from where you can run your Django application's management commands. For this, use the Terminal or Windows Command Prompt to navigate to the `bookr` directory and then create a new application named `bookr_admin`, which is going to act as the admin site for Bookr, by running the following commands:

```
python3 manage.py startapp bookr_admin
```

Once this command has been executed successfully, you should have a new directory named `bookr_admin` inside your project.

2. Now, with the default structure configured, the next step is to create a new class named `BookrAdmin`, which will extend the `AdminSite` class provided by Django to inherit the properties of the default admin site. To do this, open the `admin.py` file under the `bookr_admin` directory inside PyCharm. Once the file is open, you will see that the file already has the following code snippet present inside it:

```
from django.contrib import admin
```

Now, keeping this `import` statement as is, starting from the next line, create a new class named `BookrAdmin`, which inherits from the `AdminSite` class provided by the `admin` module you imported earlier:

```
class BookrAdmin(admin.AdminSite):
```

Inside this new `BookrAdmin` class, override the default value for the `site_header` variable, which is responsible for rendering the site header in Django's admin panel by setting the `site_header` property, as shown next:

```
    site_header = "Bookr Administration"
```

With this, the custom admin site class is now defined. To use this class, you will first create an instance of this class. This can be done as follows:

```
admin_site = BookrAdmin(name='bookr_admin')
```

3. Save the file but don't close it yet; we'll revisit it in *step 5*. Next, let's edit the `urls.py` file in the `bookr` app.

- With the custom class now defined, the next step is to modify the `urlpatterns` list to map the `/admin` endpoint in our project to the new `AdminSite` class you created. To do this, open the `urls.py` file under the `Bookr` project directory inside PyCharm and change the mapping of the `admin/` endpoint to point to our custom site:

```
from bookr_admin.admin import admin_site

urlpatterns = [
    ...,
    path('admin/', admin_site.urls)
]
```

We first imported the `admin_site` object from the `admin` module of the `bookr_admin` app. Then, we used the `urls` property of the object to map to the `admin/` endpoint in our application as follows:

```
path('admin/', admin_site.urls)
```

In this case, the `urls` property of our `admin_site` object is being automatically populated by the `admin.AdminSite` base class provided by Django's `admin` module. Once complete, your `urls.py` file should look like this: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter10/Exercise10.01/bookr/urls.py>.

- Now, with the configuration done, let's run our admin app in the browser. For this, run the following command from the root of your project directory where the `manage.py` file is located:

```
python manage.py runserver localhost:8000
```

Then, navigate to `http://localhost:8000/admin` (or `http://127.0.0.1:8000/admin`), which opens a page that resembles the following screenshot:

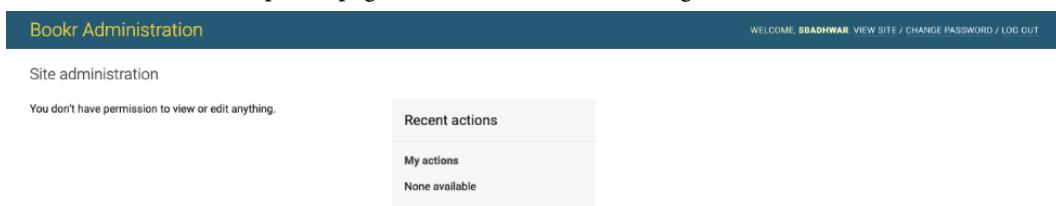


Figure 10.3 – The home page view for the custom Bookr admin site

In the preceding screenshot (Figure 10.3), you will see that Django displays a message, **You don't have permission to view or edit anything**. The issue of not having adequate permissions happens because, before now, we have not registered any models with our custom AdminSite instance. The issue also applies to the User and Groups models that are shipped along with the Django auth module. Now, let's make our custom admin site a bit more useful by registering the User model from Django's auth module.

6. To register the User model from Django's auth module, open the `admin.py` file under the `bookr_admin` directory inside PyCharm, and add the following line at the top of the file:

```
from django.contrib.auth.admin import User
```

At the end of the file, use your `BookrAdmin` instance to register this model as follows:

```
admin_site.register(User)
```

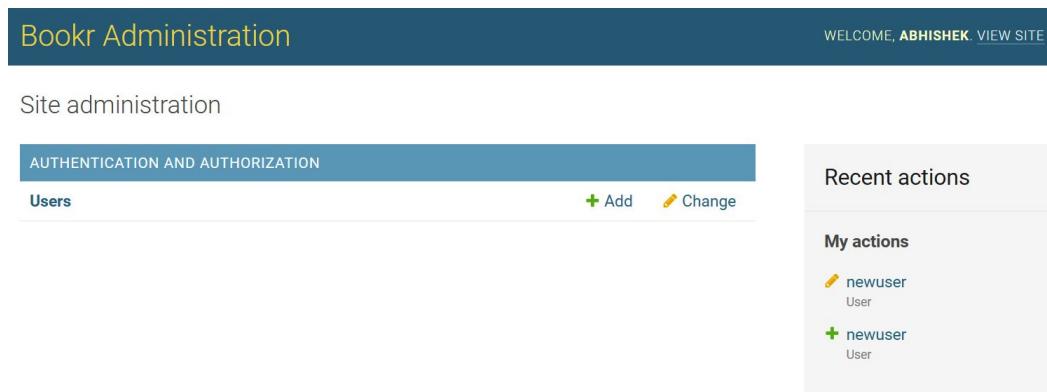
By now, your `admin.py` file should look like this:

```
from django.contrib import admin
from django.contrib.auth.admin import User

class BookrAdmin(admin.AdminSite):
    site_header = "Bookr Administration"

admin_site = BookrAdmin(name='bookr_admin')
admin_site.register(User)
```

Once this is done, reload the web server and visit `http://localhost:8000/admin`. Now, you should be able to see the User model being displayed for editing inside the admin interface, as shown here:



The screenshot shows the 'Bookr Administration' home page. At the top, there is a dark header bar with the text 'Bookr Administration' on the left and 'WELCOME, ABHISHEK. VIEW SITE /' on the right. Below the header, the page is titled 'Site administration'. A blue navigation bar at the top left contains the text 'AUTHENTICATION AND AUTHORIZATION' and a 'Users' link. To the right of the navigation bar are two buttons: a green '+' icon labeled 'Add' and a yellow pencil icon labeled 'Change'. On the right side of the page, there is a sidebar titled 'Recent actions' which lists two entries: 'newuser' (User) and another 'newuser' (User). Below the sidebar, there is a section titled 'My actions' with the same two entries.

Figure 10.4 – The home page view showing our registered models on the Bookr Administration site

With this, we just created our admin site application, and we can also now validate the fact that the custom site has a different header – **Bookr Administration**.

Overriding the default admin.site

In the previous section, after we created our own AdminSite application, we saw that we had to register models manually. This happens because most of the apps that we have built prior to our custom admin site still use the `admin.site` property to register their models, and if we want to use our AdminSite instance, we will have to update all those applications to use our instance, which can become cumbersome if there are a lot of applications inside a project.

Luckily, we can avoid this additional burden by overriding the default `admin.site` property. To do this, we first have to create a new AdminConfig class, which will override the default `admin.site` property for us, so that our application is marked as the default admin site and, hence, overrides the `admin.site` property inside our project. In the next exercise, we'll look at how we can map our custom admin site as a default admin site for an application.

Exercise 10.02 – overriding the default admin site

In this exercise, you will use the AdminConfig class to override the default admin site for your project such that you can keep on using the default `admin.site` variable to register models, override site properties, and so on:

1. Open the `admin.py` file under the `bookr_admin` directory and remove the import for the `User` model and the `BookrAdmin` instance creation, which you wrote in *step 5* of *Exercise 10.01 – creating a custom admin site for Bookr*. Once this is done, the file contents should resemble the following:

```
from django.contrib import admin

class BookrAdmin(admin.AdminSite):
    site_header = "Bookr Administration"
```

2. You will then need to create an AdminConfig class for the custom admin site, such that Django recognizes the `BookrAdmin` class as `AdminSite` and overrides the `admin.site` property. To do this, open up the `apps.py` file inside the `bookr_admin` directory and overwrite the contents of the file with the contents shown here:

```
from django.contrib.admin.apps import AdminConfig

class BookrAdminConfig(AdminConfig):
    default_site = 'bookr_admin.admin.BookrAdmin'
```

In this, we first imported the `AdminConfig` class from Django's `admin` module. This class is used to define which application should be used as a default admin site and to override the default behavior of how Django's admin site works.

For our use case, we create a class called `BookrAdminConfig`, which acts as a child class of Django's `AdminConfig` class and overrides the `default_site` property to point to our `BookrAdmin` class, which is our custom admin site:

```
default_site = 'bookr_admin.admin.BookrAdmin'
```

Once this is done, we need to set our application as an admin application inside our `Bookr` project. To achieve this, open the `settings.py` file of the `Bookr` project, and under the `INSTALLED_APPS` section, replace `'reviews.apps.ReviewsAdminConfig'` with `'bookr_admin.apps.BookrAdminConfig'`. The `settings.py` file should look like this: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter10/Exercise10.02/bookr/settings.py>.

3. With the application mapped as the admin application, the final step involves modifying the URL mapping such that the `'admin/'` endpoint uses the `admin.site` property to find the correct URL. For this, open the `urls.py` file under the `bookr` project. Consider the following entry in the `urlpatterns` list:

```
path('admin/', admin_site.urls)
```

Replace it with the following entry:

```
from django.contrib import admin

urlpatterns = [
    ...
    path('admin/', admin.site.urls)
]
```

Remember that `admin_site.urls` is a module, while `admin.site` is a Django internal property.

Once the preceding steps are complete, let's reload our web server and check whether our admin site loads by visiting `http://localhost:8000/admin`. If the website that loads looks like the one shown here, we have our own custom admin app now being used for the admin interface:

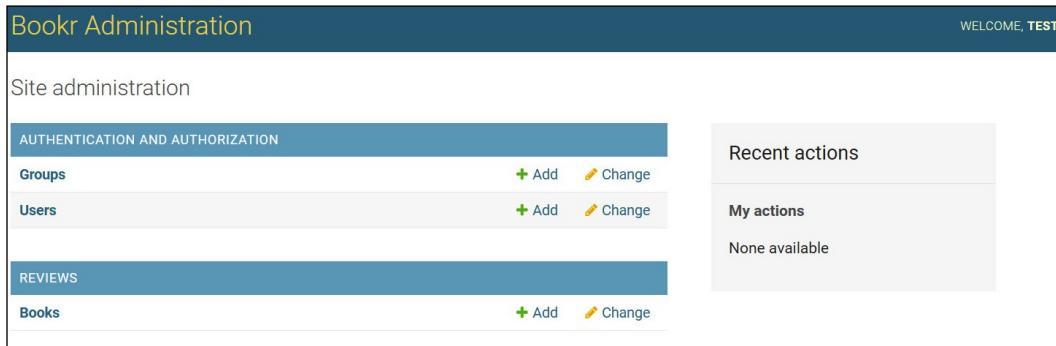


Figure 10.5 – The home page view of the custom Bookr Administration site

As you can see, once we override `admin.site` with our `admin` app, the models that were registered earlier using the `admin.site.register` property start to show up automatically.

With this, we now have a custom base template, which we can now utilize to build the remainder of our Django admin customizations. As we work through the chapter, we will discover some interesting customizations that allow us to make the admin dashboard an integrated part of our application.

Customizing admin site text using AdminSite attributes

Just as we can use the `admin.site` properties to customize the text for our Django application, we can also use the attributes exposed by the `AdminSite` class to customize these texts. In *Exercise 10.02 – overriding the default admin site*, we took a look at updating the `site_header` property of the admin site. Similarly, there are many other properties we can modify. Some of the properties that can be overridden are described as follows:

- `site_header`: Text to display at the top of every admin page (defaults to `Django Administration`).
- `site_title`: Text to display in the title bar of the browser (defaults to `Django Admin Site`).
- `site_url`: The link to use for the `View Site` option (defaults to `/`). This is overridden when the site runs on a custom path, and the redirection should take the user to the sub-path directly.
- `index_title`: This is the text that should be shown on the index page of the admin application (defaults to `Site administration`).

Note

For more information on all the `Adminsite` attributes, refer to the official Django documentation at <https://docs.djangoproject.com/en/4.1/ref/contrib/admin/#adminsite-attributes>.

If we want to override these attributes in our custom admin site, the process is very simple, as shown next:

```
class MyAdminSite(admin.AdminSite):  
    site_header = "My web application"  
    site_title = "My Django Web application"  
    index_title = "Administration Panel"
```

As we have seen in the examples so far, we have created a custom admin application for Bookr and then made it the default admin site for our project. An interesting question that arises is, since the properties that we have customized hitherto can also be customized by using the `admin.site` object directly, when should we create a custom admin application in comparison to modifying the `admin.site` properties?

As it turns out, there could be multiple reasons why someone would opt for a custom admin site – for example, they want to change the layout of the default admin site to make it align with the overall layout of their application. This is very common when you create a web application for a business where the homogeneity of the content is very important. Here is a short list of requirements that may compel a developer to go ahead and build a custom admin site, as opposed to simply modifying the properties of the `admin.site` variable:

- A need to override the index template for the admin interface
- A need to override the login or logout template
- A need to add a custom view to the admin interface

Customizing admin site templates

Just like some of the customizable common texts, such as `site_header` and `site_title`, that appear across the admin site, Django also allows us to customize the templates, which are used to render different pages on the admin site by setting certain properties in the `AdminSite` class.

These customizations can include the modification of templates that are used to render the index page, login page, model data page, and more. These customizations can be easily done by leveraging the templating system provided by Django. For example, the following code snippet shows how we can add a new template to the Django admin dashboard:

```
{% extends "admin/base_site.html" %}

{% block content %}
    <!-- Template Content -->
{% endblock %}
```

In this custom template, there are a couple of important aspects that we need to understand.

When customizing the existing Django admin dashboard by modifying how certain pages inside the dashboard appear or by adding a new set of pages to the dashboard, a developer might not want to write every single piece of HTML again from scratch to maintain the basic look and feel of the Django admin dashboard.

Usually, while customizing the admin dashboard, we want to retain the layout in which Django organizes the different elements displayed on the dashboard such that we can focus on modifying parts of a page that matter to us. This basic layout of the page, along with the common page elements, such as the page header and page footer, are defined inside the Django admin's base template, which also acts as a master template for all the pages inside the default Django admin website.

When we want to retain the way the common elements inside the Django admin pages are organized and rendered, we need to extend from this base template such that our custom template pages provide a consistent user experience in the same way as the other pages inside the Django admin dashboard. This can be done by using the template extension tags and extending the `base_site.html` template from the `admin` module provided by Django:

```
{% extends "admin/base_site.html" %}
```

Once this is done, the next task is to define our own content for the custom template. The `base_site.html` template provided by Django provides a block-based placeholder for a developer to add their own content for the template. To add this content, the developer has to put the logic for their own custom elements for the page inside the `{% block content %}` tags. This essentially overrides any content defined by the `{% block content %}` tag inside the `base_site.html` template, following the concepts of template inheritance in Django.

Now, let's look at how we can customize the template, which is used to render the logout page, once the user clicks the **Logout** button in the admin panel.

Exercise 10.03 – customizing the logout template for the Bookr admin site

In this exercise, you are going to customize the template that is used to render the logout page once the user clicks the **Logout** button on the admin site. Such overrides can come in handy when a website developed for a banking user might want to redirect them to an instruction page once they click **Logout**, in order to show them instructions to make sure that their banking session is securely closed:

1. Under the `templates` directory that you should have created in the earlier chapters, create another directory named `admin`, which will be used for storing templates for your custom admin site.

Note

Before proceeding, make sure that the `templates` directory is added to the `DIRS` list in your `settings.py` file (under `bookr/project`).

2. Now, with the directory structure setup complete and Django configured to load the templates, the next step involves writing a custom logout template that you want to render. To do this, let's create a new file named `logout.html` under the `templates/admin` directory we created in *step 1* and add the following content beneath it:

```
{% extends "admin/base_site.html" %}

{% block content %}


You have been logged out from the Admin panel.</p>
<p><a href="{% url 'admin:index' %}">Login Again</a>
or <a href="{{ site_url }}">Go to Home Page</a></p>
{% endblock %}


```

In the preceding code snippet, we are doing a couple of things. First, for our custom logout template, we are going to use the same master layout as provided by the `django.contrib.admin` module. So, consider the following:

```
{% extends "admin/base_site.html" %}
```

When we write this, Django tries to find and load the `admin/base_site.html` template inside the `templates` directory provided by the `django.contrib.admin` module.

Now, with our base template all set to be extended, the next thing we will do is try to override the HTML of the content block by executing the following command:

```
{% block content %}
...
{% endblock %}
```

The value of `admin:index` and `site_url` is provided by the `AdminSite` class automatically, based on the settings we define.

Using the value for `admin:index` and `site_url`, we will create our **Login Again** hyperlink, which, when clicked, will take the user back to the login form, and the **Go to Home Page** link, which will take the user back to the home page of the website. The file should look like this now: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter10/Exercise10.03/templates/admin/logout.html>.

3. Now, with the custom template defined, the next step is to make use of the custom template on our custom admin site. To do this, let's open the `admin.py` file under the `bookr_admin` directory and add the following field as the final value in the `BookrAdmin` class:

```
logout_template = 'admin/logout.html'
```

Save the file. It should look like this: https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter10/Exercise10.03/bookr_admin/admin.py.

4. Once all the preceding steps are complete, let's start our development server by running the following command:

```
python manage.py runserver localhost:8000
```

Then, we navigate to <http://localhost:8000/admin>.

Once you are there, try to log in and then click **Logout**. Once you are logged out, you will see the following page rendered:

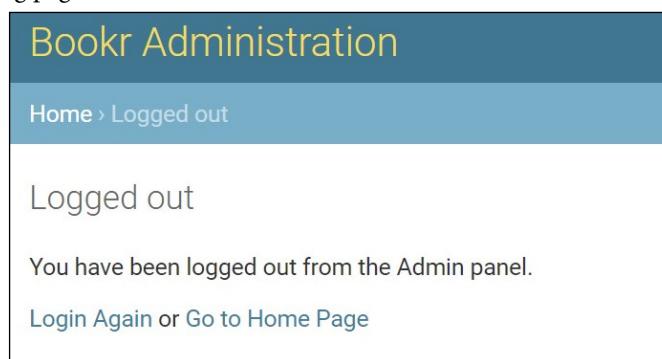


Figure 10.6 – The logout view rendered to users after clicking the Logout button

With this, we have successfully overridden our first template. Similarly, we can also override other templates inside Django's admin panel, such as the templates for the index view and the login form.

Adding views to the admin site

Just like general applications inside Django, which can have multiple views associated with them, Django allows developers to add custom views to the admin site also, which allows a developer to increase the scope of what the admin site interface can do.

The ability to add your own views to the admin site provides a lot of extensibility to the admin panel of the website, which can be leveraged for several additional use cases. For example, as we discussed at the start of the chapter, the IT team of a big organization can add a custom view to the admin site, which can then be used to monitor the health of the different IT systems at the organization, as well as provide the IT team with the ability to quickly look at any urgent alerts that need to be addressed.

Now, the next question we need to answer is, *how can we add a custom view to the admin site?*

As it turns out, adding a new view inside the admin template is quite easy and follows the same approach we used while creating views for our applications, with some minor modifications. So, let's look at how we can add a new view to our Django admin dashboard.

Creating the view function

The first step to adding a new view to the Django application is to create a view function that implements the logic to handle the view. In the previous chapters, we created the view functions inside a separate file known as `views.py`, which was used to hold all our method- and class-based views.

In the context of adding a new view to the Django admin dashboard, to create a new view, we need to define a new view function inside our custom `AdminSite` class. For example, to add a new view that renders a page showing the health of the different IT systems inside an organization, we can create a new view function named `system_health_dashboard()` inside our custom `AdminSite` class implementation, as shown in the following code snippet:

```
class SysAdminSite(admin.AdminSite):
    def system_health_dashboard(self, request):
        # View function logic
```

Inside the view function, we can perform any operations we want in order to generate a view and, finally, use that response to render a template. Inside this view function, there are some important pieces of logic we need to make sure are implemented correctly.

The first one is to set the `current_app` property for the `request` field inside the view function. This is required in order to allow Django's URL resolver inside the templates to correctly resolve the view functions for an application. To set this value inside the custom view function we just created, we need to set the `current_app` property, as shown in the following code snippet:

```
request.current_app = self.name
```

The `self.name` field is automatically populated by Django's `AdminSite` class, and we don't need to initialize it explicitly. With this, our minimal custom view implementation will appear, as shown in the following code snippet:

```
class SysAdminSite(admin.AdminSite):
    def system_health_dashboard(self, request):
        request.current_app = self.name
        # View function logic
```

Accessing common template variables

When creating a custom view function, we might want access to the common template variables, such as `site_header` and `site_title`, in order to render them correctly in the template associated with our view function. As it turns out, this is quite easy to achieve with the use of the `each_context()` method provided by the `AdminSite` class.

The `each_context()` method of the `AdminSite` class takes a single parameter, `request`, which is the current request context, and returns the template variables that are to be inserted in all the admin site templates.

For example, if we wanted to access the template variables inside our custom view function, we could implement code similar to the following code snippet:

```
def system_health_dashboard(self, request):
    request.current_app = self.name
    context = self.each_context(request)
    # view function logic
```

The value returned by the `each_context()` method is a dictionary containing the name of the variable and the associated value.

Mapping URLs for the custom view

Once the view function has been defined, the next step involves mapping this view function to a URL such that a user can access it or allow the other views to link to it. For the views defined inside `AdminSite`, this URL mapping to views is controlled by the `get_urls()` method implemented by the `AdminSite` class. The `get_urls()` method returns the `urlpatterns` list that maps to the `AdminSite` views.

If we want to add a URL mapping for our custom view, the preferred approach includes overriding the implementation of `get_urls()` in our custom `AdminSite` class and adding the URL mapping there. This approach is demonstrated in the following code snippet:

```
class SysAdminSite(admin.AdminSite):
    def get_urls(self):
        base_urls = super().get_urls(). # Get the existing
```

```
set of URLs
# Define our URL patterns for custom views
urlpatterns = [
    path("health_dashboard/",
          self.system_health_dashboard)
]
return base_urls + urlpatterns. # Return the
                                updated mapping
```

The `get_urls()` method is generally called automatically by Django, and there is no need to perform any manual processing on it.

Once this is done, the last step involves making sure that our custom admin view is only accessible through the admin site, and non-admin users should not be able to access it. Let's take a look at how that can be achieved.

Restricting custom views to the admin site

If you followed all the previous headings thoroughly, you will now have a custom `AdminSite` view ready for use. However, there is a small glitch. This view is also directly accessible to any user who is not on the admin site.

To ensure that such a situation does not arise, we need to restrict this view to the admin site. This can be achieved quite simply by wrapping our URL path inside the `admin_view()` call, as shown in the following code snippet:

```
urlpatterns = [
    self.admin_view(path("health_dashboard/",
                         self.system_health_dashboard))
]
```

The `admin_view` function makes sure the path provided to it is restricted just to the admin dashboard and that no user without admin privileges can access it.

Now, let's add a new custom view to our admin site.

Exercise 10.04 – adding custom views to the admin site

In this exercise, you will add a custom view to the admin site, which will render a user profile and will show the user the options to modify their email and add a new profile picture. To build this custom view, follow the steps described:

1. Open the `admin.py` file under the `bookr_admin` directory and add the following imports. These will be required to build our custom view inside the admin site application:

```
from django.template.response import TemplateResponse
from django.urls import path
```

2. Open the `admin.py` file under the `bookr_admin` directory and create a new method named `profile_view`, which takes in a `request` variable as its parameter, inside the `BookrAdmin` class:

```
def profile_view(self, request):
```

Next, inside the method, get the name of the current application and set that in the request context. For this, you can use the `name` property of the class, which is auto-populated by Django. To get this property and set it in your request context, you need to add the following line:

```
request.current_app = self.name
```

Once you have the application name populated to the request context, the next step is to fetch the template variables, which are required to render the contents, such as `site_title` and `site_header`, in the admin templates. For this, leverage the `each_context()` method of the `AdminSite` class, which provides the dictionary of the admin site template variables from the class:

```
context = self.each_context(request)
```

Once you have the data in place, the last step is to return a `TemplateResponse` object, which will render the custom profile template when someone visits the URL endpoint mapped to your custom view:

```
return TemplateResponse(request,
    "admin/admin_profile.html", context)
```

3. With the view function now created, the next step is to make `AdminSite` return the URLs mapping the view to a path inside `AdminSite`. To do this, you need to create a new method called `get_urls()`, which overrides the `AdminSite.get_urls()` method and returns the mapping of your new view. This can be done by first creating a new method named `get_urls()` inside the `BookrAdmin` class you have created for your custom admin site:

```
def get_urls(self):
```

Inside this method, the first thing you need to do is to get the list of the URLs that are already mapped to the admin endpoint. This is a required step; otherwise, your custom admin site will not be able to load any results associated with the model editing pages, logout page, and so on if this mapping is lost. To get this mapping, call the `get_urls()` method of the base class from which the `BookrAdmin` class is derived:

```
urls = super().get_urls()
```

Once the URLs from the base class have been captured, the next step is to create a list of URLs that map our custom view to a URL endpoint in the admin site. To do this, we will create a new list named `url_patterns` and map our `profile_view` method to the `admin_profile` endpoint. This is done by using the `path` utility function from Django, which allows us to map the view function with a string-based API endpoint path:

```
url_patterns = [
    path("admin_profile", self.profile_view)
]
return url_patterns + urls
```

Save the `admin.py` file. It should look like this: https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter10/Exercise10.04/bookr_admin/admin.py.

Note that the value returned by `get_urls()` is a non-modifiable list. In the preceding code snippet, if we try to switch the order between the merging of the `url_patterns` and `urls` objects, the method won't work and the `admin_profile` URL won't be created.

4. Now, with the `BookrAdmin` class configured for the new view, the next step is to create your template for the admin profile page. To do this, create a new file named `admin_profile.html` under the `templates/admin` directory of your project root. Inside this file, first, add an `extends` tag to make sure that you are extending from the default `admin` template:

```
{% extends "admin/index.html" %}
```

This step ensures that all of your admin template style sheets and HTML are available for use inside your custom view template. For example, without having this `extends` tag, your custom view will not show any specific content already mapped to the admin site, such as `site_header`, `site_title`, or any links to log out or go to another page.

Once the `extends` tag is added, add a `block` tag and provide it with the value of `content`. This makes sure that the code you add between the pair of `{% block content %}...{% endblock %}` segments overrides whatever value is present in the `index.html` template that comes prepackaged with the Django admin module:

```
{% block content %}
```

Inside the block tag, add the HTML required to render the profile view that was created in *step 2* of this exercise:

```
<p>Welcome to your profile, {{ username }}</p>
<p>You can do the following operations</p>
<ul>
    <li><a href="#">Change E-Mail Address</a></li>
    <li><a href="#">Add Profile Picture</a></li>
</ul>
{ % endblock %}
```

The file should now look like this: https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter10/Exercise10.04/templates/admin/admin_profile.html.

5. Now, with the preceding steps complete, reload your application server by running `python manage.py runserver localhost:8000` and then visiting `http://localhost:8000/admin/admin_profile`.

When the page opens, you can expect to see something like the following screenshot:

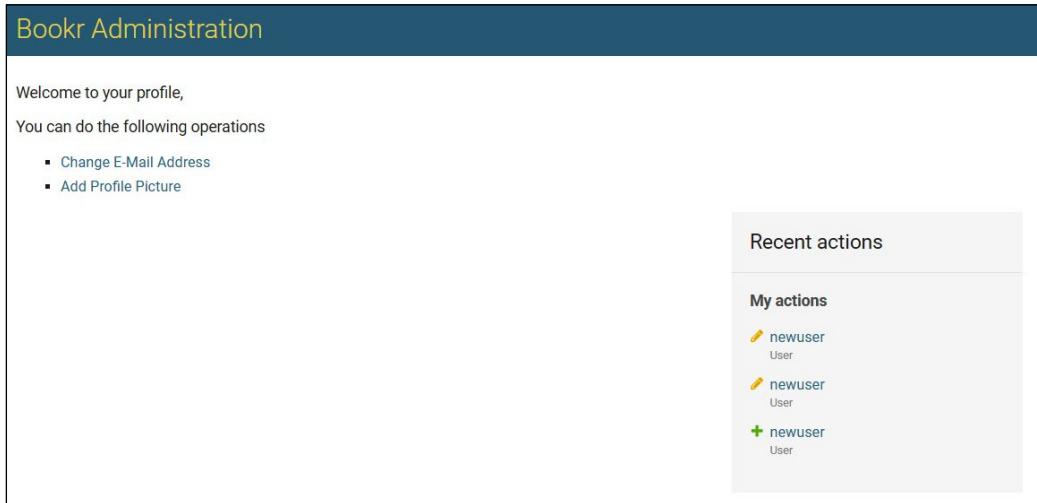


Figure 10.7 – The profile page view on the Bookr Administration site

Note

The view created so far will render just fine, irrespective of whether a user is logged into the admin application.

To make sure that this view is only accessible to logged-in admins, you need to make a small modification inside your `get_urls()` method, which you defined in *step 2* of this exercise.

Inside the `get_urls()` method, modify the `url_patterns` list to look something like the one shown here:

```
url_patterns = [
    path("admin_profile",
        self.admin_view(self.profile_view)),
]
```

In the preceding code, you wrapped your `profile_view` method inside the `admin_view()` method.

The `AdminSite.admin_view()` method causes the view to be restricted to users who are logged in. If a user who is currently not logged into the admin site tries to visit the URL directly, they will be redirected to the login page, and only in the event of a successful login will they be allowed to see the contents of our custom page.

During this exercise, we leveraged our existing understanding of writing views for Django applications, merging it with the context of the `AdminSite` class to build a custom view for our admin dashboard. With this knowledge, we can now move on and add useful functionalities to our Django admin to supercharge its usefulness.

Passing additional keys to templates using template variables

Inside the admin site, the variable values passed to the templates are done so by using template variables. These template variables are prepared and returned by the `AdminSite.each_context()` method.

Now, if there is a value that you want to pass to all the templates of your admin site, you can override the `AdminSite.each_context()` method and add the required fields to the `request` context. Let's look at an example to see how we can achieve this outcome.

Consider the `username` field, which we passed to our `admin_profile` template earlier. If we want to pass it to every template inside our custom admin site, we first need to override the `each_context()` method inside our `BookrAdmin` class, as shown here:

```
def each_context(self, request):
    context = super().each_context(request)
    context['username'] = request.user.username
    return context
```

The `each_context()` method takes a single argument (we're not considering `self` here) of the `HTTPRequest` type, which it uses to evaluate certain other values.

Now, inside our overridden `each_context()` method, we first make a call to the base class `each_context()` method to retrieve the `context` dictionary for the admin site:

```
context = super().each_context(request)
```

Once that is done, the next thing to do is to add our `username` field to `context` and set its value to the value of the `request.user.username` field:

```
context['username'] = request.user.username
```

Once this is done, the last thing that remains is to return this modified context.

Now, whenever a template is rendered by our custom admin site, the template will be passed with this additional `username` variable.

Activity 10.01 – building a custom admin dashboard with a built-in search

In this activity, you will use the knowledge gained about the different aspects of creating a custom admin site to build a custom admin dashboard for Bookr. Inside this dashboard, you will introduce the capability of allowing a user to search for books, by using either the name of the book or the name of the book publisher and allowing the user to modify or delete these book records.

The following steps will help you build a custom admin dashboard and add the ability to search for a book record by the name of the publisher:

1. Create a new application inside the Bookr project named `bookr_admin`, if not created already. This is going to store the logic for our custom admin site.
2. Inside the `admin.py` file under the `bookr_admin` directory, create a new class, `BookrAdmin`, which inherits from the `AdminSite` class of Django's `admin` module.
3. Inside the newly created `BookrAdmin` class from *step 2*, add any customizations for the site title or any other branding components of the admin dashboard.
4. Inside the `apps.py` file under the `bookr_admin` directory, create a new `BookrAdminConfig` class, and inside this new `BookrAdminConfig` class, set the `default site` attribute to the fully qualified module name for our custom admin site class, `BookrAdmin`.
5. Inside the `settings.py` file of your Django project, add the fully qualified path of the `BookrAdminConfig` class created in *step 4* as the first installed application.
6. To register the `Books` model from the `reviews` application inside Bookr, open the `admin.py` file inside the `reviews` directory and make sure that the `Books` model is registered to the admin site by using `admin.site.register(ModelClass)`.

7. To allow users to search for a book by the name of the publisher, inside the `admin.py` file of the `reviews` application, modify the `BookAdmin` class and add to it a property named `search_fields`, which contains `publisher_name` as a field.
8. To get the publisher's name correctly for the `search_fields` property, introduce a new method named `get_publisher` inside the `BookAdmin` class, which will return the name field of the publisher from the `Book` model.
9. Make sure that the `BookAdmin` class is registered as a Model admin class for the `Book` model inside our Django admin dashboard by using `admin.site.register(Book, BookModel)`.

After completing this activity, once you start the application server and visit `http://localhost:8000/admin` and navigate to the `Book` model, you should be able to search for books by using the publisher's name and, in the event of a successful search, see a page that resembles the one shown in the following screenshot:

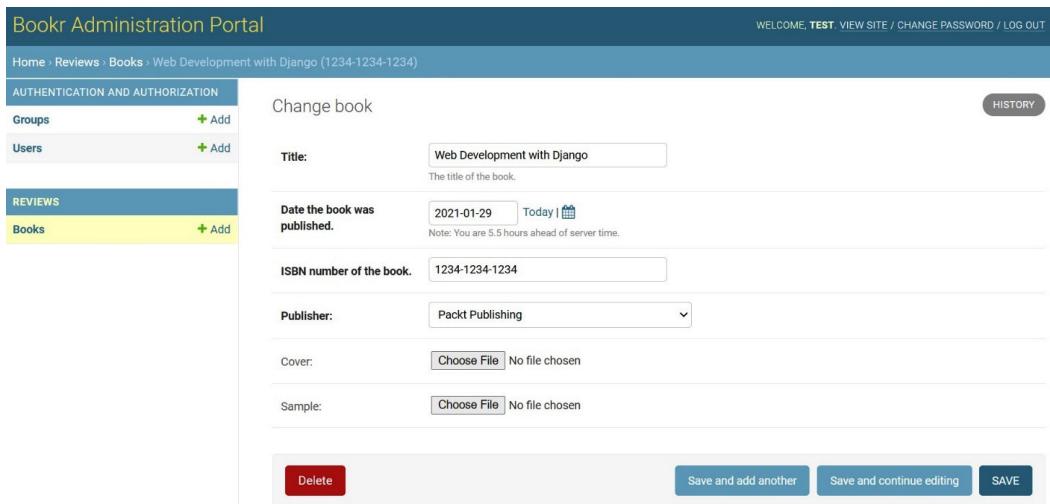


Figure 10.8 – The book editing page inside the Bookr Administration dashboard

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>

Summary

In this chapter, we took a look at how Django allows the customization of its admin site, by providing easy-to-use properties for some of the more general parts of the site, such as title fields, headings, and home links. Beyond this, we learned how to build a custom admin site by leveraging the concepts of object-oriented programming in Python and creating a child class of `AdminSite`.

This functionality was further enhanced by implementing a custom template for the logout page. We also learned how we can supercharge our admin dashboard by adding a new set of views to allow enhanced usage of the dashboard.

As we move on to the next chapter, we will get to build upon what we have learned thus far and extend that knowledge by being introduced to the concept of building our own custom tags and filters for templates, as well as the ability to build our views in an object-oriented style using the concept of class-based views.

11

Advanced Templating and Class-Based Views

In *Chapter 3, URL Mapping, Views, and Templates*, we learned how to build views and create templates in Django. Then, we learned how to use those views to render the templates we built. In this chapter, we will build upon our knowledge of developing views by using **class-based views**, allowing us to write views that can group logical methods into a single entity. This skill comes in handy when developing a view that maps to multiple HTTP request methods for the same **application programming interface (API)** endpoint. With method-based views, we may end up using a lot of the `if-else` conditions to successfully handle the different types of HTTP request methods. In contrast, class-based views allow us to define separate methods for every HTTP request method we want to handle. Then, based on the type of request received, Django takes care of calling the correct method in the class-based view.

Beyond the ability to build views based on different development techniques, Django also comes packed with a powerful templating engine. This engine allows developers to build reusable templates for their web applications. This reusability of the templating engine is further enhanced by using **template tags** and **filters**, which help easily implement commonly used features inside templates, features such as iterating over lists of data, formatting the data in a given style, extracting a piece of text from a variable to display, and overriding the content in a specific block of a template. All these features also expand the reusability of a Django template.

As we go through this chapter, we will look at how we can expand the default set of template filters and template tags provided by Django by leveraging Django's ability to define our own custom template tags and filters. These custom template tags and filters can then be used to implement some common features in a reusable fashion across our web application. For example, while building a user profile badge that can be shown in several places inside a web application, it is better to leverage the ability to write a custom template inclusion tag that just inserts the template of the badge in any of the views we desire, rather than rewriting the entire code for the badge template or by introducing additional complexity to the templates.

In this chapter, we will be covering the following topics:

- Template filters
- Custom template filters
- String filters
- Template tags
- Django views

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter11>

Template filters

While developing templates, developers often just want to change the value of a template variable before rendering it to the user. For example, consider that we are building a profile page for a Bookr user. There, we want to show the number of books the user has read. Below that, we also want to show a table listing the books they have read.

To achieve this, we can pass two separate variables from our view to the HTML template. One can be named `books_read`, which denotes the number of books read by the user. The other can be `book_list`, containing the list of names of the books read by the user, for example:

```
<span class="books_read">You have read {{ books_read }}  
books</span>  
<ul>  
    {% for book in book_list %}  
        <li>{{ book }} </li>  
    {% endfor %}  
</ul>
```

Template filters in Django are simple Python-based functions that accept a variable as an argument (and any additional data in the context of the variable), change its value as per our requirements, and then render the changed value.

Now, the same outcome from writing the previous snippet can also be obtained without the use of two separate variables by using template filters in Django, as follows:

```
<span class="books_read">You have read  
    {{ book_list|length }}</span>  
<ul>  
    {% for book in book_list %}
```

```
<li>{{ book }}</li>
{%
  endfor
%}
</ul>
```

Here, we used the built-in `length` filter provided by Django. The use of this filter causes the length of the `book_list` variable to be evaluated and returned, which is then inserted into our HTML template during rendering.

Like `length`, there are a lot of other template filters that come pre-packaged with Django and that are ready to be used. For example, the `lowercase` filter converts the text to all lowercase format, the `last` filter can be used to return the last item in the list, and the `json_script` filter can be used to output a Python object passed to the template as a JSON value wrapped in a `<script>` tag in your template.

Important note

You can refer to Django's official documentation for the complete list of template filters offered by Django at <https://docs.djangoproject.com/en/3.1/ref/templates/builtins/>.

Now, with the knowledge of how to use template filters, let's jump into understanding how we can write our own custom filters.

Custom template filters

Django supplies a lot of useful filters that we can use in our templates while we are working on our projects. But what if someone wants to format a specific piece of text and render it with different fonts? Or, say, someone wants to translate an error code to a user-friendly error message based on the mapping of the error code in the backend. In these cases, predefined filters do not suffice, and we would like to write our own filter that we can reuse across the project.

Luckily, Django supplies an easy-to-use API that we can use to write custom filters. This API provides developers with some useful decorator functions that can be used to quickly register a Python function as a custom template filter. Once a Python function is registered as a custom filter, a developer can start using the function in templates.

An instance of this template library method is required to access the filters. This instance can be created by instantiating the `Library()` class in Django from Django's `template` module, as shown here:

```
from django import template
register = template.Library()
```

Once the instance is created, we can now use the filter decorator from the template library instance to register our filters.

Creating custom template filters

There are a couple of steps we need to take to create custom template filters. Let's try to understand what these steps are and how they help us with the creation of a custom template filter in the next subsection.

Setting up the directory for storing template filters

It is important to note that when creating a custom template filter or template tag, we need to put them in the `templatetags` directory under the application directory. This requirement arises because Django is internally configured to look for custom template tags and filters when loading a web application. A failure to name the directory `templatetags` will result in Django not loading the custom template filters and tags created by us.

To create this directory, first, navigate to the application folder inside which you want to create custom template filters, and then run the following command in the terminal:

```
mkdir templatetags
```

Once the directory is created, the next step is to create a new file inside the `templatetags` directory to store the code for our custom filters. This can be done by executing the following command inside the `templatetags` directory:

```
touch custom_filter.py
```

Important note

The aforementioned command won't work on Windows. You can, however, navigate to the desired directory and create a new file using Windows Explorer.

Alternatively, this can be done by using the GUI interface provided by PyCharm.

Setting up the template library

Once the file for storing the code for the custom filter is created, we can now start working on implementing our custom filter code. For custom filters to work in Django, they need to be registered to Django's template library before they can be used inside templates. To that end, the first step is to set up an instance of the template library, which will be used to register our custom filters. For this, inside the `custom_filters.py` file we created in the previous section, we first need to import the template module from the Django project:

```
from django import template
```

Once the import is resolved, the next step is to create an instance of the template library by adding the following line of code:

```
register = template.Library()
```

The `Library` class from Django's template module is implemented as a **singleton** class that returns the same object that is only initialized once at the start of the application.

Once the template library instance is set up, we can proceed with implementing our custom filter.

Implementing the custom filter function

Custom filters inside Django are nothing more than simple Python functions that essentially take the following parameters:

- The value on which the filter is being applied (mandatory)
- Any additional parameters (zero or more) that need to be passed to the filter (optional)

These functions need to be decorated with the `filter` attribute from Django's template library instance to behave as template filters. For example, the generic implementation of a custom filter will look like the following:

```
@register.filter
def my_filter(value, arg):
    # Implementation logic of the filter
```

With this, we have learned how to implement a custom filter that can be used inside templates. In the next section, we will learn how to use these custom filters.

Using custom filters inside templates

Once the filter is created, it's simple to start using it inside our templates. To do that, the filter must first be imported into the template. This can be easily done by adding the following line to the top of the template file:

```
{% load custom_filter %}
```

When Django's templating engine parses the template files, the preceding line is automatically resolved by Django to find the correct module specified under the `templatetags` directory. As a consequence, all the filters mentioned inside the `custom_filter` module are automatically made available inside the template.

Using our custom filter inside the template is as simple as adding the following line:

```
{% some_value|generic_filter:"arg" %}
```

Equipped with this knowledge, let's now create our first custom filter.

Exercise 11.01 – Creating a custom template filter

In this exercise, we will write a custom filter named `explode`, which returns a list of strings when provided with a string and a user-supplied separator. For example, consider the following string:

```
names = "john,doe,mark,swain"
```

You will apply the following filter to this string:

```
{% names|explode:"," %}
```

The output after applying this filter should be as follows:

```
["john", "doe", "mark", "swain"]
```

Let's get started with the steps:

1. Create a new application inside the `bookr` project that you can use for demo purposes:

```
python manage.py startapp filter_demo
```

The preceding command will set up a new application inside your Django project.

2. Now, create a new directory named `templatetags` inside your `filter_demo` application directory to store the code for your custom template filters. To create the directory, run the following command from inside the `filter_demo` directory from the terminal app or Command Prompt:

```
mkdir templatetags
```

3. Once the directory is created, create a new file named `explode_filter.py` inside the `templatetags` directory.
4. Open the file and add the following lines to it:

```
from django import template

register = template.Library()
```

The preceding code creates an instance of the Django library that can be used to register our custom filter with Django.

5. Add the following code to implement the `explode` filter:

```
@register.filter
def explode(value, separator):
    return value.split(separator)
```

The `explode` filter takes two arguments; one is `value` on which the filter was used, and the second is `separator`, passed from the template to the filter. The filter will use this separator to convert the string into a list.

6. With the custom filter ready, create a template where this filter can be applied. For this, first, create a new folder named `templates` under the `filter_demo` directory and then create a new file named `index.html` inside it with the following contents:

```
<html>
<head>
    <title>Custom Filter Example</title>
<body>
    {% load explode_filter %}

    {{ names|explode:"," }}
</body>
</html>
```

In the first line, Django's template engine loads the custom filter from the `explode_filter` module so that it can be used inside the templates. To achieve this, Django will look for the `explode_filter` module under the `templatetags` directory, and if found, will load it for use.

In the next line, you pass the `names` variable passed to the template and apply the `explode` filter to it, while also passing in the `,` symbol as a separator value to the filter.

7. Now, with the template created, the next thing is to create a Django view that can render this template and pass the `name` variable to the template. For this, open the `views.py` file and add the following highlighted code:

```
from django.shortcuts import render

def index(request):
    names = "john,doe,mark,swain"
    return render(request, "index.html", {'names': names})
```

The preceding code snippet performs some basic operations. It first imports the `render` helper from the `django.shortcuts` module, which helps render the templates. Once the import is complete, it defines a new view function named `index()`, which renders `index.html`.

8. Now map the view to a URL that can then be used to render the results in the browser. To do this, create a new file named `urls.py` inside the `filter_demo` directory and add the following code to it:

```
from django.urls import path

from . import views

urlpatterns = [
    path('', views.index, name='index')
]
```

9. Add the `filter_demo` application to the project URL mapping. To this end, open `urls.py` in the `bookr` project directory and add the following highlighted line inside `urlpatterns`:

```
urlpatterns = [
    path('filter_demo/', include('filter_demo.urls')),
    ...
]
```

After the changes have been made, the `urls.py` file under the `bookr` project directory should resemble the one shown under <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter11/Exercise11.01/bookr/urls.py>.

10. Finally, add the application under the `INSTALLED_APPS` section under `settings.py` of the `bookr` project:

```
INSTALLED_APPS = [
    ...,
    'filter_demo'
]
```

This requirement arises due to the security guidelines implemented by Django, which require that the application implementing custom filters/tags needs to be added to the `INSTALLED_APPS` section. Once the changes have been made, the `settings.py` file should resemble the file at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter11/Exercise11.01/bookr/settings.py>.

11. To view whether the custom filter works, run the following command:

```
python manage.py runserver localhost:8000
```

Now, navigate to the following page in our browser: http://localhost:8000/filter_demo.

12. This page should appear as shown in *Figure 11.1*:

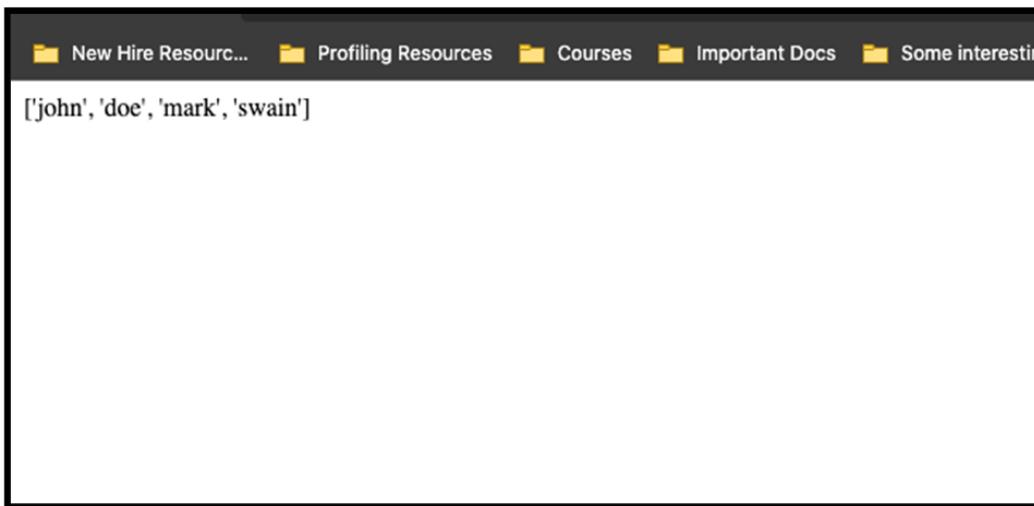


Figure 11.1 – Index page displayed by using the explode filter

With this, we saw how we can quickly create a custom filter inside Django and then use it in our templates. In the next section, let's take a look at another type of filter, namely, string filters, which work solely on string-type values.

String filters

In *Exercise 11.01, Creating a custom template filter*, we built a custom filter, which allowed us to split a provided string with a separator and generate a list from it. This filter can take any kind of variable and split it as a list of values based on a delimiter provided. But what if we wanted to restrict our filter to work only with strings and not with any other type of values, such as integers?

We can use the `stringfilter` decorator provided by Django's template library to develop filters that work only on *strings*. When the `stringfilter` decorator is used to register a Python method as a filter in Django, the framework ensures that the value being passed to the filter is converted to a string before the filter executes. This reduces any potential issues that may arise when non-string values are passed to our filter.

The steps to implement a **string filter** are similar to the ones we followed for building a custom filter with some minor changes.

Remember the `custom_filter.py` file we created in the *Setting up the directory for storing template filters* section? Let's now add a new Python function inside it that will act as our string filter.

Before we can implement a string filter, we first need to import the `stringfilter` decorator, which demarcates a custom filter function as a string filter. You can add this decorator by adding the following `import` statement inside the `custom_filters.py` file:

```
from django.template.defaultfilters import stringfilter
```

Now, to implement our custom string filter, the following syntax can be used:

```
@register.filter
@stringfilter
def generic_string_filter(value, arg):
    # Logic for string filter implementation
```

With this approach, we can build as many string filters as we want and use them just like any other filter.

Template tags

Template tags are a powerful feature of Django's templating engine. They allow developers to build powerful templates by generating HTML by evaluating certain conditions and help avoid the repetitive writing of common code.

One example where we may use template tags is the sign-up/login options in the navigation bar of a website. In this case, we can use template tags to evaluate whether the visitor on the current page is logged in. Based on that, we can render either a profile banner or a sign-up/login banner.

Tags are also a common occurrence while developing templates. For example, consider the following line of code, which we used to import the custom filters inside our templates in the previous section:

```
{% load explode_filter %}
```

This uses a template tag known as `load`, responsible for loading the `explode` filter into the template. Template tags are much more powerful compared to filters. While filters have access only to the values they are operating on, template tags have access to the context of the whole template and hence they can be used to build a lot of complex functionalities inside a template.

Let's now look at the different types of template tags supported by Django and how we can build our own custom template tags.

The types of template tags

Django mainly supports two types of template tags:

- **Simple tags:** These are the tags that operate on the variable data provided (and any additional variables to them) and render in the same template they have been called in. For example, one such use case can include rendering a custom welcome message to the user based on their username or displaying the user's last login time based on their username.
- **Inclusion tags:** These tags take in the provided data variables and generate an output by rendering another template. For example, the tag can take in a list of objects and iterate over them to generate an HTML list.

In the next sections, we will look at how we can create these different types of tags and use them in our application.

Simple tags

Simple tags provide a way for developers to build template tags that take in one or more variables from the template, process them, and return a response. The response returned from the template tag replaces the template tag definition provided inside the HTML template. These kinds of tags can be used to build several useful functionalities, for example, the parsing of dates or displaying any active alerts, if there are any, that we want to show to the user.

The simple tags can be created easily using the `simple_tag` decorator provided by the template library, by decorating the Python method, which should act as a template tag. Now, let's look at how we can implement a custom simple tag using Django's template library.

Creating a simple template tag

Creating simple template tags follows the same conventions we discussed in the *Custom template filters* section, with some subtle differences. First, let's go over understanding how template tags can be created for use in our Django templates.

Setting up the directory

Just like custom filters, custom template tags also need to be created inside the same `templatetags` directory to make them discoverable by Django's templating engine. The directory can be created either directly using the PyCharm GUI or by running the following command inside the application directory where we want to create our custom tags:

```
mkdir templatetags
```

Once this is done, we can now create a new file that will store the code for our custom template tags by using the following command:

```
touch custom_tags.py
```

Important note

The aforementioned command won't work on Windows. You can, however, create a new file using Windows Explorer.

Setting up the template library

Once the directory structure is set up and we have a file in place for keeping the code for our custom template tags, we can now start creating our template tags. But before that, we need to set up an instance of Django's template library as we did earlier. This can be done by adding the following lines of code to our `custom_tag.py` file:

```
from django import template
register = template.Library()
```

Like custom filters, the template library instance is used here to register the custom template tags for use inside Django templates.

Implementing a simple template tag

Simple template tags inside Django are Python functions that can take any number of arguments as we desire. These Python functions need to be decorated with the `simple_tag` decorator from the template library to register those functions as simple template tags. The following snippet of code shows how a simple template tag is implemented:

```
@register.simple_tag
def generic_simple_tag(arg1, arg2):
    # Logic to implement a generic simple tag
```

Next, let's use these simple tags inside the templates.

Using simple tags inside templates

Using simple tags inside Django templates is quite easy. Inside the template file, we need first to make sure that we have the tag imported inside the template by adding the following to the top of the template file:

```
{% load custom_tag %}
```

The preceding statement will load all the tags from the `custom_tag.py` file we defined earlier and make them available inside our template. Then we can use our custom simple tag by adding the following command:

```
{% custom_simple_tag "argument1" "argument2" %}
```

Wasn't that easy?! Now, let's put this knowledge into practice and create our first custom simple tag.

Exercise 11.02 – Creating a custom simple tag

In this exercise, we will create a simple tag that will take in two arguments: the first will be a greeting message, and the second will be the user's name. This tag will print a formatted greeting message. Let's get started with the steps:

1. Following on from the example shown in *Exercise 11.01*, let's reuse the same directory structure to store the code for the simple tag inside. So, first, create a new file named `simple_tag.py` under the `filter_demo` directory. Inside this file, add the following code:

```
from django import template

register = template.Library()

@register.simple_tag
def greet_user(message, username):
    return "{greeting_message}, {user}!!!"
    .format(greeting_message=message, user=username)
```

In this case, you create a new Python method, `greet_user()`, which takes in two arguments, `message`, the message to use for the greeting, and `username`, the name of the user who should be greeted. This method is then decorated with `@register.simple_tag`, indicating that this method is a simple tag and can be used as a template tag in the templates.

2. Now, create a new template that will use your simple tag. For this, create a new file named `simple_tag_template.html` under the `filter_demo/templates` directory and add the following code to it:

```
<html>
<head>
    <title>Simple Tag Template Example</title>
</head>
<body>
    {% load simple_tag %}
    {% greet_user "Hey" username %}
</body>
</html>
```

In the preceding code snippet, we just created a bare-bones HTML page that will use your custom simple tag. The semantics of loading a custom template tag is similar to that of loading a custom template filter and requires the use of a `{% load %}` tag in the template. The process will look for the `simple_tag.py` module under the `templatetags` directory and, if found, will load the tags that have been defined under the module.

The following line shows how we can use the custom template tag:

```
{% greet_user "Hey" username %}
```

In this, we first used Django's tag specifier, `{% %}`, and inside it, the first argument we passed is the name of the tag that needs to be used, followed by the first argument, `Hey`, which is the greeting message, and the second argument, `username`, which will be passed to the template from the view function.

3. With the template created, the next step involves creating a view that will render your template. For this, add the following code in the `views.py` file under the `filter_demo` directory:

```
def greeting_view(request):  
    return render(request, 'simple_tag_template.html',  
    {'username': 'jdoe'})
```

In the preceding code snippet, we created a simple function-based view, which will render your `simple_tag_template` defined in step 2 and pass the `'jdoe'` value to the `username` variable.

4. With the view created, the next step is to map it to a URL endpoint in your application. To do this, open the `urls.py` file under the `filter_demo` directory and add the following inside the `urlpatterns` list:

```
path('greet', views.greeting_view, name='greeting')
```

With this, `greeting_view` is now mapped to the `/greet` URL endpoint for your `filter_demo` application.

The final set of changes should resemble the ones at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter11/Exercise11.02>.

5. To see the custom tag in action, start your web server by running the following command:

```
python manage.py runserver localhost:8000
```

After visiting `http://localhost:8000/filter_demo/greet` in the browser, you should see the following page:

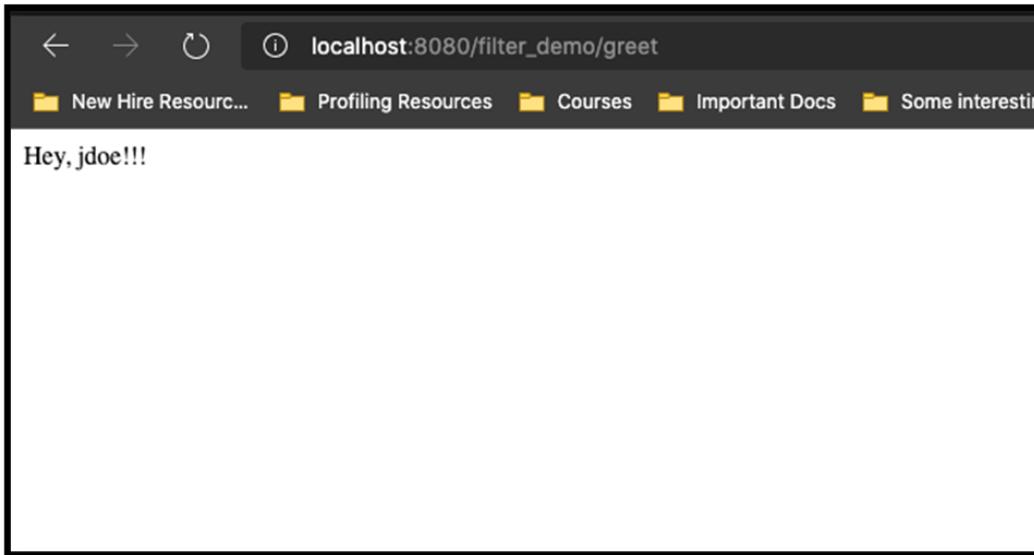


Figure 11.2 – Greeting message generated with the help of the custom simple tag

With this, we created our first custom template tag and used it successfully to render our template, as shown in *Figure 11.2*. Now, let's look at another important aspect of simple tags, which is associated with passing the context variables available in the template to the template tag.

Passing the template context in a custom template tag

In the previous exercise, we created a simple tag to which we passed two arguments, the greeting message and the username. But what if we wanted to pass a large number of variables to the tag? Or simply, what if we did not want to pass the user's username explicitly to the tag?

There are times when developers would like to have access to all the variables, and data that is present in the template to be available inside the custom tag. Fortunately for us, this is easy to implement.

Using our previous example of the `greet_user` tag, let's create a new tag named `contextual_greet_user` and see how we can pass the data available in the template directly to the tag instead of passing it manually as an argument.

The first modification we need to make is to modify our decorator to look like the following:

```
@register.simple_tag(takes_context=True)
```

With this, we tell Django that when our `contextual_greet_user` tag is used, Django should also pass it the template context, which has all the data passed from the view to the template. With this addition done, the next thing we need to do is to change our `contextual_greet_user` implementation to accept the added context as an argument. The following code shows the modified form of the `contextual_greet_user` tag, which uses our template context to render a greeting message:

```
@register.simple_tag(takes_context=True)
def contextual_greet_user(context, message):
    username = context['username']
    return "{greeting_message}, {user}"
        .format(greeting_message=message, user=username)
```

In the preceding code example, we can see how the `contextual_greet_user()` method was modified to accept the passed context as the first argument, followed by the greeting message passed by the user.

To leverage this modified template tag, all we need to do is to change our call to the `contextual_greet_user` tag inside `simple_tag_template.html` under `filter_demo` to look like this:

```
{% contextual_greet_user "Hey" %}
```

Then, when we reload our Django web application, the output at `http://localhost:8000/filter_demo/greet` should look the same as what was shown in *step 5 of Exercise 11.02 – creating a custom simple tag*.

With this, we learned how we can build a simple tag and handle passing the template context to the tag. Now, let's take a look at how we can build an inclusion tag that can be used to render data in a certain format, as described by another template.

Inclusion tags

Simple tags allow us to build tags that accept one or more input variables, do some processing on them, and return an output. This output is then inserted where the simple tag was used.

But what if we wanted to build tags that, instead of returning text output, return an HTML template, which can then be used to render the parts of the page? For example, many web applications allow users to add custom widgets to their profiles. These individual widgets can be built as an inclusion tag and rendered independently. This kind of approach keeps the code for the base page template and the individual templates separate and hence allows for easy reuse and refactoring.

Developing custom inclusion tags is a similar process to how we develop our simple tags. This involves the use of the `inclusion_tag` decorator provided by the template library. So, let's take a look at how we can do it.

Implementing inclusion tags

Inclusion tags are those tags used for rendering a template as a response to their usage inside a template. These tags can be implemented in a similar manner to how other custom template tags are implemented, with some minor modifications.

Inclusion tags are also simple Python functions that can take multiple parameters, where each parameter maps to an argument passed from the template where the tag was called. These tags are decorated using the `inclusion_tag` decorator from Django's template library. The `inclusion_tag` decorator takes a single parameter, the name of the template, which should be rendered as a response to the processing of the inclusion tag.

A generic implementation of an inclusion tag will look like the one shown in the following code snippet:

```
@register.inclusion_tag('template_file.html')
def my_inclusion_tag(arg):
    # logic for processing
    return {'key1': 'value1'}
```

Notice the return value in this case. An inclusion tag is supposed to return a dictionary of values that will be used to render the `template_file.html` file specified as an argument in the `inclusion_tag` decorator.

Using an inclusion tag inside a template

An inclusion tag can easily be used inside a template file. This can be done by first importing the tag as follows:

```
{% load custom_tags %}
```

And then we use the tag like any other tag:

```
{% my_inclusion_tag "argument1" %}
```

The response of the rendering of this tag will be a sub-template that will be rendered inside our primary template where the inclusion tag was used.

Exercise 11.03 – Building a custom inclusion tag

In this exercise, we are going to build a custom `inclusion tag`, which will render the list of books read by a user:

1. For this exercise, you will continue to use the same demo folders as in earlier exercises. First, create a new file named `inclusion_tag.py` under the `filter_demo/templatetags` directory and write the following code inside it:

```
from django import template

register = template.Library()

@register.inclusion_tag('book_list.html')
def book_list(books):
    book_list = [book_name for book_name, book_author
                 in books.items()]
    return {'book_list': book_list}
```

The `@register.inclusion_tag` decorator is used to mark the method as a custom inclusion tag. This decorator takes the template's name as an argument that should be used to render the data returned by the tag function.

After the decorator, we defined a function that implements the logic of our custom inclusion tag. This function takes a single argument called `books`. This argument will be passed from the template file and contain a list of books the reader has read (in the form of a Python dictionary). Inside the definition, we converted the dictionary into a Pythonic list of book names. The key in the dictionary is mapped to the name of the book, and the value is mapped to the author:

```
books_list = [book_name for book_name, book_author in books.
              items()]
```

Once the list is formed, the following code returns the list as a context for the template passed to the inclusion tag (in this example, `book_list.html`):

```
return {'book_list': books_list}
```

The value returned by this method will be passed by Django to the `book_list.html` template, and the contents will then be rendered.

2. Next, create the actual template, which will contain the rendering structure for the template tag. For this, create a new template file, `book_list.html`, under the `filter_demo/templates` directory, and add the following content to it:

```
<ul>
    {% for book in book_list %}
        <li>{{ book }}</li>
    {% endfor %}
</ul>
```

Here, in the new template file, we created an unordered list that will hold the list of books a user has read. Next, using the `for` template tag, we iterate over the values within `book_list` that will be provided by the custom template function:

```
{% for book in book_list %}
```

This iteration results in the creation of several list items, as defined by the following:

```
<li>{{ book }}</li>
```

The list item is generated with the contents from `book_list` and passed to the template. The `for` tag executes as many times as the number of items present in `book_list`.

- With the template defined for the `book_list` tag, modify the existing greeting template to make this tag available inside it and use it to show a list of books that the user has read. For this, modify the `simple_tag_template.html` file under the `filter_demo/templates` directory and change the code to look as follows:

```
<html>
<head>
    <title>Simple Tag Template Example</title>
</head>
<body>
    {% load simple_tag inclusion_tag %}
    {% greet_user "Hey" username %}
    <br />
    <span class="message">You have read the following
        books till date</span>
    {% book_list books %}
</body>
</html>
```

In this snippet, the first thing you did was load the `inclusion_tag` module by writing the following:

```
{% load simple_tag inclusion_tag %}
```

Once the tag is loaded, you can now use it anywhere in the template. To use it, you added the `book_list` tag in the following format:

```
{% book_list books %}
```

This tag takes a single argument, which is a dictionary of books, inside which the key is the book title, and the value of the key is the author of the book.

- With the template now modified, the final step involves passing the required data to the template. To achieve this, modify `views.py` in the `filter_demo` directory and change the greeting view function to look like this:

```
def greeting_view(request):
    books = {
```

```

        "The night rider": "Ben Author",
        "The Justice": "Don Abeman"
    }
    return render(request, 'simple_tag_template.html',
    {'username': 'jdoe', 'books': books})

```

Here, we modified the `greeting_view` function to add a dictionary of books and their authors, and then we passed it to the `simple_tag_template` context.

Once the changes have been made, the files should resemble the ones hosted at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter11/Exercise11.03>.

- With the preceding changes implemented, it's time to render the modified template. To do this, restart your Django application server by running the following command:

```
python manage.py runserver localhost:8080
```

- Navigate to `http://localhost:8080/greet`, which should now render a page similar to the following screenshot:

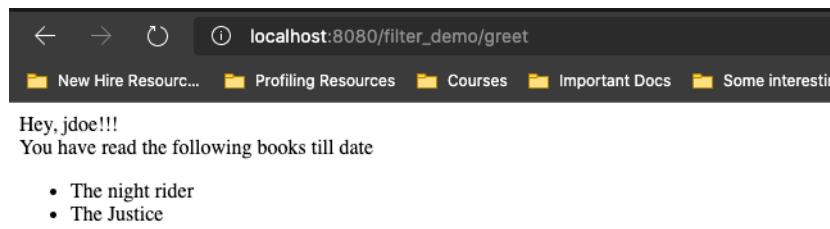


Figure 11.3 – List of books read by a user when they visit the greeting endpoint

The page shows the list of books read by a user when they visit the greeting endpoint. The list you see on the page is rendered using inclusion tags. The template for listing these books is created separately first, and then, using the inclusion tag, it is added to the page.

With this, we now have the foundations on which we can build highly complex template filters or custom tags that can be helpful in the development of the projects we want to work on.

Now, let's take a look at Django views and dive into a new territory of class-based views provided by Django to help us leverage the power of object-oriented programming that allows the reuse of code for the rendering of a view.

Django views

To recall, a view in Django is a piece of Python code that allows a request to be taken in, performs an action based on the request, and then returns a response to the user, hence forming an important part of our Django applications.

Inside Django, we have the option of building our views by following two different methodologies, one of which we have already seen in the preceding examples and is known as function-based views, while the other one, which we will be covering soon, is known as class-based views:

- **Function-based views (FBVs):** FBVs inside Django are nothing more than generic Python functions that are supposed to take an `HTTPRequest` type object as their first positional parameter and return an `HTTPResponse` type object, which corresponds to the action the view wants to perform once the request is processed by it. In the preceding exercise, `index()` and `greeting_view()` were examples of FBVs.
- **Class-based views (CBVs):** CBVs are views that closely adhere to the Python object-oriented principles and allow mapping of the view calls in a class-based representation. These views are specialized in nature, and a given type of CBV performs a specific operation. The benefits that CBVs provide include easy extensibility of the view and the reuse of code, which may turn out to be a complex task with FBVs.

Now, with the basic definitions clear and with knowledge of FBVs already in our arsenal, let's look at CBVs and see what they have in store for us.

Class-based views

Django provides different ways in which developers can write views for their applications. One way is to map a Python function to act as a view function to create FBVs. Another way of creating views is to use Python object instances (based on top of Python classes). These are known as CBVs. An important question that arises is, what is the need for a CBV when we can already create views using the FBV approach?

When creating FBVs, the idea is that, at times, we may be replicating the same logic again and again, for example, the processing of certain fields or logic for handling certain request types. Although it is completely possible to create logically separate functions that handle a particular piece of logic, the task becomes difficult to manage as the complexity of the application increases.

This is where CBVs come in handy, as they abstract away implementation of the common repetitive code that we need to write to handle certain tasks, such as the rendering of templates, while also making it easy to reuse pieces of code through the use of inheritance and mix-ins. For example, the following code snippet shows the implementation of a CBV:

```
from django.http import HttpResponse
from django.views import View

class IndexView(View):

    def get(self, request):
        return HttpResponse("Hey there!")
```

In the preceding example, we built a simple CBV by inheriting from the built-in view class, which Django provides.

Using these CBVs is also quite easy. For example, let's say we wanted to map `IndexView` to a URL endpoint in our application. In this case, all we need to do is to add the following line to our `urlpatterns` list inside the `urls.py` file of the application:

```
urlpatterns = [
    path('my_path', IndexView.as_view(), name='index_view')
]
```

In this, as we can observe, we used the `as_view()` method of the CBV we created. Every CBV implements the `as_view()` method, which allows the view class to be mapped to a URL endpoint by returning the instance of the view controller from the view class.

Django provides a couple of built-in CBVs that provide the implementation of a lot of common tasks, such as how to render a template or how to process a particular request. The built-in CBVs help avoid rewriting code from scratch when handling basic functionality, thereby enabling the reusability of code. Some of these in-built views include the following:

- `View`: This is the base class for all CBVs available in Django that allows a custom CBV to be written with all the features provided and overridable. A user can implement their own definitions for different HTTP Request methods, such as GET, POST, PUT, and DELETE, and the view will automatically delegate the call to the method responsible for handling the request based on the type of request received.
- `TemplateView`: This is a view that can be used to render a template based on the parameters for the template data provided in the URL of the call. This allows developers to easily render a template without writing any logic related to how the rendering should be handled.
- `RedirectView`: This is a view that can automatically redirect a user to the correct resource based on the request they have made.

- `DetailView`: This is a view that is mapped to a Django model and can be used to render the data obtained from the model using a template of choice.

The preceding views are just some of the built-in views that Django provides by default, and we will cover more of them as we move through the chapter.

Now, to better understand how CBVs work inside Django, let's try to build our first CBV.

Exercise 11.04 – Creating a book catalog using a CBV

In this exercise, we will create a class-based form view that will help build a book catalog. This catalog will consist of the name of the book and the name of the author of the book:

1. To get started, create a new application inside our `bookr` project and name it `book_management`. This can be done by simply running the following command:

```
python manage.py startapp book_management
```

2. Now, before building the book catalog, you first need to define a Django model that will help you store the records inside the database. To do this, open the `models.py` file under the `book_management` app you just created and define a new model named `Book`, as shown here:

```
from django.db import models

class Book(models.Model):
    name = models.CharField(max_length=255)
    author = models.CharField(max_length=50)
```

The model contains two fields, the name of the book and the name of the author.

3. With the model in place, migrate the model to your database such that you can start storing your data inside the database.
4. Once all the preceding steps are complete, add your `book_management` application to the `INSTALLED_APPS` list such that Django can discover it and you can use your model properly. For this, open the `settings.py` file under the `bookr` directory and add the following code at the final position in the `INSTALLED_APPS` section:

```
INSTALLED_APPS = [
    ...,
    'book_management'
]
```

5. After the changes have been made, the `settings.py` file should look like this: <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter11/Exercise11.04/bookr/settings.py>.

6. Migrate your model to the database by running the following two commands. These will first create a Django migrations file and then create a table in your database:

```
python manage.py makemigrations
python manage.py migrate
```

7. Now, with the database model in place, let's create a new form that we will use to capture information pertaining to the books, such as the book title, author, and ISBN. For this, create a new file named `forms.py` under the `book_management` directory and add the following code inside it:

```
from django import forms

from .models import Book
class BookForm(forms.ModelForm):
    class Meta:
        model = Book
        fields = ['name', 'author']
```

In the preceding code snippet, we first imported Django's `forms` module, which will allow us to easily create forms and also provide the form's rendering capability. The next line imports the model that will store the data for the form:

```
from django import forms
from .models import Book
```

In the next line, we created a new class named `BookForm`, which inherits from `ModelForm`. This is nothing but a class that maps the fields of the model to the form. To successfully achieve this mapping between the model and the form, we defined a new subclass named `Meta` under the `BookForm` class and set the attribute `model` to point to the `Book` model and the attribute `fields` to the list of fields that you want to display in the form:

```
class Meta:
    model = Book
    fields = ['name', 'author']
```

This allows `ModelForm` to render the correct form of HTML when expected to do so. The `ModelForm` class provides a built-in `Form.save()` method, which, when used, writes the data in the form to the database and helps avoid having to write redundant code.

8. Now that you have both your model and the form ready, go ahead and implement a view that will render the form and accept input from the user. For this, open `views.py` under the `book_management` directory and add the following lines of code to the file:

```
from django.http import HttpResponseRedirect
from django.views.generic.edit import FormView
from django.views import View
```

```
from .forms import BookForm

class BookRecordFormView(FormView):
    template_name = 'book_form.html'
    form_class = BookForm
    success_url = '/book_management/entry_success'

    def form_valid(self, form):
        form.save()
        return super(BookRecordFormView)
            .form_valid(form)

class FormSuccessView(View):
    def get(self, request, *args, **kwargs):
        return HttpResponse("Book record saved
                            successfully")
```

In the preceding code snippet, we created two major views, `BookRecordFormView`, which is also responsible for rendering the book catalog entry form, and `FormSuccessView`, which you will use to render the success message if the form data is saved successfully. Let's now look at both views individually and understand what we are doing.

First, we created a new view named the `BookRecordFormView` CBV, which inherits from `FormView`:

```
class BookRecordFormView(FormView)
```

The `FormView` class allows us to easily create views that deal with forms. To this class, we need to provide certain parameters, such as the name of the template it will render to show the form, the form class that it should use to render the form, and the success URL to redirect to when the form is processed successfully:

```
template_name = 'book_form.html'
form_class = BookForm
success_url = '/book_management/entry_success'
```

The `FormView` class also provides a `form_valid()` method, which is called when the form successfully finishes the validation. Inside the `form_valid()` method, we can decide what we want to do. For our use case, when the form validation completes successfully, we first call the `form.save()` method, which persists the data for our form into the database, and then call the base class `form_valid()` method, which will cause the form view to redirect to the successful URL in the event that form validation was a success:

```
def form_valid(self, form):
    form.save()
    return super().form_valid(form)
```

Important note

The `form_valid()` method should always return an `HttpResponse` object.

This completes the implementation of `BookRecordFormView`. The next thing we have to do is to build a view named `FormSuccessView`, which we will use to render the success message once the data is saved successfully for the book record form we just created. This is done by creating a new view class named `FormSuccessView`, which inherits from the view base class of Django CBVs:

```
class FormSuccessView(View)
```

Inside this class, we override the `get()` method, which will be rendered when the form is saved successfully. Inside the `get()` method, we render a simple success message by returning a new `HttpResponse`:

```
def get(self, request, *args, **kwargs):
    return HttpResponse("Book record saved
                        successfully")
```

9. Now, create a template that will be used to render the form. For this, create a new `templates` folder under the `book_management` directory and a new file named `book_form.html`. Add the following lines of code inside the file:

```
<html>
<head>
    <title>Book Record Insertion</title>
</head>
<body>
    <form method="POST">
        {% csrf_token %}
        {{ form.as_p }}
        <input type="submit" value="Save record" />
    </form>
</body>
</html>
```

Inside this code snippet, two important things need to be discussed.

The first is the use of the `{% csrf_token %}` tag. This tag is inserted to prevent the form from running into **cross-site request forgery (CSRF)** attacks. The `csrf_token` tag is one of the built-in template tags provided by Django to avoid such attacks. It does so by generating a unique token for every form instance that is rendered.

The second is the use of the `{{ form.as_p }}` template variable. The data for this variable is provided by our `FormView`-based view automatically. The `as_p` call causes the form fields to be rendered inside the `<p></p>` tags.

- With the CBVs now built, go ahead and map them to URLs so that you can start using them to add new book records. To do this, create a new file named `urls.py` under the `book_management` directory and add the following code to it:

```
from django.urls import path

from .views import BookRecordFormView, FormSuccessView

urlpatterns = [
    path('new_book_record',
         BookRecordFormView.as_view(),
         name='book_record_form'),
    path('entry_success', FormSuccessView.as_view(),
         name='form_success')
]
```

Most parts of the preceding snippet are similar to the ones you wrote earlier, but there is one thing different in the way we map our CBVs under the URL patterns. When using CBVs, instead of directly adding the function name, we use the class name and its `as_view` method, which maps the class object to the view. For example, to map `BookRecordFormView` as a view, we will use `BookRecordFormView.as_view()`.

- With the URLs added to our `urls.py` file, the next thing is to add our application URL mapping to our `bookr` project. To do this, open the `urls.py` file under the `bookr` application and add the following line to `urlpatterns`:

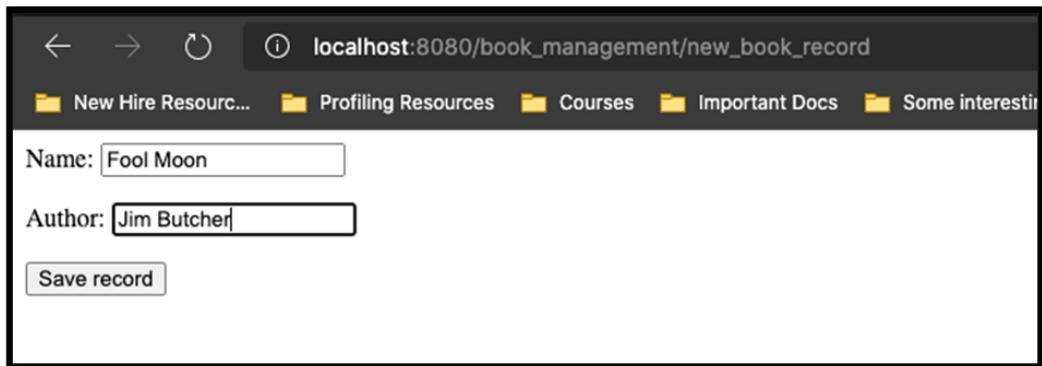
```
urlpatterns = [
    path('book_management/',
         include('book_management.urls')),
    ...
]
```

- Now, start your development server by running the following command:

```
python manage.py runserver localhost:8080
```

13. Then, visit `http://localhost:8080/book_management/new_book_record`.

If everything works fine, you will see a page as shown here:

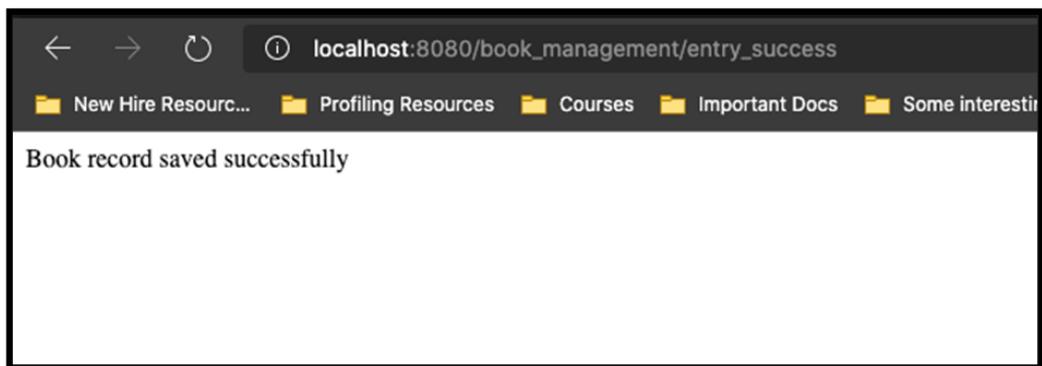


Name:

Author:

Figure 11.4 – The view for adding a new book to the database

Upon clicking **Save record**, your record will be written to the database, and the following page will appear:



Book record saved successfully

Figure 11.5 – The template is rendered when the record is successfully inserted

With this, we have created our own CBV, which allows us to save records for new books. With our knowledge of CBVs in tow, let's now look at how we can perform **create, read, update, and delete (CRUD)** operations with the help of CBVs.

CRUD operations with CBVs

While working with Django models, one of the most common patterns we run into involves creating, reading, updating, and deleting objects that are stored inside our database. The Django admin interface allows us to achieve these CRUD operations easily, but what if we wanted to build custom views to give us the same capability?

As it turns out, Django's CBVs allow us to achieve this quite easily. All we need to do is to write our custom CBVs and inherit them from the built-in base classes provided by Django. Building on our existing example of book record management, let's see how we can build CRUD-based views in Django.

The Create view

To build a view that helps in object creation, we'll need to open the `view.py` file under the `book_management` directory and add the following lines of code to it:

```
from django.views.generic.edit import CreateView
from .models import Book

class BookCreateView(CreateView):
    model = Book
    fields = ['name', 'author']
    template_name = 'book_form.html'
    success_url = '/book_management/entry_success'
```

With this, we created `CreateView` for the book resource. Before we can use it, we will need to map it to a URL. To do this, we can open the `urls.py` file and add the following entry under the `urlpatterns` list:

```
urlpatterns = [
    ...,
    path('book_record_create', BookCreateView.as_view(),
         name='book_create')
]
```

Now, when we visit http://localhost:8080/book_management/book_record_create, we will be greeted with the following page:

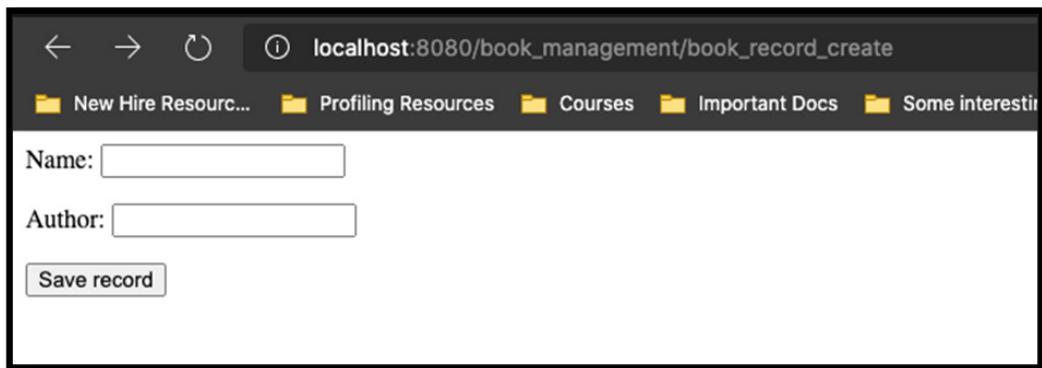


Figure 11.6 – A view to insert a new book record based on the Create view

This looks similar to the one we got when using the form view. On filling in the data and clicking **Save record**, Django will save the data to the database.

The Read view

In this view, we want to see a list of records that our database stores for the books. To achieve this, we will build a view named `DetailView`, which will render details about the book we are requesting. To build this view, we can add the following lines of code to our `views.py` file under the `book_management` directory:

```
from django.views.generic import DetailView

class BookRecordDetailView(DetailView):
    model = Book
    template_name = 'book_detail.html'
```

In the preceding code snippet, we created `DetailView`, which will help us to render the details pertaining to the book ID we are asking for. `DetailView` internally queries our database model with the book ID we provide to it and, if a record is found, renders the template with the data stored inside the record by passing it as an object variable inside the template context.

Once this is done, the next step is to create the template for our book details. For this, we'll need to create a new template file named `book_detail.html` under our `templates` directory inside the `book_management` application with the following contents:

```
<html>
<head>
```

```
<title>Book List</title>
</head>
<body>
    <span>Book Name: {{ object.name }}</span><br />
    <span>Author: {{ object.author }}</span>
</body>
</html>
```

Now, with the template created, the last thing we need to do is to add a URL mapping for the Detail view. This can be done by appending the following to the `urlpatterns` list under the `urls.py` file of the `book_management` application:

```
path('book_record_detail/<int:pk>',
      BookRecordDetailView.as_view(), name='book_detail')
```

Now, with all of this configured, if we now open `http://localhost:8080/book_management/book_record_detail/2`, we will get to see the details about our book, as shown here:

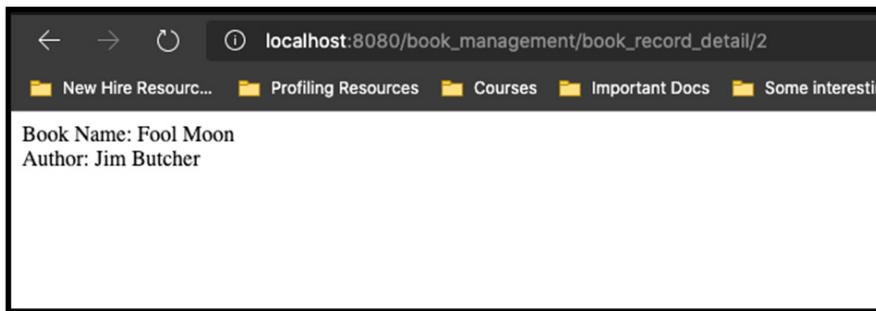


Figure 11.7 – The view rendered when we try to access a previously stored book record

With the preceding examples, we just enabled CRUD operations for our `Book` model, all while using CBVs.

The Update view

In this view, we want to update the data for a given record. To do this, we need to open the `view.py` file under the `book_management` directory and add the following lines of code to it:

```
from django.views.generic.edit import UpdateView

from .models import Book

class BookUpdateView(UpdateView):
    model = Book
```

```
fields = ['name', 'author']
template_name = 'book_form.html'
success_url = '/book_management/entry_success'
```

In the preceding code snippet, we used the built-in `UpdateView` template, which allows us to update the stored records. Therefore `fields` attribute here should take in the name of the fields that we would like to allow the user to update.

Once the view is created, the next step is to add the URL mapping. To do this, open the `urls.py` file under the `book_management` directory and add the following lines of code:

```
urlpatterns = [
    path('book_record_update/<int:pk>',
         BookUpdateView.as_view(), name='book_update')
]
```

In this example, we appended `<int :pk>` to the URL field. This signifies the field input we will have to retrieve the record for. Inside Django models, Django inserts a primary key of the integer type, which uniquely identifies the records. Inside the URL mapping, this is the field that we asked to insert.

Now, when we try to open `http://localhost:8080/book_management/book_record_update/1`, it should show us a record of the first record that we inserted into our database and allow us to edit it:

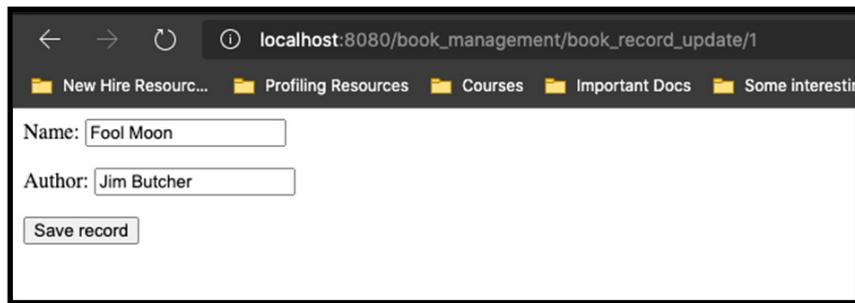


Figure 11.8 – The view displaying the book record update template based on the Update view

The Delete view

The Delete view, as the name suggests, is a view that deletes the record from our database. To implement such a view for our `Book` model, we will need to open the `views.py` file under the `book_management` directory and add the following code snippet to it:

```
from django.views.generic.edit import DeleteView
from .models import Book
```

```
class BookDeleteView(DeleteView):  
    model = Book  
    template_name = 'book_delete_form.html'  
    success_url = '/book_management/delete_success'
```

With this, we just created a Delete view for our book records. As we can see, this view uses a different template where all we would like to confirm from the user is whether they really want to delete the record or not. To achieve this, you can create a new template file, `book_delete_form.html`, and add the following code to it:

```
<html>  
<head>  
    <title>Delete Book Record</title>  
</head>  
<body>  
    <p>Delete Book Record</p>  
    <form method="POST">  
        {%- csrf_token %}  
        Do you want to delete the book record?  
        <input type="submit" value="Delete record" />  
    </form>  
</body>  
</html>
```

Important note

For the code related to the `/delete_success` endpoint, take a look at the code files under the Chapter 11 directory on the book's accompanying GitHub repository.

Then we can add a mapping for our delete view by modifying the `urlpatterns` list inside the `urls.py` file under the `book_management` directory as follows:

```
urlpatterns = [  
    ...,  
    path('book_record_delete/<int:pk>',  
         BookDeleteView.as_view(), name='book_delete')  
]
```

Now, when visiting `http://localhost:8080/book_management/book_record_delete/1`, we should be greeted with the following page:

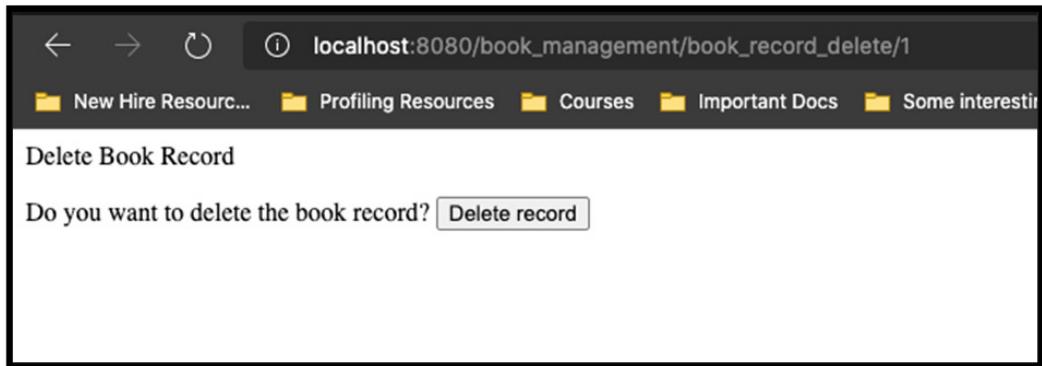


Figure 11.9 – The Delete Book Record view based on the Delete view class

Upon clicking the **Delete record** button, the record will be deleted from the database, and the **Deletion success** page will be rendered.

With our knowledge of the CRUD views and template filters and tags, let's look into applying it to solve an implementation activity.

Activity 11.01 – Rendering details on the user profile page using inclusion tags

In this activity, you will create a custom inclusion tag that helps to develop a user profile page that renders not only the users' details but also lists the books they have read.

The following steps should help you to complete this activity successfully:

1. Create a new `templatetags` directory under the `reviews` application inside the `bookr` project to provide a place where you can create your custom template tags.
2. Create a new file named `profile_tags.py`, which will store the code for your inclusion tag.
3. Inside the `profile_tags.py` file, import Django's template library and use it to initialize an instance of the template library class.
4. Import the `Review` model from the `reviews` application to fetch the reviews written by a user. This will be used to filter the reviews for the current user to render on the user profile page.
5. Next, create a new Python function named `book_list`, which will contain the logic for your inclusion tag. This function should only take a single parameter, the username of the currently logged-in user.

6. Inside the body of the `book_list` function, add the logic for fetching the reviews for this user and extract the name of the books this user has read. Assume that a user has read all those books for which they have provided a review.
7. Decorate this `book_list` function with the `inclusion_tag` decorator and provide it with a template name `book_list.html`.
8. Create a new template file named `book_list.html`, which was specified to the inclusion tag decorator in step 7. Inside this file, add code to render a list of books. This can be done by using a `for` loop construct and rendering HTML list tags for every item inside the list provided.
9. Modify the existing `profile.html` file under the `templates` directory, which will be used to render the user profile. Inside this template file, include the custom template tag and use it to render the list of books read by the user.

Once you have completed all these steps, starting the application server and visiting the user profile page should render a page that is similar to the one shown in *Figure 11.10*:

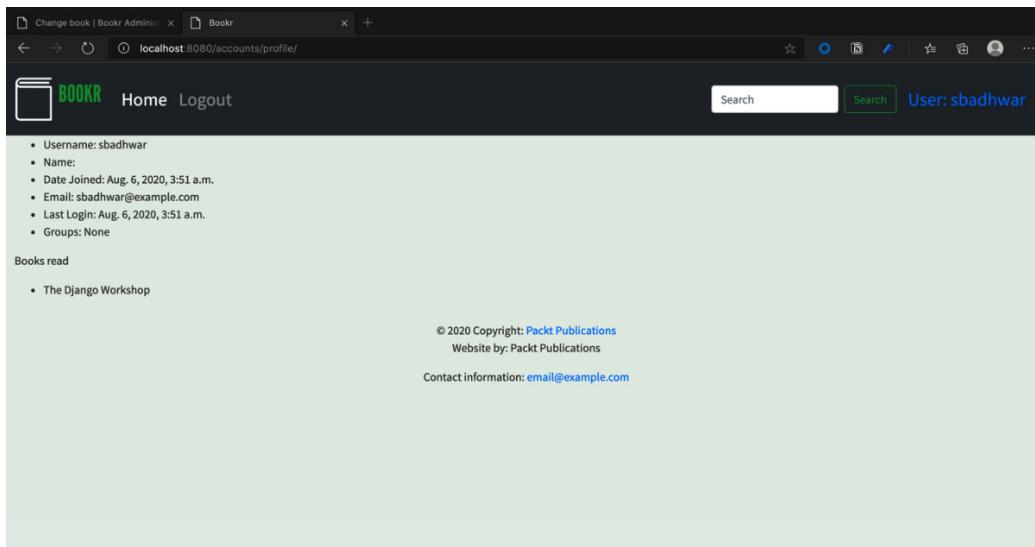


Figure 11.10 – The user profile page with the list of books read by the user

Important note

The solution to this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we learned about the advanced templating concepts in Django and understood how we can create custom template tags and filters to fit a myriad of use cases and support the reusability of components across the application. We then examined how Django provides us with the flexibility to implement FBVs and CBVs to render our responses.

While exploring CBVs, we learned how they can help us avoid code duplication and leverage the built-in CBVs to render forms that save data, help us update existing records, and implement CRUD operations on our database resources.

As we move to the next chapter, we will now utilize our knowledge of building CBVs to work on implementing REST APIs in Django. This will allow us to perform well-defined HTTP operations on our data inside our Bookr application without maintaining any state inside the application.

12

Building a REST API

In the previous chapter, we learned about templates and class-based views. These concepts greatly help expand the range of functionalities we can provide to the user on the frontend (in their web browser). However, that is not sufficient to build a modern web application. Web apps typically have the frontend built with an entirely separate library, such as ReactJS or AngularJS. These libraries provide powerful tools for building dynamic user interfaces but do not communicate directly with our backend Django code or database. The frontend code simply runs in the web browser and does not have direct access to any data on our backend server. Therefore, we need to create a way for these applications to “talk” to our backend code. One of the best ways to do this in Django is by using REST APIs.

APIs are used to facilitate interaction between different pieces of software, and they communicate using **Hypertext Transfer Protocol (HTTP)**. This is the standard protocol for communication between servers and clients and is fundamental to transferring information on the web. APIs receive requests and send responses in HTTP format.

This chapter introduces **REST APIs** and **Django REST framework (DRF)**. You will start by implementing a simple API for the Bookr project. Next, you will learn about the serialization of model instances, which is a key step in delivering data to the frontend of Django applications. Finally, you will explore different types of API views, including both functional and class-based types.

In our use case in this chapter, an API will help facilitate interaction between our Django backend and our frontend JS code. For example, imagine that we want to create a frontend application that allows users to add new books to the Bookr database. The user’s web browser would send a message (an HTTP request) to our API to say that they want to create an entry for a new book and perhaps include some details about the book in that message. Our server would send back a response to report on whether the book was successfully added or not. The web browser would then be able to display the outcome of their action to the user.

In this chapter, we will cover the following topics:

- Understanding REST APIs
- Django REST framework

By the end of this chapter, you will be able to implement custom API endpoints, including token-based authentication.

Technical requirements

All the code files of this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter12>.

Understanding REST APIs

Most modern web APIs can be classified as **REST** APIs. REST APIs are simply a type of API that focuses on communicating and synchronizing the *state* of objects between the database server and frontend client.

For example, imagine that you are updating your details on a website for which you are signed in to your account. When you go to the account details page, the web server tells your browser about the various details attached to your account. When you change the values on that page, the browser sends back the updated details to the web server and tells it to update these details on the database. If the action is successful, the website will show you a confirmation message.

This is a very simple example of what is known as **decoupled** architecture between frontend and backend systems. Decoupling allows greater flexibility and makes it easier to update or change components in your architecture. So, let's say you want to create a new frontend website. In such a case, you don't have to change the backend code at all, as long as your new frontend is built to make the same API requests as the old one.

REST APIs are *stateless*, meaning neither the client nor the server stores any states in between to do the communication. Every time a request is made, the data is processed, and a response is sent back without the protocol having to store any intermediate data. What this means is that the API is processing each request in isolation. It doesn't need to store information regarding the session itself. This is in contrast to a stateful protocol (such as **Transmission Control Protocol (TCP)**), which maintains information regarding the session during its life.

So, a **RESTful web service**, as the name suggests, is a collection of REST APIs used to carry out a set of tasks. For example, if we develop a set of REST APIs for the Bookr application to carry out a certain set of tasks, we can call it a RESTful web service. Next, we will start working on DRF.

Django REST framework

DRF is an open source Python library that can be used to develop REST APIs for a Django project. DRF has most of the necessary functionality built in to help develop APIs for any Django project. Throughout this chapter, we will be using it to develop APIs for our Bookr project. In the following section, we will install and configure DRF.

Installation and configuration

To install `djangorestframework` in the virtual environment setup along with PyCharm, follow these steps:

1. Enter the following code in your Terminal app or Command Prompt to do this:

```
pip install djangorestframework
```

2. Next, open the `settings.py` file and add `rest_framework` to `INSTALLED_APPS` as shown in the following snippet:

```
INSTALLED_APPS = [  
    "bookr_admin.apps.BookrAdminConfig",  
    "django.contrib.auth",  
    "django.contrib.contenttypes",  
    "django.contrib.sessions",  
    "django.contrib.messages",  
    "django.contrib.staticfiles",  
    "rest_framework",  
    "reviews",]
```

You are now ready to start using DRF to create your first simple API.

Functional API views

In *Chapter 3, URL Mapping, Views, and Templates*, we learned about simple functional views that take a request and return a response. We can write similar functional views using DRF. However, note that class-based views are more commonly used and will be covered next. A functional view is created by simply adding the following decorator onto a normal view, as follows:

```
from rest_framework.decorators import api_view  
  
@api_view  
def my_view(request):  
    ...
```

This decorator takes the functional view and turns it into a subclass of the `APIView` DRF. It's a quick way to include an existing view as part of your API. Using what we have learned so far, we will create a simple REST API using DRF in the next section.

Exercise 12.01 – creating a simple REST API

In this exercise, you will create your first REST API using DRF and implement an endpoint using a functional view. You will create this endpoint to view the total number of books in the database:

Note

You'll need to have DRF installed on your system to proceed with this exercise. If you haven't already installed it, make sure you refer to the *Installation and configuration* section earlier in this chapter.

1. Create `api_views.py` in the `bookr/reviews` folder.

REST API views work like Django's conventional views. We could have added the API views, along with the other views, in the `views.py` folder. However, having our REST API views in a separate file will help us maintain a cleaner code base.

2. Add the following code in `api_views.py`:

```
from rest_framework.decorators import api_view
from rest_framework.response import Response
from .models import Book

@api_view()
def first_api_view(request):
    num_books = Book.objects.count()
    return Response({"num_books": num_books})
```

The first line imports the `api_view` decorator, which will convert our functional view into one that can be used with DRF, and the second line imports `Response`, which will be used to return a response.

The `view` function returns a `Response` object containing a dictionary with the number of books in our database (see the highlighted part).

3. Open `bookr/reviews/urls.py` and import the `api_views` module. Then, add a new path to the `api_views` module in the URL patterns that we have developed throughout this book, as follows:

```
from . import views, api_views

urlpatterns = [path('api/first_api_view/',
                    api_views.first_api_view),
               ...
               ]
```

4. Start the Django service with the `python manage.py runserver` command and go to `http://127.0.0.1:8000/api/first_api_view/` to make your first API request. Your screen should appear as in *Figure 12.1*:

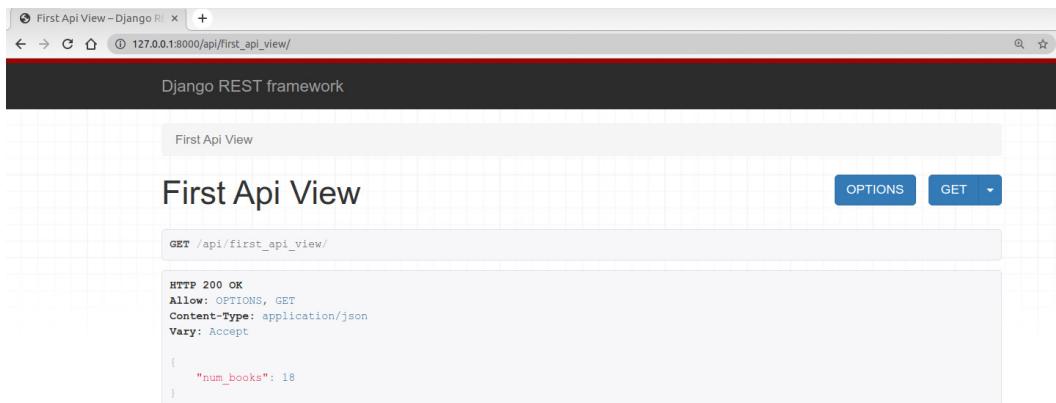


Figure 12.1 – First Api View with the number of books

Calling this URL endpoint made a default GET request to the API endpoint, which returned a JSON key-value pair ("num_books" : 0). Also, notice how DRF provides a nice interface to view and interact with the APIs.

5. We could also use the Linux `curl` (client URL) command to send an HTTP request as follows:

```
curl http://127.0.0.1:8000/api/first_api_view/
{ "num_books":18 }
```

Alternatively, if you are using Windows 10, you can make an equivalent HTTP request with `curl.exe` from Command Prompt as follows:

```
curl.exe http://127.0.0.1:8000/api/first_api_view/
```

In this exercise, we learned how to create an API view using DRF and a simple functional view. We will now look at a more elegant way to convert between information stored in the database and what gets returned by our API using serializers.

Understanding serializers

By now, we are well versed in the way Django works with data in our application. Broadly, the columns of a database table are defined in a class in `models.py`, and when we access a row of the table, we are working with an instance of that class. Ideally, we often just want to pass this object to our frontend application. For example, if we wanted to build a website that displayed a list of books in our `Bookr` app, we would want to call the `title` property of each book instance to know what string to display to the user. However, our frontend application knows nothing about Python and needs to retrieve this data through an HTTP request, which just returns a string in a specific format.

This means that any translation of information between Django and the frontend (via our API) must be done by representing the information in **JSON** format. JSON objects look similar to a Python dictionary, except there are some extra rules that constrict the exact syntax. In our previous example in *Exercise 12.01, Creating a simple REST API*, the API returned the following JSON object containing the number of books in our database:

```
{ "num_books": 0 }
```

But what if we wanted to return the full details about an actual book in our database with our API? DRF's `serializer` class helps to convert complex Python objects into formats such as JSON or XML so that they can be transmitted across the web using the HTTP protocol. The part of DRF that does this conversion is named `serializer`. Serializers also perform deserialization, which refers to converting serialized data back into Python objects, so that the data can be processed in the application. In the following section, we will implement a way to display a list of books.

Exercise 12.02 – creating an API view to display a list of books

In this exercise, you will use serializers to create an API that returns a list of all books present in the `bookr` application:

1. Create a file named `serializers.py` in the `bookr/reviews` folder. This is the file where we will place all the serializer code for the APIs.
2. Add the following code to `serializers.py`:

```
from rest_framework import serializers

class PublisherSerializer(serializers.Serializer):
    name = serializers.CharField()
    website = serializers.URLField()
    email = serializers.EmailField()
```

```
class BookSerializer(serializers.Serializer):
    title = serializers.CharField()
    publication_date = serializers.DateField()
    isbn = serializers.CharField()
    publisher = PublisherSerializer()
```

Here, the first line imports the `serializers` from the `rest_framework` module.

Following the imports, we have defined two classes, `PublisherSerializer` and `BookSerializer`. As the names suggest, they are serializers for the `Publisher` and `Book` models, respectively. Both these serializers are subclasses of `serializers.Serializer`, and we have defined field types for each serializer, such as `CharField`, `URLField`, `EmailField`, and so on.

Look at the `Publisher` model in the `bookr/reviews/models.py` file. The `Publisher` model has the `name`, `website`, and `email` attributes. So, to serialize a `Publisher` object, we need the `name`, `website`, and `email` attributes in the `serializer` class, which we have defined accordingly in `PublisherSerializer`.

Similarly, for the `Book` model, we have defined `title`, `publication_date`, `isbn`, and `publisher` as the desired attributes in `BookSerializer`. Since `publisher` is a foreign key for the `Book` model, we have used `PublisherSerializer` as the serializer for the `publisher` attribute.

3. Open `bookr/reviews/api_views.py`, remove any pre-existing code, and add the following code:

```
from rest_framework.decorators import api_view
from rest_framework.response import Response

from .models import Book
from .serializers import BookSerializer

@api_view()
def all_books(request):
    books = Book.objects.all()
    book_serializer = BookSerializer(books, many=True)
    return Response(book_serializer.data)
```

In the second line, we have imported the newly created `BookSerializer` from the `serializers` module.

We then add a functional view, `all_books` (as in the previous exercise). This view takes a query set containing all books and then serializes them using `BookSerializer`. The `serializer` class is also taking an argument, `many=True`, which indicates that the `books` object is `queryset` or a list of many objects. Remember that serialization takes Python objects and returns them in a JSON serializable format, as follows:

```
[OrderedDict([('title', 'Advanced Deep Learning with Keras'), ('publication_date', '2018-10-31'), ('isbn', '9781788629416'), ('publisher', OrderedDict([('name', 'Packt Publishing')], [('website', 'https://www.packtpub.com/'), ('email', 'info@packtpub.com')])), OrderedDict([('title', 'Hands-On Machine Learning for Algorithmic Trading'), ('publication_date', '2018-12-31'), ('isbn', '9781789346411'), ('publisher', OrderedDict([('name', 'Packt Publishing'), ('website', 'https://www.packtpub.com/'), ('email', 'info@packtpub.com')]))) ...
```

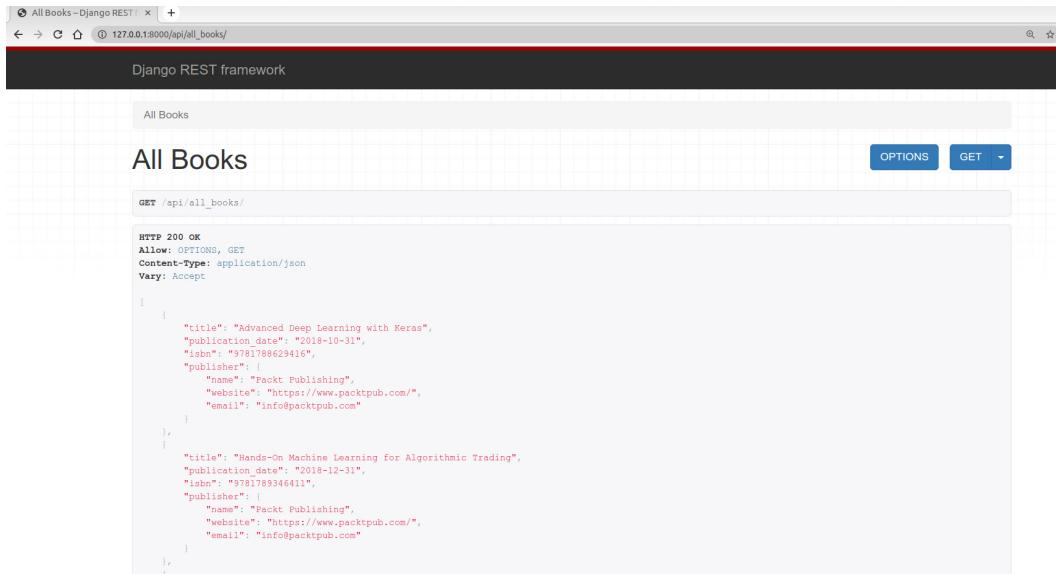
4. Open `bookr/reviews/urls.py`, remove the previous example path for `first_api_view`, and add the `all_books` path, as shown in the following code:

```
from django.urls import path
from . import views, api_views

urlpatterns = [
    path(
        'api/all_books/',
        api_views.all_books, name='all_books'
    ),
    ...
]
```

This newly added path calls the `all_books` view function when it comes across the `api/all_books/` path in the URL.

- Once all the code is added, run the Django server with the `python manage.py runserver` command and navigate to `http://127.0.0.1:8000/api/all_books/`. You should see something similar to *Figure 12.2*:



```
HTTP 200 OK
Allow: OPTIONS, GET
Content-Type: application/json
Vary: Accept

[
    {
        "title": "Advanced Deep Learning with Keras",
        "publication_date": "2018-10-31",
        "isbn": "9781788629416",
        "publisher": {
            "name": "Packt Publishing",
            "website": "https://www.packtpub.com/",
            "email": "info@packtpub.com"
        }
    },
    {
        "title": "Hands-on Machine Learning for Algorithmic Trading",
        "publication_date": "2018-12-31",
        "isbn": "9781789346411",
        "publisher": {
            "name": "Packt Publishing",
            "website": "https://www.packtpub.com/",
            "email": "info@packtpub.com"
        }
    }
]
```

Figure 12.2: List of books shown in the `all_books` endpoint

The preceding screenshot shows that the list of all books is returned upon calling the `/api/all_books` endpoint. And with that, you have successfully used a serializer to return data efficiently in your database with the help of a REST API.

Till now, we have been focusing on functional views. However, you will now learn that class-based views are more commonly used in DRF and will make your life much easier.

Class-based API views and generic views

Similar to what we learned in *Chapter 11, Advanced Templating and Class-Based Views*, we can also write class-based views for REST APIs. Class-based views are the preferred way of writing views among developers, as a lot can be achieved by writing very little code.

Just as with conventional views, DRF offers a set of generic views that makes writing class-based views even simpler. Generic views are designed considering some of the most common operations needed while creating APIs. Some of the generic views offered by DRF are `ListAPIView`, `RetrieveAPIView`, and so on. In *Exercise 12.02, Creating an API view to display a list of books*, our functional view was responsible for creating `queryset` of the objects and then calling the serializer. Equivalently, we could use `ListAPIView` to do the same thing:

```
class AllBooks(ListAPIView):
    queryset = Book.objects.all()
    serializer_class = BookSerializer
```

Here, `queryset` of objects is defined as a class attribute. Passing `queryset` through to `serializer` is handled by methods on `ListAPIView`.

Model serializers

In *Exercise 12.02, Creating an API view to display a list of books*, our serializer was defined as follows:

```
class BookSerializer(serializers.Serializer):
    title = serializers.CharField()
    publication_date = serializers.DateField()
    isbn = serializers.CharField()
    publisher = PublisherSerializer()
```

However, our model for Book looks like this (note how similar the definitions of the model and serializer appear to be):

```
class Book(models.Model):  
    """A published book."""  
  
    title = models.CharField(max_length=70,  
        help_text="The title of the book.")  
    publication_date = models.DateField(  
        verbose_name="Date the book was published.")  
    isbn = models.CharField(max_length=20,  
        verbose_name="ISBN number of the book.")  
    publisher = models.ForeignKey(Publisher, on_delete=models.  
CASCADE)  
    contributors = models.ManyToManyField("Contributor",  
        through="BookContributor")
```

We would prefer not to specify that the title must be `serializers.CharField()`. It would be easier if the serializer just looked at how `title` was defined in the model and could figure out what serializer field to use.

This is where model serializers come in. They provide shortcuts to create serializers by utilizing the definition of the fields on the model. Instead of specifying that `title` should be serialized using `CharField`, we just tell the model serializer we want to include `title`, and it uses the `CharField` serializer because the `title` field on the model is also `CharField`.

For example, suppose we wanted to create a serializer for the `Contributor` model in `models.py`. Instead of specifying the types of serializers that should be used for each field, we can give it a list of the field names and let it figure out the rest:

```
from rest_framework import serializers

from .models import Contributor

class ContributorSerializer(serializers.ModelSerializer):

    class Meta:
        model = Contributor
        fields = ['first_names', 'last_names', 'email']
```

In the following exercise, we will see how we can use a model serializer to avoid the duplication of code in the preceding classes.

Exercise 12.03 – creating class-based API views and model serializers

In this exercise, you will create class-based views to display a list of all books while using model serializers:

1. Open the `bookr/reviews/serializers.py` file, remove any pre-existing code, and replace it with the following code:

```
from rest_framework import serializers
from .models import Book, Publisher

class PublisherSerializer(
    serializers.ModelSerializer):

    class Meta:
```

```

        model = Publisher
        fields = ['name', 'website', 'email']

    class BookSerializer(serializers.ModelSerializer):
        publisher = PublisherSerializer()

        class Meta:
            model = Book
            fields = ['title', 'publication_date', 'isbn',
                      'publisher']

```

Here, we have included two model serializer classes, `PublisherSerializer` and `BookSerializer`. Both these classes inherit the `serializers.ModelSerializer` parent class. We do not need to specify how each field gets serialized; instead, we can simply pass a list of field names, and the field types are inferred from the definition in `models.py`.

Although mentioning the field inside `fields` is sufficient for the model serializer, under certain special cases, such as this one, we may have to customize the field since the `publisher` field is a foreign key. Hence, we must use `PublisherSerializer` to serialize the `publisher` field.

2. Next, open `bookr/reviews/api_views.py`, remove any pre-existing code, and add the following code:

```

from rest_framework import generics

from .models import Book
from .serializers import BookSerializer


class AllBooks(generics.ListAPIView):
    queryset = Book.objects.all()
    serializer_class = BookSerializer

```

Here, we use the DRF class-based `ListAPIView` instead of a functional view. This means that the list of books is defined as a class attribute, and we do not have to write a function that directly handles the request and calls the serializer. The book serializer from the previous step is also imported and assigned as an attribute of this class.

Open the `bookr/reviews/urls.py` file and modify the `/api/all_books` API path to include the new class-based view as follows:

```

urlpatterns = [
    path(
        'api/all_books/',
        api_views.AllBooks.as_view(),
        name='all_books'
    ),

```

```
...
```

```
]
```

Since we are using a class-based view, we have to use the class name along with the `as_view()` method.

- Once all the preceding modifications are completed, wait till the Django service restarts or start the server with the `python manage.py runserver` command, and then open the API at `http://127.0.0.1:8000/api/all_books/` in the web browser. You should see something like *Figure 12.3*:

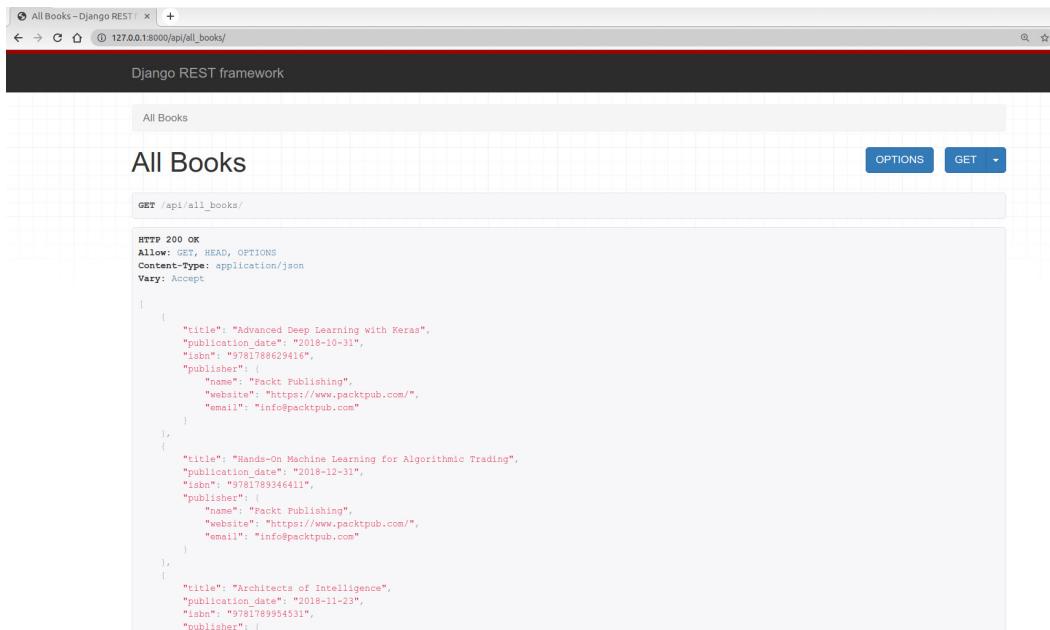


Figure 12.3: List of books shown in the `all_books` endpoint

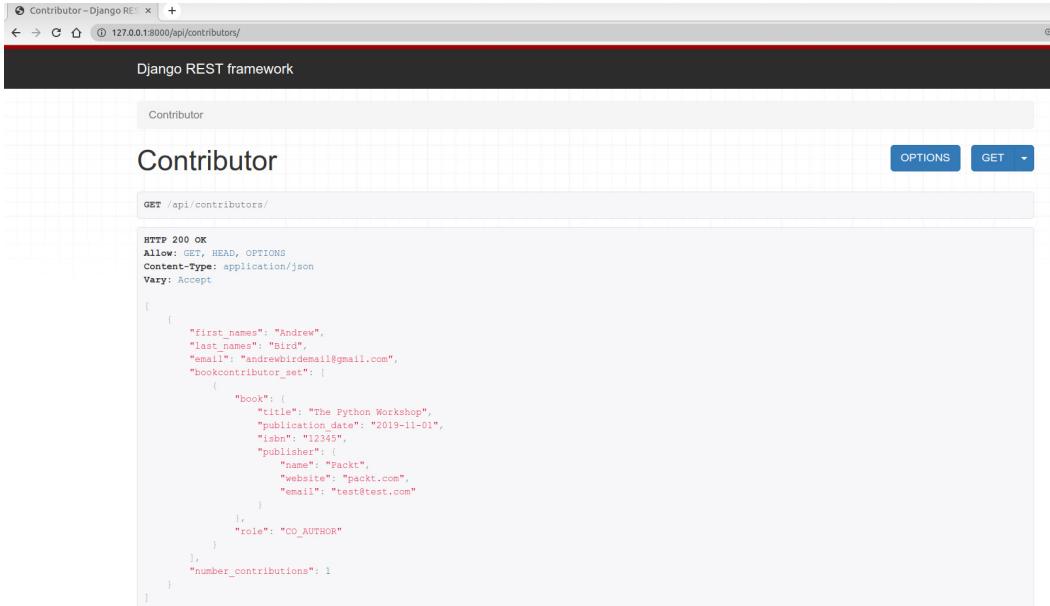
Similar to what we saw in *Exercise 12.02, Creating an API view to display a list of books*, this is a list of all books present in the book review application. In this exercise, we used model serializers to simplify our code, and the generic class-based `ListAPIView` to return a list of the books in our database. In the next section, you will implement an API endpoint to return a list of top contributors.

Activity 12.01 – creating an API endpoint for a top contributors page

Imagine that your team decides to create a web page that displays the top contributors (authors, coauthors, and editors) in your database. They decide to enlist the services of an external developer to create an app in ReactJS. To integrate with the Django backend, the developer will need an endpoint that provides the following:

- A list of all contributors in the database
- For each contributor, a list of all books they contributed to
- For each contributor, the number of books they contributed to
- For each book they contributed to, their role in the book

The final API view should look like this:



The screenshot shows a browser window with the title 'Contributor - Django REST'. The address bar shows '127.0.0.1:8000/api/contributors/'. The page content is a JSON response from the Django REST framework. The response is a list of contributors, each with their first name, last name, email, and a list of books they contributed to. Each book entry includes the title, publication date, ISBN, and publisher information. The response is paginated with a 'next' link.

```
HTTP 200 OK
Allow: GET, HEAD, OPTIONS
Content-Type: application/json
Vary: Accept

[{"id": 1, "first_name": "Andrew", "last_name": "Bird", "email": "andrewbird@gmail.com", "bookcontributor_set": [{"book": {"title": "The Python Workshop", "publication_date": "2019-11-01", "isbn": "12345", "publisher": {"name": "Packt", "website": "packt.com", "email": "test@test.com"}}, {"role": "CO_AUTHOR"}]}, {"id": 2, "first_name": "John", "last_name": "Doe", "email": "john.doe@example.com", "bookcontributor_set": [{"book": {"title": "Python Programming", "publication_date": "2020-05-01", "isbn": "98765", "publisher": {"name": "O'Reilly", "website": "oreilly.com", "email": "john.oreilly@example.com"}}, {"role": "AUTHOR"}]}], "number_contributions": 1}
```

Figure 12.4: The top contributors endpoint

To perform this task, execute the following steps:

1. Add a method to the `Contributor` class that returns the number of contributions made.
2. Add `ContributionSerializer`, which serializes the `BookContribution` model.
3. Add `ContributorSerializer`, which serializes the `Contributor` model.
4. Add `ContributorView`, which uses `ContributorSerializer`.
5. Add a pattern to `urls.py` to enable access to `ContributorView`.

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Simplifying the code using ViewSets

We have seen how we can optimize our code and make it more concise using class-based generic views. **ViewSets** and **routers** help us further simplify our code. As the name indicates, ViewSets are a set of views represented in a single class. For example, we used the `AllBooks` view to return a list of all books in the application and the `BookDetail` view to return the details of a single book. Using ViewSets, we could combine both these classes into a single class.

DRF also provides a class named `ModelViewSet`. This class not only combines the two views mentioned in the preceding discussion (`list` and `detail`) but also allows you to create, update, and delete model instances. The code needed to implement all this functionality could be as simple as specifying the serializer and `queryset`. For example, a view that allows you to manage all these actions for your user model could be defined as tersely as the following:

```
class UserViewSet(viewsets.ModelViewSet):  
    serializer_class = UserSerializer  
    queryset = User
```

Lastly, DRF provides a `ReadOnlyModelViewSet` class. This is a simpler version of the preceding `ModelViewSet`. It is identical, except that it only allows you to list and retrieve specific users. You cannot create, update, or delete records. In the following section, we will learn about routers.

URL configuration using routers and Viewsets

Routers, when used along with a viewset, automatically create the required URL endpoints for the viewset. This is because a single viewset is accessed at different URLs. For example, in the preceding `UserViewSet`, you would access a list of users at the `/api/users/` URL, and a specific user record at the `/api/users/123` URL, where `123` is the primary key of that user record. Here is a simple example of how you might use a router in the context of the previously defined `UserViewSet`:

```
from rest_framework import routers
router = routers.SimpleRouter()
router.register(r'users', UserViewSet)
urlpatterns = router.urls
```

Now, let's try combining the concepts of routers and ViewSets in a simple exercise.

Exercise 12.04 – using ViewSets and routers

In this exercise, we will combine the existing views to create a viewset and create the required routing for the viewset:

1. Open the `bookr/reviews/serializers.py` file, remove the pre-existing code, and add the following code snippet:

```
from django.contrib.auth.models import User
from django.utils import timezone
from rest_framework import serializers
from rest_framework.exceptions import
    NotAuthenticated, PermissionDenied
from .models import Book, Publisher, Review
from .utils import average_rating
class PublisherSerializer(
    serializers.ModelSerializer):
```

You can find the complete code snippet at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter12/Exercise12.04/bookr/reviews/serializers.py>.

Here, we added two new fields to `BookSerializer`, namely `reviews` and `rating`. The interesting thing about these fields is that the logic behind them is defined as a method on the serializer itself. This is why we use the `serializers.SerializerMethodField` type to set the `serializer` class attributes.

2. Open the `bookr/reviews/api_views.py` file, remove the pre-existing code, and add the following:

```
from rest_framework import viewsets
from rest_framework.pagination import (
    LimitOffsetPagination)

from .models import Book, Review
from .serializers import (BookSerializer,
    ReviewSerializer)

class BookViewSet(viewsets.ReadOnlyModelViewSet):
    queryset = Book.objects.all()  serializer_class
    = BookSerializer

class ReviewViewSet(viewsets.ModelViewSet):
    queryset = Review.objects.order_by('-
        date_created')
    serializer_class = ReviewSerializer
    pagination_class = LimitOffsetPagination
    authentication_classes = []
```

Here, we have removed the `AllBook` and `BookDetail` views and replaced them with `BookViewSet` and `ReviewViewSet`. In the first line, we import the `ViewSets` module from `rest_framework`. The `BookViewSet` class is a subclass of `ReadOnlyModelViewSet`, which ensures that the views are only used for the GET operation.

3. Next, open the `bookr/reviews/urls.py` file, remove the first two URL patterns starting with `api/`, and then add the following (highlighted) code:

```
from django.urls import path, include
from rest_framework.routers import DefaultRouter
from . import views, api_views
router = DefaultRouter()
router.register(r'books', api_views.BookViewSet)
router.register(r'reviews', api_views.ReviewViewSet)

urlpatterns = [
    path(
        'api/', include((router.urls, 'api'))
    ),
    path(
```

```
'books/' , views.book_list,
      name='book_list') ,  
  
path(
    'books/<int:pk>/' , views.book_detail,
      name='book_detail'
) ,  
path(
    'books/<int:book_pk>/reviews/new/' ,
      views.review_edit, name='review_create'
) ,  
path(
    'books/<int:book_pk>/reviews/<int:review_pk>/' ,
      views.review_edit, name='review_edit'
) ,  
path(
    'books/<int:pk>/media/' , views.book_media,
      name='book_media'
) ,  
path(
    'publishers/<int:pk>/' ,
      views.publisher_edit,
      name='publisher_detail'
) ,  
path(
    'publishers/new/> , views.publisher_edit,
      name='publisher_create'
) ]
```

Here, we have combined the `all_books` and `book_detail` paths into a single path called `books`. We have also added a new endpoint under the `reviews` path, which we will need in a later chapter.

We start by importing the `DefaultRouter` class from `rest_framework.routers`. Then, we create a `router` object using the `DefaultRouter` class and then register the newly created `BookViewSet` and `ReviewViewSet`, as can be seen from the highlighted code. This ensures that `BookViewSet` is invoked whenever the API has the `/api/books` path.

- Save all the files, and once the Django service restarts (or you start it manually with the `python manage.py runserver` command), go to the `http://127.0.0.1:8000/api/books/` URL to get a list of all the books. You should see the following view in the API explorer:

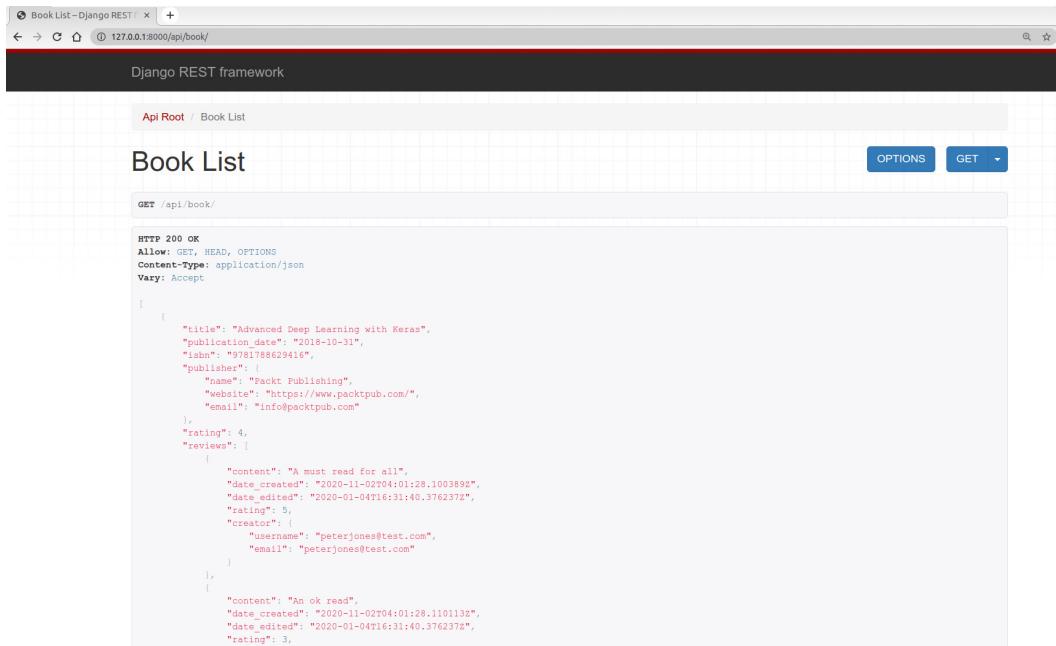


Figure 12.5: Book list at the /api/books path

- You can also access the details for a specific book using the `http://127.0.0.1:8000/api/books/1/` URL. In this case, it will return details for the book with a primary key of 1 (if it exists in your database):

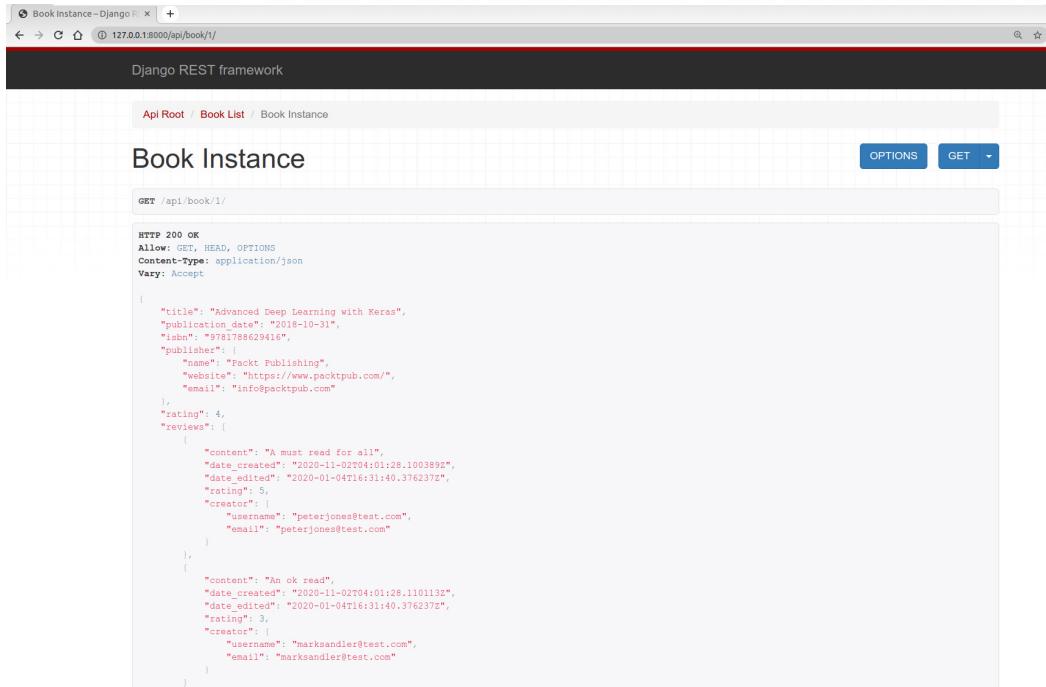


Figure 12.6: Book details for “Advanced Deep Learning with Keras”

In this exercise, we saw how we can use ViewSets and routers to combine the list and detail views into a single viewset. Using ViewSets makes our code more consistent and idiomatic, making it easier to collaborate with other developers. This becomes particularly important when integrating with a separate frontend application. In the next section, we will learn in brief about different types of authentication and implement token-based authentication for the book review application.

Implementing authentication

As we learned in *Chapter 9, Sessions and Authentication*, it is important to authenticate the users of our application. It is good practice to only allow those users who have registered in the application to log in and access information from the application. Similarly, for REST APIs, we also need to design a way to authenticate and authorize users before any information is passed on. For example, suppose Facebook’s website makes an API request to get a list of all comments for a post. If they did not have authentication on this endpoint, you could use it to get comments for any post you want programmatically. They obviously don’t want to allow this, so some sort of authentication needs to be implemented.

There are different authentication schemes, such as **basic authentication**, **session authentication**, **token authentication**, **remote user authentication**, and various third-party authentication solutions. For the scope of this chapter and for our Bookr application, we will use token authentication.

Note

For further reading on all the authentication schemes, please refer to the official documentation at <https://www.django-rest-framework.org/api-guide/authentication>.

Token-based authentication

Token-based authentication works by generating a unique token for a user in exchange for the user's username and password. Once the token is generated, it will be stored in the database for further reference and will be returned to the user upon every login.

This token is unique to the user, and they can then use it to authorize every API request they make. Token-based authentication eliminates the need to pass the username and password on every request. It is much safer and is best suited to client-server communication, such as a JavaScript-based web client interacting with the backend application via REST APIs.

An example of this would be a ReactJS or AngularJS application interacting with a Django backend via REST APIs.

The same architecture can be used if you are developing a mobile application to interact with the backend server via REST APIs, for instance, an Android or iOS application interacting with a Django backend via REST APIs.

Exercise 12.05 – implementing token-based authentication for Bookr APIs

In this exercise, you will implement token-based authentication for the bookr application's APIs:

1. Open the `bookr/settings.py` file and add `rest_framework.authtoken` to `INSTALLED_APPS`:

```
INSTALLED_APPS = [
    'django.contrib.admin',
    'django.contrib.auth',
    'django.contrib.contenttypes',
    'django.contrib.sessions',
    'django.contrib.messages',
    'django.contrib.staticfiles',
    'rest_framework',
    'rest_framework.authtoken',
    'reviews']
```

2. Since the auth token app has associated database changes, run the `migrate` command in the command line/terminal as follows:

```
python manage.py migrate
```

3. Open the `bookr/reviews/api_views.py` file, remove any pre-existing code, and replace it with the following:

```
from django.contrib.auth import authenticate
from rest_framework import viewsets
from rest_framework.authentication import
    TokenAuthentication
from rest_framework.authtoken.models import Token
from rest_framework.pagination import
    LimitOffsetPagination
from rest_framework.permissions import IsAuthenticated
from rest_framework.response import Response
from rest_framework.status import HTTP_404_NOT_FOUND,
    HTTP_200_OK
from rest_framework.views import APIView
```

You can find the complete code for this file at https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter12/Exercise12.05/bookr/reviews/api_views.py.

Here, we have defined a view called `Login`. The purpose of this view is to allow a user to get (or create if it does not already exist) a token that they can use to authenticate with the API.

We override the `post` method of this view because we want to customize the behavior when a user sends us data (that is, their login details). First, we use the `authenticate` method from Django's `auth` library to check whether the username and password are correct. If they are correct, then we will have a `user` object. If not, we return an `HTTP 404` error. If we do have a valid `user` object, then we simply get or create a token and return it to the user.

4. Next, let's add the authentication class to `BookViewSet`. This means that when a user tries to access this viewset, it will require them to authenticate using token-based authentication. Note that it's possible to include a list of different accepted authentication methods, not just one. We also add the `permissions_classes` attribute, which just uses DRF's built-in class that checks to see whether the given user has permission to view the data in this model:

```
class BookViewSet(viewsets.ReadOnlyModelViewSet):
    queryset = Book.objects.all()
    serializer_class = BookSerializer
    authentication_classes = [TokenAuthentication]
    permission_classes = [IsAuthenticated]
```

Note

The preceding code (highlighted) won't match the code you see on GitHub as we'll be modifying it later in step 9.

5. Open the `bookr/reviews/urls.py` file and add the following path into URL patterns:

```
path(
    'api/login', api_views.Login.as_view(),
    name=>login>
)
```

6. Save the file and wait for the application to restart, or start the server manually with the `python manage.py runserver` command. Then access the application using the `http://127.0.0.1:8000/api/login` URL. Your screen should appear as follows:

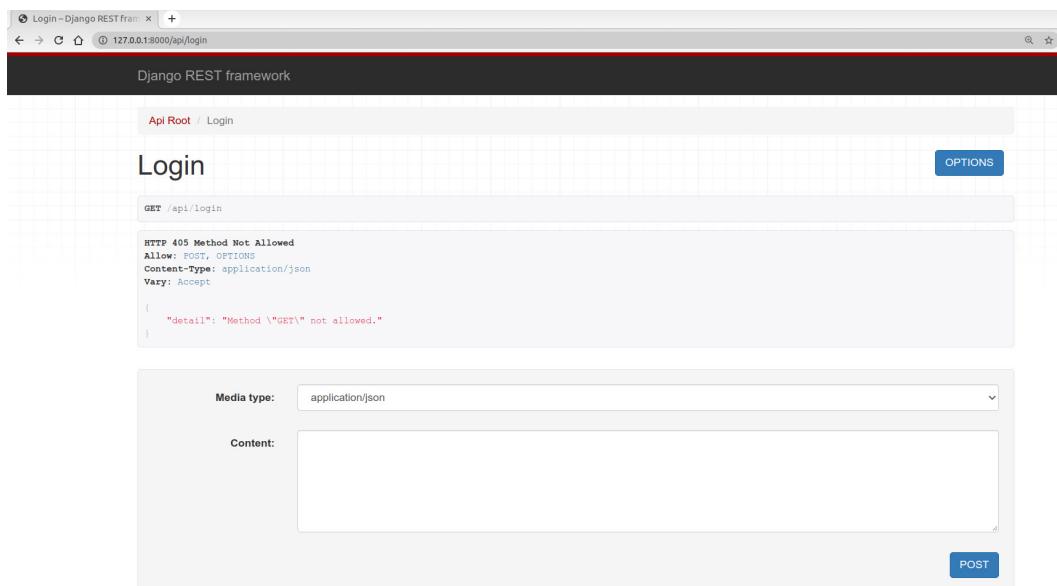


Figure 12.7: Login page

The API at `/api/login` is a POST-only message; hence, Method GET not allowed is displayed.

7. Next, enter the following snippet in the **Content** field and click on **POST**:

```
{  
    "username": "Peter",  
    "password": "testuserpassword"  
}
```

You will need to replace this with an actual username and password for your account in the database. Now you can see the token generated for the user. This is the token we need to use to access BookSerializer:

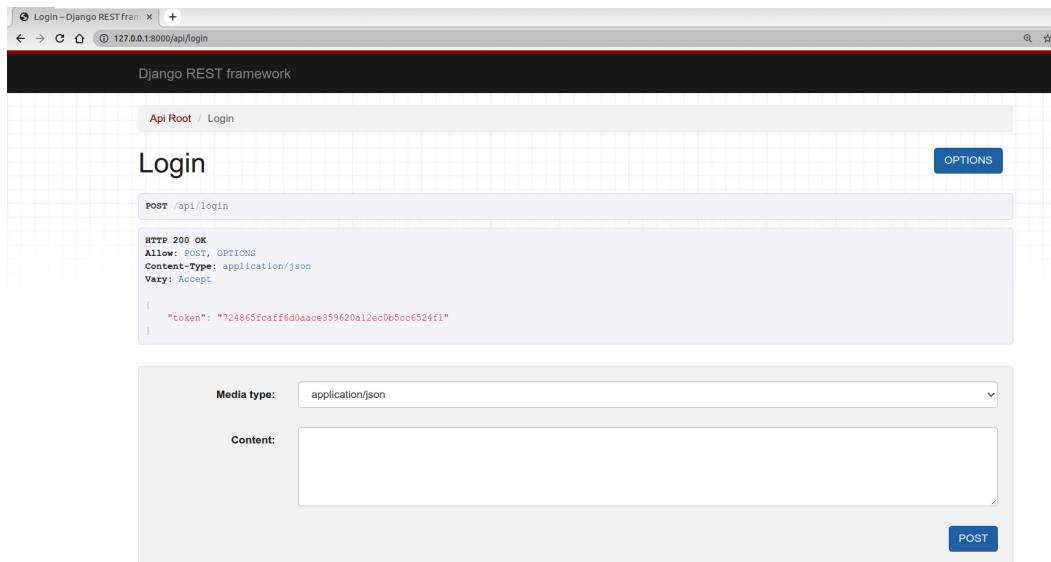


Figure 12.8: The token generated for the user

8. Try to access the list of books using the API that we previously created at <http://127.0.0.1:8000/api/books/>. Note that you are now not allowed to access it. This is because this viewset now requires you to use your token to authenticate.

The same API can be accessed using `curl` on the command line:

```
curl -X GET http://127.0.0.1:8000/api/books/  
  
{ "detail": "Authentication credentials were not provided." }
```

Since the token was not provided, the Authentication credentials were not provided message is displayed:



Figure 12.9: Message saying that the authentication details weren't provided

Note that if you're using Windows 10, replace `curl` in the preceding command with `curl.exe` and execute it from Command Prompt.

To pass the Authorization token (obtained in step 7) as a header, you can use the following command (Windows users can replace `curl` with `curl.exe`):

```
curl -X GET http://127.0.0.1:8000/api/books/ -H "Authorization: Token cd5dafa1d4361fd1502652d266eed3dcdb55c64f1"
```

Note

Before pasting this command, make sure you've replaced the token with the one you got when you ran step 7 of this exercise. It will be different from the one we have shown here.

The preceding command should now return the list of books:

```
[{"title": "Advanced Deep Learning with Keras", "publication_date": "2018-10-31", "isbn": "9781788629416", "publisher": {"name": "Packt Publishing", "website": "https://www.packtpub.com/"}, "email": "info@packtpub.com"}, {"rating": 4, "reviews": [{"content": "A must read for all", "date_created": "... (truncated)"}]
```

This operation ensured that only an existing user of the application could access and fetch the collection of all books.

9. Before moving on, set the authentication and permission classes on `BookViewSet` to an empty string. Future chapters will not utilize these authentication methods, and we will assume for the sake of simplicity that an unauthenticated user can access our API:

```
class BookViewSet(viewsets.ReadOnlyModelViewSet):  
    queryset = Book.objects.all()  
    serializer_class = BookSerializer  
    authentication_classes = []  
    permission_classes = []
```

In this exercise, we implemented token-based authentication in our Bookr app. We created a login view that allows us to retrieve the token for a given authenticated user. This then enabled us to make API requests from the command line by passing through the token as a header in the request.

Overall, in this section, we learned about different types of authentication, and then we learned in detail about token authentication and authorization while working with REST APIs.

Summary

This chapter introduced REST APIs, a fundamental building block in most real-world web applications. These APIs facilitate communication between the backend server and the web browser, so they are central to your growth as a Django web developer. We learned how to serialize data in our database so that it can be transmitted via an HTTP request. Next, we learned about the various options DRF gives us to simplify the code we write, taking advantage of the existing definitions of the models themselves. We also covered ViewSets and routers and saw how they can be used to condense code even further by combining the functionality of multiple views. Finally, we learned about authentication and authorization and implemented token-based authentication for the book review app.

In the next chapter, we will extend Bookr's functionality for its users by learning how to generate CSVs, PDFs, and other binary file types.

13

Generating CSV, PDF, and Other Binary Files

So far, we have learned about the various aspects of the Django framework and explored how we can build web applications using Django with all the features and customizations we want.

Let's say that, while building a web application, we need to do some analysis and prepare some reports. We may need to analyze user demographics about how the platform is being used or generate data that can be fed into machine learning systems to find patterns. We want our website to display some of the results of our analysis in a tabular format and other results as detailed graphs and charts. Furthermore, we want to allow our users to export the reports and peruse them further in applications such as Jupyter Notebook and Excel.

As we work our way through this chapter, we will learn how to bring these ideas to fruition and implement functionality in our web application that allows us to export records into structured formats such as tables through the use of **Comma-Separated Value (CSV)** files or Excel files. We will also learn how to allow our users to generate visual representations of the data we have stored inside our web application and export it in a PDF format so that it can be distributed easily for quick reference.

Let's start our journey by learning how to work with CSV files in Python. Learning this skill will help us create functionality that allows our readers to export our data for further analysis.

In this chapter, we will be covering the following topics:

- Working with CSV files inside Python
- Working with Python's csv module
- Working with Excel files in Python
- Working with PDF files in Python
- Playing with graphs in Python
- Integrating visualizations with Django

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter13>

Working with CSV files inside Python

There are several reasons we may need to export the data in our application. One of the reasons may involve analyzing that data – for example, we may need to understand the demographics of users registered on the application or extract patterns of application usage. We may also need to find out how our application is working for users to design future improvements. Such use cases require data to be in a format that can be easily consumed and analyzed. Here, the CSV file format comes to the rescue.

CSV is a handy file format that can be used to quickly export data from an application in a row-and-column format. CSV files usually have data separated by simple delimiters, which are used to differentiate one column from another, and newlines, which are used to indicate the start of a new record (or row) inside the table.

Python has great support for working with CSV files in its standard library, thanks to the `csv` module. This support enables the reading, parsing, and writing of CSV files. Let's take a look at how we can leverage the `csv` module provided by Python to work on CSV files and read and write data from them.

Working with Python's `csv` module

The `csv` module from Python allows us to interact with files that are in CSV format, which is nothing but a text file format – that is, the data stored inside the CSV files is human-readable.

The `csv` module requires that the file is opened before the methods supplied by the `csv` module can be applied. Let's take a look at how we can start with the very basic operation of reading data from CSV files.

Reading data from a CSV file

Reading data from CSV files is quite easy and consists of the following steps:

First, we must open the file by running the following script:

```
csv_file = open('path to csv file')
```

Here, we are reading the file using the Python `open()` method, which requires the fully qualified name of the file. This should be opened so that it can be passed as the parameter to the method.

Then, we must read the data from the `file` object using the `csv` module's `reader` method:

```
import csv
csv_data = csv.reader(csv_file)
```

In the first line, we imported the `csv` module, which contains the set of methods required to work on CSV files.

With the file opened, the next step is to create a CSV `reader` object by using the `csv` module's `reader` method. This method takes in the `file` object as returned by the `open()` call and uses the `file` object to read the data from the CSV file:

```
csv_reader = csv.reader(csv_file)
```

The data read by the `reader()` method is returned as a list of a list, where every sublist is a new record and every value inside the list is a value for the specified column. Generally, the first record in the list is referred to as a header, which denotes the different columns that are present inside the CSV file, but it is not necessary to have a `header` field inside a CSV file.

Once the data has been read by the `csv` module, we can iterate over this data to perform any kind of operation we may desire. This can be done as follows:

```
for csv_record in csv_data:  
    # do something
```

Once the processing is done, we can close the CSV file simply by using the `close()` method in Python's file handler object:

```
csv_file.close()
```

Now, let's look at our first exercise, where we will implement a simple module that helps us read a CSV file and output its contents on our screen.

Exercise 13.0 – reading a CSV file with Python

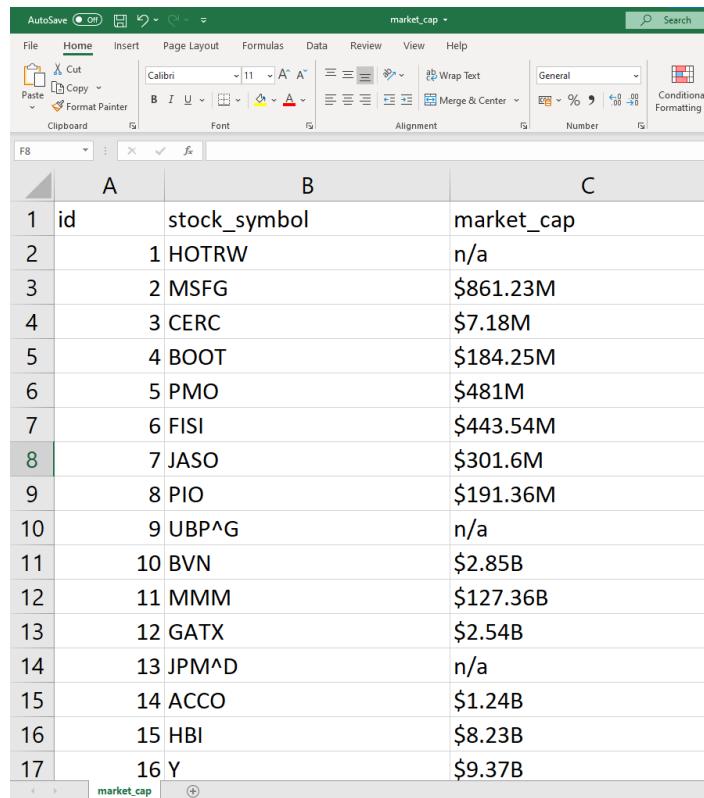
In this exercise, we will read and process a CSV file inside Python using Python's built-in `csv` module. The CSV file contains fictitious market data of several NASDAQ-listed companies. Let's get started with the steps:

1. First, download the `market_cap.csv` file from the GitHub repository for this book by clicking the following link: https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/blob/main/Chapter13/Exercise13.01/market_cap.csv.

Important note

The CSV file consists of randomly generated data and does not correspond to any historical market trends.

2. Once the file has been downloaded, open it and take a look at its contents. You will realize that the file contains a set of comma-separated values with each different record on its own line:



	A	B	C
1	id	stock_symbol	market_cap
2	1	HOTRW	n/a
3	2	MSFG	\$861.23M
4	3	CERC	\$7.18M
5	4	BOOT	\$184.25M
6	5	PMO	\$481M
7	6	FISI	\$443.54M
8	7	JASO	\$301.6M
9	8	PIO	\$191.36M
10	9	UBP^G	n/a
11	10	BVN	\$2.85B
12	11	MMM	\$127.36B
13	12	GATX	\$2.54B
14	13	JPM^D	n/a
15	14	ACCO	\$1.24B
16	15	HBI	\$8.23B
17	16	Y	\$9.37B

Figure 13.1 – Contents of the market_cap CSV file

3. Once the file has been downloaded, you can write the first piece of code. For this, create a new file named `csv_reader.py` in the same directory where the CSV file was downloaded and add the following code to it:

```
import csv

def read_csv(filename):
    """Read and output the details of CSV file."""
    try:
        with open(filename, newline='') as csv_file:
            csv_reader = csv.reader(csv_file)
            for record in csv_reader:
                print(record)
    except (IOError, OSError) as file_read_error:
```

```
print("Unable to open the csv file. Exception:  
{}".format(file_read_error))  
  
if __name__ == '__main__':  
    read_csv('market_cap.csv')
```

Let's try to understand what you just implemented in the preceding snippet of code.

After importing the `csv` module, to keep the code modular, you created a new method named `read_csv()` that takes in a single parameter – the filename to read the data from:

```
try:  
    with open(filename, newline='') as csv_file:
```

Now, if you are not familiar with the approach of opening the file shown in the preceding snippet, this is also known as the **try-with-resources** approach. In this case, any block of code that is encapsulated in the scope of the `with` block will have access to the `file` object, and once the code exits the scope of the `with` block, the file will be closed automatically.

Important note

It is a good habit to encapsulate file I/O operations within a `try-except` block since file I/O can fail for several reasons and showing stack traces to the users is not a good option.

The `reader()` method returns a `reader` object over which you can iterate to access the values, as we saw in the *Reading data from a CSV file* section:

```
for record in csv_reader:  
    print(record)
```

Once this is done, you must write the entry point method, from which your code will begin executing, by calling the `read_csv()` method and passing the name of the CSV file to read:

```
if __name__ == '__main__':  
    read_csv('market_cap.csv')
```

4. With this, you are done and ready to parse your CSV file. You can do this by running your Python file in a Terminal or Command Prompt, as shown here:

```
python3 csv_reader.py
```

Once the code executes, you should see the following output:

```
|sbadhwar@sbadhwar-mn1 Exercise13.01 % python3 csv_reader.py
[['id', 'stock_symbol', 'market_cap'],
 ['1', 'HOTRW', 'n/a'],
 ['2', 'MSFG', '$861.23M'],
 ['3', 'CERC', '$7.18M'],
 ['4', 'BOOT', '$184.25M'],
 ['5', 'PMO', '$481M'],
 ['6', 'FISI', '$443.54M'],
 ['7', 'JASO', '$301.6M'],
 ['8', 'PIO', '$191.36M'],
 ['9', 'UBP^G', 'n/a'],
 ['10', 'BVN', '$2.85B'],
 ['11', 'MMM', '$127.36B'],
 ['12', 'GATX', '$2.54B'],
 ['13', 'JPM^D', 'n/a'],
 ['14', 'ACCO', '$1.24B'],
 ['15', 'HBI', '$8.23B'],
 ['16', 'Y', '$9.37B'],
 ['17', 'DHX', '$128.99M'],
 ['18', 'AGNCB', '$9.12B'],
 ['19', 'BT', '$37.72B'],
 ['20', 'GGP^A', 'n/a'],
 ['21', 'XRAY', '$14.6B'],
 ['22', 'ANCB', '$62.74M'],
 ['23', 'MLP', '$386.9M'],
 ['24', 'AEG', '$10.51B'],
 ['25', 'CELGZ', 'n/a']]
```

Figure 13.2: Output from the CSV reader program

With this, you now know how to read CSV file contents. Also, as you can see from the output of *Exercise 13.01 – reading a CSV file with Python*, the output for individual rows is represented in the form of a list.

Now, let's look at how we can use the Python `csv` module to create new CSV files.

Writing to CSV files using Python

In the previous section, we explored how we can use the `csv` module in Python to read the contents of CSV-formatted files. Now, let's learn how we can write CSV data to files.

Writing CSV data follows a similar approach as reading from a CSV file, with some minor differences. The following steps outline the process of writing data to CSV files:

Open the file in writing mode by running the following script:

```
csv_file = open('path to csv file', 'w')
```

Obtain a CSV writer object, which can help us write data that is correctly formatted in CSV format. This can be done by calling the `writer()` method of the `csv` module, which returns a `writer` object, which can be used to write CSV format-compatible data to a CSV file:

```
csv_writer = csv.writer(csv_file)
```

Once the `writer` object is available, we can start writing the data. This is facilitated by the `write_row()` method of the `writer` object. The `write_row()` method takes in a list of values that it writes to the CSV file. The list itself indicates a single row and the values inside the list indicate the values of columns:

```
record = ['value1', 'value2', 'value3']
csv_writer.writerow(record)
```

If you want to write multiple records in a single call, you can also use the `writerows()` method of the CSV writer. The `writerows()` method behaves similarly to the `writerow()` method but takes a list of lists and can write multiple rows in one go:

```
records = [['value11', 'value12', 'value13'], ['value21',
'value22', 'value23']]
csv_writer.writerows(records)
```

Once the records have been written, we can close the CSV file:

```
csv_file.close()
```

Now, let's apply what we've learned to the next exercise and implement a program that will help us write values to CSV files.

Exercise 13.02 – generating a CSV file using Python's csv module

In this exercise, you will use the Python `csv` module to create new CSV files:

1. Create a new file named `csv_writer.py`, inside which you will write the code for the CSV writer. Inside this file, add the following code:

```
import csv

def write_csv(filename, header, data):
    """Write the provided data to the CSV file.

    :param str filename: The name of the file to which
        the data should be written
    :param list header: The header for the columns in
        csv file
    :param list data: The list of list mapping the
        values to the columns
    """

```

```
try:
    with open(filename, 'w') as csv_file:
        csv_writer = csv.writer(csv_file)
        csv_writer.writerow(header)
        csv_writer.writerows(data)
except (IOError, OSError) as csv_file_error:
    print("Unable to write the contents to csv
          file. Exception:
          {}".format(csv_file_error))
```

With this code, you should now be able to create new CSV files easily. Now, going step by step, let's understand what we are trying to do in this code.

We defined a new method called `write_csv()`, which takes three parameters: the name of the file to which the data should be written (`filename`), the list of column names that should be used as headers (`header`), and a list of a list that contains the data that needs to be mapped to individual columns (`data`):

```
def write_csv(filename, header, data):
```

Now, with the parameters in place, the next step is to open the file to which the data needs to be written and map it to an object:

```
with open(filename, 'w') as csv_file:
```

Once the file is opened, we perform three main steps:

- First, we obtain a new CSV writer object by using the `writer()` method from the `csv` module and pass it to the file handler that holds a reference to our opened file:

```
csv_writer = csv.writer(csv_file)
```

- The next step involves using the CSV writer's `writerow()` method to write our dataset's header fields into the file:

```
csv_writer.writerow(header)
```

- Once we have written the header, the last step is to write the data to the CSV file for the individual columns that are present. For this, we can use the `csv` module's `writerows()` method to write multiple rows at once:

```
csv_writer.writerows(data)
```

Important note

While the preceding steps outline the verbose version of writing to a CSV file to provide better clarity of the underlying working of how CSV files should be structured, we could have merged the step of writing the header and data into a single line of code by having the header list as the first element of the data list and calling the `writerows()` method with the data list as a parameter.

- Once we have created the methods that can write the provided data to a CSV file, we must write the code for the entry point call, and inside it, set up the values for the header, data, and filename fields, and finally call the `write_csv()` method that we defined earlier:

```
if __name__ == '__main__':
    header = ['name', 'age', 'gender']
    data = [['Richard', 32, 'M'], ['Mumzil', 21, 'F'],
            ['Melinda', 25, 'F']]
    filename = 'sample_output.csv'
    write_csv(filename, header, data)
```

- Now, with the code in place, execute the file we just created and see whether it creates the CSV file. To execute it, run the following command:

```
python3 csv_writer.py
```

Once the execution finishes, you will see that a new file has been created in the same directory as the one in which you executed the command. When you open the file, the contents should resemble what you can see in the following figure:

	A	B	C	D	E	F
1	name	age	gender			
2						
3	Richard		32 M			
4						
5	Mumzil		21 F			
6						
7	Melinda		25 F			
8						
9						
10						

Figure 13.3: Output from the CSV writer `sample_output.csv`

Now, we are well equipped to read and write the contents of CSV files.

With this exercise, we have learned how to write data to a CSV file. Now, it is time to look at some enhancements that can make reading and writing data to CSV files as a developer more convenient.

A better way to read and write CSV files

Now, there is one important thing that needs to be taken care of. As you may recall, the data read by the CSV reader usually maps values to a list. Now, if you want to access the values of individual columns, you need to use list indexes to access them. This is not natural and causes a higher degree of coupling between the program responsible for writing the file and the one responsible for reading the file. For example, what if the writer program shuffled the order of the rows? In this case, you now have to update the reader program to make sure it identifies the correct rows. So, the question arises, do we have a better way to read and write values that, instead of using list indexes, uses column names while preserving the context?

The answer to this is yes, and the solution is provided by another set of CSV modules known as `DictReader` and `DictWriter`, which provide the functionality of mapping objects in a CSV file to `dict`, rather than to a list.

This interface is easy to implement. Let's revisit the code you wrote in *Exercise 13.01 – reading a CSV file with Python*. If you wanted to parse the code as `dict`, the implementation of the `read_csv()` method would need to be changed, as shown here:

```
def read_csv(filename):
    """Read and output the details of CSV file."""
    try:
        with open(filename, newline='') as csv_file:
            csv_reader = csv.DictReader(csv_file)
            for record in csv_reader:
                print(record)
    except (IOError, OSError) as file_read_error:
        print("Unable to open the csv file. Exception: {}".
              format(file_read_error))
```

As you will notice, we only changed `csv.reader()` to `csv.DictReader()`, which should represent individual rows in the CSV file as `OrderedDict`. You can also verify this by making this change and executing the following command:

```
python3 csv_reader.py
```

This should result in the following output:

```
{'id': '1', 'stock_symbol': 'HOTRW', 'market_cap': 'n/a'}
{'id': '2', 'stock_symbol': 'MSFG', 'market_cap': '$861.23M'}
{'id': '3', 'stock_symbol': 'CERC', 'market_cap': '$7.18M'}
{'id': '4', 'stock_symbol': 'BOOT', 'market_cap': '$184.25M'}
{'id': '5', 'stock_symbol': 'PMO', 'market_cap': '$481M'}
{'id': '6', 'stock_symbol': 'FISI', 'market_cap': '$443.54M'}
{'id': '7', 'stock_symbol': 'JASO', 'market_cap': '$301.6M'}
{'id': '8', 'stock_symbol': 'PIO', 'market_cap': '$191.36M'}
{'id': '9', 'stock_symbol': 'UBPG', 'market_cap': 'n/a'}
{'id': '10', 'stock_symbol': 'BVN', 'market_cap': '$2.85B'}
{'id': '11', 'stock_symbol': 'MMM', 'market_cap': '$127.36B'}
{'id': '12', 'stock_symbol': 'GATX', 'market_cap': '$2.54B'}
{'id': '13', 'stock_symbol': 'JPM^D', 'market_cap': 'n/a'}
```

Figure 13.4: Output with DictReader

As you can see in the preceding figure, the individual rows are mapped as key-value pairs in the dictionary. To access these individual fields in rows, we can use this:

```
print(record.get('stock_symbol'))
```

That should give us the value of the `stock_symbol` field from our individual records.

Similarly, you can also use the `DictWriter()` interface to operate on CSV files as dictionaries. To see this, let's take a look at the `write_csv()` method from *Exercise 13.02 – generating a CSV file using Python's csv module*, and modify it as follows:

```
def write_csv(filename, header, data):
    """Write the provided data to the CSV file.

    :param str filename: The name of the file to which the
        data should be written
    :param list header: The header for the columns in csv
        file
    :param list data: The list of dicts mapping the values
        to the columns
    """
    try:
        with open(filename, 'w') as csv_file:
            csv_writer = csv.DictWriter(csv_file,
                                         fieldnames=header)
            csv_writer.writeheader()
            csv_writer.writerows(data)
    except (IOError, OSError) as csv_file_error:
        print("Unable to write the contents to csv file.
              Exception: {}".format(csv_file_error))
```

In the preceding code, we replaced `csv.writer()` with `csv.DictWriter()`, which provides a dictionary-like interface to interact with CSV files. `DictWriter()` also takes in a `fieldnames` parameter, which is used to map the individual columns in a CSV file before writing.

Next, to write this header, call the `writeheader()` method, which writes the `fieldname` header to the CSV file.

The final call involves the `writerows()` method, which takes in a list of dictionaries and writes them to the CSV file. For the code to work correctly, you also need to modify the data list to resemble the one shown here:

```
data = [{ 'name': 'Richard', 'age': 32, 'gender': 'M' },
        { 'name': 'Mumzil', 'age': 21, 'gender': 'F' },
        { 'name': 'Melinda', 'age': 25, 'gender': 'F' }]
```

With this, you have enough knowledge to work with CSV files inside Python.

With this section, we have learned about how to work with CSV files in Python using the standard library-provided `csv` module.

Since we are talking about how to deal with tabular data, specifically reading and writing it to files, let's take a look at one of the more well-known file formats used by one of the most popular tabular data editors – Microsoft Excel.

Working with Excel files in Python

Microsoft Excel is a world-renowned software in the field of bookkeeping and tabular record management. Similarly, the XLSX file format that was introduced with Excel has seen rapid and widespread adoption and is now supported by all the major product vendors.

You will find that Microsoft Excel and its XLSX format are used quite a lot in the marketing and sales departments of many companies. Let's say, for one such company's marketing department, you are building a web portal in Django that keeps track of the products purchased by users. It also displays data about the purchases, such as the time of purchase and the location where the purchase was made. The marketing and sales teams are planning to use this data to generate leads or to create relevant advertisements.

The marketing and sales teams also use Excel quite a lot. You might want to export the data available inside your web application in XLSX format, which is native to Excel.

Soon, we will look at how we can make our website work with the XLSX format. But before that, let's quickly take a look at the usage of binary file formats.

Binary file formats for data exports

Until now, we have worked mainly with textual data and how we can read and write it from text files. But often, text-based formats are not enough. For example, imagine you want to export an image or a graph. How will you represent an image or a graph as text, and how will you read and write to these images?

To help us in these situations, binary file formats come into the picture, which can help us read and write to and from a rich and diverse set of data. All commercial operating systems provide native support for working with both text and binary file formats, and it comes as no surprise that Python provides one of the most versatile implementations to work on binary data files. A simple example of this is the `open` command, which you can use to state the format you would like to open a file in:

```
file_handler = open('path to file', 'rb')
```

Here, `b` indicates binary.

Starting from this section, we will now be dealing with how we can work on binary files and use them to represent and export data from our Django web application. The first of the formats we are going to look at is the XLSX file format made popular by Microsoft Excel.

So, let's dive into handling XLSX files with Python.

Working with XLSX files using the `XlsxWriter` package

In this section, we will learn more about the XLSX file format and understand how we can work with it using the `XlsxWriter` package.

XLSX files

XLSX files are binary files that are used to store tabular data. These files can be read by any software that implements support for this format. The XLSX format arranges data into two logical partitions:

- **Workbooks:** Each XLSX file is called a workbook and is supposed to contain datasets related to a particular domain. In *Figure 13.5, Examplefile.xlsx* is a workbook (1):

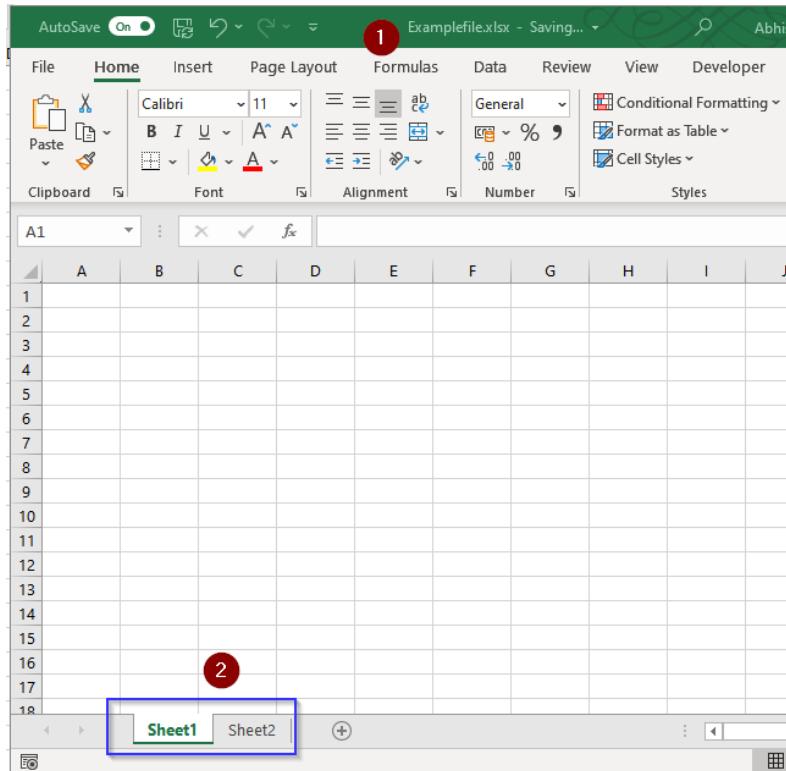


Figure 13.5: Workbooks and worksheets in Excel

- **Worksheets:** Inside each workbook, there can be one or more worksheets, which are used to store data about different but logically related datasets in a tabular format. In *Figure 13.5*, **Sheet1** and **Sheet2** are two worksheets (2).

When working with XLSX format, these are the two units that we generally work on. If you know about relational databases, you can think of workbooks as databases and worksheets as tables.

Now, let's try to understand how we can start working on XLSX files inside Python.

The XlsxWriter Python package

Python does not provide native support for working with XLSX files through its standard library. But thanks to the vast community of developers within the Python ecosystem, it is easy to find several packages that can help us manage our interaction with XLSX files. One popular package in this category is **XlsxWriter**.

`XlsxWriter` is an actively maintained package by the developer community that provides support for interacting with XLSX files. The package provides a lot of useful functionalities and supports the creation and management of workbooks as well as worksheets in individual workbooks. You can install it by running the following command in a Terminal or Command Prompt:

```
pip install XlsxWriter
```

Once installed, you can import the `xlsxwriter` module as follows:

```
import xlsxwriter
```

So, let's look at how we can start creating XLSX files with the support of the `XlsxWriter` package.

Creating a workbook

To start working on XLSX files, we first need to create them. An XLSX file is also known as a workbook and can be created by calling the `Workbook` class from the `xlsxwriter` package, as follows:

```
workbook = xlsxwriter.Workbook(filename)
```

The call to the `Workbook` class opens a binary file, specified with the `filename` argument, and returns an instance of `workbook` that can be used to further create worksheets and write data.

Creating a worksheet

Before we can start writing data to an XLSX file, we first need to create a worksheet. This can be done easily by calling the `add_worksheet()` method of the `workbook` object we obtained in the previous section:

```
worksheet = workbook.add_worksheet()
```

The `add_worksheet()` method creates a new worksheet, adds it to the workbook, and returns an object mapping the worksheet to a Python object, through which we can write data to the worksheet.

Writing data to the worksheet

Once a reference to the worksheet is available, we can start writing data to it by calling the `write` method of the `worksheet` object, as shown here:

```
worksheet.write(row_num, col_num, col_value)
```

As you can see, the `write()` method takes three parameters: a row number (`row_num`), a column number (`col_num`), and the data that belongs to the `[row_num, col_num]` pair, as represented by `col_value`. This call can be repeated to insert multiple data items into the worksheet.

Writing data to the workbook

Once all the data has been written, to finalize the written datasets and cleanly close the XLSX file, you must call the `close()` method on the workbook:

```
workbook.close()
```

This method writes any data that may be in the file buffer and finally closes the workbook. Now, let's use this knowledge to implement our own code, which will help us write data to an XLSX file.

Important note

It's not possible to cover all the methods and features the `XlsxWriter` package provides in this chapter. For more information, you can read the official documentation: <https://xlsxwriter.readthedocs.io/contents.html>.

Exercise 13.03 – creating XLSX files in Python

In this exercise, we will use the `XlsxWriter` package to create a new Excel (XLSX) file and add data to it from Python:

1. Install the `XlsxWriter` package on your system by running the following command in your Terminal app or Command Prompt:

```
pip install XlsxWriter
```

Once the command finishes, you will have the package installed on your system.

With the package installed, we can start writing the code that will create the Excel file.

2. Create a new file named `xlsx_demo.py` and add the following code to it:

```
import xlsxwriter

def create_workbook(filename):
    """Create a new workbook on which we can work."""
    workbook = xlsxwriter.Workbook(filename)
    return workbook
```

In the preceding code snippet, we created a new function that will assist us in creating a new workbook in which we can store your data. Once we have created a new workbook, the next step is to create a worksheet that provides us with the tabular format needed for us to organize the data to be stored inside the XLSX workbook.

3. With the workbook created, create a new worksheet by adding the following code snippet to your `xlsx_demo.py` file:

```
def create_worksheet(workbook):  
    """Add a new worksheet in the workbook."""  
    worksheet = workbook.add_worksheet()  
    return worksheet
```

In the preceding code snippet, we created a new worksheet using the `add_worksheet()` method of the `workbook` object provided by the `XlsxWriter` package. This worksheet will be used to write the data for the objects.

4. The next step is to create a helper function that can assist us in writing the data to the worksheet in a tabular format defined by the row and column numbering. For this, add the following snippet of code to your `xlsx_writer.py` file:

```
def write_data(worksheet, data):  
    """Write data to the worksheet."""  
    for row in range(len(data)):  
        for col in range(len(data[row])):  
            worksheet.write(row, col, data[row][col])
```

In the preceding code snippet, we created a new function named `write_data()` that takes two parameters: the `worksheet` object to which the data needs to be written and the `data` object represented by a list of lists that needs to be written to the worksheet. The function iterates over the data passed to it and then writes the data to the row and column it belongs to.

With all the core methods now implemented, we can add the method that can help close the `workbook` object cleanly so that the data is written to the file without any file corruption happening. For this, implement the following code snippet in the `xlsx_demo.py` file:

```
def close_workbook(workbook):  
    """Close an opened workbook."""  
    workbook.close()
```

The last step is to integrate all the methods we have implemented in the previous steps. For this, create a new entry point method in your `xlsx_demo.py` file:

```
if __name__ == '__main__':  
    data = [['John Doe', 38], ['Adam Cuvver', 22],  
            ['Stacy Martin', 28], ['Tom Harris', 42]]  
    workbook = create_workbook('sample_workbook.xlsx')  
    worksheet = create_worksheet(workbook)  
    write_data(worksheet, data)  
    close_workbook(workbook)
```

In the preceding code snippets, we created a dataset that we want to write to the XLSX file in the form of a list of lists. Once that was done, we obtained a new `workbook` object, which will be used to create an XLSX file. Inside this `workbook` object, we created a worksheet to organize our data in a row-and-column format, wrote the data to the worksheet, and closed the `workbook` to persist the data to the disk.

- Now, let's see whether the code you wrote works the way it is expected to work. For this, run the following command:

```
python3 xlsx_demo.py
```

Once the command has finished executing, you will see that a new file called `sample_workbook.xlsx` has been in the directory where the command was executed. To verify whether it contains the correct results, open this file with either Microsoft Excel or Google Sheets and view its contents. It should resemble what you can see here:

	A	B	C	D	E	F	G	H	I	J	K	L	M
1	John Doe	38											
2	Adam Cuvver	22											
3	Stacy Martin	28											
4	Tom Harris	42											
5													
6													
7													
8													
9													
10													
11													
12													
13													
14													
15													

Figure 13.6: Excel sheet generated using `xlsxwriter`

With the help of the `xlsxwriter` module, you can also apply formulas to your columns. For example, if you wanted to add another row that shows the average age of the people in the spreadsheet, you can do that simply by modifying the `write_data()` method, as shown here:

```
def write_data(worksheet, data):
    """Write data to the worksheet."""
    for row in range(len(data)):
        for col in range(len(data[row])):
            worksheet.write(row, col, data[row][col])
    worksheet.write(len(data), 0, "Avg. Age").
    # len(data) will give the next index to write to
    avg_formula = "=AVERAGE(B{}:B{})".format(
        1, len(data))
    worksheet.write(len(data), 1, avg_formula)
```

In the preceding code snippet, we added an additional `write` call to the worksheet and used the `AVERAGE` function provided by Excel to calculate the average age of the people in the worksheet.

With this, you now know how we can generate Microsoft Excel-compatible XLSX files using Python and how to export tabular content that's easily consumable by the different teams in your organization.

Now, let's cover another interesting file format that is widely used across the world: PDF.

Working with PDF files in Python

Portable Document Format or **PDF** is one of the most common file formats in the world. You must have encountered PDF documents at some point. These documents can include business reports, digital books, and more.

Do you remember encountering websites that have a button that reads **Print page as PDF**? A lot of websites for government agencies readily provide this option, which allows you to print the web page directly as a PDF. So, the question arises, how can we do this for our web app? How should we add the option to export certain content as a PDF?

Over the years, a huge community of developers has contributed a lot of useful packages to the Python ecosystem. One of those packages can help us achieve PDF file generation.

Converting web pages into PDFs

Sometimes, we may run into situations where we want to convert a web page into a PDF. For example, we may want to print a web page to store it as a local copy. This also comes in handy when trying to print a certificate that is natively displayed as a web page.

To help us in such efforts, we can leverage a simple library known as `weasyprint`, which is maintained by a community of Python developers and allows web pages to be quickly and easily converted into PDFs. So, let's take a look at how we can generate a PDF version of a web page.

Exercise 13.04 – generating a PDF version of a web page in Python

In this exercise, we will generate a PDF version of a website using Python. We will use a community-contributed Python module known as `weasyprint` that will help us generate the PDF. Let's get started:

1. To make the code in the upcoming steps work correctly, install the `weasyprint` module on your system by running the following command:

```
pip install weasyprint
```

Important note

`weasyprint` depends on the `cairo` library. If you haven't installed it yet, using `weasyprint` might raise an error with a message stating `libcairo-2.dll` file not found. To resolve this error, follow the steps mentioned in the `weasyprint` installation link at <https://weasyprint.readthedocs.io/en/stable/install.html#windows>.

- With the package now installed, create a new file named `pdf_demo.py` that will hold the PDF generation logic. Inside this file, write the following code:

```
from weasyprint import HTML

def generate_pdf(url, pdf_file):
    """Generate PDF version of the provided URL."""
    print("Generating PDF...")
    HTML(url).write_pdf(pdf_file)
```

Now, let's try to understand what this code does. In the first line, we imported the `HTML` class from the `weasyprint` package, which you installed in *step 1*:

```
from weasyprint import HTML
```

This `HTML` class provides us with a mechanism through which we can read the HTML content of a website if we have its URL.

In the next step, we created a new method named `generate_pdf()` that takes in two parameters – the URL that should be used as the source URL for generating the PDF and the `pdf_file` parameter, which takes in the name of the file to which the document should be written:

```
def generate_pdf(url, pdf_file):
```

Next, we passed the URL to the `HTML` class object we imported earlier. This caused the URL to be parsed by the `weasyprint` library and caused its HTML content to be read. Once this was done, we called the `write_pdf()` method of the `HTML` class object and provided the name of the file to which the content should be written:

```
HTML(url).write_pdf(pdf_file)
```

3. After this, write the entry point code that sets up the URL (for this exercise, we will use the text version of the **National Public Radio (NPR)** website) that should be used for your demo and the filename that should be used to write the PDF content to. Once that has been set, the code calls the `generate_pdf()` method to generate the content:

```
if __name__ == '__main__':
    url = 'http://text.npr.org'
    pdf_file = 'demo_page.pdf'
    generate_pdf(url, pdf_file)
```

4. Now, to see the code in action, run the following command:

```
python3 pdf_demo.py
```

Once the command finishes executing, you will have a new PDF file named `demo_page.pdf` that has been saved in the same directory where the command was executed. When you open the file, it should resemble what you can see here:

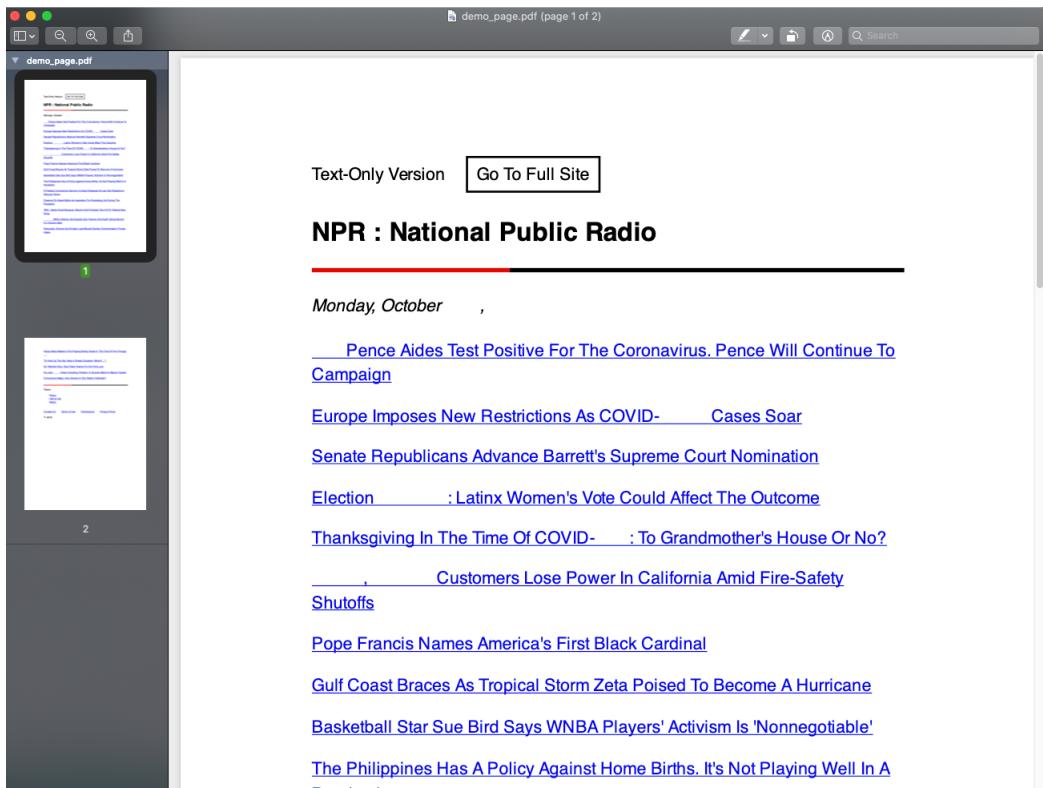


Figure 13.7: Web page converted into a PDF using weasyprint

In the PDF file that was generated, we can see that the content seems to lack the formatting that the actual website has. This has happened because the `weasyprint` package reads the HTML content but does not parse the attached CSS stylesheets for the page, so the page formatting is lost.

`weasyprint` also makes it quite easy to change the formatting of a page. This can be done simply by introducing the `stylesheets` parameter to the `write_pdf()` method. A simple modification to our `generate_pdf()` method is shown here:

```
from weasyprint import CSS, HTML

def generate_pdf(url, pdf_file):
    """Generate PDF version of the provided URL."""
    print("Generating PDF...")
    css = CSS(string='body{ font-size: 8px; }')
    HTML(url).write_pdf(pdf_file, stylesheets=[css])
```

Now, when the preceding code is executed, we will see that the font size for all the text inside the HTML body content of the page has a size of 8px in the printed PDF version.

Important note

The `HTML` class in `weasyprint` is also capable of taking any local files as well as raw HTML string content and using those files to generate PDFs. For further information, please visit the `weasyprint` documentation at <https://weasyprint.readthedocs.io>.

So far, we have learned how we can generate different types of binary files with Python, which can help us export our data in a structured manner or help us print PDF versions of our pages. Next, we will learn how to generate graph representations of our data using Python.

Playing with graphs in Python

Graphs are a great way to visually represent data that changes within a specific dimension. We come across graphs quite frequently in our day-to-day lives, be it weather charts for a week, stock market movements, or student performance report cards.

Similarly, graphs can come in quite handy when we are working with our web applications. For Bookr, we can use graphs as a visual medium to show how many books users read each week. Alternatively, we can show them the popularity of a book over time based on how many readers were reading the given book at a specific time. Now, let's look at how we can generate plots with Python and have them show up on our web pages.

Generating graphs with `plotly`

Graphs can come in quite handy when trying to visualize patterns in the data maintained by our applications. There are a lot of Python libraries that support developers in generating static or interactive graphs.

For this book, we will use `plotly`, a community-supported Python library that generates graphs and renders them on web pages. `plotly` is particularly interesting to us due to its ease of integration with Django.

To install it on your system, you can run the following command:

```
pip install plotly
```

Now that you've done that, let's take a look at how we can generate a graph visualization using `plotly`.

Setting up a figure

Before we can start generating a graph, we need to initialize a `plotly Figure` object, which essentially acts as a container for our graph. A `plotly Figure` object is quite easy to initialize; it can be done by using the following code snippet:

```
from plotly.graph_objs import graphs
figure = graphs.Figure()
```

The `Figure()` constructor from the `graph_objs` module of the `plotly` library returns an instance of the `Figure` graph container, inside which a graph can be generated. Once the `Figure` object is in place, we must generate a plot.

Generating a plot

A plot is a visual representation of a dataset. This plot could be a scatter plot, a line graph, a chart, and so on. For example, to generate a scatter plot, you can use the following code snippet:

```
scatter_plot = graphs.Scatter(x_axis_values, y_axis_values)
```

The `Scatter` constructor takes in the values for the `X`-axis and `Y`-axis and returns an object that can be used to build a scatter plot. Once the `scatter_plot` object has been generated, the next step is to add this plot to our `Figure`. This can be done as follows:

```
figure.add_trace(scatter_plot)
```

The `add_trace()` method is responsible for adding a plotting object to the figure and generating its visualization inside the figure.

Rendering a plot on a web page

Once the plot has been added to the figure, it can be rendered on a web page by calling the `plot()` method from the `offline` plotting module of the `plotly` library. This is shown in the following code snippet:

```
from plotly.offline import plot
visualization_html = plot(figure, output_type='div')
```

The `plot()` method takes two primary parameters: the first is the figure that needs to be rendered and the second is the HTML tag of the container inside which the figure HTML will be generated. The `plot` method returns fully integrated HTML that can be embedded in any web page or made a part of the template to render a graph.

Now, with this understanding of how graph plotting works, let's try a hands-on exercise where we will generate a graph for our sample dataset.

Exercise 13.05 – generating graphs in Python

In this exercise, we will generate a graph plot using Python. It will be a scatter plot that represents two-dimensional data:

1. For this exercise, we will be using the `plotly` library. To use this library, we need to install it on the system by running the following command:

```
pip install plotly
```

2. With the library now installed, create a new file named `scatter_plot_demo.py` and add the following `import` statements to it:

```
from plotly.offline import plot
import plotly.graph_objs as graphs
```

3. Once the imports have been sorted, create a method named `generate_scatter_plot()` that takes in two parameters – the values for the X-axis and the values for the Y-axis:

```
def generate_scatter_plot(x_axis, y_axis):
    """Generate a scatter plot for the provided x and
    y-axis values."""

```

4. Inside this method, first, create an object to act as a container for the graph:

```
figure = graphs.Figure()
```

5. Once the container for the graph has been set up, create a new `Scatter` object with the values for the `X`-axis and `Y`-axis and add it to the graph's `Figure` container:

```
scatter = graphs.Scatter(x=x_axis, y=y_axis)
figure.add_trace(scatter)
```

6. Once the scatter plot is ready and has been added to the figure, the last step is to generate the HTML, which can be used to render this plot inside a web page. To do this, call the `plot()` method and pass the graph container object to it. Then, render the HTML inside an HTML `div` tag:

```
return plot(figure, output_type='div')
```

The complete `generate_scatter_plot()` method should look like this now:

```
def generate_scatter_plot(x_axis, y_axis):
    figure = graphs.Figure()
    scatter = graphs.Scatter(x=x_axis, y=y_axis)
    figure.add_trace(scatter)
    return plot(figure, output_type='div')
```

7. Once the HTML for the plot has been generated, it needs to be rendered somewhere. For this, create a new method named `generate_html()`, which will take in the plot HTML as its parameter and render an HTML file consisting of the plot:

```
def generate_html(plot_html):
    """Generate an HTML page for the provided plot."""
    html_content = "<html><head><title>Plot Demo
    </title></head><body>{ }</body></html>".
    format(plot_html)
    try:
        with open('plot_demo.html', 'w') as plot_file:
            plot_file.write(html_content)
    except (IOError, OSError) as file_io_error:
        print("Unable to generate plot file.
        Exception: {}".format(file_io_error))
```

8. Once the method has been set up, the last step is to call it. For this, create a script entry point that will set up the values for the `X`-axis list and the `Y`-axis list and then call the `generate_scatter_plot()` method. With the value returned by the method, make a call to the `generate_html()` method, which will create an HTML page consisting of the scatter plot:

```
if __name__ == '__main__':
    x = [1,2,3,4,5]
    y = [3,8,7,9,2]
    plot_html = generate_scatter_plot(x, y)
    generate_html(plot_html)
```

- With the code in place, run the file and see what output is generated. To run the code, execute the following command:

```
python3 scatter_plot_demo.py
```

Once the execution completes, there will be a new `plot_demo.html` file in the same directory in which the script was executed. Upon opening the file, you should see the following:



Figure 13.8: Graph generated in the browser using plotly

With this, we have generated our first scatter plot, where different points are connected by a line.

In this exercise, we used the `plotly` library to generate a graph that can be rendered inside a browser for your readers to visualize data.

Now, you know how you can work with graphs in Python and how to generate HTML pages from them.

But as a web developer, how you can use these graphs in Django? Let's find out.

Integrating plotly with Django

The graphs generated by `plotly` are quite easy to embed in Django templates. Since the `plot()` method returns a fully contained HTML that can be used to render a graph, we can use the HTML returned as a template variable in Django and pass it as-is. The Django templating engine will then take care of adding this generated HTML to the final template before it is shown in the browser.

Some sample code for doing this is shown here:

```
def user_profile(request):
    username = request.user.get_username()
    scatter_plot_html = scatter_plot_books_read(username)
    return render(request, 'user_profile.html',
                  context={'plt_div': scatter_plot_html})
```

The preceding code will cause the `{ { plt_div } }` content used inside the template to be replaced by the HTML stored inside the `scatter_plot_demo` variable, and the final template to render the scatter plot of the number of books read per week.

Now, let's take a deep dive into how to handle adding visualizations to Django and enhancing the experience of our applications with plots and graphs.

Integrating visualizations with Django

In the preceding sections, you learned how data can be read and written in different formats that cater to the different needs of users. But how can we use what we've learned to integrate with Django?

For example, in Bookr, we might want to allow the user to export a list of books that they have read or visualize their book-reading activity over a year. How can that be done? The next exercise focuses on this aspect. You will learn how the components that we have seen so far can be integrated into Django web applications.

Exercise 13.06 – visualizing a user's reading history on the user's profile page

In this exercise, we will modify the user's profile page so that the user can visualize their book-reading history when they visit their profile page on Bookr.

Let's look at how this can be done:

1. To start integrating the ability to visualize the reading history of the user, you need to install the `plotly` library. To do this, run the following command in your Terminal:

```
pip install plotly
```

2. Once the library has been installed, the next step is to write the code that will fetch the total books read by the user as well as the books read by the user on a per-month basis. For this, create a new file named `utils.py` under the `bookr` application directory and add the required imports, which will be used to fetch the book-reading history of the user from the `Review` model of the `reviews` application:

```
import datetime

from django.db.models import Count
from reviews.models import Review
```

3. Next, create a new utility method named `get_books_read_by_month()` that takes in the username of the user for whom the reading history needs to be fetched.
4. Inside this method, we will query the `Review` model and return a dictionary of books read by the user on a per-month basis:

```
def get_books_read_by_month(username):
    """Get the books read by the user on per month
    basis.

    :param: str The username for which the books needs
           to be returned

    :return: dict of month wise books read
    """
    current_year = datetime.datetime.now().year

    books = Review.objects.filter(
        creator_username__contains=username,
        date_created__year=current_year)
        .values('date_created__month')
        .annotate(book_count=Count('book_title'))
    return books
```

Now, let's examine the following query, which is responsible for fetching the results of books read this year every month:

```
Review.objects.filter(
    creator_username__contains=username,
    date_created__year=current_year)
    .values('date_created__month')
    .annotate(book_count=Count('book_title'))
```

This query can be broken down into the following components:

- **Filtration**

```
Review.objects.filter(  
    creator__username__contains=username,  
    date_created__year=current_year)
```

Here, you filter the review records to choose all the records that belong to the current user, as well as the current year. The `year` field can be easily accessed from our `date_created` field by appending `__year`.

- **Projection**

Once the review records have been filtered, you are not interested in all the fields that might be there. What you are mainly interested in is the month and the number of books read each month. For this, use the `values()` call to select only the `month` field from the `date_created` attribute of the `Review` model on which you are going to run the group by operation.

- **Group By**

Here, you select the total number of books read in a given month. This is done by applying the `annotate` method to the `QuerySet` instance returned by the `values()` call.

5. Once you have the utilities file in place, you must write the view function, which is going to help in showing the books-read-per-month plot on the user's profile page. For this, open the `views.py` file under the `bookr` directory and start by adding the following imports to it:

```
from plotly.offline import plot  
import plotly.graph_objects as graphs  
  
from .utils import get_books_read_by_month
```

6. Once you have done this, you must modify the view function that renders the profile page. Currently, the profile page is being handled by the `profile()` method inside the `views.py` file. Modify the method so that it resembles the one shown here

```
@login_required  
def profile(request):  
    user = request.user  
    permissions = user.get_all_permissions()  
    # Get the books read in different months this year  
    books_read_by_month =  
        get_books_read_by_month(user.username)  
  
    # Initialize the Axis for graphs, X-Axis is  
    # months, Y-axis is books read  
    months = [i+1 for i in range(12)]  
    books_read = [0 for _ in range(12)]
```

```

# Set the value for books read per month on Y-Axis
for num_books_read in books_read_by_month:
    list_index =
        num_books_read['date_created_month'] - 1
    books_read[list_index] =
        num_books_read['book_count']

# Generate a scatter plot HTML
figure = graphs.Figure()
scatter = graphs.Scatter(x=months, y=books_read)
figure.add_trace(scatter)
figure.update_layout(xaxis_title="Month",
    yaxis_title="No. of books read")
plot_html = plot(figure, output_type='div')
# Add to template
return render(request, 'profile.html',
    {'user': user, 'permissions': permissions,
     'books_read_plot': plot_html})

```

In this method, you did a couple of things. First, you called the `get_books_read_by_month()` method and provided it with the username of the currently logged-in user. This method returns the list of books read by a given user on a per-month basis in the current year:

```

books_read_by_month =
    get_books_read_by_month(user.username)

```

The next thing you did was pre-initialize the *X*-axis and *Y*-axis for the graph with some default values. For this visualization, use the *X*-axis to display months and the *Y*-axis to display the number of books read.

Now, since you already know that a year is going to have only 12 months, pre-initialize the *X*-axis with a value between 1 and 12:

```

months = [i+1 for i in range(12)]

```

For the books read, initialize the *Y*-axis with all the 12 indexes set to 0, as follows:

```

books_read = [0 for _ in range(12)]

```

Now, with the pre-initialization done, fill in some actual values for the books read per month. For this, iterate upon the list you got as a result of the call made to `get_books_read_by_month(user.username)` and extract the month and the book count for the month from it.

Once the book count and month have been extracted, the next step is to assign the `book_count` value to the `books_read` list at the month index:

```

for num_books_read in books_read_by_month:
    list_index =
        num_books_read['date_created_month'] - 1

```

```
books_read[list_index] =  
    num_books_read['book_count']
```

Now, with the values for the axes set, generate a scatter plot using the `plotly` library:

```
figure = graphs.Figure()  
scatter = graphs.Scatter(x=months, y=books_read)  
figure.add_trace(scatter)  
figure.update_layout(xaxis_title="Month",  
                      yaxis_title="No. of books read")  
plot_html = plot(figure, output_type='div')
```

Once the HTML for the plot has been generated, pass it to the template using the `render()` method so that it can be visualized on the profile page:

```
return render(request, 'profile.html',  
{'user': user, 'permissions': permissions,  
'books_read_plot': plot_html})
```

7. With the view function created, the next step is to modify the template to render this graph. For this, open the `profile.html` file under the `templates` directory and add the following highlighted code to the file, just before the last `{% endblock %}` statement:

```
{% extends "base.html" %}  
  
{% block title %}Bookr{% endblock %}  
  
{% block heading %}Profile{% endblock %}  
  
{% block content %}  
    <ul>  
        <li>Username: {{ user.username }} </li>  
        <li>Name: {{ user.first_name }} {{  
            user.last_name }}</li>  
        <li>Date Joined: {{ user.date_joined }} </li>  
        <li>Email: {{ user.email }}</li>  
        <li>Last Login: {{ user.last_login }}</li>  
        <li>Groups: {{ groups }}{% if not groups  
            %}None{% endif %}</li>  
    </ul>  
    {% autoescape off %}  
        {{ books_read_plot }}  
    {% endautoescape %}  
  
{% endblock %}
```

This code snippet adds the `books_read_plot` variable that's passed in the `view` function to be used inside our HTML template. Also, note that `autoescape` is set to `off` for this variable. This is required because this variable contains HTML generated by the `plotly` library. If you allow Django to escape the HTML, you will only see raw HTML on the profile page and not a graph visualization.

With this, you have successfully integrated the plot into the application.

8. To try the visualization, run the following command and then navigate to your user profile by visiting `http://localhost:8080/accounts/profile`

```
python manage.py runserver localhost:8080
```

You should see a page that resembles the one shown here:

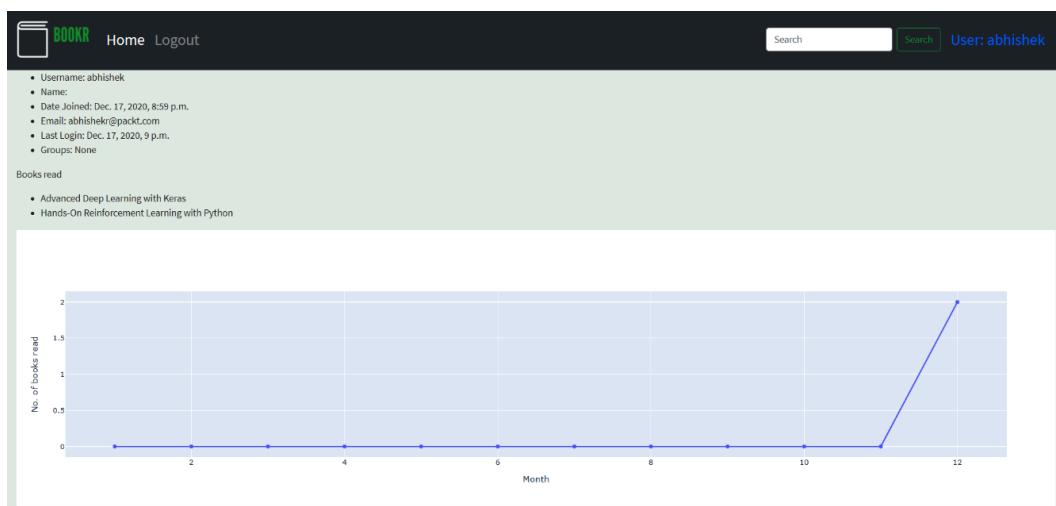


Figure 13.9: User book reading history scatter plot

In this exercise, you learned how to integrate a plotting library with Django to visualize the reading history of a user. Similarly, Django allows you to integrate any generic Python code into a web application, with the only constraint being that the data generated as a result of the integration should be transformed into a valid HTTP response that can be handled by any standard HTTP-compatible tool, such as a web browser or command-line tools such as `curl`.

Activity 13.01 – exporting the books read by a user as an XLSX file

In this activity, you will implement a new API endpoint inside Bookr that will allow your users to export and download a list of books they have read as an XLSX file:

1. Install the `XlsxWriter` library.
2. Inside the `utils.py` file you created under the `bookr` application, create a new function that will help in fetching the list of books that have been read by the user.
3. Inside the `views.py` file under the `bookr` directory, create a new view function that will allow the user to download their reading history in XLSX format.
4. To create an XLSX file inside the view function, first, create a `BytesIO`-based in-memory file that can be used to store the data from the `XlsxWriter` library.
5. Read the data stored inside the in-memory file using the `getvalue()` method of the temporary file object.
6. Finally, create a new `HttpResponse` instance with the `'application/vnd.ms-excel'` content type header, and then write the data obtained in *step 5* to the response object.
7. With the response object prepared, return the response object from the view function.
8. With the view function ready, map it to a URL endpoint that can be visited by a user to download their book-reading history.

Once you have mapped the URL endpoint, start the application and log in to it with your user account. Once you've done this, visit the URL endpoint you just created, and if upon visiting the URL endpoint your browser starts to download an Excel file, you have successfully completed the activity.

Important note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we looked at how we can deal with binary files and how Python's standard library, which comes pre-loaded with the necessary tools, allows us to handle commonly-used file formats such as CSV. We then moved on to learning how to read and write CSV files in Python using Python's `csv` module. Later, we worked with the `XlsxWriter` package, which allows us to generate Microsoft Excel-compatible files right from our Python environment without us having to worry about the internal formatting of the file.

The second half of this chapter was dedicated to learning how to use the `weasyprint` library to generate PDF versions of HTML pages. This skill can come in handy when we want to provide our users with an easy option to print the HTML version of our page with any added CSS styling of our choosing. The last section of this chapter discussed how we can generate interactive graphs in Python and render them as HTML pages that can be viewed inside the browser using the `plotly` library.

In the next chapter, we will look at how we can test the different components we have been implementing in the previous chapters to make sure that code changes do not break our website's functionality.

14

Testing Your Django Applications

In the preceding chapters, we focused on building our web application in Django by writing different components such as database models, views, and templates. We did all that to provide our users with an interactive application where they can create a profile and write reviews for the books they have read.

Apart from building and running the application, there is another important aspect of making sure that the application code works the way we expect it to work. This is ensured by a technique called **testing**. In testing, we run the different parts of the web application and check whether the output of the executed component matches the output we expected. If the output matches, we say that the component was tested successfully, while if the output doesn't match, we say that the component failed to work as intended.

In this chapter, as we go through the different sections, we will learn why testing is important, the different ways to test a web application, and how we can build a strong testing strategy that will help us ensure that the web application we build is robust.

In this chapter, we will cover the following topics:

- Importance of testing
- Automation testing
- Testing in Django
- Testing Django models
- Testing Django views
- Django request factory
- Test case classes in Django

Let's start our journey by first learning about the importance of testing.

Technical requirements

The complete code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter14>

Importance of testing

Making sure that an application works the way it was designed to work is an important aspect of the development efforts because otherwise, our users might keep on encountering weird behaviors that will usually drive away their engagement from the application.

The efforts we put into testing help us ensure that the different kinds of problems that we intend to solve are indeed being solved correctly. Imagine a case where a developer is building an online event scheduling platform. On this platform, users can schedule events on their calendars as per their local time zone. Now, what if in this platform, users can schedule events as expected, but due to a bug, the events are scheduled in an incorrect time zone? It is such issues that tend to drive many users away.

That's why a lot of enterprise companies spend a huge amount of money making sure that the applications they are building have undergone thorough testing. That way, they ensure that they do not release a buggy product or a product that is far away from satisfying user requirements.

In brief, testing helps us achieve the following goals:

- Ensure that the components of the application work according to specifications
- Ensure interoperability on different infrastructure platforms, if an application can be deployed on a different operating system such as Linux or Windows, and more
- Helps reduce the probability of introducing a bug while refactoring the application code

Now, a common assumption many people have about testing is that they have to test all the components manually as they are developed to make sure each component works according to its specifications and repeat this exercise every time a change is made or a new component is added to the application. While this is true, this doesn't provide a complete picture of testing. Testing as a technique has grown very powerful with time and as a developer, you can reduce a huge amount of testing effort by implementing **automated test cases**. So, what are these automated test cases? Or, in other words, what is **automation testing**? Let's find out.

Automation testing

Testing the whole application repeatedly when a single component is modified can turn out to be a challenging task, and even more so if that application consists of a large code base. The size of the code base could be due to the sheer number of features or the complexity of the problem it solves.

As we develop applications, it is important to make sure that the changes being made to these applications can be tested easily so that we can verify whether something is breaking. That's where the concept of automation testing comes in handy. The focus of automation testing is to write tests as code so that the individual components of an application can be tested in isolation as well as testing their interaction with each other.

In light of this, it is important for us to define the different kinds of automation tests that can be done for the applications.

Automation testing can be broadly categorized into five different types:

- **Unit testing:** In this type of testing, the individual isolated units of code are tested. For example, a unit test can target a single method or a single isolated API. This kind of testing is performed to make sure the basic units of the application work according to their specification.
- **Integration testing:** In this type of testing, the individual isolated units of code are merged to form a logical grouping. Once this grouping is formed, testing is performed on this logical group to make sure that the group works in the way it is expected.
- **Functional testing:** In this kind of testing, the overall functionality of the different components of the application is tested. This may include different APIs, user interfaces, and more.
- **Smoke testing:** In this kind of testing, the stability of the deployed application is tested to make sure that the application continues to remain functional as the users interact with it, without causing a crash.
- **Regression testing:** This kind of testing is done to make sure that the changes being made to the application do not degrade the previously built functionality of the application.

As we can see, testing is a big domain that takes time to master; entire books have been written on this topic. To make sure we highlight the important aspects of testing, we are going to focus on the aspect of unit testing in this chapter.

Testing in Django

Django is a feature-packed framework that aims to make web application development rapid. It also provides a full-featured way of testing applications.

Django provides a well-integrated module that allows application developers to write unit tests for their applications. This module is based on the Python `unittest` library, which ships with most of the Python distributions pre-packaged.

So, let's get started by understanding how we can write basic test cases in Django and how to leverage the framework-provided modules to test our application code.

Implementing test cases

While working on implementing mechanisms for testing your code, the first thing you need to understand is how this implementation can be logically grouped so that modules that are closely related to each other are tested in one logical unit.

This can be simplified by implementing a **test case**. A test case is nothing more than a logical unit that groups together tests that are related so that all the common logic to initialize the environment for the test cases can be combined in the same place, hence avoiding duplication of work while implementing application testing code.

Unit testing in Django

Now that we understand tests, let's take a look at how we can do unit testing inside Django. In the context of Django, a unit test consists of two major parts:

- A `TestCase` class, which wraps the different test cases that are grouped for a given module
- An actual test case that needs to be executed to test the flow of a particular component

The class implementing a unit test should inherit from the `TestCase` class provided by Django's `test` module. By default, Django provides a `tests.py` file in every application directory, which can be used to store the test cases for the application module.

Once these unit tests have been written, they can be executed easily by running them directly using the provided `test` command in `manage.py`, as follows:

```
python manage.py test
```

Now, let's look at how assertions can help us build our unit tests and check for the expected behavior.

Utilizing assertions

An important part of writing tests is validating if the test passed or failed. Generally, to implement such decisions inside a testing environment, we utilize something known as **assertions**.

Assertions are a common concept in software testing. They take in two operands and validate whether the value of the operand on the **left-hand side (LHS)** matches the value of the operand on the **right-hand side (RHS)**. If the value of the LHS matches the value of the RHS, an assertion is considered to be successful, whereas if the values differ, the assertion is considered to have failed.

An assertion that evaluates to `False` essentially causes a test case to be evaluated as a failure, which is then reported to the user.

Assertions in Python are quite easy to implement, and they use a simple keyword called `assert`. For example, the following code snippet shows a very simple assertion:

```
assert 1 == 1
```

The preceding assertion takes in a single expression, which evaluates to `True`. If this assertion was part of a test case, the test would have succeeded.

Now, let's see how we can implement test cases using the Python `unittest` library. Doing so is quite easy and can be accomplished with some easy-to-follow steps:

1. Import the `unittest` module, which allows us to build the test cases:

```
import unittest
```

2. Once the module has been imported, we can create a class whose name starts with `Test`, and which inherits from the `TestCase` class provided by the `unittest` module:

```
class TestMyModule(unittest.TestCase):  
    def test_method_a(self):  
        assert <expression>
```

Only if the `TestMyModule` class inherits the `TestCase` class will Django be able to run it automatically with full integration with the framework. Once the class has been defined, we can implement a new method inside the class named `test_method_a()`, which validates an assertion. And that's it.

An important part to note here is the naming scheme for the test cases and test functions. The test cases being implemented should be prefixed with the name `Test` so that the test execution modules can detect them as valid test cases and execute them. The same rule applies to naming testing methods.

3. Once the test case has been written, it can be simply executed by running the following command:

```
python manage.py test
```

With our basic understanding of implementing test cases clarified, let's write a very simple unit test to get a feel of how the unit testing framework behaves inside Django.

Exercise 14.01 – writing a simple unit test

In this exercise, we will write a simple unit test to understand how the Django unit testing framework works and use this knowledge to implement our first test case that validates a simple expression:

1. To get started, open the `tests.py` file under the `reviews` application of the Bookr project. By default, this file will contain only a single line that imports Django's `TestCase` class from the test module:

```
from django.test import TestCase
```

2. Add the following lines of code to the `tests.py` file you just opened:

```
class TestSimpleComponent(TestCase):  
    def test_basic_sum(self):  
        assert 1+1 == 2
```

Here, we created a new class named `TestSimpleComponent`, which inherits from the `TestCase` class provided by Django's test module. The `assert` statement will compare the expression on the LHS (`1+1`) with the one on the RHS (`2`).

3. Once you have written the test case, navigate back to the project folder and run the following command:

```
python manage.py test
```

The following output should be generated:

```
% ./manage.py test  
Creating test database for alias 'default'...  
System check identified no issues (0 silenced).  
.----  
----  
Ran 1 test in 0.001s  
  
OK  
Destroying test database for alias 'default'...
```

The preceding output signifies that Django's test runner executed 1 test case, which successfully passed the evaluation.

4. With the test case confirmed to be working and passing, we will now try to add another assertion at the end of the `test_basic_sum()` method, as shown in the following code snippet:

```
assert 1+1 == 3
```

- With the `assert` statement added in *step 4*, execute the test cases by running the following command from the project folder:

```
python manage.py test
```

At this point, you will notice Django reporting that the execution of the test cases has failed.

With this, you understand how test cases can be written in Django and how the assertions can be used to validate whether the output generated from your method calls under test is correct or not.

Now that we have basic knowledge of assertions, let's deep dive into the different types of assertions that are present and their usage.

Types of assertions

In *Exercise 14.01 – writing a simple unit test*, we briefly encountered assertions when we came across the following `assert` statement:

```
assert 1+1 == 2
```

These assertion statements are simple and use the Python `assert` keyword. A few different types of assertions are possible that can be tested inside a unit test while using the `unittest` library. Let's look at those:

- `assertNone`: This assertion is used to check whether an expression evaluates to `None` or not. For example, this type of assertion can be used in cases where a query to a database returns `None` because no records were found for the specified filtering criteria.
- `assertInstance`: This assertion is used to validate if a provided object evaluates to an instance of the provided type. For example, this assertion can be used to validate if the value returned by a method indeed is of a specific type, such as `List`, `Dict`, `Tuple`, and more.
- `assertEquals`: This is a very basic function that takes in two arguments and checks whether the arguments provided to it are equal in value or not. This can be useful when you plan to compare the values of data structures that do not guarantee ordering.
- `assertRaises`: This method is used to validate whether the name of the method provided to it when called raises a specified exception or not. This is helpful when we are writing test cases where a code path that raises an exception needs to be tested. As an example, this kind of assertion can be useful when we want to make sure an exception is raised by a method performing a database query if the database connection is not yet established.

These were just a small set of useful assertions that we can make in our test cases. The `unittest` module on top of which Django's testing library is built provides a lot more assertions that can be tested for.

Ever wondered what to do when our unit tests require some lists to be pre-populated or databases to be pre-initialized? The next section will cover how to pre-test this setup and how it can help us avoid introducing repetitive code that must be executed before every test case.

Performing pre-test setup and cleanup after every test case run

Sometimes, while writing test cases, we may need to perform some repetitive tasks, such as setting up some variables that will be required for the test. Once the test is over, we will want to clean up all the changes we've made to the test variables so that any new test starts with a fresh instance.

Luckily, the `unittest` library provides us with a useful way to automate our repetitive efforts of setting up the environment before every test case runs and cleaning it up after the test case is finished. This can be achieved using the two methods that we can implement in our `TestCase`.

The first is the `setUp()` method. This method is called before the execution of every test method inside the `TestCase` class. It implements the code required to set up the test case's environment before the test executes. This method can be a good place to set up any local database instance or test variables that may be required for the test cases.

Important note

The `setUp()` method is only valid for test cases written inside the `TestCase` class.

For example, the following example illustrates a simple definition of how the `setUp()` method is used inside a `TestCase` class:

```
class MyTestCase(unittest.TestCase):  
    def setUp(self):  
        # Do some initialization work  
    def test_method_a(self):  
        # code for testing method A  
    def test_method_b(self):  
        # code for testing method B
```

In the preceding example, when we try to execute the test cases, the `setUp()` method we defined here will be called every time before a test method executes. In other words, the `setUp()` method will be called before the execution of the `test_method_a()` call and then it will be called again before the `test_method_b()` call is called.

The second is the `tearDown()` method. This method is called once the test function finishes execution and cleans up the variables and their values once the test case's execution has finished. This method is executed regardless of whether the test case evaluates to `True` or `False`. An example of using the `tearDown()` method is shown here:

```
class MyTestCase(unittest.TestCase):  
    def setUp(self):  
        # Do some initialization work  
    def test_method_a(self):  
        # code for testing method A  
    def test_method_b(self):  
        # code for testing method B  
    def tearDown(self):  
        # perform cleanup
```

In the preceding example, the `tearDown()` method will be called every time a test method finishes execution – that is, once `test_method_a()` finishes executing and again once after `test_method_b()` finishes executing.

Now that we are aware of the different components of writing test cases, let's look at how we can test the different aspects of a Django application using the provided test framework.

Testing Django models

Models in Django are object-based representations of how data will be stored inside the database of an application. They provide methods that can help us validate the data input provided for a given record, as well as perform any processing on the data before it is inserted into the database.

It is as easy to test models in Django as it is to create them. Now, let's look at how Django models can be tested using the Django Test framework.

Exercise 14.02 – testing Django models

In this exercise, we will create a new Django model and write test cases for it. The test case will validate whether your model can correctly insert and retrieve the data from the database. These kinds of test cases that work on database models can turn out to be useful in cases where a team of developers is collaborating on a large project and the same database model may be modified by multiple developers over time. Implementing test cases for the database models allows developers to pre-emptively identify potentially breaking changes that they may inadvertently introduce as a part of their work. Let's get started with the steps:

1. Create a new application that you will use for the exercises in this chapter. For this, run the following command, which will set up a new application for your use case:

```
python manage.py startapp bookr_test
```

2. To make sure the `bookr_test` application behaves the same way as any other application in the Django project, add this application to our `INSTALLED_APPS` section of the `Bookr` project. To do this, open the `settings.py` file in your `Bookr` project and append the following code to the `INSTALLED_APPS` list:

```
INSTALLED_APPS = [  
    ...,  
    ...,  
    'bookr_test',  
]
```

3. Now, with the application setup complete, create a new database model that you will use for testing purposes. For this exercise, we are going to create a new model named `Publisher` that will store the details about the book publisher in our database. To create the model, open the `models.py` file under the `bookr_test` directory and add the following code to it:

```
from django.db import models  
  
class Publisher(models.Model):  
    """A company that publishes books."""  
    name = models.CharField(  
        max_length=50, help_text="The name of the  
        Publisher.")  
    website = models.URLField(  
        help_text="The Publisher's website.")  
    email = models.EmailField(  
        help_text="The Publisher's email address.")  
  
    def __str__(self):  
        return self.name
```

In the preceding code snippet, you created a new class named `Publisher` that inherits from the `Model` class of the Django's `models` module, defining the class as a Django model that will be used to store data about the publisher:

```
class Publisher(models.Model)
```

Inside this model, we have added three fields, which will act as the properties of the model:

- `name`: The name of the publisher
- `website`: The website belonging to the publisher
- `email`: The email address of the publisher

Once we did this, we created a class method called `__str__()` that defines what the string representation of the `Model` class will look like.

4. Now, with the model created, you need to migrate this model before you can run a test on it. To do this, run the following commands:

```
python manage.py makemigrations
python manage.py migrate
```

5. With the model now set up, write the test case with which you are going to test the model you created in *step 3*. For this, open the `tests.py` file under the `bookr_test` directory and add the following code to it:

```
from django.test import TestCase
from .models import Publisher

class TestPublisherModel(TestCase):
    """Test the publisher model."""
    def setUp(self):
        self.p = Publisher(name='Packt',
                           website='www.packt.com',
                           email='contact@packt.com')

    def test_create_publisher(self):
        self.assertIsNotNone(self.p, Publisher)

    def test_str_representation(self):
        self.assertEquals(str(self.p), "Packt")
```

In the preceding code snippet, there are a couple of things worth exploring.

At the start, after importing the `TestCase` class from the Django `test` module, you imported the `Publisher` model from the `bookr_test` directory, which is going to be used for testing.

Once the required libraries have been imported, we must create a new class named `TestPublisherModel` that inherits the `TestCase` class and is used for grouping the unit tests related to the Publisher model:

```
class TestPublisherModel(TestCase):
```

Inside this class, we defined a couple of methods. First, we defined a new method named `setUp()` and added the `Model` object creation code to it so that the `Model` object is created every time a new test method is executed inside this test case. This `Model` object is stored as a class member so that it can be accessed inside other methods without a problem:

```
def setUp(self):
    self.p = Publisher(name='Packt',
                       website='www.packt.com',
                       email='contact@packt.com')
```

The first test case validates whether the `Model` object for the Publisher model was created successfully or not. To do this, we created a new method named `test_create_publisher()` inside which we check whether the `Model` object created points to an object of the Publisher type. If this `Model` object was not created successfully, your test will fail:

```
def test_create_publisher(self):
    self.assertIsInstance(self.p, Publisher)
```

If you check carefully, you will see that we are using the `assertIsInstance()` method of the `unittest` library here to assert whether the `Model` object belongs to the Publisher type or not.

The next test validates whether the string representation of the `Model` is the same as what we expected it to be. From the code definition, the string representation of the Publisher model should output the name of the publisher. To test this, we created a new method named `test_str_representation()` and checked whether the generated string representation of the model matches the one we are expecting:

```
def test_str_representation(self):
    self.assertEqual(str(self.p), "Packt")
```

To perform this validation, we used the `assertEquals` method of the `unittest` library, which validates whether the two values provided to it are equal or not.

6. With the test cases now in place, you can run them to check what happens. To run these test cases, run the following command:

```
python manage.py test
```

Once the command finishes executing, you will see an output that resembles the one shown here:

```
% python manage.py test
Creating test database for alias 'default'...
System check identified no issues (0 silenced).
..
-----
Ran 2 tests in 0.002s

OK
Destroying test database for alias 'default'...
```

As you can see from the preceding output, the test cases are executed successfully, hence validating that the operations, such as creating a new Publisher object and its string representation when fetched, are being done correctly.

With this exercise, we saw how we can write test cases for our Django models easily and validate their functionality by creating objects, retrieving them, and representing them.

Also, notice the following important line in the output from *Exercise 14.02 – testing Django models*:

```
"Destroying test database for alias 'default'..."
```

This happens because when there are test cases that require the data to be persisted inside a database, instead of using the production database, Django creates a new empty database for the test cases, which it uses to persist the value for the test case.

Now that we know how to test models in Django, in the next section, we will learn how to test views in Django.

Testing Django views

Views in Django control the rendering of the HTTP response for users based on the URL they visit in a web application. In this section, we will learn how to test views inside Django. Imagine you are working on a website where a lot of **application programming interface (API)** endpoints are required. An interesting question to ask is, how will you be able to validate every new endpoint? If done manually, you will have to first deploy the application every time a new endpoint is added, then manually visit the endpoint in the browser to validate if it is working fine or not. Such an approach may work out when the number of endpoints is few but may become extremely cumbersome if there are hundreds of endpoints.

Django provides a very comprehensive way of testing application views. This happens by using a testing client class provided by Django's `test` module. This class can be used to visit URLs mapped to the views and capture the output generated by visiting the URL endpoint. Then, we can use the captured output to test whether the URLs are generating a correct response or not. This client can be used by importing the `Client` class from the Django `test` module and then initializing it, as shown in the following snippet:

```
from django.test import Client

c = Client()
```

The `Client` object supports several methods that can be used to simulate the different HTTP calls a user can make – namely, `GET`, `POST`, `PUT`, `DELETE`, and others. An example of making such a request looks like this:

```
response = c.get('/welcome')
```

The response generated by the view is then captured by the client and gets exposed as a `response` object, which can then be queried to validate the output of the view.

With this knowledge, let's look at how we can write test cases for our Django views.

Exercise 14.03 – writing unit tests for Django views

In this exercise, we will use the Django test client to write a test case for our Django view, which will be mapped to a specific URL. These test cases will help you validate whether your `View` function generates the correct response when visited using its mapped URL. Let's get started with the steps:

1. For this exercise, you are going to use the `bookr_test` application that was created in *step 1 of Exercise 14.02 – testing Django models*. To get started, open the `views.py` file under the `bookr_test` directory and add the following code to it:

```
from django.http import HttpResponse

def greeting_view(request):
    """Greet the user."""
    return HttpResponse("Hey there, welcome to Bookr!
        Your one stop place to review books.")
```

Here, you have created a simple Django view that will be used to greet the user with a welcome message whenever they visit an endpoint mapped to the provided view.

- Once the view has been created, you need to map this view to a URL endpoint that can then be visited in a browser or a test client. To do this, open the `urls.py` file under the `bookr_test` directory and add the highlighted code to the `urlpatterns` list:

```
from django.urls import path
from . import views

urlpatterns = [
    path('test/greeting', views.greeting_view,
         name='greeting_view')
]
```

In the preceding code snippet, we mapped `greeting_view` to the '`test/greeting`' endpoint for the application by setting the path in the `urlpatterns` list.

- Once this path has been set up, we need to make sure that it is also identified by our project. For this, we need to add this entry to the *Bookr* project's URL mapping. To achieve that, open the `urls.py` file in the `bookr` directory and append the following highlighted line to the end of the `urlpatterns` list:

```
urlpatterns = [
    ...,
    ...,
    path('', include('bookr_test.urls'))
]
```

- Once the view has been set up, validate that it works correctly. Do this by running the following command:

```
python manage.py runserver localhost:8080
```

Then, visit `http://localhost:8080/test/greeting` in your web browser. Once the page opens, you should see the following text you added in *step 1* to the greeting view being displayed in the browser:

```
Hey there, welcome to Bookr! Your one stop place to review
books.
```

5. Now, we are ready to write the test cases for `greeting_view`. For this exercise, we are going to write a test case that checks whether, on visiting the `/test/greeting` endpoint, you get a successful result or not. To implement this test case, open the `tests.py` file under the `bookr_test` directory and add the following code to the end of the file:

```
from django.test import TestCase, Client

class TestGreetingView(TestCase):
    """Test the greeting view."""
    def setUp(self):
        self.client = Client()

    def test_greeting_view(self):
        response = self.client.get('/test/greeting')
        self.assertEqual(response.status_code, 200)
```

In the preceding code snippet, we defined a test case that helps validate whether the greeting view is working fine or not.

This is done by first importing Django's test client, which allows you to test views mapped to the URLs by making calls to them and analyzing the generated response:

```
from django.test import TestCase, Client
```

Once the import has been done, we must create a new class named `TestGreetingView` that will group the test cases related to the greeting view we created in *step 2*:

```
class TestGreetingView(TestCase) :
```

Inside this test case, we defined two methods, `setUp()` and `test_greeting_view()`. The `test_greeting_view()` method implements your test case. Inside this, we are first making an HTTP GET call to the URL that is mapped to the greeting view and storing the response generated by the view inside the `response` object created:

```
    response = self.client.get('/test/greeting')
```

Once this call finishes, you will have its HTTP response code, contents, and headers available inside the `response` variable. Next, with the following code, we make an assertion, validating whether the status code generated by the call matches the status code for successful HTTP calls (HTTP 200):

```
        self.assertEqual(response.status_code, 200)
```

With this, we are now ready to run the tests.

- With the test case written, look at what happens when you run the test case. This can be done by running the following command:

```
python manage.py test
```

Once the command executes, you can expect to see an output like the one shown in the following snippet:

```
% python manage.py test
Creating test database for alias 'default'...
System check identified no issues (0 silenced).
...
-----
-----
Ran 3 tests in 0.006s

OK
Destroying test database for alias 'default'...
```

As you can see from the output, your test cases executed successfully, hence validating that the response generated by the `greeting_view()` method is as per your expectations.

In this exercise, we learned how we can implement a test case for a Django view function and use Django's `TestClient` to assert that the output generated by the view function matches the one the developer should see.

In the next section, we will learn how to test views that have authentication enabled.

Testing views with authentication

In the previous example, we looked at how we can test views inside Django. An important thing to note about this view is that the view we created could be accessed by anyone and is not protected by any authentication or login checks. Now, imagine a case where a view should only be accessible if the user is logged in. For example, consider implementing a view function that renders a profile page of a registered user of our web application. To make sure that only logged-in users can view the profile page for their account, you might want to restrict the view to logged-in users only.

With this, we now have an important question: *how can we test views that require authentication?*

Luckily, Django's test client provides this functionality, through which we can log in to our views and then run tests over them. This result can be achieved using Django's test client `login()` method. When this method is called, Django's test client performs an authentication operation against the service. If the authentication succeeds, it stores the login cookie internally, which it can then use for further test runs. The following code snippet shows how you can set up Django's test client to simulate a logged-in user:

```
login = self.client.login(username='testuser',  
password='testpassword')
```

The `login` method requires a username and password for the test user that we are going to test with, as will be shown in the next exercise. So, let's take a look at how we can test a flow that requires user authentication.

Exercise 14.04 – writing test cases to validate authenticated users

In this exercise, we will write test cases for views that require the user to be authenticated. As part of this, we will validate the output generated by the view method when a user who is not logged in tries to visit the page and when a user who is logged in tries to visit the page mapped to the view function. Let's get started with the steps:

1. For this exercise, we are going to use the `bookr_test` application that you created in *step 1 of Exercise 14.02 – testing Django models*. To get started, open the `views.py` file under the `bookr_test` application and add the following code to it:

```
from django.http import HttpResponseRedirect  
from django.contrib.auth.decorators import  
login_required
```

2. Once the preceding code snippet has been added, create a new function called `greeting_view_user()` at the end of the file, as shown in the following code snippet:

```
@login_required  
def greeting_view_user(request):  
    """Greeting view for the user."""  
    user = request.user  
    return HttpResponseRedirect("Welcome to Bookr!  
    {username}".format(username=user))
```

Here, we have created a simple Django view that will be used to greet the logged-in user with a welcome message whenever they visit an endpoint mapped to the provided view.

3. Once this view has been created, we need to map this view to a URL endpoint that can then be visited in a browser or a test client. To do this, open the `urls.py` file under the `bookr_test` directory and add the highlighted code to it:

```
from django.urls import path
from . import views

urlpatterns = [
    path('test/greet_user', views.greeting_view_user,
         name='greeting_view_user')
]
```

In the preceding code snippet, we mapped `greeting_view` to the '`test/greet_user`' endpoint for the application by setting the path in the `urlpatterns` list:

```
path('test/greet_user', views.greeting_view_user,
      name='greeting_view_user')
```

If you have followed the previous exercises, this URL should already be set up for detection in the project and no further steps are required to configure the URL mapping.

4. Once the view has been set up, the next thing we need to do is validate whether it works correctly. To do this, run the following command:

```
python manage.py runserver localhost:8080
```

Then, visit `http://localhost:8080/test/greet_user` in your web browser.

If you are not logged in already, by visiting the preceding URL, you will be redirected to the login page for the project.

5. Now, write the test cases for `greeting_view_user`, which check whether, on visiting the `/test/greet_user` endpoint, you get a successful result or not. To implement this test case, open the `tests.py` file under the `bookr_test` directory and add the following code to it:

```
from django.contrib.auth.models import User

class TestLoggedInGreetingView(TestCase):
    """Test the greeting view for the authenticated
    users."""
    def setUp(self):
        test_user =
            User.objects.create_user(
                username='testuser',
                password='test@#628password')
        test_user.save()
        self.client = Client()

    def test_user_greeting_not_authenticated(self):
```

```
response = self.client.get('/test/greet_user')
self.assertEqual(response.status_code, 302)

def test_user_authenticated(self):
    login = self.client.login(username='testuser',
                             password='test@#628password')
    response = self.client.get('/test/greet_user')
    self.assertEqual(response.status_code, 200)
```

In the preceding code snippet, we implemented a test case that checks the views that have authentication enabled before their content can be seen.

Here, first, we imported the required classes and methods that will be used to define the test case and initialize a testing client:

```
from django.test import TestCase, Client
```

The next thing you require is the `User` model from Django's `auth` module:

```
from django.contrib.auth.models import User
```

This model is required because, for the test cases requiring authentication, we will need to initialize a new test user. Next up, we created a new class named `TestLoggedInGreetingView`, which wraps your tests related to the `greeting_user` view (which requires authentication). Inside this class, you defined three methods called `setUp()`, `test_user_greeting_not_authenticated()`, and `test_user_authenticated()`. The `setUp()` method is used to initialize a test user that you will use for authentication. This is a required step because a test environment inside Django is a completely isolated environment that doesn't use data from your production application, so all the required models and objects are to be instantiated separately inside the test environment.

Then, we created the test user and initiated the test client using the following code:

```
test_user = User.objects.create_user(
    username='testuser', password='test@#628password')
test_user.save()
self.client = Client()
```

Next, we wrote the test case for the `greet_user` endpoint when the user is not authenticated. Here, you should expect Django to redirect the user to the login endpoint. This redirect can be detected by checking the HTTP status code of the response, which should be set to HTTP 302, indicating a redirect operation:

```
def test_user_greeting_not_authenticated(self):  
    response = self.client.get('/test/greet_user')  
    self.assertEqual(response.status_code, 302)
```

Next, we wrote another test case to check whether the `greet_user` endpoint renders successfully when the user is authenticated. To authenticate the user, you must first call the `login()` method of the test client and perform authentication by providing the username and password of the test user you created in the `setUp()` method, as follows:

```
login = self.client.login(username='testuser',  
                         password='test@#628password')
```

Once the login has been completed, we must make an HTTP GET request to the `greet_user` endpoint and validate whether the endpoint generates a correct result or not by checking the HTTP status code of the returned response:

```
response = self.client.get('/test/greet_user')  
self.assertEqual(response.status_code, 200)
```

6. With the test cases written, it's time to check how they run. For this, run the following command:

```
python manage.py test
```

Once the execution finishes, we can expect to see a response that resembles the following one:

```
% python manage.py test  
Creating test database for alias 'default'...  
System check identified no issues (0 silenced).  
.....  
-----  
-----  
Ran 5 tests in 0.366s  
  
OK  
Destroying test database for alias 'default'...
```

As we can see from the preceding output, our test cases have passed successfully, validating that the view we created generates the desired response of redirecting the user if the user is unauthenticated to the website, and allowing the user to see the page if the user is authenticated.

In this exercise, we implemented a test case where we test what output is generated by a view function based on whether the user is authenticated or not.

In the next section, we will explore how to create requests in Django unit tests without the use of a Django-provided test client, but by providing raw request data.

Django RequestFactory

So far, we have been using Django's test client to test the views that we have created for our application. The test client class simulates a browser and uses this simulation to make calls to the required APIs. But what if we didn't want to use the test client and its associated simulation of being a browser, but rather wanted to test the view functions directly bypassing the request parameter? How can we do that?

To help us in such cases, we can leverage the `RequestFactory` class provided by Django. The `RequestFactory` class helps us provide the `request` object that we can pass to our view functions to evaluate whether they are working. The following object for the `RequestFactory` class can be created by instantiating the class as follows:

```
factory = RequestFactory()
```

The factory object created only supports HTTP methods such as `get()`, `post()`, `put()`, and others to simulate a call to any URL endpoint. Now, let's learn how to modify the test case that we wrote in *Exercise 14.04 – writing test cases to validate authenticated users*, so that it uses `RequestFactory`.

Exercise 14.05 – using RequestFactory to test views

In this exercise, we will use `RequestFactory` to test view functions in Django:

1. For this exercise, we will use the existing `greeting_view_user` view function that we created earlier in *step 1 of Exercise 14.04 – writing test cases to validate authenticated users*, which looks like this:

```
@login_required
def greeting_view_user(request):
    """Greeting view for the user."""
    user = request.user
    return HttpResponse("Welcome to Bookr!
        {username}".format(username=user))
```

2. Next, modify the existing test case, `TestLoggedInGreetingView`, defined inside the `tests.py` file under the `bookr_test` directory. Open the `tests.py` file and make the following changes:

- First, we need to add the following import to use the `RequestFactory` class inside the test cases:

```
from django.test import RequestFactory
```

- The next thing we need is an import for the `AnonymousUser` class from Django's auth module and the `greeting_view_user` view method from the `views` module. This is required to test the view functions with a simulated user who is not logged in. This can be done by adding the following code:

```
from django.contrib.auth.models import AnonymousUser
from .views import greeting_view_user
```

3. Once the `import` statements have been added, modify the `setUp()` method of the `TestLoggedInGreetingView` class and change its contents so that it resembles the following:

```
def setUp(self):
    self.test_user =
        User.objects.create_user(username='testuser',
                               password='test@#628password')
    self.test_user.save()
    self.factory = RequestFactory()
```

In this method, we first created a user object and stored it as a class member so that we can use it later in the tests. Once the user object was created, we instantiated a new instance of the `RequestFactory` class to test our view function.

4. With the `setUp()` method now defined, modify the existing tests so that they use the request factory instance. For the test for a non-authenticated call to the view function, modify the `test_user_greeting_not_authenticated` method so that it has the following contents:

```
def test_user_greeting_not_authenticated(self):
    request = self.factory.get('/test/greet_user')
    request.user = AnonymousUser()
    response = greeting_view_user(request)
    self.assertEqual(response.status_code, 302)
```

In this method, we first created a request object using the `RequestFactory` instance we defined in the `setUp()` method. Once we had done that, we assigned an `AnonymousUser()` instance to the `request.user` property. Assigning the `AnonymousUser()` instance to the property makes the view function think that the user making the request is not logged in:

```
request.user = AnonymousUser()
```

Once we did this, we made a call to the `greeting_view_user()` view method and passed to it the request object you created. Once the call is successful, we must capture the output of the method in the `response` variable using the following code:

```
response = greeting_view_user(request)
```

For the unauthenticated user, we expect to get a redirect response that can be tested by checking the HTTP status code of the response, as follows:

```
self.assertEqual(response.status_code, 302)
```

5. Now, go ahead and modify the other method, `test_user_authenticated()`, similarly by using the `RequestFactory` instance, as follows:

```
def test_user_authenticated(self):  
    request = self.factory.get('/test/greet_user')  
    request.user = self.test_user  
    response = greeting_view_user(request)  
    self.assertEqual(response.status_code, 200)
```

As you can see, most of the code resembles what you wrote in the `test_user_greeting_not_authenticated` method, with a small change – in this method, instead of using `AnonymousUser` for the `request.user` property, you use the `test_user` that you created in the `setUp()` method:

```
request.user = self.test_user
```

With the changes done, it's time to run the tests.

6. To run the tests and validate whether `RequestFactory` works as expected or not, run the following command:

```
python manage.py test
```

Once the command executes, you can expect to see an output that resembles the one shown here:

```
% python manage.py test  
Creating test database for alias 'default'...  
System check identified no issues (0 silenced).  
.....  
-----  
-----  
Ran 5 tests in 0.229s
```

```
OK
Destroying test database for alias 'default'...
```

As we can see from the output, the test cases written by us have passed successfully, thus validating the behavior of the `RequestFactory` class.

With this exercise, we learned to write test cases for view functions by leveraging `RequestFactory` and passing the request object directly to the view function, rather than simulating a URL visit using the test client approach, thus allowing for more direct testing.

In the next section, we will test class-based views.

Testing class-based views

In the previous exercise, we saw how we can test views defined as methods. But what about class-based views? How can we test those?

As it turns out, it is quite easy to test class-based views. For example, if we have a class-based view defined with the name `ExampleClassView` (`View`), to test the view, all we need to do is to use the following syntax:

```
response = ExampleClassView.as_view()(request)
```

It is as simple as that.

A Django application generally consists of several different components that can either work in isolation, such as models, and some other components that need to interact with the URL mapping and other parts of the framework to work. Testing these different components may require some steps that are common to only those components. For example, when testing a model, first, we may want to create certain objects of the `Model` class before we start testing, or for views, we may want to initialize a test client with user credentials first.

As it turns out, Django also provides some other classes based on the `TestCase` class that can be used to write test cases of specific types about the type of component being used. Let's look at the different classes provided by Django.

Test case classes in Django

Beyond the base `TestCase` class provided by Django, which can be used to define a multitude of test cases for different components, Django also provides some specialized classes derived from the `TestCase` class.

These classes are used for specific types of test cases based on the capabilities they provide to the developer. Let's look at the different classes provided by Django.

The `SimpleTestCase` class

This class is derived from the `TestCase` class provided by Django's `test` module and should be used when writing simple test cases that test the view functions. Usually, this class is not preferred when your test case involves making database queries. The class also provides a lot of useful features, such as these:

- Ability to check for exceptions raised by a view function
- Ability to test form fields
- A built-in test client
- Ability to verify a redirect by a view function
- Can match the equality of two HTML, JSON, or XML outputs generated by the view functions

Now that we have a basic idea of what a `SimpleTestCase` class is, let's move on and try to understand another type of test case class that helps in writing test cases about the interaction with databases.

The `TransactionTestCase` class

This class is derived from the `SimpleTestCase` class and should be used when writing test cases that involve interaction with the database, such as database queries, model object creations, and more.

The class provides the following added features:

- Ability to reset the database to a default state before a test case runs
- Can skip tests based on database features, which can be quite handy if the database being used for testing does not support all the features of your production database

The `LiveServerTestCase` class

This class is like the `TransactionTestCase` class, but with a small difference that the test cases written in the class, instead of using the default test client, use a live server created by Django.

This ability to run the live server for testing comes in handy when writing test cases that test for the rendered web pages and any interaction with them, which is not possible while using the default test client.

Such test cases can leverage tools such as Selenium, which can be used to build interactive test cases that modify the state of the rendered page by interacting with it.

Modularizing test code

In the previous exercises, we saw how we can write test cases for different components of our project. But an important aspect to note is, till now, we have written the test cases for all the components in a single file.

This approach is okay when the application does not have a lot of views and models, but this can become problematic as our application grows because now, our single `tests.py` file will be hard to maintain.

To avoid running into such scenarios, we should try to modularize our test cases so that the test cases for models are kept separately from test cases related to the views and other properties. To achieve this modularization, all we need to do is follow two simple steps:

1. Create a new directory named `tests` inside our application directory by running the following command:

```
mkdir tests
```

2. Create a new empty file named `__init__.py` inside our `tests` directory by running the following command:

```
touch __init__.py
```

This `__init__.py` file is required by Django to correctly detect the `tests` directory we created as a module and not a regular directory.

Once you have completed the preceding steps, you can go ahead and create new testing files for the different components in your application. For example, to write test cases for your models, you can create a new file named `test_models.py` inside the `tests` directory and add any associated code for your model testing inside this file.

Also, you don't need to take any other additional steps to run your tests. The same command will work perfectly fine for your modular testing code base as well:

```
python manage.py test
```

With this, we know how we can write test cases for our projects. So, how about we test this knowledge and write test cases for the Bookr project we are working on?

Activity 14.01 – testing models and views in Bookr

In this activity, you will implement test cases for the Bookr project. You will implement test cases to validate the functioning of the models created inside the *reviews* application of the Bookr project, and then you will implement a simple test case for validating the *index view* inside the *reviews* application.

The following steps will help you work through this activity:

1. Create a directory named `tests` inside the *reviews* application directory so that all our test cases for the reviews application can be modularized.
2. Create an empty `__init__.py` file so that the directory is considered not as a general directory, but rather a Python module directory.
3. Create a new file named `test_models.py` for implementing the code for testing the models. Inside this file, import the models you want to test.
4. Inside the `test_models.py` file, create a new class that inherits from the `TestCase` class of the `django.tests` module and implements methods to validate the process of creating and reading the `Model` objects.
5. To test the view function, create a new file named `test_views.py` inside the `tests` directory that was created in *step 1*.
6. Inside the `test_views.py` file, import the `test Client` class from the `django.tests` module and the `index` view function from the *reviews* application's `views.py` file.
7. Inside the `test_views.py` file you created in *step 5*, create a new `TestCase` class and implement methods to validate the `index` view.
8. Inside the `TestCase` class you created in *step 7*, create a new function called `setUp()`, inside which you should initialize an instance of `RequestFactory`, which will be used to create a request object that can be directly passed to the view function for testing.
9. Once you have completed the previous steps are written the test cases, run them by executing `python manage.py` to validate that the test cases pass.

Upon completing this activity, all test cases should pass successfully.

Important note

The solution for this activity can be found at on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we learned how to write test cases for different components of our web application project with Django. We learned why testing plays a crucial role in the development of any web application and the different types of testing techniques that are employed in the industry to make sure the application code they ship is stable and bug-free.

Then, we looked at how we can use the `TestCase` class provided by Django's `test` module to implement our unit tests, which can be used to test the models as well as views. We also looked at how we can use Django's `test` client to test our view functions, some of which require the user to be authenticated. We also glanced over another approach of using `RequestFactory` to test method views and class-based views.

We concluded this chapter by understanding the predefined classes provided by Django and where they should be used, while also looking at how we can modularize our testing code base to make it appear clean.

As we move to the next chapter, we will try to understand how we can make our Django application more powerful by integrating third-party libraries into our project. This functionality will be used to implement third-party authentication into our Django application and allow the users to log in to the application using popular services such as Google Sign-In, Facebook Sign-In, and others.

15

Django Third-Party Libraries

Because Django has been around since 2007, there is a rich ecosystem of third-party libraries that can be plugged into an application to give it extra features. So far, we have learned a lot about Django and used many of its features, including database models, URL routing, templating, forms, and more. We used these Django tools directly to build a web app, but now we will look at how to leverage the work of others to quickly add even more advanced features to our own apps. We have alluded to apps for storing files (in *Chapter 5, Serving Static Files*, we mentioned an app, `django-storages`, that can store our static files in a CDN), but in addition to file storage, we can also use apps to plug into third-party authentication systems, integrate with payment gateways, customize how our settings are built, modify images, build forms more easily, debug our site, use different types of databases, and much more. Chances are that if you want to add a certain feature, an app exists for it.

We don't have space to cover every app in this chapter, so we'll just focus on four that provide useful features across many different types of apps. `django-configurations` allows you to configure your Django settings using classes and take advantage of inheritance to simplify settings for different environments. This works in tandem with `dj-database-urls` to specify your database connection setting using just a URL. The *Django Debug Toolbar* lets you get extra information to help with debugging, right in your browser. The last app we'll look at is `django-crispy-forms`, which provides extra CSS classes to make forms look nicer, as well as making them easier to configure using just Python code.

For each of these libraries, we will cover its installation, basic setup, and use, mostly as they apply to Bookr. They also have more configuration options to further customize to fit your application. Each of these apps can be installed with `pip`.

We will also briefly introduce `django-allauth`, which allows a Django application to authenticate users against third-party providers (such as Google, GitHub, Facebook, and Twitter). We won't cover its installation and setup in detail but will provide some examples to help you configure it.

In this chapter, we will cover the following topics:

- Environment variables
- `django-configurations`
- `dj-database-urls`
- The Django Debug Toolbar
- `django-crispy-forms`
- `django-allauth`

Technical requirements

All the code files for this chapter can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter15>.

Environment variables

When we create a program, we often want the user to be able to configure some of its behavior. For example, say you have a program that connects to a database and saves all the records it finds into a file. Normally, it would probably print out just a *success* message to the terminal, but you might also want to run it in *debug mode*, which also makes it print out all the SQL statements it is executing.

There are many ways of configuring a program like this. For example, you could have it read from a configuration file. But in some cases, the user may quickly want to run the Django server with a particular setting on (say, debug mode), and then run the server again with the same setting off. Having to change the configuration file every time can be inconvenient. In this case, we can read from an *environment variable*. Environment variables are key-value pairs that can be set in your operating system and then read by a program. There are several ways they can be set:

- Your shell (terminal) can read variables from a profile script when it starts, then each program will have access to these variables.
- You can set a variable inside a terminal, and it will be made available to any programs that start subsequently. In Linux and macOS, this is done with the `export` command; Windows uses the `set` command. Any variables you set in this way override those in the profile script, but only for the current session. When you close the terminal, the variables are lost.
- You can set environment variables at the same time as running a command in a terminal. These will only persist for the program being run, and they override exported environment variables and those read from a profile script.
- You can set environment variables inside a running program, and they will be available only inside the program (or for programs that your program starts). Environment variables set in this way will override all the other methods we have just set.

These might sound complicated, but we will explain them with a short Python script and show how variables can be set in the last three ways (the first method depends on what shell you use). The script will also show how environment variables are read.

Environment variables are available in Python using the `os.environ` variable. This is a dictionary-like object that can be used to access environment variables by name. It is safest to access values using the `get` method, just in case they are not set. It also provides a `setdefault` method, which allows you to set a value only if it is not set (that is, it doesn't overwrite an existing key).

Here is the example Python script that reads environment variables:

```
import os

# This will set the value since it's not already set
os.environ.setdefault('UNSET_VAR', 'UNSET_VAR_VALUE')

# This value will not be set since it's already passed
# in from the command line
os.environ.setdefault('SET_VAR', 'SET_VAR_VALUE')

print(f"UNSET_VAR: {os.environ.get('UNSET_VAR', '')}")
print(f"SET_VAR: {os.environ.get('SET_VAR', '')}")

# All these values were provided from the shell in some way
print(f"HOME: {os.environ.get('HOME', '')}")
print(f"VAR1: {os.environ.get('VAR1', '')}")
print(f"VAR2: {os.environ.get('VAR2', '')}")
print(f"VAR3: {os.environ.get('VAR3', '')}")
print(f"VAR4: {os.environ.get('VAR4', '')}")
```

We then set up our shell by setting some variables. In Linux or macOS, we use `export` (note there is no output from these commands):

```
$ export SET_VAR="Set Using Export"
$ export VAR1="Set Using Export"
$ export VAR2="Set Using Export"
```

In Windows, we will use the `set` command in the command line, as follows:

```
set SET_VAR="Set Using Export"
set VAR1="Set Using Export"
set VAR2="Set Using Export"
```

In Linux and macOS, we can also provide environment variables by setting them before the command (the actual command is just `python3 env_example.py`):

```
$ VAR2="Set From Command Line" VAR3="Also Set From Command Line"  
python3 env_example.py
```

Important note

Note that the preceding command will not work on Windows. For Windows, the environment variables must be set before execution and cannot be passed in at the same time.

The output from this preceding command is as follows:

```
UNSET_VAR:UNSET_VAR_VALUE  
SET_VAR:Set Using Export  
HOME: /Users/ben  
VAR1: Set Using Export  
VAR2: Set From Command Line  
VAR3: Also Set From Command Line  
VAR4:
```

This is how the environment variables are interpreted by the Python script:

- When the script runs `os.environ.setdefault('UNSET_VAR', 'UNSET_VAR_VALUE')`, the value is set inside the script, since no value for `UNSET_VAR` was set by the shell. The value that is output is the one set by the script itself.
- When `os.environ.setdefault('SET_VAR', 'SET_VAR_VALUE')` is executed, the value is not set, since one was provided by the shell. This was set with the `export SET_VAR="Set Using Export"` command.
- The value for `HOME` was not set by any of the commands that were run – this is one provided by the shell. It is the user's home directory. This is just an example of an environment variable that a shell normally provides.
- `VAR1` was set by `export` and was not overridden when executing the script.
- `VAR2` was set by `export` but was subsequently overridden when executing the script.
- `VAR3` was only set when executing the script.
- `VAR4` was never set – we use the `get` method to access it to avoid a `KeyError`.

Now that environment variables have been covered, we can return to discussing the changes that need to be made to `manage.py` to support `django-configurations`.

django-configurations

One of the main considerations when deploying a Django application to production is how to configure it. As you have seen throughout this book, the `settings.py` file is where all your Django configuration is defined. Even third-party apps have their configuration in this file. You have already seen this in *Chapter 12, Building a REST API*, when working with the Django REST framework.

There are many ways to provide different configurations and switch between them in Django. If you have begun working on an existing application that already has a specific method of switching between configurations in development and production environments, then you should probably keep using that method.

When we release Bookr onto a product web server, in *Chapter 17, Deploying a Django Application (Part 1 – Server Setup)*, we will need to switch to a production configuration, and that's when we will use `django-configurations`.

To install `django-configurations`, use `pip3` as follows:

```
pip3 install django-configurations
```

Important note

For Windows, you can use `pip` instead of `pip3` in the preceding command.

The output will be as follows:

```
Collecting django-configurations
  Downloading django_configurations-2.4-py3-none-any.whl
    (17 kB)
Requirement already satisfied: django>=3.2 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django-configurations) (4.0)
Requirement already satisfied: sqlparse>=0.2.2 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django>=3.2->django-configurations) (0.4.2)
Requirement already satisfied: asgiref<4,>=3.4.1 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django>=3.2->django-configurations) (3.4.1)
Requirement already satisfied: backports.zoneinfo; python_version < "3.9" in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django>=3.2->django-configurations) (0.2.1)
Installing collected packages: django-configurations
Successfully installed django-configurations-2.4
```

`django-configurations` changes your `settings.py` file so that all the settings are read from a class you define, which will be a subclass of `configurations.Configuration`. Instead of the settings being global variables inside `settings.py`, they will be attributes of the class you define. By using this class-based method, we can take advantage of object-oriented paradigms, most

notably inheritance. Settings, defined in a class, can inherit settings from another class. For example, the production settings class can inherit the development settings class and just override some specific settings – such as forcing DEBUG to be `False` in production.

We can illustrate what needs to be done to the settings file by just showing the first few settings in the file. A standard Django `settings.py` file normally starts like this (comment lines have been removed):

```
from pathlib import Path

BASE_DIR = Path( file ).resolve().parent.parent
SECRET_KEY = "django-insecure-c!...-pz6s^_k6w#eo7&"
DEBUG = True
# The rest of the settings are not shown
```

To convert the settings into django-configurations, first import `Configuration` from `configurations`. Then, define a `Configuration` subclass. Finally, indent all the settings to be under the class. In PyCharm, this is as simple as selecting all the settings and pressing `Tab` to indent them all.

After doing this, your `settings.py` file will look like this:

```
from pathlib import Path
from configurations import Configuration

class Dev(Configuration):
    BASE_DIR = Path(__file__).resolve().parent.parent
    SECRET_KEY = "django-insecure-c!...&xxf*skpgkey6%$ld-
    pz6s^_k6w#eo7&"
    DEBUG = True
    ...
    # All other settings indented in the same manner
```

To have different configurations (different sets of settings), you can just extend your configuration classes and override the settings that should differ.

For example, one variable that needs overriding in production is `DEBUG`; it should be `False` (for security and performance reasons). A `Prod` class can be defined that extends `Dev` and sets `DEBUG`, like this:

```
class Dev(Configuration):
    DEBUG = True
    ...
    # Other settings truncated
```

```
class Prod(Dev):
    DEBUG = False
    # no other settings defined since we're only overriding
    DEBUG
```

Of course, you can override other production settings too, not just DEBUG. Usually, for security, you would also redefine SECRET_KEY and ALLOWED_HOSTS, and to configure Django to use your production database, you'd set the DATABASES value too. Any Django setting can be configured as you choose.

If you try to execute `runserver` (or other management commands) now, you will get an error because Django doesn't know how to find the `settings.py` file when the settings files are laid out like this:

```
django.core.exceptions.ImproperlyConfigured: django-configurations
settings importer wasn't correctly installed. Please use one of the
starter functions to install it as mentioned in the docs: https://
django-configurations.readthedocs.io/
```

We need to make some changes to the `manage.py` file before it starts to work again.

manage.py changes

There are two lines that need to be added/changed in `manage.py` to enable `djangocfg`. First, we need to define a default environment variable that tells the Django configuration which Configuration class it should load.

This line should be added in the `main()` function to set the default value for the `DJANGO_CONFIGURATION` environment variable:

```
os.environ.setdefault('DJANGO_CONFIGURATION', 'Dev')
```

This sets the default to `Dev` – the name of the class we defined. As we saw in our example script, if this value is already defined, it won't be overwritten. This will allow us to switch between configurations using an environment variable.

The second change is to swap the `execute_from_command_line` function with one that `djangocfg` provides. Consider the following line:

```
from django.core.management import execute_from_command_line
```

This line is changed as follows:

```
from configurations.management import execute_from_command_line
```

From now on, `manage.py` will work as it did before, except it now prints out which Configuration class it's using when it starts (*Figure 15.1*):

```
[(bookr) ~ bookr python manage.py runserver
django-configurations version 2.4, using configuration Dev
Watching for file changes with StatReloader
Performing system checks...

System check identified no issues (0 silenced).
December 31, 2022 - 02:18:39
Django version 4.0, using settings 'bookr.settings'
Starting development server at http://127.0.0.1:8000/
Quit the server with CONTROL-C.
```

Figure 15.1 – django-configurations using configuration Dev

In the second line, you can see the `django-configurations` output that uses the `Dev` class for settings.

Configuration from environment variables

As well as switching between Configuration classes using environment variables, `django-configurations` lets us give values for individual settings using environment variables. It provides Value classes that will automatically read values from the environment. We can define defaults if no values are provided. Since environment variables are always strings, the different Value classes are used to convert from a string into the specified type.

Let's look at this in practice with a few examples. We will allow `DEBUG`, `ALLOWED_HOSTS`, `TIME_ZONE`, and `SECRET_KEY` to be set with environment variables as follows:

```
from configurations import Configuration, values

class Dev(Configuration):
    DEBUG = values.BooleanValue(True)
    ALLOWED_HOSTS = values.ListValue([])
    TIME_ZONE = values.Value('UTC')
    SECRET_KEY = "django-insecure-c!...-pz6s^_k6w#eo7&"
    ...
    # Other settings truncated

class Prod(Dev):
    DEBUG = False
    SECRET_KEY = values.SecretValue()
    # no other settings are present
```

We'll explain the settings one at a time:

- In Dev, DEBUG is read from an environment variable and cast to a Boolean value. The values `yes`, `y`, `true`, and `1` become `True`; the values `no`, `n`, `false`, and `0` become `False`. This allows us to run with DEBUG off, even on a development machine, which can be useful in some cases (for example, testing a custom exception page rather than Django's default one). In the Prod configuration, we don't want DEBUG to accidentally become `True`, so we set it statically.
- ALLOWED_HOSTS is required in production. It is a list of hosts for which Django should accept requests.
- The `ListValue` class will convert a comma-separated string into a Python list.
For example, the `www.example.com`, `example.com` string is converted into `["www.example.com", "example.com"]`.
- TIME_ZONE accepts just a string value, so it is set using the `Value` class. This class just reads the environment variable and does not transform it at all.
- SECRET_KEY is statically defined in the Dev configuration; it can't be changed with an environment variable. In the Prod configuration, it is set with `SecretValue`. This is like `Value` in that it is just a string setting; however, it does not allow a default. If a default is set, then an exception is raised. This is to ensure you don't ever put a secret value into `settings.py`, since it might be accidentally shared (for example, uploaded to GitHub). Note that since we do not use `SECRET_KEY` for Dev in production, we don't care if it's leaked.

By default, `django-configurations` expects the `DJANGO_` prefix for each environment variable. For example, to set DEBUG, use the `DJANGO_DEBUG` environment variable, and to set ALLOWED_HOSTS, use `DJANGO_ALLOWED_HOSTS`.

Now that we've introduced `django-configurations` and the changes that need to be made to a project to support it, let's add it to Bookr and make those changes. In the next exercise, we will install and set up `django-configurations` in Bookr.

Exercise 15.01 – Django configurations setup

In this exercise, we will install `django-configurations` using pip, and then update `settings.py` to add a Dev and Prod configuration. We'll then make the necessary changes to `manage.py` to support the new configuration style, and test that everything is still working. Let's get started with the steps:

1. In a terminal, make sure you have activated the `bookr` virtual environment, and then run the following command to install `django-configurations` using pip3:

```
pip3 install django-configurations
```

Note

For Windows, you can use pip instead of pip3 in the preceding command.

The install process will run, and you should have an output like *Figure 15.2*:

```
(bookr) ~ bookr pip3 install django-configurations
Collecting django-configurations
  Using cached django_configurations-2.4-py3-none-any.whl (17 kB)
Collecting django>=3.2
  Downloading Django-4.1.4-py3-none-any.whl (8.1 MB)
                                         8.1/8.1 MB 6.9 MB/s eta 0:00:00
Collecting sqlparse>=0.2.2
  Using cached sqlparse-0.4.3-py3-none-any.whl (42 kB)
Collecting asgiref<4,>=3.5.2
  Downloading asgiref-3.6.0-py3-none-any.whl (23 kB)
Installing collected packages: sqlparse, asgiref, django, django-configurations
Successfully installed asgiref-3.6.0 django-4.1.4 django-configurations-2.4 sqlparse-0.4.3
```

Figure 15.2 – django-configurations installation with pip

2. In PyCharm, open `settings.py` inside the `bookr` package. Underneath the existing `os` import, import `Configuration` and `values` from `configurations`, like this:

```
from configurations import Configuration, values
```

3. After the imports but before your first setting definition (the line that sets the `BASE_DIR` value), add a new `Configuration` subclass, called `Dev`:

```
class Dev(Configuration):
```

4. Now, we need to move all the existing settings so that they are attributes of the `Dev` class rather than global variables. In PyCharm, this is as simple as selecting all the settings and then pressing the `Tab` key to indent them. After doing this, your settings should look as follows:

```
from pathlib import Path
from configurations import Configuration, values

class Dev(Configuration):
    # Build paths inside the project like this: BASE_DIR / 'subdir'.
    BASE_DIR = Path(__file__).resolve().parent.parent

    # Quick-start development settings - unsuitable for production
    # See https://docs.djangoproject.com/en/4.0/howto/deployment/checklist/

    # SECURITY WARNING: keep the secret key used in production secret!
    SECRET_KEY = "django-insecure-c1ed4dqj0ous$1*xzf&xxf*skgkey6%$1d-pz6s^_k6w#eo7&"

    # SECURITY WARNING: don't run with debug turned on in production!
    DEBUG = True
```

Figure 15.3 – The new `Dev` configuration

5. After indenting the settings, we will change some of the settings to be read from environment variables. First, change DEBUG to be read as BooleanValue. It should default to True. Consider this line:

```
DEBUG = True
```

Then, change it to this:

```
DEBUG = values.BooleanValue(True)
```

This will automatically read DEBUG from the DJANGO_DEBUG environment variable and convert it into a Boolean. If the environment variable is not set, then it will default to True.

6. Also, convert ALLOWED_HOSTS to be read from an environment variable, using the values.ListValue class. It should default to [] (an empty list). Consider the following line:

```
ALLOWED_HOSTS = []
```

Change it to this:

```
ALLOWED_HOSTS = values.ListValue([])
```

ALLOWED_HOSTS will be read from the DJANGO_ALLOWED_HOSTS environment variable, and default to an empty list.

7. Everything you have done so far has involved adding/changing attributes on the Dev class. Now, at the end of the same file, add a Prod class that inherits from Dev. It should define two attributes, DEBUG = True and SECRET_KEY = values.SecretValue(). The completed class should look like this:

```
class Prod(Dev):  
    DEBUG = False  
    SECRET_KEY = values.SecretValue()
```

Save `settings.py`.

8. If we try to run any management command now, we will receive an error that django-configurations is not set up properly. We need to make some changes to `manage.py` to make it work again. Open `manage.py` in the `bookr` project directory.

Consider the line that reads as follows:

```
os.environ.setdefault('DJANGO_SETTINGS_MODULE',  
                      'bookr.settings')
```

Under it, add this line:

```
os.environ.setdefault('DJANGO_CONFIGURATION', 'Dev')
```

This will set the default configuration to the Dev class. It can be overridden by setting the DJANGO_CONFIGURATION environment variable (for example, to Prod).

9. Two lines below the line from the previous step, you should already have the following `import` statement:

```
from django.core.management import execute_from_command_line
```

Change this to the following:

```
from configurations.management import execute_from_command_line
```

This will make the `manage.py` script use the Django configuration's `execute_from_command_line` function, instead of the Django built-in one.

Save `manage.py`.

10. Start the Django dev server. If it begins without error, you can be confident that the changes you made have worked. To be sure, check that the pages load in your browser. Open `http://127.0.0.1:8000/` and try browsing around the site. Everything should look and feel as it did before:

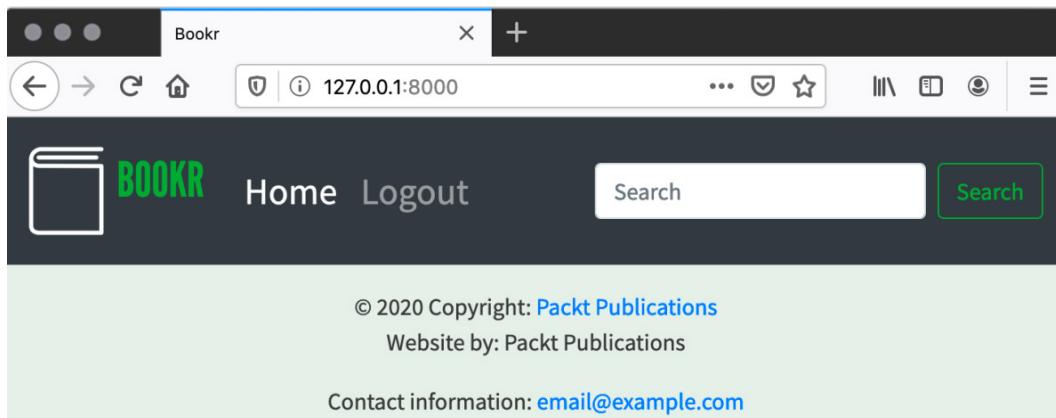


Figure 15.4 – The Bookr site should look and feel as it did before

In this exercise, we installed `django-configurations` and refactored our `settings.py` file to use its `Configuration` class to define our settings. We added the `Dev` and `Prod` configurations and made `DEBUG`, `ALLOWED_HOSTS`, and `SECRET_KEY` settable with environment variables. Finally, we updated `manage.py` to use the Django configuration's `execute_from_command_line` function, which enabled the use of this new `settings.py` format.

In the next section, we will cover `dj-database-url`, a package that makes it possible to configure your Django database settings using URLs.

dj-database-url

`dj-database-url` is another app that helps with the configuration of your Django application. Specifically, it allows you to set the database (that your Django app connects to) using a URL instead of a dictionary of configuration values. As you can see in your existing `settings.py` file, the `DATABASES` setting contains a couple of items and gets more verbose when using a different database that has more configuration options (for username, password, and so on). We can instead set these from a URL, which can contain all these values.

The URL's format will differ slightly, depending on whether you are using a local SQLite database or a remote database server. To use SQLite on disk (as Bookr is currently), the URL is like this:

```
sqlite:///<path>
```

Note there are three slashes present. This is because SQLite doesn't have a hostname, so this is like a URL being like this:

```
<protocol>://<hostname>/<path>
```

That is, the URL has a blank hostname. All three slashes are, therefore, together.

To build a URL for a remote database server, the format is usually like this:

```
<protocol>://<username>:<password>@<hostname>:<port>/  
<database_name>
```

For example, to connect to a PostgreSQL database called `bookr_django` on the host, `db.example.com`, on port 5432, with the username `bookr` and the password `b00ks`, the URL would be like this:

```
postgres://bookr:b00ks@db.example.com:5432/bookr_django
```

Now that we've seen the format for URLs, let's look at how we can actually use them in our `settings.py` file. First, `dj-database-url` must be installed using `pip3`:

```
pip3 install dj-database-url
```

Note

For Windows, you can use `pip` instead of `pip3` in the preceding command.

The output is as follows:

```
Collecting dj-database-url
  Downloading https://files.pythonhosted.org/packages/d4/
  a6/4b8578c1848690d0c307c7c0596af2077536c9ef2a04d42b00fabaa
  7e49d/dj_database_url-0.5.0-py2.py3-none-any.whl
Installing collected packages: dj-database-url
Successfully installed dj-database-url-0.5.0
```

Now, `dj_database_url` can be imported into `settings.py`, and the `dj_database_url.parse` method can be used to transform the URL into a dictionary that Django can use. We can use its return value to set the `default` (or other) item in the `DATABASES` dictionary:

```
import dj_database_url

DATABASES = {'default': dj_database_url.parse(
    'postgres://bookr:b00ks@db.example.com:5432/bookr_django'
)}
```

Alternatively, for our SQLite database, we can utilize the `BASE_DIR` setting, as we are already, and include it in the URL:

```
import dj_database_url

DATABASES = {'default': dj_database_url.parse(
    f'sqlite:///{{BASE_DIR}}/db.sqlite3')}
```

After parsing, the `DATABASES` dictionary is similar to what we defined before. It includes some redundant items that do not apply to an SQLite database (`USER`, `PASSWORD`, `HOST`, and so on), but Django will ignore them:

```
DATABASES = {
    'default': {
        'NAME': '/Users/ben/bookr/bookr/db.sqlite3',
        'USER': '',
        'PASSWORD': '',
        'HOST': '',
        'PORT': '',
        'CONN_MAX_AGE': 0,
        'ENGINE': 'django.db.backends.sqlite3'
    }
}
```

This method of setting the database connection information is not that useful, since we are still statically defining the data in `settings.py`. The only difference is that we are using a URL instead

of a dictionary. `dj-database-url` can also automatically read the URL from an environment variable. This will allow us to override these values by setting them in the environment.

To read the data from the environment, use the `dj_database_url.config` function, like this:

```
import dj_database_url

DATABASES = {'default': dj_database_url.config()}
```

The URL is automatically read from the `DATABASE_URL` environment variable.

We can improve on this by also providing a `default` argument to the `config` function. This is the URL that will be used by default if one is not specified in an environment variable:

```
import dj_database_url

DATABASES = {
    'default': dj_database_url.config(
        default=f'sqlite:///{BASE_DIR}/db.sqlite3')}
```

This way, we can specify a default URL that can be overridden by an environment variable in production.

We can also specify the environment variable that the URL is read from by passing in the `env` argument – this is the first positional argument. In this way, you could read multiple URLs for different database settings:

```
import dj_database_url

DATABASES = {
    'default': dj_database_url.config(
        default=f'sqlite:///{BASE_DIR}/db.sqlite3'),
    'secondary': dj_database_url.config(
        'DATABASE_URL_SECONDARY',
        default=f'sqlite:///{BASE_DIR}/db-secondary.sqlite3')
}
```

In this example, the `default` item's URL is read from the `DATABASE_URL` environment variable, and `secondary` is read from `DATABASE_URL_SECONDARY`.

`django-configurations` also provides a `config` class that works in tandem with `dj_database_url`: `DatabaseURLValue`. This differs slightly from `dj_database_url`. `config` in that it generates the entire `DATABASES` dictionary, including the `default` item. For example, consider the following code:

```
import dj_database_url

DATABASES = {'default': dj_database_url.config() }
```

This code is equivalent to the following:

```
from configurations import values

DATABASES = values.DatabaseURLValue()
```

Do not write `DATABASES['default'] = values.DatabaseURLValue()`, as your dictionary will be doubly nested.

If you need to specify multiple databases, you will need to fall back to `dj_database_url.config` directly, rather than using `DatabaseURLValue`.

Like other `values` classes, `DatabaseURLValue` takes a default value as its first argument. You might also want to use the `environment_prefix` argument and set it to `DJANGO` so that the environment variable being read is consistently named with the others. A full example of using `DatabaseURLValue` would, therefore, be like this:

```
DATABASES = values.DatabaseURLValue(
    f'sqlite:///{BASE_DIR}/db.sqlite3',
    environ_prefix='DJANGO')
```

By setting the `environment_prefix` like this, we can set the database URL using the `DJANGO_DATABASE_URL` environment variable (rather than just `DATABASE_URL`). This means it is consistent with other environment variable settings that also start with `DJANGO_`, such as `DJANGO_DEBUG` or `DJANGO_ALLOWED_HOSTS`.

Note that even though we are not importing `dj-database-url` in `settings.py`, `django-configurations` uses it internally, so it still must be installed.

In the next exercise, we will configure Bookr to use `DatabaseURLValue` to set its database configuration. It will be able to read from an environment variable and fall back to a default we specify.

Exercise 15.02 – dj-database-url and setup

In this exercise, we will install `dj-database-url` using `pip3`. Then, we will update Bookr's `settings.py` to configure the `DATABASE` setting using a URL, which is read from an environment variable. Let's get started with the steps:

1. In a terminal, make sure you have activated the `bookr` virtual environment, and then run the following command to install `dj-database-url` using `pip3`:

```
pip3 install dj-database-url
```

The installation process will run, and you should have output similar to this:

```
(bookr) ~ bookr pip install dj-database-url
Collecting dj-database-url
  Using cached dj_database_url-1.2.0-py3-none-any.whl (7.1 kB)
Requirement already satisfied: Django>=3.2 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from dj-database-url) (4.0)
Requirement already satisfied: backports.zoneinfo; python_version < "3.9" in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from Django>=3.2->dj-database-url) (0.2.1)
Requirement already satisfied: asgiref<4,>=3.4.1 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from Django>=3.2->dj-database-url) (3.4.1)
Requirement already satisfied: sqlparse==0.2.2 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from Django>=3.2->dj-database-url) (0.4.2)
Installing collected packages: dj-database-url
Successfully installed dj-database-url-1.2.0
```

Figure 15.5 – `dj-database-url` installation with `pip`

2. In PyCharm, open `settings.py` in the `bookr` package directory. Scroll down to find where the `DATABASES` attribute is being defined. Replace it with the `values.DatabaseURLValue` class. The first argument (default value) should be the URL to the SQLite database: `f'sqlite:/// {BASE_DIR}/db.sqlite3'`. Also, pass in `environ_prefix`, set to `DJANGO`. After completing this step, you should be setting the attribute like this:

```
DATABASES = values.DatabaseURLValue(
    f'sqlite:/// {BASE_DIR}/db.sqlite3',
    environ_prefix='DJANGO'
)
```

Save `settings.py`.

3. Start the Django dev server. As with *Exercise 15.01 – Django configurations setup*, if it starts fine, you can be confident that your change was successful. To be sure, open `http://127.0.0.1:8000/` in a browser and check that everything looks and behaves as it did before. You should visit a page that queries from the database (such as the `Books List` page) and check that a list of books is displayed:

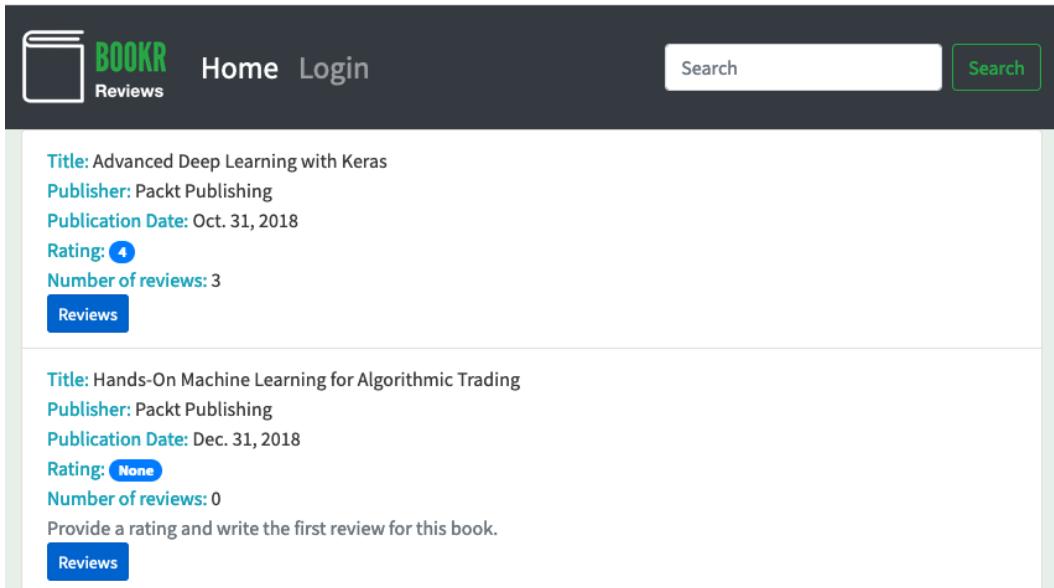


Figure 15.6 – Bookr pages with database queries still working

In this exercise, we updated our `settings.py` to determine its `DATABASES` setting from a URL specified in an environment variable. We used the `values.DatabaseURLValue` class to automatically read the value, and provided a default URL. We also set the `environ_prefix` argument to `DJANGO` so that the environment variable name is `DJANGO_DATABASE_URL`, which is consistent with other settings.

In the next section, we will take a tour of the Django Debug Toolbar, an app that helps you debug your Django applications through the browser.

The Django Debug Toolbar

The Django Debug Toolbar is an app that displays debug information about a web page right in your browser. It includes information about what SQL commands were run to generate the page, the request and response headers, how long the page took to render, and more. These can be useful in the following situations:

- *A page is taking a long time to load – maybe it is running too many database queries.* You can see whether the same queries are being run multiple times, in which case you could consider caching. Otherwise, some queries can be sped up by adding an index to the database.
- *You want to determine why a page is returning the wrong information.* Your browser may have sent headers you did not expect, or maybe some headers from Django are incorrect.

- *Your page is slow because it is spending time on non-database code* – you can profile the page to see what functions are taking the longest.
- *The page looks incorrect*. You can see what templates Django rendered. There might be a third-party template that is being rendered unexpectedly. You can also check all the settings that are being used (including the built-in Django ones that we are not setting). This can help to pinpoint a setting that is incorrect and causing the page to not behave correctly.

We'll explain how to use the Django Debug Toolbar to see this information. Before diving into how to set up the Django Debug Toolbar and use it, let's take a quick look at it. The toolbar is shown on the right of the browser window and can be toggled open and closed to display information:

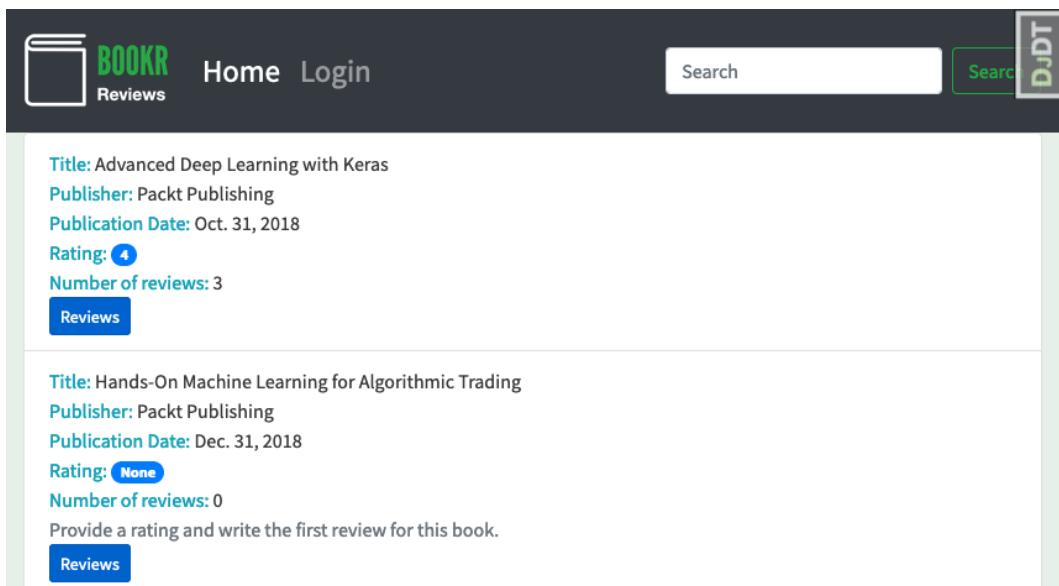


Figure 15.7 – The Django Debug Toolbar closed

The preceding figure shows the Django Debug Toolbar in its closed state. Note the toggle bar in the top-right corner of the window. Clicking the toolbar opens it:

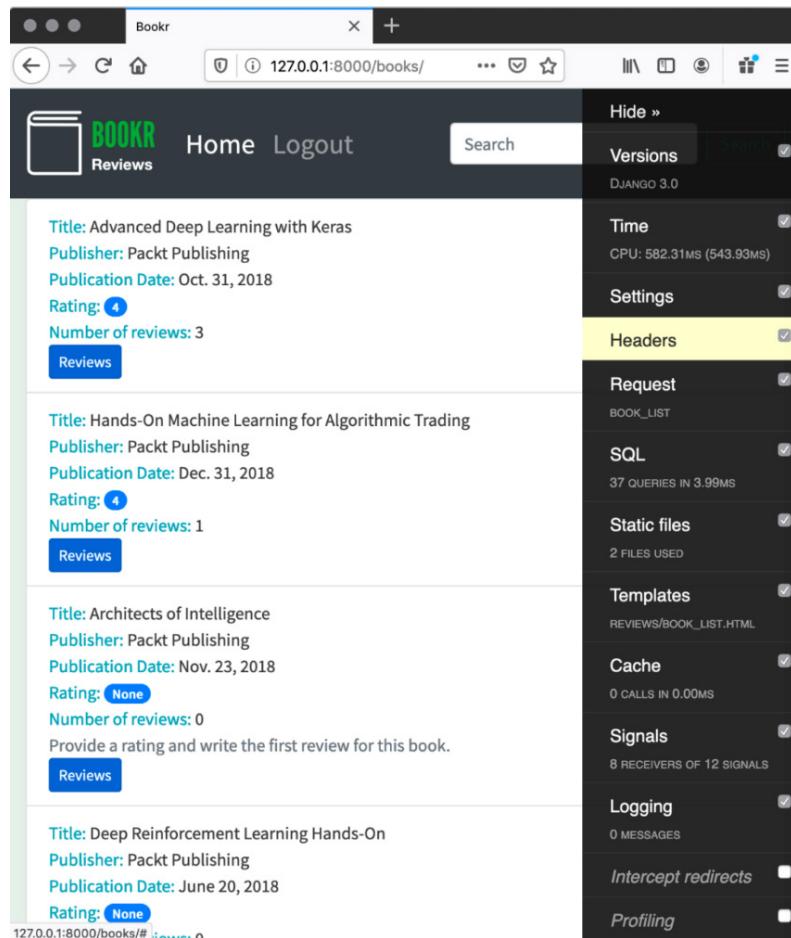


Figure 15.8 – The Django Debug Toolbar open

Figure 15.8 shows the Django Debug Toolbar open.

Installing the Django Debug Toolbar is done using pip:

```
pip3 install django-debug-toolbar
```

Note

For Windows, you can use pip instead of pip3 in the preceding command.

Then, there are a few steps to set it up, mostly by making changes to `settings.py`:

1. Add `debug_toolbar` to the `INSTALLED_APPS` settings list.
2. Add `debug_toolbar.middleware.DebugToolbarMiddleware` to the `MIDDLEWARE` settings list. It should be done as early as possible; for Bookr, it can be the first item on the list. This is the middleware that all requests and responses pass through.
3. Add '`127.0.0.1`' to the `INTERNAL_IPS` settings list (this setting may have to be created). The Django Debug Toolbar will only show IP addresses listed here.
4. Add the Django Debug Toolbar URLs to the base `urls.py` file. We want to add this mapping only if we are in `DEBUG` mode:

```
path('__debug__/', include(debug_toolbar.urls))
```

In the next exercise, we will go through these steps in detail.

Once the Django Debug Toolbar is installed and set up, any page you visit will show the **DjDT** sidebar (you can open or close it using the **DjDT** menu). When it's open, you'll be able to see another set of sections that you can click on to get more information.

Each panel has a checkbox next to it, this allows you to enable or disable the collection of that metric. Each metric that is collected will slightly slow down the page load (although, usually, this is not noticeable). If you find that one metric collection is slow, you can turn it off here. In the following points, we will go over each panel of this toolbar:

- The first is **History**, which shows the history of requests made in the browser. You can *switch* to previous requests in order to see the stats associated with them in the Django Debug Toolbar (*Figure 15.9*):

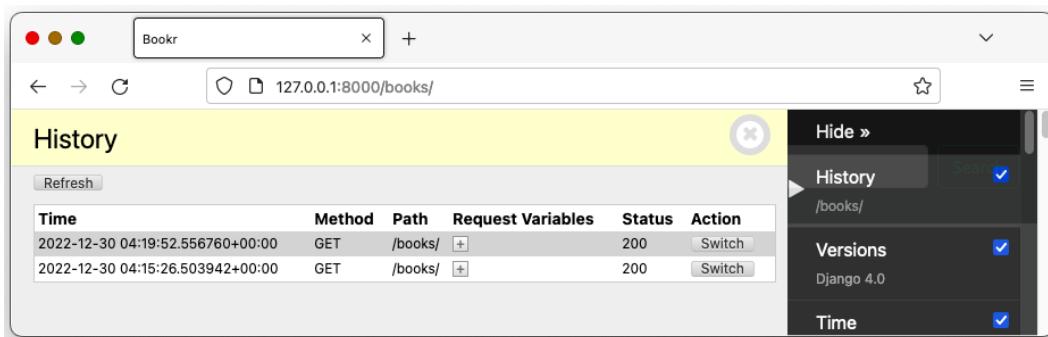
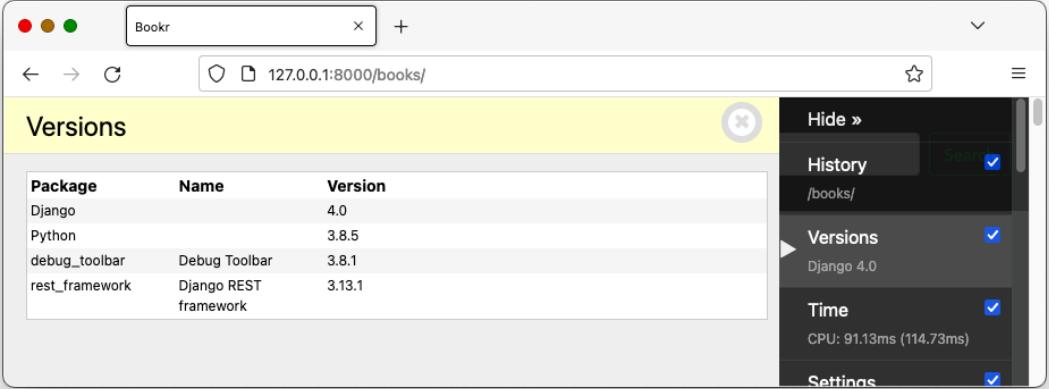


Figure 15.9 – The DjDT History panel (screenshot cropped for brevity)

- The second panel is **Versions**, which shows the version of Django running. You can click it to open a large **Versions** display, which will also show the version of Python and the Django Debug Toolbar (*Figure 15.10*):



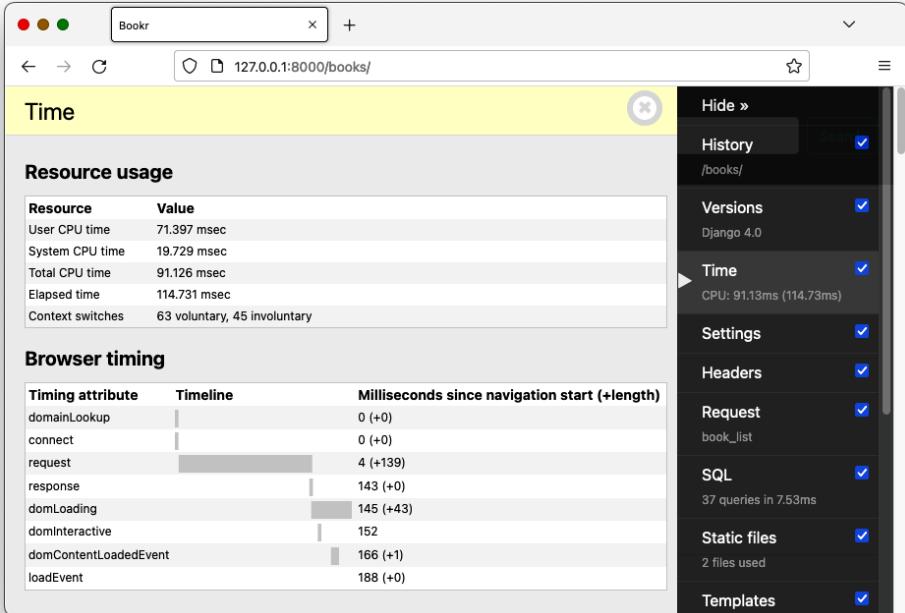
The screenshot shows a web browser window for 'Bookr' at '127.0.0.1:8000/books/'. The main content area is titled 'Versions' and displays a table of installed packages:

Package	Name	Version
Django		4.0
Python		3.8.5
debug_toolbar	Debug Toolbar	3.8.1
rest_framework	Django REST framework	3.13.1

To the right is a sidebar with checkboxes for 'History', 'Versions' (which is checked), 'Time' (which is checked), and 'Settings'.

Figure 15.10 – The DjDT Versions panel (screenshot cropped for brevity)

- The third panel is **Time**, which shows how long it took to process the request. It is broken down into system time and user time as well (*Figure 15.11*):



The screenshot shows a web browser window for 'Bookr' at '127.0.0.1:8000/books/'. The main content area is titled 'Time' and displays two sections: 'Resource usage' and 'Browser timing'.

Resource usage table:

Resource	Value
User CPU time	71.397 msec
System CPU time	19.729 msec
Total CPU time	91.126 msec
Elapsed time	114.731 msec
Context switches	63 voluntary, 45 involuntary

Browser timing table:

Timing attribute	Timeline	Milliseconds since navigation start (+length)
domainLookup		0 (+0)
connect		0 (+0)
request	██████████	4 (+139)
response	██████████	143 (+0)
domLoading	██████████	145 (+43)
domInteractive	██████████	152
domContentLoadedEvent	██████████	166 (+1)
loadEvent	██████████	188 (+0)

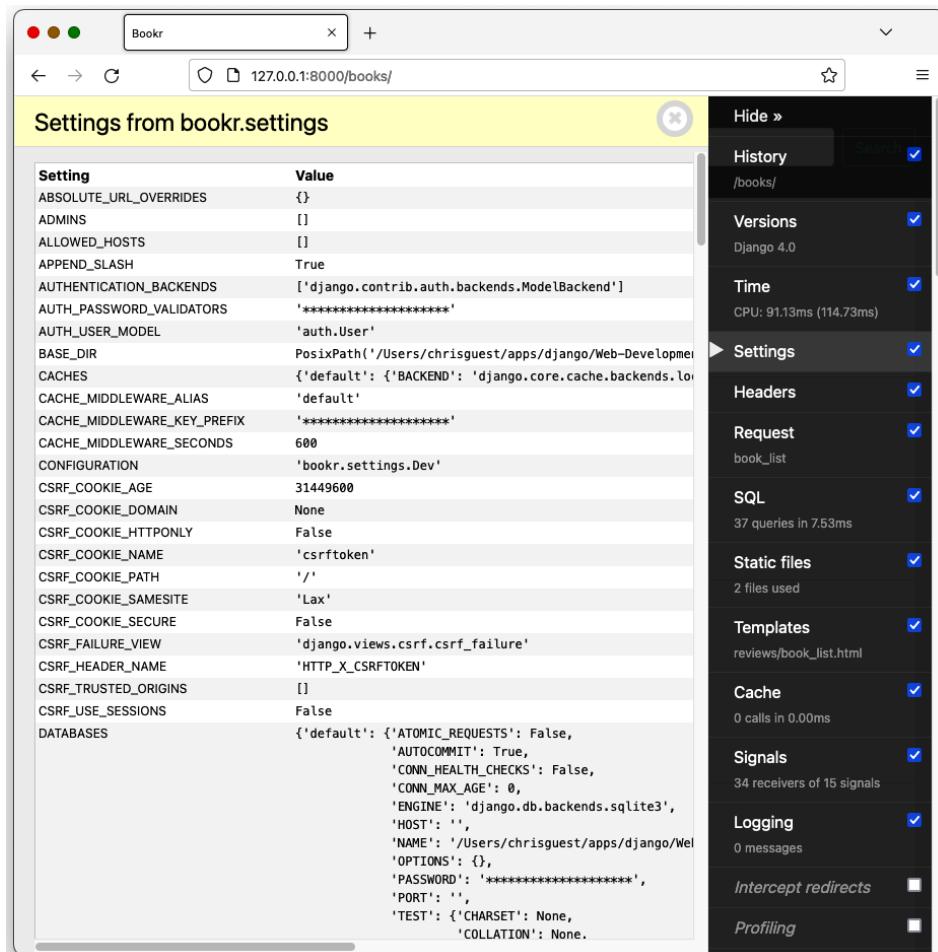
To the right is a sidebar with checkboxes for 'History', 'Versions' (which is checked), 'Time' (which is checked), 'Settings', 'Headers', 'Request', 'SQL', 'Static files', and 'Templates'.

Figure 15.11 – The DjDT Time panel

The differences between system time and user time are beyond the scope of this book; basically, system time is time spent in the kernel (for example, doing network or file reading/writing) and user time is code that is outside the operating system kernel (this includes the code you've written in Django, Python, and so on).

Also shown is the time spent in the browser, such as the time taken to get a request and how long it took to render a page.

- The fourth panel, **Settings**, shows all the settings your application is using (Figure 15.12):



The screenshot shows the Django Debug Toolbar (DjDT) Settings panel. The main area displays a table of settings and their values, with the first few rows listed below:

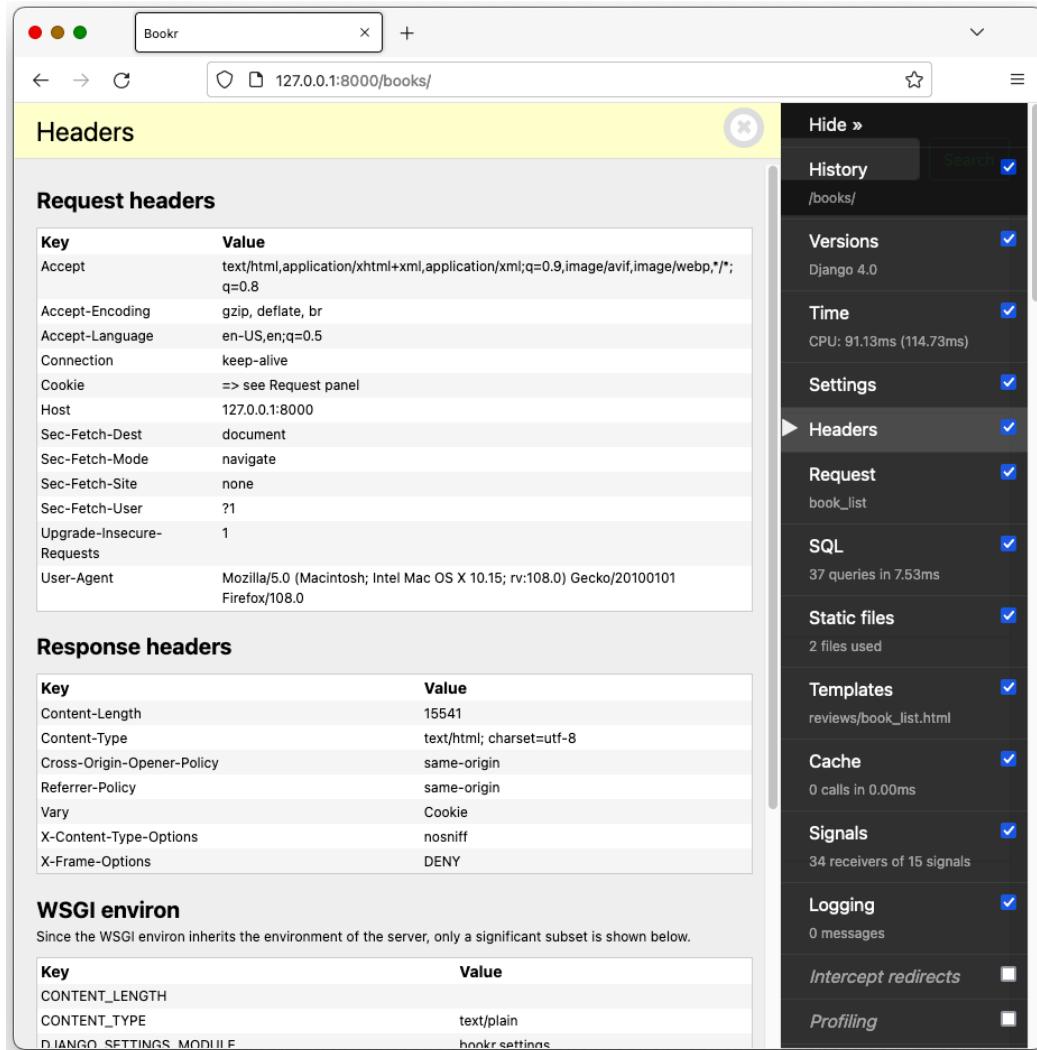
Setting	Value
ABSOLUTE_URL_OVERRIDES	{}
ADMINS	[]
ALLOWED_HOSTS	[]
APPEND_SLASH	True
AUTHENTICATION_BACKENDS	['django.contrib.auth.backends.ModelBackend']
AUTH_PASSWORD_VALIDATORS	'*****'
AUTH_USER_MODEL	'auth.User'
BASE_DIR	PosixPath('/Users/chrisguest/apps/django/Web-Development/Bookr/')
CACHES	{'default': {'BACKEND': 'django.core.cache.backends.locmem.LocMemCache', 'LOCATION': 'locmem://'}}
CACHE_MIDDLEWARE_ALIAS	'default'
CACHE_MIDDLEWARE_KEY_PREFIX	'*****'
CACHE_MIDDLEWARE_SECONDS	600
CONFIGURATION	'bookr.settings.Dev'
CSRF_COOKIE_AGE	31449600
CSRF_COOKIE_DOMAIN	None
CSRF_COOKIE_HTTPONLY	False
CSRF_COOKIE_NAME	'csrftoken'
CSRF_COOKIE_PATH	'/'
CSRF_COOKIE_SAMESITE	'Lax'
CSRF_COOKIE_SECURE	False
CSRF_FAILURE_VIEW	'django.views.csrf.csrf_failure'
CSRF_HEADER_NAME	'HTTP_X_CSRFTOKEN'
CSRF_TRUSTED_ORIGINS	[]
CSRF_USE_SESSIONS	False
DATABASES	{'default': {'ATOMIC_REQUESTS': False, 'AUTOCOMMIT': True, 'CONN_HEALTH_CHECKS': False, 'CONN_MAX_AGE': 0, 'ENGINE': 'django.db.backends.sqlite3', 'HOST': '', 'NAME': '/Users/chrisguest/apps/django/Web-Development/Bookr/db.sqlite3', 'OPTIONS': {}, 'PASSWORD': '*****', 'PORT': '', 'TEST': {'CHARSET': None, 'COLLATION': None}}}

The right sidebar contains a list of other panels with checkboxes, including History, Versions, Time, Settings, Headers, Request, SQL, Static files, Templates, Cache, Signals, Logging, Intercept redirects, and Profiling. The 'Settings' panel is currently selected.

Figure 15.12 – The DjDT Settings panel

This is useful because it shows both your settings from `settings.py` and the default Django settings.

- The fifth panel is **Headers** (Figure 15.13). It shows the headers of the request the browser made and the response headers that Django has sent:



The screenshot shows the DjDT Headers panel, which is the fifth panel in the Django Debug Toolbar. The panel is divided into two main sections: **Request headers** and **Response headers**.

Request headers:

Key	Value
Accept	text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*;q=0.8
Accept-Encoding	gzip, deflate, br
Accept-Language	en-US,en;q=0.5
Connection	keep-alive
Cookie	=> see Request panel
Host	127.0.0.1:8000
Sec-Fetch-Dest	document
Sec-Fetch-Mode	navigate
Sec-Fetch-Site	none
Sec-Fetch-User	?1
Upgrade-Insecure-Requests	1
User-Agent	Mozilla/5.0 (Macintosh; Intel Mac OS X 10.15; rv:108.0) Gecko/20100101 Firefox/108.0

Response headers:

Key	Value
Content-Length	15541
Content-Type	text/html; charset=utf-8
Cross-Origin-Opener-Policy	same-origin
Referrer-Policy	same-origin
Vary	Cookie
X-Content-Type-Options	nosniff
X-Frame-Options	DENY

WSGI environ:

Since the WSGI environ inherits the environment of the server, only a significant subset is shown below.

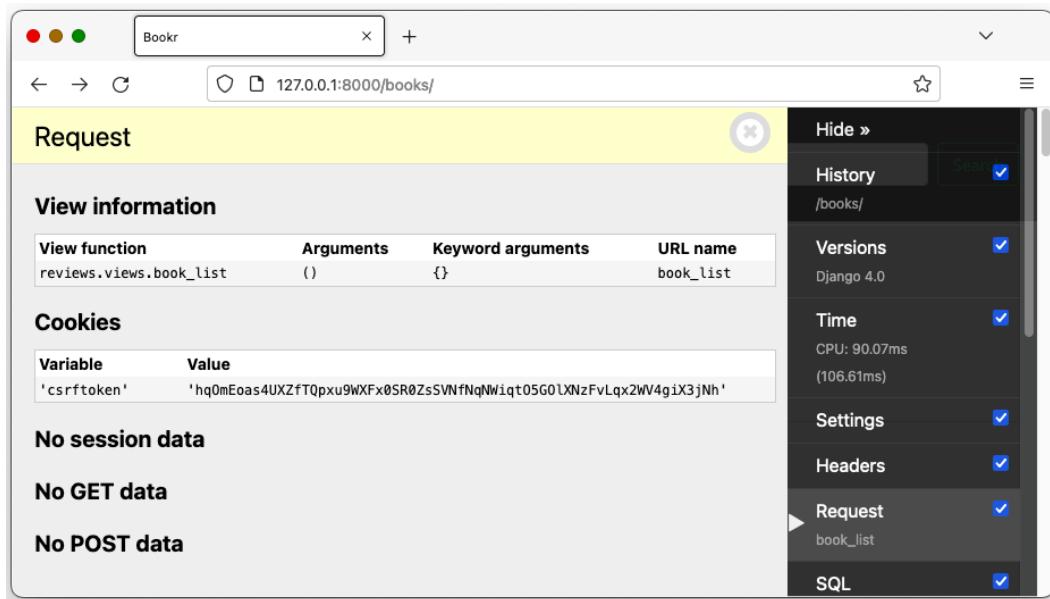
Key	Value
CONTENT_LENGTH	
CONTENT_TYPE	text/plain
DJANGO_SETTINGS_MODULE	bookr.settings

Toolbar Sidebar:

- History (checked)
- Versions (checked)
- Time (checked)
- Settings (checked)
- Headers (checked)
- Request (checked)
- SQL (checked)
- Static files (checked)
- Templates (checked)
- Cache (checked)
- Signals (checked)
- Logging (checked)
- Intercept redirects (unchecked)
- Profiling (unchecked)

Figure 15.13 – The DjDT Headers panel

- The sixth panel, **Request**, shows the view that generated the response, and the arguments and keyword arguments (args and kwargs parameters) that it was called with (Figure 15.14). You can also see the name of the URL used in URL map:



The screenshot shows the 'Request' panel of the Django Debug Toolbar. The panel is divided into several sections:

- View information:** Shows the view function as `reviews.views.book_list`, arguments as `()`, keyword arguments as `{}`, and URL name as `book_list`.
- Cookies:** Shows a table with one row: 'Variable' `'csrf_token'` and 'Value' `'hq0mEoas4UXZfTQpxu9WFX0SR0ZsSVNfNqNWiqt05G0lXNzFvLqx2WV4giX3jNh'`.
- No session data:** A section indicating there is no session data.
- No GET data:** A section indicating there is no GET data.
- No POST data:** A section indicating there is no POST data.

On the right side, there is a sidebar with checkboxes for various panels: History (checked), Versions (checked), Time (checked), Settings (checked), Headers (checked), Request (checked), and SQL (checked). The 'History' panel is currently active, showing the URL `/books/` and the Django version `Django 4.0`. The 'Time' panel shows CPU time as `90.07ms` and total time as `(106.61ms)`.

Figure 15.14 – The DjDT Request panel (some panels not shown for brevity)

It also shows the request's cookies, the information stored in the session (sessions were introduced in *Chapter 8, Media Serving and File Upload*), as well as the `request.GET` and `request.POST` data.

- The seventh panel, **SQL**, shows all the SQL database queries that were executed when building the response (Figure 15.15):

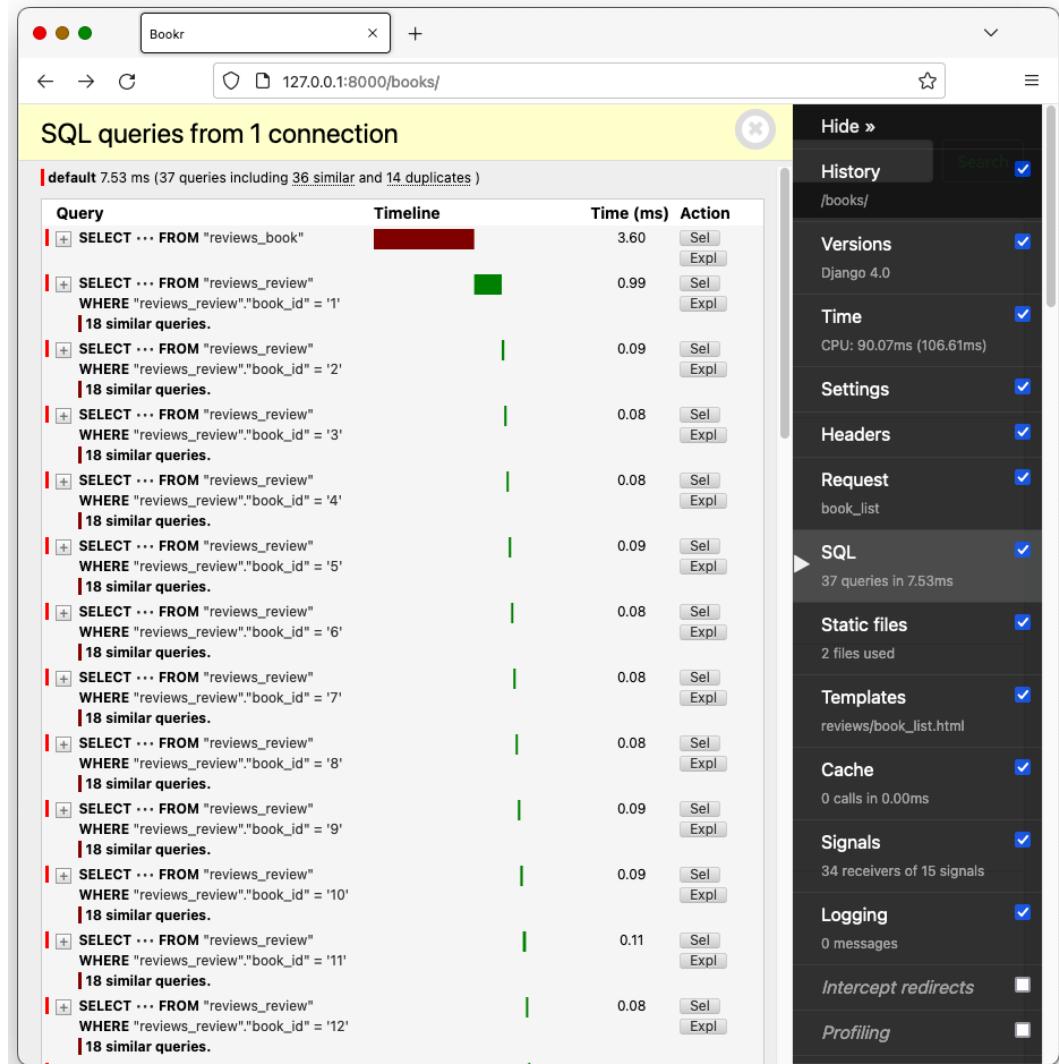
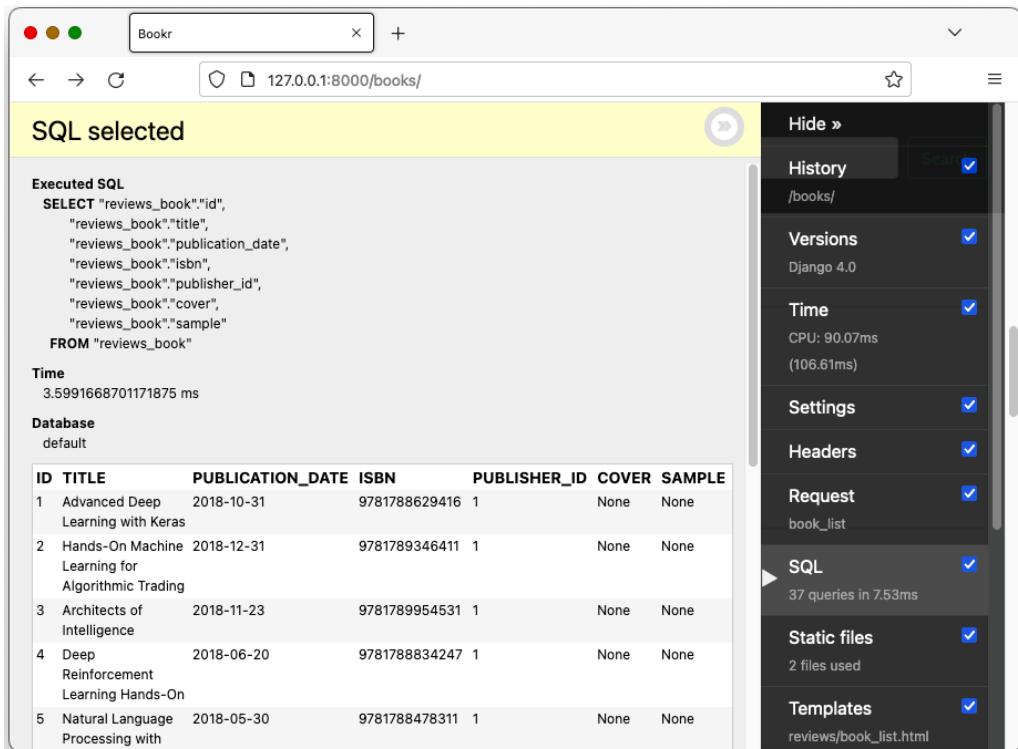


Figure 15.15 – The DjDT SQL panel

You can see how long each query took to execute and in what order they were executed. It also flags similar and duplicate queries so that you can potentially refactor your code to remove them.

Each `SELECT` query displays two action buttons, `Sel`, short for *select*, and `Expl`, short for *explain*. These do not show up for `INSERT`, `UPDATE`, or `DELETE` queries.

The `Sel` button shows the `SELECT` statement that was executed and all the data that was retrieved for the query (Figure 15.16):



SQL selected

Executed SQL

```
SELECT "reviews_book"."id",
    "reviews_book"."title",
    "reviews_book"."publication_date",
    "reviews_book"."isbn",
    "reviews_book"."publisher_id",
    "reviews_book"."cover",
    "reviews_book"."sample"
FROM "reviews_book"
```

Time

3.5991668701171875 ms

Database

default

ID	TITLE	PUBLICATION_DATE	ISBN	PUBLISHER_ID	COVER	SAMPLE
1	Advanced Deep Learning with Keras	2018-10-31	9781788629416	1	None	None
2	Hands-On Machine Learning for Algorithmic Trading	2018-12-31	9781789346411	1	None	None
3	Architects of Intelligence	2018-11-23	9781789954531	1	None	None
4	Deep Reinforcement Learning Hands-On	2018-06-20	9781788834247	1	None	None
5	Natural Language Processing with	2018-05-30	9781788478311	1	None	None

Hide >

History /books/

Versions Django 4.0

Time CPU: 90.07ms (106.61ms)

Settings

Headers

Request book_list

SQL 37 queries in 7.53ms

Static files 2 files used

Templates reviews/book_list.html

Figure 15.16 – The DjDT SQL selected panel

The `Expl` button shows the `EXPLAIN` query for the `SELECT` query (Figure 15.17):

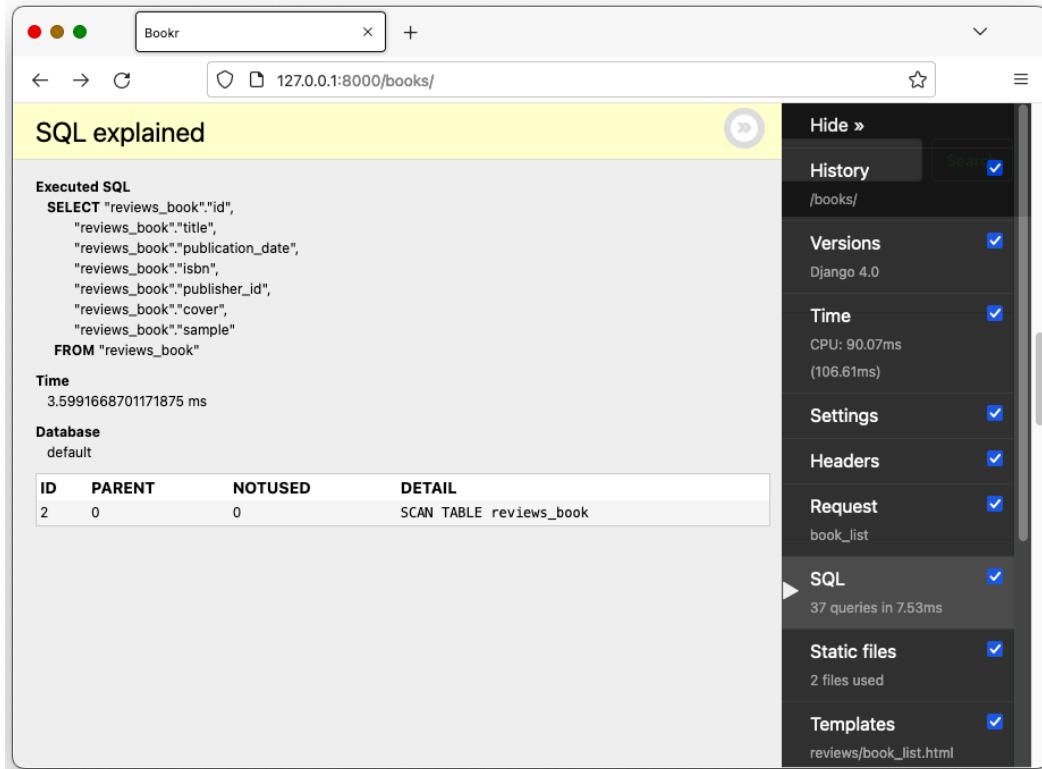


Figure 15.17 – The DjDT SQL explained panel (some panels not shown for brevity)

EXPLAIN queries are beyond the scope of the book, but they basically show how a database tried to execute the `SELECT` query – for example, what database indexes were used. You might find that a query does not use an index and you can, therefore, get faster performance by adding one.

- The eighth panel is **Static files**, and it shows you which static files were loaded in this request (Figure 15.18). It also shows you all the static files that are available and how they would be loaded (that is, which static file finder found them). The **Static files** panel's information is like the information you can get from the `findstatic` management command:

Static files (171 found, 2 used)

Static file path
/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/djdt/bookr/static

Static file apps
django.contrib.admin
rest_framework
reviews
debug_toolbar

Static files

main.css	Location
main.css	/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/djdt/bookr/static/main.css

reviews/logo.png	Location
reviews/logo.png	/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/djdt/bookr/reviews/static/reviews/logo.png

django.contrib.staticfiles.finders.FileSystemFinder (2 files)

Path	Location
main.css	/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/djdt/bookr/static/main.css
logo.png	/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/djdt/bookr/static/logo.png

django.contrib.staticfiles.finders.AppDirectoriesFinder (169 files)

Path	Location
admin/css/widgets.css	/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/bookr/lib/python3.8/site-packages/django/contrib/admin/static/admin/css/widgets.css
admin/css/login.css	/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/bookr/lib/python3.8/site-packages/django/contrib/admin/static/admin/css/login.css
admin/css/dashboard.css	/Users/chrisguest/apps/django/Web-Development-with-Django-Second-Edition/Chapter15/bookr/lib/python3.8/site-packages/django/contrib/admin/static/admin/css/dashboard.css

History
/books/

Versions
Django 4.0

Time
CPU: 90.07ms (106.61ms)

Settings

Headers

Request
book_list

SQL
37 queries in 7.53ms

Static files
2 files used

Templates
reviews/book_list.html

Cache
0 calls in 0.00ms

Signals
34 receivers of 15 signals

Logging
0 messages

Intercept redirects

Figure 15.18 – The DjDT Static files panel

- The ninth panel, **Templates**, shows information about the templates that were rendered (Figure 15.19):

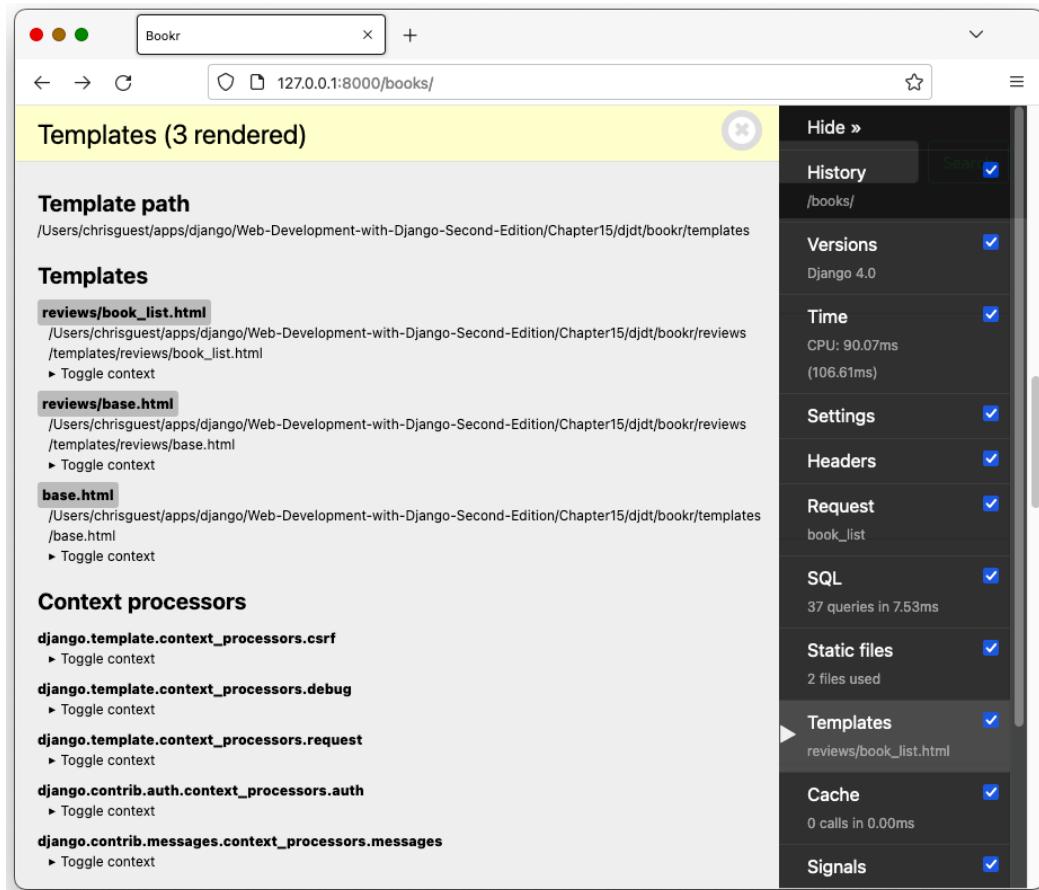


Figure 15.19 – The DjDT Templates panel

It shows the paths the templates were loaded from and the inheritance chain.

- The 10th panel, **Cache**, shows information about data fetched from Django's cache:

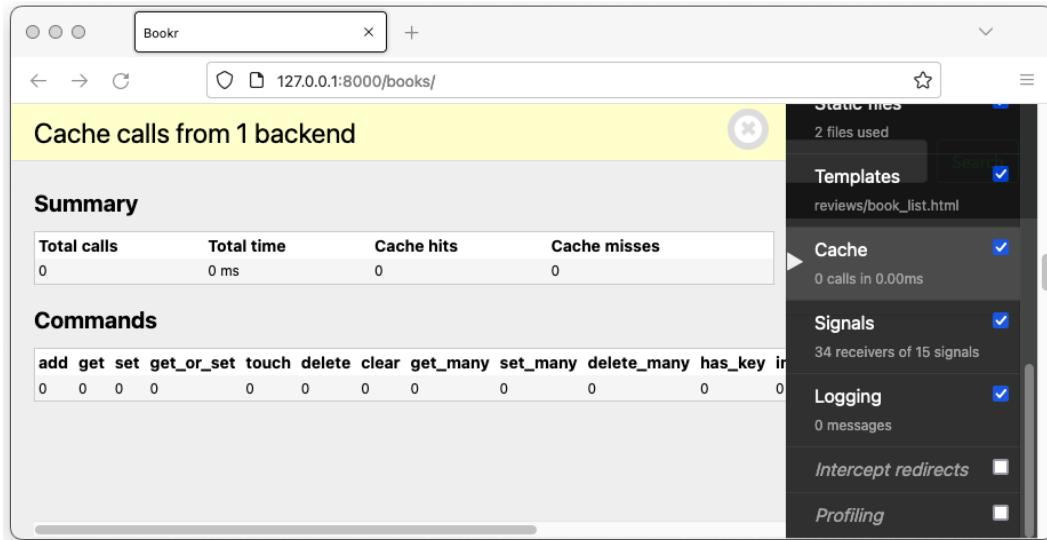
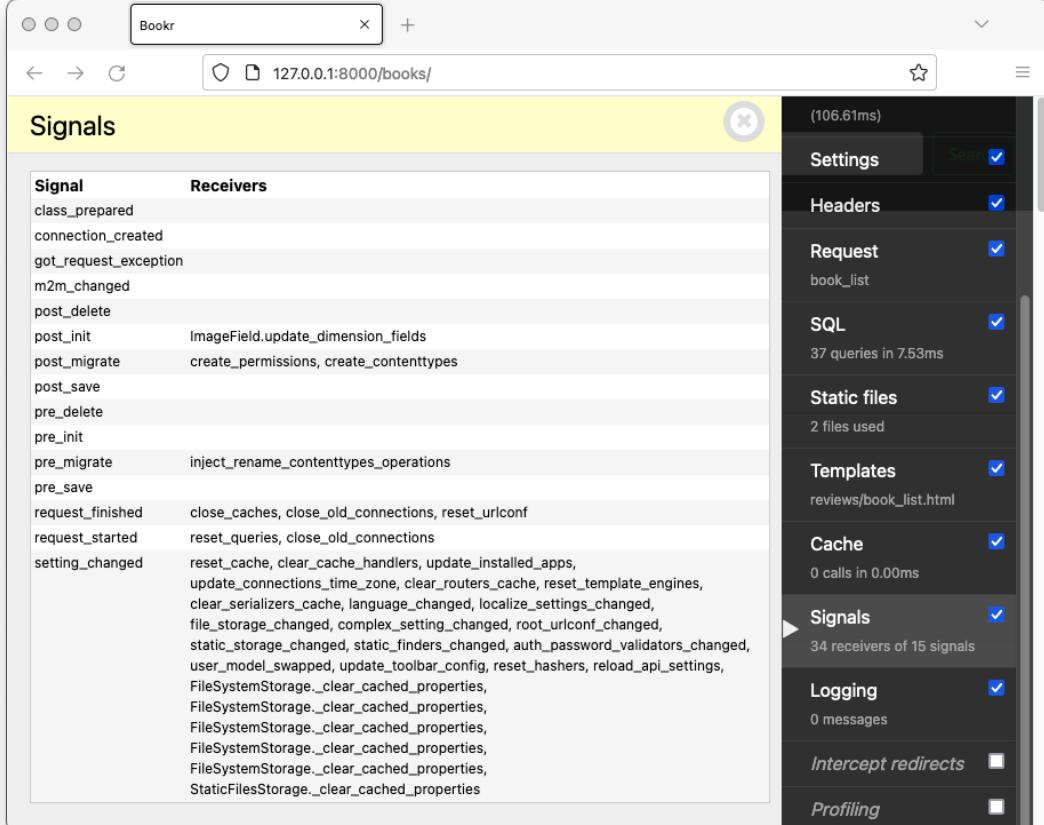


Figure 15.20 – The DjDT Cache panel (some panels not shown for brevity)

Since we aren't using caching in Bookr, this section is blank. If we were, we would be able to see how many requests to the cache had been made, and how many of those requests were successful in retrieving items. We would also see how many items had been added to the cache. This can give you an idea about whether you are using the cache effectively or not. If you are adding a lot of items to the cache but not retrieving any, then you should reconsider what data you are caching. Conversely, if you have a lot of cache misses (a miss is when you request data that is not in the cache), then you should be caching more data than you are already.

- The 11th panel is **Signals**, which shows information about Django signals (Figure 15.21):



The screenshot shows the DjDT (Django DevTools) interface for a Django application named 'Bookr'. The 'Signals' panel is active, displaying a table of signals and their corresponding receivers. The table has two columns: 'Signal' and 'Receivers'. The 'Signal' column lists various Django signal names, and the 'Receivers' column lists the functions or classes that handle these signals. For example, the 'post_init' signal has a receiver 'ImageField.update_dimension_fields', and the 'pre_save' signal has a receiver 'create_permissions, create_contenttypes'. The 'Signals' panel also includes a sidebar with various performance and configuration metrics, such as 'Settings' (106.61ms), 'Headers' (37 queries in 7.53ms), 'Request' (book_list), 'SQL' (2 files used), 'Static files' (Templates), 'Cache' (0 calls in 0.00ms), 'Signals' (34 receivers of 15 signals), 'Logging' (0 messages), and 'Intercept redirects' (Profiling).

Signal	Receivers
class_prepared	
connection_created	
got_request_exception	
m2m_changed	
post_delete	
post_init	ImageField.update_dimension_fields
post_migrate	create_permissions, create_contenttypes
post_save	
pre_delete	
pre_init	
pre_migrate	inject_rename_contenttypes_operations
pre_save	
request_finished	close_caches, close_old_connections, reset_urlconf
request_started	reset_queries, close_old_connections
setting_changed	reset_cache, clear_cache_handlers, update_installed_apps, update_connections_time_zone, clear_routers_cache, reset_template_engines, clear_serializers_cache, language_changed, localize_settings_changed, file_storage_changed, complex_setting_changed, root_urlconf_changed, static_storage_changed, static_finders_changed, auth_password_validators_changed, user_model_swapped, update_toolbar_config, reset_hashers, reload_api_settings, FileSystemStorage._clear_cached_properties, FileSystemStorage._clear_cached_properties, FileSystemStorage._clear_cached_properties, FileSystemStorage._clear_cached_properties, FileSystemStorage._clear_cached_properties, FileSystemStorage._clear_cached_properties

Figure 15.21 – The DjDT Signals panel (some panels not shown for brevity)

While we won't cover signals in this book, suffice to say, they are like events that you can hook into to execute functions when Django does something – for example, if a user is created, sending them a welcome email. This section shows which signals were sent and which functions received them.

- The 12th panel, **Logging**, shows log messages that were generated by your Django app (*Figure 15.22*):

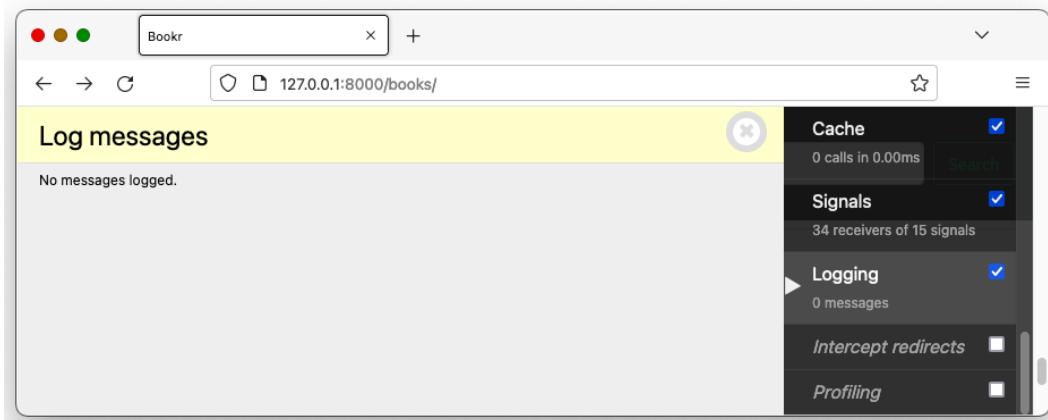


Figure 15.22 – The DjDT Logging panel

Since no log messages were generated in this request, this panel is empty.

The next option, **Intercept redirects**, is not a section with data. Instead, it lets you toggle redirect interception. If your view returns a redirect, it will not be followed. Instead, a page like *Figure 15.22* is displayed:

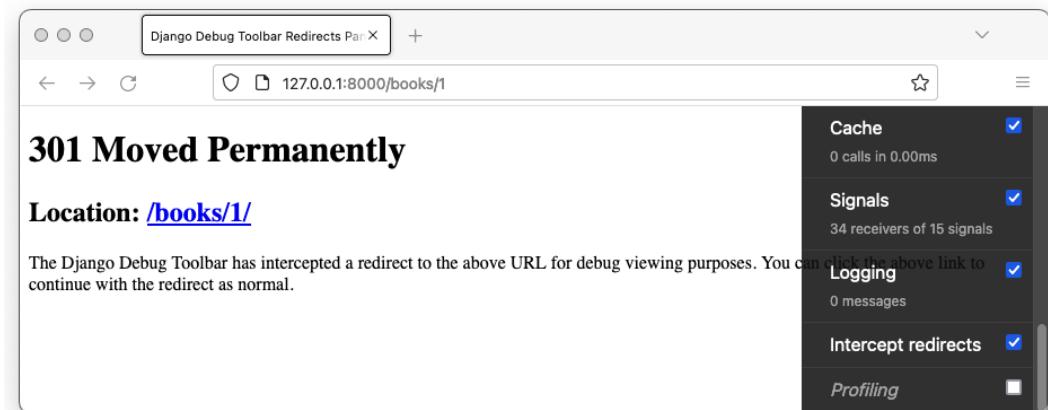


Figure 15.23 – A redirect that DjDT has intercepted

This allows you to open the Django Debug Toolbar for the view that generated the redirect; otherwise, you'd only be able to see the information for the view that you were redirected to.

- The final panel is **Profiling**. This is off by default, as profiling can slow down your response quite a lot. Once it is turned on, you must refresh the page to generate the profiling information (shown in *Figure 15.24*):

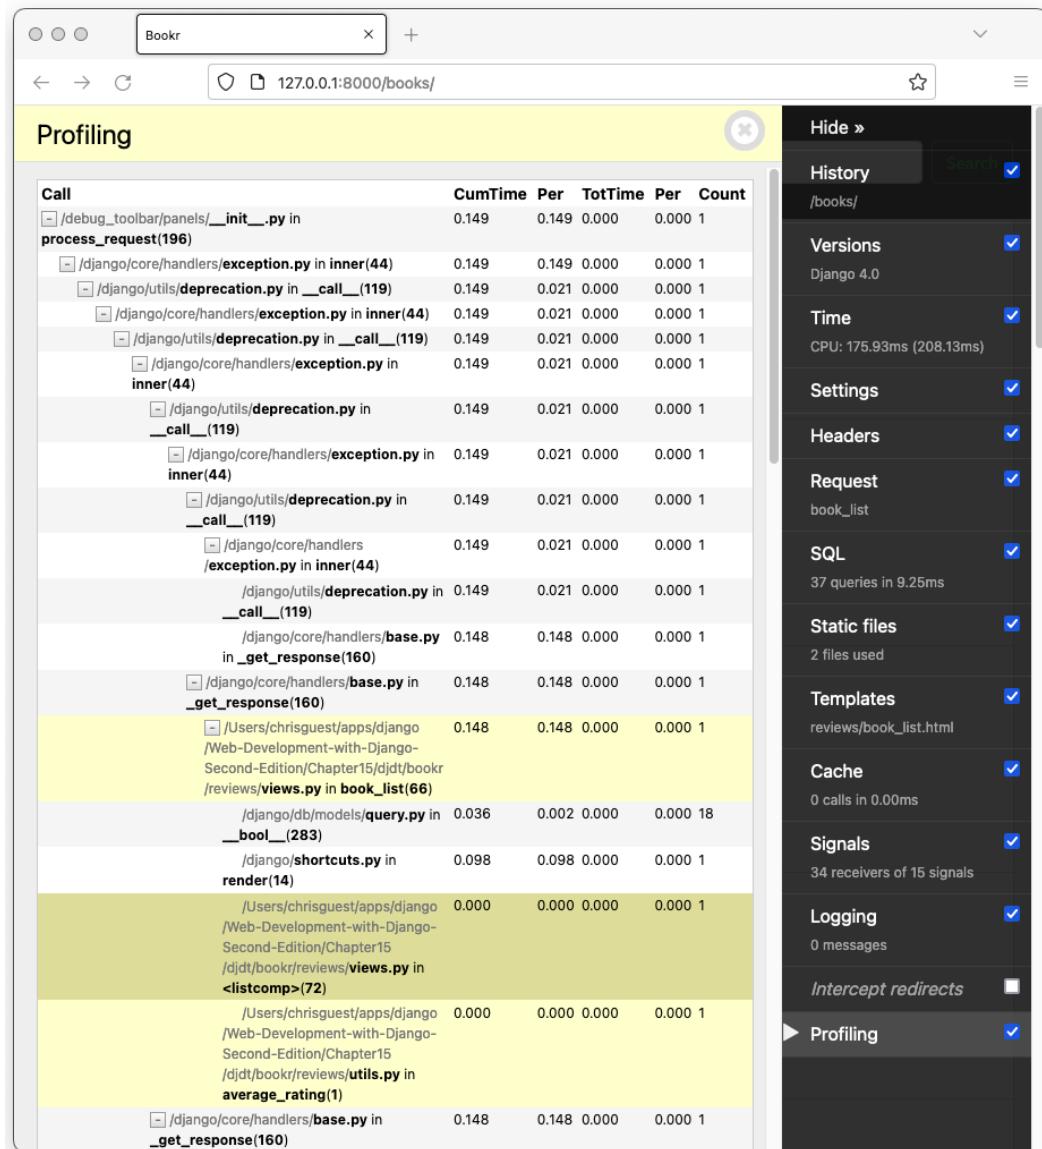


Figure 15.24 – The DjDT Profiling panel

The information shown here is a breakdown of how long each function call in your response took. The left of the page shows a stack trace of all the calls performed. On the right are columns with timing data. The columns are as follows:

- **CumTime:** The cumulative amount of time spent in the function and any sub-functions it calls
- **Per:** The cumulative time divided by the number of calls (Count)
- **TotTime:** The amount of time spent in this function but not in any sub-function it calls
- A second **Per:** The total time divided by the number of calls (Count)
- **Count:** The number of calls of this function

This information can help you determine how to speed up your app. For example, it can be easier to speed up a function that is called 1,000 times by a small fraction than to optimize a large function that is only called once. Any more in-depth tips on how to speed up your code are beyond the scope of this book.

We will reinforce what we have learned about the Django Debug Toolbar by installing and configuring it in the following exercise.

Exercise 15.03 – setting up the Django Debug Toolbar

In this exercise, we will add the Django Debug Toolbar settings by modifying the `INSTALLED_APPS`, `MIDDLEWARE`, and `INTERNAL_IPS` settings. Then, we'll add the `debug_toolbar.urls` map to the `bookr` package's `urls.py` file. Finally, we will load a page with the Django Debug Toolbar in a browser and use it:

1. In a terminal, make sure you have activated the `bookr` virtual environment, and then run the following command to install the Django Debug Toolbar using `pip3`:

```
pip3 install django-debug-toolbar
```

Note

For Windows, you can use `pip` instead of `pip3` in the preceding command.

The install process will run, and you should have output similar to *Figure 15.25*:

```
(bookr) ✘ bookr pip3 install django-debug-toolbar
Looking in indexes: https://pypi.org/simple, https://token:****@gitlab.com/api/v4/projects/28933584/packages/pypi/simple, https://token:****@gitlab.com/api/v4/projects/20903871/packages/pypi/simple
Collecting django-debug-toolbar
  Downloading django_debug_toolbar-3.8.1-py3-none-any.whl (221 kB)
    ██████████ | 221 kB 250 kB/s
Requirement already satisfied: sqlparse>=0.2 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django-debug-toolbar) (0.4.2)
Requirement already satisfied: django>=3.2.4 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django-debug-toolbar) (4.0)
Requirement already satisfied: asgiref<4,>=3.4.1 in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django>=3.2.4->django-debug-toolbar) (3.4.1)
Requirement already satisfied: backports.zoneinfo; python_version < "3.9" in /Users/ben/.virtualenvs/bookr/lib/python3.8/site-packages (from django>=3.2.4->django-debug-toolbar) (0.2.1)
Installing collected packages: django-debug-toolbar
Successfully installed django-debug-toolbar-3.8.1
```

Figure 15.25 – django-debug-toolbar installation with pip

2. Open `settings.py` in the `bookr` package directory. Add `debug_toolbar` to the `INSTALLED_APPS` setting:

```
INSTALLED_APPS = [
    ...
    'debug_toolbar',
]
```

This will allow Django to find the Django Debug Toolbar's static files.

3. Add `debug_toolbar.middleware.DebugToolbarMiddleware` to the `MIDDLEWARE` setting – it should be the first item in the list:

```
MIDDLEWARE = [
    'debug_toolbar.middleware.DebugToolbarMiddleware',
...]
```

This will route requests and responses through `DebugToolbarMiddleware`, allowing the Django Debug Toolbar to inspect the request and insert its HTML into the response.

4. The final setting to add is the `127.0.0.1` address to `INTERNAL_IPS`. You will not yet have an `INTERNAL_IPS` setting defined, so add this as a setting:

```
INTERNAL_IPS = ['127.0.0.1']
```

This will make the Django Debug Toolbar only show up on the developer's computer. You can now save `settings.py`.

5. We now need to add the Django Debug Toolbar URLs. Open `urls.py` in the `bookr` package directory. We already have an `if` condition that checks for `DEBUG` mode and then adds the media URL, like so:

```
if settings.DEBUG:
    urlpatterns += static(settings.MEDIA_URL,
        document_root=settings.MEDIA_ROOT)
```

We will also add `include` to `debug_toolbar.urls` inside this `if` statement; however, we will add it to the start of `urlpatterns`, rather than appending it to the end. Add this code inside the `if` statement:

```
import debug_toolbar

urlpatterns = [
    path('__debug__/',
         include(debug_toolbar.urls)),
] + urlpatterns
```

Save `urls.py`.

6. Start the Django dev server if it is not already running and navigate to `http://127.0.0.1:8000`. You should see the Django Debug Toolbar open. If it is not open, click the **DjDT** toggle button at the top-right to open it:



Figure 15.26 – The DjDT toggle shown in the corner

7. Try going through some of the panels and visiting different pages to see what information you can find out. Try also turning on **Intercept redirects** and then creating a new book review. After submitting the form, you should see the intercepted page, rather than being redirected to the new review (*Figure 15.27*):



Figure 15.27 – The redirect intercept page after submitting a new review

You can then click the **Location** link to go to the page that it was being redirected to.

8. You can also try turning on **Profiling** and seeing which functions are being called a lot and which are taking up most of the rendering time.
9. Once you are finished experimenting with the Django Debug Toolbar, turn off **Intercept redirects and Profiling**.

In this exercise, we installed and set up the Django Debug Toolbar by adding settings and URL maps. We then saw it in action and examined the useful information it can give us, including how to work with redirects and see profiling information.

In the next section, we will look at the `django-crispy-forms` app, which will let us reduce the amount of code needed to write forms.

django-crispy-forms

In Bookr, we are using the Bootstrap CSS framework. It provides styles that can be applied to forms using CSS classes. Since Django is independent of Bootstrap, when we use Django forms, it does not even know that we are using Bootstrap and so has no idea of what classes to apply to form widgets.

`django-crispy-forms` acts as an intermediary between Django forms and Bootstrap forms. It can take a Django form and render it with the correct Bootstrap elements and classes. It not only supports Bootstrap but also other frameworks, such as **Uni-Form** and **Foundation** (although Foundation support must be added using a separate package, `crispy-forms-foundation`).

Its installation and setup are quite simple. Once again, it is installed with `pip3`:

```
pip3 install django-crispy-forms
```

Note

For Windows, you can use `pip` instead of `pip3` in the preceding command.

Then, there are just a couple of settings changes. First, add `crispy_forms` to your `INSTALLED_APPS`. Then, you need to tell `django-crispy-forms` what framework you are using so that it loads the correct templates. This is done with the `CRISPY_TEMPLATE_PACK` setting. In our case, it should be set to `bootstrap4`:

```
CRISPY_TEMPLATE_PACK = 'bootstrap4'
```

`django-crispy-forms` has two main modes of operation, either as a filter or a template tag. The former is easier to drop into an existing template. The latter allows more configuration options and moves more of the HTML generation into the `Form` class. We'll look at both of these in order.

The crispy filter

The first method to render a form with `django-crispy-forms` is by using the `crispy` template. First, the filter must be loaded into the template. The library name is `crispy_forms_tags`:

```
{% load crispy_forms_tags %}
```

Then, instead of rendering a form with the `as_p` method (or another method), use the `crispy` filter. Consider the following line:

```
{{ form.as_p }}
```

Replace it with this:

```
{{ form|crispy }}
```

Here's a quick *before-and-after* showing the `Review Create` form. None of the rest of the HTML has been changed, apart from the form rendering. *Figure 15.28* shows the standard Django form:

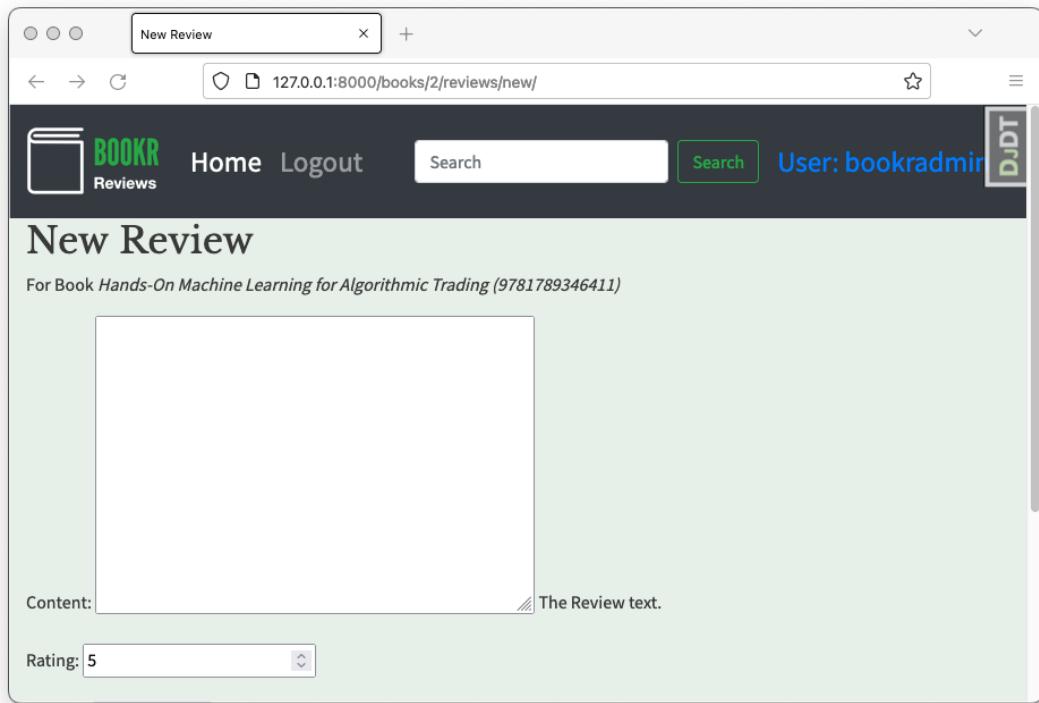


Figure 15.28 – The Review Create form with default styling

Figure 15.29 shows the form after `django-crispy-forms` has added the Bootstrap classes:

Figure 15.29 – The Review Create form with Bootstrap classes added by `django-crispy-forms`

When we integrate `django-crispy-forms` into Bookr, we will not use this method; however, it is worth knowing about because of how easy it is to drop it into your existing templates.

The crispy template Tag

The other method of rendering a form with `django-crispy-forms` is with the use of the `crispy` template tag. To use it, the `crispy_forms_tags` library must first be loaded into the template (as we did in the previous section). Then, the form is rendered like this:

```
{% crispy form %}
```

How does this differ from the `crispy` filter? The `crispy` template tag will also render the `<form>` element and the `{% csrf_token %}` template tag for you. So, consider, for example, that you used it like this:

```
<form method="post">
  {% csrf_token %}
  {% crispy form %}
</form>
```

The output for this would be as follows:

```
<form method="post" >
<input type="hidden" name="csrfmiddlewaretoken" value="...">
<form method="post">
<input type="hidden" name="csrfmiddlewaretoken" value="...">
    ... form fields ...
</form>
</form>
```

That is, the form and CSRF token fields are duplicated. In order to customize the `<form>` element that is generated, `django-crispy-forms` provides a `FormHelper` class, which can be set as a `Form` instance's `helper` attribute. It is the `FormHelper` instance that the `crispy` template tag uses to determine what attributes `<form>` should have.

Let us look at `ExampleForm` with a helper added. First, import the required modules:

```
from django import forms
from crispy_forms.helper import FormHelper
```

Next, define a form:

```
class ExampleForm(forms.Form):
    example_field = forms.CharField()
```

We could instantiate a `FormHelper` instance and then set it to the `form.helper` attribute (for example, in a view), but it's usually more useful to just create and assign it inside the form's `__init__` method. We haven't created a form with an `__init__` method yet, but it's no different from any other Python class:

```
def __init__(self, *args, **kwargs):
    super().__init__(*args, **kwargs)
```

Next, we set the `helper` and `form_method` for the helper (which is then rendered in the form's HTML):

```
self.helper = FormHelper()
self.helper.form_method = 'post'
```

Other attributes can be set on the helper, such as `form_action`, `form_id`, and `form_class`. We don't need to use these in Bookr though. We also do not need to manually set `enctype` on the form or its helper, as the `crispy` form tag will automatically set this to `multipart/form-data` if the form contains file upload fields.

If we tried to render the form now, we wouldn't be able to submit it, as there's no submit button (remember that we added submit buttons to our forms manually; they are not part of the Django form). `django-crispy-forms` also includes layout helpers that can be added to the form. They will be rendered after the other fields. We can add a submit button like this – first, import the `Submit` class:

```
from crispy_forms.layout import Submit
```

Note

`django-crispy-forms` does not properly support using a `<button>` input to submit a form, but for our purposes, `<input type="submit">` is functionally identical.

We then instantiate it and add it to the helper's inputs in a single line:

```
self.helper.add_input(Submit("submit", "Submit"))
```

The first argument to the `Submit` constructor is its *name*, and the second is its *label*.

`django-crispy-forms` is aware that we are using Bootstrap and will automatically render the button with the `btn btn-primary` classes.

The advantage of using a `crispy` template tag and `FormHelper` is that it means there is only one place where attributes and the behavior of the form are defined. We are already defining all the form fields in a `Form` class; this allows us to define the other attributes of the form in the same place. We can change a form from a GET submission to a POST submission easily here. The `FormHelper` instance will then automatically know that it needs to add a CSRF token to its HTML output when rendered.

We'll put all this into practice in the next exercise, where you will install `django-crispy-forms`, update `SearchForm` to utilize a form helper, and then render it using the `crispy` template tag.

Exercise 15.04 – using Django Crispy Forms with SearchForm

In this exercise, you will install `django-crispy-forms` and then convert `SearchForm` to be usable with the `crispy` template tag. This will be done by adding an `__init__` method and building a `FormHelper` instance inside it:

1. In a terminal, make sure you have activated the `bookr` virtual environment, and then run this command to install `django-crispy-forms` using `pip3`:

```
pip3 install django-crispy-forms
```

Note

For Windows, you can use `pip` instead of `pip3` in the preceding command.

The installation process will run, and you should have output similar to *Figure 15.30*:

```
(bookr) ➔ bookr pip install django-crispy-forms
Collecting django-crispy-forms
  Using cached django_crispy_forms-1.14.0-py3-none-any.whl (133 kB)
Installing collected packages: django-crispy-forms
Successfully installed django-crispy-forms-1.14.0
```

Figure 15.30 – django-crispy-forms installation with pip

2. Open `settings.py` in the `bookr` package directory, and then add `crispy_forms` to your `INSTALLED_APPS` setting:

```
INSTALLED_APPS = [...  
    'reviews',  
    'debug_toolbar',  
    'crispy_forms',  
]
```

This will allow Django to find the required templates.

3. While in `settings.py`, add a new setting for `CRISPY_TEMPLATE_PACK` – its value should be `bootstrap4`. This should be added as an attribute in the `Dev` class:

```
CRISPY_TEMPLATE_PACK = 'bootstrap4'
```

This lets `django-crispy-forms` know that it should be using the templates designed for Bootstrap version 4 when rendering forms. You can now save and close `settings.py`.

4. Open the `reviews` app's `forms.py` file. First, we need to add two imports to the top of the file – `FormHelper` from `crispy_forms.helper`, and `Submit` from `crispy_forms.layout`:

```
from crispy_forms.helper import FormHelper  
from crispy_forms.layout import Submit
```

5. Next, add an `__init__` method to `SearchForm`. It should accept `*args` and `**kwargs` as arguments. Then, call the `super().__init__` method with them:

```
class SearchForm(forms.Form):  
    ...  
    def __init__(self, *args, **kwargs):  
        super().__init__(*args, **kwargs)
```

This will simply pass through whatever arguments are provided to the superclass constructor.

6. Still inside the `__init__` method, set `self.helper` to an instance of `FormHelper`. Then, set the helper's `form_method` to `get`. Finally, create an instance of `Submit`, passing in an empty string as the name (first argument), and `Search` as the button label (second argument). Add this to the helper with the `add_input` method:

```
    self.helper = FormHelper()
    self.helper.form_method = "get"
    self.helper.add_input(Submit("", "Search"))
```

You can save and close `forms.py`.

7. In the `reviews` app's `templates` directory, open `search-results.html`. At the start of the file, after the `extends` template tag, use a `load` template tag to load `crispy_forms_tags`:

```
{% load crispy_forms_tags %}
```

8. Locate the existing `<form>` in the template. It should look like this:

```
<form>
    {{ form.as_p }}
    <button type="submit" class="btn btn-primary">
        Search</button>
</form>
```

You can delete the entered `<form>` element and replace it with a `crispy` template tag:

```
{% crispy form %}
```

This will use the `django-crispy-forms` library to render the form, including the `<form>` element and submit button. After making this change, this portion of the template should look like *Figure 15.31*:

```
{% block content %}
<h2>Search for Books</h2>
{% crispy form %}
{% if form.is_valid and search_text %}
    <h3>Search Results for <em>{{ search_text }}</em></h3>
```

Figure 15.31 – `search-results.html` after replacing `<form>` with the `crispy` form renderer

You can now save `search-results.html`.

9. Start the Django dev server if it is not already running and go to `http://127.0.0.1:8000/book-search/`. You should see the book search form as shown in *Figure 15.32*:

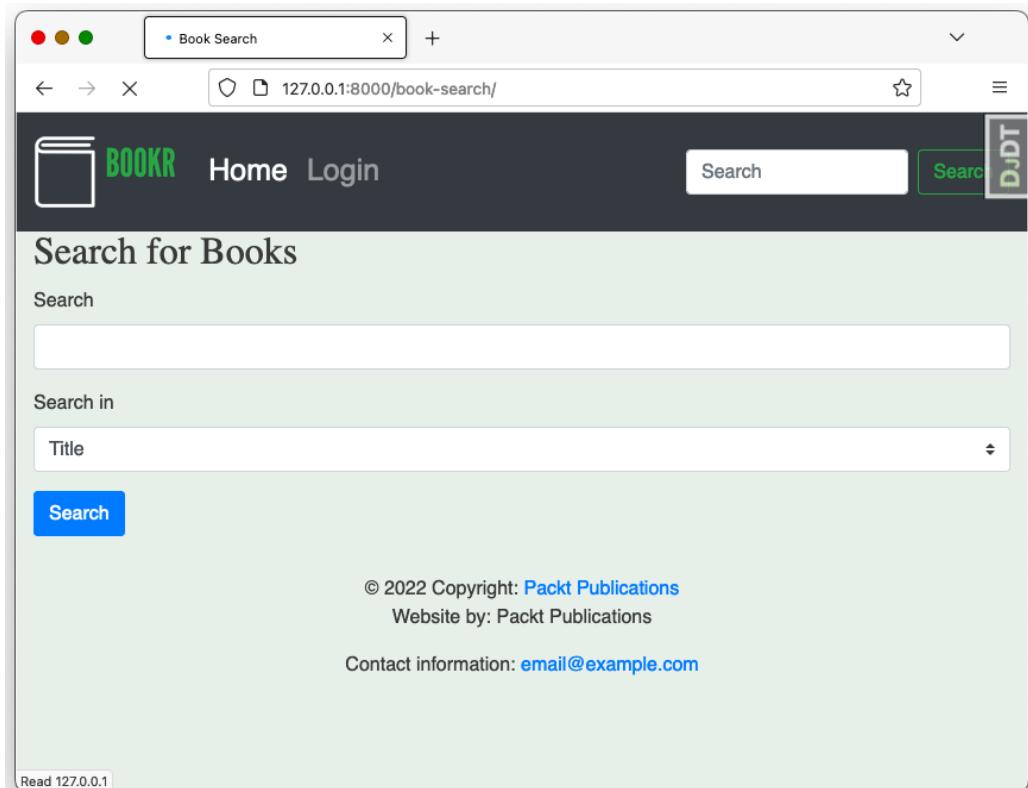


Figure 15.32 – The book search form rendered with django-crispy-forms

You should be able to use the form in the same manner as you did before (*Figure 15.33*):

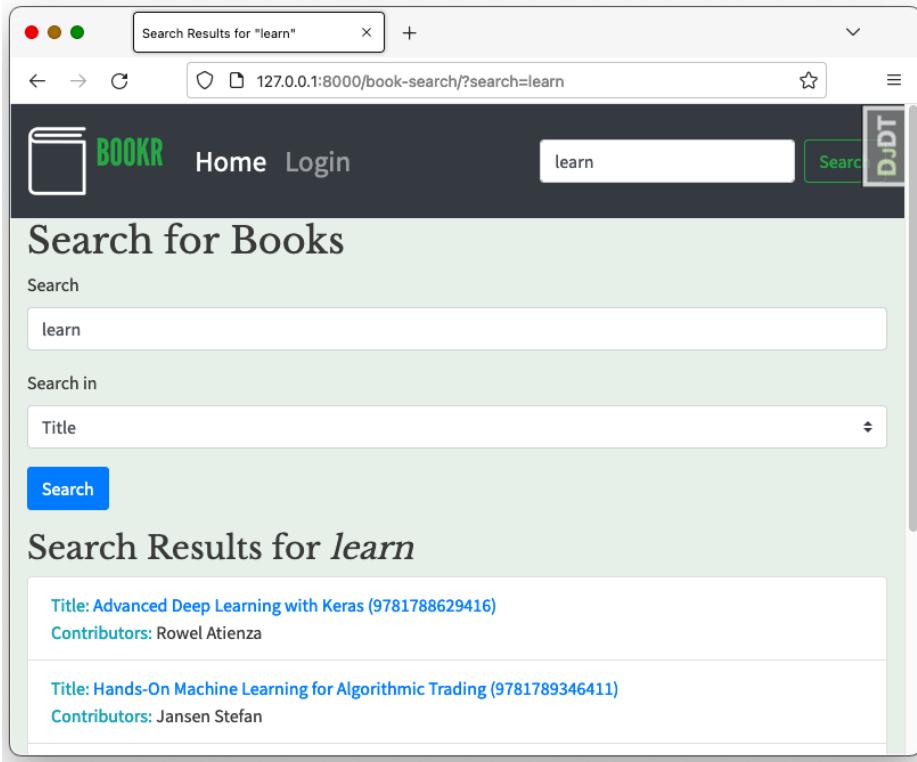


Figure 15.33 – Performing a search with the updated search form

Try viewing the source of the page in your web browser to see the rendered output. You will see that the `<form>` element has been rendered with the `method="get"` attribute, as we specified to `FormHelper` in *step 5*. Note also that `django-crispy-forms` has not inserted a CSRF token field – it knows that one is not required for a form submitted using GET.

In this exercise, we installed `django-crispy-forms` using `pip3` (`pip` for Windows) and then configured it in `settings.py` by adding it to `INSTALLED_APPS` and defining `CRISPY_TEMPLATE_PACK` we wanted to use (in our case, `bootstrap4`). We then updated the `SearchForm` class to use a `FormHelper` instance to control the attributes on the form and added a submit button using the `Submit` class. Finally, we changed the `search-results.html` template to use the `crispy` template tag to render the form, which allowed us to remove the `<form>` element we were using before and simplify form generation by moving all the form-related code into Python code (instead of it being partially in HTML and partially in Python).

Aside from giving a Django website a modern look, we also need a way to give our site standards of authentication that are convenient and trusted by users.

django-allauth

When browsing websites, you have probably seen buttons that allow you to log in using another website's credentials – for example, using your GitHub login:

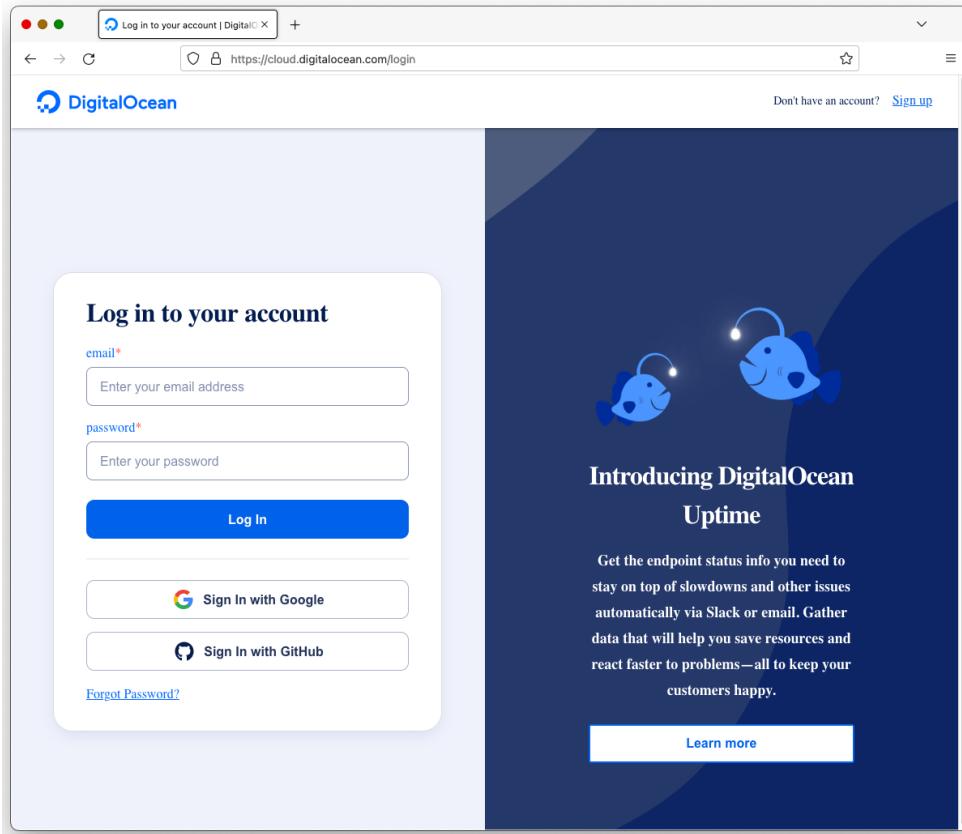


Figure 15.34 – The sign-in form with options to log in with Google or GitHub

Before we explain the process, let us introduce the terminology we will be using:

- **Requesting site:** The site the user is trying to log in to.
- **Authentication provider:** The third-party provider that the user is authenticating to (for example, Google or GitHub).

- **Authentication application:** This is something the creators of the requesting site set up on the authentication provider. It determines what permissions the requesting site will have with the authentication provider. For example, the requesting application can get access to your GitHub username but won't have permission to write to your repositories. The user can stop the requesting site from accessing your information from the authentication provider by disabling access to the authentication application.

The process is generally the same regardless of which third-party sign-in option you choose. First, you will be redirected to the authentication provider site and asked to authorize the authentication application to access your account (*Figure 15.35*):

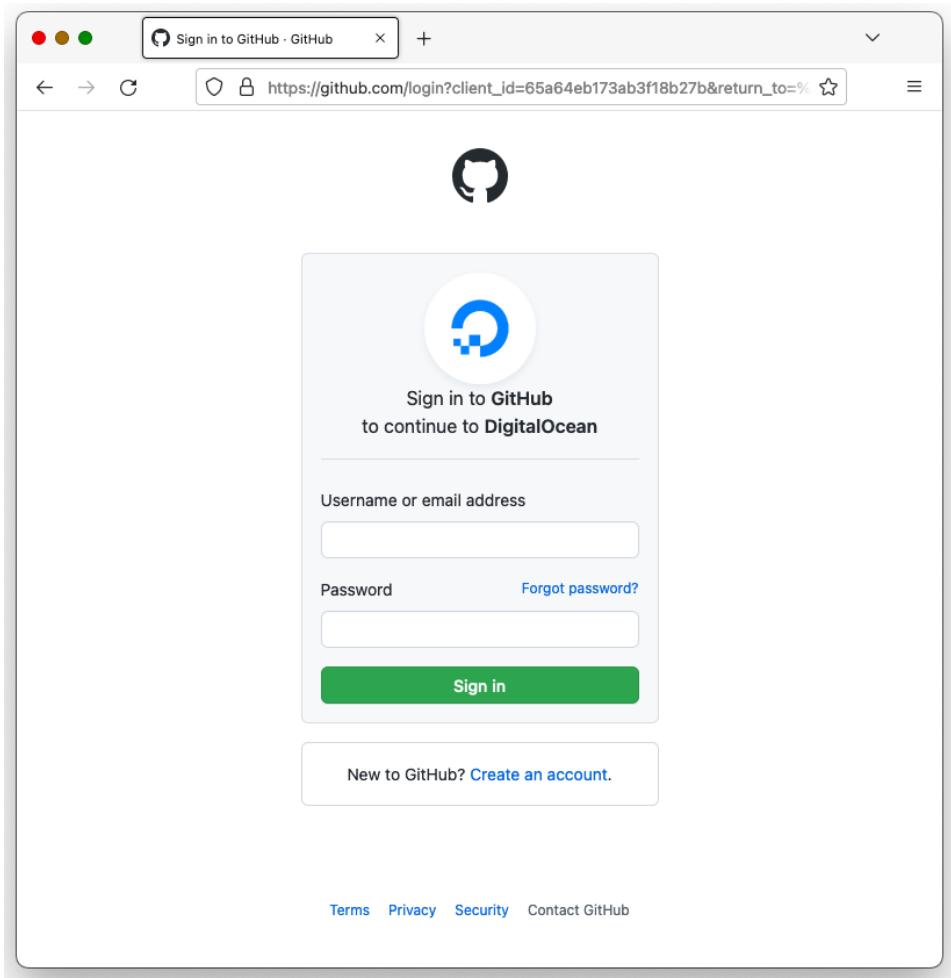


Figure 15.35 – The authentication provider authorization screen

After you authorize the authentication application, the authentication provider will redirect you back to the requesting site. The URL that you are redirected to will contain a secret token that the requesting site can use to request your user information in the backend. This allows the requesting site to verify who you are by communicating directly with the authentication provider. After validating your identity using a token, the requesting site can redirect you to your content. This flow is illustrated in a sequence diagram in *Figure 15.36*:

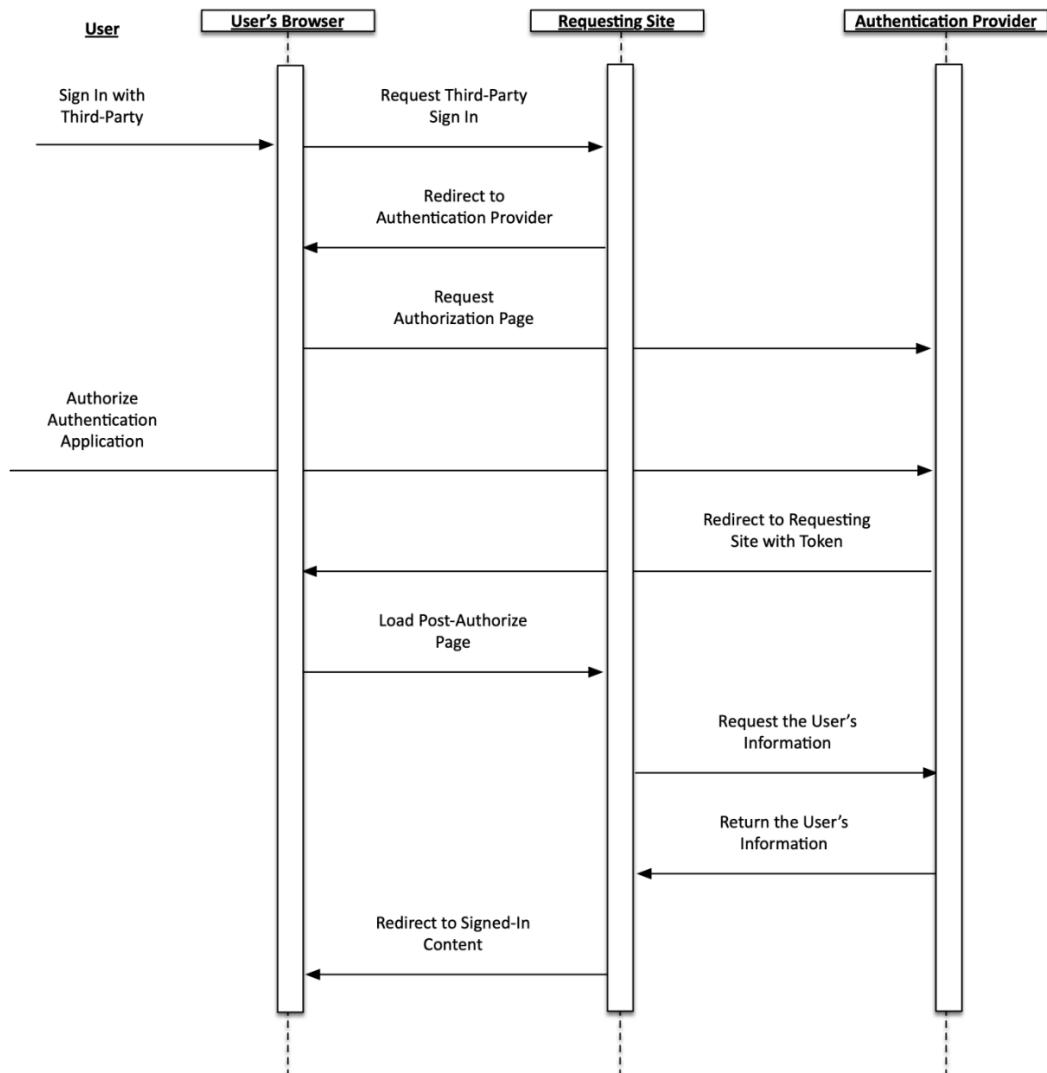


Figure 15.36 – Third-party authentication flow

Now that we have introduced authenticating using a third-party service, we can discuss `django-allauth`. This is an app that easily plugs your Django application into a third-party authentication service, including Google, GitHub, Facebook, and Twitter. In fact, at the time of writing, `django-allauth` supports over 75 authentication providers.

The first time a user authenticates to your site, `django-allauth` will create a standard Django `User` instance for you. It also knows how to parse the callback/redirect URL that the authentication provider loads after the end user authorizes the authentication application.

`django-allauth` adds three models to your application:

- `SocialApplication`: This stores the information used to identify your authentication application. The information you enter will depend on the provider, who will give you a *client ID*, *secret key*, and (optionally) a *key*. Note that these are the names that `django-allauth` uses for these values and they will differ based on the provider. We will give some examples of these values later in this section. `SocialApplication` is the only one of the `django-allauth` models that you will create yourself; `django-allauth` creates the others automatically when a user authenticates.
- `SocialApplicationToken`: This contains the values needed to identify a Django user to the authentication provider. It contains a *token* and (optionally) a *token secret*. It also contains a reference to the `SocialApplication` model that created it and the `SocialAccount` model to which it applies.
- `SocialAccount`: This links a Django user to the provider (for example, Google or GitHub) and stores extra information that the provider may have given.

Since there are so many authentication providers, we will not cover how to set them all up, but we will give a short instruction on setup and how to map the auth tokens from the providers to the right fields in `SocialApplication`. We will do this for the two auth providers we have been mentioning throughout the chapter – Google and GitHub.

django-allauth installation and setup

Like the other apps in this chapter, `django-allauth` is installed with `pip3`:

```
pip3 install django-allauth
```

Note

For Windows, you can use `pip` instead of `pip3` in the preceding command.

We then need a few settings changes. `django-allauth` requires the `django.contrib.sites` app to run, so it needs to be added to `INSTALLED_APPS`. Then, a new setting needs to be added to define a `SITE_ID` value for our site. We can just set this to 1 in our `settings.py` file:

```
INSTALLED_APPS = [
    "bookr_admin.apps.BookrAdminConfig",
    "django.contrib.sites", # this entry added
    "django.contrib.auth",
    # The rest of the values are truncated
]

SITE_ID = 1
```

Note

It is possible to have a single Django project hosted on multiple hostnames and have it behave differently on each, but also have content shared across all the sites. We don't need to use `SITE_ID` anywhere else in our project, but one instance must be set here. You can read more about the `SITE_ID` settings at <https://docs.djangoproject.com/en/3.0/ref/contrib/sites/>.

We also need to add `allauth` and `allauth.socialaccount` to `INSTALLED_APPS`:

```
INSTALLED_APPS = [
    # The rest of the values are truncated
    "allauth",
    "allauth.socialaccount",
]
```

Then, each provider we want to support must also be added to the list of `INSTALLED_APPS` – for example, consider the following snippet:

```
INSTALLED_APPS = [
    # The rest of the values are truncated
    "allauth.socialaccount.providers.github",
    "allauth.socialaccount.providers.google",
]
```

After all of this is done, we need to burn the `migrate` management command to create the `djangallauth` models:

```
python3 manage.py migrate
```

Once this is done, new social applications can be added through the Django admin interface (*Figure 15.37*):

Add social application

Provider: GitHub

Name:

Client id:
App ID, or consumer key

Secret key:
API secret, client secret, or consumer secret

Key:
Key

Sites:

Available sites ?

Filter

example.com

Chosen sites ?

+ Remove all

Choose all ?
Hold down "Control", or "Command" on a Mac, to select more than one.

Figure 15.37 – Adding a social application

To add a social application, select a provider (this list will only show those in the `INSTALLED_APPS` list), enter a name (it can just be the same as the provider), and enter the client ID from the provider's website (we will go into detail on this soon). You may also need a secret key and a key. Select the site the secret key and key should apply to. (If you only have one `Site` instance, then its name does not matter – just select it. The site name can be updated in the `Sites` section of Django admin. You can also add more sites there.)

We will now look at the tokens used by our three example providers.

GitHub auth setup

A new GitHub application can be set up under your GitHub profile. During development, your callback URL for the application should be set to `http://127.0.0.1:8000/accounts/github/login/callback/` and updated with the real hostname when you deploy to production. After creating the app, it will provide a client ID and client secret. These are your client ID and secret key respectively in `django-allauth`.

Google auth setup

The creation of a Google application is done through your Google Developers console. The authorized redirect URI should be set to `http://127.0.0.1:8000/accounts/google/login/callback/` during development and updated after production deployment. The app's Client ID is also `Client id` in `django-allauth`, and the app's Client secret is `Secret key`.

Initiating authentication with `django-allauth`

To initiate authentication through a third-party provider, you first need to add the `django-allauth` URLs in your URL maps. Somewhere inside your `urlpatterns` is one of your `urls.py` files, including `allauth.urls`:

```
urlpatterns = [path('allauth', include('allauth.urls')),]
```

You will then be able to initiate a login using URLs such as `http://127.0.0.1:8000/allauth/github/login/?process=login`, `http://127.0.0.1:8000/allauth/google/login/?process=login`, and so on. `django-allauth` will handle all the redirects for you and then create/authenticate the Django user when they return to the site. You can have buttons on your login page with text such as `Login with GitHub` or `Login with Google` that links to these URLs.

Other `django-allauth` features

Other than authentication with third-party providers, `django-allauth` can also add some useful features that Django does not have built in. For example, you can configure it to require an email address for a user, and have the user verify their email address by clicking a confirmation link they receive before they log in. `django-allauth` can also handle generating a URL for a password reset that is emailed to the user. You can find the documentation for `django-allauth` that explains these features, and more, at <https://django-allauth.readthedocs.io/en/stable/overview.html>.

Now that we have covered the first four third-party apps in depth and given a brief overview of `django-allauth`, you can undertake the activity for this chapter. In this activity, we will refactor the `ModelForm` instances we are using to use the `CrispyFormHelper` class.

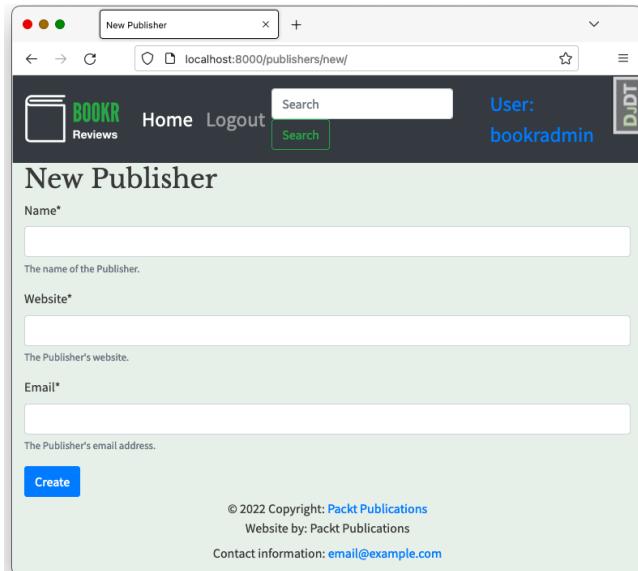
Activity 15.01 – using FormHelper to update forms

In this activity, we will update the `ModelForm` instances (`PublisherForm`, `ReviewForm`, and `BookMediaForm`) to use the `CrispyFormHelper` class. Using `FormHelper`, we can define the text of the `Submit` button inside the `Form` class itself. We can then move the `<form>` rendering logic out of the `instance-form.html` template and replace it with a `crispy` template tag.

These steps will help you complete the activity:

1. Create an `InstanceForm` class that subclasses `forms.ModelForm`. This will be the base of the existing `ModelForm` classes.
2. In the `__init__` method of `InstanceForm`, set a `FormHelper` instance on `self`.
3. Add a `Submit` button to `FormHelper`. If the form is instantiated with `instance`, then the button text should be `Save`; otherwise, it should be `Create`.
4. Update `PublisherForm`, `ReviewForm`, and `BookMediaForm` to extend from `InstanceForm`.
5. Update the `instance-form.html` template so that `<form>` is rendered using the `crispy` template tag. The rest of `<form>` can be removed.
6. In the `book_media` view, the `is_file_upload` context item is no longer required.

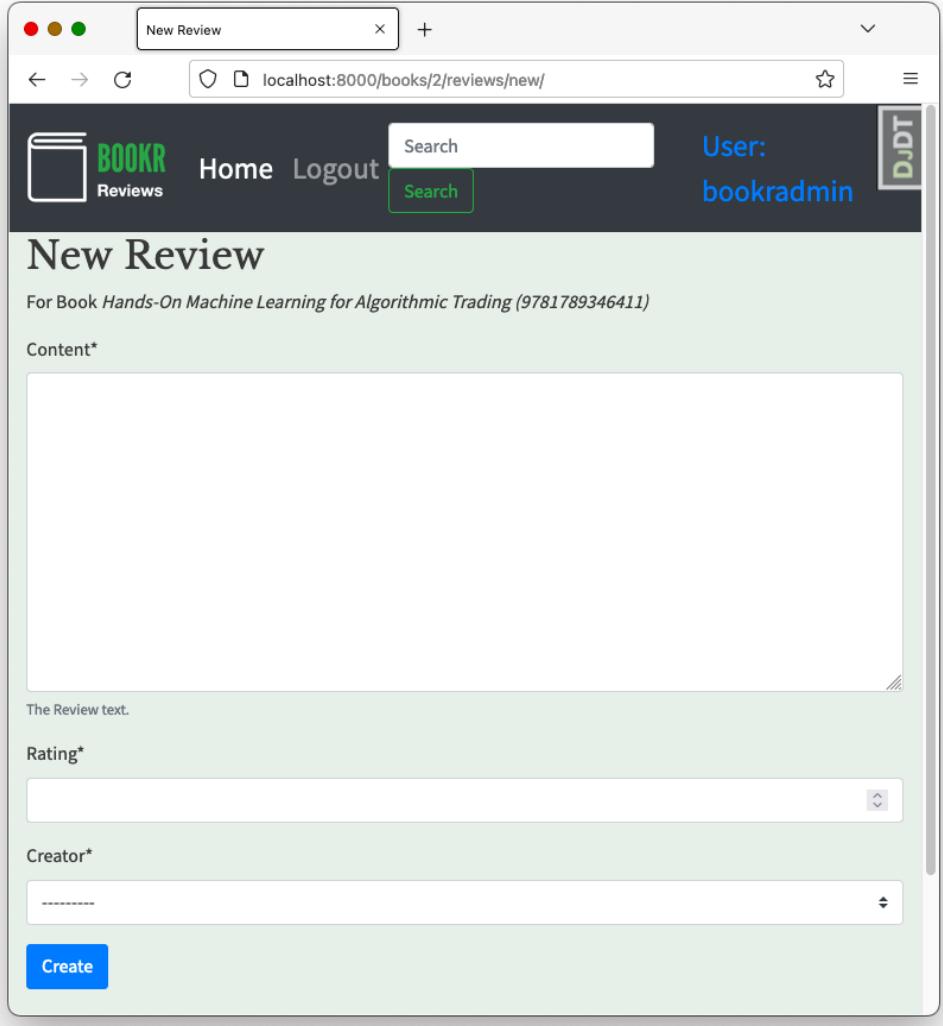
When you are finished, you should see the forms rendered with Bootstrap themes. *Figure 15.38* shows the **New Publisher** page:



The screenshot shows a web browser window with the title 'New Publisher'. The URL in the address bar is 'localhost:8000/publishers/new/'. The page has a dark header with a 'BOOKR Reviews' logo, a 'Home' link, a 'Logout' link, a search bar, and a 'User: bookadmin' label. The main content area is titled 'New Publisher'. It contains three form fields: 'Name*' with a placeholder 'The name of the Publisher.', 'Website*' with a placeholder 'The Publisher's website.', and 'Email*' with a placeholder 'The Publisher's email address.'. Below these fields is a blue 'Create' button. At the bottom of the page, there is a copyright notice: '© 2022 Copyright: Packt Publications' and 'Website by: Packt Publications'. The contact information is 'Contact information: email@example.com'.

Figure 15.38 – The New Publisher page

Figure 15.39 shows the **New Review** page:



The screenshot shows a web browser window with the title 'New Review' and the URL 'localhost:8000/books/2/reviews/new/'. The page is for creating a new review for the book 'Hands-On Machine Learning for Algorithmic Trading' (9781789346411). The interface includes a sidebar with a 'BOOKR Reviews' logo, a 'Search' bar, and a user profile for 'User: bookradmin'. The main content area is titled 'New Review' and contains fields for 'Content*', 'Rating*', and 'Creator*'. A 'Create' button is at the bottom.

New Review

localhost:8000/books/2/reviews/new/

BOOKR Reviews

Home Logout

Search

User: bookradmin

New Review

For Book *Hands-On Machine Learning for Algorithmic Trading* (9781789346411)

Content*

The Review text.

Rating*

Creator*

Create

Figure 15.39 – The New Review form

Finally, the book media page is displayed in *Figure 15.40*:

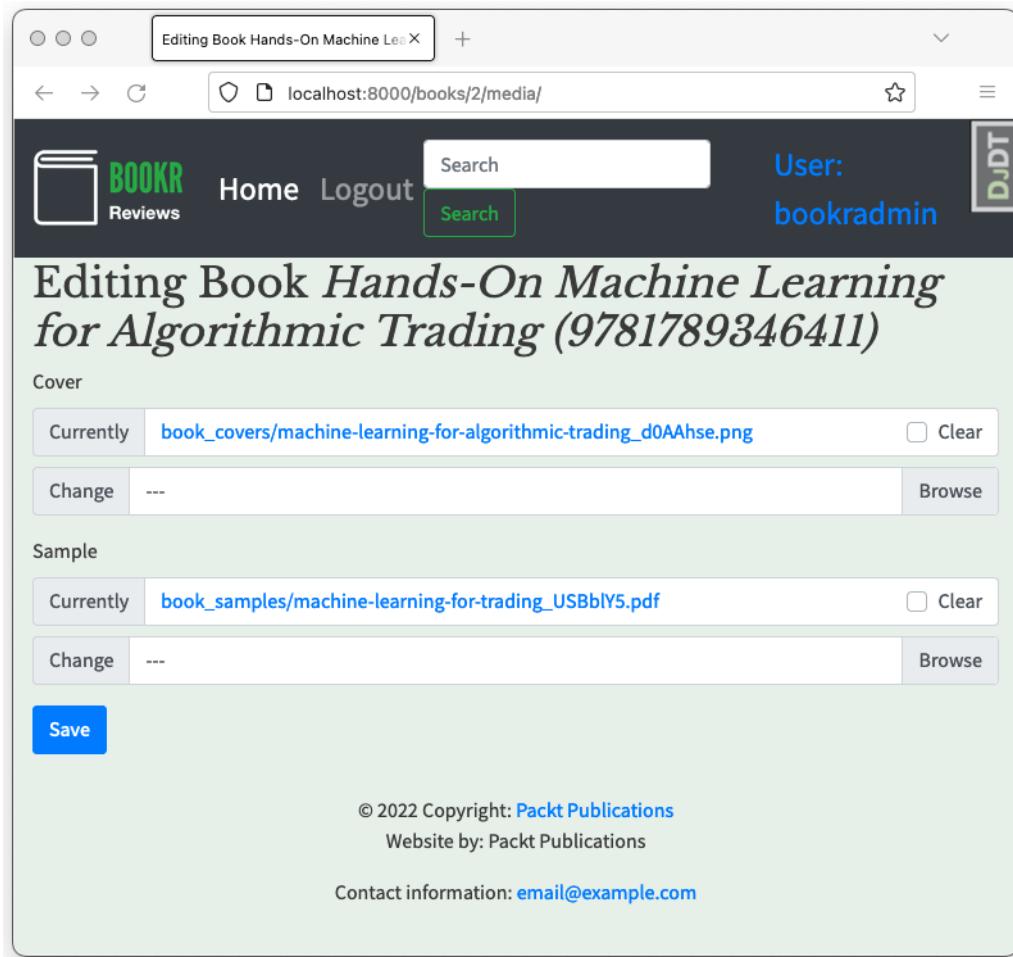


Figure 15.40 – The book media page

Note that the form still behaves fine and allows file uploads. `django-crispy-forms` has automatically added the `enctype="multipart/form-data"` attribute to `<form>`. You can verify this by viewing the page source.

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we introduced five third-party Django apps that can enhance your website. We installed and set up `django-configuration`, which allowed us to easily switch between different settings and change them using environment variables. `dj-database-url` also helped with settings, allowing us to make database settings changes using URLs. We saw how the Django Debug Toolbar could help us see what our app was doing and help us debug problems we had with it. `django-crispy-forms` can not only render our forms using the Bootstrap CSS but also lets us save code by defining their behavior as part of the form class itself. We briefly looked at `django-allauth` and saw how it can be integrated into third-party authentication providers. In the activity for this chapter, we updated our `ModelForm` instances to use the `django-crispy-forms` `FormHelper` and removed some logic from the template by using the `crispy` template tag.

In the next chapter, we will look at how to integrate the React JavaScript framework into a Django application.

16

Using a Frontend JavaScript Library with Django

Django is a great tool for building the backend of an application. You have seen how easy it is to set up the database, route URLs, and render templates. Without using JavaScript, though, when those pages are rendered to the browser, they are static and do not provide any form of interaction. By using JavaScript, your pages can be transformed into applications that are fully interactive in the browser.

This chapter will provide a brief introduction to JavaScript frameworks and how to use them with Django. While it won't be a deep dive into how to build an entire JavaScript application from scratch (that would be a book in itself), we will give enough of an introduction so that you can add interactive components to your own Django application. In this chapter, we will primarily be working with the React framework. Even if you do not have any JavaScript experience, we will introduce enough about it so that you will be comfortable writing your own React components by the end of this chapter. In *Chapter 12, Building a REST API*, you built a REST API for Bookr. We will interact with that API using JavaScript to retrieve data. We will enhance Bookr by showing some review previews on the main page that are dynamically loaded and can be paged through.

Technical requirements

The code for the exercises and activities in this chapter can be found in this book's GitHub repository at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/Chapter16>.

JavaScript frameworks

These days, real-time interactivity is a fundamental part of web applications. While simple interactions can be added without a framework (developing without a framework is often called *Vanilla JS*), as your web application grows, it can be much easier to manage with the use of a framework. Without a framework, you would need to do all these things yourself:

- Manually define the database schema
- Convert data from HTTP requests into native objects
- Write form validation
- Write SQL queries to save data
- Construct HTML to show a response

Compare this to what Django provides. Its **Object Relational Mapping (ORM)**, automatic form parsing and validation, and templating drastically reduce the amount of code you need to write. JavaScript frameworks bring similar time-saving enhancements to JavaScript development. Without them, you would have to manually update the HTML elements in the browser as your data changes. Let's take a simple example: showing the count of the number of times a button has been clicked. Without a framework, you would have to do the following:

1. Assign a handler to the button click event.
2. Increment the variable that stores the count.
3. Locate the element containing the click count display.
4. Replace the element's text with the new click count.

When using a framework, the button count variable is bound to the display (HTML), so the process you have to code is as follows:

1. Handle the button click.
2. Increment the variable.

The framework takes care of automatically re-rendering the number display. This is just a simple example, though; as your application grows, the disparity in complexity between the two approaches expands. Several JavaScript frameworks are available, each with different features, and some are supported and used by large companies. Some of the most popular are React (<https://reactjs.org>), Vue (<http://vuejs.org>), Angular (<https://angularjs.org>), Ember (<https://emberjs.com>), and Backbone.js (<https://backbonejs.org>).

In this chapter, we will use React, as it is easy to drop into an existing web application and allows *progressive enhancement*. This means that rather than having to build your application from scratch, targeting React, you can simply apply it to certain parts of the HTML that Django generates; for

example, a single text field that automatically interprets Markdown and shows the result without reloading the page. We will also cover some of the features that Django offers that can help integrate several JavaScript frameworks.

Several different levels of JavaScript can be incorporated into a web application. *Figure 16.1* shows our current stack with no JavaScript (note that the following diagrams do not show requests to the server):

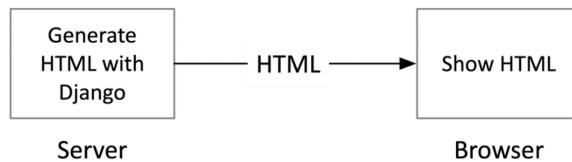


Figure 16.1 – Current stack

You can base your entire application on JavaScript using **Node.js** (a server-side JavaScript interpreter), which would replace Python and Django in the stack. *Figure 16.2* shows how this might look:

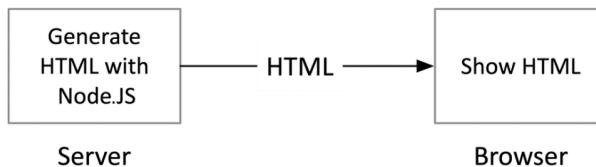


Figure 16.2 – Using Node.js to generate HTML

Or, you can have your frontend and templates entirely in JavaScript and just use Django to act as a REST API to provide data to render. *Figure 16.3* shows this stack:

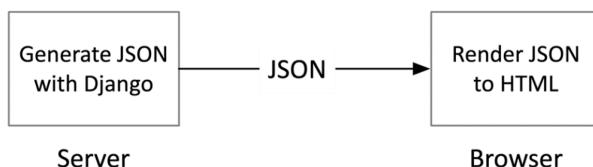


Figure 16.3 – Sending JSON from Django and rendering it in the browser

The final approach is progressive enhancement, which is (as mentioned) what we will be using. In this way, Django is still generating the HTML templates, and React sits on top of this to add interactivity:

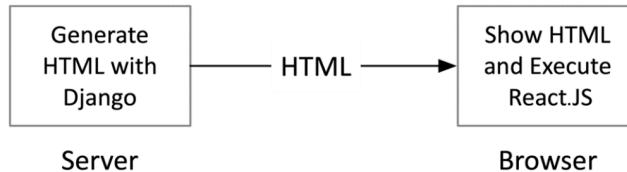


Figure 16.4 – HTML generated with Django with React providing progressive enhancement

Note that it is common to use multiple techniques together. For example, Django may generate the initial HTML to which React is applied in the browser. The browser can then query Django for JSON data to be rendered using React.

To gain an understanding of these approaches, we will first take a brief look at how we can use JavaScript in our web applications.

An introduction to JavaScript

In this section, we will briefly introduce some basic JavaScript concepts. This includes data structures such as variables, constants, arrays, and objects and callables such as functions, classes, and methods. Different operators will be covered as we introduce them.

This will give us the requisite knowledge to begin a brief overview of the React framework.

Loading JavaScript

JavaScript can either be inline in an HTML page or included in a separate JavaScript file. Both methods use the `<script>` tag. With inline JavaScript, the JavaScript code is written directly inside the `<script>` tags in an HTML file; for example, like this:

```
<script>
// comments in JavaScript can start with //
/* Block comments are also supported. This comment is
   multiple lines and doesn't end until we use a star then
   slash:/
let a = 5; // declare the variable a, and set its value to 5
console.log(a); // print a (5) to the browser console
</script>
```

Note that the `console.log` function prints out data to the browser console that is visible in **Developer Tools** of your browser:

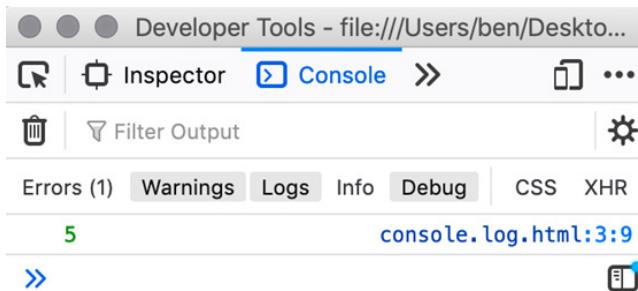


Figure 16.5 – The result of the `console.log(a)` call – 5 is printed to the browser console

We could also put the code into its own file (we would not include the `<script>` tags in the standalone file). We then load it into the page using the `<script>` tag's `src` attribute, as we saw in *Chapter 5, Serving Static Files*:

```
<script src="{% static 'file.js' %}"></script>
```

The source code, whether inline or included, will be executed as soon as the browser loads the `<script>` tag.

Variables and constants

Unlike in Python, variables in JavaScript must be declared using either the `var`, `let`, or `const` keywords:

```
var a = 1; // variable a has the numeric value 1
let b = 'a'; // variable b has the string value 'a'
const pi = 3.14; // assigned as a constant and can't be redefined
```

Just like in Python, though, a type for a variable does not need to be declared. You will notice that the lines end with semicolons. JavaScript does not require lines to be terminated with semicolons – they are optional. However, some style guides enforce their use. You should try to stick with a single convention for any project.

You should use the `let` keyword to declare a variable. Variable declarations are scoped. For example, a variable declared with `let` inside a `for` loop will not be defined outside the loop. In this example, we'll loop through and sum the multiples of 10 to 90 and then print the result to `console.log`. You'll notice we can access variables declared at the function level inside the `for` loop, but not the other way around:

```
let total = 0;
for (let i=0; i< 10; i++){ // variable i is scoped to the loop
    let toAdd = i * 10; // variable toAdd is also scoped
    total += toAdd; /* we can access total since it's in
    the outer scope */
```

```
}

console.log(total); // prints 450
console.log(toAdd); /* throws an exception as the variable is not
                     declared in the outer scope */
console.log(i); /* this code is not executed since an
                  exception was thrown the line before,
                  but it would also generate the same
                  exception */
```

`const` is for constant data and cannot be redefined. That does not mean that the object it points to cannot be changed, though. For example, you couldn't do this:

```
const pi = 3.1416;
pi = 3.1; /* raises exception since const values can't be
            reassigned */
```

The `var` keyword is required by older browsers that don't support `let` or `const`. Only 1% of browsers these days don't support those keywords, so throughout the rest of the chapter, we will only use `let` or `const`. Like `let`, variables declared with `var` can be reassigned; however, they are scoped at the function level only.

JavaScript supports several different types of variables, including strings, arrays, objects (which are like dictionaries), and numbers. We will cover arrays and objects in their own sections now.

Arrays

Arrays are defined similarly to how they are in Python, with square brackets. They can contain different types of data, just like with Python:

```
const myThings = [1, 'foo', 4.5];
```

Another thing to remember with the use of `const` is that it prevents reassigning the constant but does not prevent changing the variable or object being pointed to. For example, we would not be allowed to do this:

```
myThings = [1, 'foo', 4.5, 'another value'];
```

However, you could update the contents of the `myThings` array by using the `push` method (like Python's `list.append`) to append a new item:

```
myThings.push('another value');
```

Objects

JavaScript objects are like Python dictionaries, providing a key-value store. The syntax to declare them is similar as well:

```
const o = {foo: 'bar', baz: 4};
```

Note that, unlike Python, JavaScript object/dictionary keys do not need to be quoted when creating them – unless they contain special characters (spaces, dashes, dots, and more).

The values from `o` can be accessed either with item access or attribute access:

```
o.foo; // 'bar'  
o['baz']; // 4
```

Also note that since `o` was declared as a constant, we cannot reassign it, but we can alter the object's attributes:

```
o.anotherKey = 'another value' // this is allowed
```

Functions

There are a few different ways to define functions in JavaScript. We will look at three. You can define them using the `function` keyword:

```
function myFunc(a, b, c) {  
    if (a == b)  
        return c;  
    else if (a > b)  
        return 0;  
    return 1;  
}
```

All arguments to a function are optional in JavaScript; that is, you could call the preceding function like this: `myFunc()`, and no error would be raised (at least during call time). The `a`, `b`, and `c` variables would all be the `undefined` special type. This would cause issues in the logic of the function. The `undefined` special type is kind of like `None` in Python – although JavaScript also has `null`, which is more similar to `None`. Functions can also be defined by assigning them to a variable (or constant):

```
const myFunc = function(a, b, c) {  
    // function body is implemented the same as above  
}
```

We can also define functions using an arrow syntax. For example, we can also define `myFunc` like this:

```
const myFunc = (a, b, c) => {
    // function body as above
}
```

This is more common when defining functions as part of an object, for example:

```
const o = {
  myFunc: (a, b, c) => {
    // function body
  }
}
```

In this case, it would be called like this:

```
o.myFunc(3, 4, 5);
```

We will return to the reasons for using arrow functions after introducing classes.

Classes and methods

Classes are defined with the `class` keyword. Inside a `class` definition, methods are defined without the `function` keyword. The JavaScript interpreter can recognize the syntax and tell that it is a method. Here is an example class, which takes a number to add (through `toAdd`) when instantiated. That number will be added to whatever is passed to the `add` method, and the result returned:

```
class Adder {
    // A class to add a certain value to any number

    // this is like Python's init method
    constructor (toAdd) {
        /* "this" is like "self" in Python
           it's implicit and not manually passed into every
           method */
        this.toAdd = toAdd;
    }

    add (n) {
        // add our instance's value to the passed in number
        return this.toAdd + n;
    }
}
```

Classes are instantiated with the `new` keyword. Other than that, their usage is very similar to classes in Python:

```
const a = new Adder(5);
console.log(a.add(3)); // prints "8"
```

Arrow functions

Now that we've introduced the `this` keyword, we can return to the purpose of arrow functions. Not only are they shorter to write, but they also preserve the context of `this`. Unlike `self` in Python, which always refers to a specific object because it is passed into methods, the object that `this` refers to can change based on context. Usually, it is due to the nesting of functions, which is common in JavaScript. Let's look at two examples. First, an object with a function called `outer`. This `outer` function contains an `inner` function. We refer to `this` in both the `inner` and `outer` functions:

```
const o1 = {
  outer: function() {
    console.log(this); // "this" refers to o1
    const inner = function() {
      console.log(this); /* "this" refers to the
                           "window" object */
    }
    inner();
  }
}
```

Note

The preceding code example refers to the `window` object. In JavaScript, `window` is a special global variable that exists in each browser tab and represents information about that tab. It is an instance of the `window` class. Some examples of the attributes that `window` has are `document` (which stores the current HTML document), `location` (which is the current location shown in the tab's address bar), and `outerWidth` and `outerHeight` (which represent the width and height of the browser window respectively). For example, to print the current tab's location to the browser console, you would write `console.log(window.location)`.

Inside the `outer` function, `this` refers to `o1` itself, whereas inside the `inner` function, `this` refers to the `window` (an object containing information about the browser window).

Compare this to defining the `inner` function using arrow syntax:

```
const o2 = {
  outer: function() {
    console.log(this); // refers to o2
    const inner = () => {
```

```
        console.log(this); // also refers to o2
    }
    inner();
}
}
```

When we use arrow syntax, `this` is consistent and refers to `o2` in both cases. Now that we have had a very brief introduction to JavaScript, let's introduce React.

Further reading

Covering all the concepts of JavaScript is beyond the scope of this book. For a complete, hands-on course on JavaScript, you can always refer to *The JavaScript Workshop* book: <https://courses.packtpub.com/courses/javascript>.

With this brief review of core JavaScript programming concepts complete, we will commence a brief overview of React programming.

Working with React

React allows you to build applications using components. Each component can *render* itself by generating HTML to be inserted on the page.

A component may also keep track of its own *state*. If it does track its own state, when the state changes, the component will automatically re-render itself. This means if you have an action method that updates a state variable on a component, you don't need to then figure out whether the component needs to be redrawn; React will do this for you. A web app should track its own state so that it doesn't need to query the server to find out how it needs to update to display data.

Data is passed between components using properties, or *props* for short. The method of passing properties looks kind of like HTML attributes, but there are some differences, which we will cover later in the chapter. Properties are received by a component in a single `props` object.

To illustrate with an example, you might build a shopping list app with React. You would have a component for the list container (`ListContainer`), and a component for a list item (`ListItem`). `ListContainer` would be instantiated multiple times, once for each item on the shopping list. The container would hold a state containing a list of the items' names. Each item name would be passed to the `ListContainer` instances as a *prop*. Each `ListContainer` would then store the item's name and an `isBought` flag in its own state. As you click an item to mark it off the list, `isBought` would be set to `true`. Then, React would automatically call `render` on `ListContainer` to update the display.

There are a few different methods of using React with your application. If you want to build a deep and complex React application, you should use `npm` (Node Package Manager, a tool for managing

Node.js applications) to set up a React project. Since we will use React to enhance some of our pages, we can just include the React framework code using a `<script>` tag:

```
<script crossorigin
src="https://unpkg.com/react@16/umd/react.development.js">
</script>
<script crossorigin
src="https://unpkg.com/react-dom@16/umd/react-dom.development.js">
</script>
```

Note

The `crossorigin` attribute is for security and means cookies or other data cannot be sent to the remote server. This is necessary when using a public **content delivery network (CDN)** such as `unpkg.com` in case a malicious script has been hosted there by someone.

These should be placed on a page that you want to add React to just before the closing `</body>` tag. The reason for putting the tags here instead of in `<head>` of the page is that the script might want to refer to HTML elements on the page. If we put the script tag in the head, it will be executed before the page elements are available (as they come after).

Note

The links to the latest React versions can be found at <https://legacy.reactjs.org/docs/cdn-links.html>.

Components

There are two ways to build a component in React: with functions or with classes. Regardless of the approach, to get displayed on a page, the component must return some HTML elements to display. A functional component is a single function that returns elements, whereas a class-based component will return elements from its `render` method. Functional components cannot keep track of their own state.

React is like Django in that it automatically escapes HTML in strings that are returned from `render`. To generate HTML elements, you must construct them using their tag, the attributes/properties they should have, and their content. This is done with the `React.createElement` function. A component will return a React element, which may contain sub-elements.

Let us look at two implementations of the same component, first as a function and then as a class. The functional component takes `props` as an argument. This is an object containing the properties that are passed to it. The following function returns an `h1` element:

```
function HelloWorld(props) {
  return React.createElement('h1', null, 'Hello, ' +
```

```
    props.name + '!')  
}
```

Note that it is conventional for the function to have an uppercase first character.

While a functional component is a single function that generates HTML, a class-based component must implement a `render` method to do this. The code in the `render` method is the same as in the functional component, with one difference: the class-based component accepts the `props` object in its constructor, and then the `render` (or other) methods can refer to `props` using `this.props`. Here is the same `HelloWorld` component implemented as a class:

```
class HelloWorld extends React.Component {  
  render() {  
    return React.createElement('h1', null, 'Hello, ' +  
      this.props.name + '!');  
  }  
}
```

When using classes, all components extend from the `React.Component` class. Class-based components have an advantage over functional components; they encapsulate the handling actions/event and their own state. For simple components, using the functional style means less code. For more information on components and properties, see <https://reactjs.org/docs/components-and-props.html>.

Whichever method you choose to define a component, it is used in the same way. In this chapter, we will only be using class-based components.

We first need to add a place for React to render this component onto an HTML page. Normally, this is done using `<div>` with an `id` attribute. The following is an example of this:

```
<div id="react_container"></div>
```

Note that `id` does not have to be `react_container`; it just needs to be unique for the page. Then, in the JavaScript code, after defining all your components, they are rendered on the page using the `ReactDOM.render` function. This takes two arguments: the root React element (not the component) and the HTML element in which it should be rendered.

We would use it like this:

```
const container = document.getElementById('react_container');  
const componentElement = React.createElement(HelloWorld,  
  {name: 'Ben'});  
ReactDOM.render(componentElement, container);
```

Note that the `HelloWorld` component (class/function) itself is not being passed to the `render` function; it is wrapped in a `React.createElement` call to instantiate it and transform it into an element.

As you might have guessed from its name, the `document.getElementById` function locates an HTML element in the document and returns a reference to it.

The final output in the browser when the component is rendered is like this:

```
<h1>Hello, Ben!</h1>
```

Let's look at a more advanced example component. Note that since `React.createElement` is such a commonly used function, it's common to alias to a shorter name, such as `e`: that's what the first line of this example does.

This component displays a button and has an internal state that keeps track of how many times the button was clicked. First, let's look at the component class in its entirety:

```
const e = React.createElement;

class ClickCounter extends React.Component {
  constructor(props) {
    super(props);
    this.state = { clickCount: 0 };
  }

  render() {
    return e(
      'button', // the element name
      {onClick: () => this.setState({
        clickCount: this.state.clickCount + 1 })}, //element
      props
      this.state.clickCount // element content
    );
  }
}
```

Some things to note about the `ClickCounter` class are as follows:

- The `props` argument is an object (dictionary) of attribute values that have been passed to the component when it is used in HTML, as shown here:

```
<ClickCounter foo="bar" rex="baz" />
```

The `props` dictionary would contain the `foo` key with a `bar` value and the `rex` key with the `baz` value.

- `super(props)` calls the super class's `constructor` method and passes the `props` variable. This is analogous to the `super()` method in Python.
- Each React class has a `state` variable, which is an object. `constructor` can initialize it. The state should be changed using the `setState` method rather than being manipulated directly. When it is changed, the `render` method will be automatically called to redraw the component.

The `render` method returns a new HTML element using the `React.createElement` function (remember, the `e` variable was aliased to this function). In this case, the arguments to `React.createElement` will return a `<button>` element with a click handler and with the `this.state.clickCount` text content. Essentially, it will return an element like this (when `clickCount` is 0):

```
<button onClick="this.setState(...)">  
  0  
</button>
```

The `onClick` function is set as an anonymous function with arrow syntax. This is similar to having a function as follows (although not quite the same since it's in a different context):

```
const onClick = () => {  
  this.setState({clickCount: this.state.clickCount + 1})  
}
```

Since the function is only one line, we can also remove one set of wrapping braces, and we end up with this:

```
{ onClick: () => this.setState({  
  clickCount: this.state.clickCount + 1}) }
```

We covered how to place `ClickCounter` onto a page earlier in this section, something like this:

```
ReactDOM.render(e(ClickCounter),  
  document.getElementById('react_container'));
```

The following screenshot shows the counter in the button when the page loads:

Furthermore, `DjDt` refers to the debug toolbar that we learned about in the *Django Debug Toolbar* section in *Chapter 15, Django Third-Party Libraries*:



Figure 16.6 – Button with 0 for the count

After clicking the button a few times, the button looks as shown in *Figure 16.7*:



Figure 16.7 – Button after clicking seven times

Now, just to demonstrate how *not* to write the `render` function, we'll look at what happens if we just return HTML as a string, like this:

```
render() {
    return '<button>' + this.state.clickCount + '</button>'
```

Now the rendered page looks as shown in *Figure 16.8*:

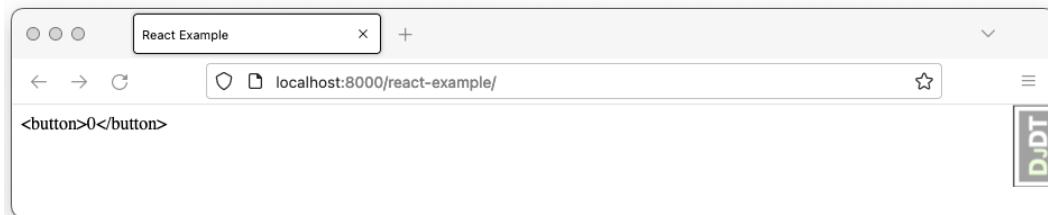


Figure 16.8 – Returned HTML rendered as a string

This shows React's automatic escaping of HTML in action. Now that we have had a brief intro to JavaScript and React, let's add an example page to Bookr so you can see it in action.

Exercise 16.01 – setting up a React example

In this exercise, we will create an example view and template to use with React. Then we will implement the `ClickCounter` component. At the end of the exercise, you will be able to interact with it with the `ClickCounter` button:

1. In PyCharm, go to **New | File** inside the project's `static` directory. Name the new file `react-example.js`.
2. Putting this code inside it will define the React component, then render it into the `react_container` `<div>` that we will be creating:

```
const e = React.createElement;

class ClickCounter extends React.Component {
    constructor(props) {
        super(props);
        this.state = { clickCount: 0 };
    }

    render() {
        return e(
            'button', { onClick: () => this.setState({
                clickCount: this.state.clickCount + 1
            })
        },
        this.state.clickCount
    );
}
}

ReactDOM.render(e(ClickCounter),
    document.getElementById('react_container'));
```

You can now save `react-example.js`.

3. Go to **New | HTML File** inside the project's templates directory:

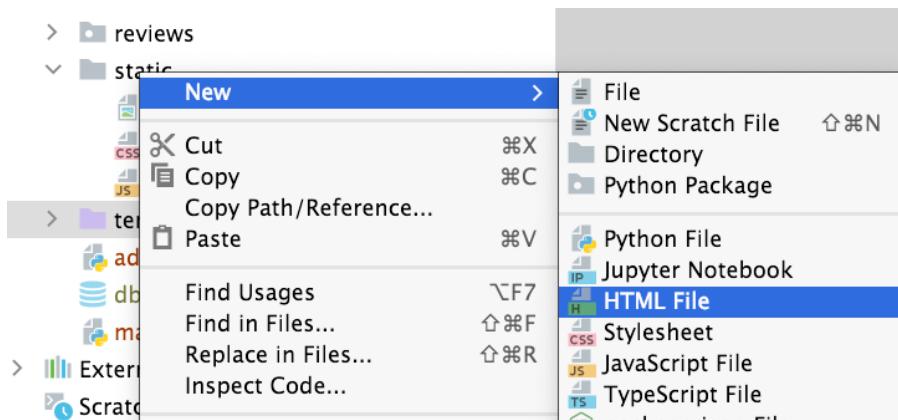


Figure 16.9 – Create a new HTML file

4. Name the new file `react-example.html`:

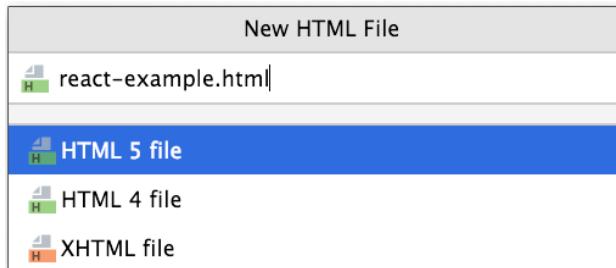


Figure 16.10 – Name the file `react-example.html`

You can change the title inside the `<title>` element to *React Example*, but that is not necessary for this exercise.

5. `react-example.html` is created with some HTML boilerplate as we have seen before. Add the following `<script>` tags to include React just before the closing `</body>` tag:

```
<script crossorigin
src="https://unpkg.com/react@16/umd/react.development.js">
</script>
<script crossorigin src="https://unpkg.com/
react-dom@16/umd/react-dom.development.js"></script>
```

6. The `react-example.js` file will be included using a `<script>` tag, and we need to generate the script path using the `static` template tag. First, load the static template library at the start of the file by adding the following on the second line:

```
{% load static %}
```

The first few lines of your file will look like *Figure 16.11*:

```
<!DOCTYPE html>
{% load static %}
<html lang="en">
<head>
```

Figure 16.11 – The load static template tag included

Then, just before the closing `</body>` tag, but after the `<script>` tags that were added in *step 5*, add this script tag to include `react-example.js`:

```
<script src="{% static 'react-example.js' %}"></script>
```

7. We now need to add the containing `<div>` that React will render into. Add this element after the opening `<body>` tag:

```
<div id="react_container"></div>
```

You can save `react-example.html`.

8. Now we'll add a view to render the template. Open the `reviews` app's `views.py` file and add a `react_example` view at the end of the file:

```
def react_example(request):
    return render(request, "react-example.html")
```

In this simple view, we just render the `react-example.html` template with no context data.

9. Finally, we need to map a URL to the new view. Open the `bookr` package's `urls.py` file. Add this map to the `urlpatterns` variable:

```
path('react-example/',
      reviews.views.react_example)
```

You can save and close `urls.py`.

10. Start the Django dev server if it's not already running, then go to `http://127.0.0.1:8000/react-example/`. You should see the `ClickCount` button rendered as in *Figure 16.12*:



Figure 16.12 – The ClickCount button

Try clicking the button a few times and watch the counter increment.

In this example, we created our first React component, then added a template and view to render it. We included the React framework source from a CDN. In the next section, we will introduce **JavaScript XML (JSX)**, which combines templates and code into a single file that can simplify our code.

JSX – a JavaScript syntax extension

It can be quite verbose to define each element using the `React.createElement` function – even when we alias to a shorter variable name. The verbosity is exacerbated when we start building larger components.

When using React, we can use JSX to build the HTML elements instead. JSX stands for JavaScript XML since both JavaScript and XML are written in the same file. For example, consider the following code in which we are creating a button using the `render` method:

```
return React.createElement('button', { onClick: ... },
  'Button Text')
```

Instead of this, we can return its HTML directly, as follows:

```
return <button onClick={...}>Button Text</button>;
```

Note that the HTML is not quoted and returned as a string. That is, we are not doing this:

```
return '<button onClick={...}>Button Text</button>';
```

Since JSX is an unusual syntax (a combination of HTML and JavaScript in a single file), we need to include another JavaScript library before it can be used: Babel (<https://babeljs.io>). This is a library that can *transpile* code between different versions of JavaScript. You can write code using the latest syntax and have it *transpiled* (a combination of translate and compile) into a version of code that older browsers can understand.

Babel can be included with a `<script>` tag like this:

```
<script crossorigin
src="https://unpkg.com/babel-standalone@6/babel.min.js"></script>
```

This should be included on the page after your other React-related script tags but before you include any files containing JSX.

Any JavaScript source code that includes JSX must have the `type="text/babel"` attribute added:

```
<script src="path/to/file.js" type="text/babel"></script>
```

This is so Babel knows to parse the file rather than just treat it as plain JavaScript.

Note

Note that using Babel in this way can be slow for large projects. It is designed to be used as part of the build process in an npm project and to have your JSX files transpiled ahead of time (rather than in real time, as we are doing here). The npm project setup is beyond the scope of this book. For our purposes and with the small amount of JSX we are using, using Babel will be fine.

JSX uses braces to include JavaScript data inside HTML, similar to Django's double braces in templates. JavaScript inside braces will be executed. We'll now look at how to convert our button creation example to JSX. Our `render` method can be changed to this:

```
render() {
  return <button onClick={() =>this.setState({
    clickCount: this.state.clickCount + 1
  )}>
    {this.state.clickCount}
  </button>;
}
```

Note that the `onClick` attribute has no quotes around its value; instead, it is wrapped in braces. This is passing the JavaScript function that is defined inline to the component. It will be available in that component's `props` dictionary that is passed to the `constructor` method. For example, imagine that we had passed it like this:

```
onClick="() =>this.setState..."
```

In such a case, it would be passed to the component as a string value and thus would not work.

We are also rendering the current value of `clickCount` as the content of the button. JavaScript could be executed inside these braces too. To show the click count plus one, we could do this:

```
{this.state.clickCount + 1}
```

In the next exercise, we will include Babel in our template and then convert our component to use JSX.

Exercise 16.02 – JSX and Babel

In this exercise, we will implement JSX in our component to simplify our code. To do this, we need to make a couple of changes to the `react-example.js` and `react-example.html` files to switch to JSX to render `ClickCounter`:

1. In PyCharm, open `react-example.js` and change the `render` method to use JSX instead by replacing it with the following code. You can refer to *step 2* from *Exercise 16.01 – setting up a React example*, where we defined this method:

```
render() {
  return <button onClick={() => this.setState({
    clickCount: this.state.clickCount + 1
  })}
  >
  {this.state.clickCount}
  </button>;
}
```

2. We can now treat `ClickCounter` as an element itself. In the `ReactDOM.render` call at the end of the file, you can replace the first argument, `e(ClickCounter)`, with a `<ClickCounter/>` element, like this:

```
ReactDOM.render(<ClickCounter/>,
  document.getElementById('react_container'));
```

3. Since we're no longer using the `React.createElement` function that we created in *step 2* of *Exercise 16.01 – setting up a React example*, we can remove the alias we created by deleting the first line:

```
const e = React.createElement;
```

You can save and close the file.

4. Open the `react-example.html` template. You need to include the Babel library JavaScript. Add this code between the React `script` elements and the `react-example.js` element:

```
<script crossorigin
src="https://unpkg.com/babelstandalone@6/babel.min.js">
</script>
```

5. Add a `type="text/babel"` attribute to the `react-example.html <script>` tag:

```
<script src="% static 'react-example.js' %"
type="text/babel"></script>
```

Save `react-example.html`.

6. Start the Django dev server if it is not already running and go to `http://127.0.0.1:8000/react-example/`. You should see the same button as we had before (*Figure 16.12*). When clicking the button, you should see the count increment as well.

In this exercise, we did not change the behavior of the `ClickCounter` React component. Instead, we refactored it to use JSX. This makes it easier to write the component's output directly as HTML and cut down on the amount of code we need to write. In the next section, we will look at passing properties to a JSX React component.

JSX properties

Properties on JSX-based React components are set in the same way as attributes on a standard HTML element. The important thing to remember is whether you are setting them as a string or a JavaScript value.

Let's look at some examples using the `ClickCounter` component. Say that we want to extend `ClickCounter` so that a `target` number can be specified. When the target is reached, the button should be replaced with the `Well done, <name>!` text. These values should be passed into `ClickCounter` as properties.

When using variables, we have to pass them as JSX values:

```
let name = 'Ben'
let target = 5;

ReactDOM.render(
  <ClickCounter name={name} target={target}/>,
  document.getElementById('react_container'));
```

We can mix and match the method of passing the values too. This is also valid:

```
ReactDOM.render(<ClickCounter name="Ben" target={5}/>,
  document.getElementById('react_container'));
```

In the next exercise, we will update `ClickCounter` to read these values from properties and change its behavior when the target is reached. We will pass these values in from the Django template.

Exercise 16.03 – React component properties

In this exercise, we will modify `ClickCounter` to read the values of `target` and `name` from its `props`. We will pass these in from the Django view and use the `escapejs` filter to make the `name` value safe for use in a JavaScript string. When you are finished, you will be able to click on the button until it reaches a target, and then see a `Well done` message. Let's get started with the steps:

1. In PyCharm, open the `reviews` app's `views.py`. We will modify the `react_example` view's `render` call to pass through a context containing `name` and `target`, like this:

```
return render(request, "react-example.html",
  {"name": "Ben", "target": 5})
```

You can use your own name and pick a different target value if you like. Save `views.py`.

2. Open the `react-example.js` file. We will update the `state` setting in the `constructor` method to set the `name` and `target` from `props`, like this:

```
constructor(props) {
  super(props);
  this.state = { clickCount: 0, name: props.name,
    target: props.target};
}
```

3. Change the behavior of the `render` method to return `Well done, <name>!` once `target` has been reached. Add this `if` statement inside the `render` method:

```
if (this.state.clickCount === this.state.target) {
  return <span>Well done, {this.state.name}!</span>;
}
```

4. To pass the values in, move the `ReactDOM.render` call into the template so that Django can render that piece of code. Cut this `ReactDOM.render` line from the end of `react-example.js`:

```
ReactDOM.render(<ClickCounter/>,
  document.getElementById('react_container'));
```

We will paste it into the template file in *step 6*. `react-example.js` should now only contain the `ClickCounter` class. Save and close the file.

5. Open `react-example.html`. After all the existing `<script>` tags (but before the closing `</body>` tag), add opening and closing `<script>` tags with the `type="text/babel"` attribute. Inside them, we need to assign the Django context values that were passed to the template as JavaScript variables. Altogether, you should be adding this code:

```
<script type="text/babel">
    let name = "{{ name | escapejs }}";
    let target = {{ target }};
</script>
```

The first assigns the `name` variable with the `name` context variable. We use the `escapejs` template filter; otherwise, we could generate invalid JavaScript code if our `name` had a double quote in it. The second value, `target`, is assigned from `target`. This is a number, so it does not need to be escaped.

Note

Due to the way Django escapes the values for JavaScript, `name` cannot be passed directly to the component property like this:

```
<ClickCounter name="{{ name | escapejs }}"/>
```

The JSX will not unescape the values correctly and you will end up with escape sequences.

However, you could pass the numerical value `target` in like this:

```
<ClickCounter target="{{ target }}"/>
```

Also, be aware of the spacing between the Django braces and JSX braces. In this book, we will first assign all properties to variables, then pass them to the component for consistency.

6. Underneath these variable declarations, paste in the `ReactDOM.render` call that you copied from `react-example.js`. Then, add the `target={{ target }}` and `name={{ name }}` properties to `ClickCounter`. Remember, these are the JavaScript variables being passed in, not the Django context variables – they just happen to have the same name. The `<script>` block should now look like this:

```
<script type="text/babel">
    let name = "{{ name | escapejs }}";
    let target = {{ target }};
    ReactDOM.render(
        <ClickCounter name={{ name }} target={{ target }}/>,
        document.getElementById('react_container'));
</script>
```

You can save `react-example.html`.

7. Start the Django dev server if it is not already running, then go to `http://127.0.0.1:8000/react-example/`. Try clicking the button a few times – it should increment until you click it the target number of times. Then, it will be replaced with the `Well done, <name>!` text. See *Figure 16.13* for how it should look after you've clicked it enough times:



Figure 16.13 – The Well done message

In this exercise, we passed data to a React component using `props`. We escaped the data when assigning it to a JavaScript variable using the `escapejs` template filter. In the next section, we will cover how to fetch data over HTTP using JavaScript.

Further reading

For a more detailed, hands-on course on React, you can always refer to *The React Workshop*: <https://courses.packtpub.com/courses/react>.

JavaScript promises

Many JavaScript functions are implemented asynchronously to prevent blocking on long-running operations. They work by returning immediately but then invoking a callback function when a result is available. The object these types of functions return is a `Promise` object. Callback functions are provided to the `Promise` object by calling its `then` method. When the function finishes running, it will either `resolve` the `Promise` (call the `success` function) or `reject` it (call the `failure` function).

We will illustrate the wrong and right ways of using promises. Consider a hypothetical long-running function that performs a big calculation called `getResult`. Instead of returning the result, it returns `Promise`. You would not use it like this:

```
const result = getResult();
console.log(result); // incorrect, this is a Promise
```

Instead, it should be invoked like this, with a callback function passed to `then` on the returned `Promise`. We will assume that `getResult` can never fail, so we only provide it with a `success` function for the `resolve` case:

```
const promise = getResult();
promise.then((result) => {
    console.log(result); /* this is called when the
                           Promise resolves*/
});
```

Normally, you wouldn't assign the returned `Promise` to a variable. Instead, you'd chain the `then` call to the function call. We'll show this in the next example, along with a failure callback (assume `getResult` can now fail). We'll also add some comments illustrating the order in which the code executes:

```
getResult().then(
    (result) => {
        // success function
        console.log(result);
        // this is called 2nd, but only on success
    },
    () => {
        // failure function
        console.log("getResult failed");
        // this is called 2nd, but only on failure
    }
)

// this will be called 1st, before either of the callbacks
console.log("Waiting for callback");
```

Now that we've introduced promises, we can look at the `fetch` function, which makes HTTP requests. It is asynchronous and works by returning promises.

The `fetch` function

Most browsers (95%) support a function called `fetch`, which allows you to make HTTP requests. It uses an asynchronous callback interface with promises.

The `fetch` function takes two arguments. The first is the URL to make the request from, and the second is an object (dictionary) with settings for the request. For example, consider this:

```
const promise = fetch("http://www.google.com", { ...settings });
```

The settings are things such as the following:

- `method`: The request HTTP method (GET, POST, and more).
- `headers`: Another object (dictionary) of HTTP headers to send.
- `body`: The HTTP body to send (for the POST/PUT requests).
- `credentials`: By default, `fetch` does not send any cookies. This means your requests will act like you are not authenticated. To have it set cookies in its requests, this should be set to the `same-origin` or `include` values.

Let's look at it in action with a simple request:

```
fetch('/api/books/', {
  method: 'GET',
  headers: {
    Accept: 'application/json'
  }
}).then((resp) => {
  console.log(resp)
})
```

The preceding code will fetch from `/api/book-list/`, and then call a function that logs the request to the browser's console using `console.log`.

Figure 16.14 shows the console output in Firefox for the preceding response:

```
▼ Response
  ▶ body: ReadableStream { locked: false }
  ▶ bodyUsed: false
  ▶ headers: Headers { }
  ▶ ok: true
  ▶ redirected: false
  ▶ status: 200
  ▶ statusText: "OK"
  ▶ type: "basic"
  ▶ url: "http://127.0.0.1:8000/api/books/"
  ▶ <prototype>: ResponsePrototype { clone: clone(), arrayBuffer: arrayBuffer(), blob: blob(), ... }
```

Figure 16.14 – Response output in the console

As you can see, the console contains little useful information as it stands. We need to decode the response before we can work with it. We can use the `json` method on the response object to decode the response body to a JSON object. This also returns `Promise`, so we will ask to get the JSON, then work with the data in our callback. The full code block to do that looks like this:

```
fetch('/api/books/', {
  method: 'GET',
  headers: {
    Accept: 'application/json'
  }
}).then((resp) => {
  return resp.json(); /* doesn't return JSON, returns a
                        Promise */
}).then((data) => {
  console.log(data);
});
```

This will log the decoded object that was in JSON format to the browser console. In Firefox, the output looks like *Figure 16.15*:

```
▼ (18) [...]
  ▶ 0: Object { title: "Advanced Deep Learning with Keras", publication_date: "2018-10-31", isbn: "9781788629416", ... }
  ▶ 1: Object { title: "Hands-On Machine Learning for Algorithmic Trading", publication_date: "2018-12-31", isbn: "9781789346411", ... }
  ▶ 2: Object { title: "Architects of Intelligence", publication_date: "2018-11-23", isbn: "9781789954531", ... }
  ▶ 3: Object { title: "Deep Reinforcement Learning Hands-On", publication_date: "2018-06-20", isbn: "978178834247", ... }
  ▶ 4: Object { title: "Natural Language Processing with TensorFlow", publication_date: "2018-05-30", isbn: "9781788478311", ... }
  ▶ 5: Object { title: "Hands-On Reinforcement Learning with Python", publication_date: "2018-06-27", isbn: "978178836524", ... }
  ▶ 6: Object { title: "Brave New World", publication_date: "2006-10-18", isbn: "9780060850524", ... }
  ▶ 7: Object { title: "The Grapes of Wrath", publication_date: "2006-03-28", isbn: "9780143039433", ... }
  ▶ 8: Object { title: "For Whom The Bell Tolls", publication_date: "2019-07-16", isbn: "9781476787770", ... }
  ▶ 9: Object { title: "To Kill A Mocking Bird", publication_date: "2002-01-01", isbn: "9780060935467", ... }
  ▶ 10: Object { title: "The Great Gatsby", publication_date: "2004-09-30", isbn: "9780743273565", ... }
  ▶ 11: Object { title: "The Catcher in the Rye", publication_date: "2001-01-30", isbn: "9780316769174", ... }
  ▶ 12: Object { title: "Fahrenheit 451", publication_date: "2012-01-10", isbn: "9781451673319", ... }
```

Figure 16.15 – The decoded book list output to the console

In *Exercise 16.04 – fetching and rendering books*, we will write a new React component that will fetch a list of books and then render each one as a list item (``). Before that, we need to learn about the JavaScript `map` method and how to use it to build HTML in React.

The JavaScript map method

Sometimes, we want to execute the same piece of code (JavaScript or JSX) multiple times for different input data. In this chapter, it will be most useful to generate JSX elements with the same HTML tags but different content. In JavaScript, the `map` method iterates over the target array and then executes a callback function for each element in the array. Each of these elements is then added to a new array, which is then returned. For example, this short snippet uses `map` to double each number in the `numbers` array:

```
const numbers = [1, 2, 3];
const doubled = numbers.map((n) => {
    return n * 2;
});
```

The `doubled` array now contains the `[2, 4, 6]` values.

We can also create a list of JSX values using this method. The only thing to note is that each item in the list must have a unique `key` property set. In this next short example, we are transforming an array of numbers into `` elements. We can then use them inside ``. Here is an example `render` function to do this:

```
render() {
    const numbers = [1, 2, 3];
    const listItems = numbers.map((n) => {
        return <li key={n}>{n}</li>;
    });
    return <ul>{listItems}</ul>
}
```

When rendered, this will generate the following HTML:

```
<ul>
  <li>1</li>
  <li>2</li>
  <li>3</li>
</ul>
```

In the next exercise, we will build a React component with a button that will fetch the list of books from the API when it is clicked. The list of books will then be displayed.

Exercise 16.04 – fetching and rendering books

In this exercise, we will create a new component named `BookDisplay` that renders an array of books inside ``. The books will be retrieved using `fetch`. To do this, we add the React component to the `react-example.js` file. Then we pass the URL of the book list to the component inside the Django template:

1. In PyCharm, open `react-example.js`, which you previously used in *step 9* of *Exercise 16.03 – React component properties*. You can delete the entire `ClickCounter` class.
2. Create a new class called `BookDisplay` that uses `extends` to extend from `React.Component`.
3. Then, add a `constructor` method that takes `props` as an argument. It should call `super(props)` and then set its state like this:

```
this.state = { books: [], url: props.url,
                fetchInProgress: false};
```

This will initialize `books` as an empty array, read the API URL from the passed-in `url` property, and set a `fetchInProgress` flag to `false`. The code of your `constructor` method should be like this:

```
constructor(props) {
  super(props);
  this.state = { books: [], url: props.url,
                fetchInProgress: false };
}
```

4. Next, add a `doFetch` method. You can copy and paste this code to create it:

```
doFetch() {
  if (this.state.fetchInProgress)
    return;

  this.setState({ fetchInProgress: true })

  fetch(this.state.url, {
    method: 'GET',
    headers: {
      Accept: 'application/json'
    }
  })
  .then((response) => {
    return response.json();
  })
  .then((data) => {
    this.setState({ books: data })
  })
  .catch((error) => {
    console.error(error)
  })
}
```

```
        }).then((data) => {
            this.setState({ fetchInProgress: false,
                books: data })
        })
    }
```

First, with the `if` statement, we check whether `fetch` has already been started. If so, we return from the function. Then, we use `setState` to update the state, setting `fetchInProgress` to `true`. This will both update our button display text and stop multiple requests from being run at once. We then use `fetch` to fetch `this.state.url` (which we will pass in through the template later in the exercise). The response is retrieved with the `GET` method, and we only want to use `Accept` to accept a JSON response. After we get a response, we then return its JSON using the `json` method. This returns `Promise`, so we use another `then` to handle the callback when the JSON is parsed. In that final callback, we set the state of the component, with `fetchInProgress` going back to `false` and the `books` array being set to the decoded JSON data.

5. Next, create the `render` method. You can copy and paste this code too:

```
render() {
    const bookListItems = this.state.books.map((book) =>
    {
        return <li key={ book.pk }>{ book.title }</li>;
    })

    const buttonText = this.state.fetchInProgress  ?
        'Fetch in Progress' : 'Fetch';

    return <div>
<ul>{ bookListItems }</ul>
<button onClick={ () =>this.doFetch() } disabled={ this.state.fetchInProgress }>{buttonText}</button>
</div>;
}
```

This uses the `map` method to iterate over the array of books in `state`. We generate `` for each book, using the book's `pk` as the `key` instance for the list item. The content of `` is the book's title. We define a `buttonText` variable to store (and update) the text that the button will display. If we currently have a `fetch` operation running, then this will be **Fetch in Progress**. Otherwise, it will be **Fetch**. Finally, we return `<div>` that contains all the data we want. The content of `` is the `bookListItems` variable (the array of `` instances). It also contains a `<button>` instance added in a similar way to the previous exercises. The `onClick` method calls the `doFetch` method of the class. We can use `disabled` to make the button disabled (that is, the user can't click the button) if there is a fetch in progress. We set the button text to the `buttonText` variable we created earlier. You can now save and close `react-example.js`.

6. Open `react-example.html`. We need to replace the `ClickCounter` render (from *Exercise 16.03 – React component properties*) with a `BookDisplay` render. Delete the `name` and `target` variable definitions. We will instead render `<BookDisplay>`. Set the `url` property as a string and pass in the URL to the book list API, using the `{% url %}` template tag to generate it. The `ReactDOM.render` call should then look like this:

```
ReactDOM.render
  (<BookDisplay url="{% url 'api:book-list' %}" />,
  document.getElementById('react_container'));
```

You can now save and close `react-example.html`.

7. Start the Django dev server if it's not already running, then visit `http://127.0.0.1:8000/react-example/`. You should see a single **Fetch** button on the page (*Figure 16.16*):



Figure 16.16 – The book Fetch button

After clicking the **Fetch** button, it should become disabled and have its text changed to **Fetch in Progress**, as we can see here:



Figure 16.17 – Fetch in Progress

Once the fetch is complete, you should see the list of books rendered as follows:

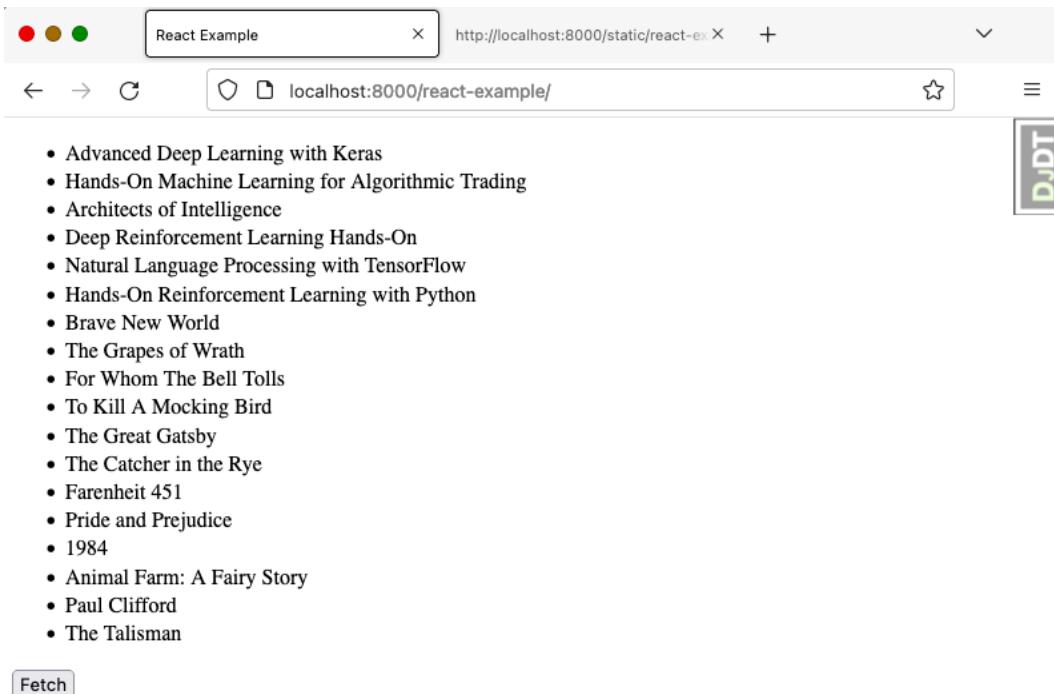


Figure 16.18 – The book fetch complete

This exercise was a chance to integrate React with the Django REST API you built in *Chapter 12, Building a REST API*. We built a new component (BookDisplay) with a call to `fetch` to get a list of books. We used the JavaScript `map` method to transform the book array to some `` elements. As we had seen before, we used `button` to trigger `fetch` when it was clicked. We then provided the book list API URL to the React component in the Django template. Later, we saw a list of books in Bookr that were loaded dynamically using the REST API.

Before we move on to the activity for this chapter, we will talk about some considerations for other JavaScript frameworks when working with Django.

The `verbatim` template tag

We have seen that when using React, we can use JSX interpolation values in Django templates. This is because JSX uses single braces to interpolate values, and Django uses double braces. It should work fine as long as there are spaces between the JSX and Django braces.

Other frameworks, such as Vue, also use double braces for variable interpolation. What that means is if you had a Vue component's HTML in your template, you might try to interpolate a value like this:

```
<h1>Hello, {{ name }}!</h1>
```

Of course, when Django renders the template, it will interpolate the `name` value before the Vue framework gets a chance to render.

We can use the `verbatim` template tag to have Django output the data exactly as it appears in the template without performing any rendering or variable interpolation. Using it with the previous example is simple, as shown here:

```
{% verbatim %}
<h1>Hello, {{ name }}!</h1>
{% endverbatim %}
```

Now when Django renders the template, the HTML between the template tags will be output exactly as it is written, allowing Vue (or another framework) to take over and interpolate the variables itself. Many other frameworks separate their templates into their own files, which should not conflict with Django's templates.

Many JavaScript frameworks are available, and which one you ultimately decide to use will depend on your own opinion or what your company/team uses. If you do run into conflicts, the solution will depend on your particular framework. The examples in this section should help lead you in the right direction.

We have now covered most things you will need to integrate React (or other JavaScript frameworks) with Django. In the next activity, you will implement these learnings to fetch the most recent reviews on Bookr.

Activity 16.01 – reviews preview

In this activity, we will update the Bookr main page to fetch the six most recent reviews and display them. The user will be able to click buttons to go forward to the next six reviews and then back to the previous ones.

These steps will help you complete the activity:

1. First, we can clean up some code from previous exercises. If you like, you can take backups of these files to preserve them for later reference. Alternatively, you can use the GitHub versions too, for future reference. Delete the `react_example` view, the `react-example` URL, the `react-example.html` template, and the `react-example.js` file.
2. Create a `recent-reviews.js` static file.
3. Create two components, a `ReviewDisplay` component that displays the data for a single review and a `RecentReviews` component that handles fetching the review data and displaying a list of the `ReviewDisplay` components.

First, create the `ReviewDisplay` class. In its constructor, you should read `review` being passed in through `props` and assign it to the state.

4. The `render` method of `ReviewDisplay` should return JSX HTML like this:

```
<div className="col mb-4">
  <div className="card">
    <div className="card-body">
      <h5 className="card-title">{ BOOK_TITLE }</h5>
      <strong>({ REVIEW_RATING })</strong>
      </h5>
      <h6 className="card-subtitle mb-2 text-muted">
        CREATOR_EMAIL</h6>
      <p className="card-text">REVIEW_CONTENT</p>
    </div>
    <div className="card-footer">
      <a href={'/books/' + BOOK_ID + '/' } className="card-link">View Book</a>
    </div>
  </div>
</div>
```

However, you should replace the `BOOK_TITLE`, `REVIEW_RATING`, `CREATOR_EMAIL`, `REVIEW_CONTENT`, and `BOOK_ID` placeholders with their proper values from the `review` that the component has fetched.

Note

Note that when working with JSX and React, `class` of an element is set with the `className` attribute, not `class`. When it's rendered as HTML, it becomes `class`.

5. Create another React component called `RecentReviews`. Its `constructor` method should set up `state` with the following keys/values:
 - `reviews: []` (empty list)
 - `currentUrl: props.url`
 - `nextUrl: null`
 - `previousUrl: null`
 - `loading: false`
6. Implement a method to download the reviews from the REST API. Call it `fetchReviews`. It should return immediately if `state.loading` is `true`. Then, it should set the `loading` property of `state` to `true`.
7. Implement `fetch` in the same way as you did in *Exercise 16.04 – fetching and rendering books*. It should follow the same pattern of requesting `state.currentUrl` and then getting the JSON data from the response. Then, set the following values in `state`:
 - `loading: false`
 - `reviews: data.results`
 - `nextUrl: data.next`
 - `previousUrl: data.previous`
8. Implement a `componentDidMount` method. This is a method that is called when React has loaded the component onto the page. It should call the `fetchReviews` method.
9. Create a `loadNext` method. If `nextUrl` in `state` is `null`, it should return immediately. Otherwise, it should set `state.currentUrl` to `state.nextUrl`, then call `fetchReviews`.
10. Similarly, create a `loadPrevious` method; however, this should set `state.currentUrl` to `state.previousUrl`.
11. Implement the `render` method. If the state is `loading`, then it should return the `Loading...` text inside an `<h5>` element.
12. Create two variables to store the `previousButton` and `nextButton` HTML. They should both have the `btn btn-secondary` class, and the next button should also have the `float-right` class. They should have the `onClick` attributes set to call the `loadPrevious` or `loadNext` methods. They should have their `disabled` attributes set to `true` if the respective `previousUrl` or `nextUrl` attributes are `null`. The button text should be `Previous` or `Next`.

13. Iterate over the reviews using the `map` method and store the result in a variable. Each review should be represented by a `ReviewDisplay` component with the `key` attribute set to the `pk` review and `review` set to the `Review` class. If there are no reviews (`reviews.length === 0`), then the variable instead should be an `<h5>` element with the **No reviews to display** content.
14. Finally, return all the content wrapped in the `<div>` elements, like this:

```
<div>
<div className="row row-cols-1 row-cols-sm-2 row-cols-md-3">
  { reviewItems }
</div>
<div>
  {previousButton}
  {nextButton}
</div>
</div>
```

The `className` we are using here will display each review preview in one, two, or three columns depending on the screen size.

15. Next, edit `base.html`. You will add all the new content inside the `content` block so that it will not be displayed on the non-main pages that override this block. Add an `<h4>` element with the `Recent Reviews` content.
16. Add a `<div>` element for React to render into. Make sure you give it a unique `id`.
17. Include the `<script>` tags to include React, React DOM, Babel, and the `recent-reviews.js` file. These four tags should be similar to what you had in *Exercise 16.04 – fetching and rendering books*.
18. The last thing to add is another `<script>` tag containing the `ReactDOM.render` call code. The root component being rendered is `RecentReviews`. It should have a `url` attribute set to the `url="{'% url value 'api:review-list' %}?limit=6"`. This does a URL lookup for `ReviewViewSet` and then appends a page size argument of 6, limiting the number of reviews that are retrieved to a maximum of 6.

Once you have completed these steps, you should be able to navigate to `http://127.0.0.1:8000/` (the main Bookr page) and see a page like this:

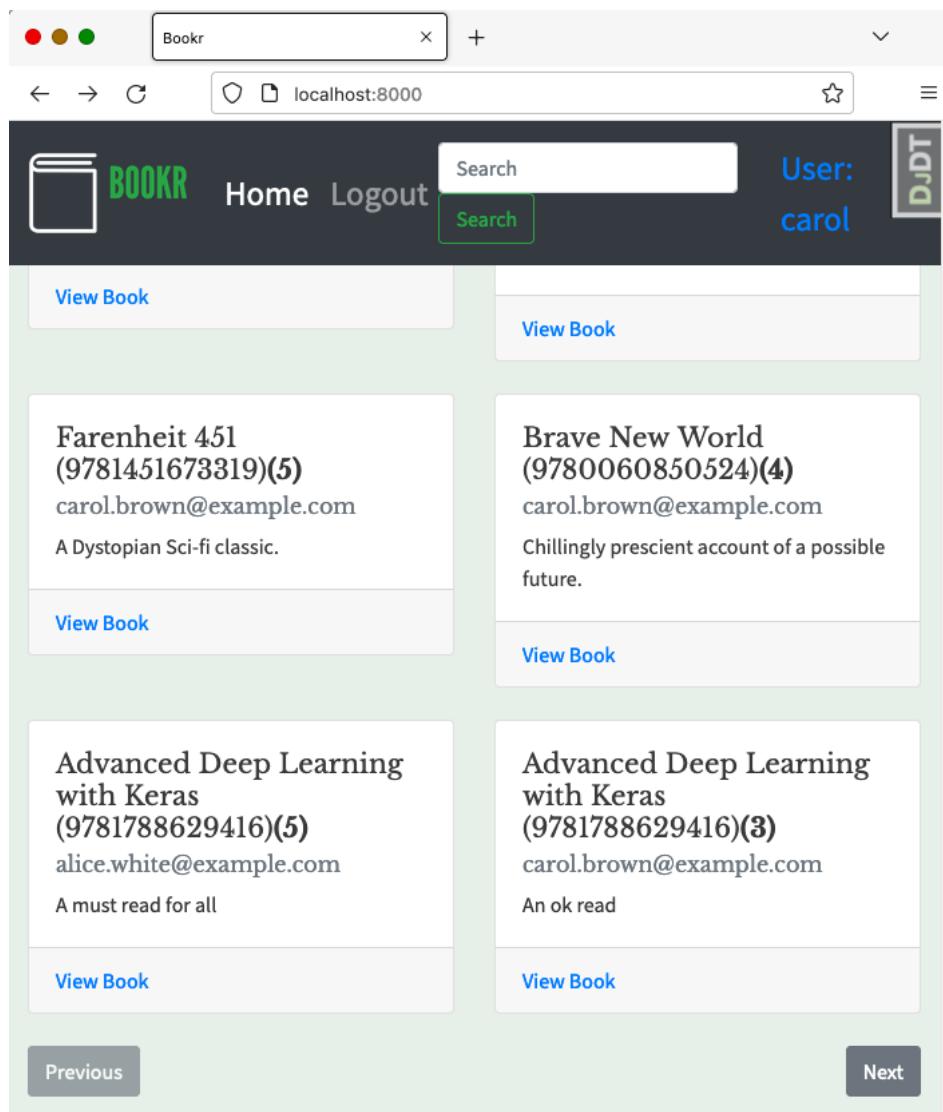


Figure 16.19 – Completed reviews preview

In the screenshot, the page has been scrolled to show the **Previous/Next** buttons. Notice the **Previous** button is disabled because we are on the first page.

If you click **Next**, you should see the next page of reviews. If you click **Next** enough times (depending on how many reviews you have), you will eventually reach the last page, and then the **Next** button will be disabled:

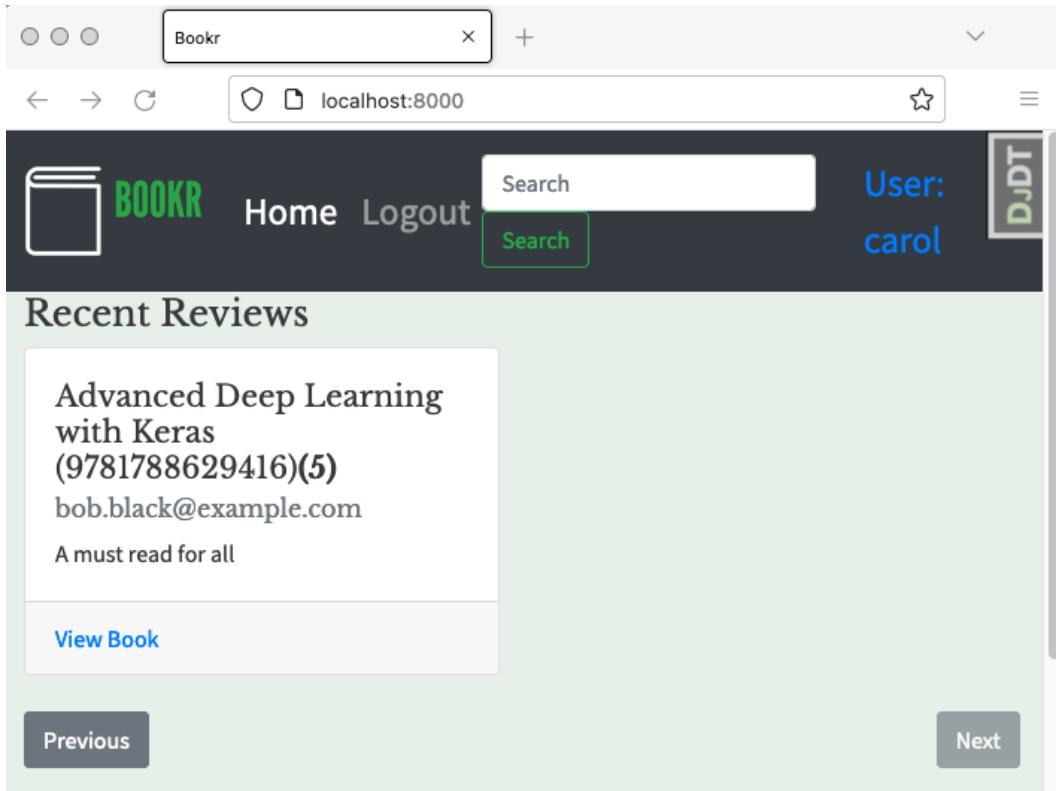


Figure 16.20 – The Next button is disabled

If you have no reviews, then you should see the **No reviews to display** message:

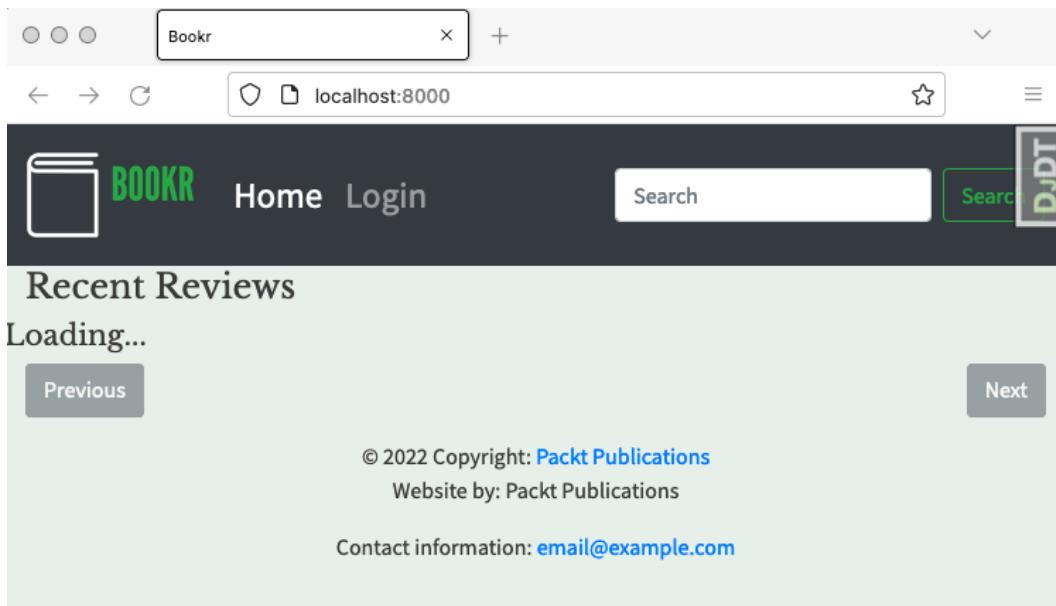


Figure 16.21 – The No reviews to display text

While the page is loading the reviews, you should see the **Loading...** text; however, it will probably only display for a split second since the data is being loaded from your own computer:

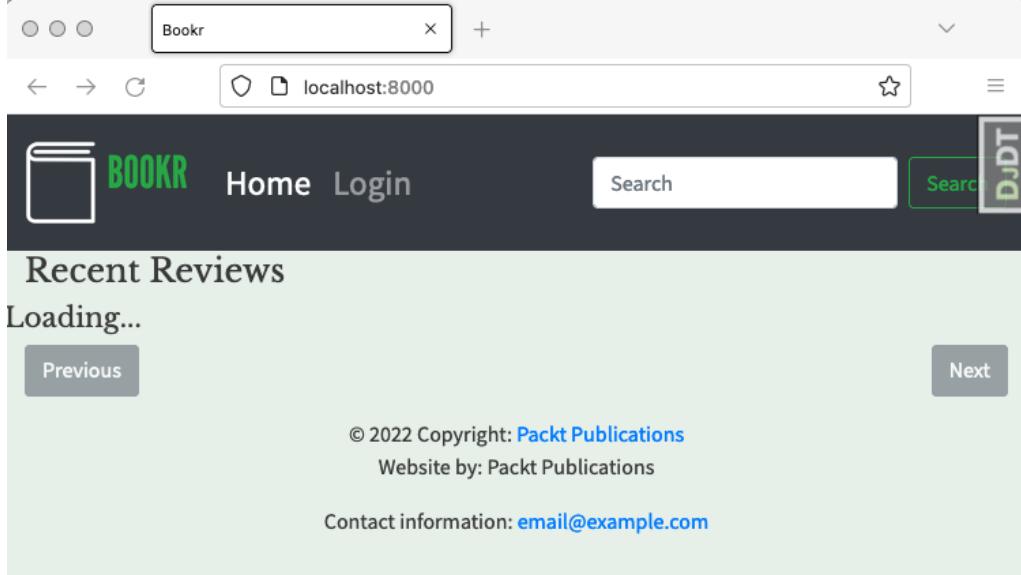


Figure 16.22 – The Loading... text

Note

The solution for this activity can be found on <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we introduced JavaScript frameworks and described how they work with Django to enhance templates and add interactivity. We introduced the JavaScript language and covered some of its main features, variable types, and classes. We then introduced the concepts behind React and how it builds HTML by using components. We built a React component using just JavaScript and the `React.createElement` function. After that, we introduced JSX and saw how it made the development of components easier by letting you directly write HTML in your React components. The concepts of promises and the `fetch` function were introduced, and we saw how to get data from a REST API using `fetch`. The chapter finished with an exercise that retrieved reviews from Bookr using the REST API and rendered them to the page in an interactive component.

In the next chapter, we will look at how to deploy our Django project to a production web server.

Index

A

Add user page

creating 152

Admin app login template

repurposing 438-440

admin files

discovering, in Django 472, 473

admin module 473

admin site

custom views, adding 488-491

custom views, restricting 487

templates, customizing 481, 482

text, customizing with AdminSite

attributes 480, 481

views, adding 485

admin Site

customizing 472

AdminSite

class 473, 474

object, customizing 183-187

object, examining from Python shell 179

subclassed 180-183

AdminSite attributes

reference link 481

used, for customizing admin

site text 480, 481

AdminSite class 473, 474

Angular

reference link 680

API endpoint

creating, for top contributors page 544, 545

API view

creating 536-539

application programming

interface (API) 531, 603

Argon2 441

arrays 684

Arrow Functions 687

assertions 594

types 597

utilizing 594, 595

authentication

implementing 550

initiating, with django-allauth 673

token-based authentication 551

authentication application 668

authentication app's templates

implementing 433-437

authentication app's views

implementing 433-437

authentication-based content

conditional blocks, using in

templates 452, 453

authentication data

used, for enhancing templates 450

authentication decorators 444-447

adding, to views 447-450

authentication provider 667**authentication redirection 444-447****authentication schemes**

reference link 551

automated test cases 592**automation testing 592**

functional testing 593

integration testing 593

regression testing 593

smoke testing 593

unit testing 593

B**Babel 698-700**

reference link 697

Backbone.js

reference link 680

base template

login links, toggling 450-452

logout links, toggling 450-452

basic authentication 551**BCrypt/SHA256 441****binary file formats**

for data exports 569

book catalog

creating, with CBV 517-522

book cover

displaying 425, 426

Bookr

custom admin site, creating 475-478

models and views, testing 618

most recent reviews, fetching on 713-718

Bookr admin site

logout template, customizing 483, 484

Bookr APIs

used, for implementing token-based authentication 551-556

Bookr application 1**Book Search page**

session storage, using 467, 468

book search scaffold 61, 62**Book Search view**

implementing, with Django Form 314-318

Bootstrap 140

using, for template styling 140, 141

browsers' sent values

security and trust 379-381

built-in search

using, to build custom admin dashboard 492, 493

bulk create operations 115

multiple records, creating 116

bulk update operations 115

used, for updating multiple records 117

C**cache invalidation**

file serving 243-245

Cascading Style Sheets (CSS) 140, 211**class-based API views**

creating 541-543

class-based views (CBVs) 124, 125, 515, 516, 539, 540

in-built views 516

testing 615

used, for creating book catalog 517-522

used, for performing CRUD operations 523

classes, JavaScript 686, 687

-
- code**
 simplifying, with ViewSets 545
- code debugger** 56-60
- collectstatic management command**
 static file, collecting for production 235, 236
 using 233, 234
- comma-separated values (CSV)** 120
- comments** 135
- common template variables**
 accessing 486
- complex lookups**
 performing, with Q objects 117
- complex query**
 used, for performing Q objects 118
- components**
 building, in React 689-693
- conditional blocks**
 using, for authentication-based
 content in templates 452, 453
- constants** 684
- content delivery network**
 (CDN) 222, 366, 689
- context processor** 371, 372
- create, read, update, and delete (CRUD)** 65
- crispy filter** 659, 660
- crispy template Tag** 660-662
- Cross-Site Request Forgery**
 (CSRF) 272, 431, 520
 security 272-275
- CRUD operations**
 Django admin app, using for 151, 152
- CRUD operations, performing**
with CBVs 523
- Create view 523, 524
 Delete view 526-528
 Read view 524
 Update view 525, 526
- CRUD operations, with Django admin app**
 Add user page 152, 153
 delete 158, 159
 retrieve 154, 155
 update 156, 157
 user account, creating 154
- CSS enhancements** 254-256
- CSV file**
 data, reading 566-568
 data, reading from 558, 559
 data, writing 566-568
 generating, with Python's csv
 module 563-566
 Python, used for writing to 562, 563
 reading, with Python 559-562
- CSV files, in Python**
 working with 558
- custom admin dashboard**
 building, with built-in search 492, 493
- custom admin site**
 creating, for Bookr 475-478
- custom filter**
 function, implementing 499
 using, in templates 499
- custom inclusion tag**
 building 512-515
- custom simple tag**
 creating 507-509
 template context, passing 509, 510
- custom template filter** 497, 498
 creating 498, 500-503
- custom view**
 adding, to admin site 488-491
 restricting, to admin site 487
 URLs, mapping 486, 487

D

data

- reading, from CSV file 558, 559
- storing, in sessions 461, 462

database records

- populating 120, 121

databases

- using 66

data exports

- binary file formats 569

date_hierarchy filters

- adding 196-199

date_list_filter

- adding 196-199

debugger 55, 56

decoupled architecture 532

default admin site

- overriding 478-480

delete operation 75

directory

- setting up, to store template filters 498

Django 472, 679

- admin files, discovering 472, 473
- AdminSite class 473, 474
- password storage 441
- plotly, integrating with 583
- read by user books, exporting
 - as XLSLX file 589
- used, for serving uploaded files 397, 398
- users and groups, adding through
 - admin app 160-165
- users and groups, managing 159, 160
- users and groups, modifying
 - through admin app 160-165
- user's reading history, visualizing on
 - user's profile page 583-588

views 515

- visualizations, integrating with 583

Django admin app

- interface, customizing 177
- using, for CRUD operations 151, 152

django-allauth 667-670

- authentication, initiating with 673
- features 673
- GitHub auth setup 673
- Google auth setup 673
- installing, with pip3 670, 671
- reference link 673
- setting up 671, 672
- SocialAccount 670
- SocialApplication 670
- SocialApplicationToken 670

Django app 3, 4, 19-21, 77

- creating 4-6

django-configurations 625-627

- configuring, from environment variables 628, 629

- installing, with pip 629

- manage.py, modifying 627

- setting up 629-632

Django Cookie Consent

- reference link 455

Django Cookie Law

- reference link 455

django-crispy-forms 658

- crispy filter 659, 660
- crispy template Tag 660-662
- installing, with pip3 658
- using, with SearchForm 662-666

Django Debug Toolbar 638-655

- installing, with pip3 656
- setting up 655-658
- usage scenarios 638, 639

Django development server 3, 19

Django form

- initial values, adding 336-338
- placeholders, adding 336-338
- used, for implementing Book Search view 314-318
- using, for file uploads 385-391
- using, for image uploads 391-397

Django Forms library 283, 284

- building and rendering 300-304
- form, defining 284-295
- form, rendering in template 295-300

Django migration 78-80**Django models**

- creating 80, 81, 340, 342
- editing 340, 342
- field options 82-84
- model methods 92, 93
- ModelForm class 342-346
- primary keys 84-86
- publisher, creating 346-352
- publisher, editing 346-352
- publisher form, integrating 352-355
- publisher form, styling 352-355
- relationships 86
- review_edit view, adding 356-362
- review model, adding 90, 91
- review ModelForm, creating 356-362
- reviews app, migrating 93-95
- testing 599-603

Django migrations

- creating 80, 81
- field options 82-84
- model methods 92, 93
- primary keys 84-86
- relationships 86
- review model, adding 90, 91
- reviews app, migrating 93-95

Django ORM

- database configuration 76, 77
- Django applications, creating 76, 77
- exploring 75, 76

Django project 2-4

- creating 4-6

Django project structure

- Django apps 19-21
- Django development server 19
- exploring 16, 17
- myproject directory 18
- PyCharm setup 21

Django REST Framework (DRF) 531, 532

- configuring 533
- functional API views 533
- installing 533

Django's database CRUD operations 95

- across relationships, querying 111
- all() method, using to retrieve object set 105
- create() method, using 97, 98
- create() method, using for many-to-many relationship 102
- create operation, performing 96, 97
- delete() method, using 115
- field lookups, filtering 106, 107
- filter() method, using to retrieve object 106
- foreign key relationship, querying
 - with object instance 112
- foreign keys, querying 111
- get() method, using to retrieve object 103
- many-to-many relationship, querying
 - with field lookup 112, 113
- many-to-many relationship, querying with object 113
- many-to-many relationship, querying with set() method 113
- many-to-many relationship, querying with add() method 101

model name, querying 111
object, creating with foreign key 98
object, returning with `get()` method 104
objects by filtering, retrieving 105
object, retrieving by `exclude()` method 108
object, retrieving with `order_by()` method 108-111
pattern matching, using for filtering operation 107
read operation 103
records, creating for many-to-many relationship 100, 101
records, creating for many-to-one relationship 98, 99
`set()` method, using for many-to-many relationship 102
`update()` method, using 114

Django settings

exploring 40, 41
referring 42

Django Storages

reference link 249

Django template language 134

comments 135
filters 135
template tags 135
template variable 134

Django templates

Book Details view, implementing 144
Bootstrap navigation bar, adding 142, 143
inheritance 139
list of books and reviews, displaying 136-138
styling, with Bootstrap 140, 141
template inheritance, adding 142, 143
using, to display greeting message 132-134
working with 130-132

Django views 30

test cases, writing, to validate authenticated users 608-611
testing 603, 604
testing, with authentication 607, 608
unit tests, writing 604-607

dj-database-url 633-636

installing, with pip 637
setting up 637, 638

double-underscore lookup 106

E

Elasticsearch 201

Ember

reference link 680

environment variable 622-624

interpreted, by Python script 624

error debugging 52

error handling 52

Excel files, in Python

working with 568

exception

generating 54, 55
viewing 54, 55

F

fetch function 704, 706

FieldFile

custom storage engines 403
custom storage engines, storing 405, 406
PIL images, writing to `ImageField` 406-408
reading from 404
working with 403

field options

`help_text` 82
`max_length` 82

-
- field validation** 320
 clean method 322, 323
 custom validators 320, 321
 field clean method, implementing 327-335
 form clean method, implementing 327-335
 validation function, implementing 327-335
- file**
 downloading 381-384
 uploading 381-384
- FileField**
 model, creating with 409-414
- files**
 storing, on model instances 398-401
- file saving**
 handling 419-421
- file serving**
 for cache invalidation 243-245
- FileSystemFinder** 229, 230
 static directory project, serving 230-233
- file uploads**
 Django forms, using for 385-391
 HTML forms, using for 376, 377
 ModelForm, using for 414-416
- filters** 135
- findstatic command** 239
 used, for finding files 240-242
- Flask framework** 1
- flexbox** 464
- form** 260-262
 building, in HTML 264-272
 data, accessing in View 276
 defining 284-295
 element 262, 263
 inputs, types 263
 query string, need for 282, 283
 rendering, in template 295, 297-300
 security, with Cross-Site Request Forgery Protection 272-275
- selecting, between GET or POST request 280-282
 validating 305, 307
 working, with POST data in view 276-280
- FormHelper**
 using, to update forms 674, 676
- forms**
 built-in field validation 311, 312
 extra field validation, adding 312-314
 validating, in view 308-310
- Foundation** 658
- frameworks, JavaScript** 680-682
- functional API views** 533
- functional component** 689
- functional testing** 593
- function-based views (FBVs)** 124, 515
- functions** 685, 686

G

- generic views** 539, 540
- GET method**
 working with 36-38
- GET request**
 versus POST request 280-282
- get_urls() method** 486
- GET values**
 exploring 38-40
- GitHub auth**
 setting up 673
- global logo**
 adding 256, 258
- Google auth**
 setting up 673
- graphs, in Python**
 generating 580-582
 playing with 578

graphs, with plotly
figure, setting up 579
generating 579
plot, generating 579
plot, rendering on web page 580

H

Help Desk user group 160

HTML forms

building 264-272
using, for file uploads 376, 377

HTML templates

base template, creating 43-46
finding, in app directories 43
template, rendering in view 47, 48
template, rendering with render function 47
templates directory, creating 43-46
variables, rendering in template 49
variables, using 49-51

HTTP 11-15

HTTP 403 Forbidden response 397

HTTP request

processing 15, 16

HttpRequest attributes

GET 30
headers 30
method 30
path 31
POST 30

HTTP response status codes

reference link 15

HyperText Markup Language (HTML) 8, 211

Hypertext Transfer Protocol (HTTP) 531

I

i18n module 437

image and PDF books

uploading 421-424

ImageField

model, creating with 409-414

images

storing, on model instances 402, 403

image uploads

Django forms, using for 391-397
ModelForm, using for 414-416

inclusion tags 505, 510

implementing 511
used, for rendering details on user profile page 528, 529
using, in template 511

initial values

adding, to Django form 336-338
adding, to OrderForm class 338-340

integration testing 593

J

JavaScript 682

classes 686, 687
constants 684
frameworks 680-682
functions 685, 686
loading 682, 683
map method 707
methods 686
objects 685
promises 703, 704
variables 683

JavaScript Object Notation

(JSON) 456, 457, 536

JavaScript XML (JSX) 697-700

properties 700, 701

Juggler application 120**L****list display fields 188, 190-194****login links**

toggling, in base template 450-452

logout links

toggling, in base template 450, 451

logout template

customizing, for Bookr admin site 483, 484

M**ManifestFilesStorage storage engine**

exploring 245-249

many-to-many relationship 88, 89**many-to-one relationship 86, 87****map method 707****media files**

serving, in development 366

media serving

configuring 367-370

settings 366

media storage

configuring 367-370

media uploads

settings 366

MEDIA_URL

using, in templates 371-376

methods 686**middleware 430****middleware modules 431, 432**

purpose 431

migration 78**model 7****ModelAdmin classes**

customizing 187, 188, 206, 208, 209

date_hierarchy filters, adding 196-199

date list_filter, adding 196-199

display decorator 194, 195

fields, excluding and grouping 202-206

filter 195, 196

list display fields 188-194

search bar 199-202

ModelForm 342

using, for file and image uploads 416-419

ModelForm class 342-346**model methods 92, 93****model serializers 540, 541**

creating 541-543

models, with Django admin app

Book change page 172-175

Change publisher page 168-172

deletion behavior 175-177

foreign keys 175-177

Lists, modifying 167, 168

registering 165

reviews model, registering 166, 167

Model View Controller (MVC) 7**Model View Template (MVT) 2, 7, 65**

best practices 8, 10

multi-field validation 323-326**multiple records**

creating, with bulk create operations 116

updating, with bulk update operations 117

myproject directory 18**N****National Public Radio (NPR) 577****Node.js 681****non-field errors 325**

O

Object Relational Mapping (ORM) 7, 65, 680
objects 685
one-to-one relationship 90
OrderForm class
 initial values, adding 338-340
 placeholders, adding 338-340

P

password storage
 in Django 441
PBKDF2/SHA256 441
PDF files, in Python
 working with 575
PDF version
 generating, in Python 575-578
pickle module 456
Pillow
 image, resizing with 393
placeholders
 adding, to Django form 336-338
 adding, to OrderForm class 338-340
plotly
 integrating, with Django 583
Portable Document Format (PDF) 575
 web pages, converting into 575
POST data, in view
 working with 276-280
POST method
 content-length 13
 content-type 13
 working with 36-38
POST request
 versus GET request 280-282

primary keys 84-86
profile page 441
 adding 442-444
project management application
 models, creating 120
Promise object 703, 704
publisher
 creating 346-352
 editing 346-352
publisher form
 integrating 352-355
 styling 352-355
PyCharm 2, 459
 project, setting up 21-29
 setting up 21
Pylons framework 1
Python
 PDF version, generating 575-578
 used, for reading CSV file 559-562
 used, for writing to CSV files 562, 563
 values, retrieving 305, 307
 XLSX files, creating 572-575
Python exceptions 52
 ImportError 52
 IndentationError 52
 IndexError 53
 KeyError 53
 NameError 53
 SyntaxError 52
 TypeError 53
Python Imaging Library (PIL) 391
Python's csv module
 used, for generating CSV file 563-566
 working with 558
Python shell
 AdminSite object, examining from 179

Q

Q objects

- queryset object, verifying 118, 119
- used, for performing complex lookups 117
- used, for performing complex query 118

QueryDict objects

- exploring 38-40
- working with 36-38

queryset object

- verifying 118, 119

R

React

- components, building 689-693
- reference link 680
- working with 688, 689

React component

- properties 701-703

React example

- books, fetching 708-711
- books, rendering 708-711
- setting up 694-696

regression testing 593

relationships 86

- many-to-many relationship 88, 89
- many-to-one relationship 86, 87
- one-to-one relationship 90

remote user authentication 551

representational state transfer (REST) 532

RequestFactory 612

- using, to test views 612-615

requesting site 667

request.user object 442

REST APIs 531, 532

- creating 534, 535

RESTful web service 532

retrieve 154

review_edit view

- adding 356-362

review model

- adding 90, 91

review ModelForm

- creating 356-362

reviews app

- migrating 93-95

Reviews logo

- adding 252, 254

Routers 545

- URL configuration, using 546

- using 546-550

S

Sample Link

- displaying 425, 426

SearchForm

- Django Crispy Forms, using with 662-666

serializers 536

- class-based views 539, 540
- generic views 539, 540

session authentication 551

session engine 454, 455

session engine, settings

- cached sessions 455
- cookie-based sessions 455
- file-based sessions 455

session key

- examining 457-461

sessions 454

- cookie content, flagging 455

- data, storing 461, 462

- recently viewed books, storing 462-466

session storage

- using, for Book Search page 467, 468

- setUp() method** 598
- simple tags** 505
- using, inside templates 506
- simple template tag**
- creating 505
 - directory, setting up 505
 - implementing 506
 - template library, setting up 506
- singleton class** 499
- sites framework**
- reference link 671
- site welcome screen**
- creating 60, 61
- site-wide Django admin**
- customizing 177, 179
- site-wide Django admin, customizations**
- AdminSite object, customizing 183-187
 - AdminSite object, examining
 - from Python shell 179
 - AdminSite, subclassing 180-183
- smoke testing** 593
- soft delete** 157
- SQL command**
- delete operation 75
- static file finder** 214, 215
- AppDirectoriesFinder 215
 - file, serving from app directory 218-222
 - namespacing 216, 217
 - request 215
- STATICFILES_DIRS prefixed mode** 236-239
- static file serving** 213, 214
- static template tag**
- used, for generating static URLs 222-226
 - using 226-228
- static URLs**
- generating, with static template tag 222-226
- storage engine** 249-251
- CSS enhancements 254-256
 - global logo, adding 256, 258
 - Reviews logo, adding 252, 254
- string filter** 503, 504
- Structured Query Language (SQL)** 7
- superuser account**
- creating 149-151
- T**
- tearDown() method** 599
- template context**
- passing, in custom simple tag 509, 510
- template filters** 496, 497
- directory, setting up to store 498
- template library**
- setting up 498
- templates** 8
- enhancing, with authentication data 450
 - media, referring in 408, 409
- template tags** 135, 504
- template tags, types**
- inclusion tags 505
 - simple tags 505
- template variable** 134
- using, to pass additional keys
 - to templates 491, 492
- test case classes** 615
- LiveServerTestCase class 616
 - SimpleTestCase class 616
 - test code, modularizing 617
 - TransactionTestCase class 616
- test cases** 594
- implementing 594
 - pre-test setup and cleanup,
 - performing 598, 599

testing 591

- automation testing 592
- functional testing 593
- goals, achieving with 592
- importance 592
- in Django 593
- integration testing 593
- regression testing 593
- smoke testing 593

token authentication 551**token-based authentication** 551

- implementing, for Bookr APIs 551-556

top contributors page

- used, for creating API endpoint 544, 545

Transmission Control Protocol (TCP) 532**try-with-resources** 561**U****Uni-Form** 658**unit test**

- writing 596

unit testing 593

- implementing 594

update 156**uploaded files**

- working with, in view 378, 379

URL configuration 125-127

- simple function-based view, implementing 127-129

- with Routers 546

- with Viewsets 546

URL mapping 127

- exploring 31, 32

- for custom view 486, 487

- parameters 282, 283

- view, writing 32-36

user profile page

- details, rendering with inclusion tags 528, 529
- user's reading history, visualizing on 583-588

V**Vanilla JS** 680**variables** 683**verbatim template tag** 712**view function**

- creating 485

views 8

- adding, to admin site 485

ViewSets 545

- URL configuration, using 546

- used, for simplifying code 545

- using 546-550

visualizations

- integrating, with Django 583

Vue

- reference link 680

W**web page**

- converting, into PDFs 575

- PDF version, generating in Python of 575-578

- plot, rendering on 580

workbook

- creating 571

- data, writing to 572

worksheet

- creating 571

- data, writing to 571

X

XLSX file

books read by a user, exporting as 589

XLSX files 569, 570

creating, in Python 572-575

data, writing to workbook 572

data, writing to worksheet 571

workbook, creating 571

worksheet, creating 571

XlsxWriter package, used for
working with 569

XlsxWriter Python package 570, 571

XlsxWriter package

used, for working with XLSX files 569

XlsxWriter Python package 570, 571



www.packtpub.com

Subscribe to our online digital library for full access to over 7,000 books and videos, as well as industry leading tools to help you plan your personal development and advance your career. For more information, please visit our website.

Why subscribe?

- Spend less time learning and more time coding with practical eBooks and Videos from over 4,000 industry professionals
- Improve your learning with Skill Plans built especially for you
- Get a free eBook or video every month
- Fully searchable for easy access to vital information
- Copy and paste, print, and bookmark content

Did you know that Packt offers eBook versions of every book published, with PDF and ePub files available? You can upgrade to the eBook version at packtpub.com and as a print book customer, you are entitled to a discount on the eBook copy. Get in touch with us at customercare@packtpub.com for more details.

At www.packtpub.com, you can also read a collection of free technical articles, sign up for a range of free newsletters, and receive exclusive discounts and offers on Packt books and eBooks.

Other Books You May Enjoy

If you enjoyed this book, you may be interested in these other books by Packt:



Becoming an Enterprise Django Developer

Michael Dinder

ISBN: 978-1-80107-363-9

- Use Django to develop enterprise-level apps to help scale your business
- Understand the steps and tools used to scale up a proof-of-concept project to production without going too deep into specific technologies
- Explore core Django components and how to use them in different ways to suit your app's needs
- Find out how Django allows you to build RESTful APIs
- Extract, parse, and migrate data from an old database system to a new system with Django and Python
- Write and run a test using the built-in testing tools in Django



Django 4 for the Impatient

Greg Lim , Daniel Correa

ISBN: 978-1-80324-583-6

- Understand and implement Django Apps' basic structure, including URLs, views, templates, and models
- Add bootstrap to improve the aesthetics of the site
- Create your own custom pages and have different URLs to route to them
- Navigate between pages by adding a header bar to all pages
- Work with databases and models
- Explore the powerful built-in admin interface with Django
- Use Django's powerful, built-in authentication system
- Deploy your Django project on the internet for the world to use

Packt is searching for authors like you

If you're interested in becoming an author for Packt, please visit authors.packtpub.com and apply today. We have worked with thousands of developers and tech professionals, just like you, to help them share their insight with the global tech community. You can make a general application, apply for a specific hot topic that we are recruiting an author for, or submit your own idea.

Share Your Thoughts

Hi,

We are Ben Shaw, Saurabh Badhwar, Bharath Chandra K S, and Chris Guest, authors of *Web Development with Django – Second Edition*. I really hope you enjoyed reading this book and found it useful for increasing your productivity and efficiency in Django.

It would really help us (and other potential readers!) if you could leave a review on Amazon sharing your thoughts on *Web Development with Django – Second Edition*.

Go to the link below to leave your review:

<https://packt.link/r/1800560788>

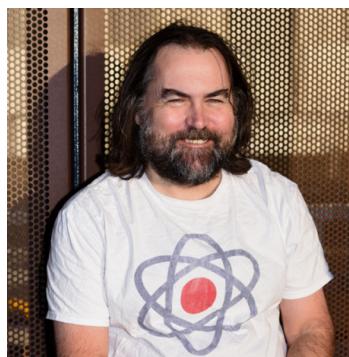
Your review will help us to understand what's worked well in this book, and what could be improved upon for future editions, so it really is appreciated.

Best Wishes,

,



Saurabh Badhwar



Chris Guest



Bharath Chandra K S

Download a free PDF copy of this book

Thanks for purchasing this book!

Do you like to read on the go but are unable to carry your print books everywhere? Is your eBook purchase not compatible with the device of your choice?

Don't worry, now with every Packt book you get a DRM-free PDF version of that book at no cost.

Read anywhere, any place, on any device. Search, copy, and paste code from your favorite technical books directly into your application.

The perks don't stop there, you can get exclusive access to discounts, newsletters, and great free content in your inbox daily

Follow these simple steps to get the benefits:

1. Scan the QR code or visit the link below



<https://packt.link/free-ebook/9781803230603>

2. Submit your proof of purchase
3. That's it! We'll send your free PDF and other benefits to your email directly

